

B-Trees

CMSC 420

Improving our index

- **B-Tree**: A generalization of a 2-3 tree.
 - Parameter p (for p ointer) controls the **fan-out (#children)** of every node.
 - For efficiency purposes, usually nodes allocated with **static array of size p** that contain pointers to subtrees.
 - Any **non-root** node can have between $\left\lceil \frac{p}{2} \right\rceil - 1$ and $p - 1$ **keys inclusive**.
 - The root itself is the only node allowed to have a **single key (2 children)**
 - Just **allowed**. Not necessarily **“has to have”**.

Next-level quiz

- **Quiz:** To separate p subtrees I need the following #separators...

Next-level quiz

- **Quiz:** To separate p subtrees I need the following #separators...

$$p - 1$$


$$p$$

$$p + 1$$

Something
else (what?)

Next-level quiz

- **Quiz:** To separate p subtrees I need the following #separators...



So, **all B-Tree nodes** have $(p - 1)$ 4-byte separators and p 8-byte pointers, for a total memory footprint of $12p - 4 = \mathcal{O}(p)$ **each**!

Insertion / Deletion into B-Trees

- Generalization of insertion /deletion in 2-3 trees.
 - Difference: **arrays** of keys and references **inside every node**
- Some new constraints, though:
 1. Any **non-root** node can have between $\left\lceil \frac{p}{2} \right\rceil - 1$ and $p - 1$ keys.
 - Exception: root **can have 2 children! (we'll see why)**
 2. If a key is inserted into a node with $p - 1$ keys, we have an **overflow**.
 3. If a key is deleted from a leaf node with $\left\lceil \frac{p}{2} \right\rceil - 1$ keys, we have an **underflow**.


What's with the constraint on the root?

- Suppose we have $p = 8$ and the following “stub” B-Tree consisting of just the root...

30, 36, 52, 67, 79, 90, 98

What's with the constraint on the root?


- Suppose we have $p = 8$ and the following “stub” B-Tree consisting of just the root...



- And now I want to insert the key 34 in my tree.

What's with the constraint on the root?

- Suppose we have $p = 8$ and the following “stub” B-Tree consisting of just the root...



- And now I want to insert the key 34 in my tree.



What's with the constraints on root and leaves?

- Suppose we have $p = 8$ and the following “stub” B-Tree consisting of just the root...




- And now I want to insert the key 34 in my tree.
- Solution: Split the root!



What's with the constraints on root and leaves?

- Suppose we have $p = 8$ and the following “stub” B-Tree consisting of just the root...



- And now I want to insert the key 34 in my tree.
- Solution: Split the root!
 - Left child: holds $4 = \left\lceil \frac{p}{2} \right\rceil \geq \left\lceil \frac{p}{2} \right\rceil - 1$ keys which is our lower bound
 - Right child: holds $3 = \left\lceil \frac{p}{2} \right\rceil \geq \left\lceil \frac{p}{2} \right\rceil - 1$, our lower bound



Dealing with overflows

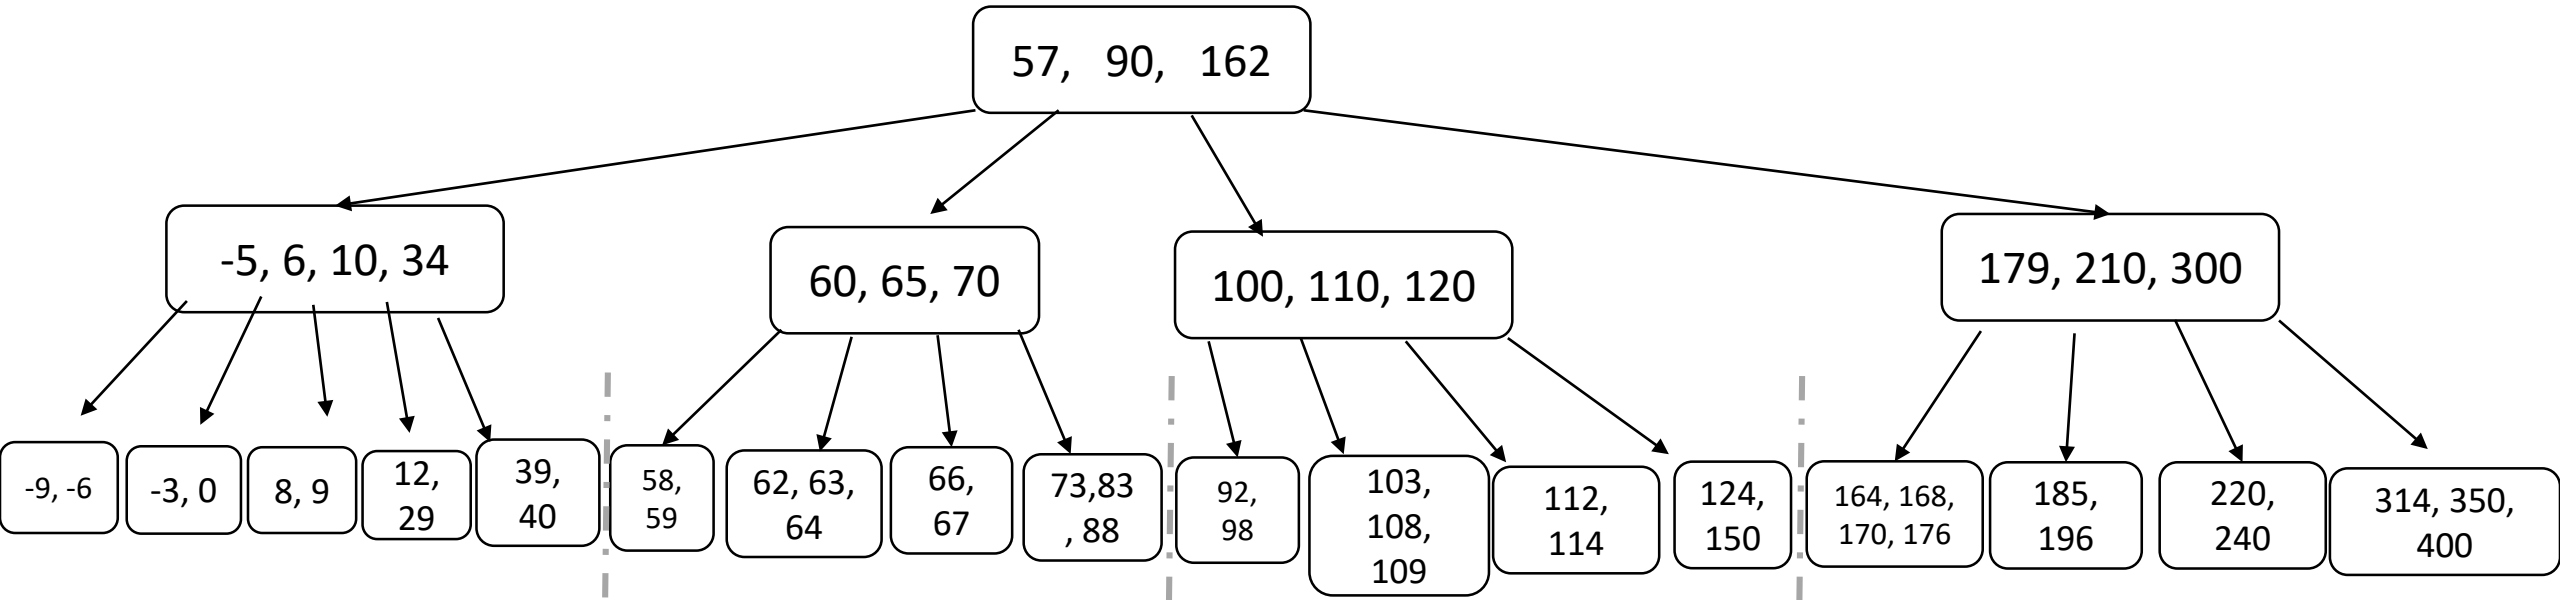
- Unlike in 2-3 trees, where splitting is practically the only option allowed, in B-trees, we will always attempt **key rotations** first.
 - To avoid asking the heap for **large contiguous storage!**

Dealing with overflows

- Unlike in 2-3 trees, where splitting is practically the only option allowed, in B-trees, we will always attempt **key rotations** first.
 - To avoid asking the heap for **large contiguous storage!**
- Idea: Look at your sibling nodes (≤ 2) and locate **at least one** that would **not overflow** if you were to **add a key to it**
 - If you find one such sibling, **rotate the key “closest” to the sibling node to that node through the parent.**
 - This is always either the **smallest** or the **largest** key of the current node

In all examples, $p = 5$

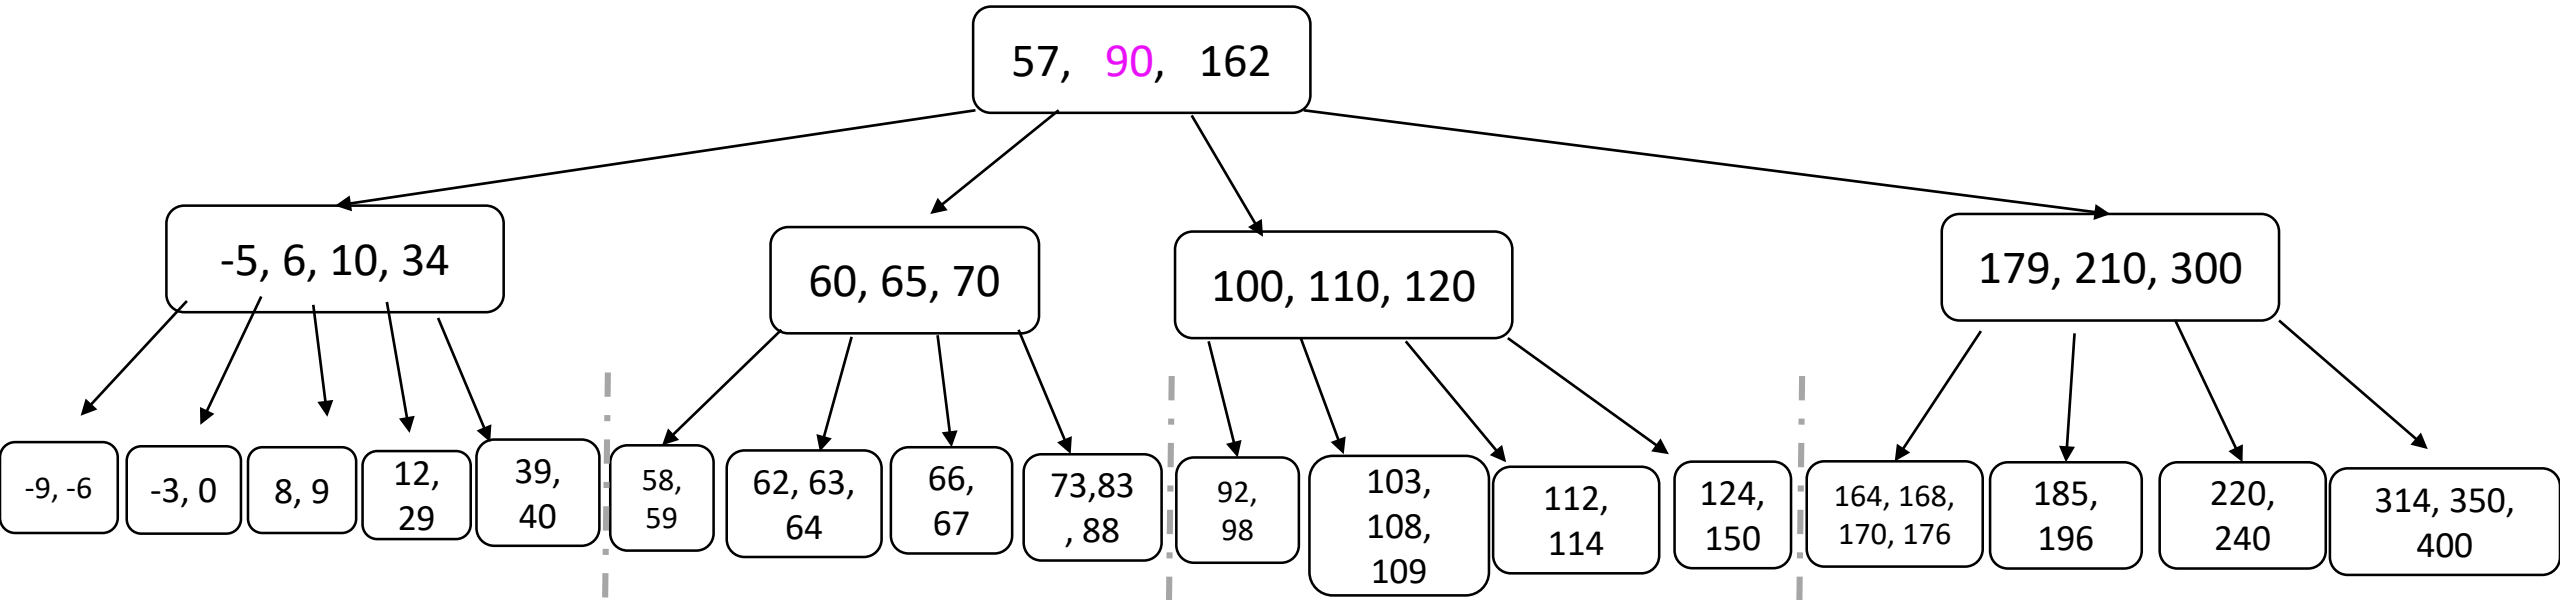
Example #0: No overflows



- Task: Insert 119

In all examples, $p = 5$

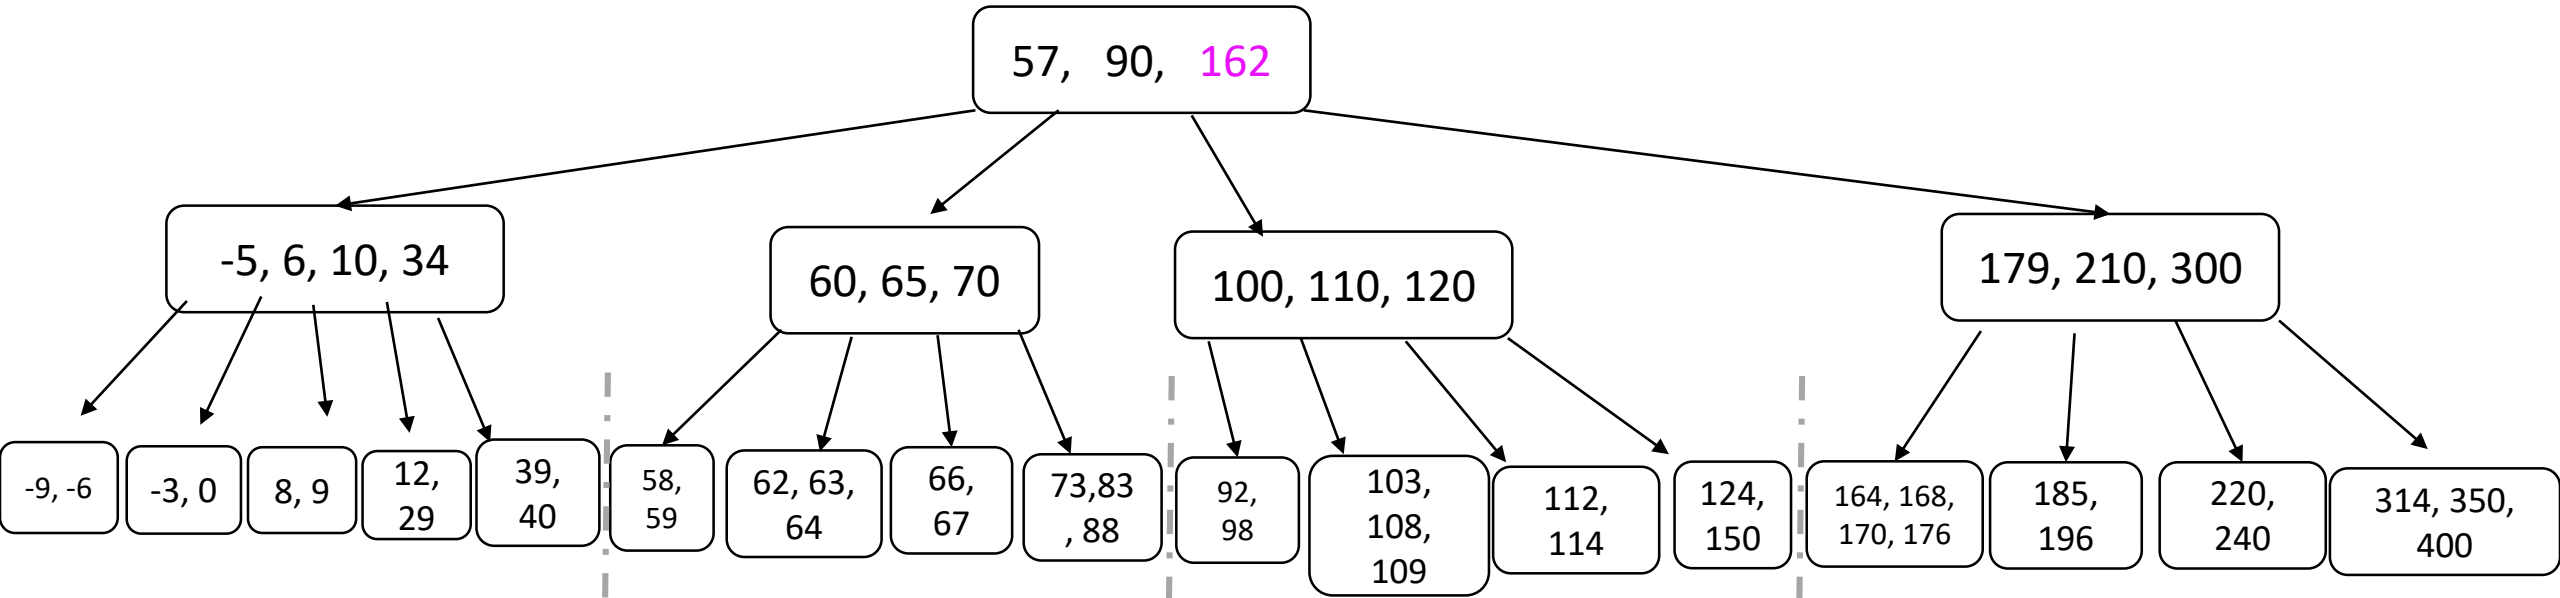
Example #0: No overflows



- Task: Insert 119

In all examples, $p = 5$

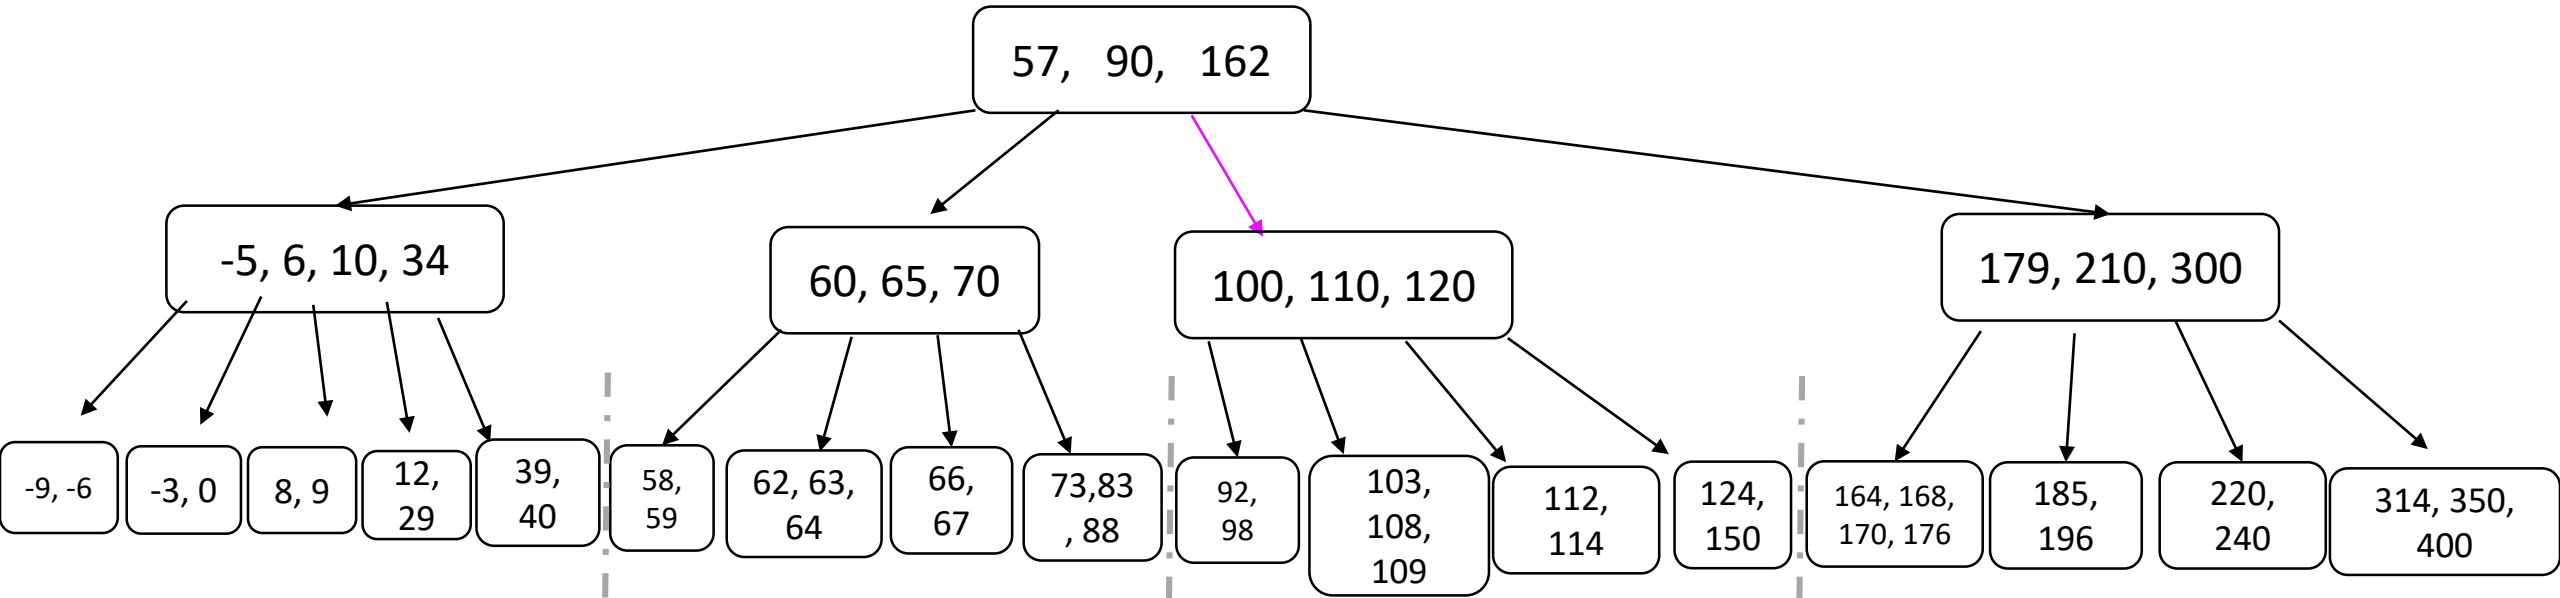
Example #0: No overflows



- Task: Insert 119

In all examples, $p = 5$

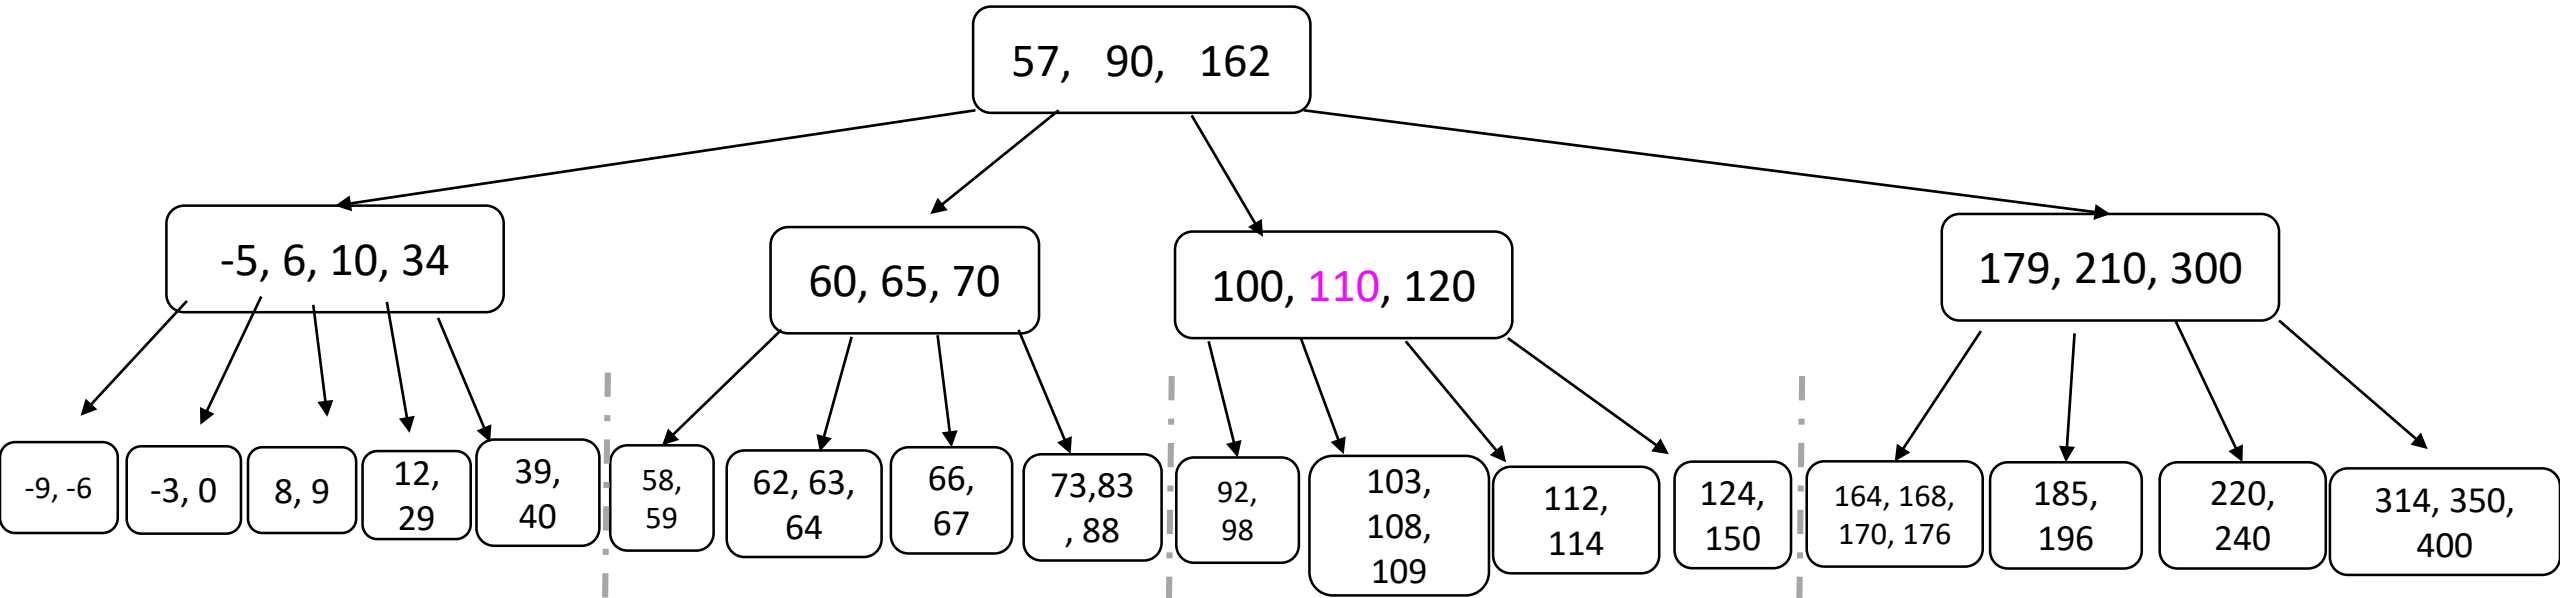
Example #0: No overflows



- Task: Insert 119

In all examples, $p = 5$

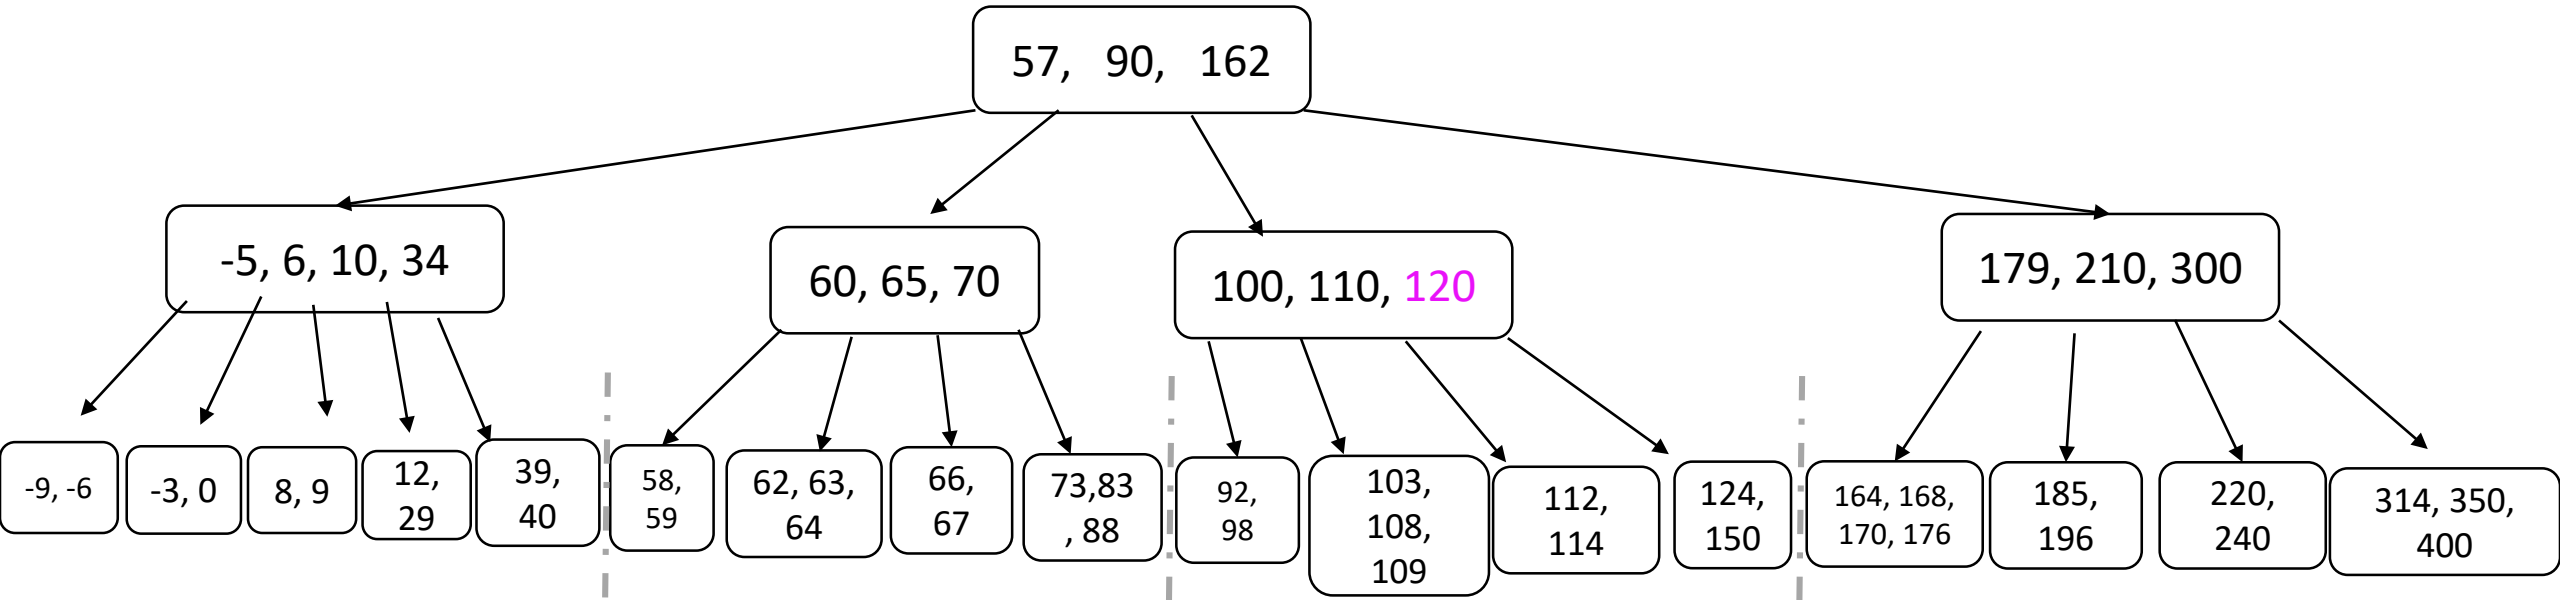
Example #0: No overflows



- Task: Insert 119

In all examples, $p = 5$

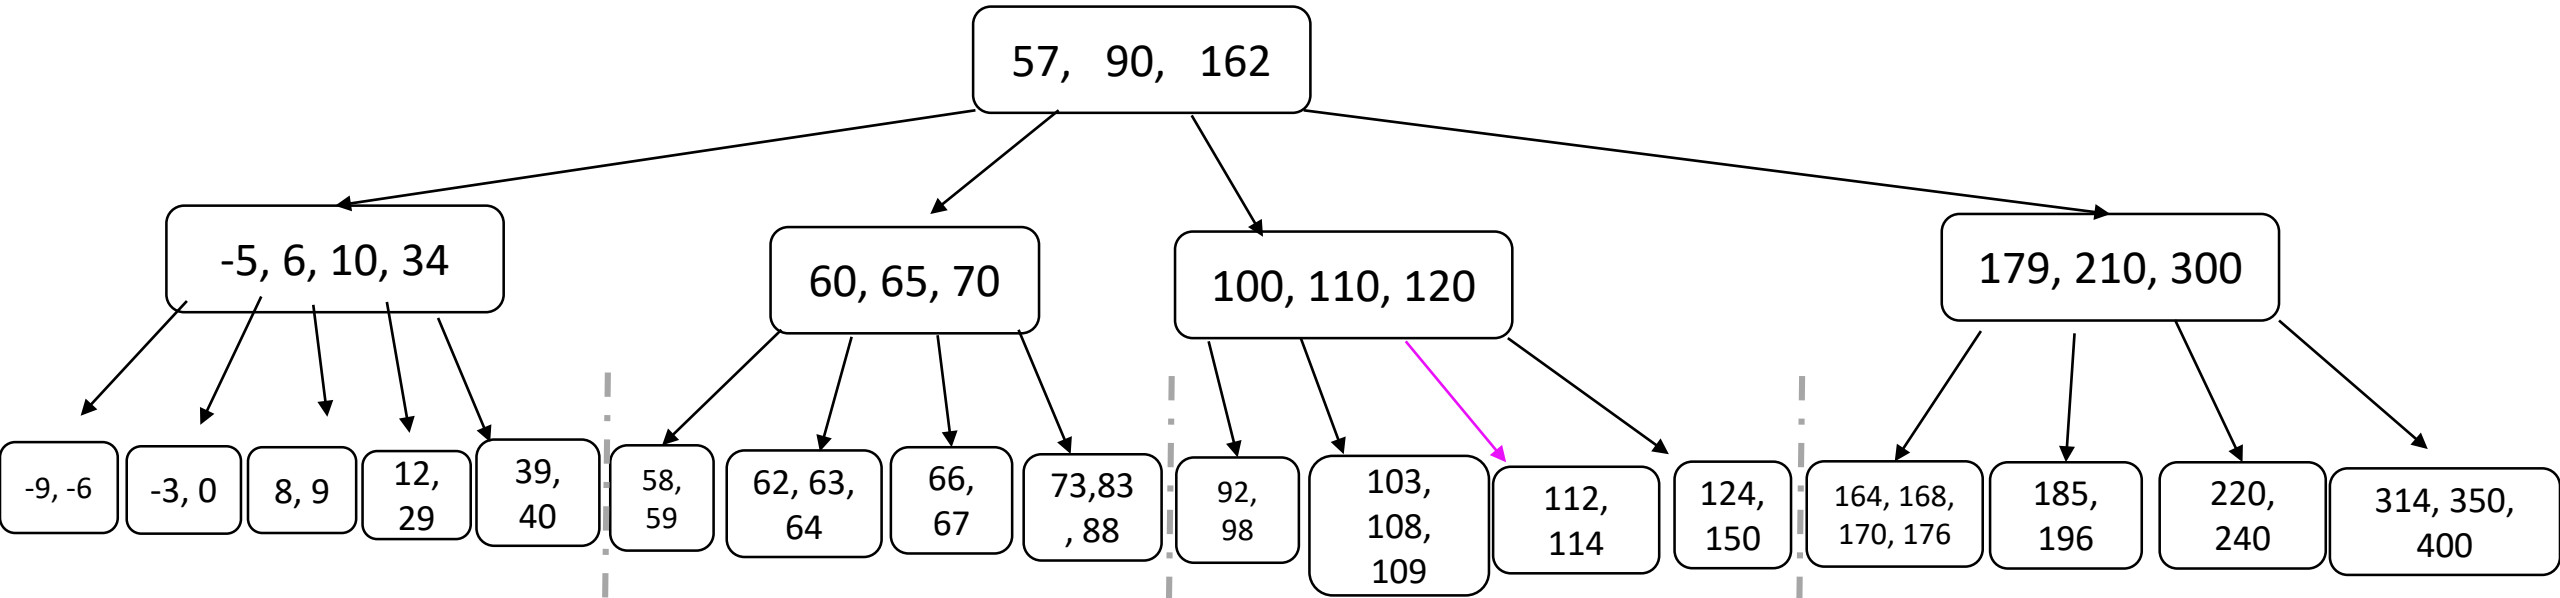
Example #0: No overflows



- Task: Insert 119

In all examples, $p = 5$

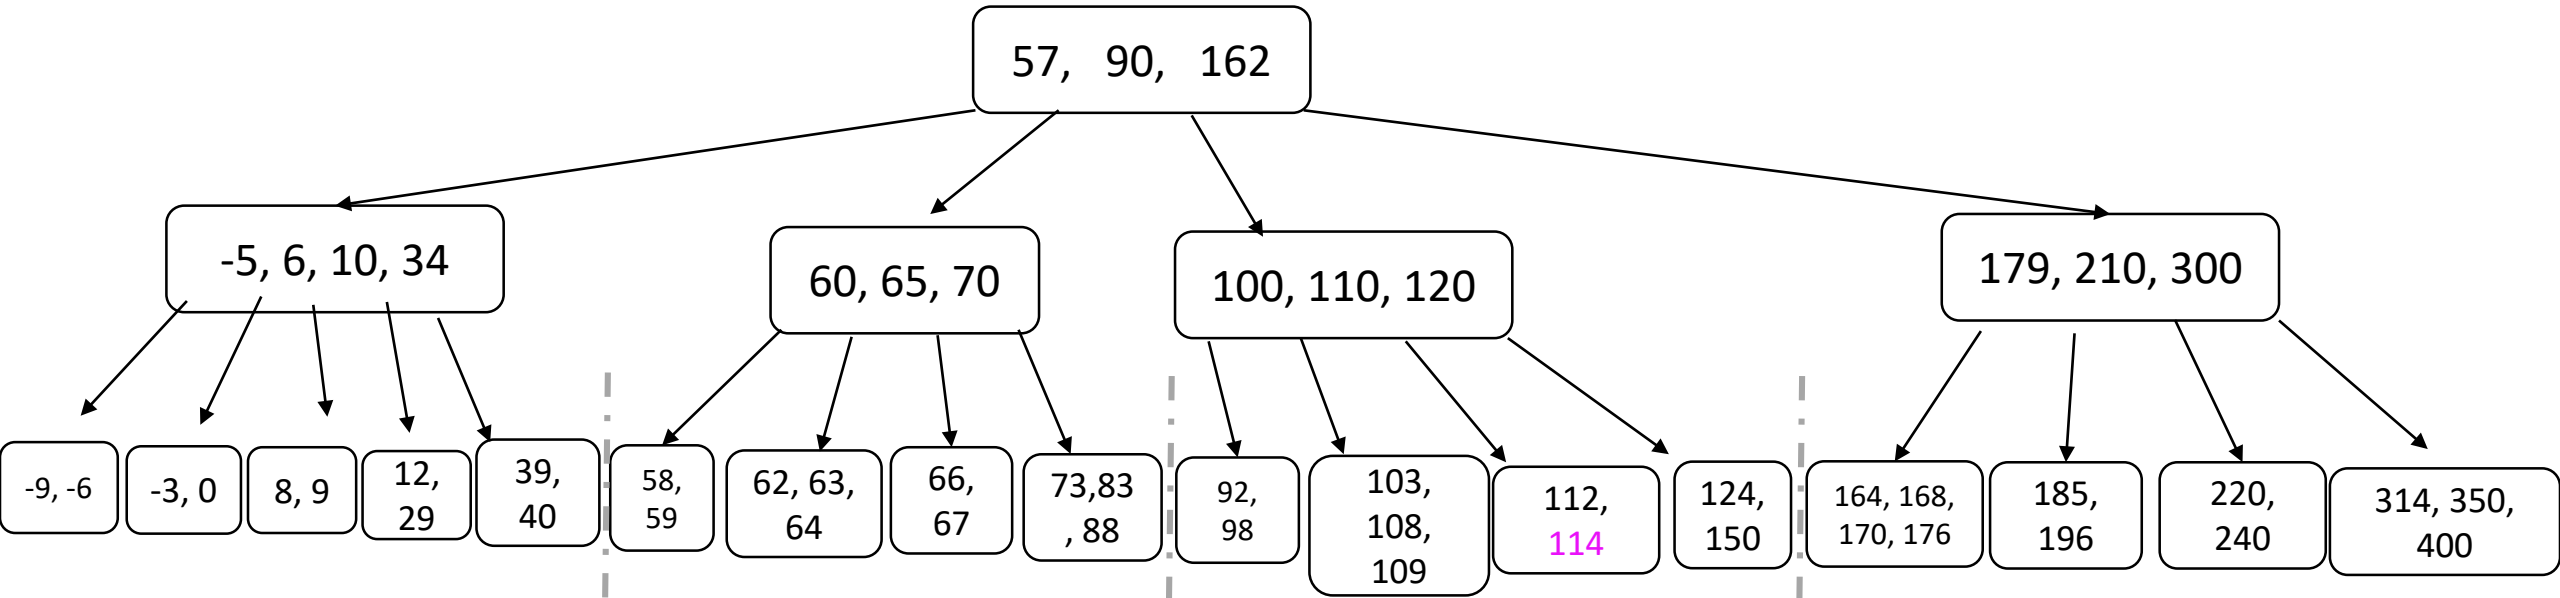
Example #0: No overflows



- Task: Insert 119

In all examples, $p = 5$

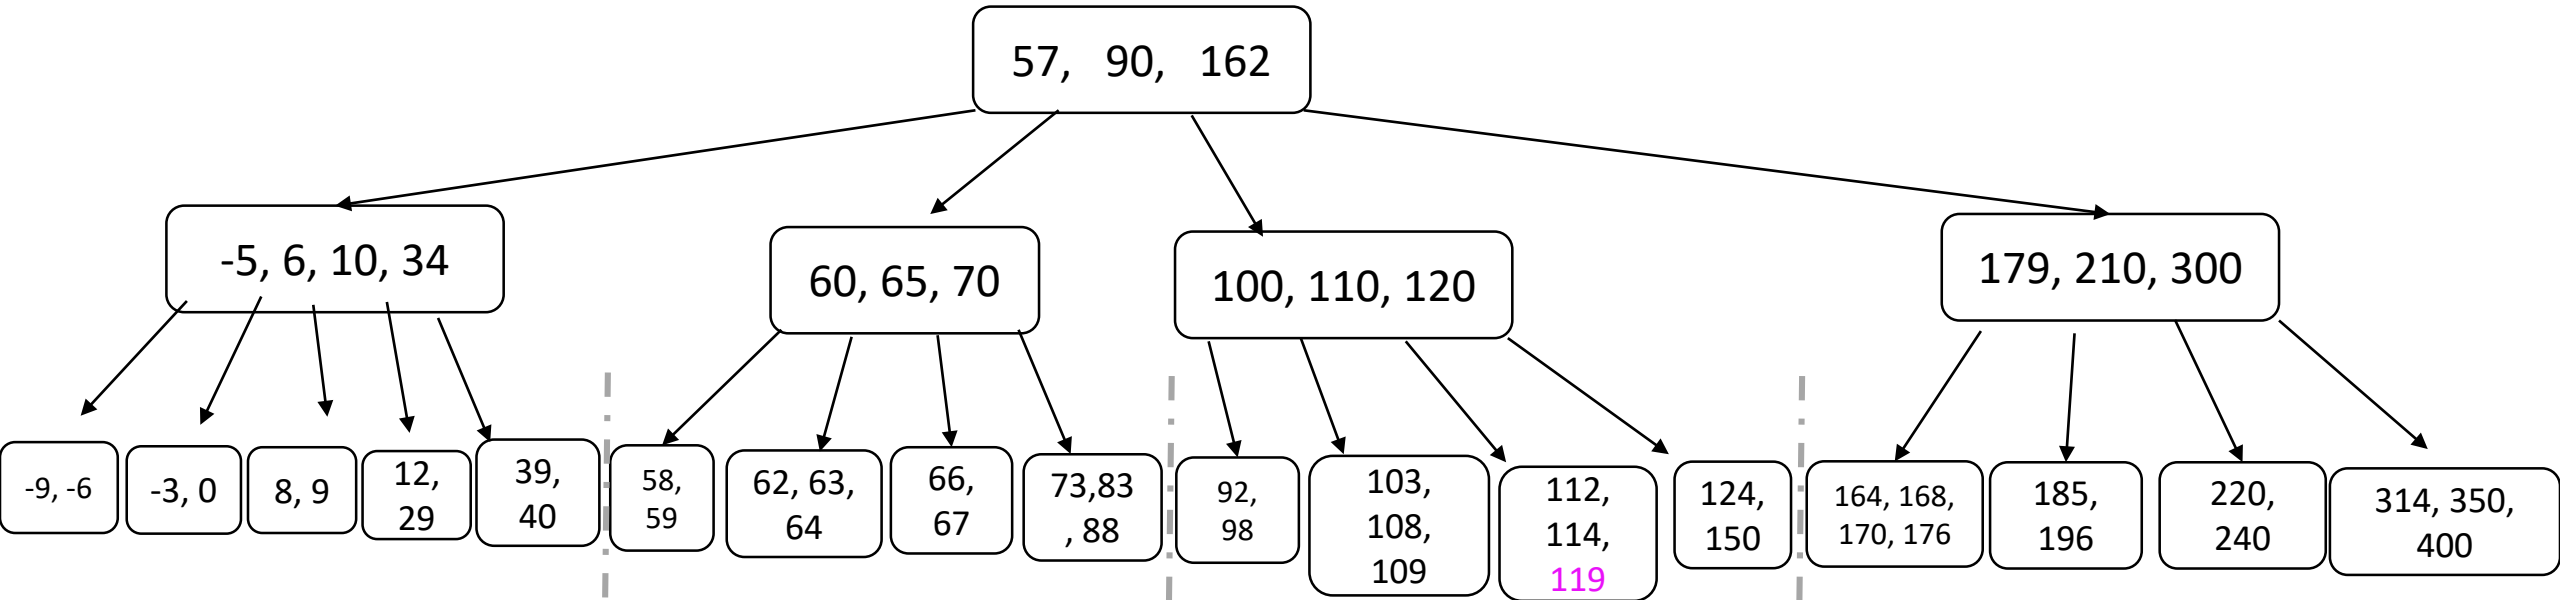
Example #0: No overflows



- Task: Insert 119

In all examples, $p = 5$


Example #0: No overflows




- Task: Insert 119

In all examples, $p = 5$

Example #0: No overflows

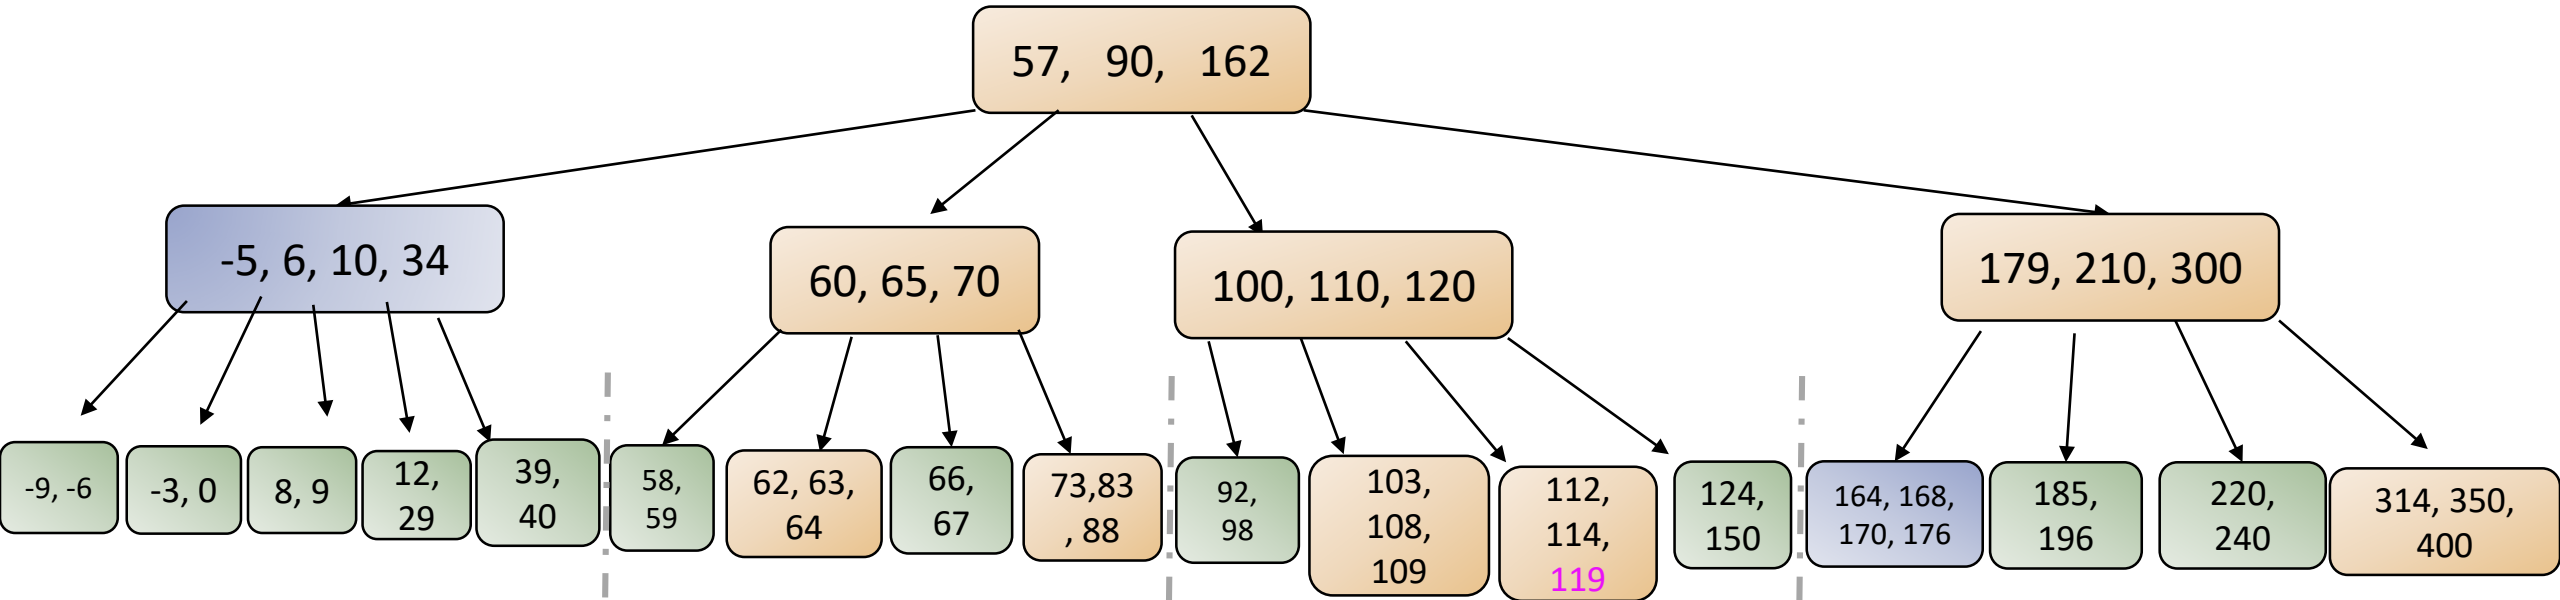


- Task: Insert 119
- No overflow at the leaf, so no problem



In all examples, $p = 5$

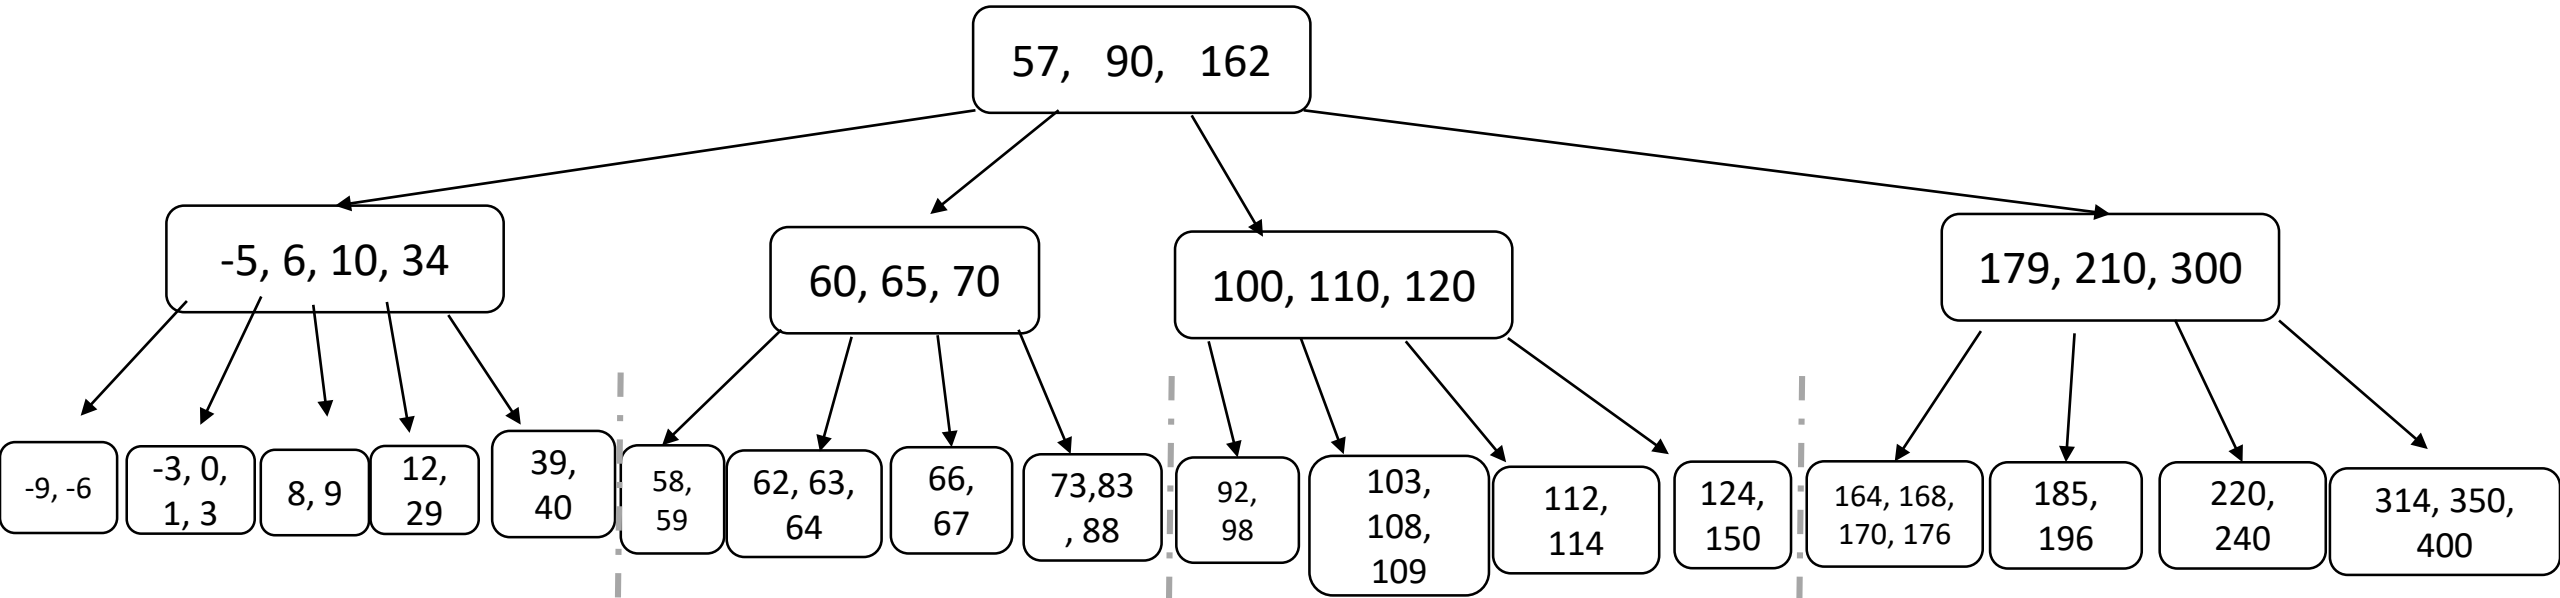
Example #0: No overflows



- **Side notes:**

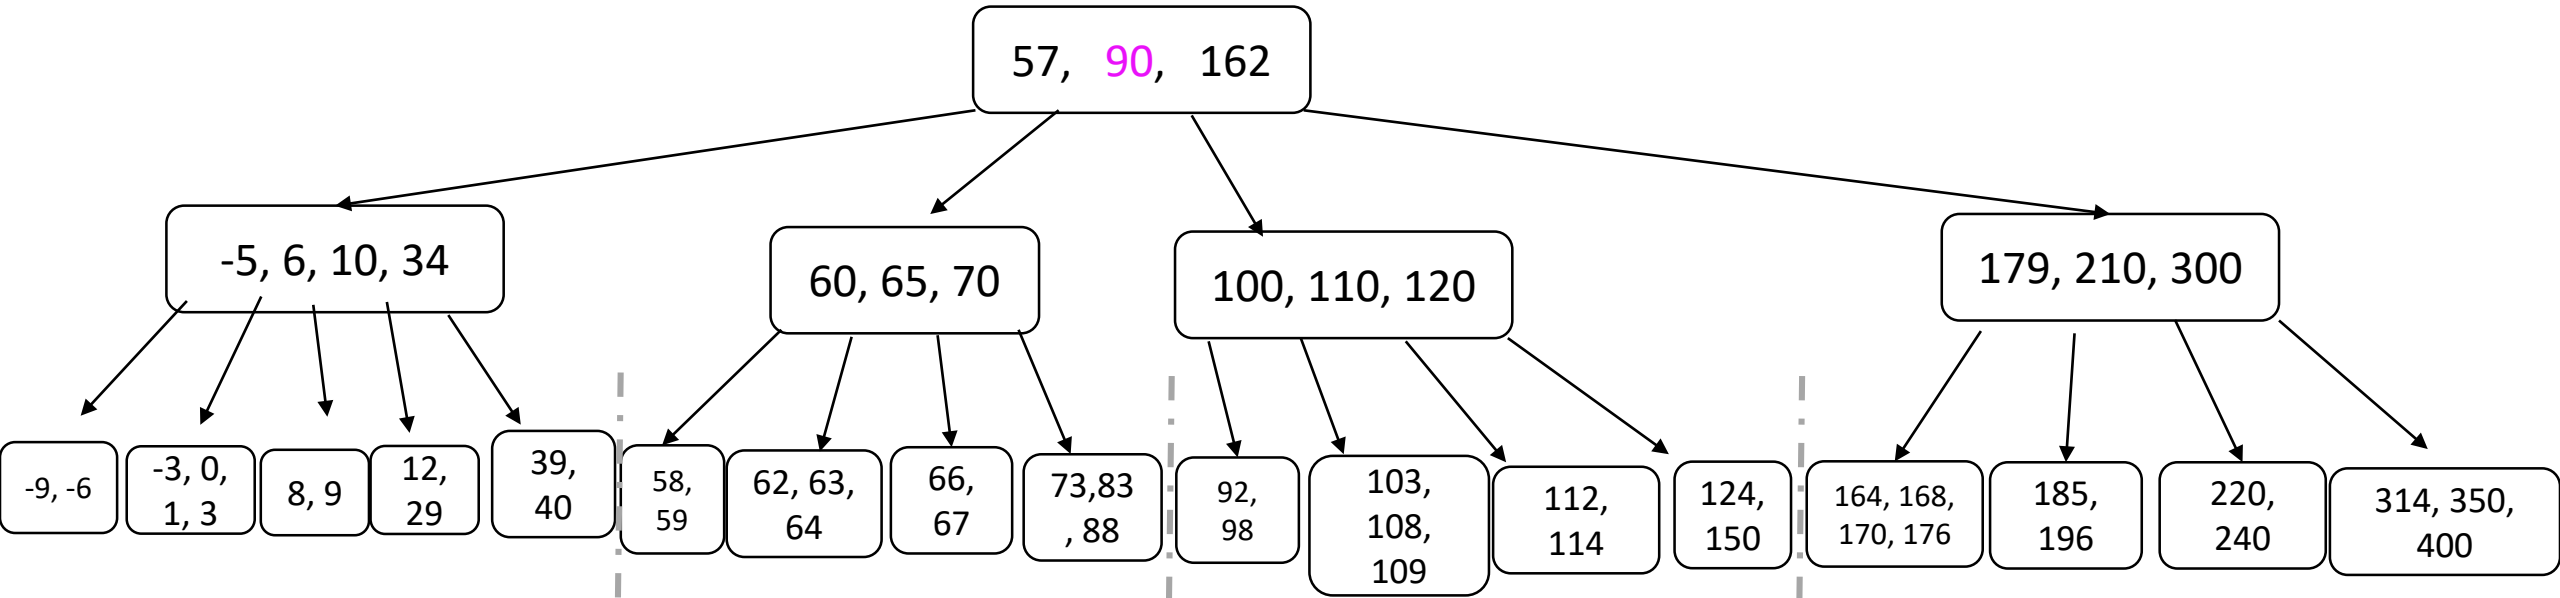
1. Only 2 nodes with maximum amount of $p - 1 = 4$ keys: That's ok.
2. Bunch of leaf nodes holding only $2 = \left\lceil \frac{p}{2} \right\rceil - 1$ keys; also ok
3. Bunch of inner and leaf nodes holding $3 = \left\lceil \frac{p}{2} \right\rceil$ keys; also ok

Example #1: Pick either sibling



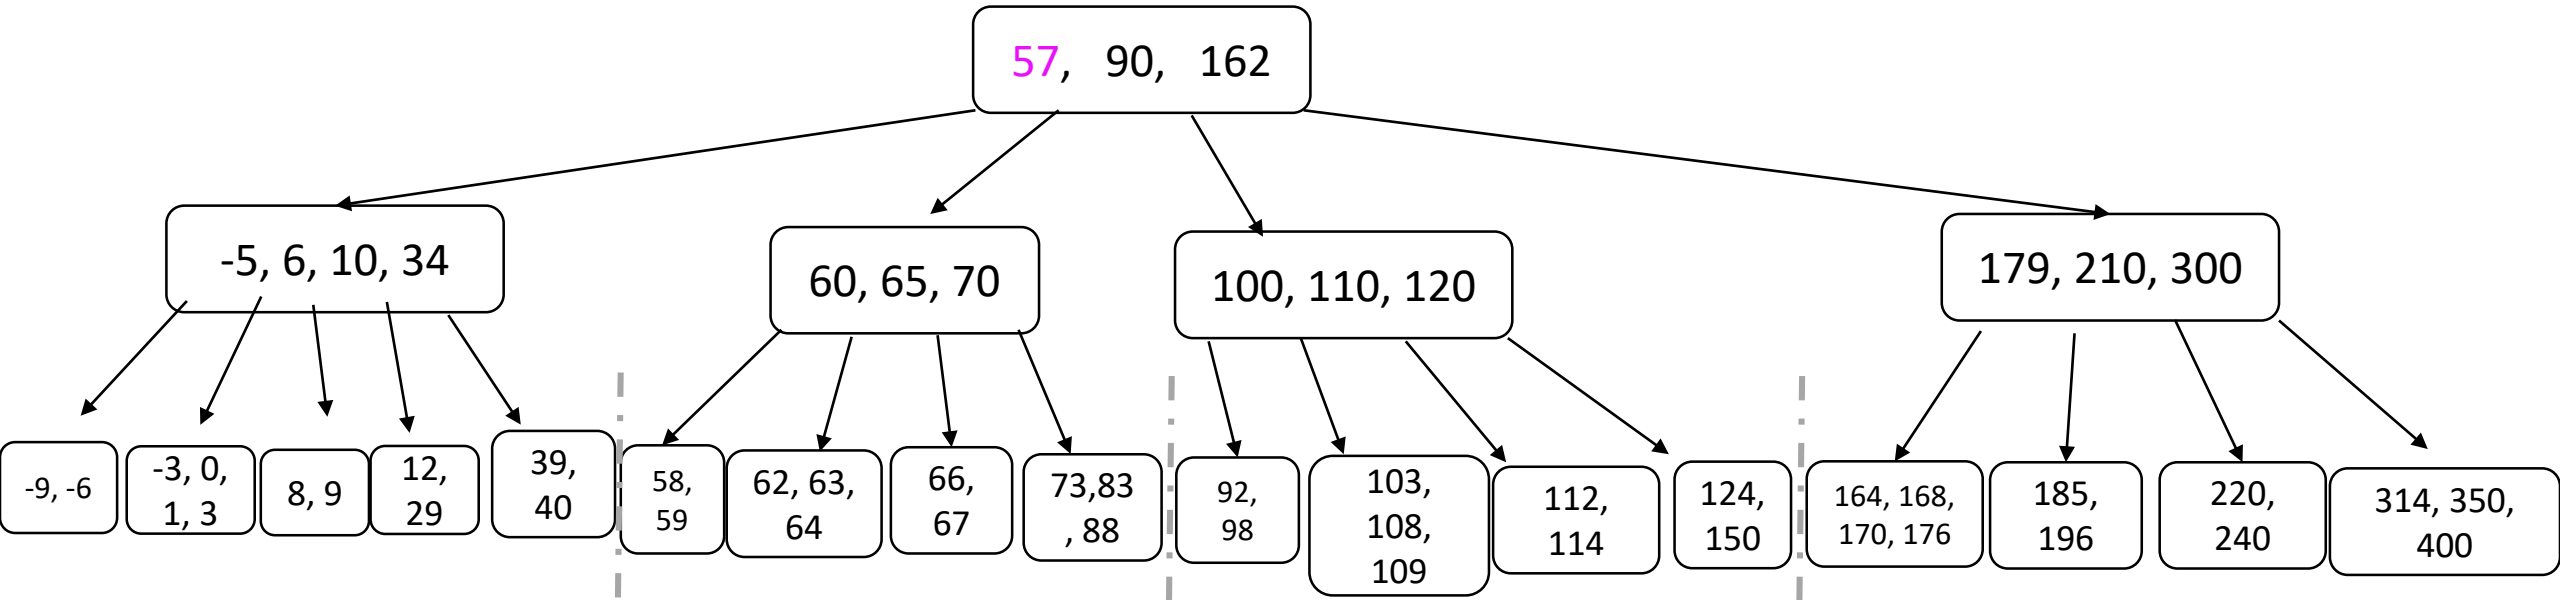
- Task: Insert -1

Example #1: Pick either sibling



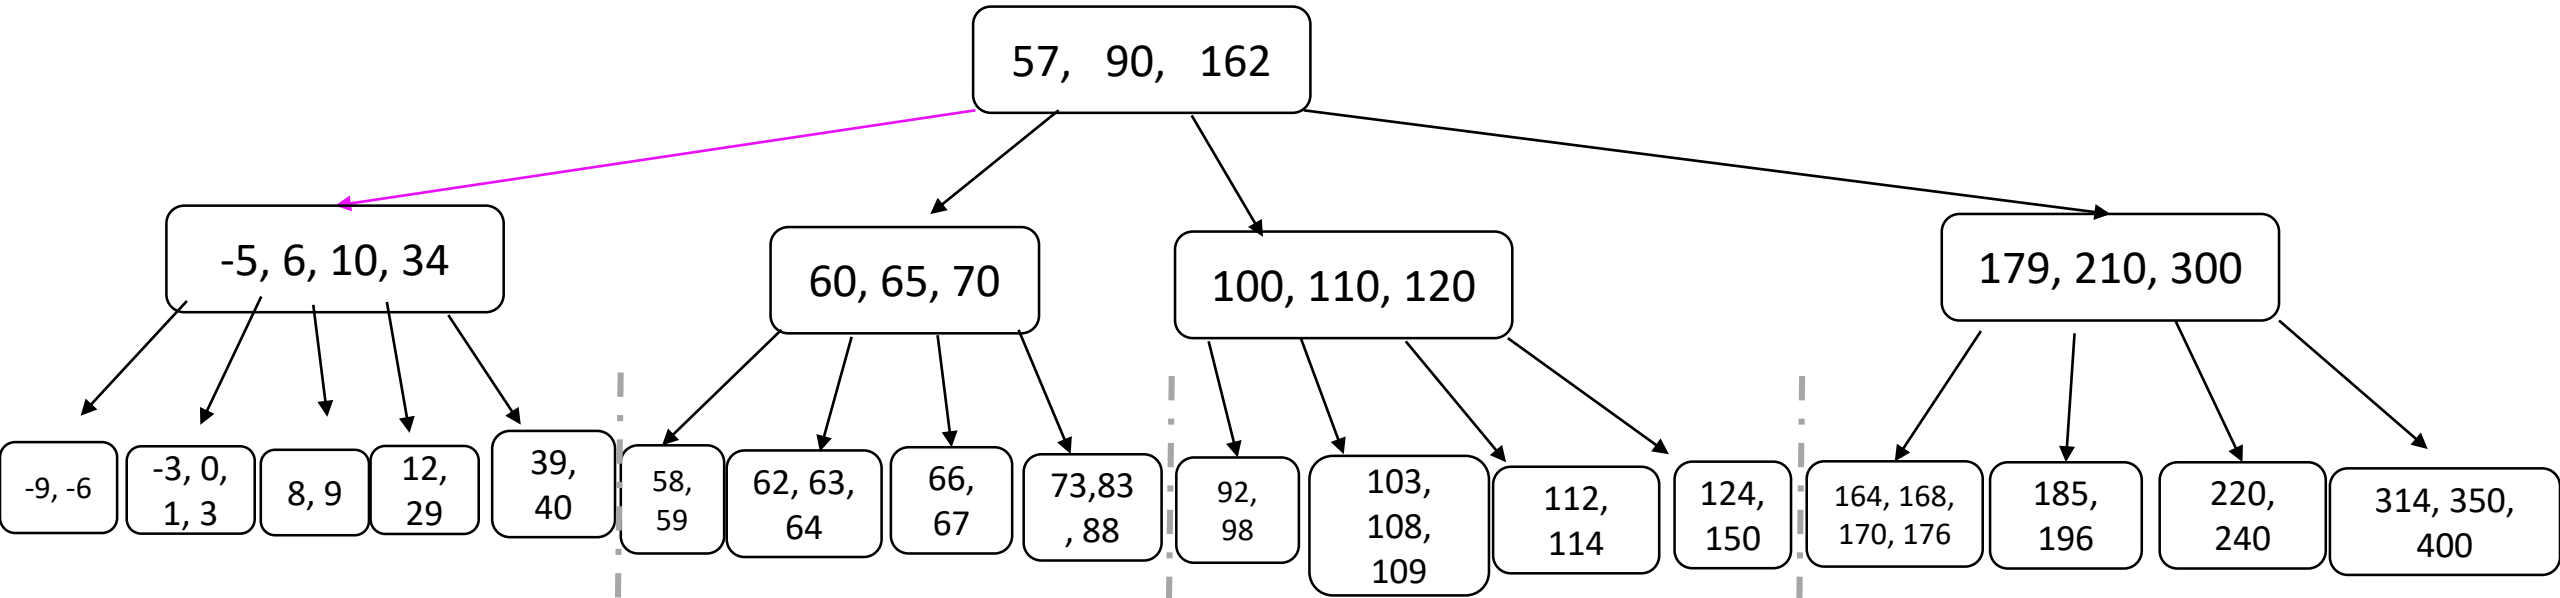
- Task: Insert -1

Example #1: Pick either sibling



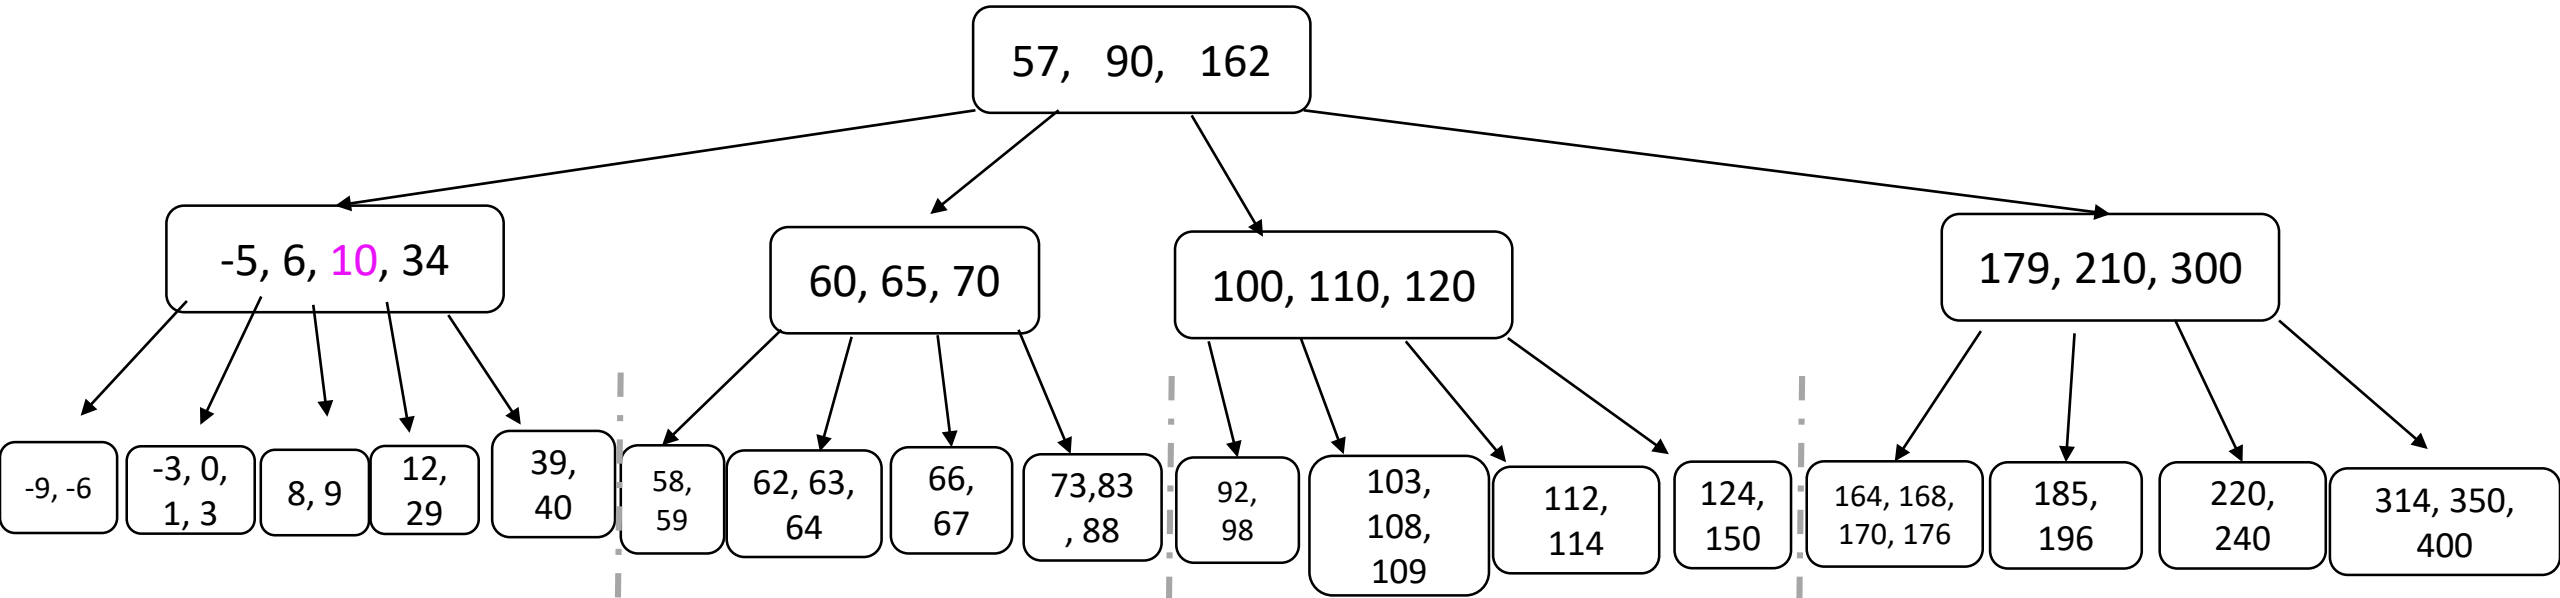
- Task: Insert -1

Example #1: Pick either sibling



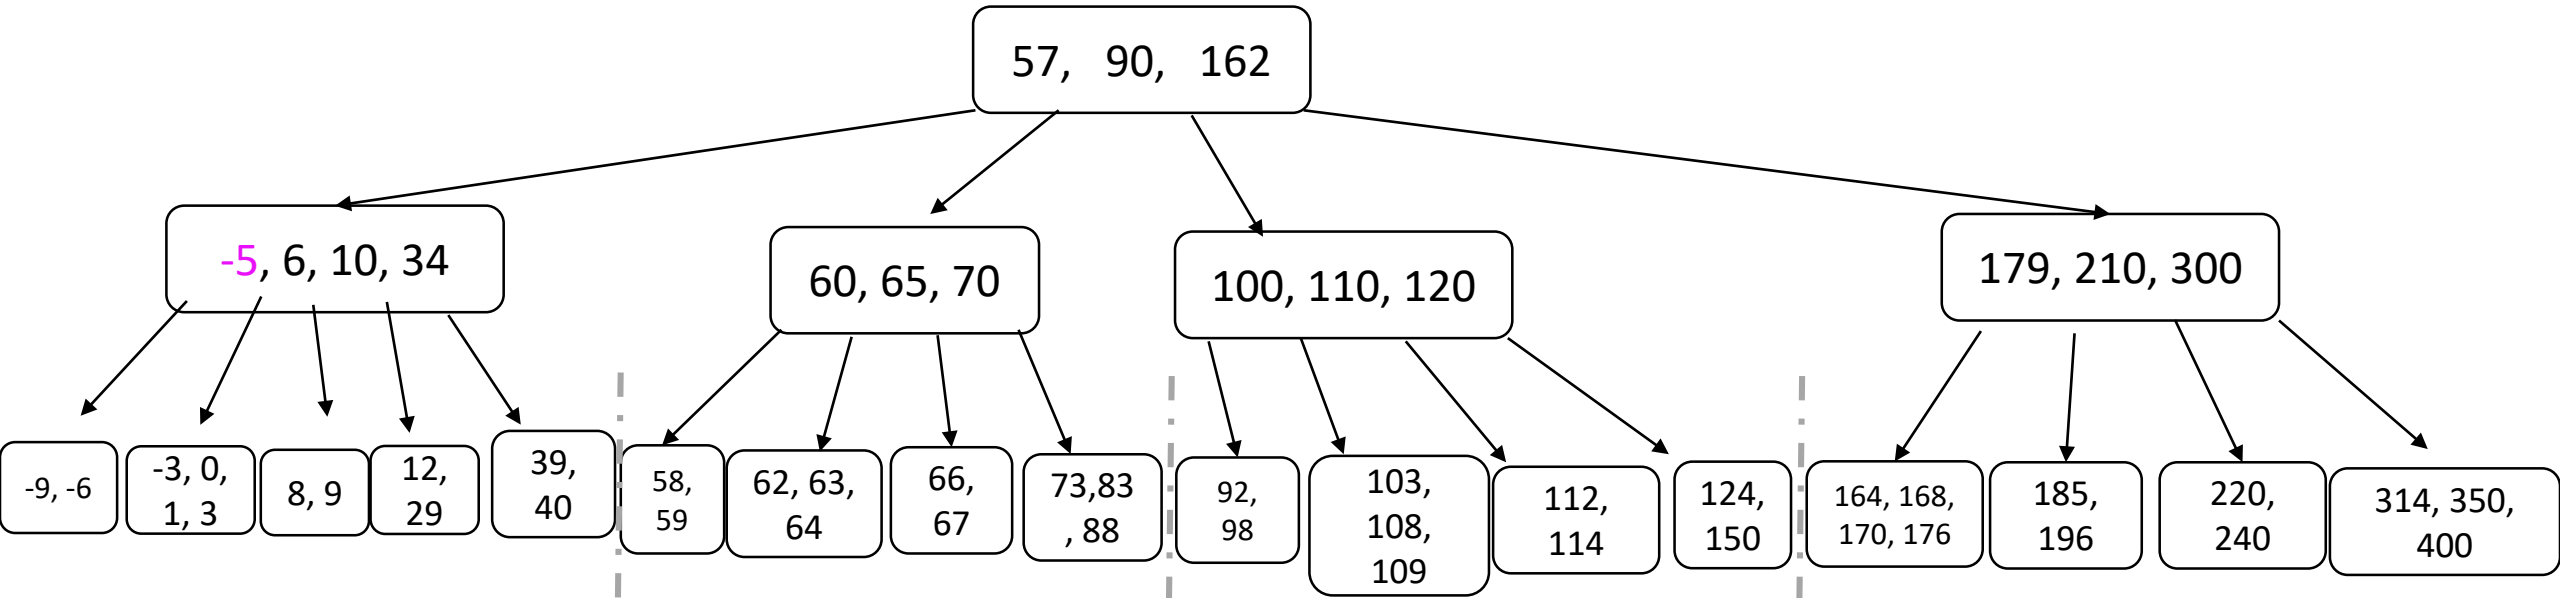
- Task: Insert -1

Example #1: Pick either sibling



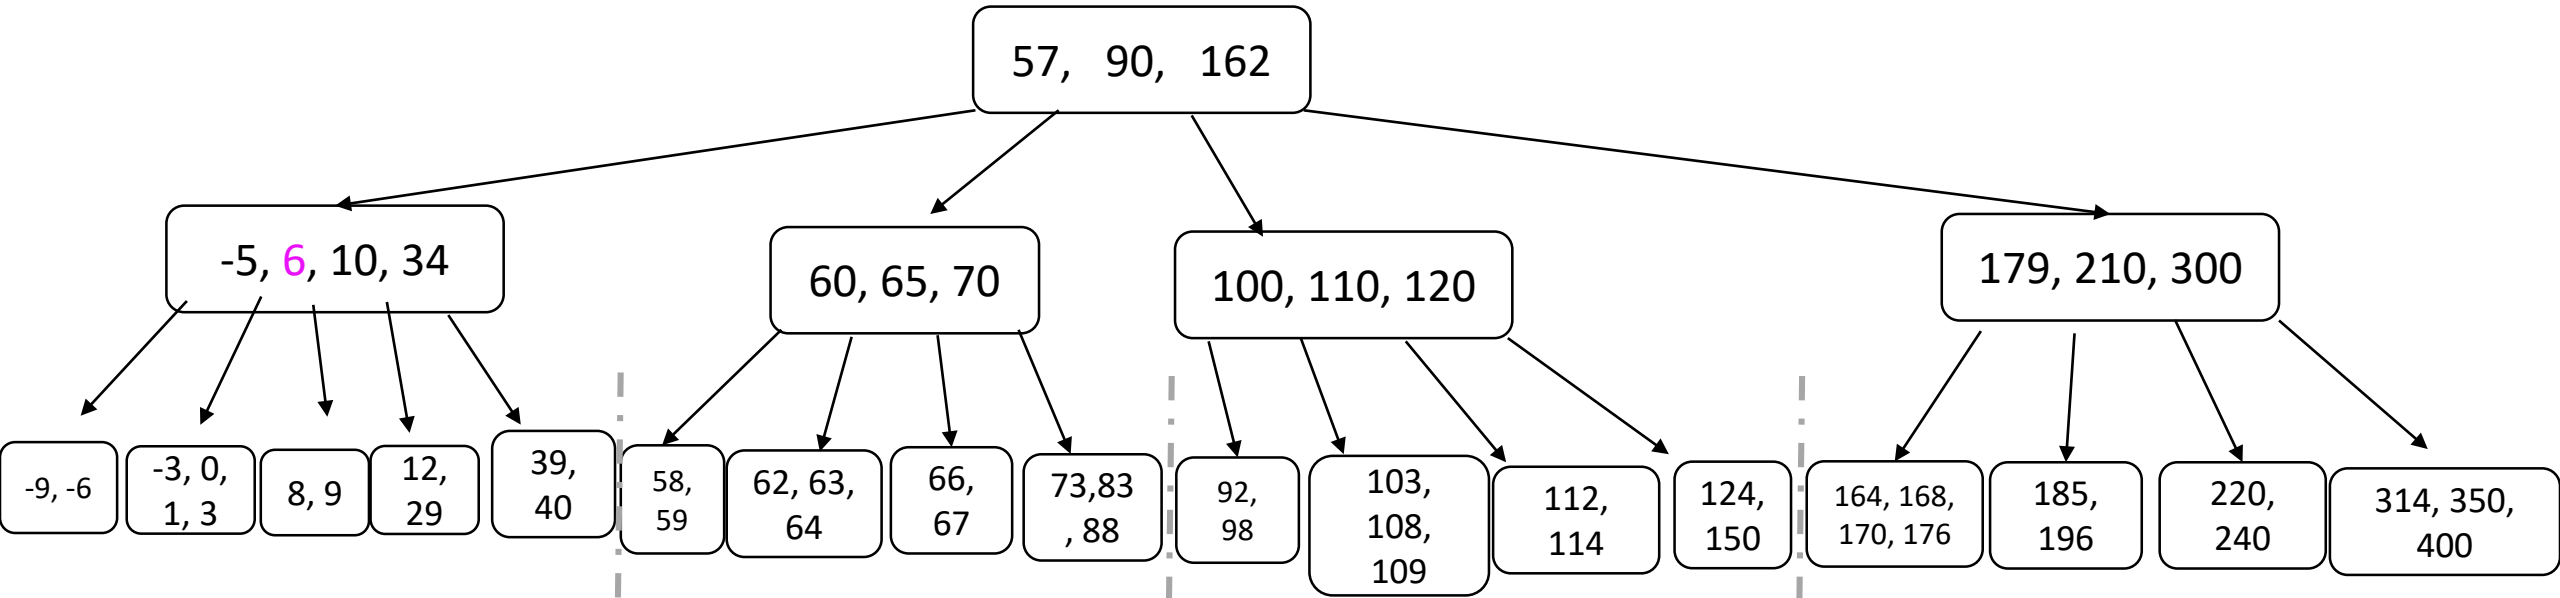
- Task: Insert -1

Example #1: Pick either sibling



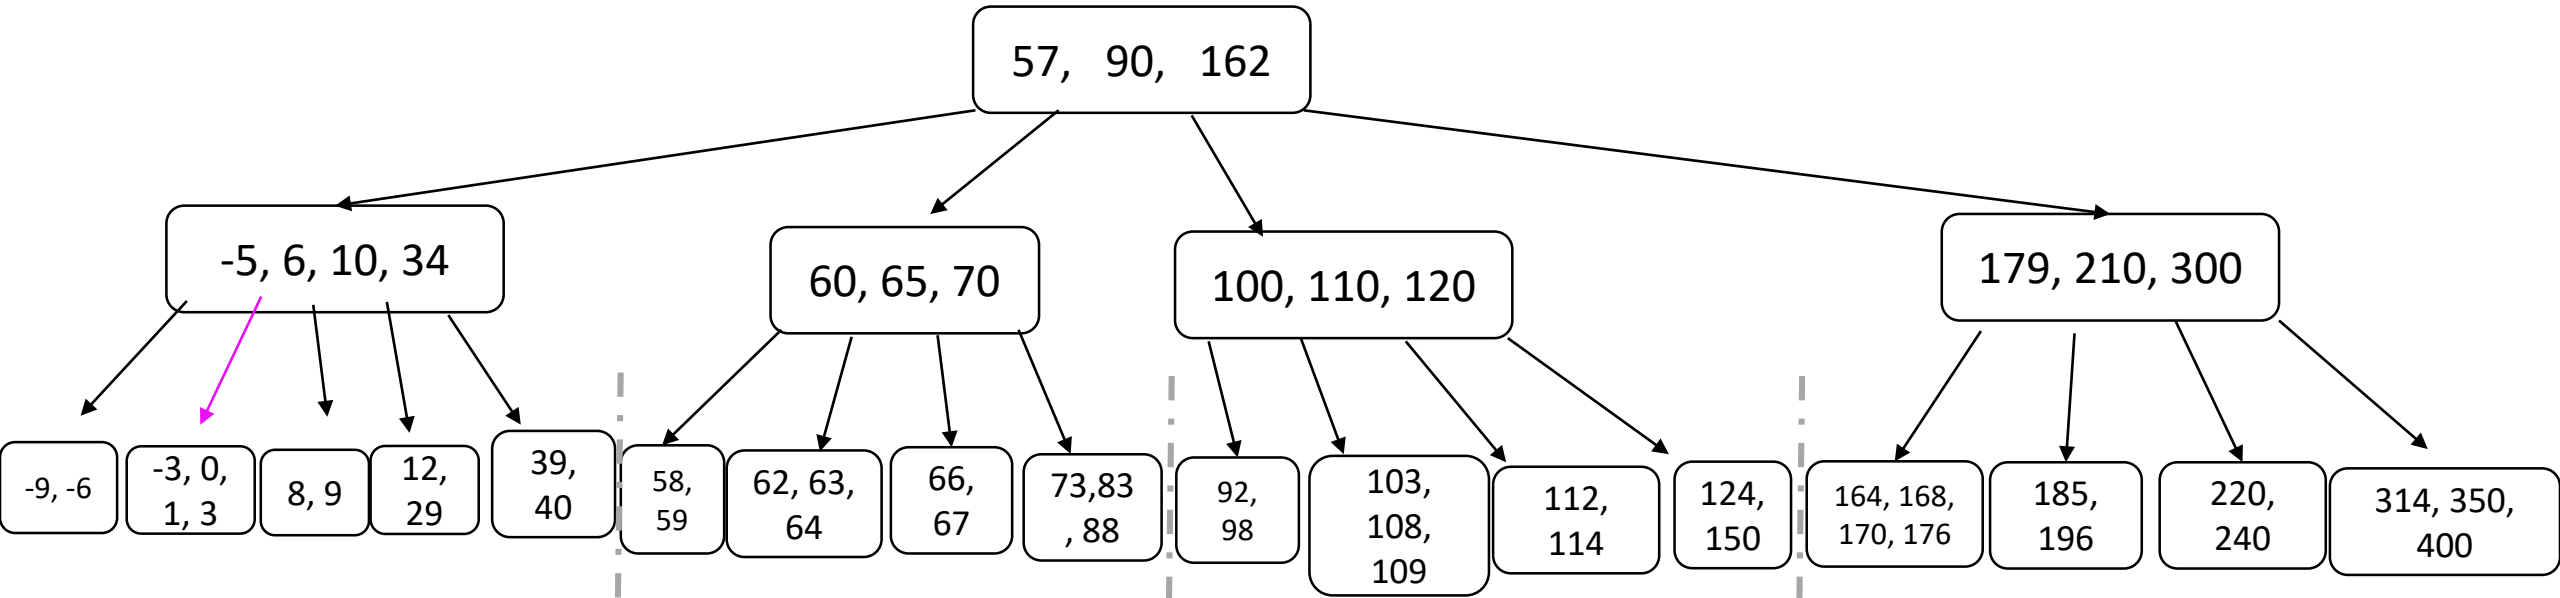
- Task: Insert -1

Example #1: Pick either sibling



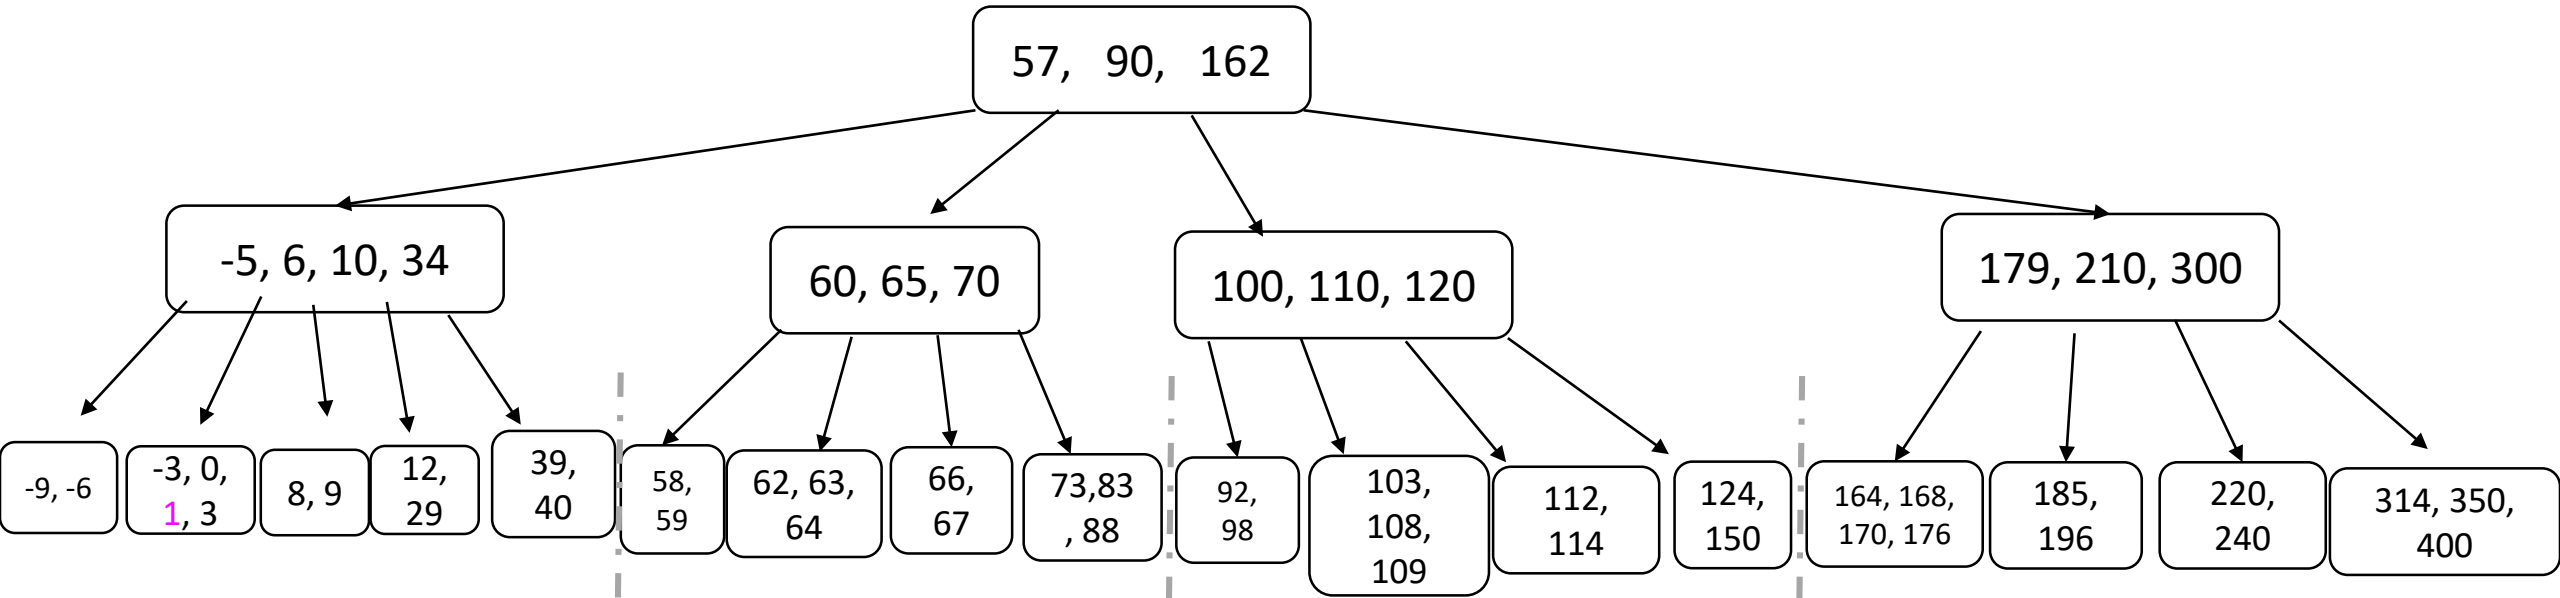
- Task: Insert -1

Example #1: Pick either sibling



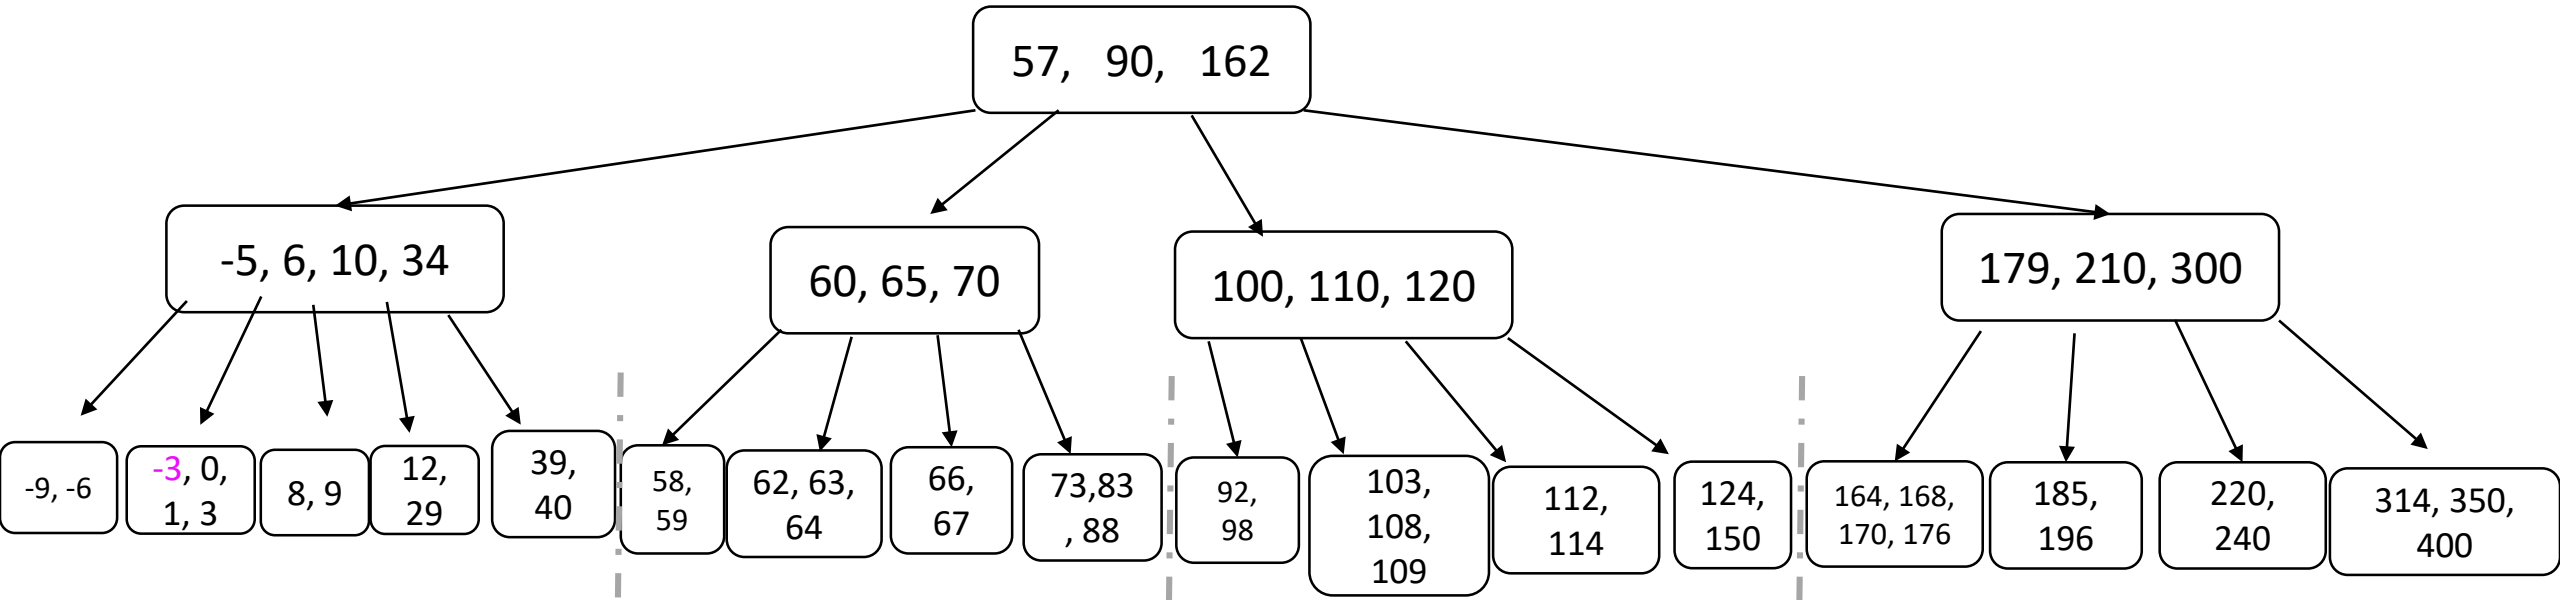
- Task: Insert -1

Example #1: Pick either sibling



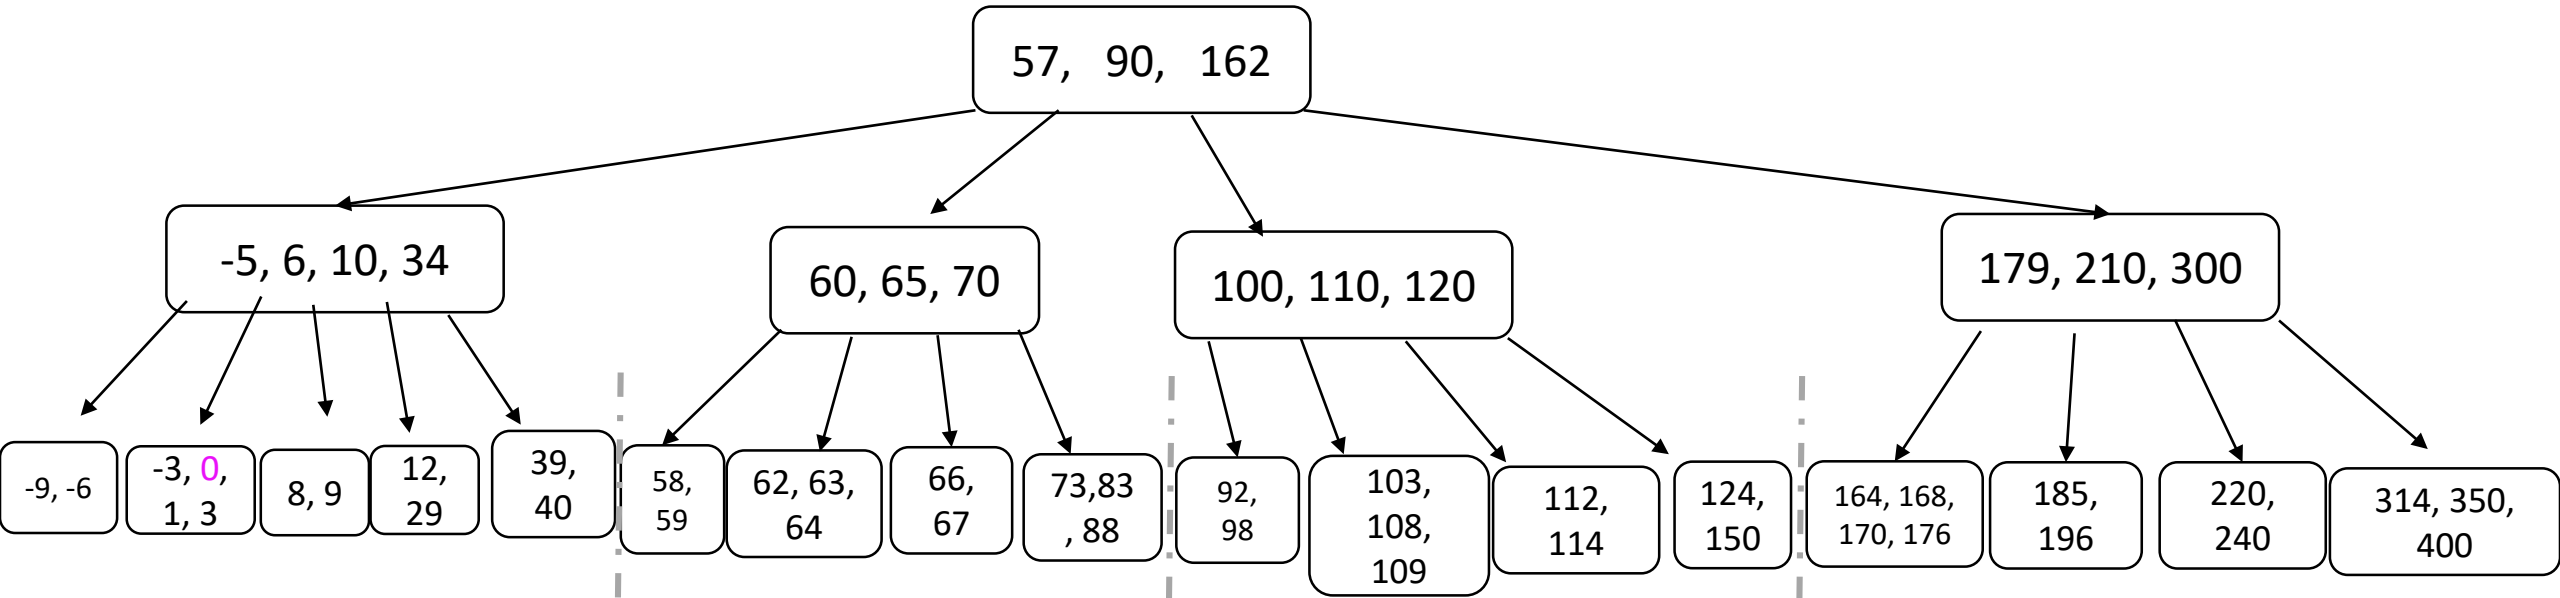
- Task: Insert -1

Example #1: Pick either sibling



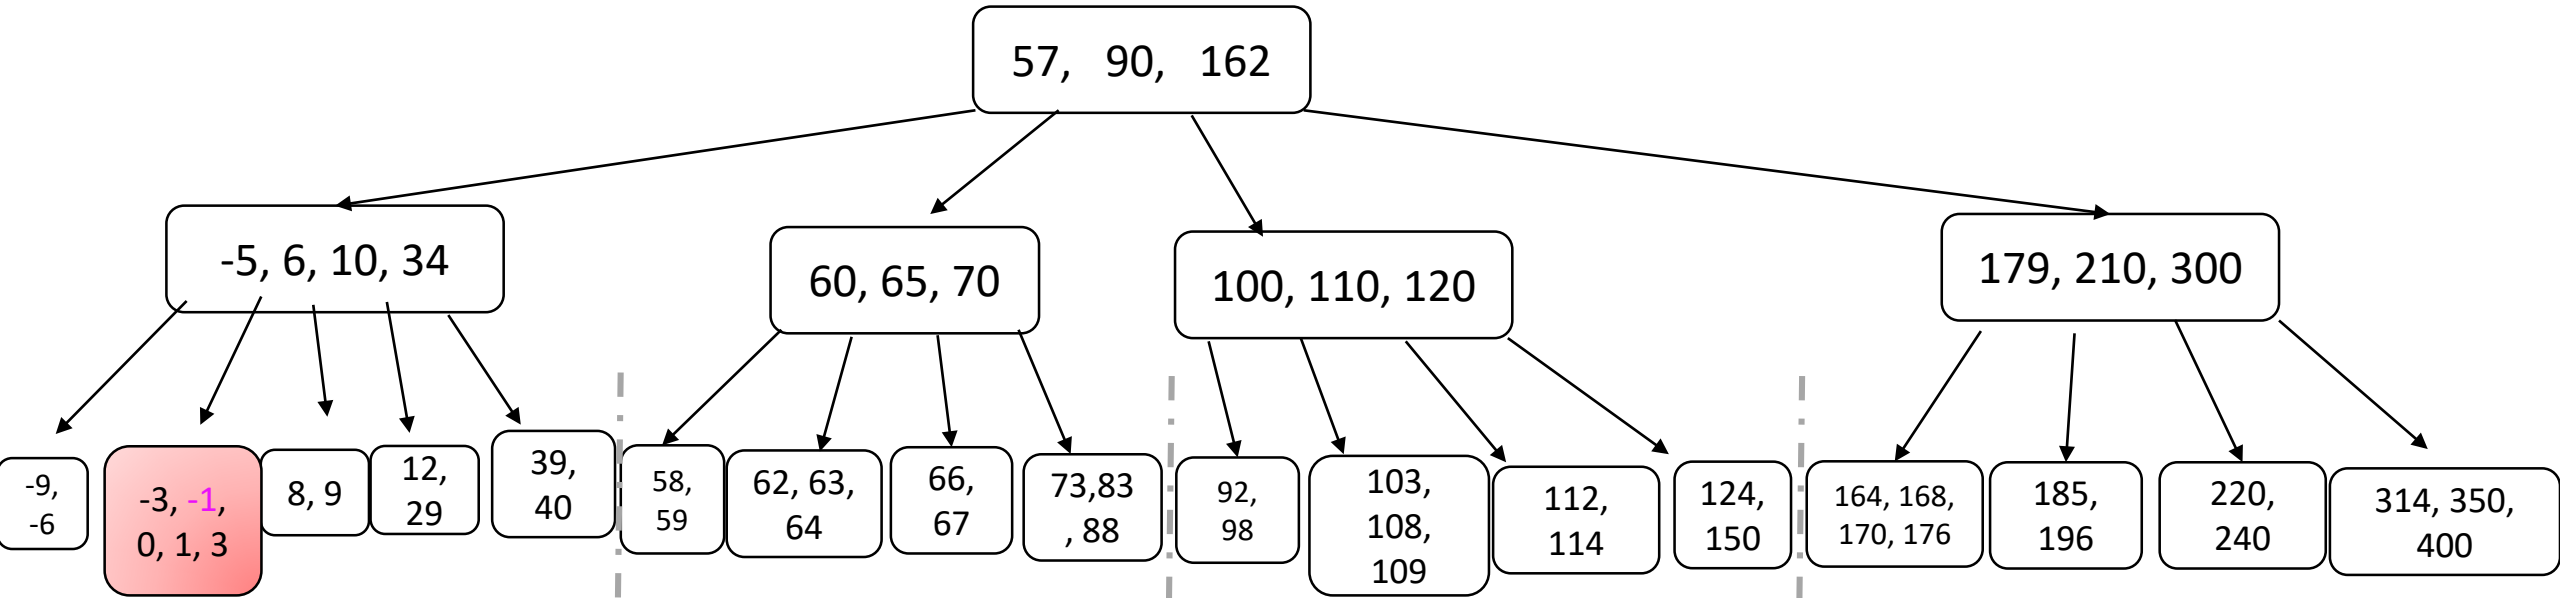
- Task: Insert -1

Example #1: Pick either sibling



- Task: Insert -1

Example #1: Pick either sibling

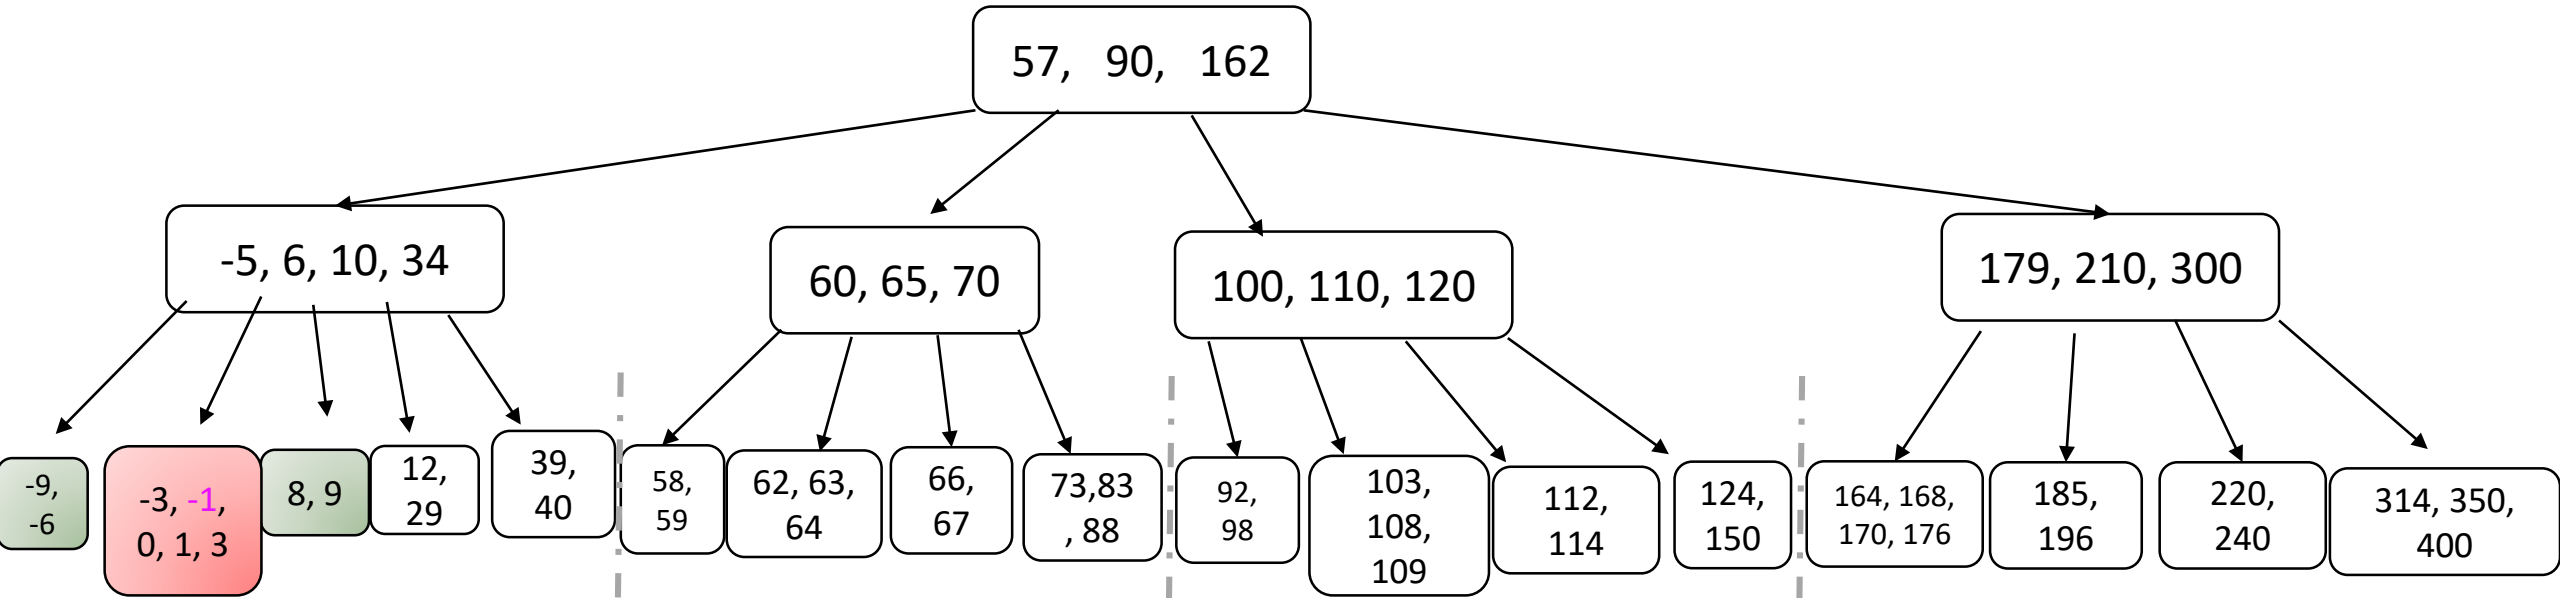


- Task: Insert -1

Overflow ☹️



Example #1: Pick either sibling

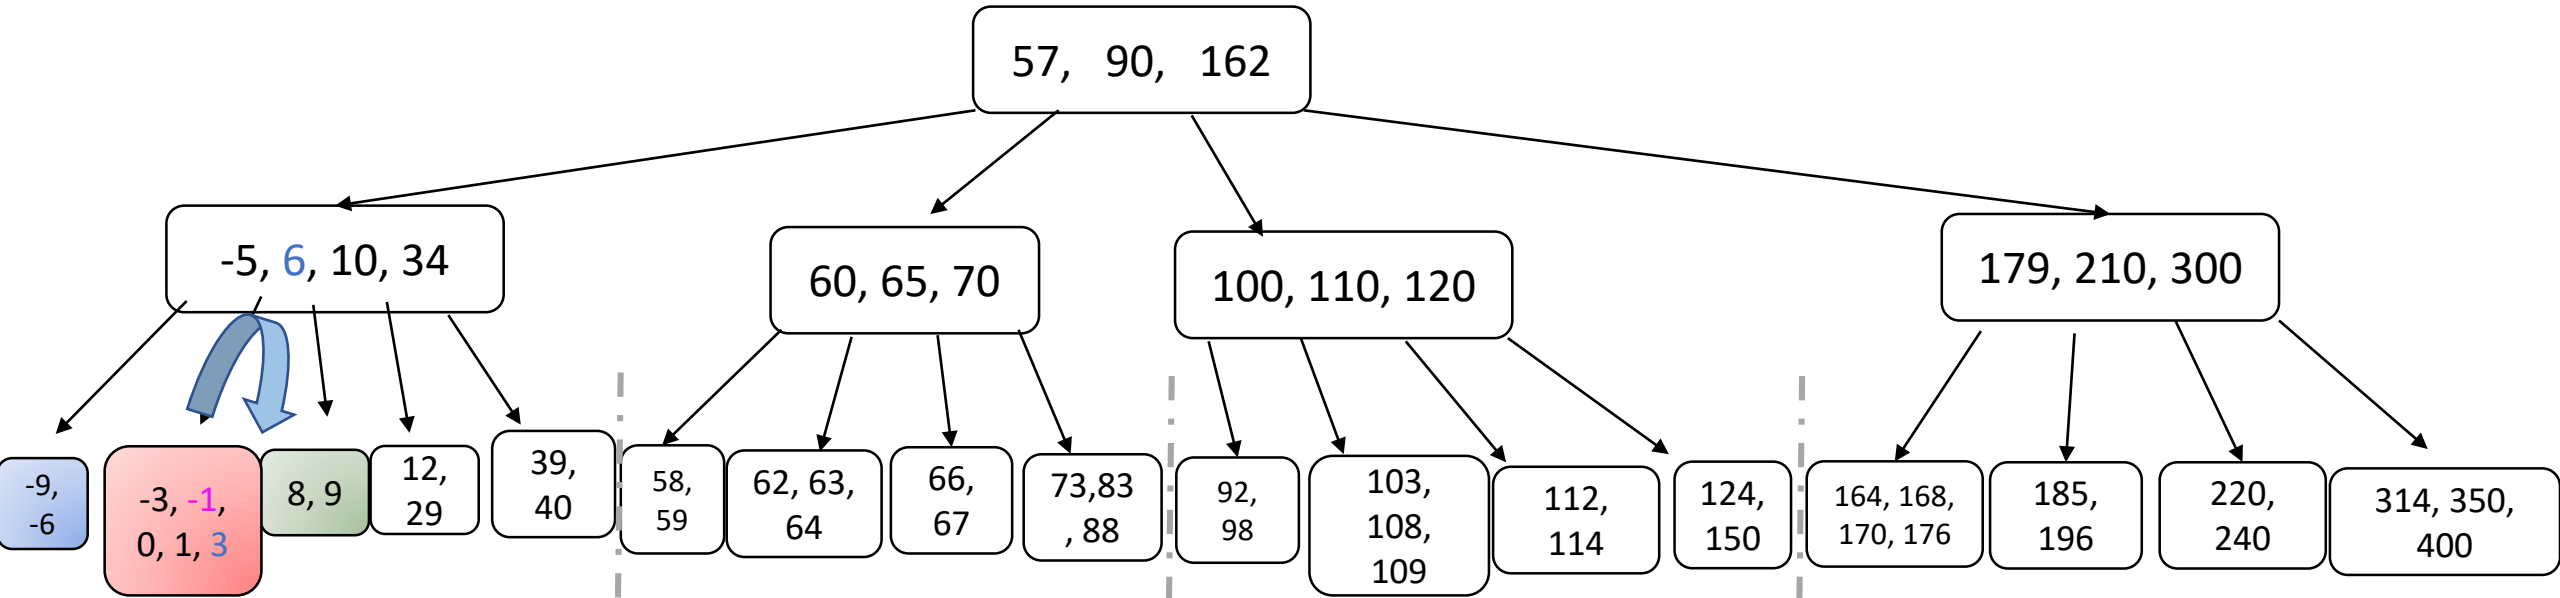


- Task: Insert -1

No worries! I can rotate to either sibling 😊



Example #1: Pick either sibling

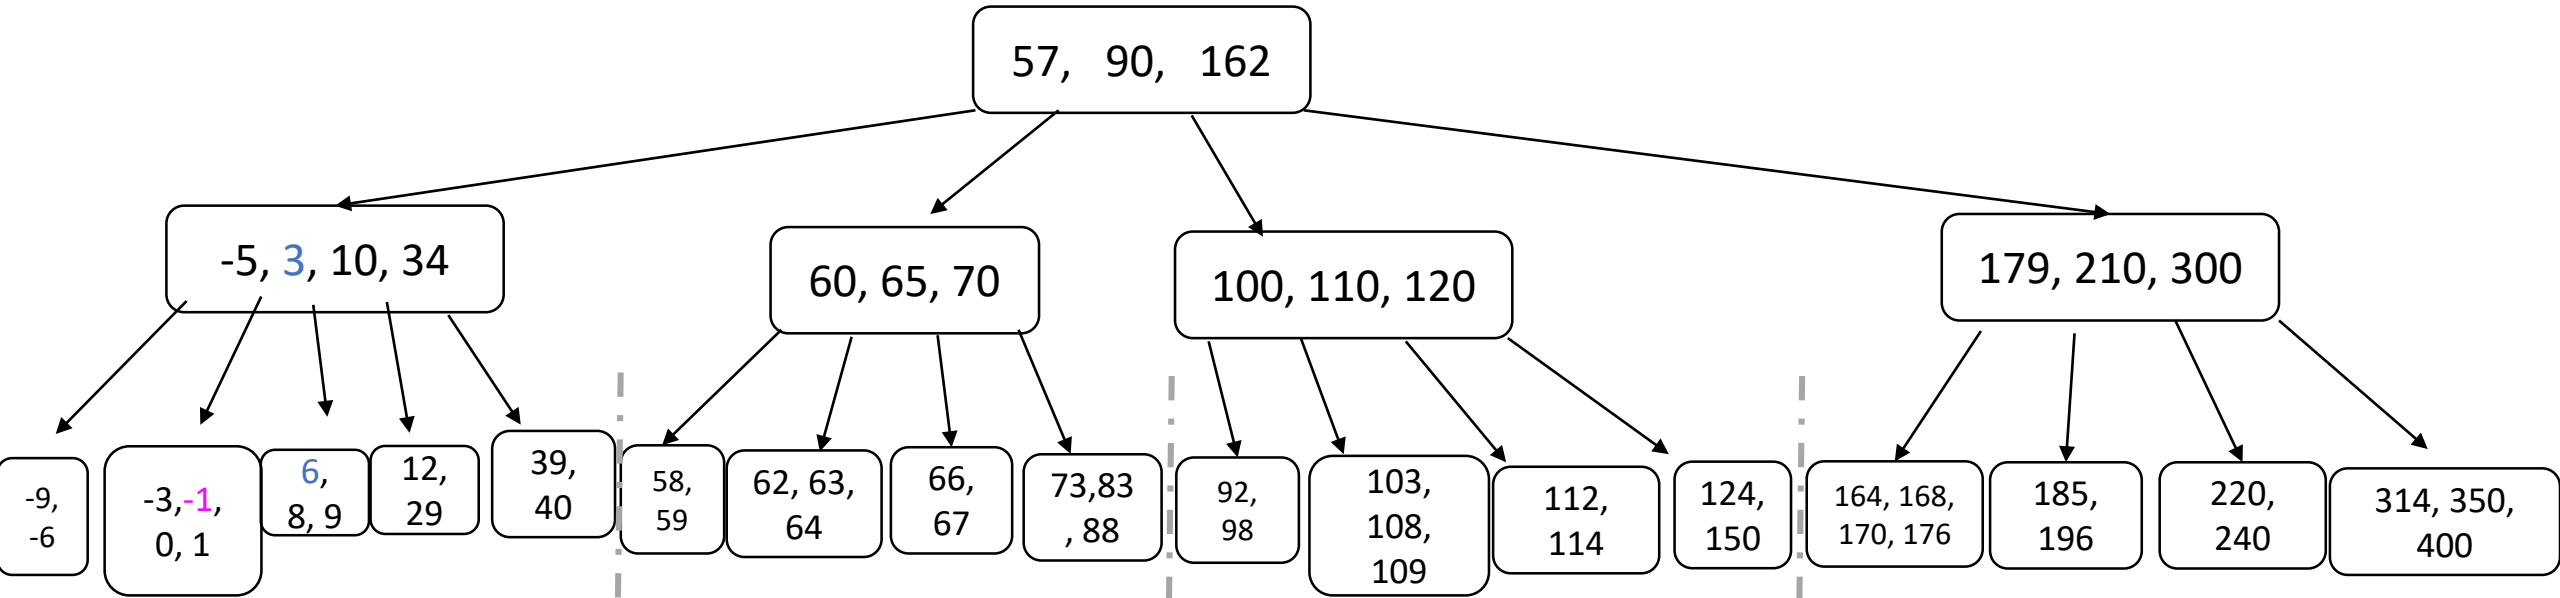


- Task: Insert -1

E.g pick right
sibling, rotate 3
rightwards

(This is a remnant from an older semester: In Fall 19, we make the convention of rotation rightwards first)

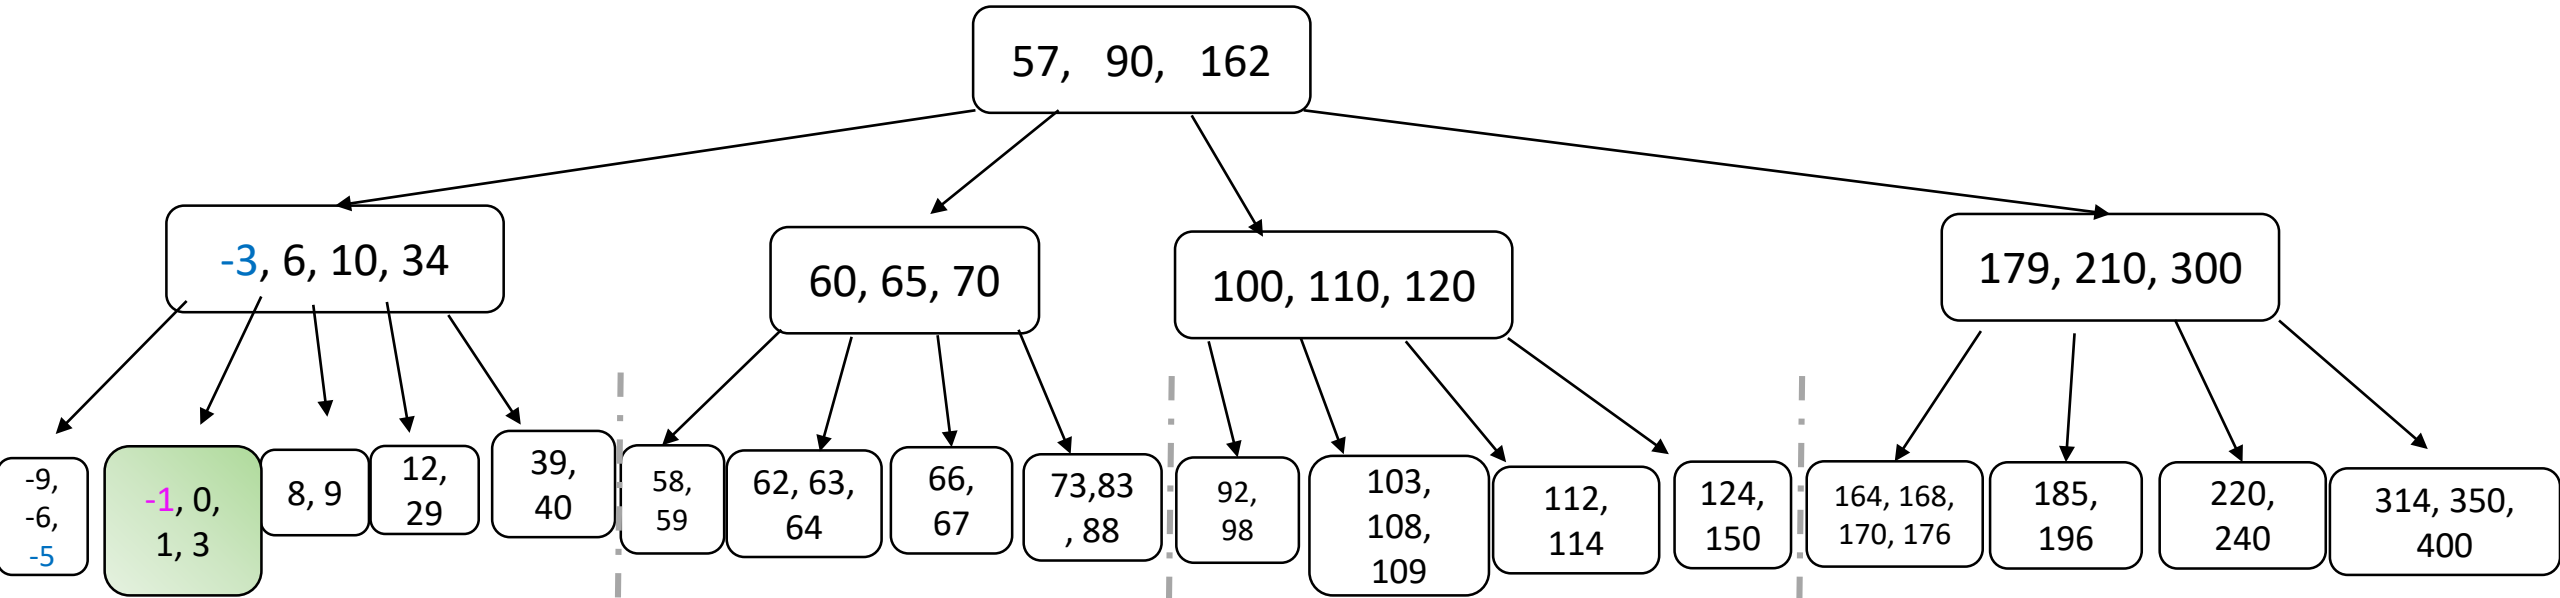
Example #1: Pick either sibling



- Task: Insert -1

Careful: This means that the parent's key (6) is the one that lands on the sibling node... 3 goes to the parent!

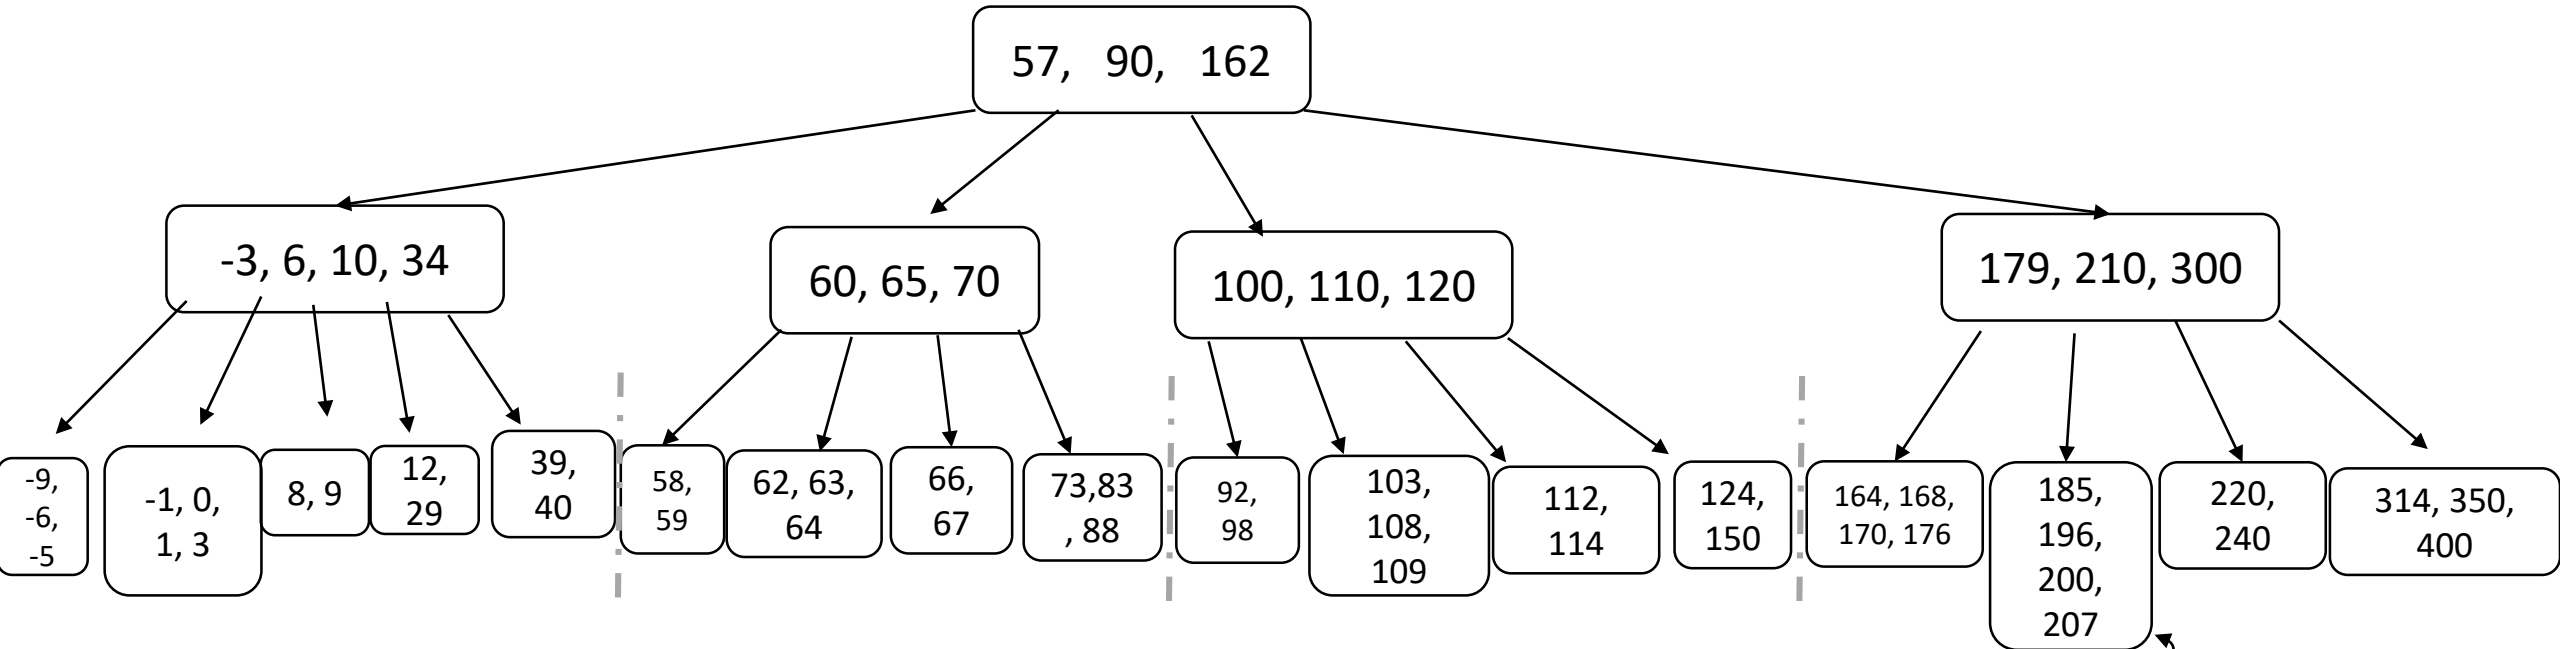
Example #1: Pick either sibling



- Task: Insert -1

Overflow fixed, key inserted 😊

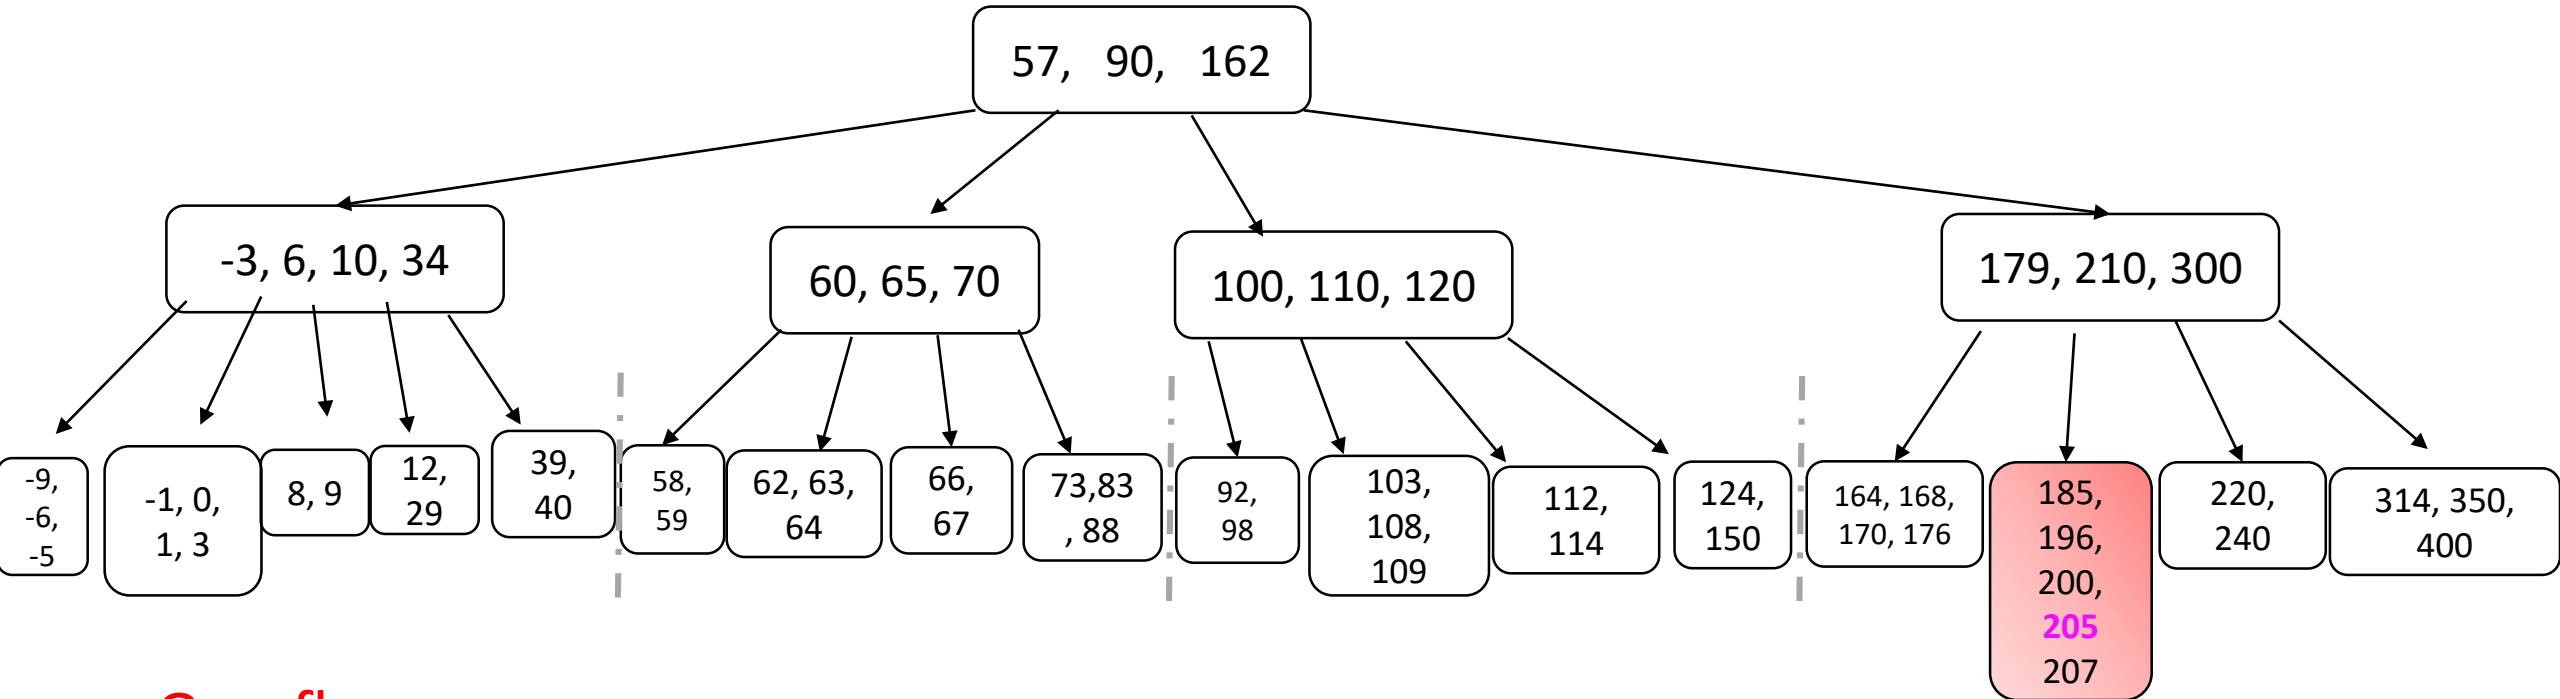
Example #2: Only one sibling with space!



- Task: Insert 205

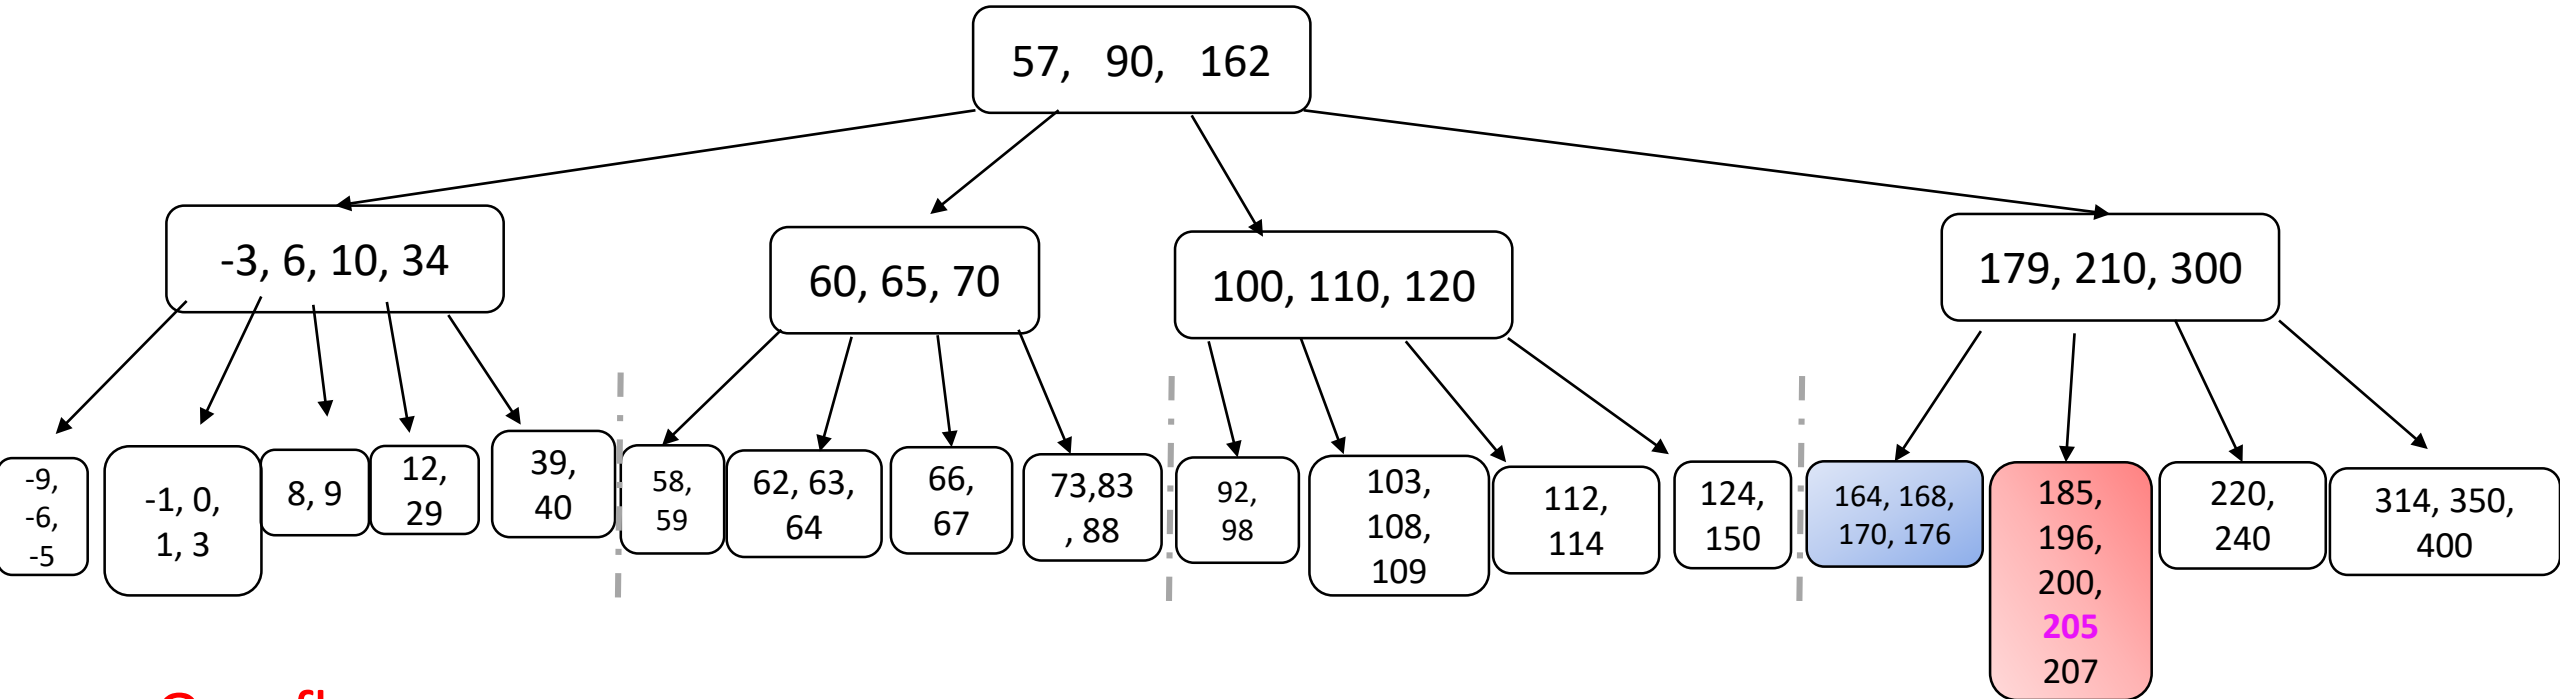
Let's add some more keys here to make this insertion interesting....

Example #2: Only one sibling with space!



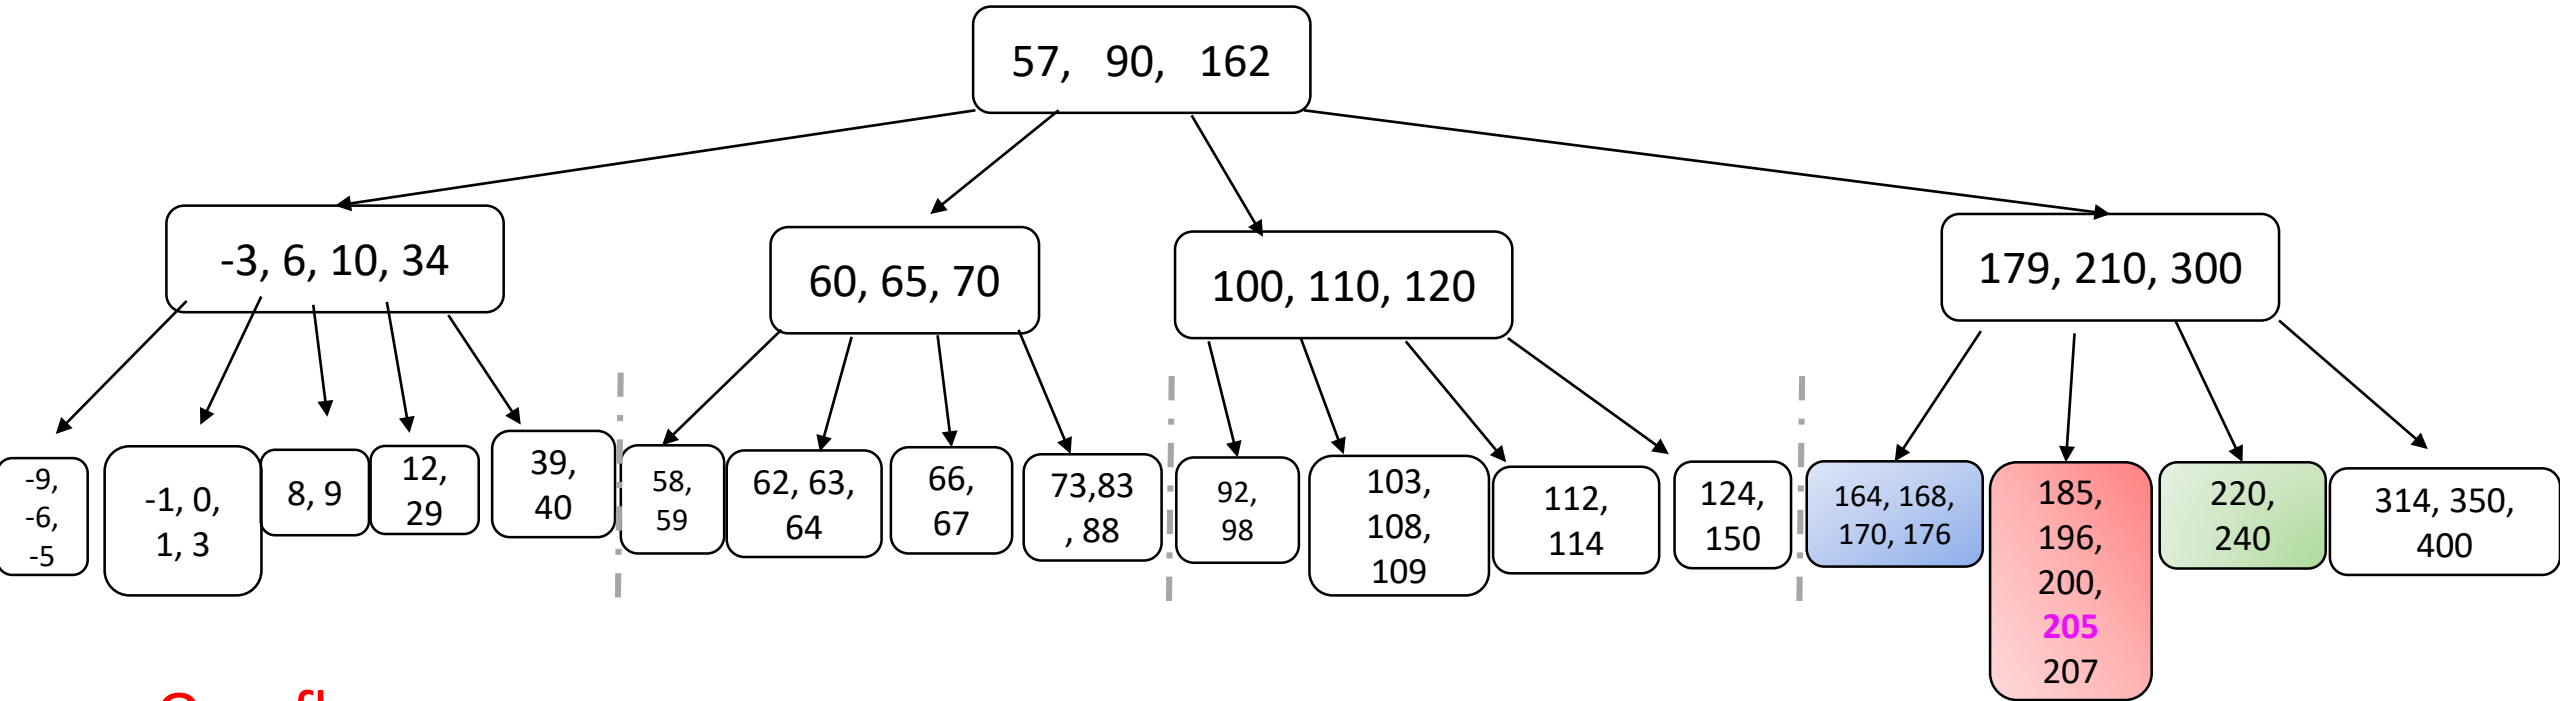
- Overflow...

Example #2: Only one sibling with space!



- Overflow...
- Can't rotate 185 leftwards because left sibling full ☹️

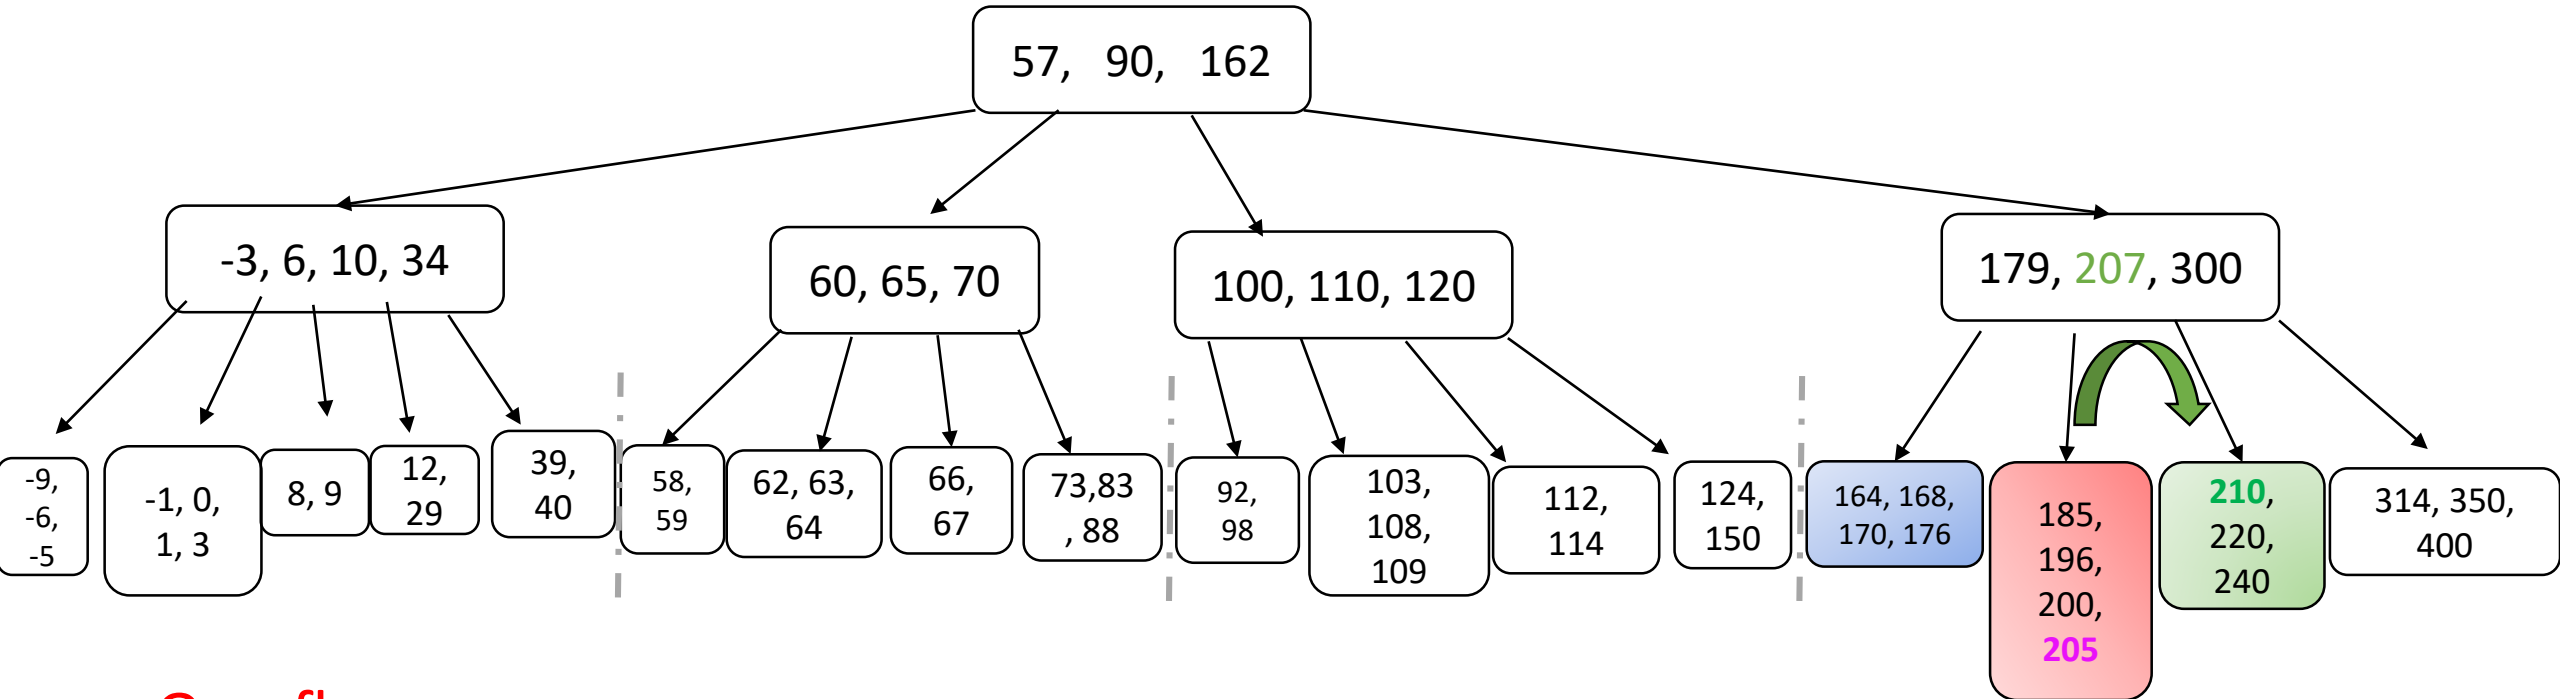
Example #2: Only one sibling with space!



- Overflow...
- Can't rotate 185 leftwards because left sibling full 😞
- But can rotate 207 rightwards because right sibling non-full 😊

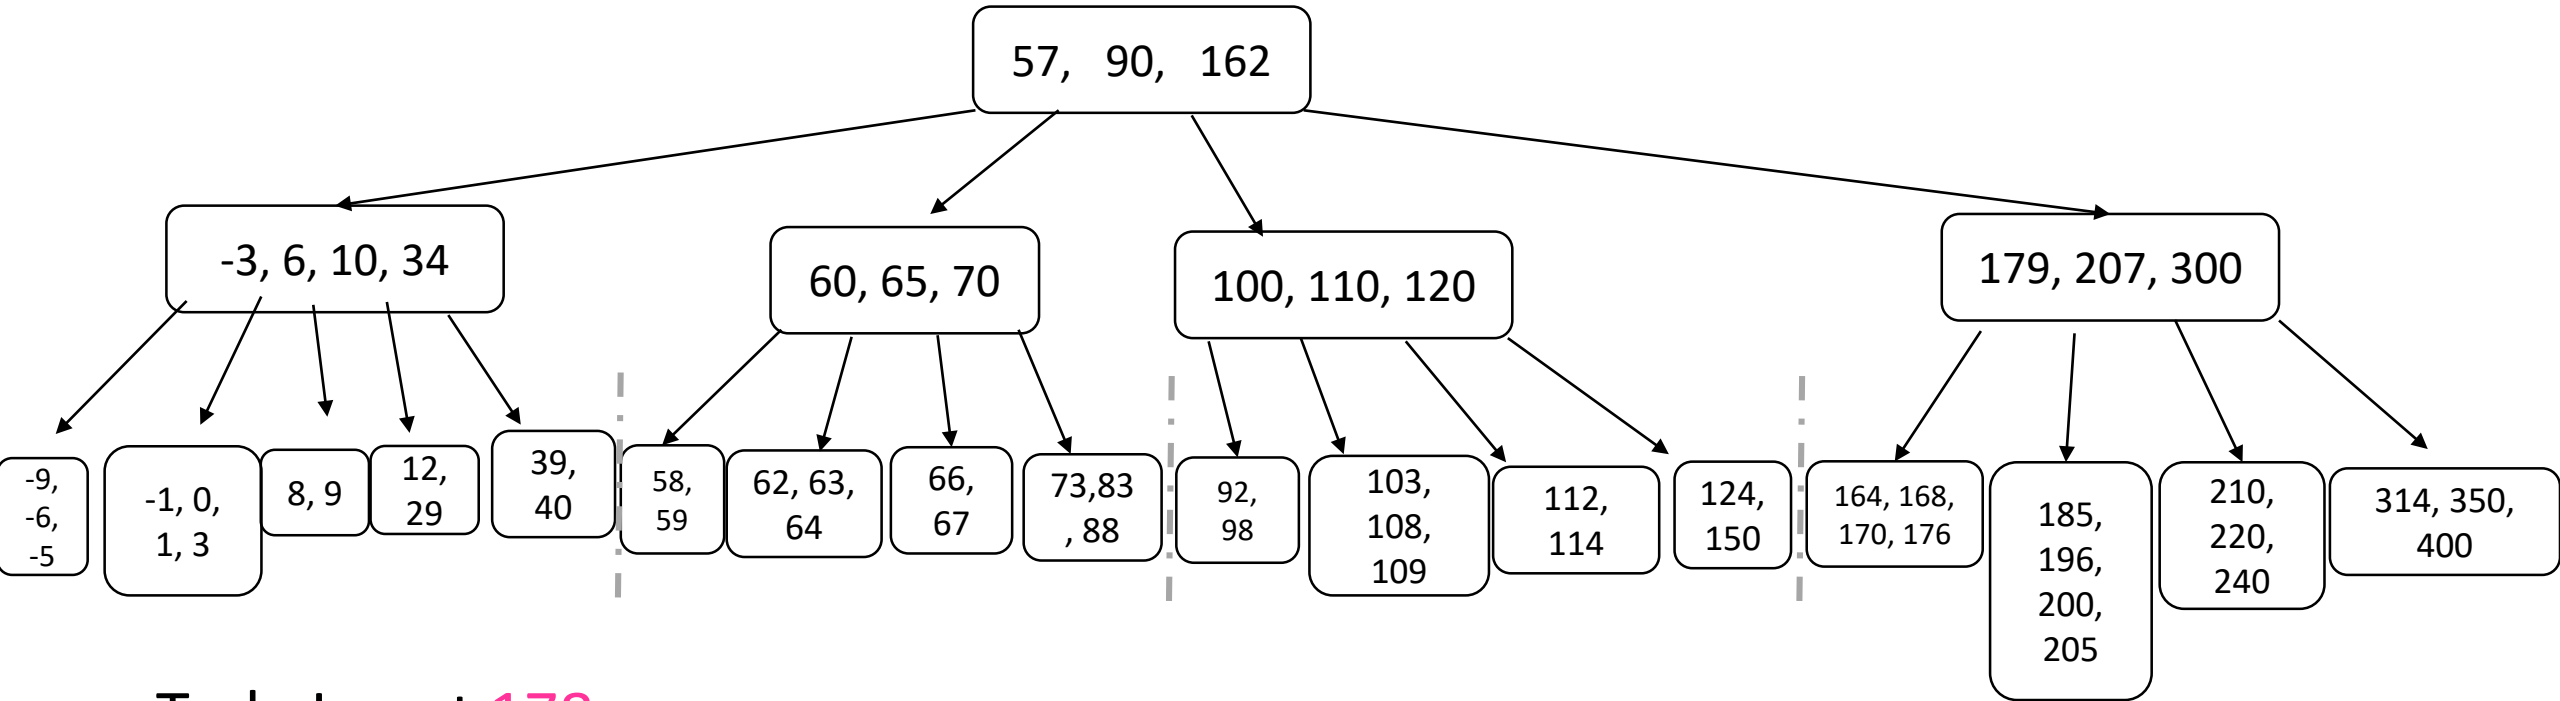
Once again, for Fall 19, we would make the convention of going right first (so we would check it first).

Example #2: Only one sibling with space!



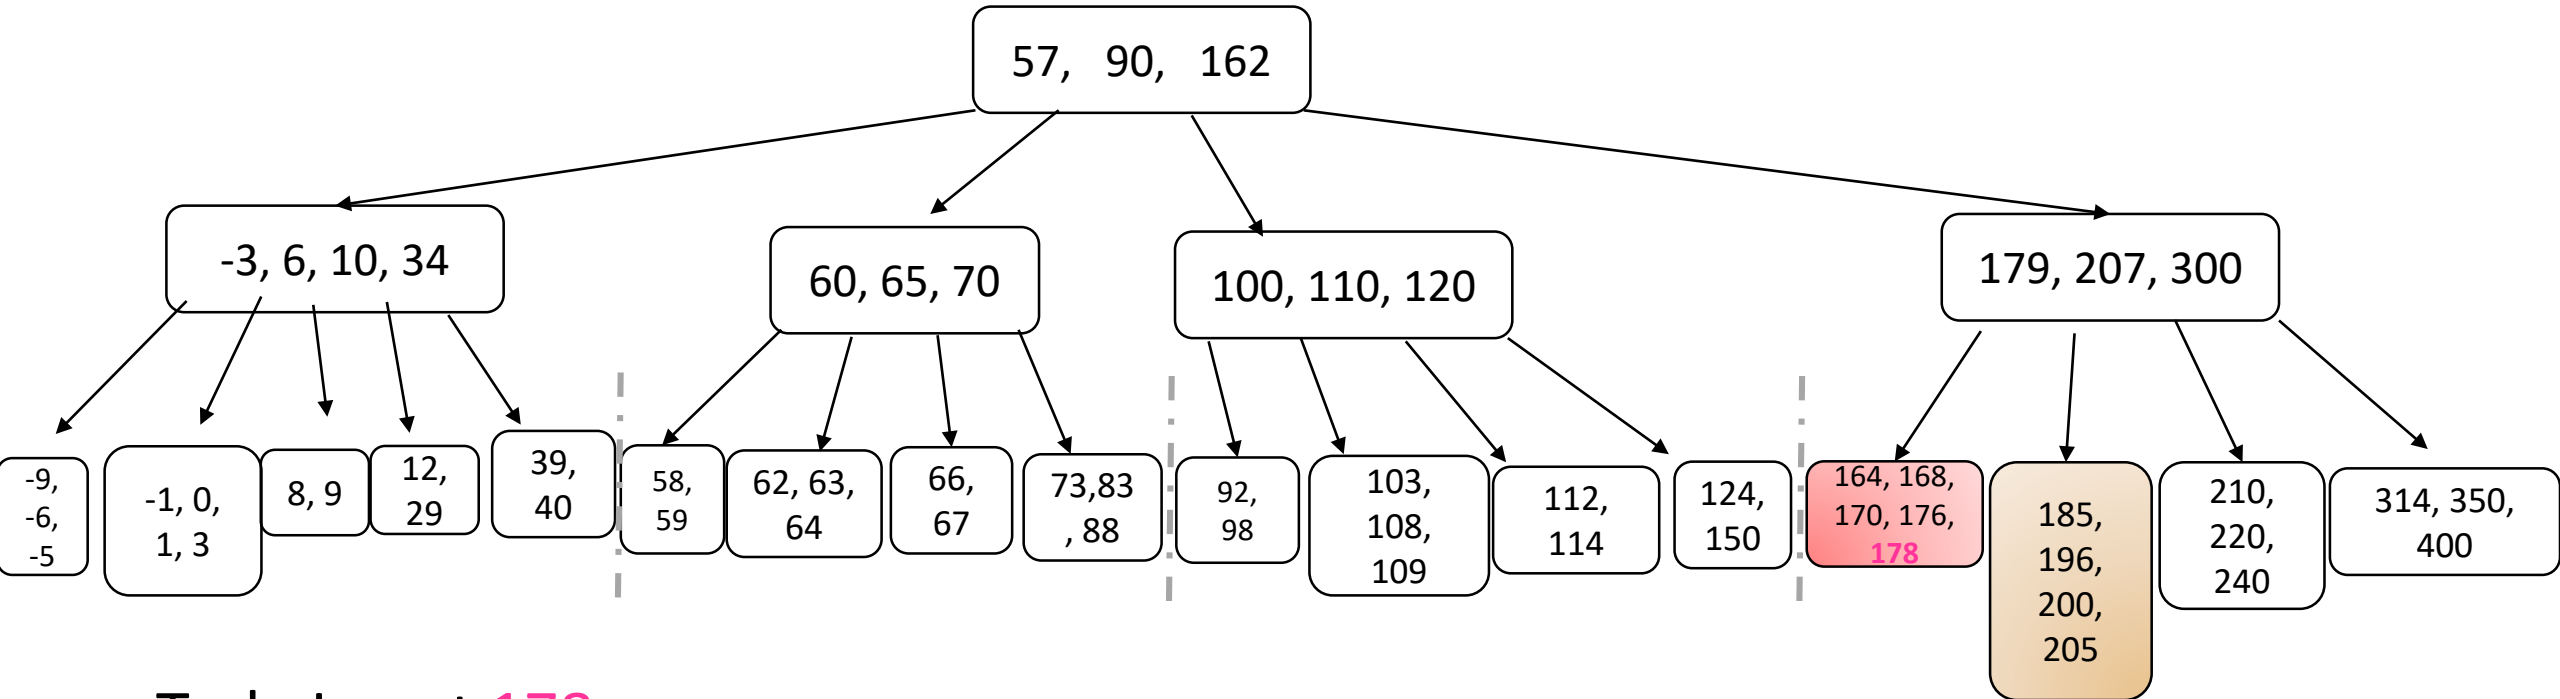
- Overflow...
- Can't rotate 185 leftwards because left sibling full ☹️
- But can rotate 207 rightwards because right sibling non-full 😊

Example #3: Examine first sibling node!



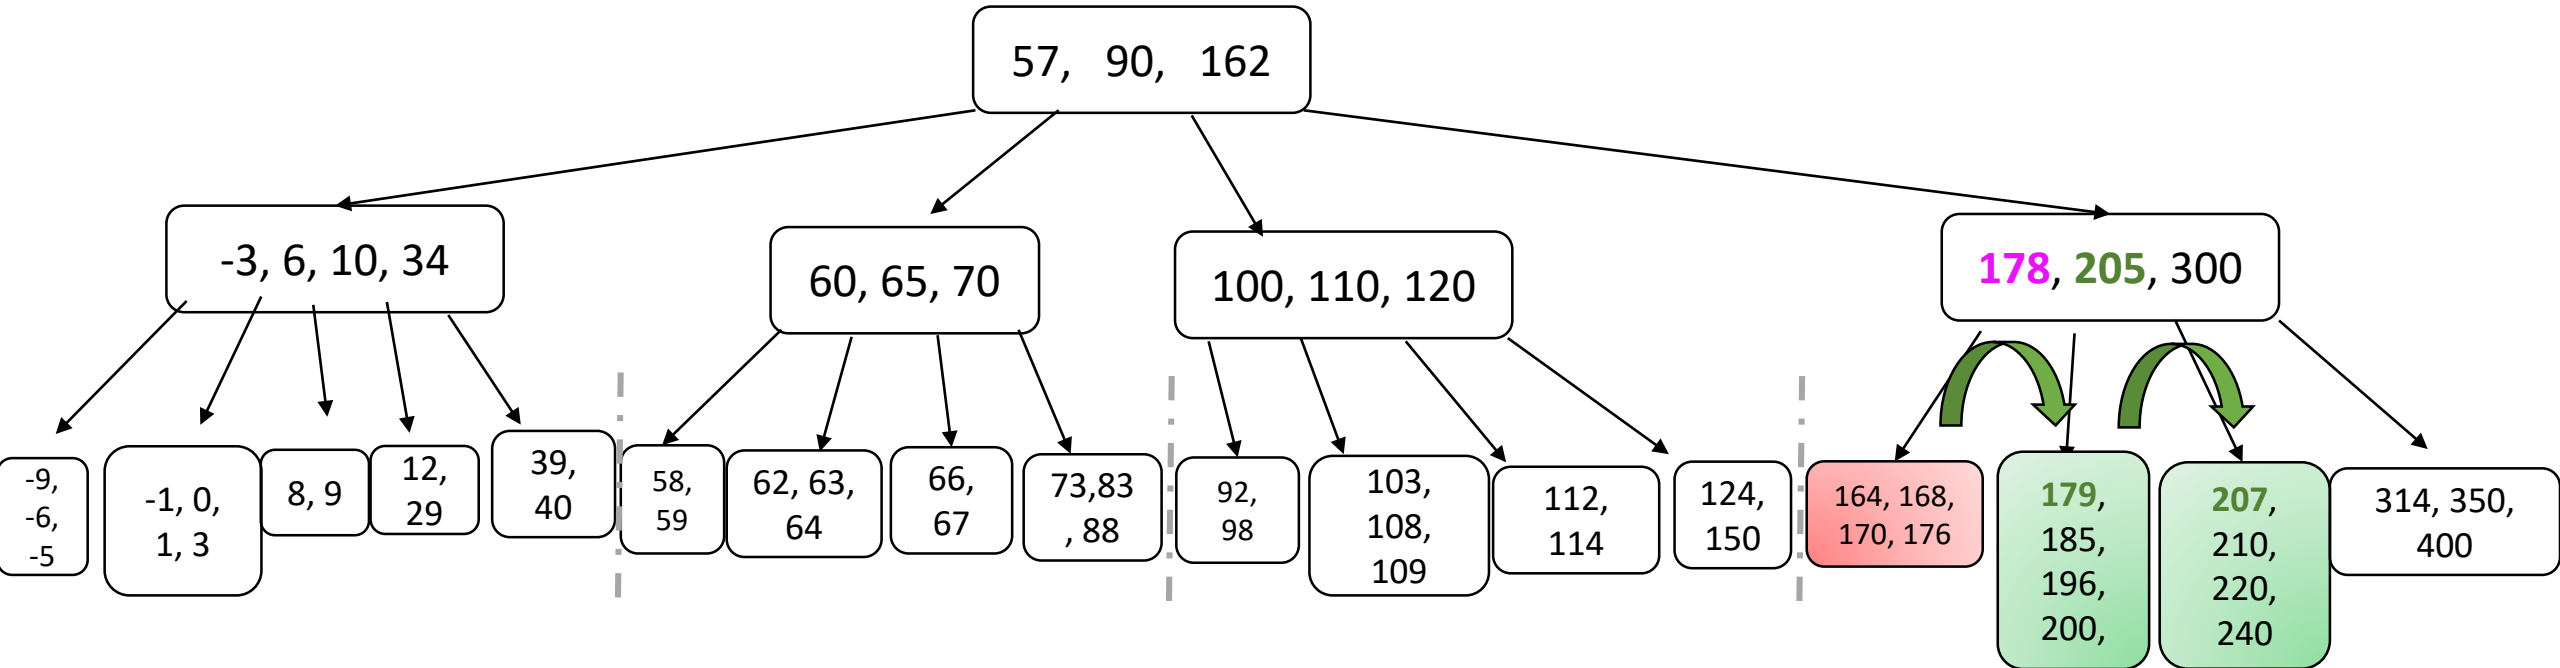
- Task: Insert 178

Example #3: Examine first sibling node!



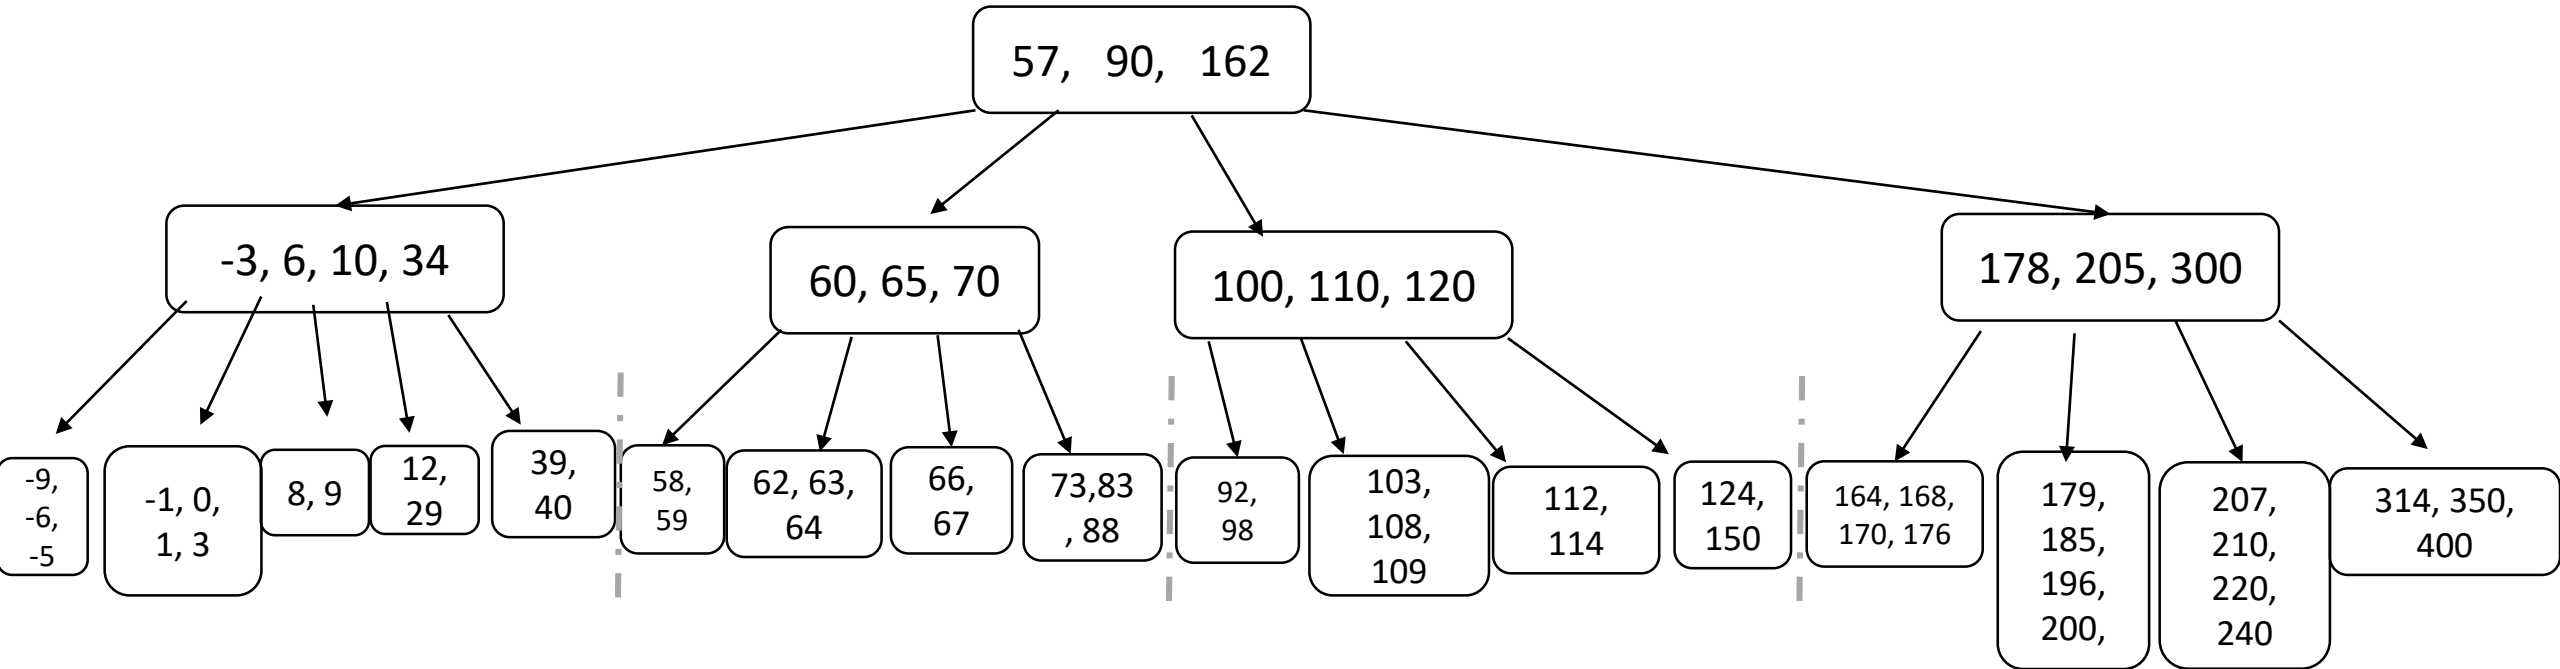
- Task: Insert 178
- Overflow... 😞
- Only sibling also full 😞

Example #3: Examine first sibling node!



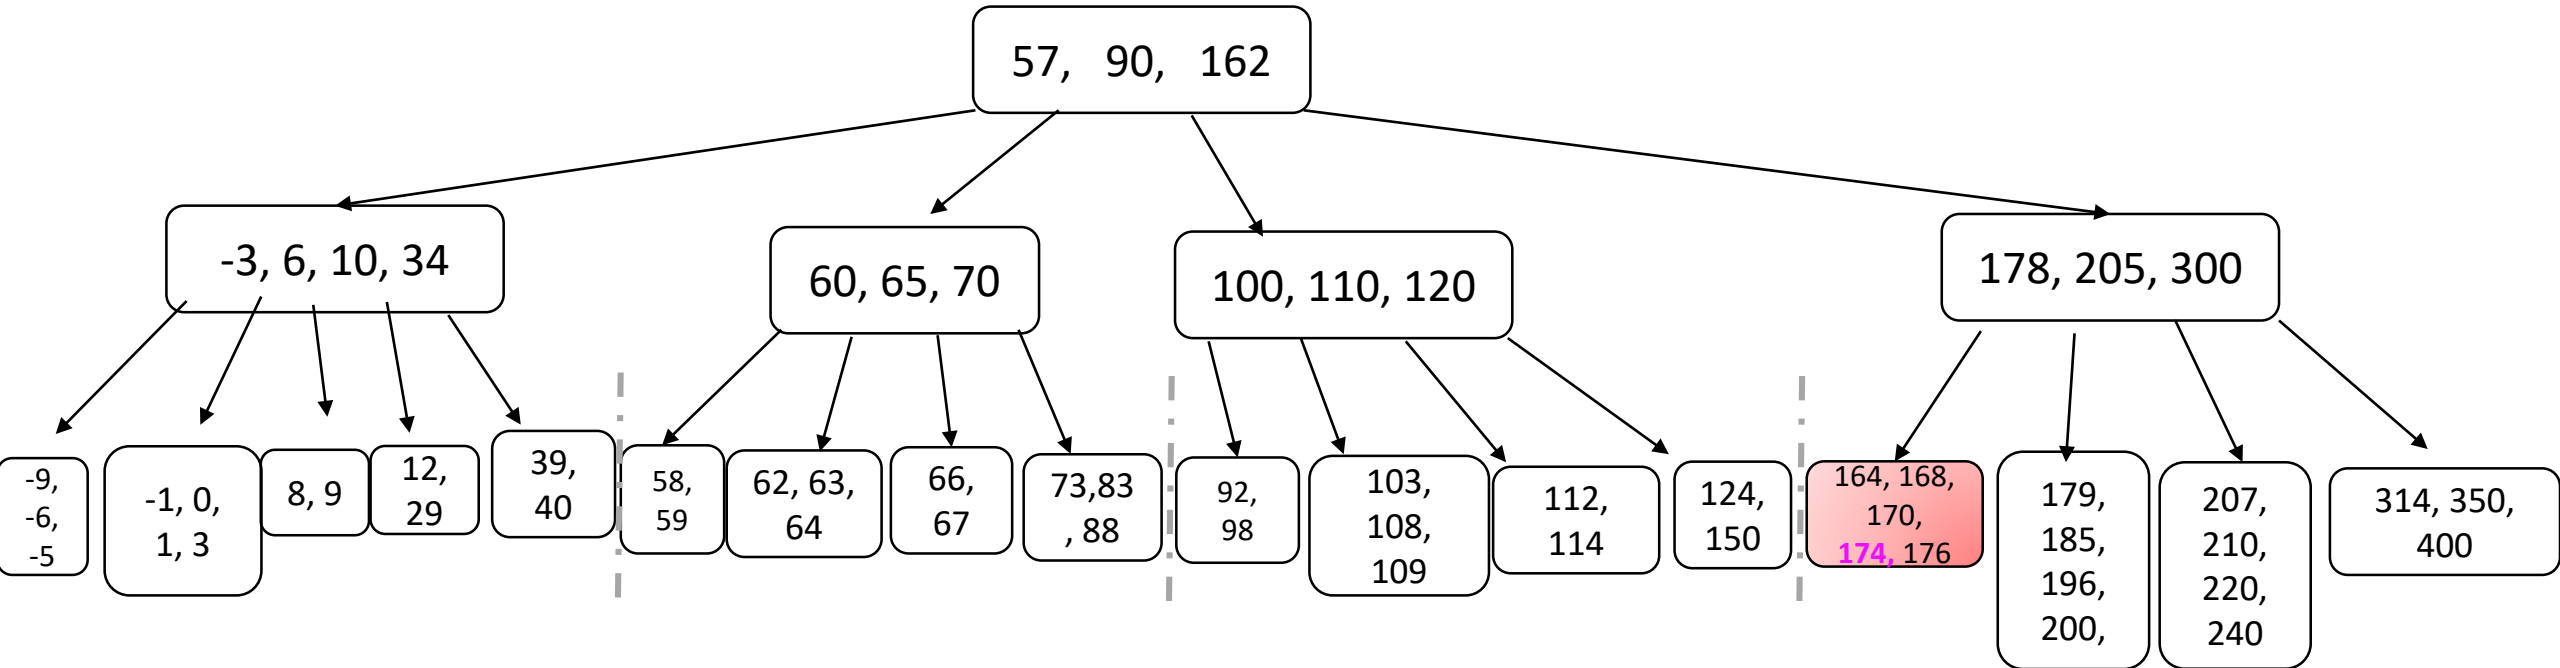
- Task: Insert 178
- Overflow... ☹️
- Only sibling also full ☹️
- Solution: Rotate 205 rightwards and 178 (new key!) rightwards! 😊

Example #3: Examine first sibling node!



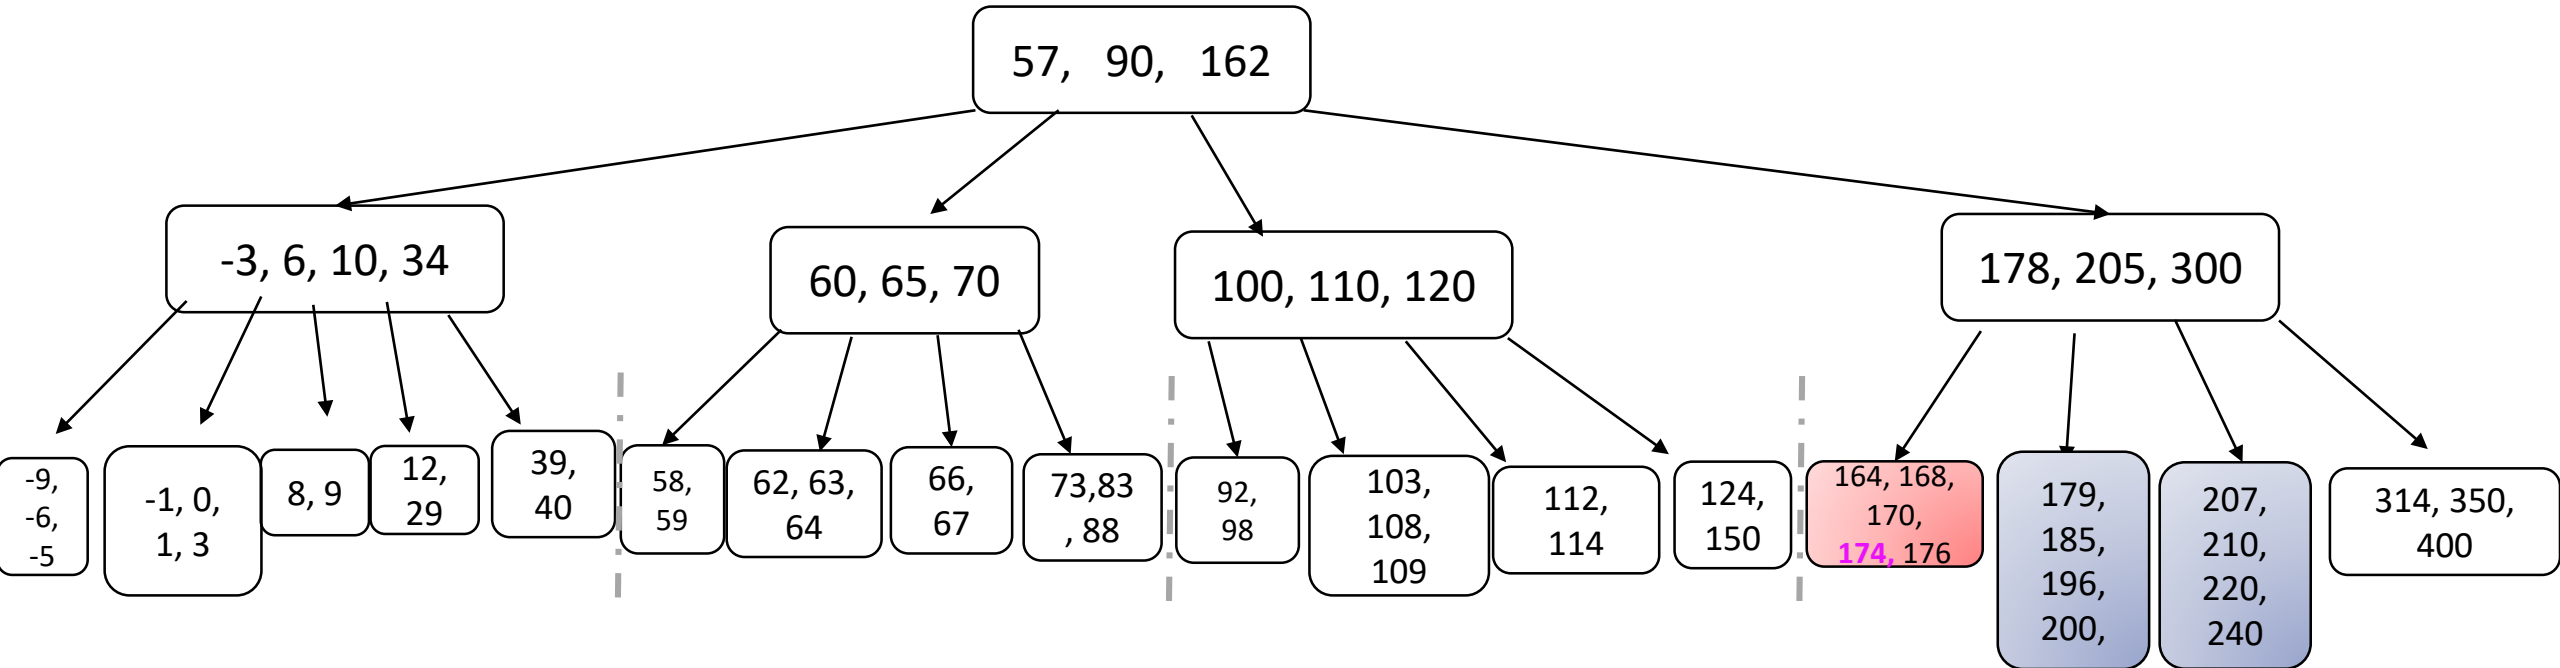
- Task: Insert 174

Example #3: Examine first sibling node!



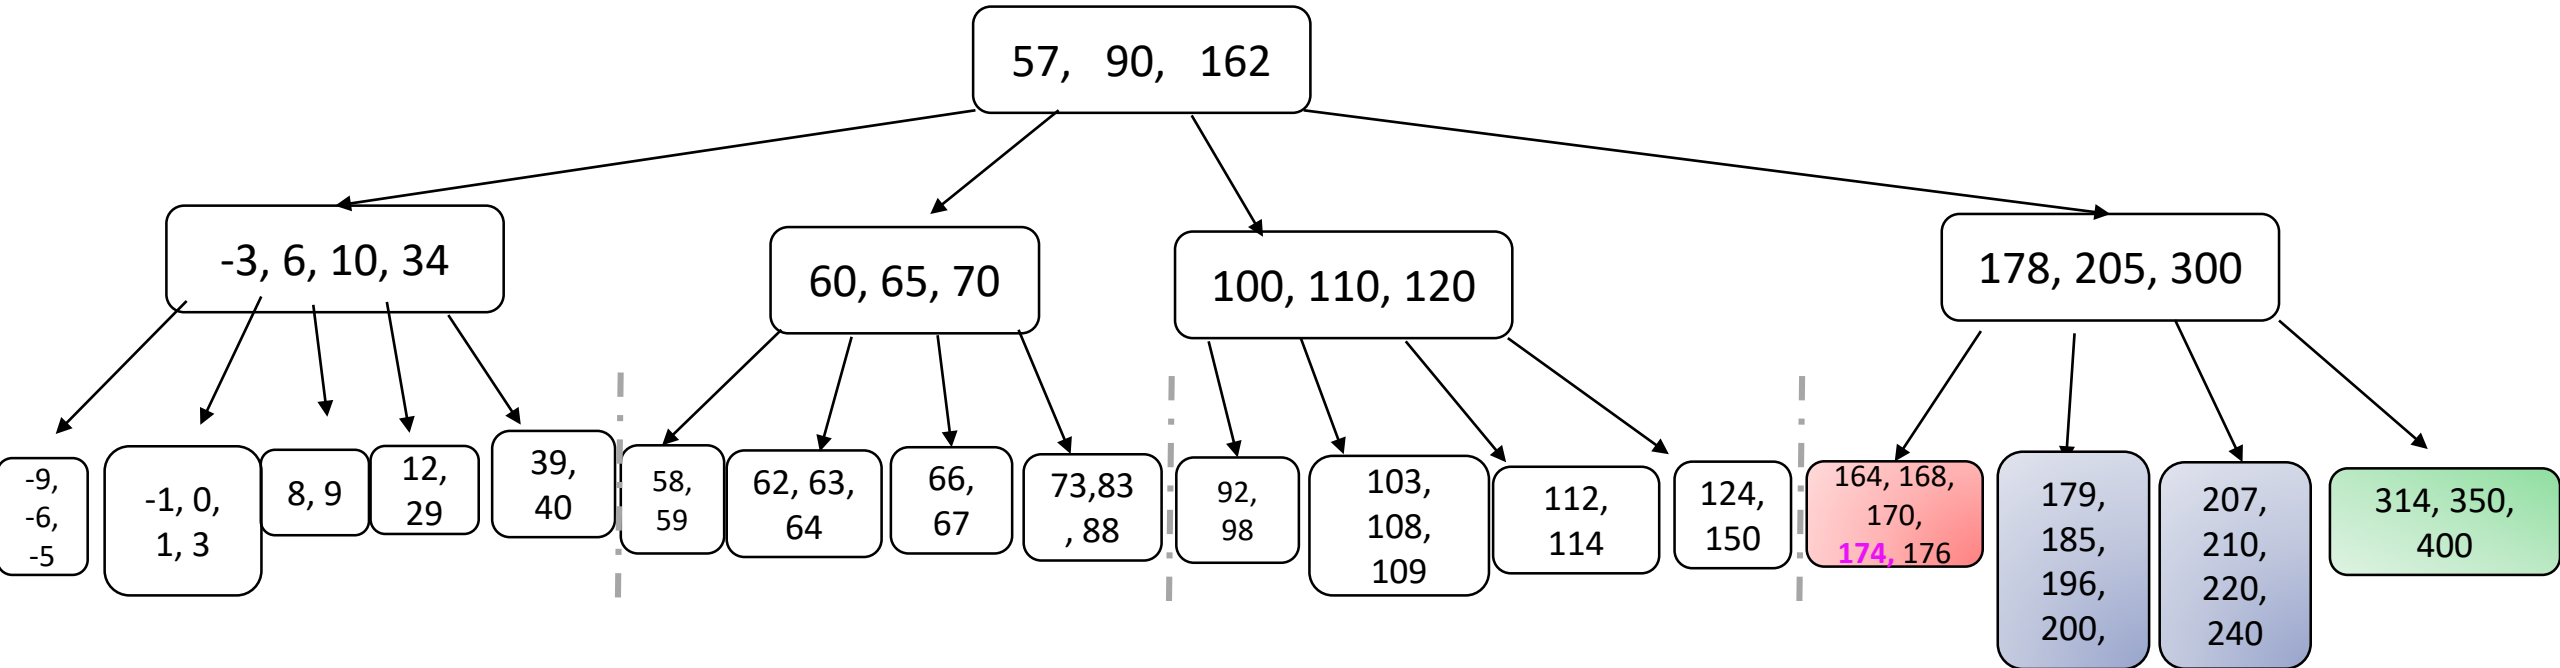
- Task: Insert 174
- Overflow 😞

Example #3: Examine first sibling node!



- Task: Insert 174
- Overflow 😞
- 1st and 2nd sibling don't have space to rotate keys to 😞

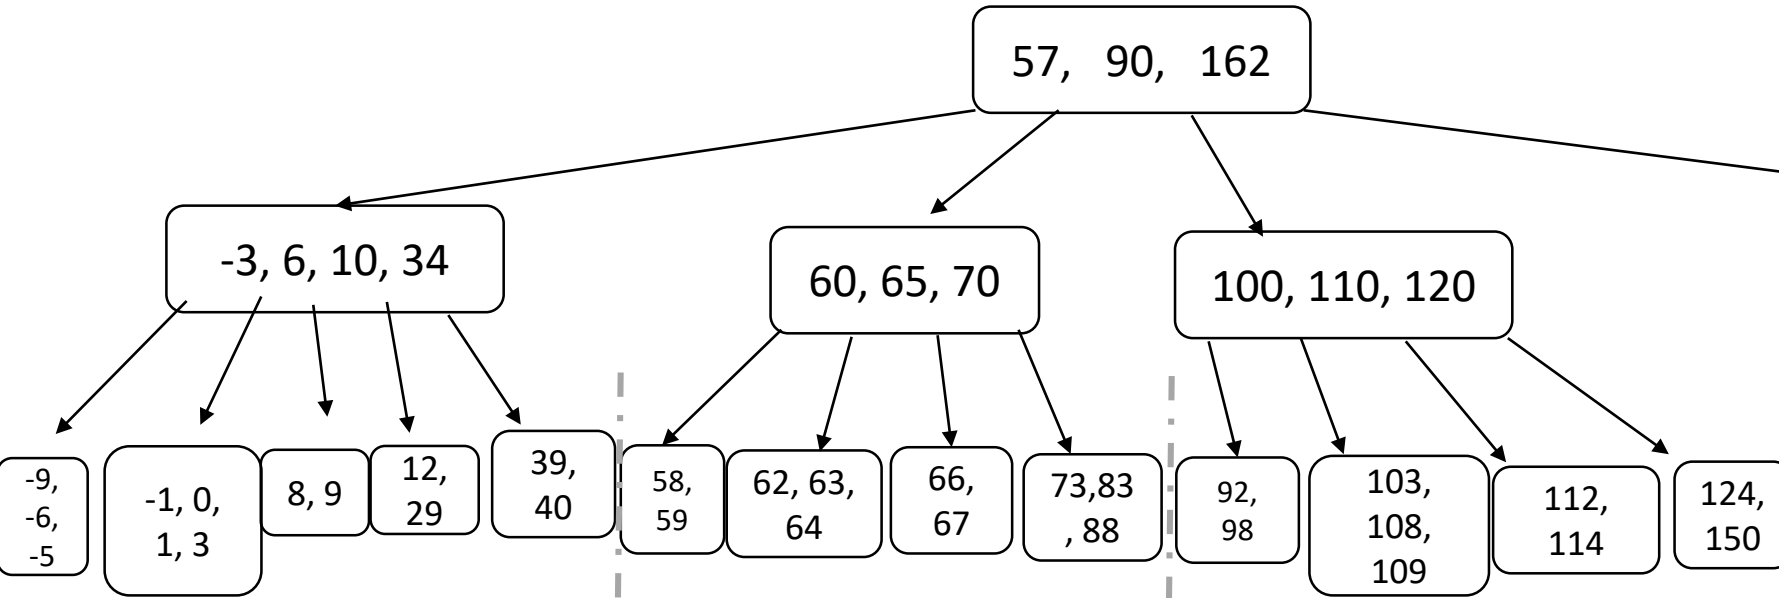
Example #4: Examine second sibling node!



- Task: Insert 174
- Overflow 😞
- 1st and 2nd sibling don't have space to rotate keys to 😞
- But 3rd sibling does! 😊



Example #4: Examine second sibling node!



- Task: Insert 174
- Overflow 😞
- 1st and 2nd sibling don't have space to rotate keys to 😞
- But 3rd sibling does! 😊



WE INTERRUPT YOUR
REGULARLY
SCHEDULED PROGRAM
FOR AN
INVESTIGATION INTO
SIBLING SEARCH!

So, should I always rotate as far away as I can?

So, should I always rotate as far away as I can?

- **UNCLEAR.**

So, should I always rotate as far away as I can?

- **UNCLEAR.**
- In the 70s and 80s, **yes, absolutely**
 - **Memory expensive:** inner nodes in disk as well means we want to minimize height.

So, should I always rotate as far away as I can?

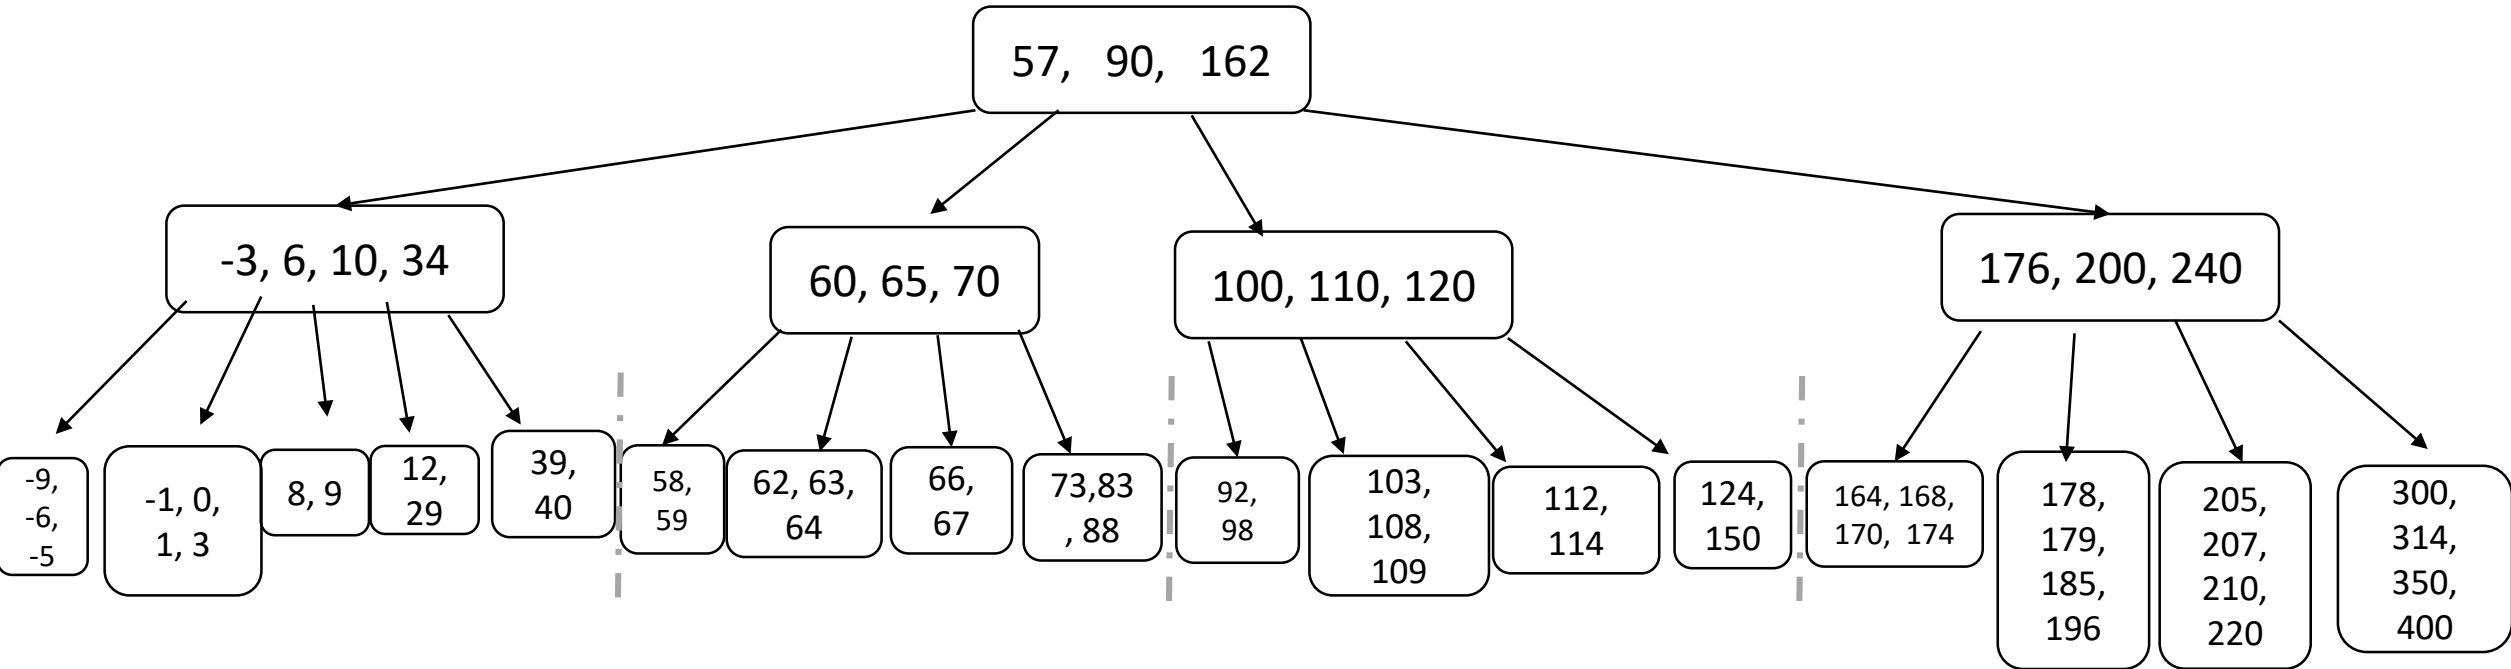
- **UNCLEAR.**

- In the 70s and 80s, **yes, absolutely**
 - **Memory expensive:** inner nodes in disk as well means we want to minimize height.
- Nowadays, can fit the index (B-Tree) in memory.
 - So might not care about **adding to the height that much**

So, should I always rotate as far away as I can?

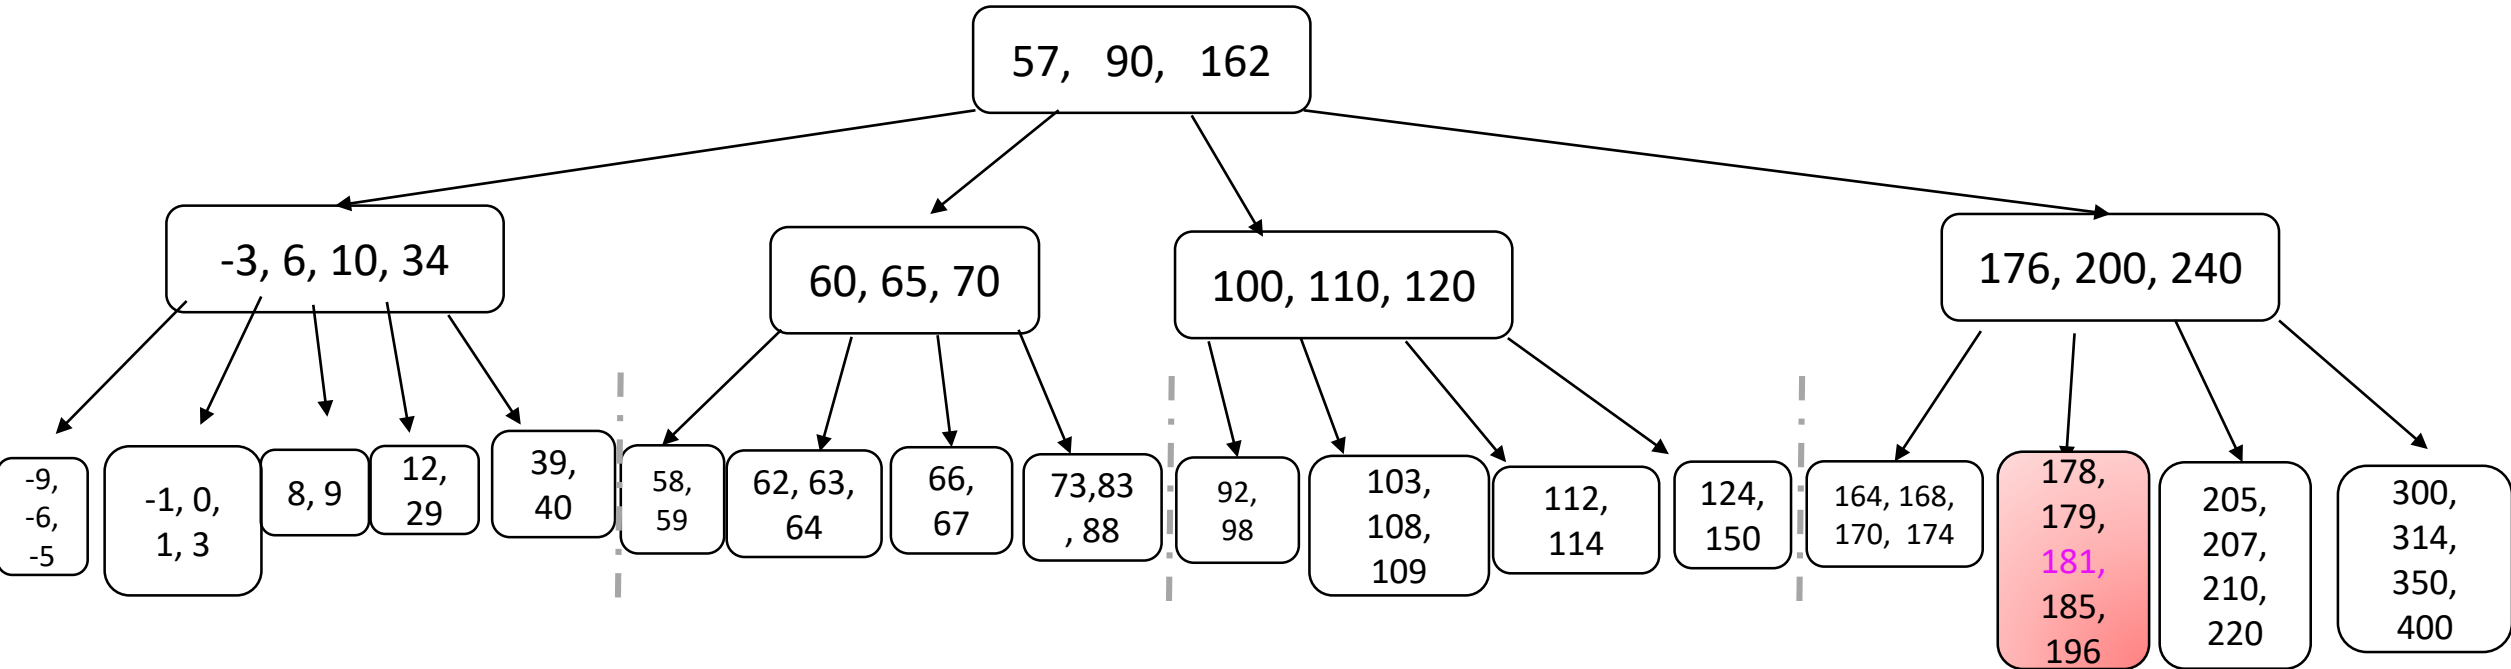
- **UNCLEAR.**
- In the 70s and 80s, **yes, absolutely**
 - **Memory expensive:** inner nodes in disk as well means we want to minimize height.
- Nowadays, can fit the index (B-Tree) in memory.
 - So might not care about **adding to the height that much**
- **Our solution (for exams, projects, etc):** supply an `int` parameter *hops* to either **insertion method** or **B-Tree constructor** to indicate **how far away** you are willing to search for key space when you overflow.
 - *hops* = `-1` can be used to indicate “as far away as you want”.

Example #5: Splitting, parent has space



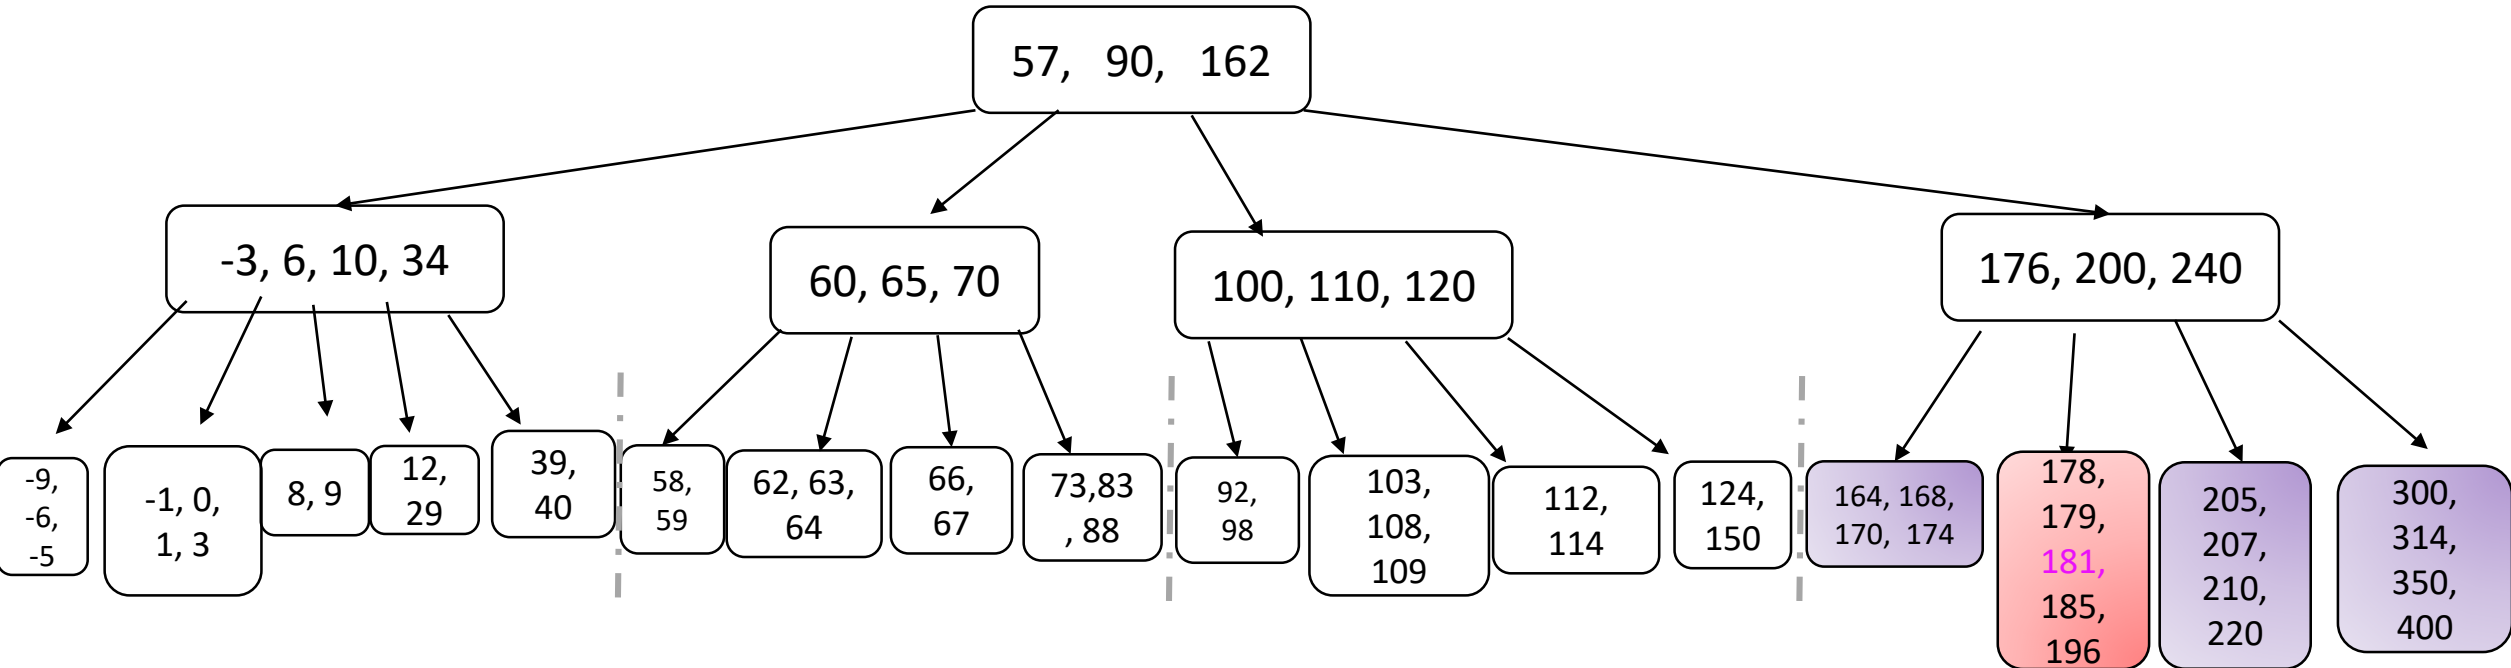
- Task: Insert 181

Example #5: Splitting, parent has space



- Task: Insert 181
- Overflow 😞

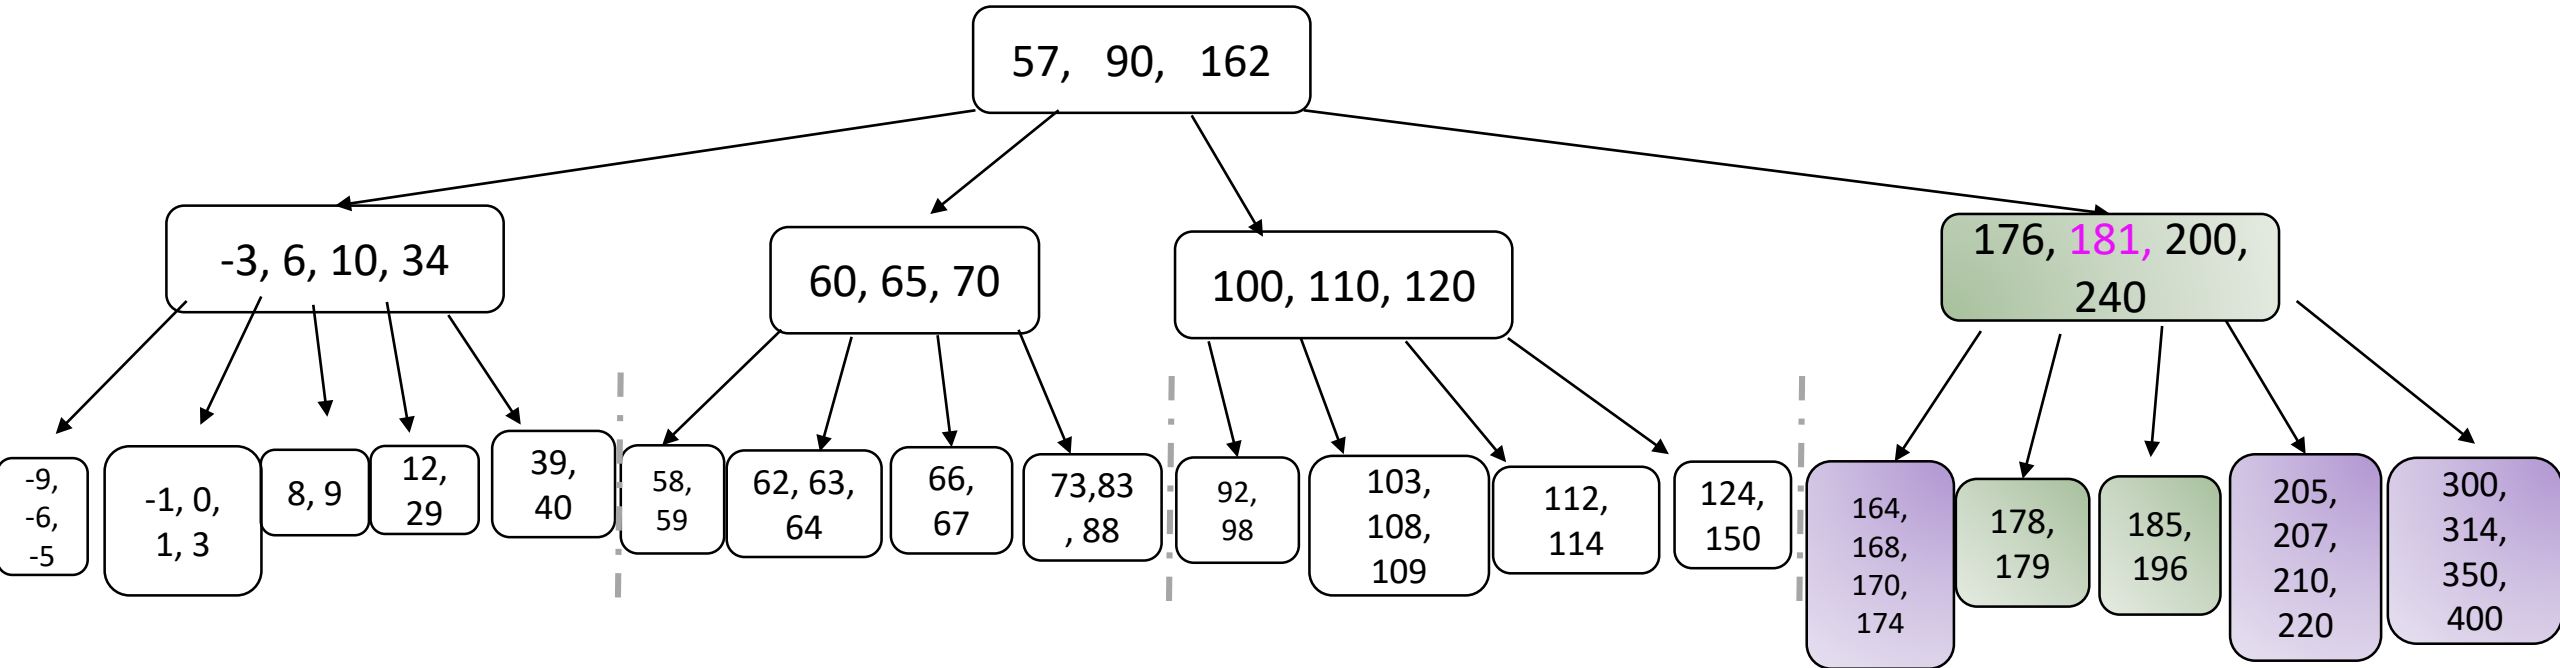
Example #5: Splitting, parent has space




- Task: Insert **181**
- **Overflow** 😞
- No sibling or cousin has space for rotating keys to! 😞



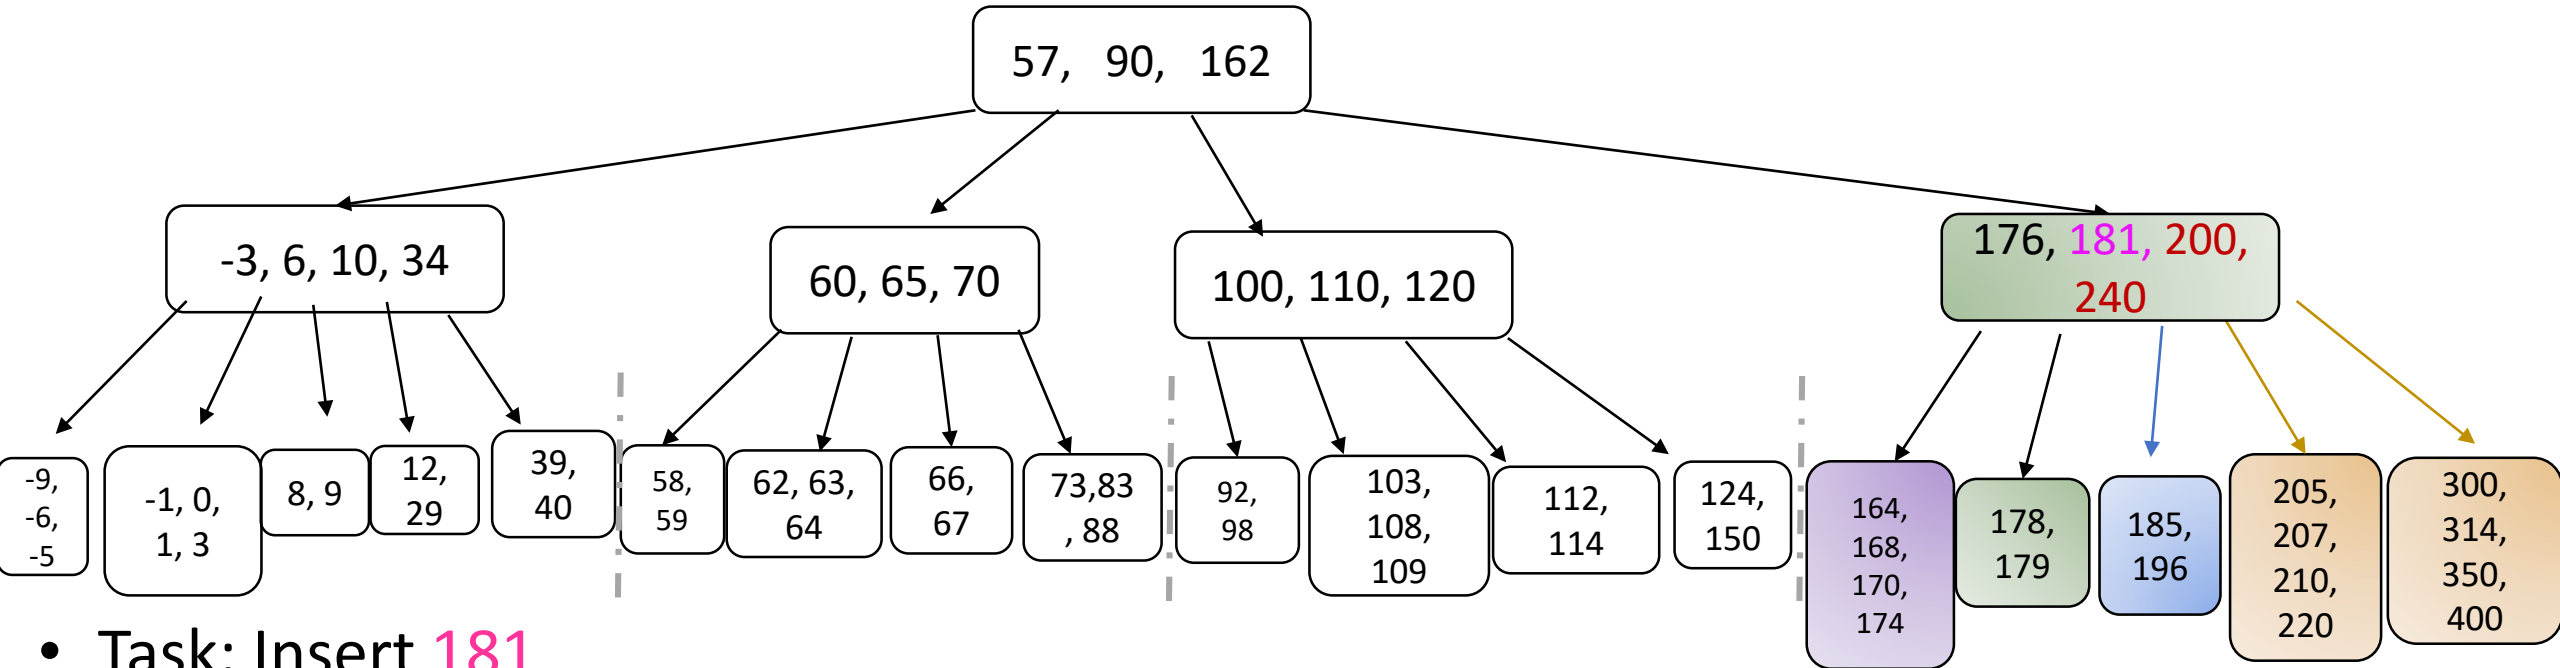
Example #5: Splitting, parent has space



- Task: Insert **181**
- **Overflow** 😞
- No sibling or cousin has space for rotating keys to! 😞
- Solution: Split the current node, and elevate middle key to the parent 😊

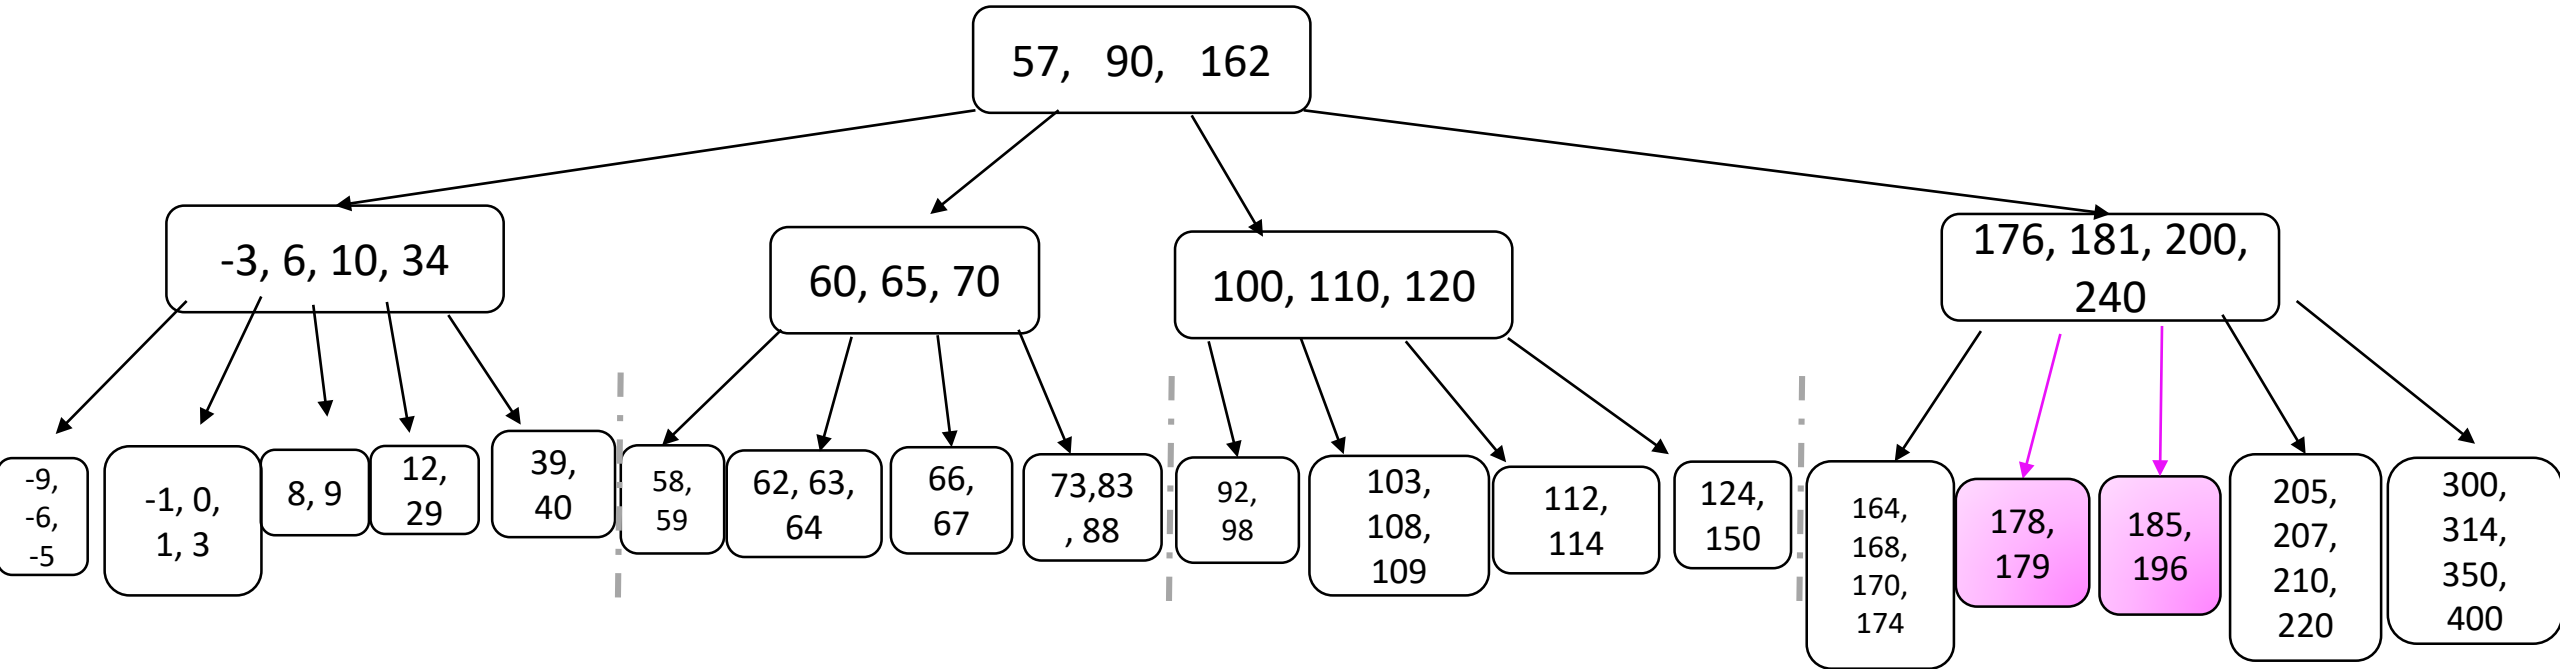


Example #5: Splitting, parent has space



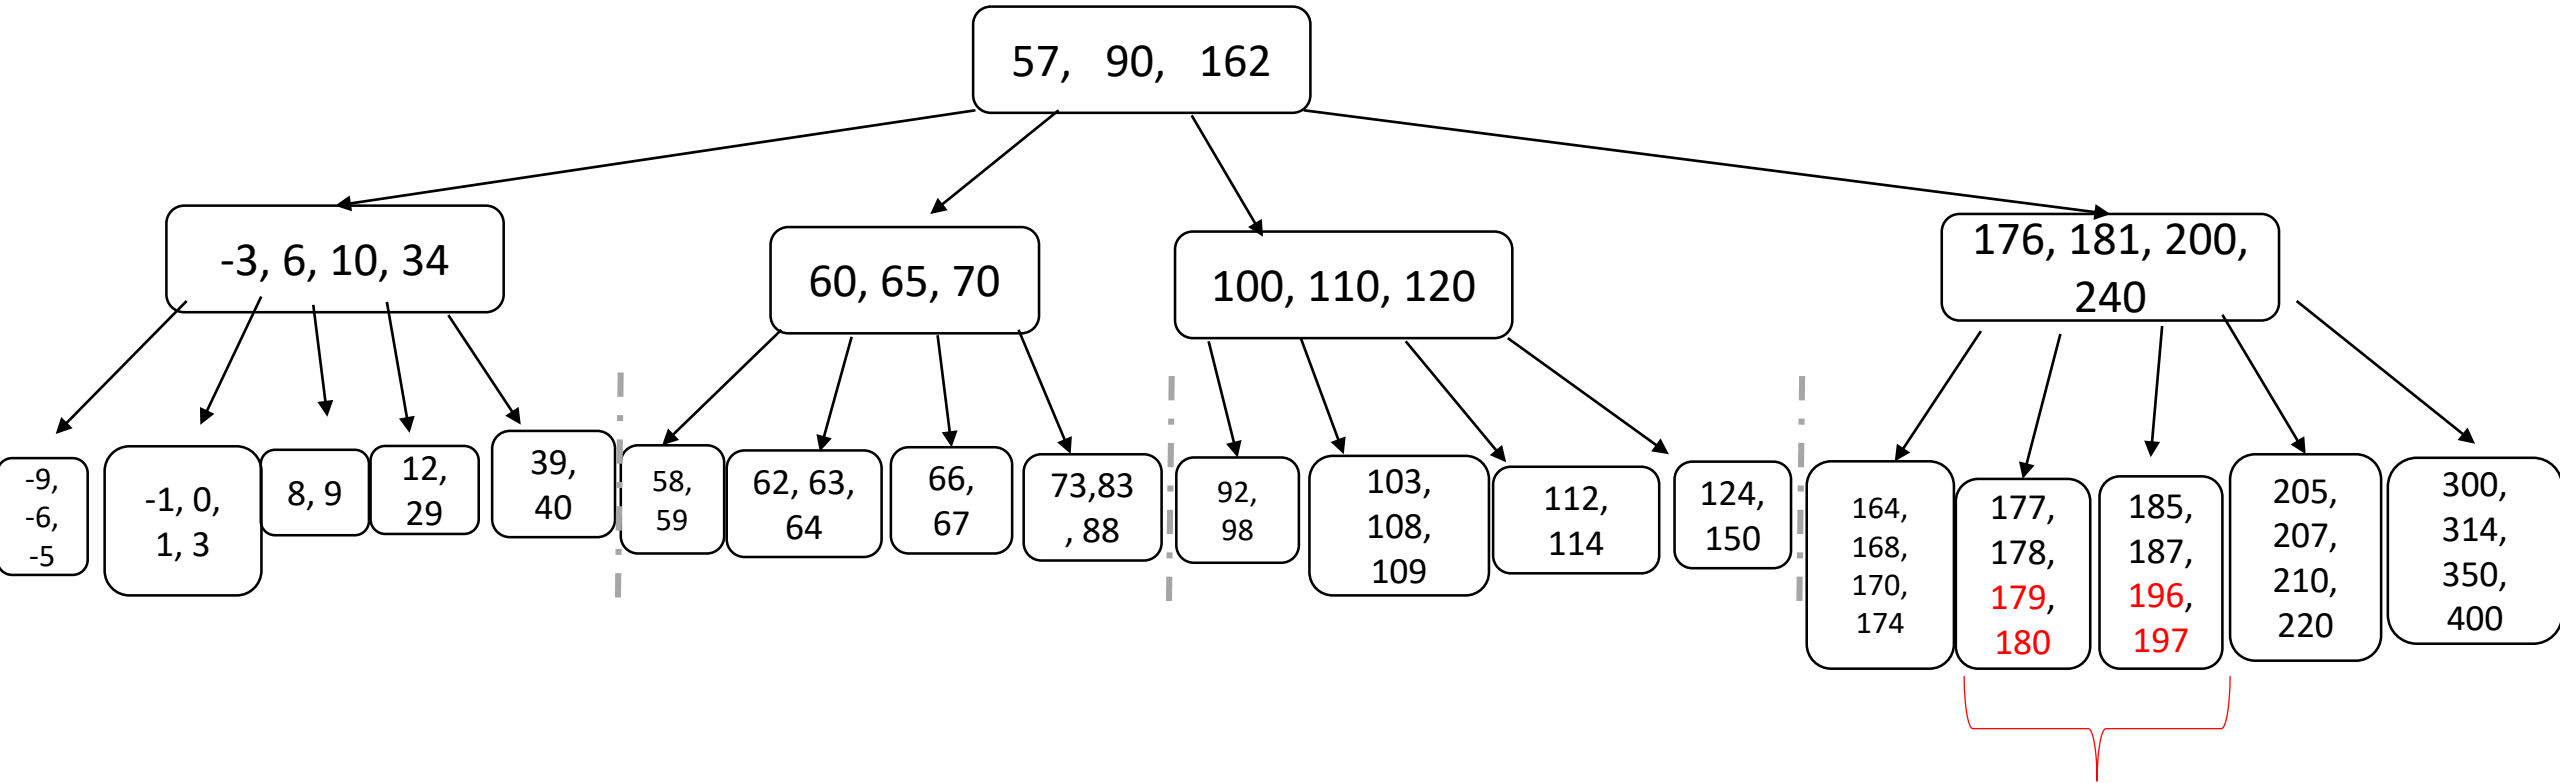
- Task: Insert 181
- Overflow 😞
- No sibling or cousin has space for rotating keys to! 😞
- Solution: Split the current node, and elevate middle key to the parent 😊
 - Parent has 1 additional child subtree now.

Interesting notes about splitting



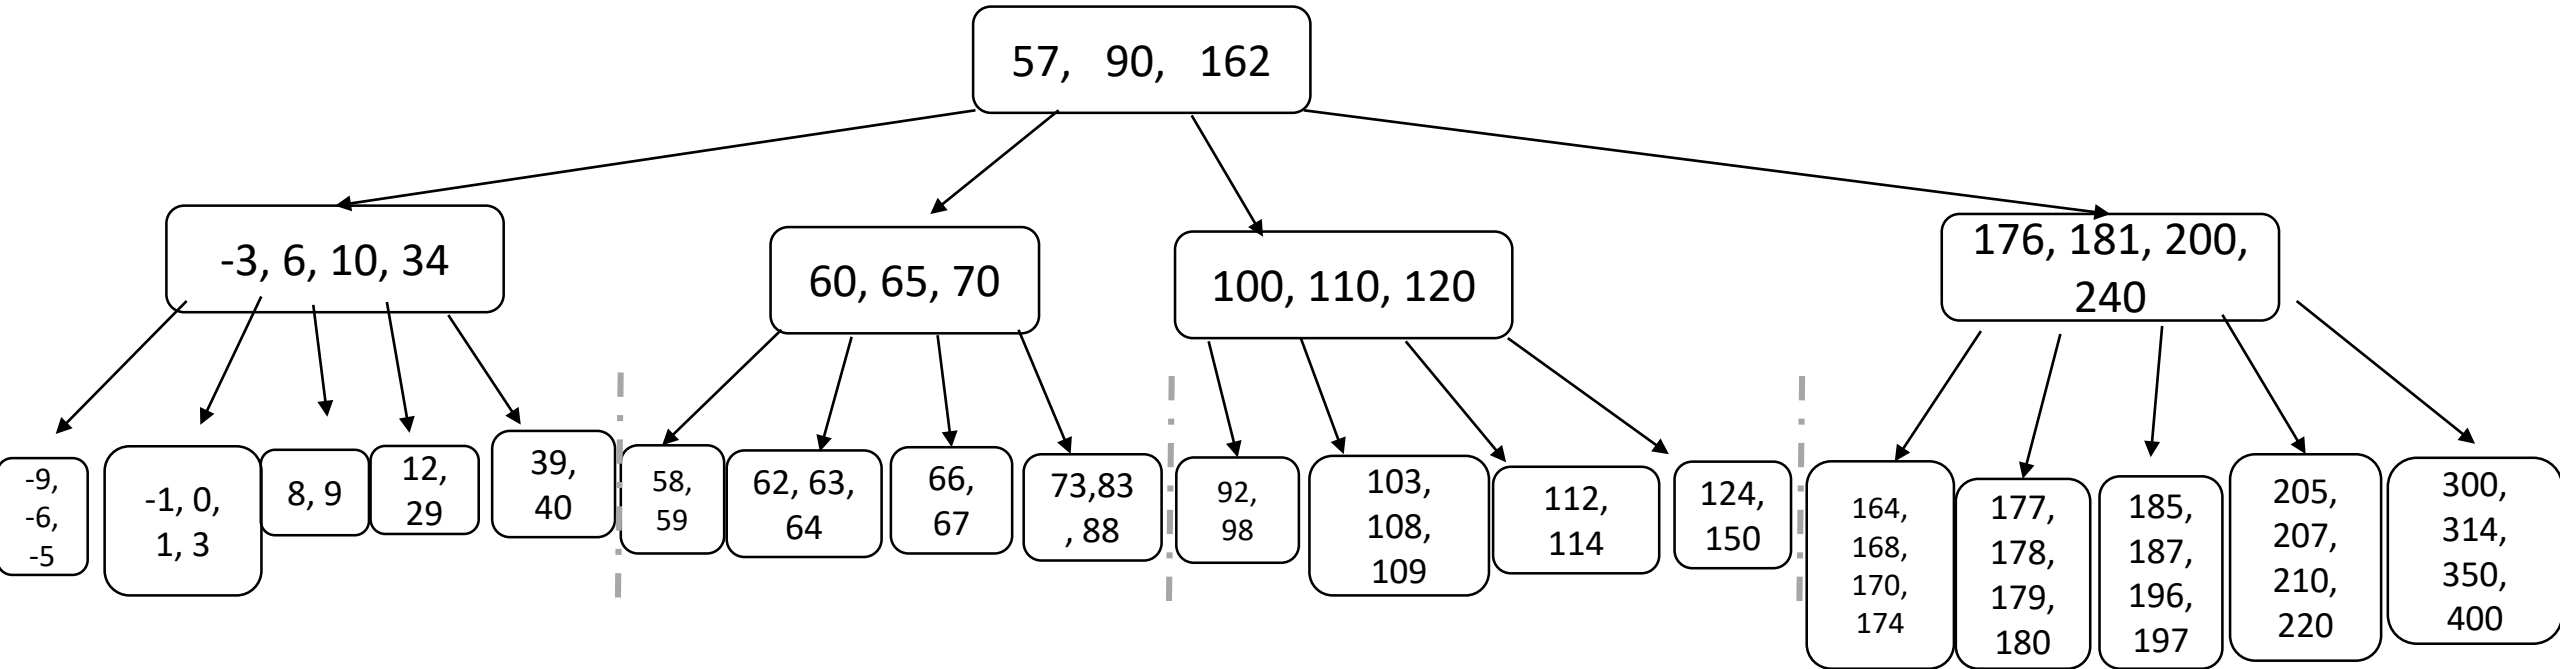
- 4 key insertions in the range [177, 199] will be solvable with **key rotations to the left or right!**
- Point: when a splitting in a certain range occurs, it's unlikely to happen again in the near future (*locality of reference*)

Example #6: Splitting, parent overflows and rotates



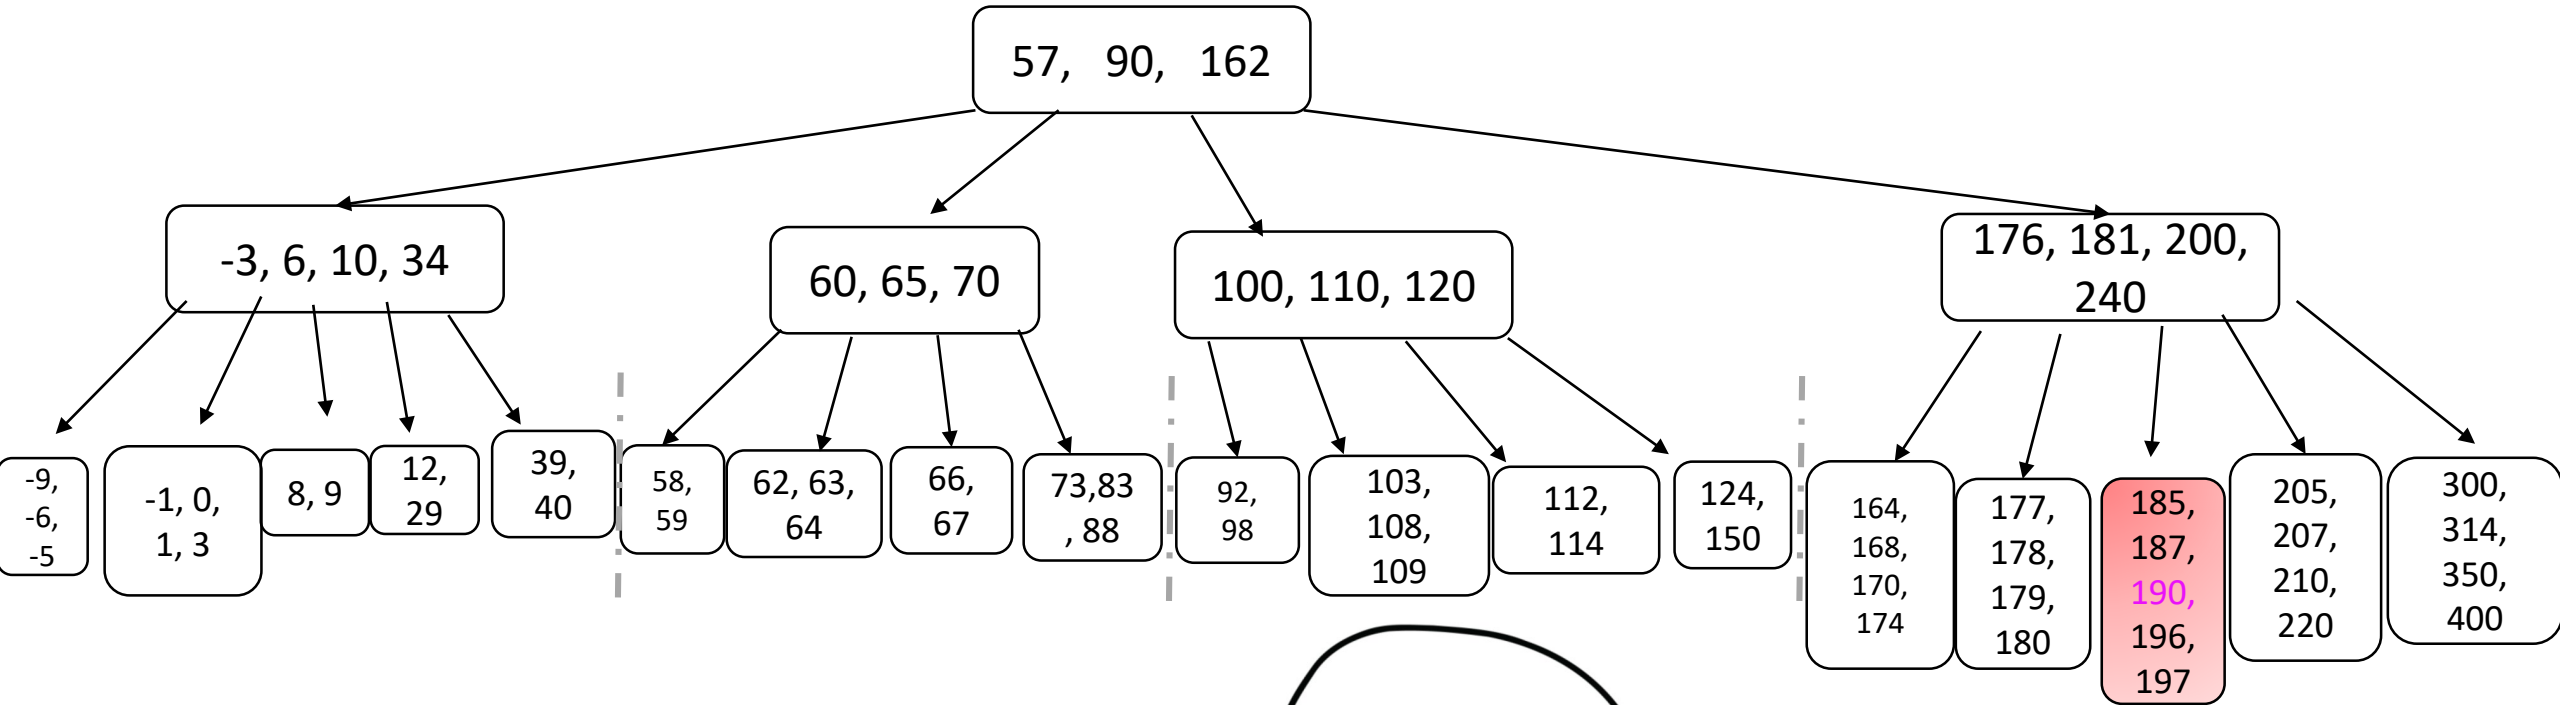
Added some keys here to make example interesting...

Example #6: Splitting, parent overflows and rotates



- Task: Insert 190

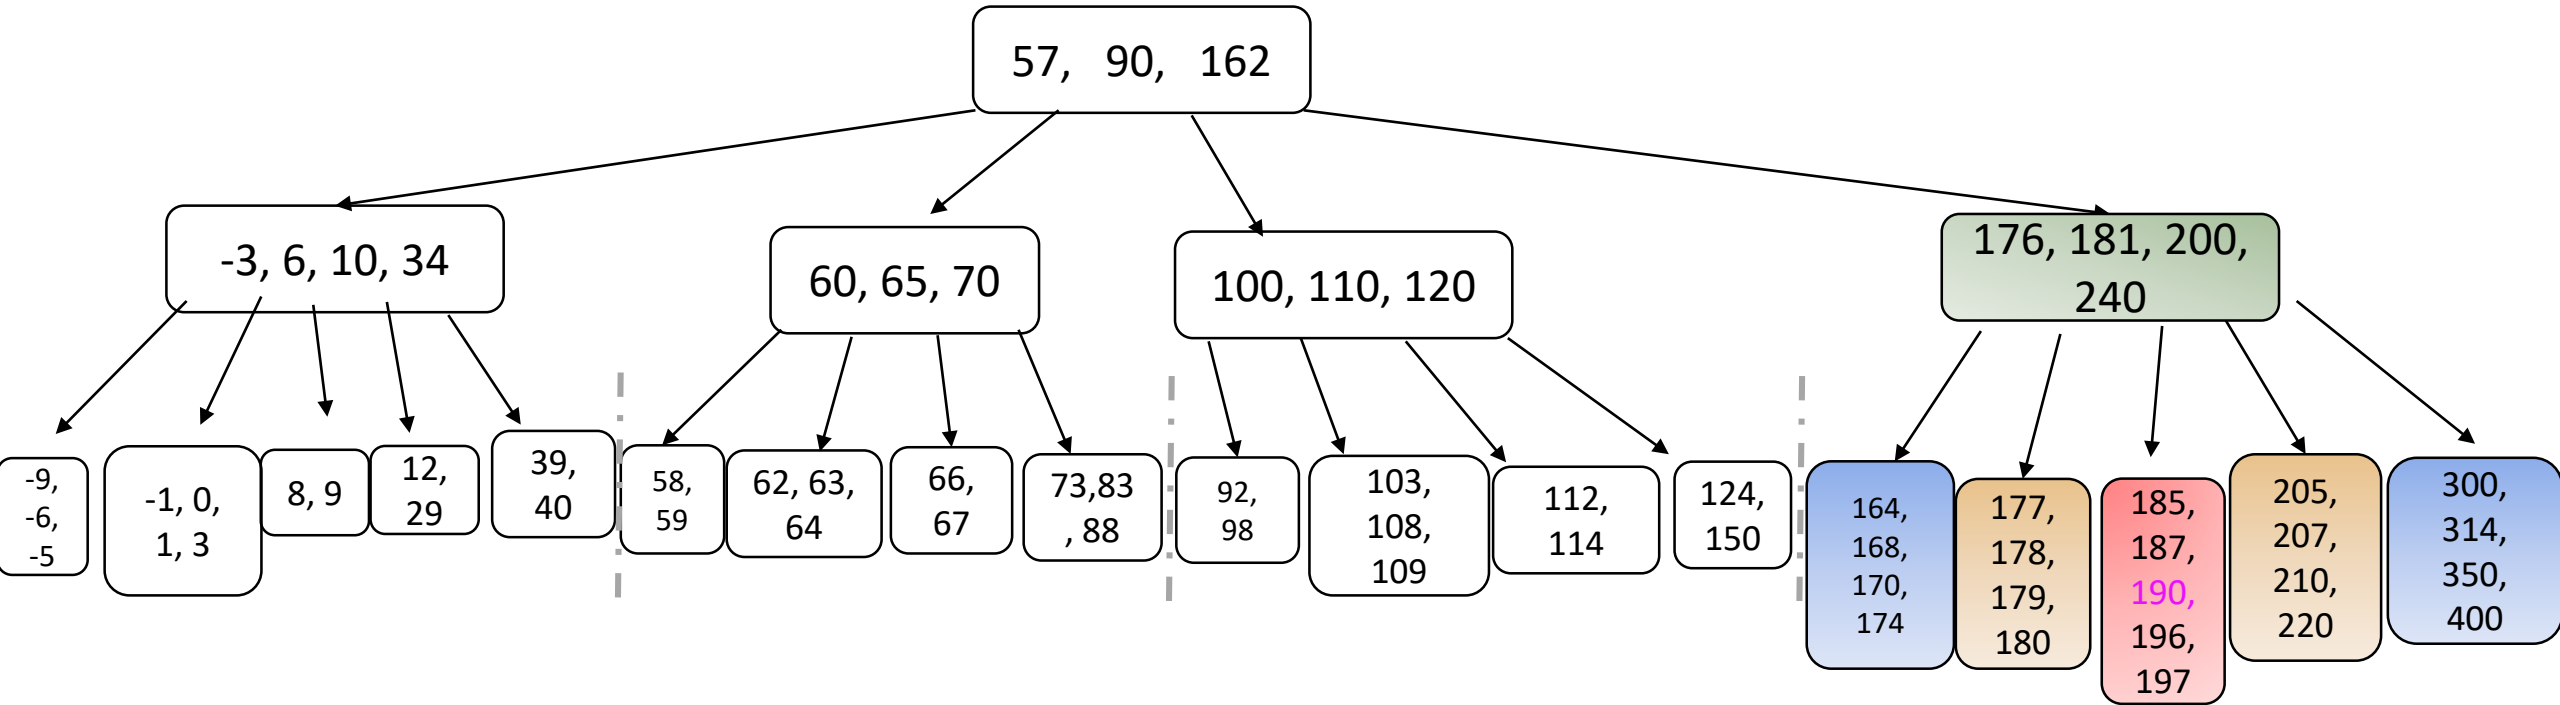
Example #6: Splitting, parent overflows and rotates



- Task: Insert 190
- Overflow 😞



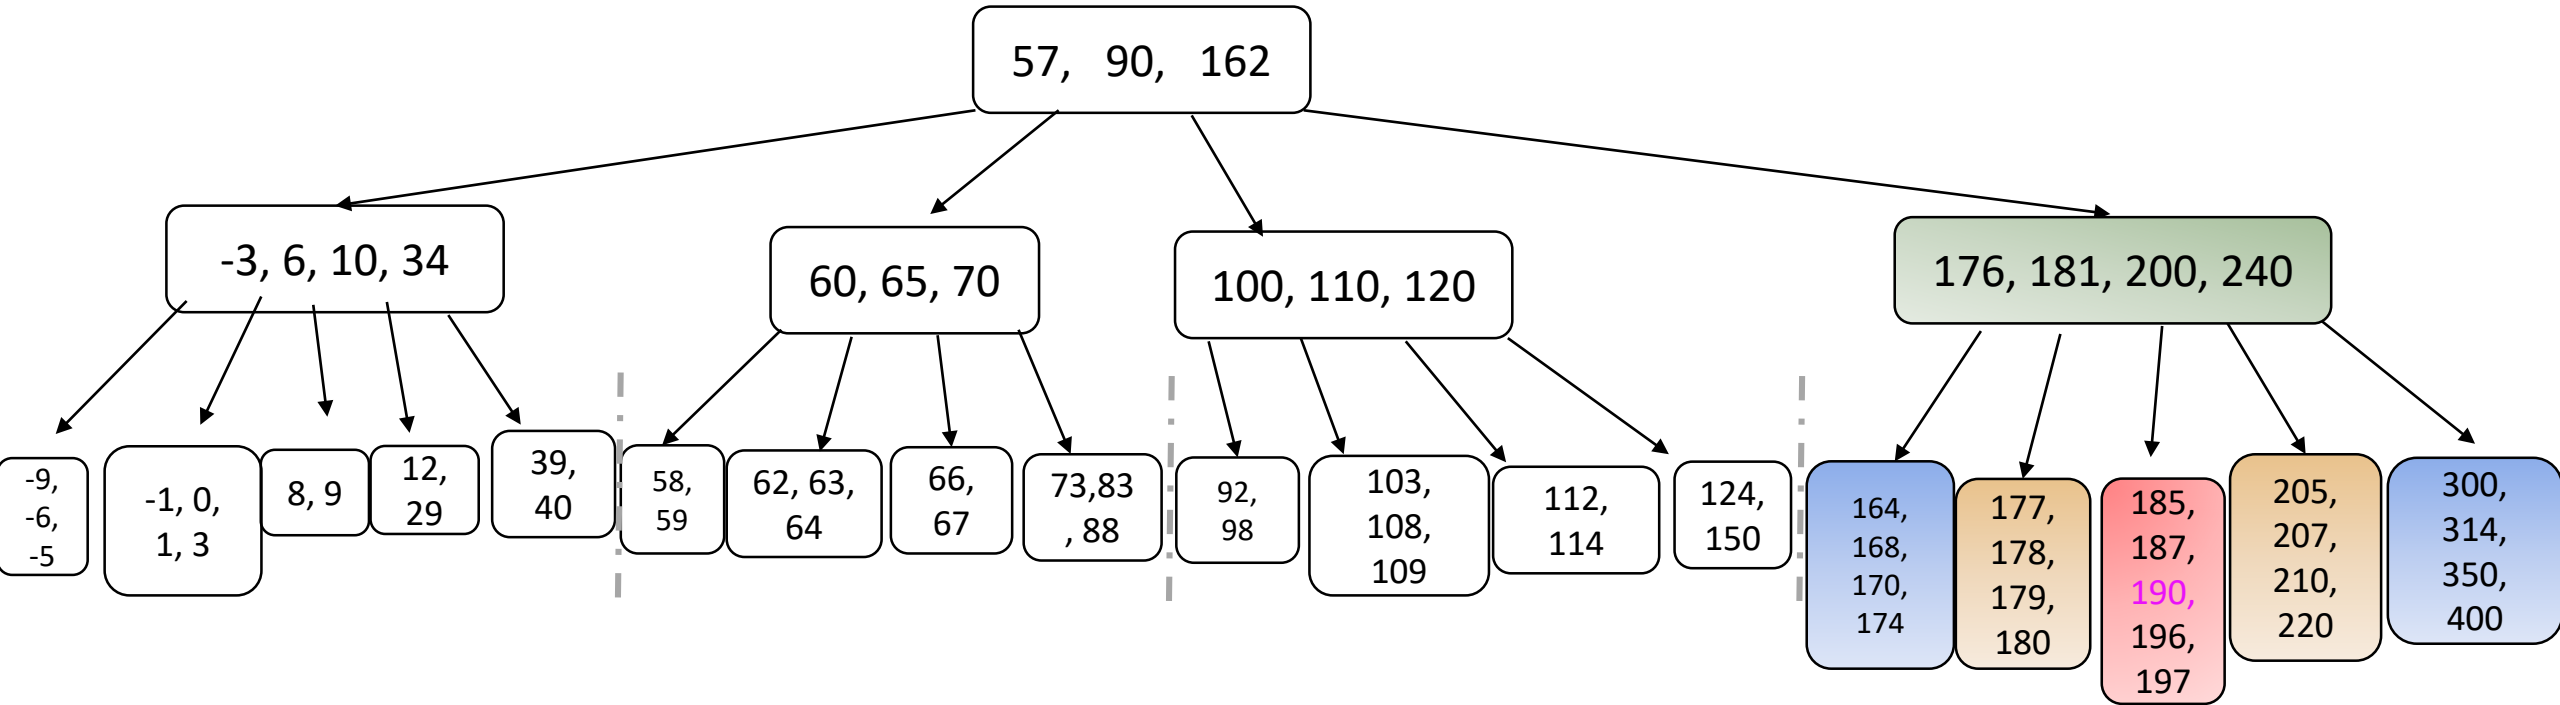
Example #6: Splitting, parent overflows and rotates



- Task: Insert 190
- Overflow ☹️
- 1st Siblings, 2nd siblings and parent all full! ☹️

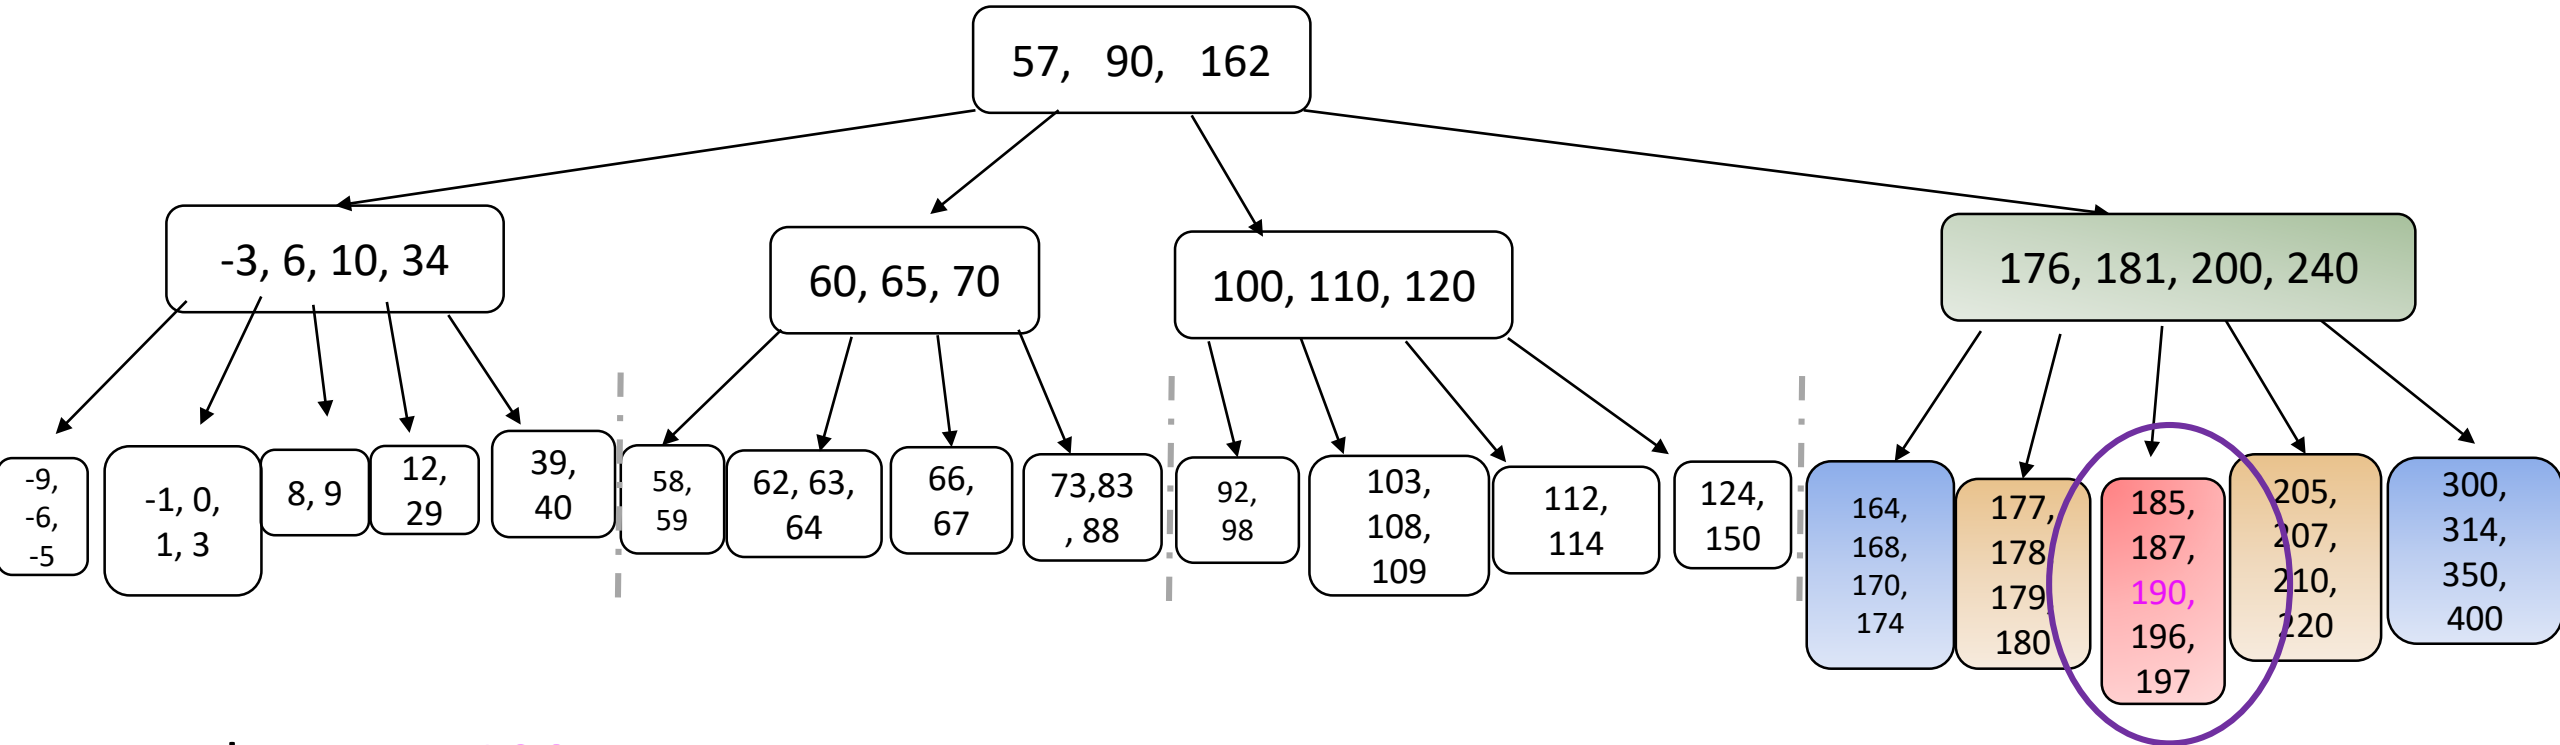


Example #6: Splitting, parent overflows and rotates



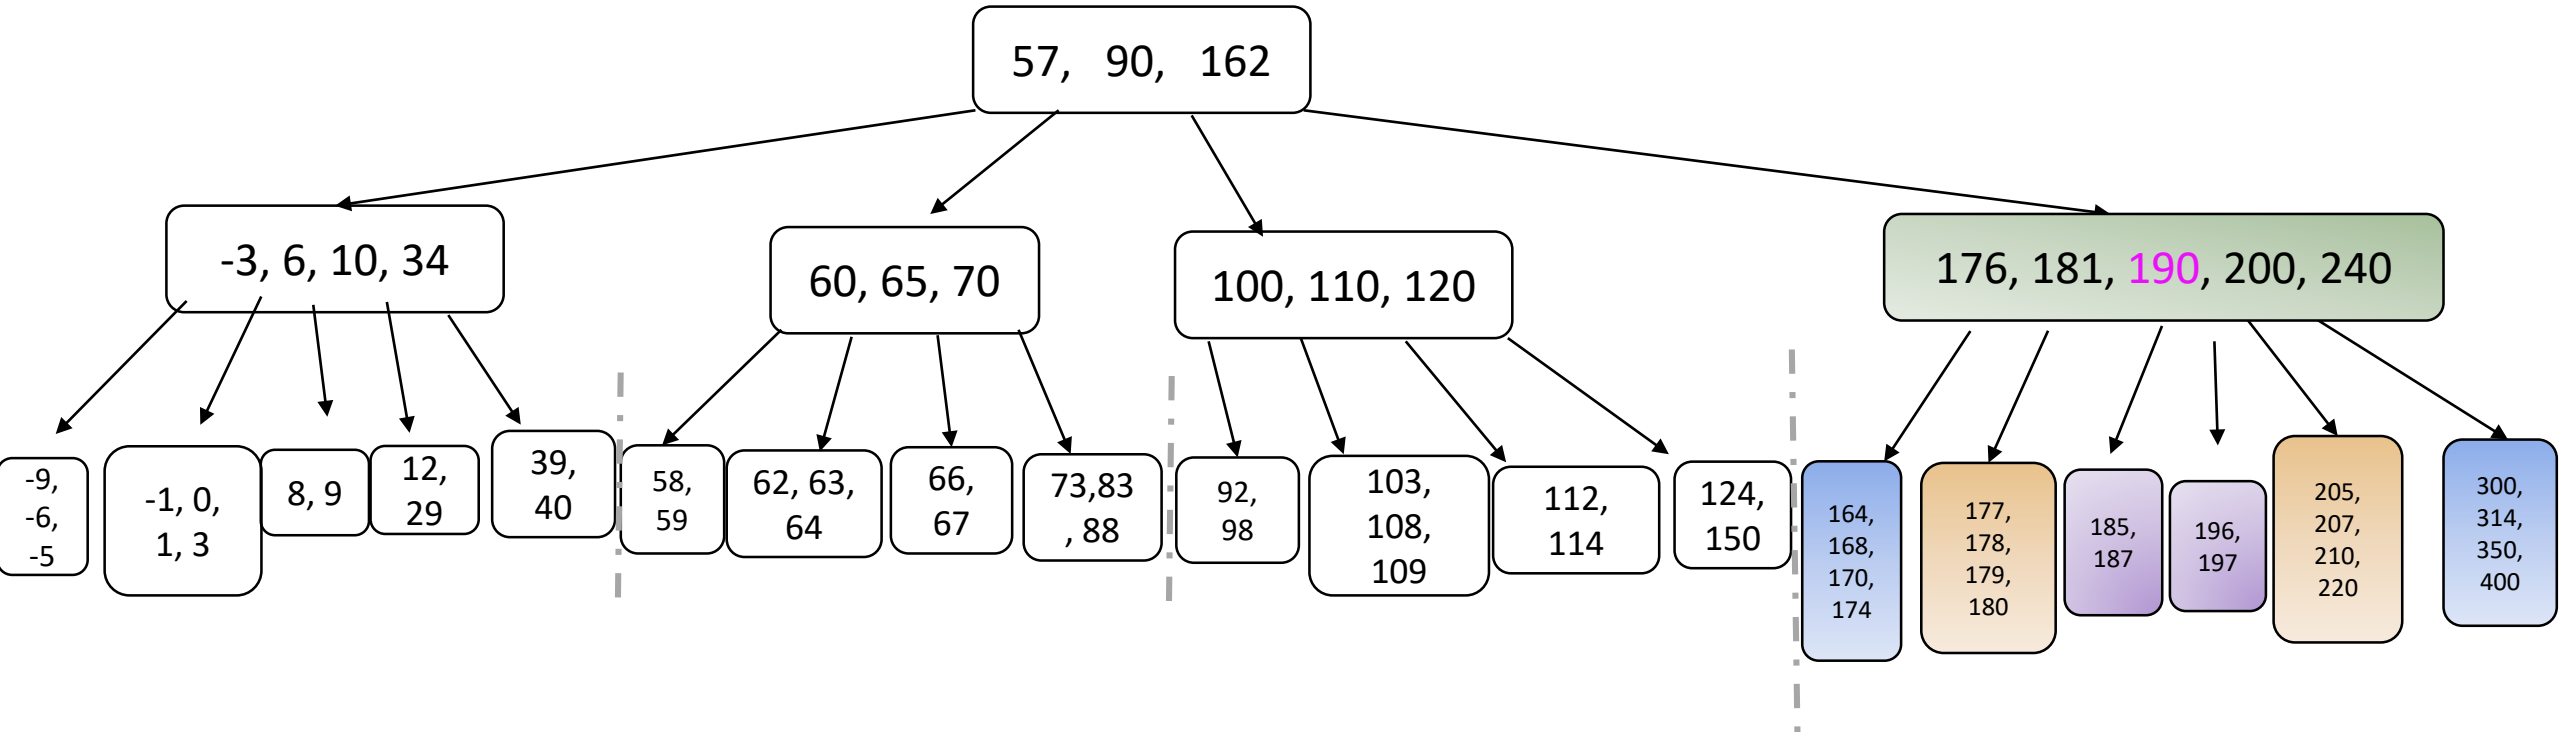
- Task: Insert 190
- Overflow 😞
- 1st Siblings, 2nd siblings and parent all full! 😞
- Solution: Split **node**, elevating 190 to **parent** and then rotate **176** to the left!

Example #6: Splitting, parent overflows and rotates



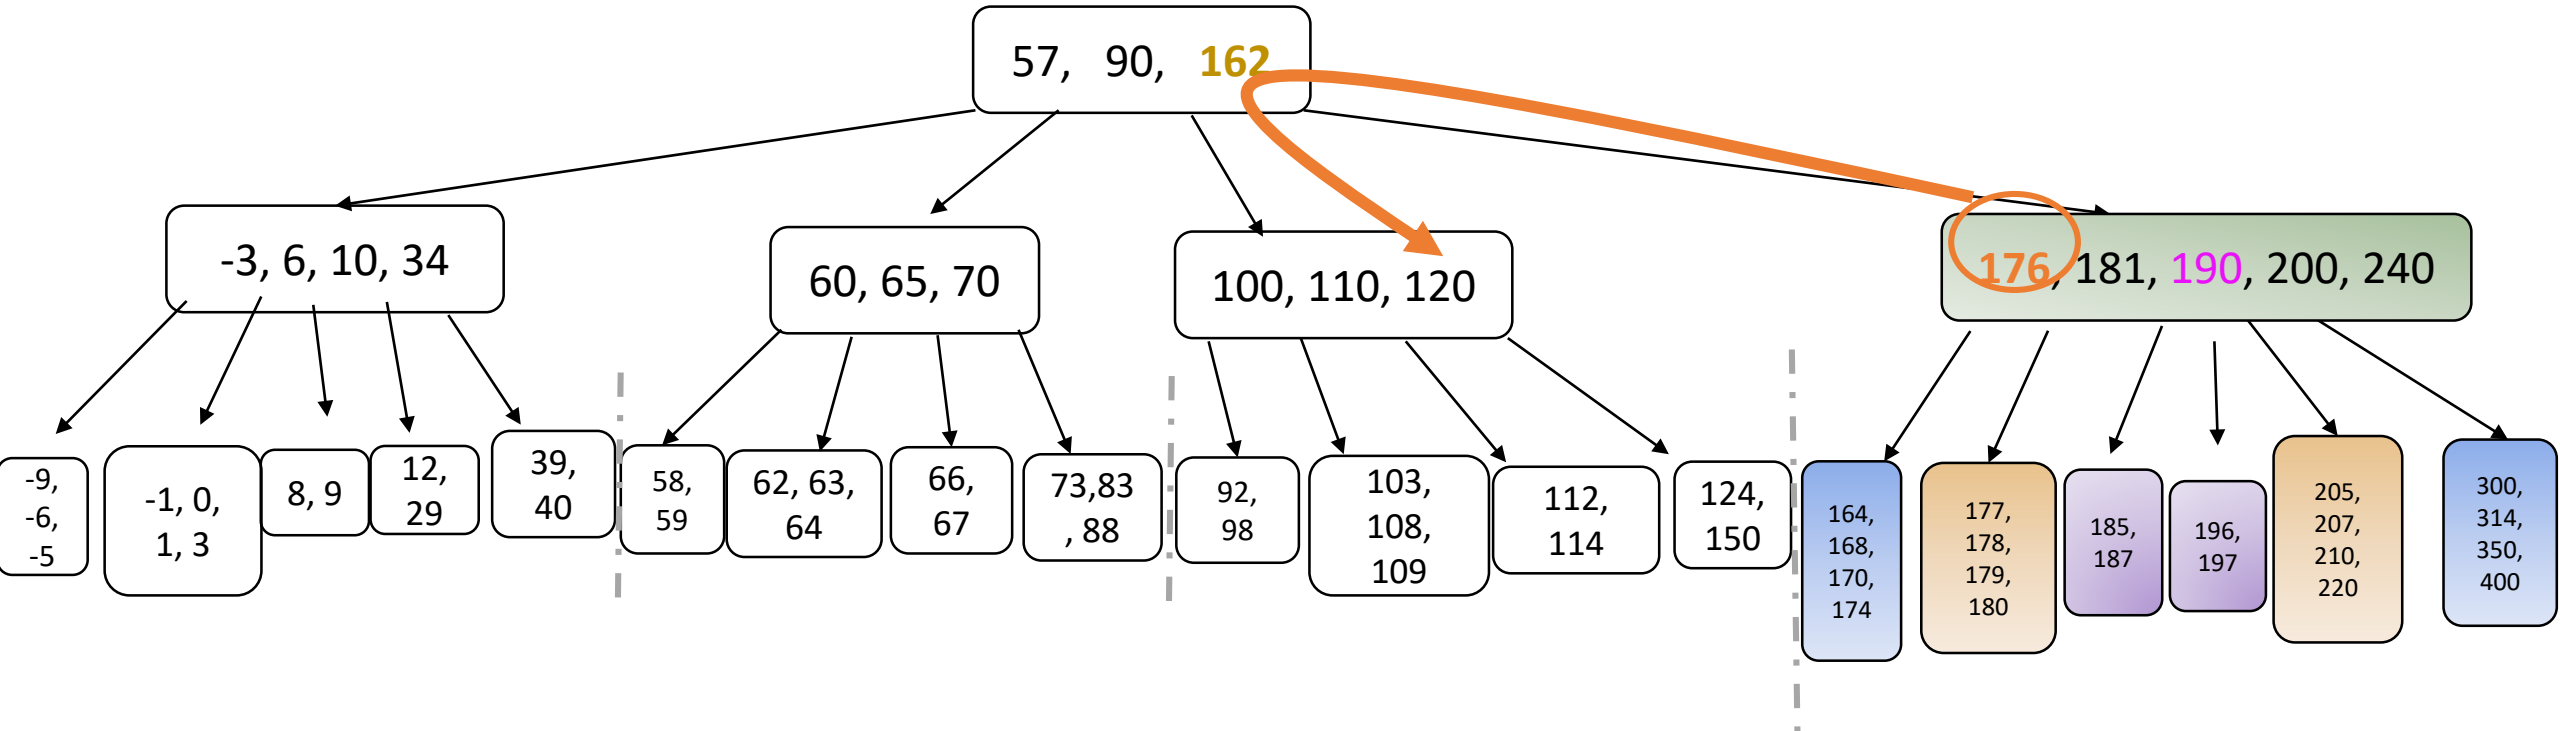
- Task: Insert 190
- Overflow 😞
- 1st Siblings, 2nd siblings and parent all full! 😞
- Solution: Split node, elevating 190 to parent and then rotate 176 to the left!

Example #6: Splitting, parent overflows and rotates



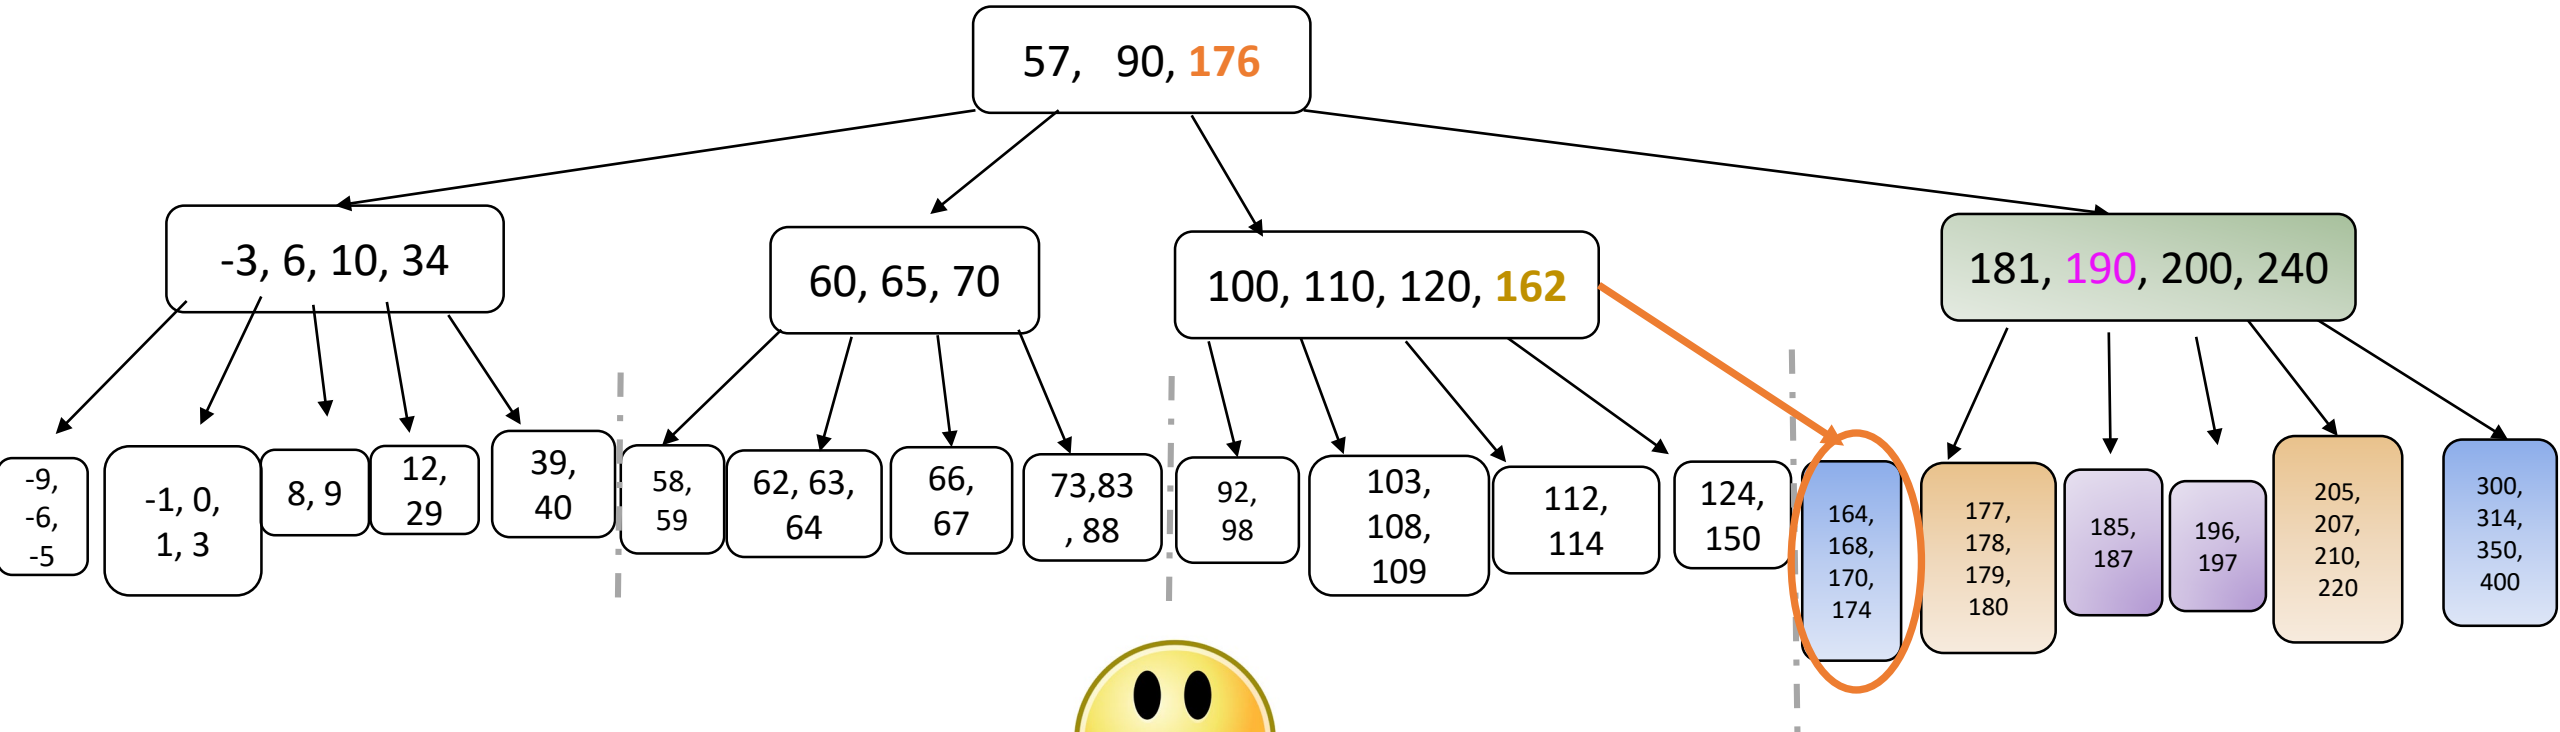
- Task: Insert 190
- Overflow 😞
- 1st Siblings, 2nd siblings and parent all full! 😞
- Solution: Split **node**, elevating 190 to **parent** and then rotate **176** to the left!

Example #6: Splitting, parent overflows and rotates



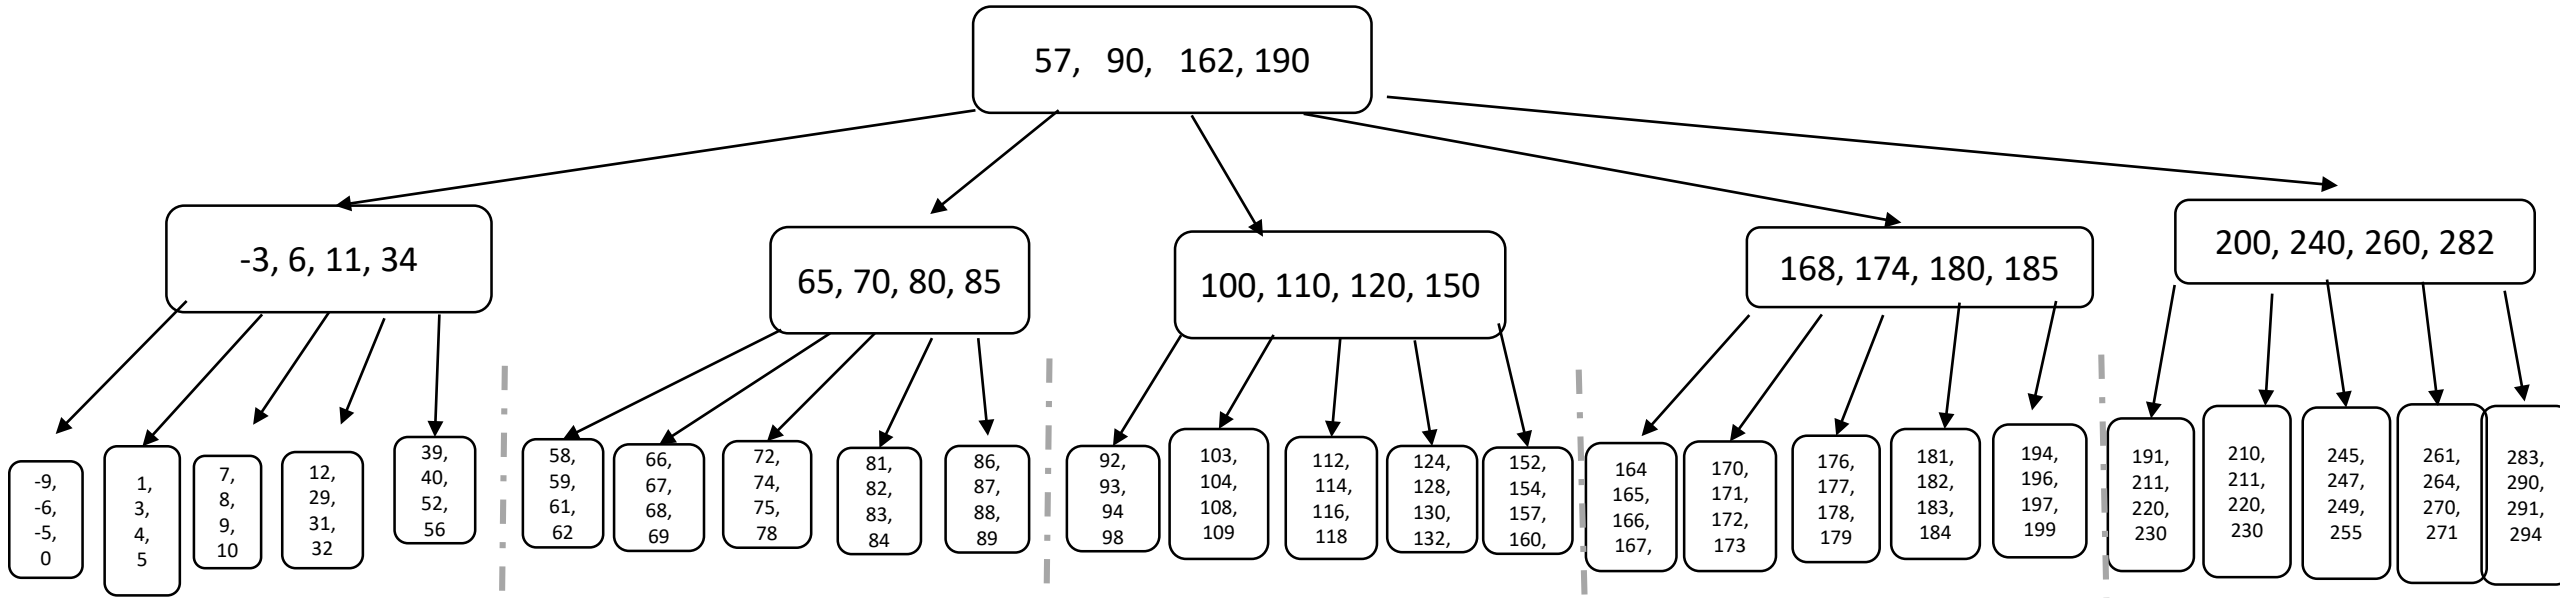
- Task: Insert 190
- Overflow 😞
- 1st Siblings, 2nd siblings and parent all full! 😞
- Solution: Split **node**, elevating 190 to **parent** and then rotate 176 to the left!

Example #6: Splitting, parent overflows and rotates



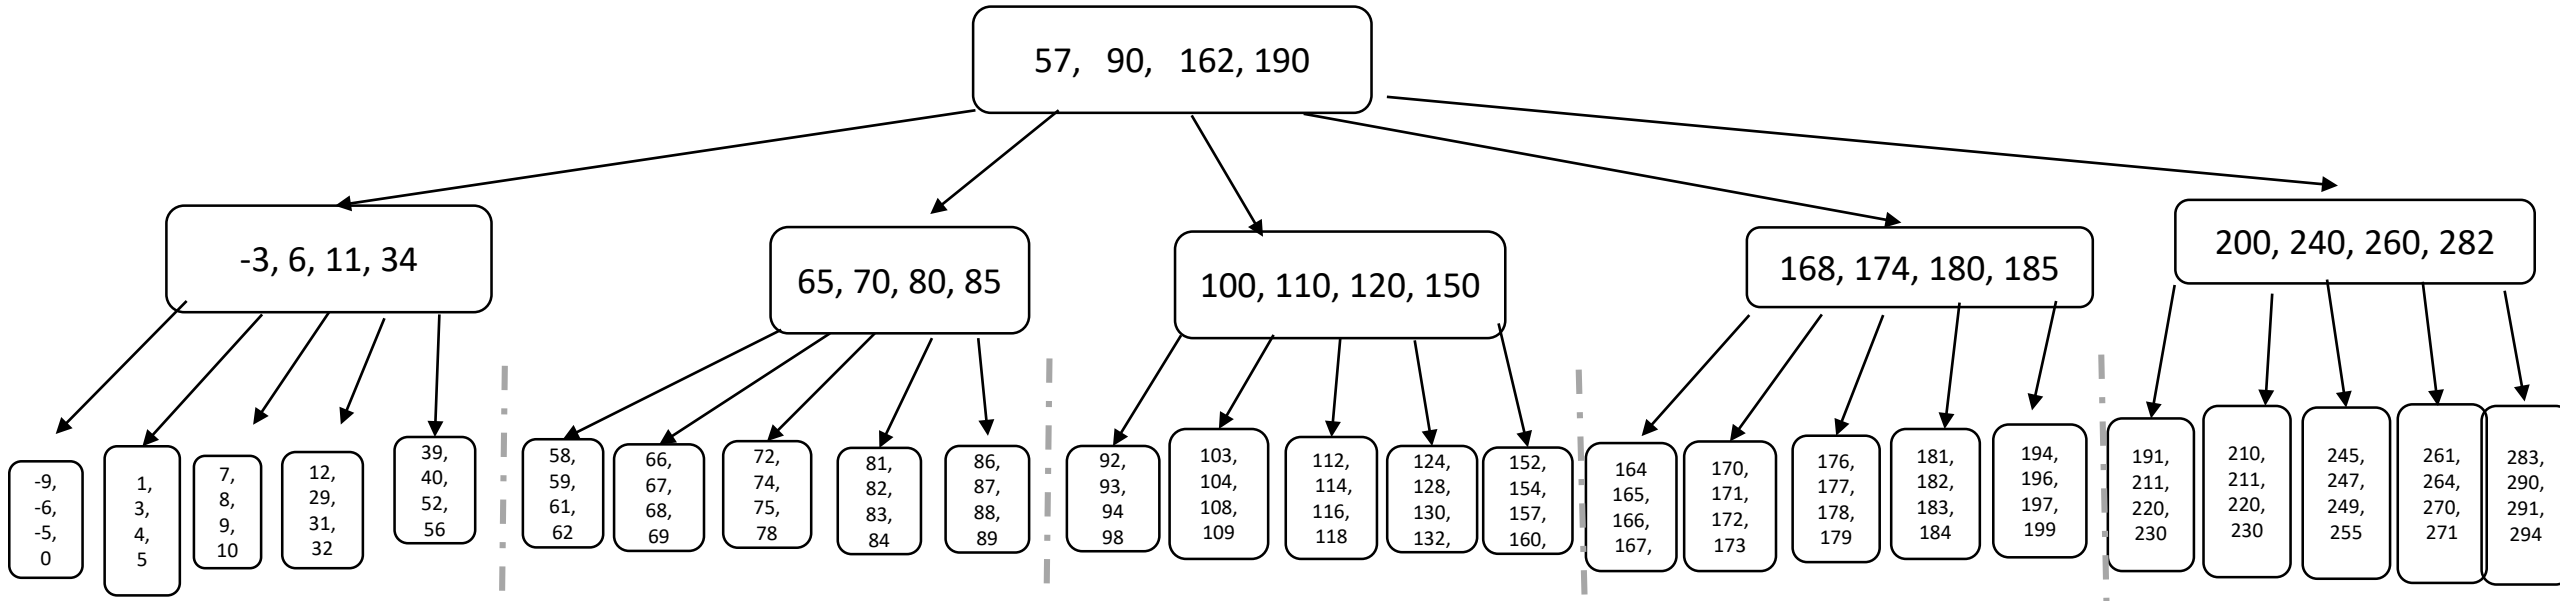
- Task: Insert 190
 - Overflow 😞
 - 1st Siblings, 2nd siblings and parent all full! 😞
 - Solution: Split **node**, elevating 190 to **parent** and then rotate 176 to the left!
- Don't forget to re-distribute subtrees when rotating keys above the leaf level!**

Example #7: Worst-case, splitting the root



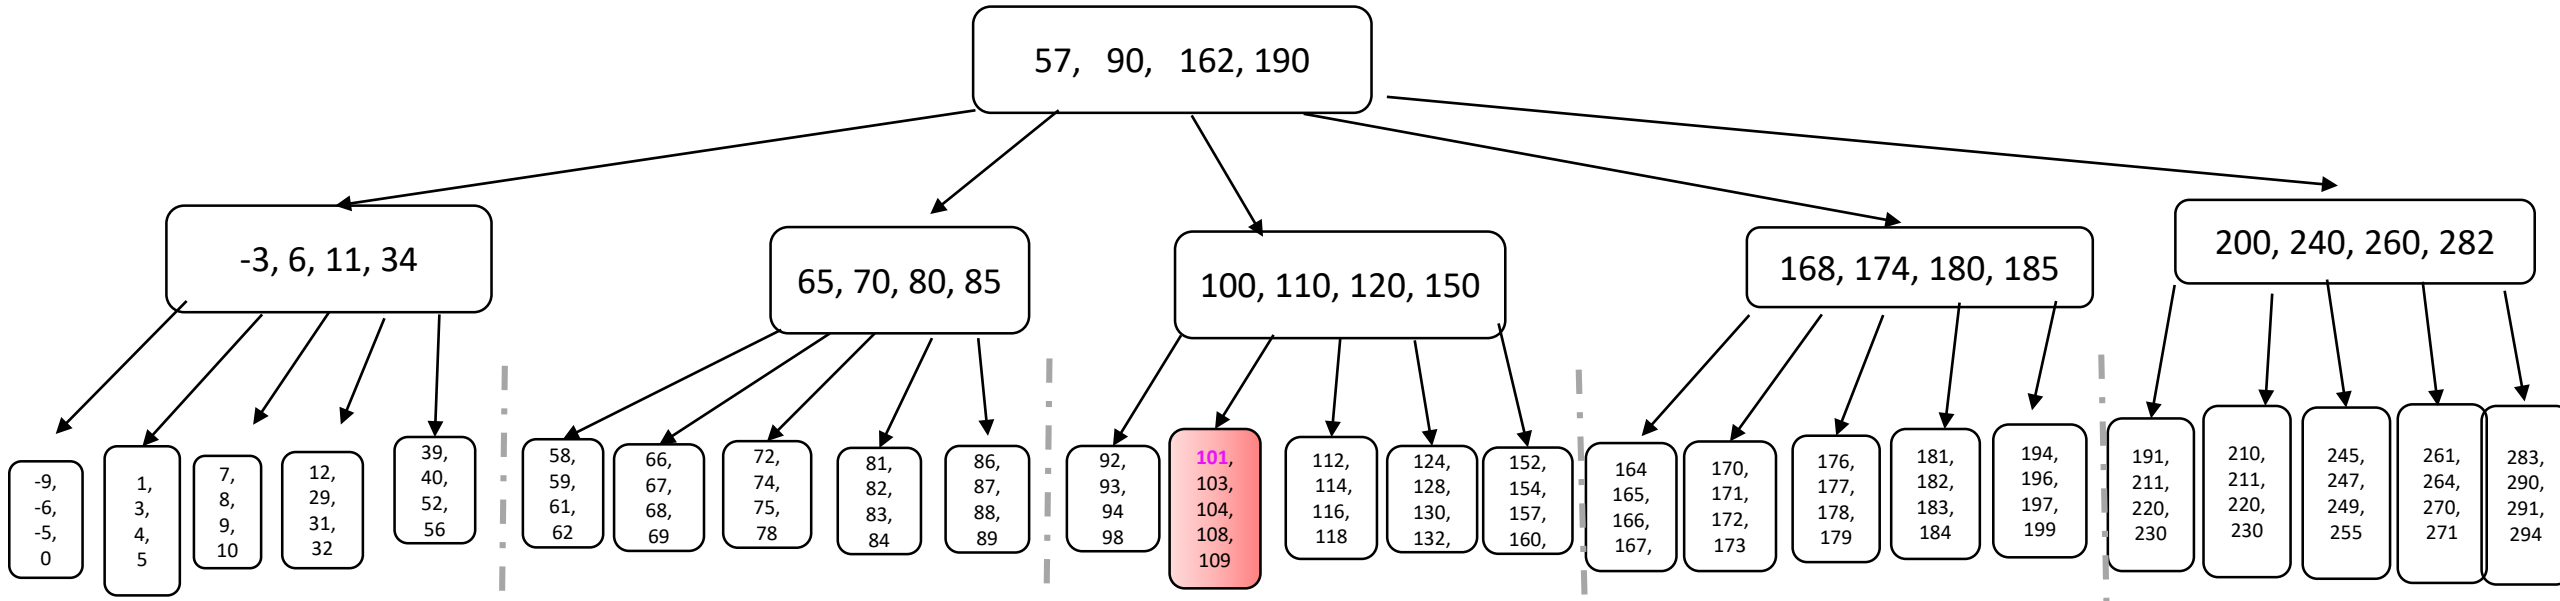
- Added a bunch of keys to make the tree filled to the brim!
- $4 + 4^2 + 4^3 = 20 + 64 = 84$ keys with $h = 2$ and $p = 5$.

Example #7: Worst-case, splitting the root



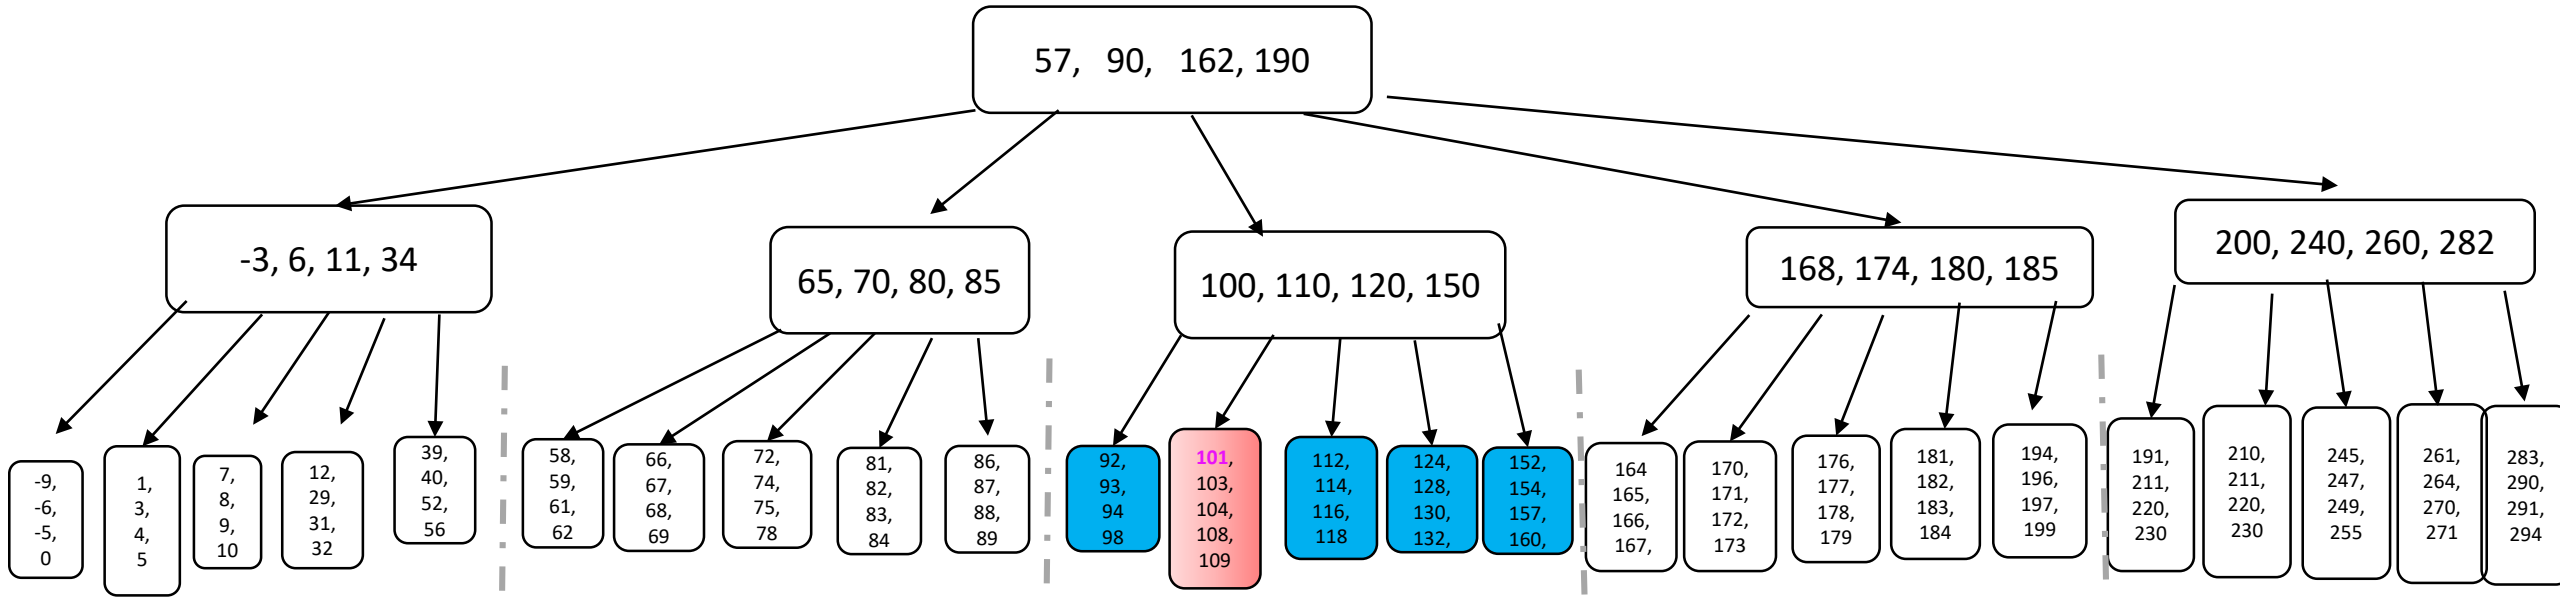
- Task: Insert 101

Example #7: Worst-case, splitting the root




- Task: Insert 101
- Overflow 😞

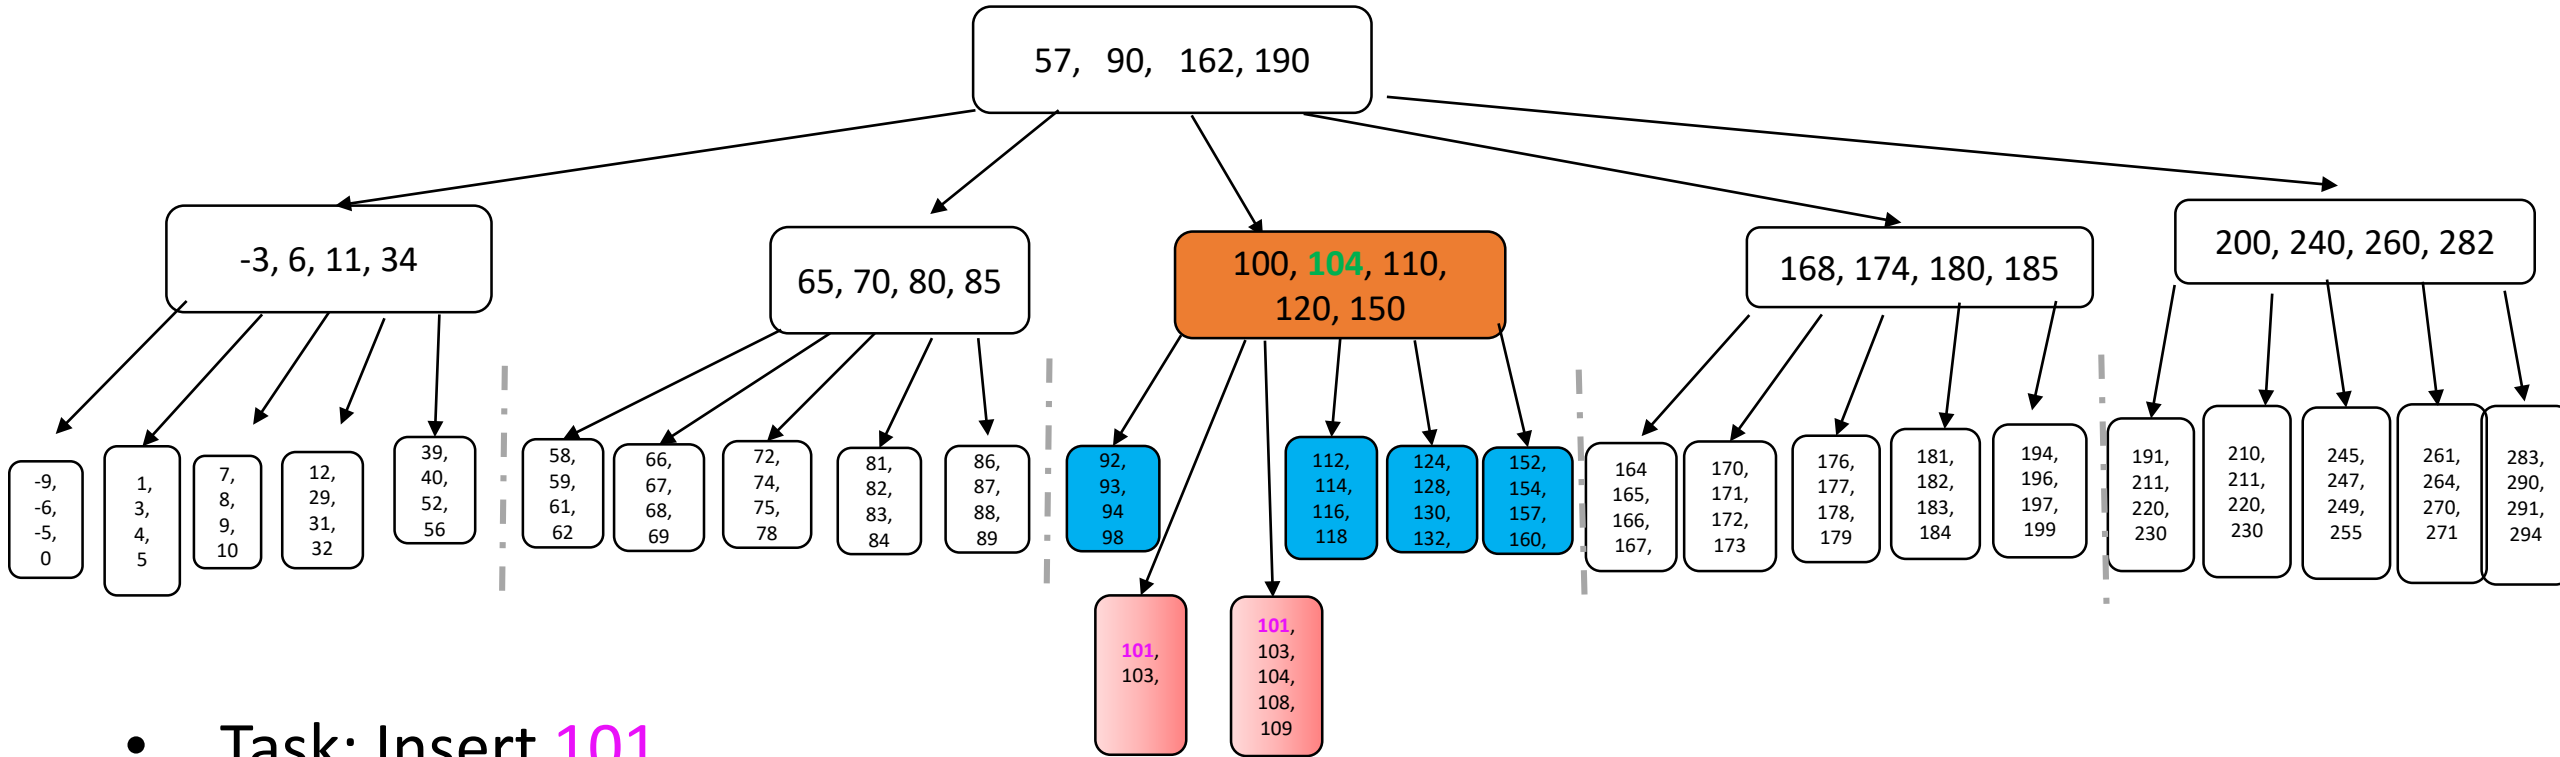
Example #7: Worst-case, splitting the root



- Task: Insert 101
- Overflow 😞
- No possible key rotation to siblings! 😞



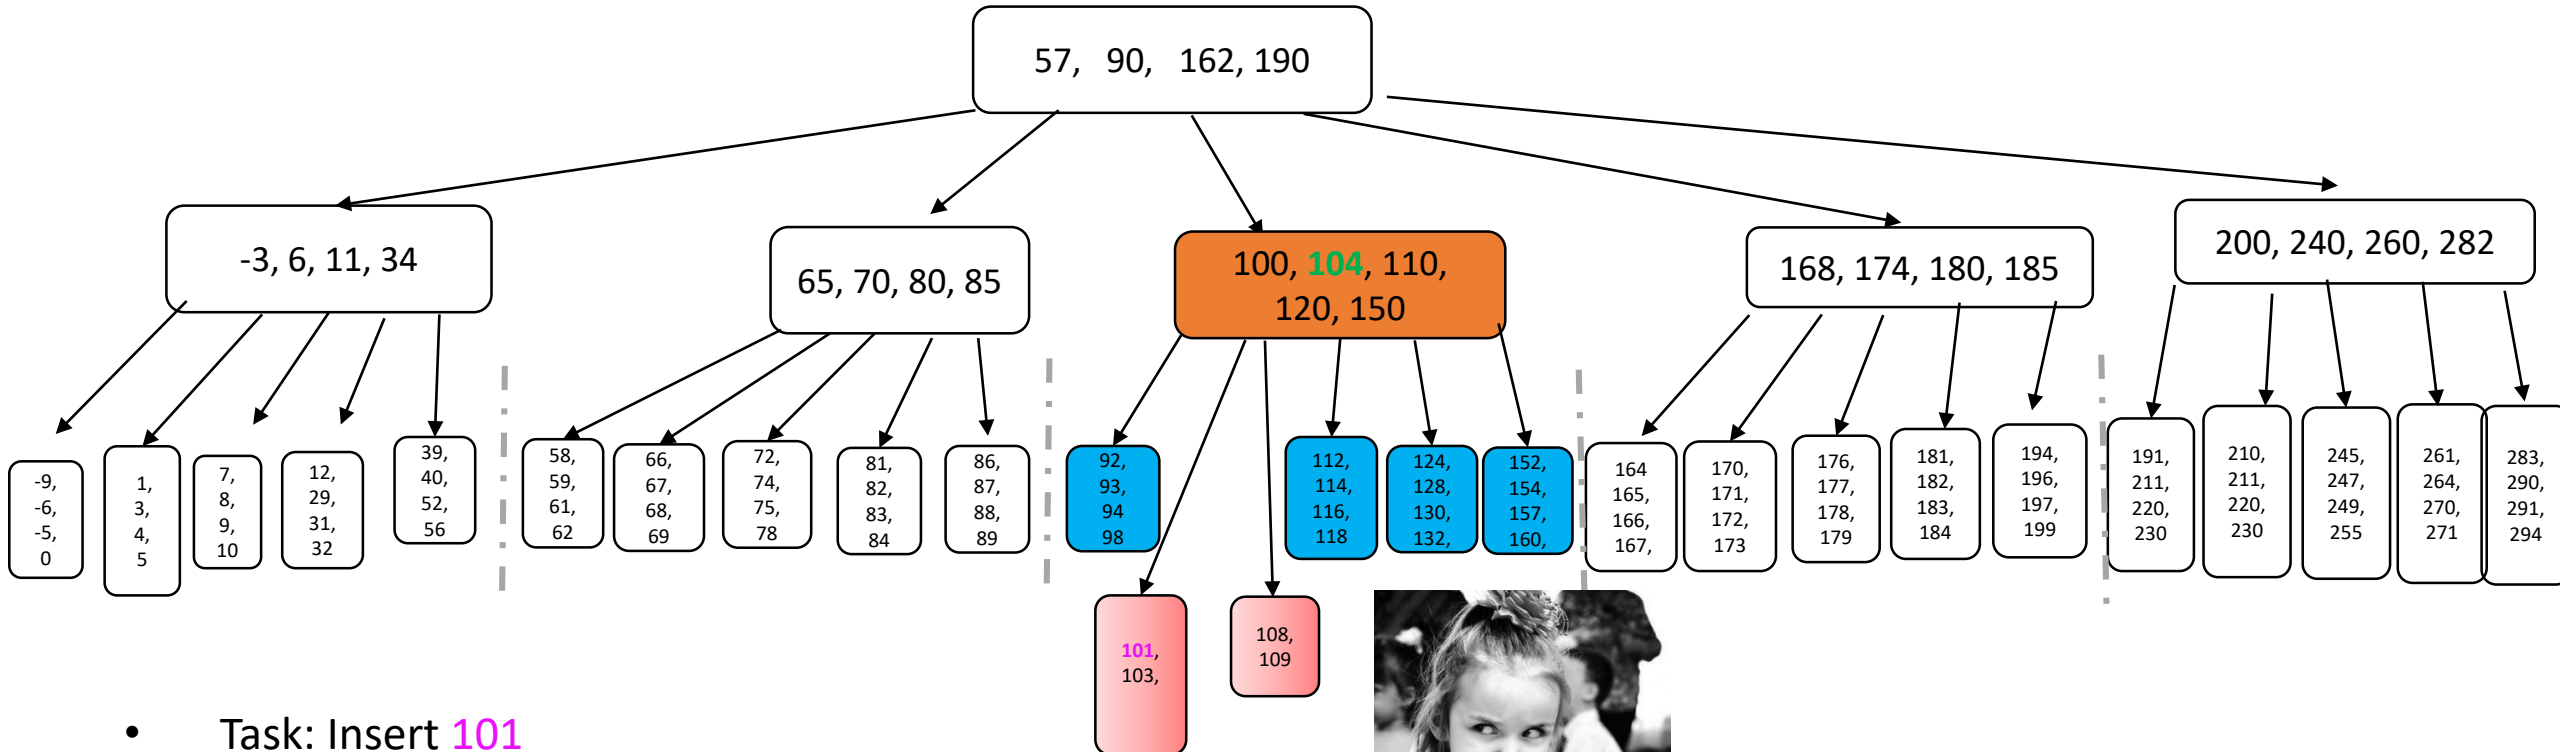
Example #7: Worst-case, splitting the root



- Task: Insert 101
- Overflow 😞
- No possible key rotation anywhere in the tree! 😞
- Split and elevate 104 (middle key) to parent...



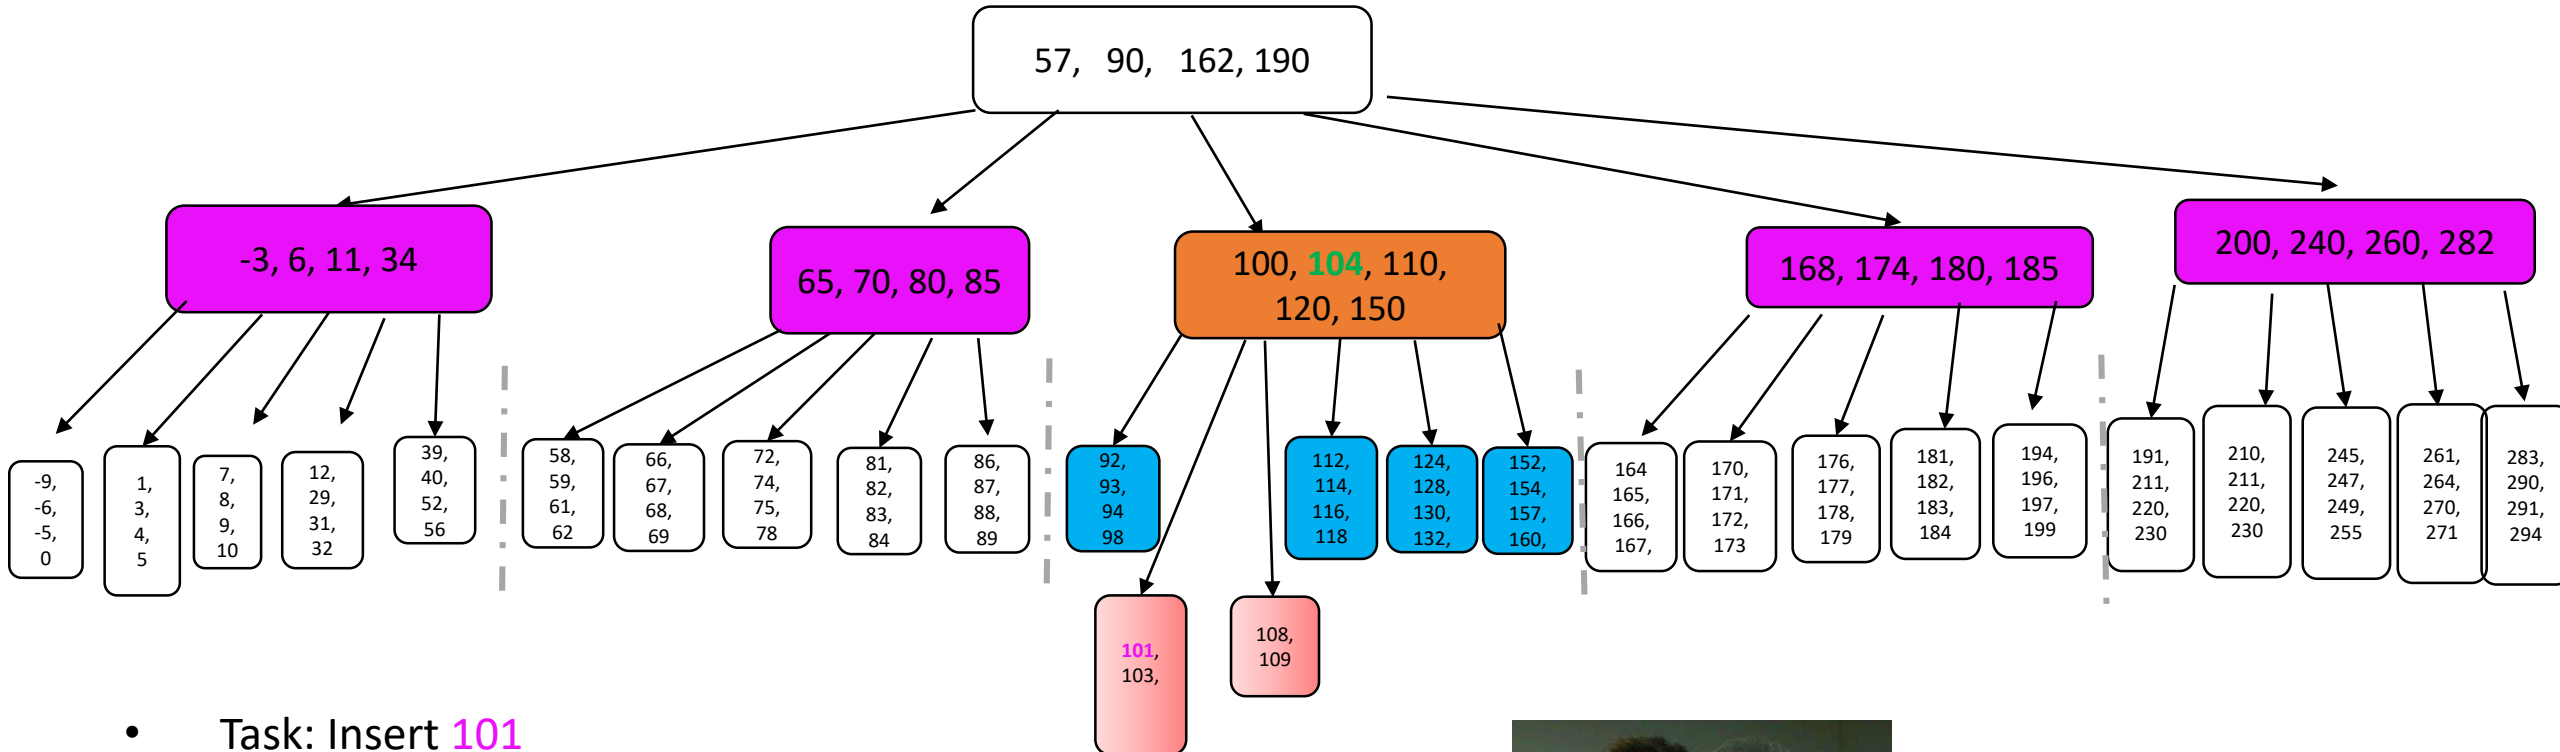
Example #7: Worst-case, splitting the root



- Task: Insert 101
- Overflow 😞
- No possible key rotation anywhere in the tree! 😞
- Split and elevate 104 (middle key) to parent...
- Which also overflows!



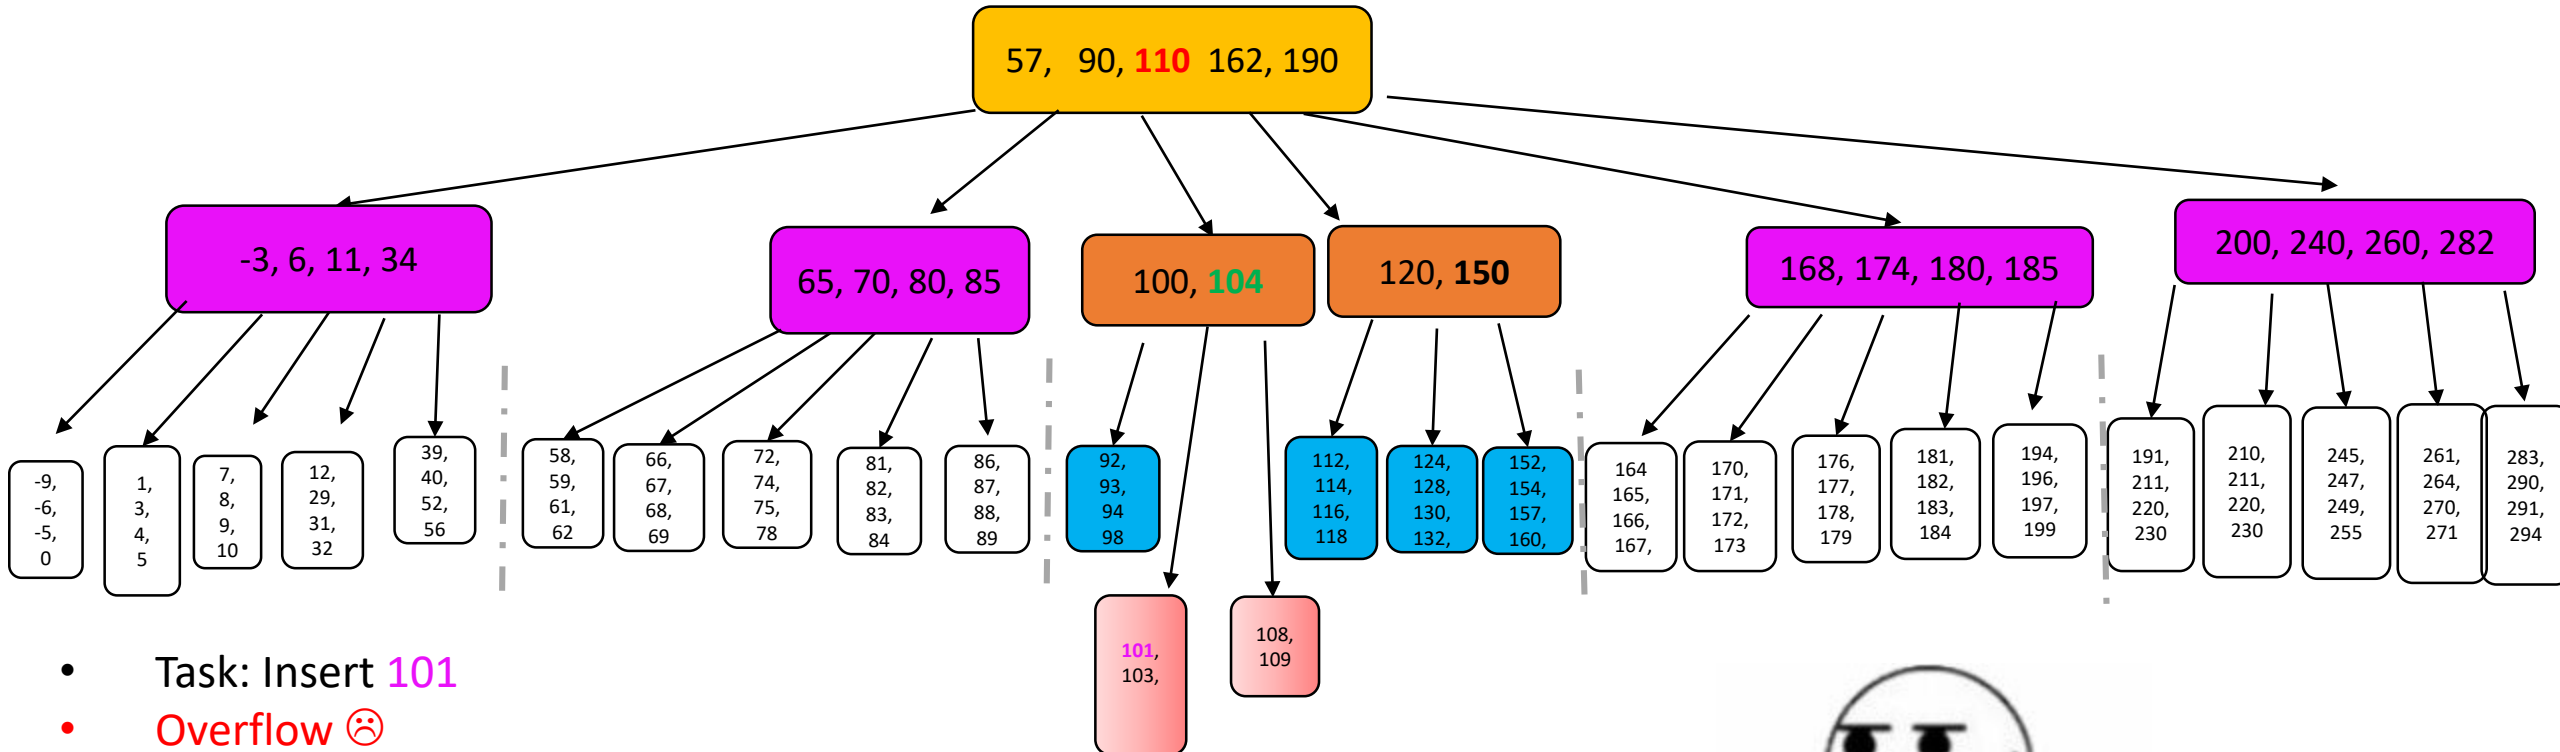
Example #7: Worst-case, splitting the root



- Task: Insert 101
- Overflow 😞
- No possible key rotation anywhere in the tree! 😞
- Split and elevate 104 (middle key) to parent...
- Which also overflows!
- AND ITS OWN SIBLINGS CAN'T HELP WITH ROTATIONS!



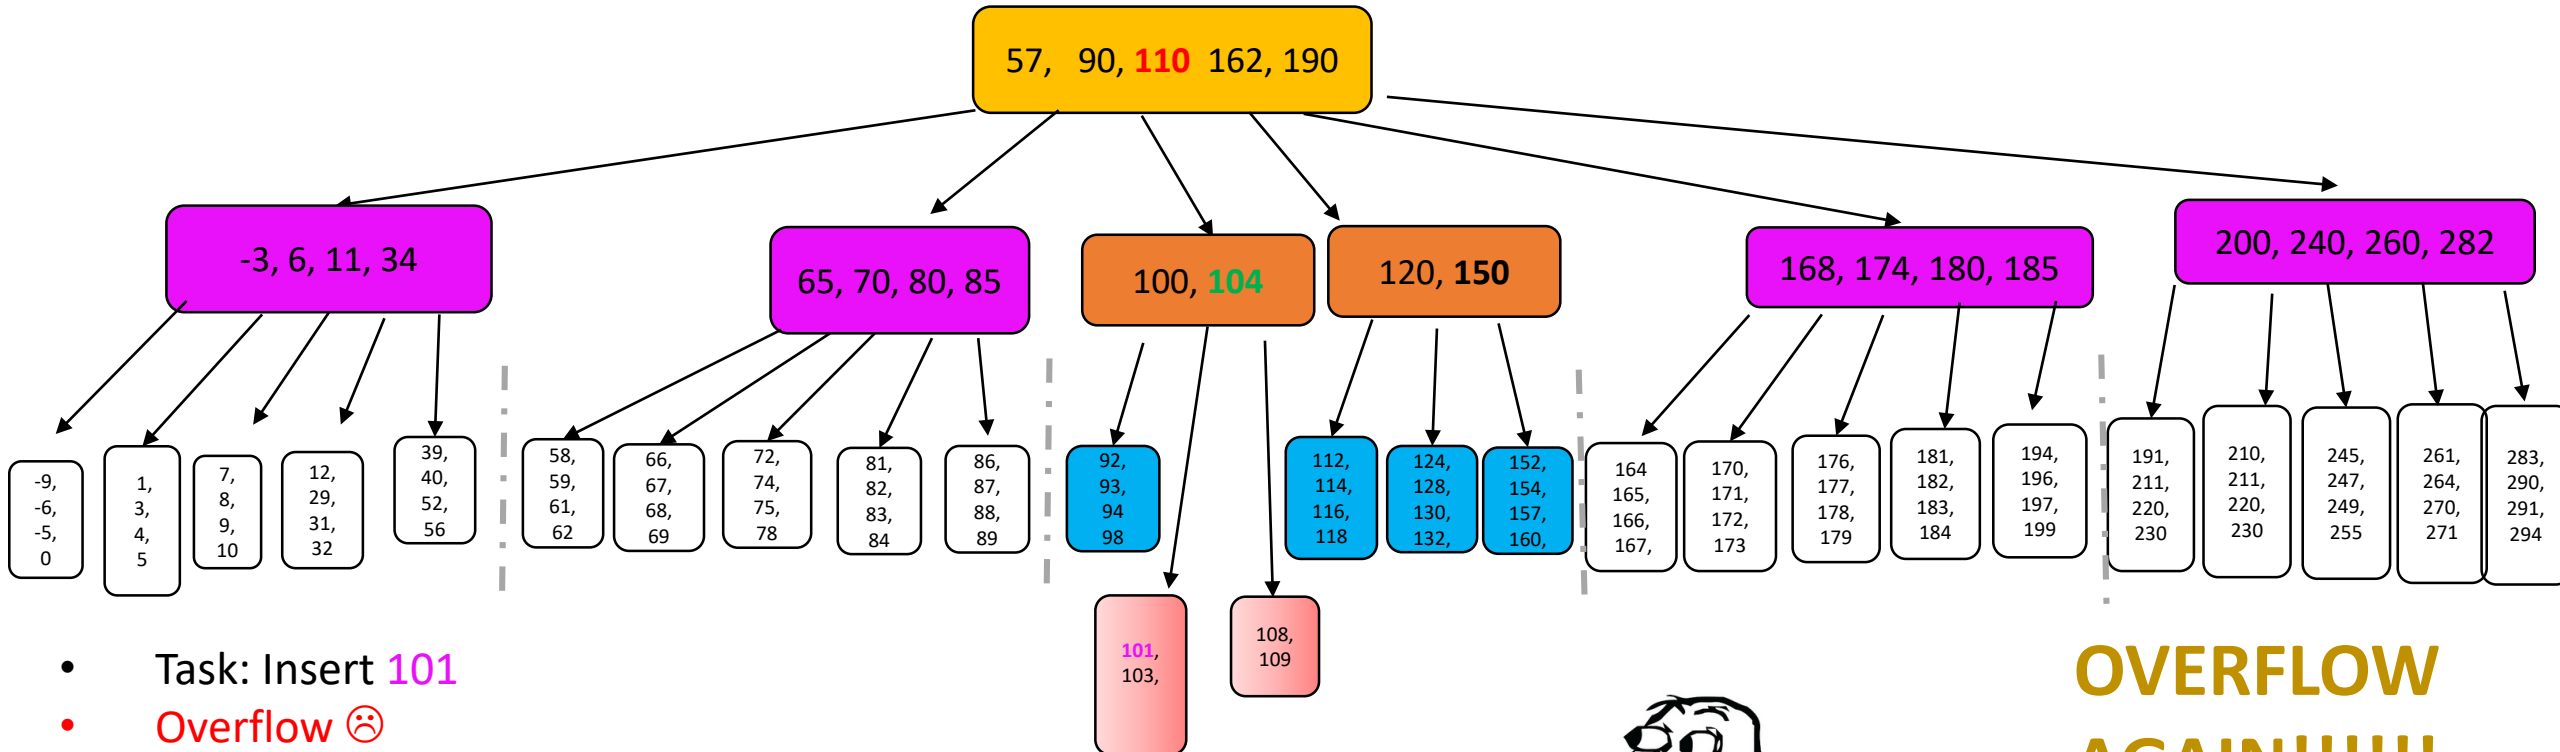
Example #7: Worst-case, splitting the root




- Task: Insert 101
- Overflow ☹️
- No possible key rotation anywhere in the tree! ☹️
- Split and elevate 104 (middle key) to parent...
- Which also overflows!
- AND ITS OWN SIBLINGS CAN'T HELP WITH ROTATIONS!
- Split current node and elevate 110 (middle key) to the parent (the root).



Example #7: Worst-case, splitting the root

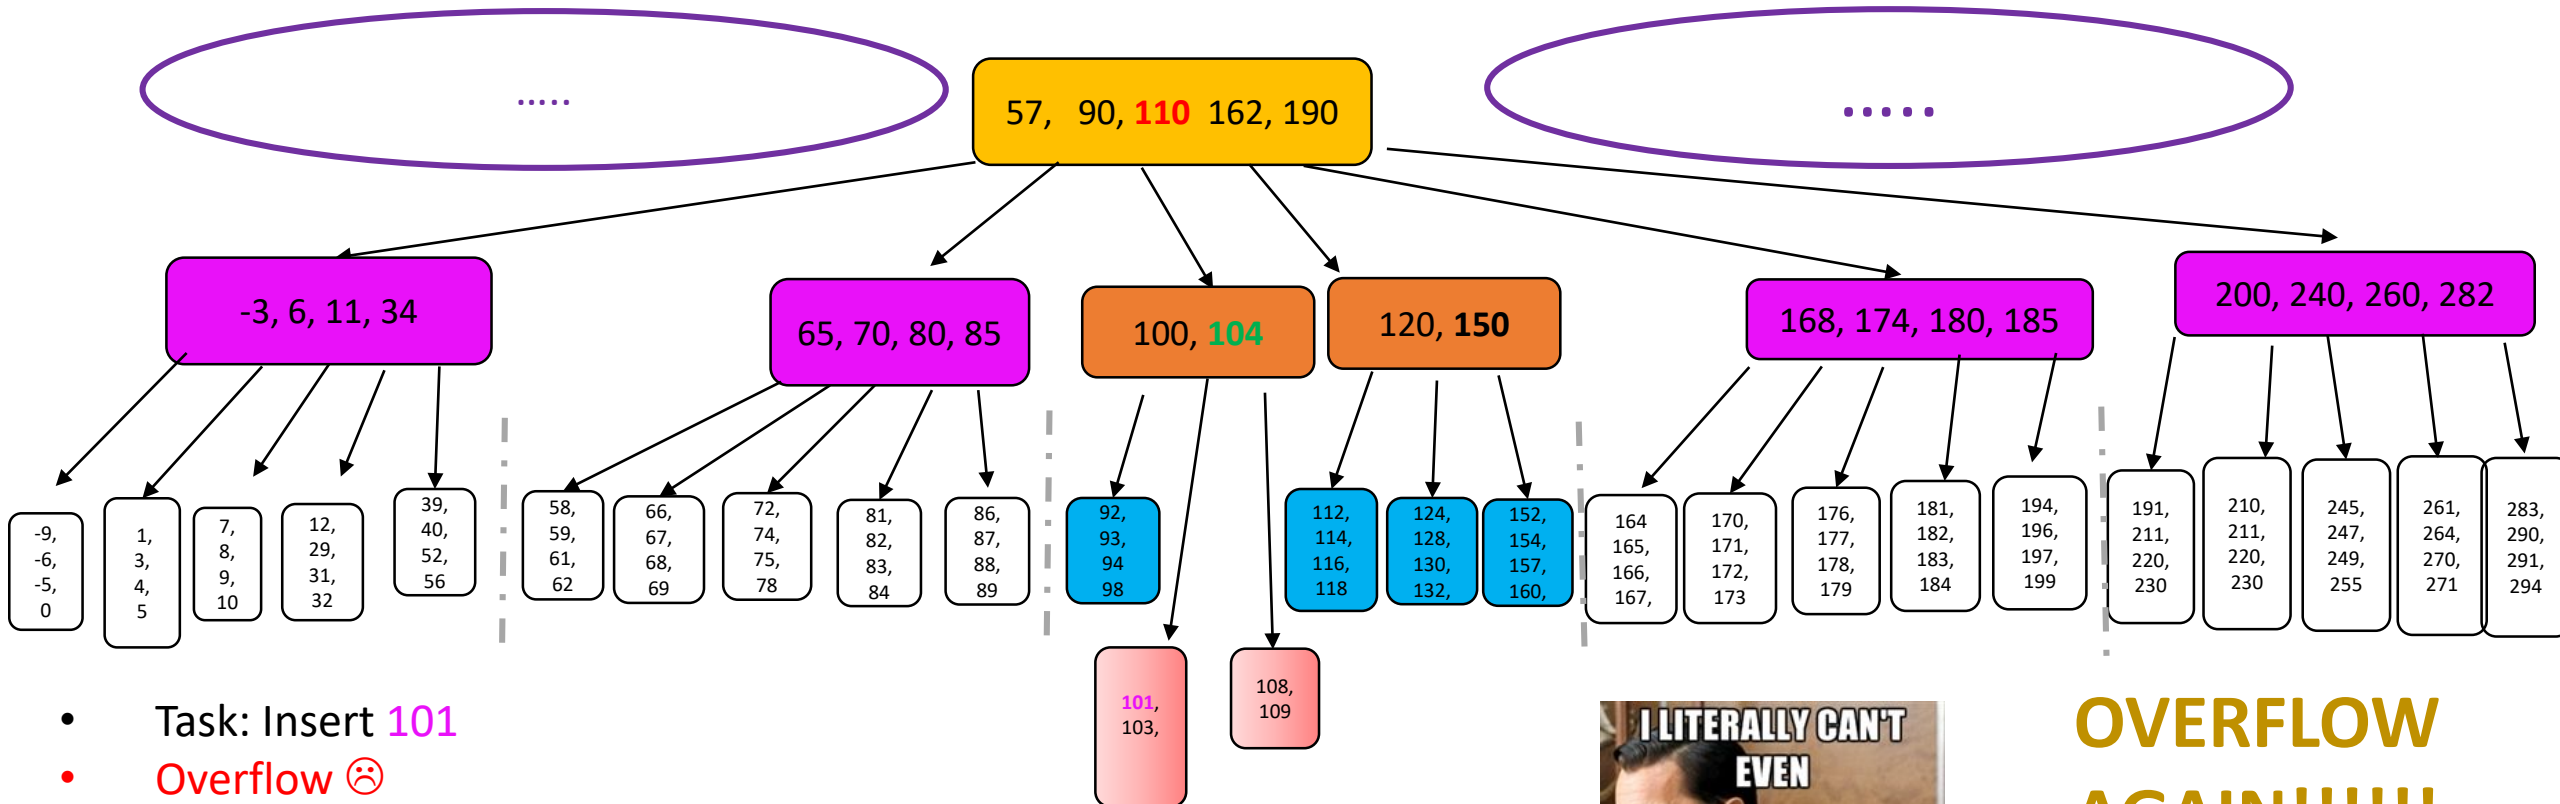


- Task: Insert 101
- Overflow 😞
- No possible key rotation anywhere in the tree! 😞
- Split and elevate 104 (middle key) to parent...
- Which also overflows!
- AND ITS OWN SIBLINGS CAN'T HELP WITH ROTATIONS!
- Split current node and elevate 110 (middle key) to the parent (the root).




**OVERFLOW
AGAIN!!!!!!**

Example #7: Worst-case, splitting the root



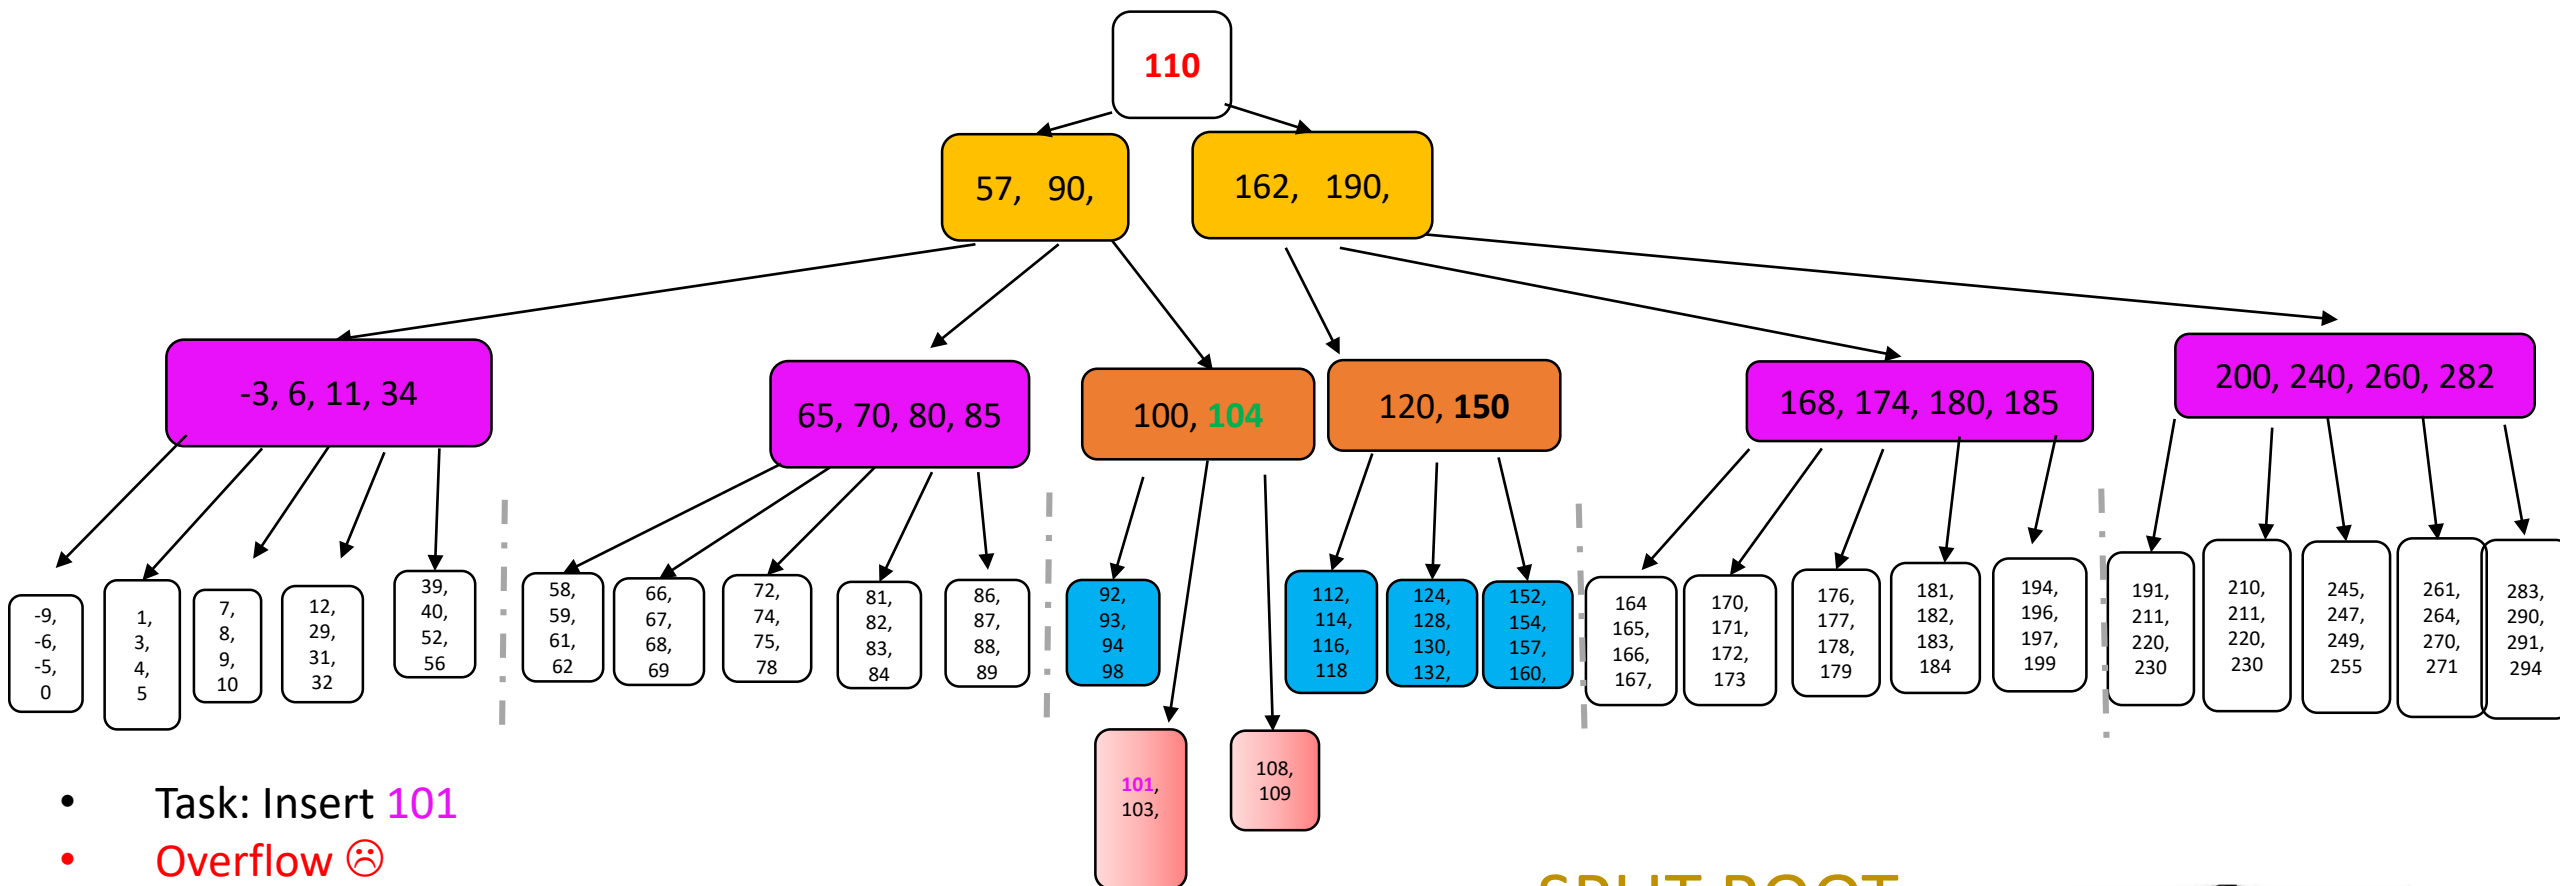
- Task: Insert 101
- Overflow 😞
- No possible key rotation anywhere in the tree! 😞
- Split and elevate 104 (middle key) to parent...
- Which also overflows!
- AND ITS OWN SIBLINGS CAN'T HELP WITH ROTATIONS!
- Split current node and elevate 110 (middle key) to the parent (the root).



**OVERFLOW
AGAIN!!!!!!**

**NO SIBLINGS
TO ROTATE TO!**

Example #7: Worst-case, splitting the root



- Task: Insert 101
- Overflow ☹️
- No possible key rotation anywhere in the tree! ☹️
- Split and elevate 104 (middle key) to parent...
- Which also overflows!
- AND ITS OWN SIBLINGS CAN'T HELP WITH ROTATIONS!
- Split current node and elevate 110 (middle key) to the parent (the root).

SPLIT ROOT,
ELEVATING 110,
AND DONE.



Dealing with underflows

- Since deleting any key will always boil down to deleting some key from the leaf level, underflows can only occur at the leaf level

Dealing with underflows

- Since deleting any key will always boil down to deleting some key from the leaf level, underflows can only occur at the leaf level
- We will once again attempt key rotations first.

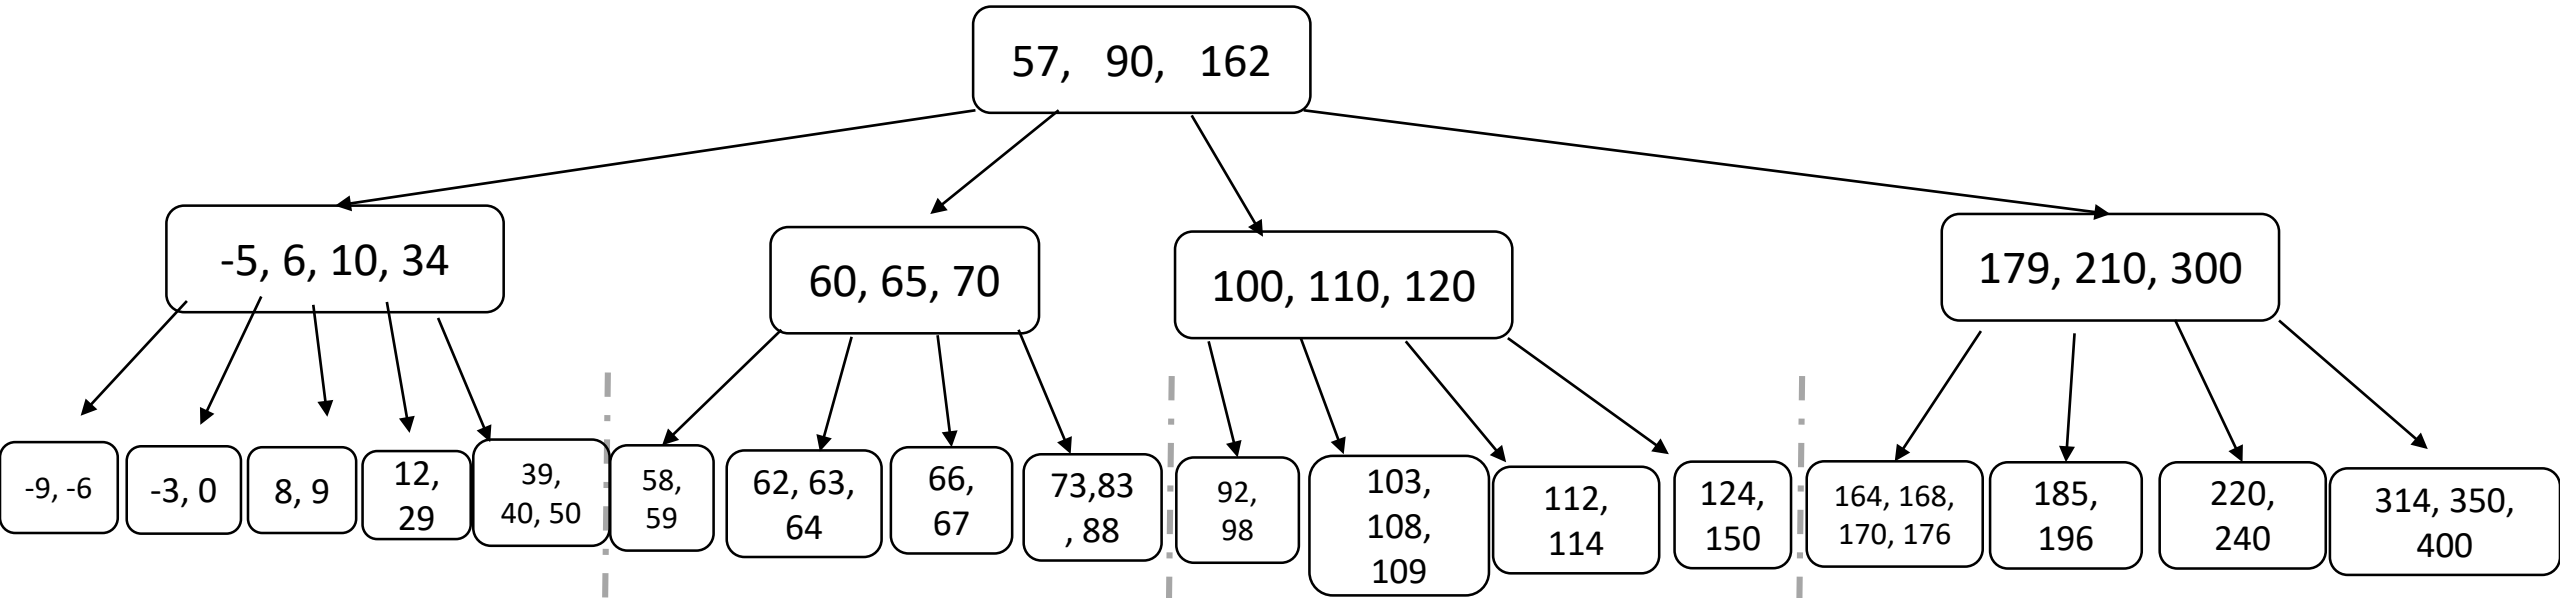
Dealing with underflows

- Since deleting any key will always boil down to deleting some key from the leaf level, underflows can only occur at the leaf level
- We will once again attempt key rotations first.
- Idea: Look at your sibling nodes (≤ 2) and locate at least one that would not underflow if you were to take a key from it
 - If you find one such sibling, rotate the key “closest” to the current node from that sibling through the parent.
 - This is always either the smallest or the largest key of the sibling

Dealing with underflows

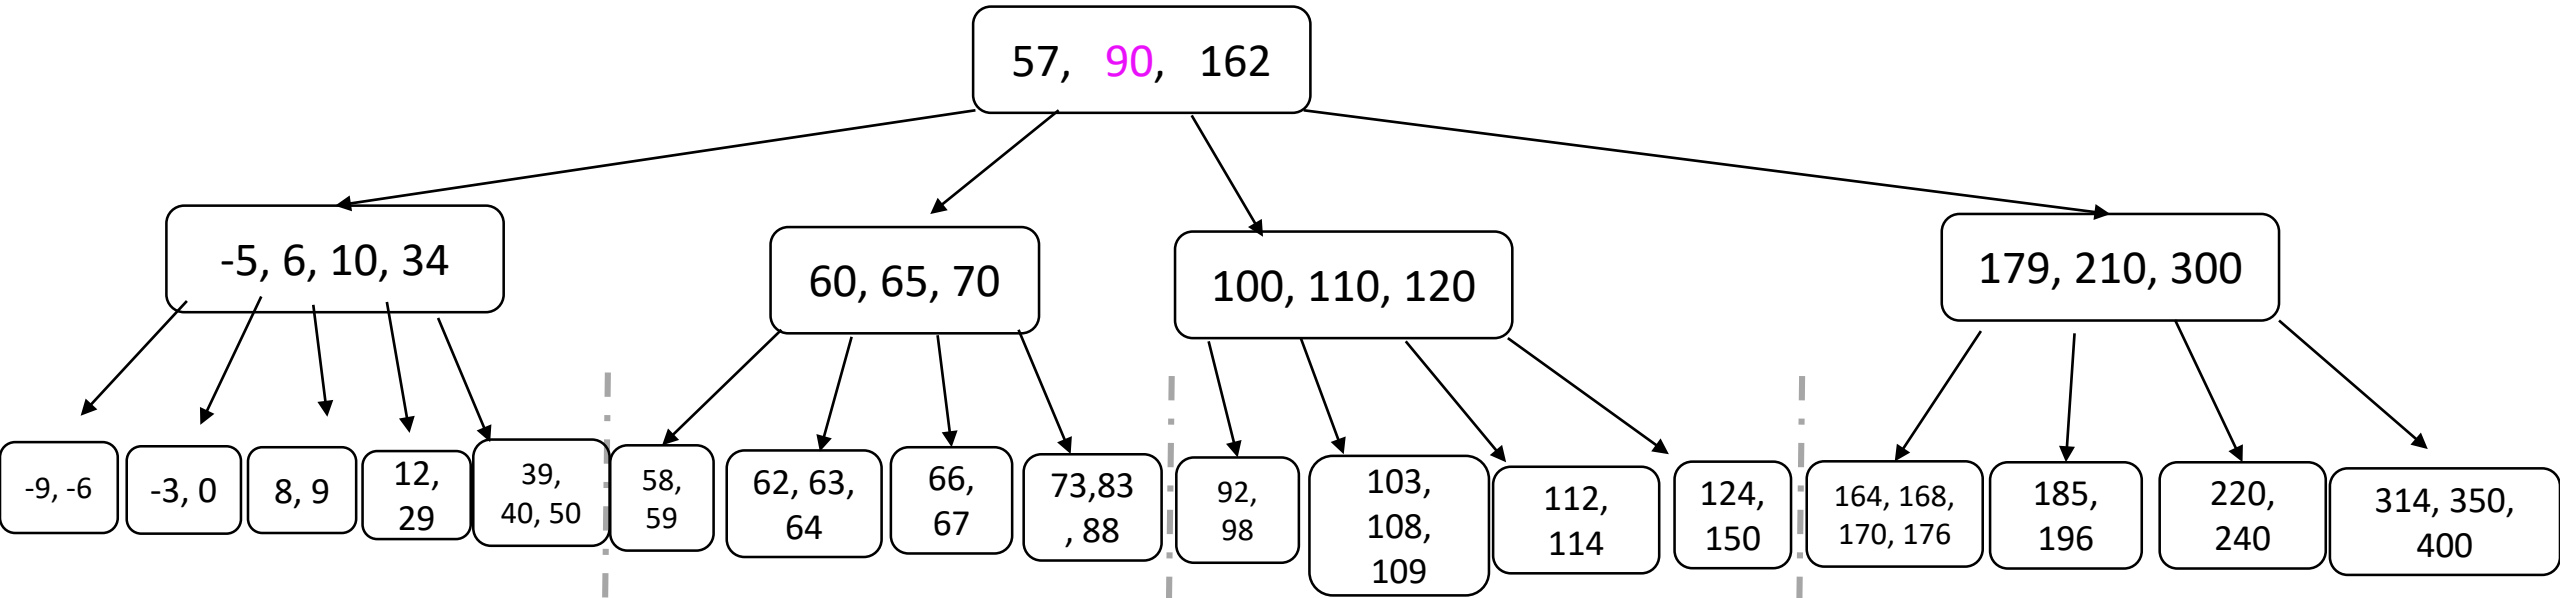
- Since deleting any key will always boil down to deleting some key from the leaf level, underflows can only occur at the leaf level
- We will once again attempt key rotations first.
- Idea: Look at your sibling nodes (≤ 2) and locate at least one that would not underflow if you were to take a key from it
 - If you find one such sibling, rotate the key “closest” to the current node from that sibling through the parent.
 - This is always either the smallest or the largest key of the sibling
- Familiar generalization to “further” siblings applies here too.
- If all fails, then merge current node with parent key
 - This may propagate an underflow in the parent, which should be dealt with recursively...

Example #0: Leaf deletion, no underflows



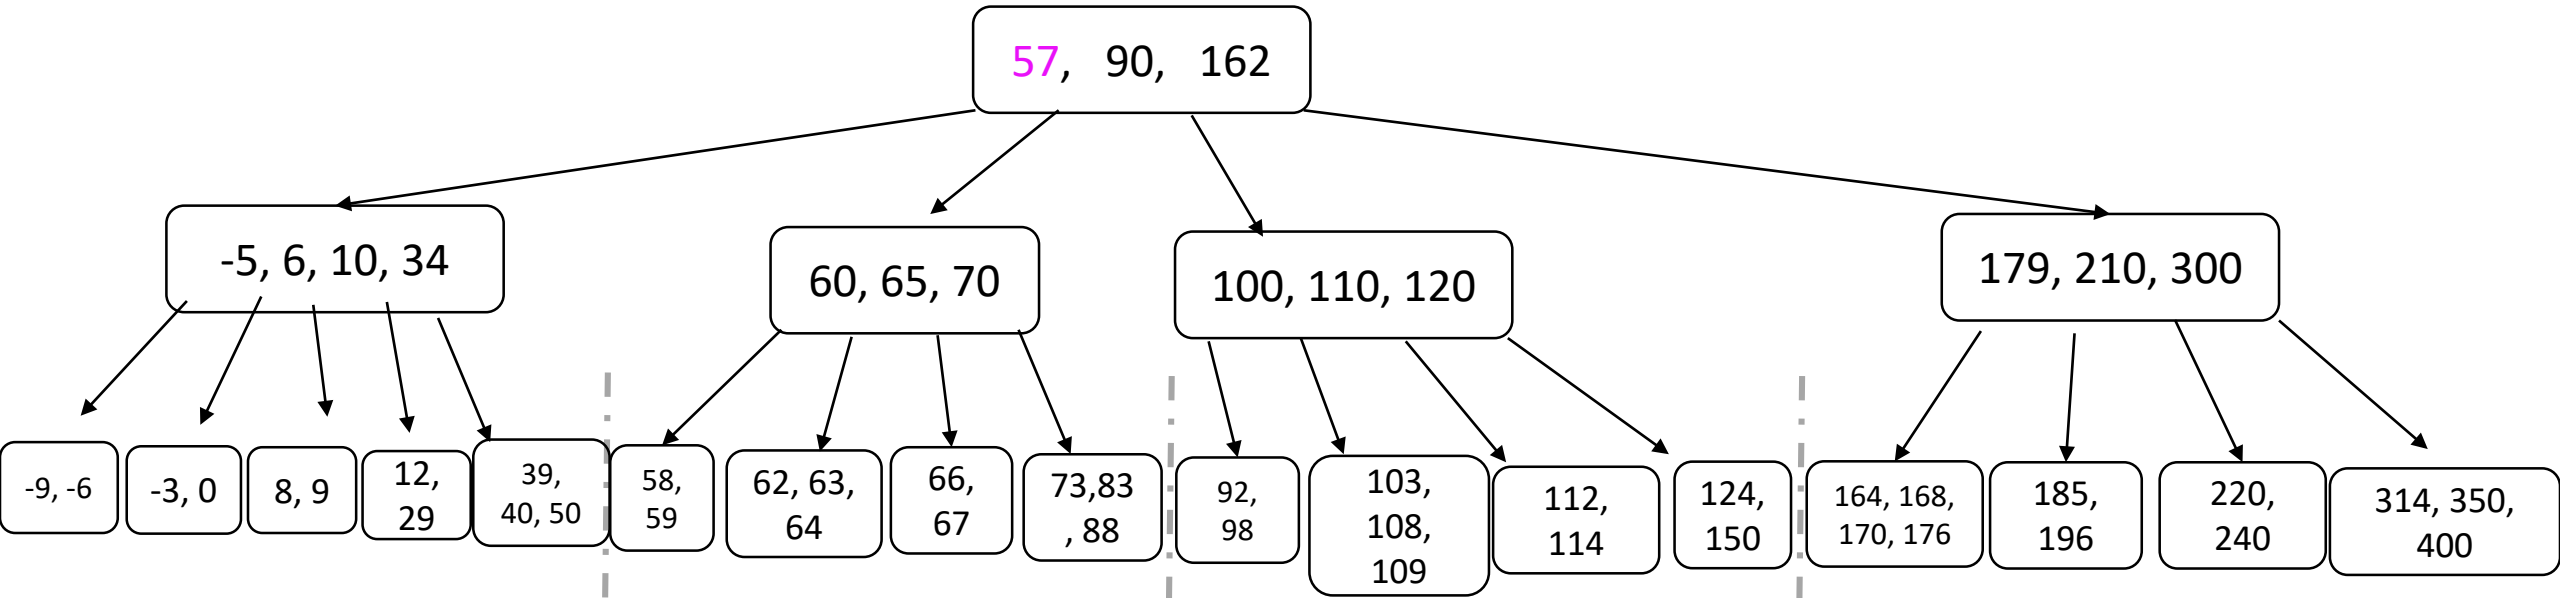
- Task: delete 62

Example #0: Leaf deletion, no underflows



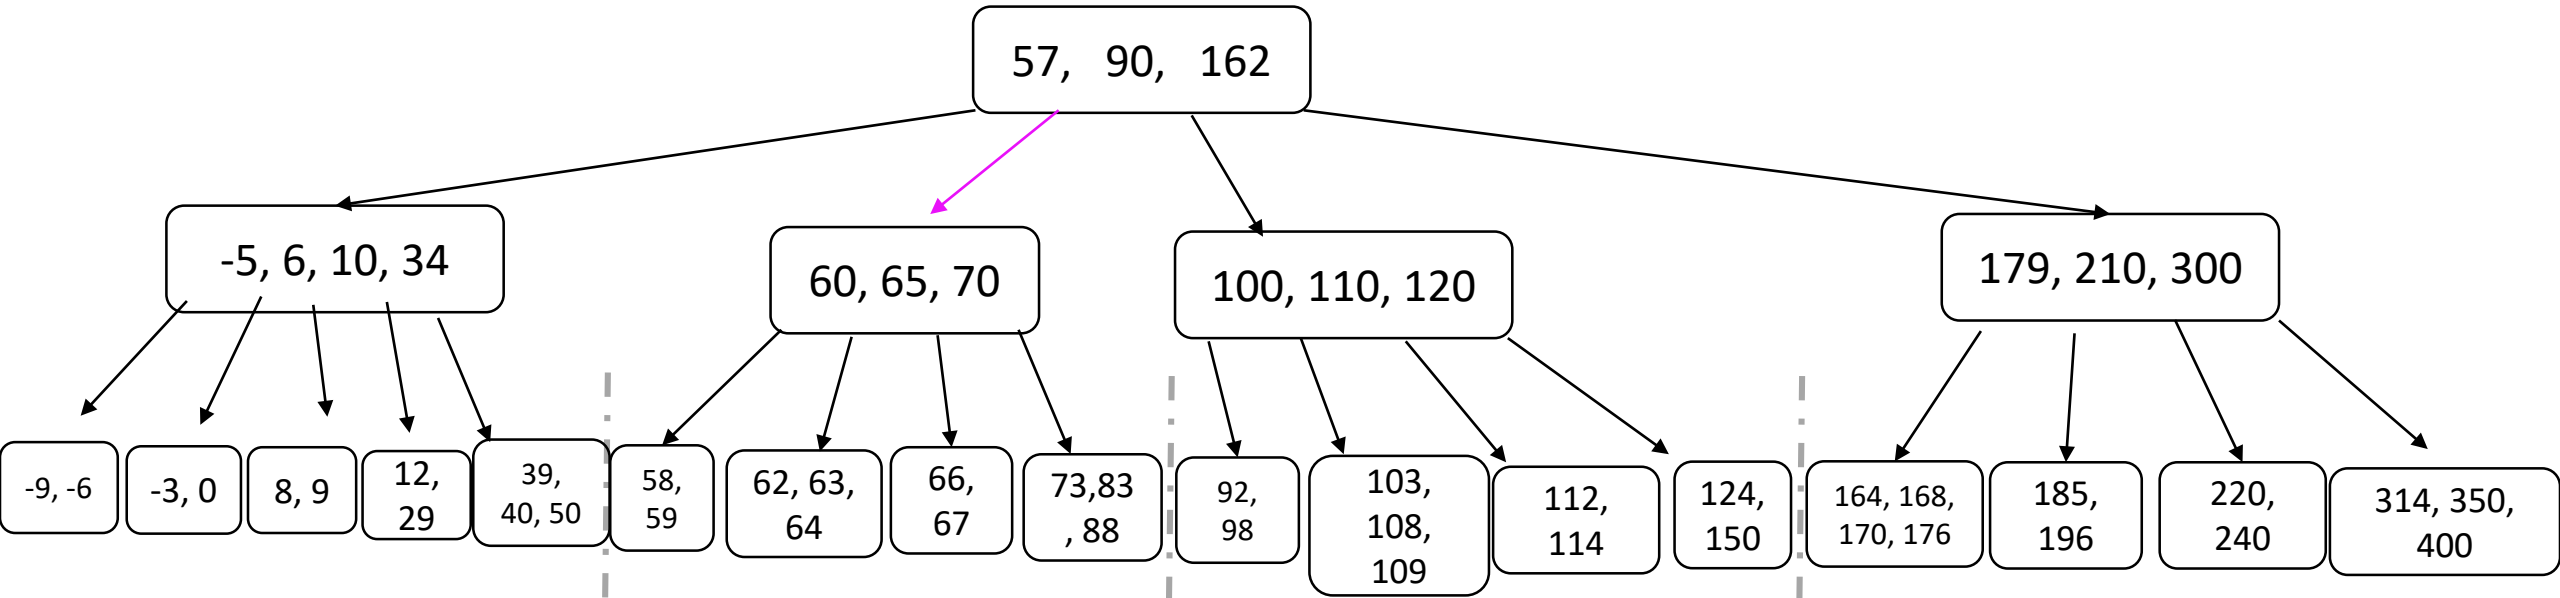
- Task: delete 62

Example #0: Leaf deletion, no underflows



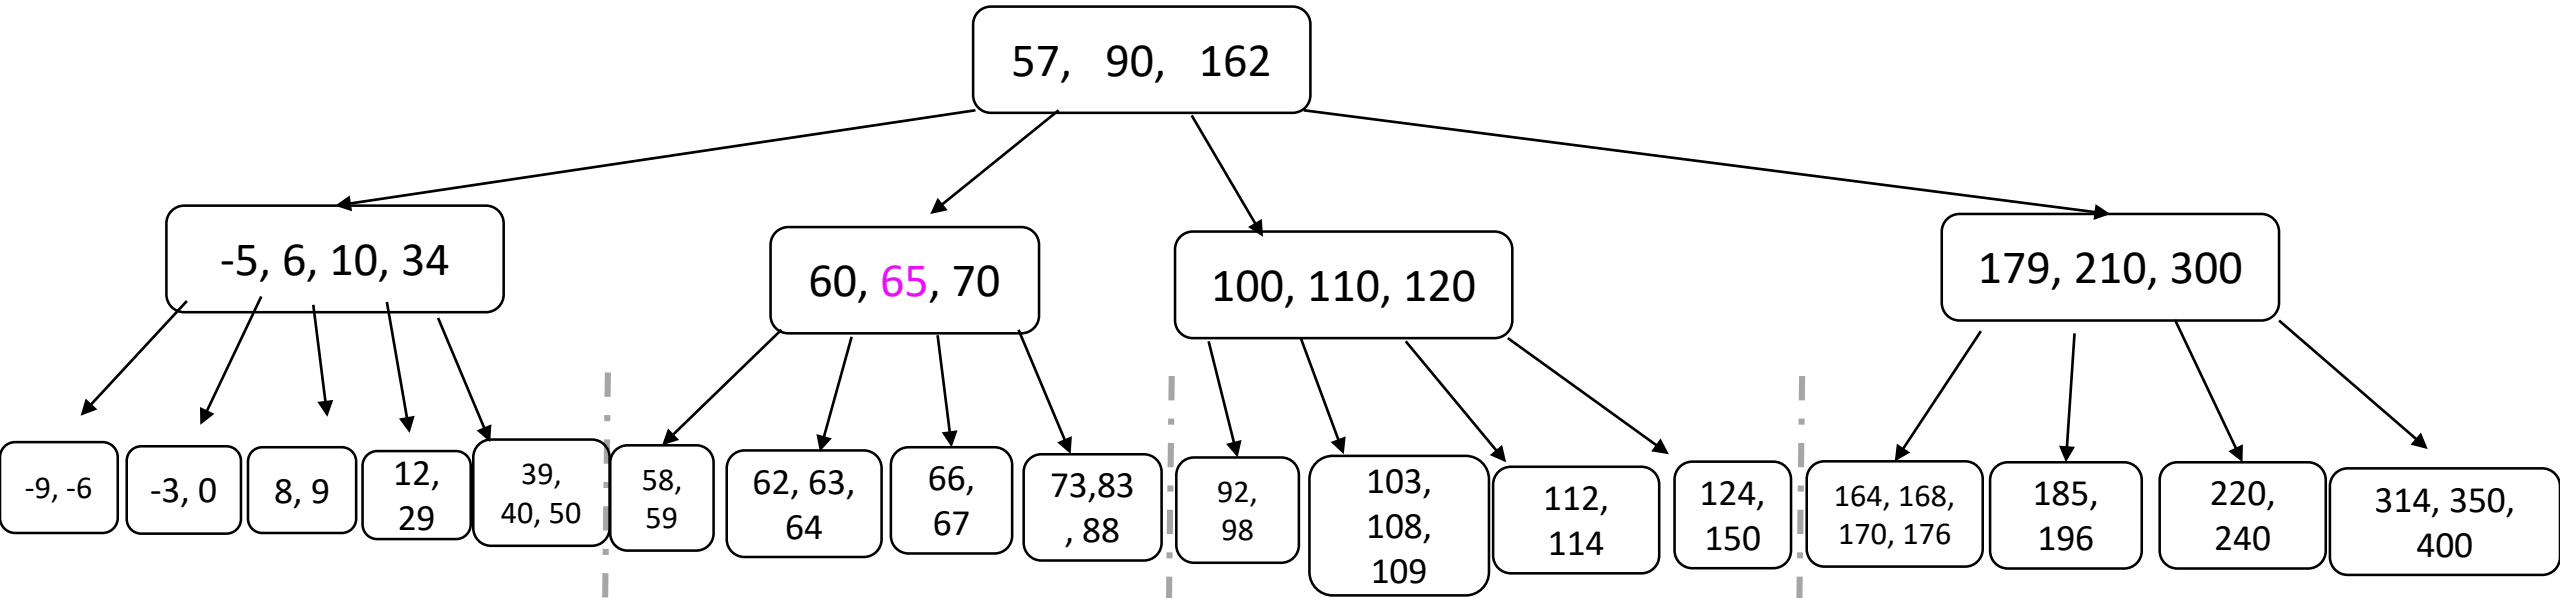
- Task: delete 62

Example #0: Leaf deletion, no underflows



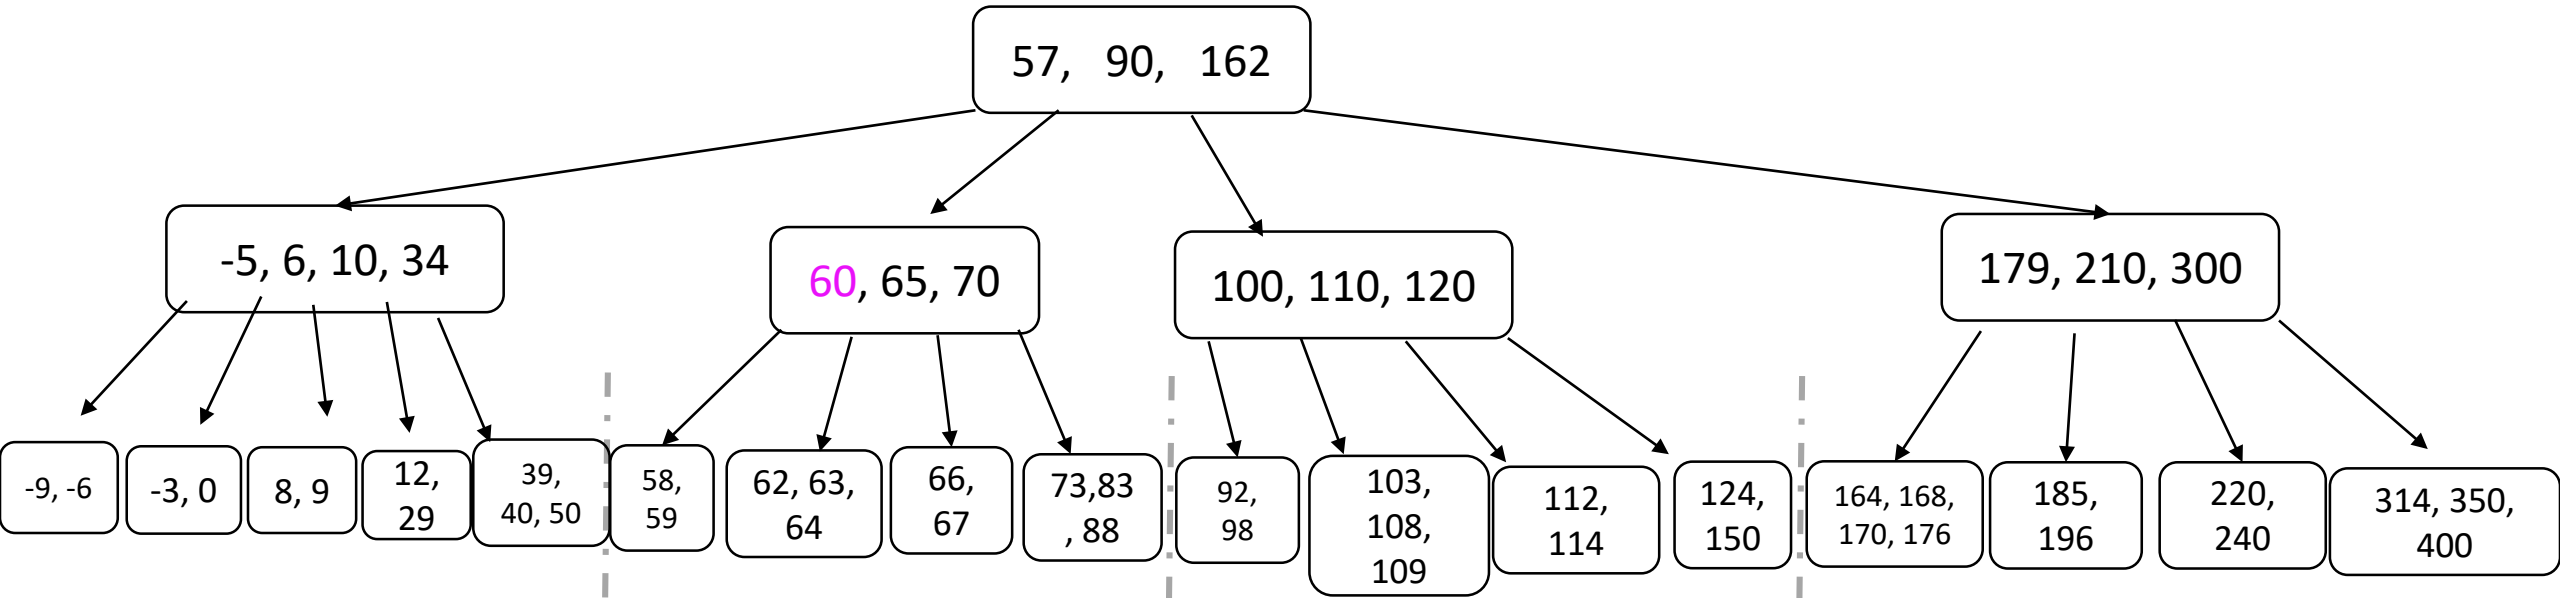
- Task: delete 62

Example #0: Leaf deletion, no underflows



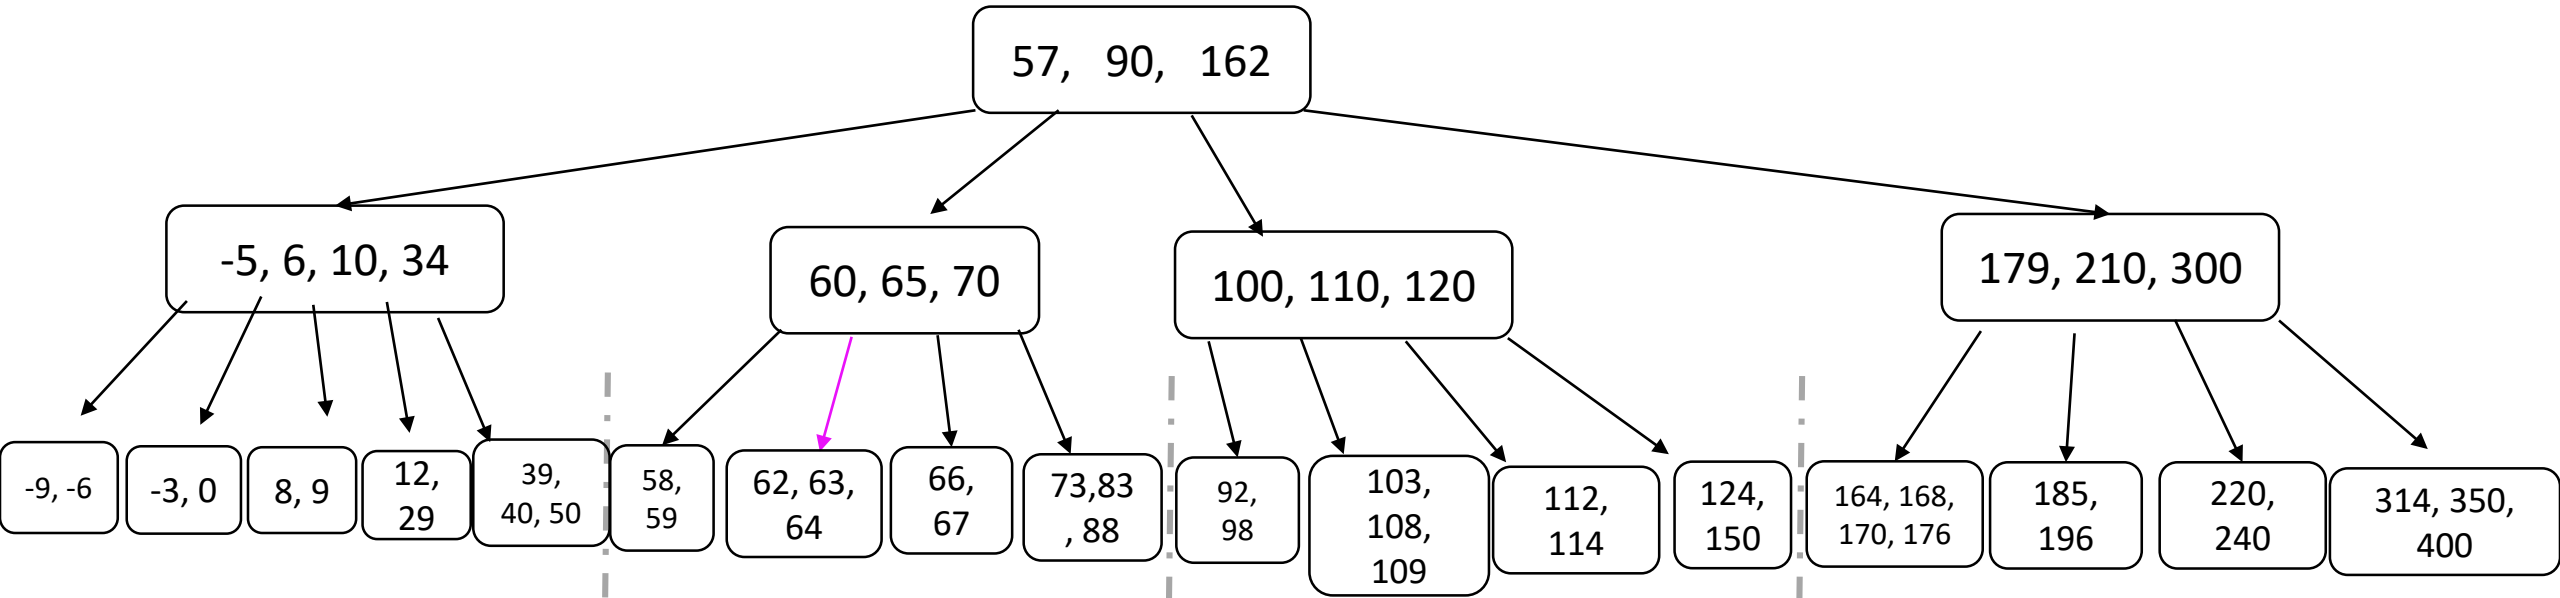
- Task: delete **62**

Example #0: Leaf deletion, no underflows



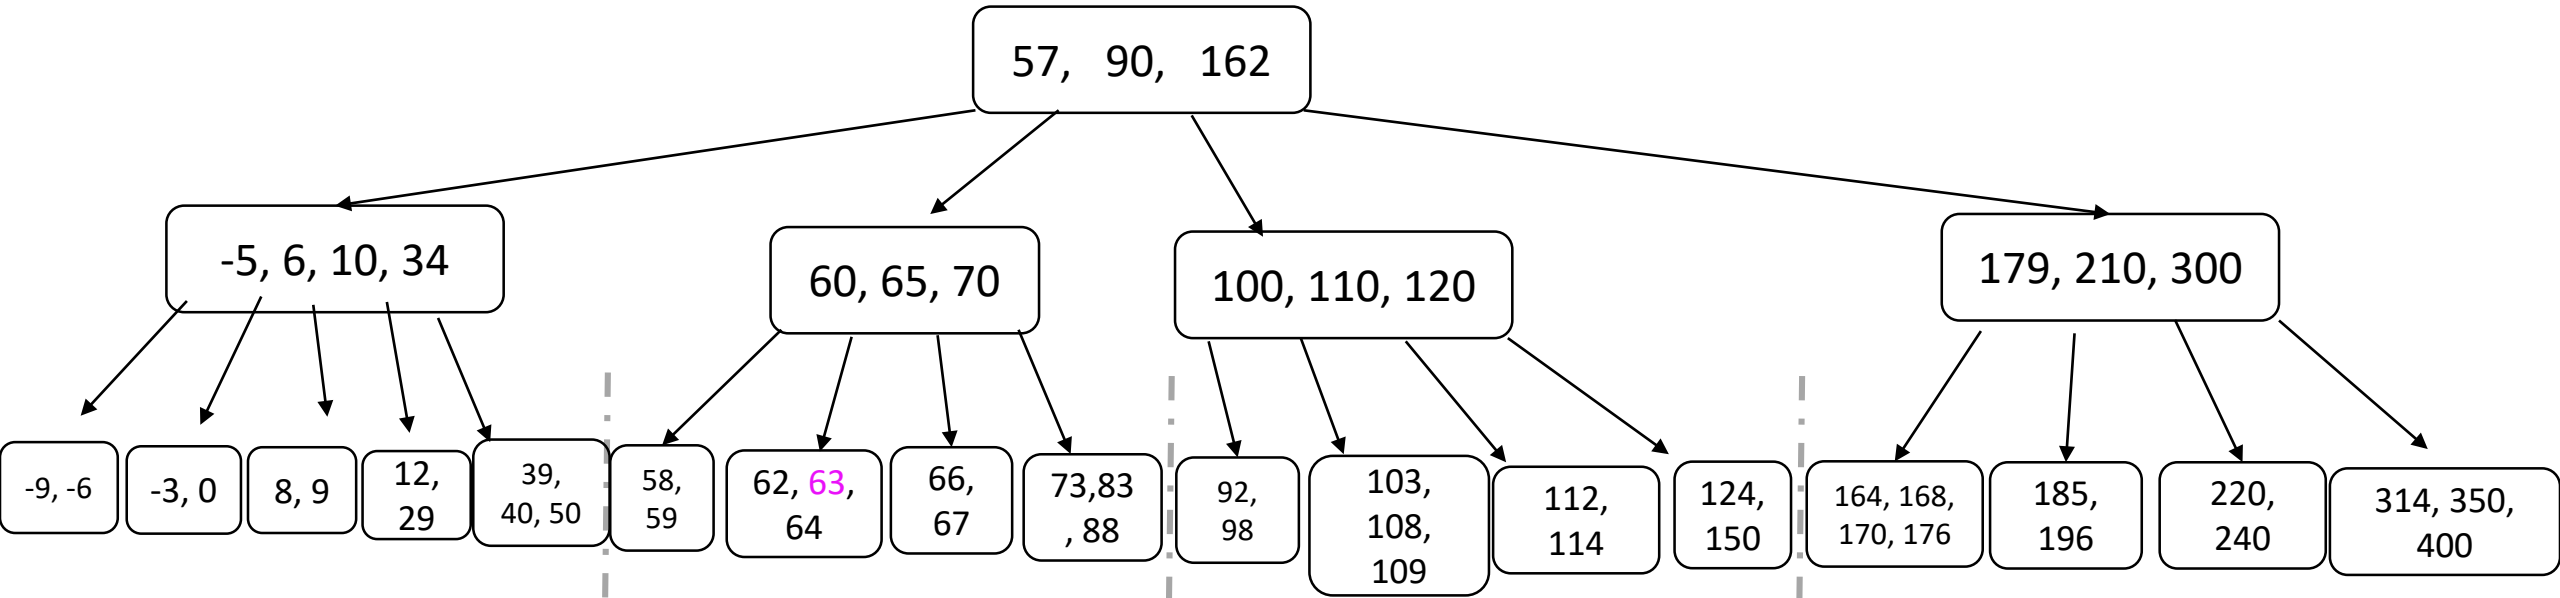
- Task: delete 62

Example #0: Leaf deletion, no underflows



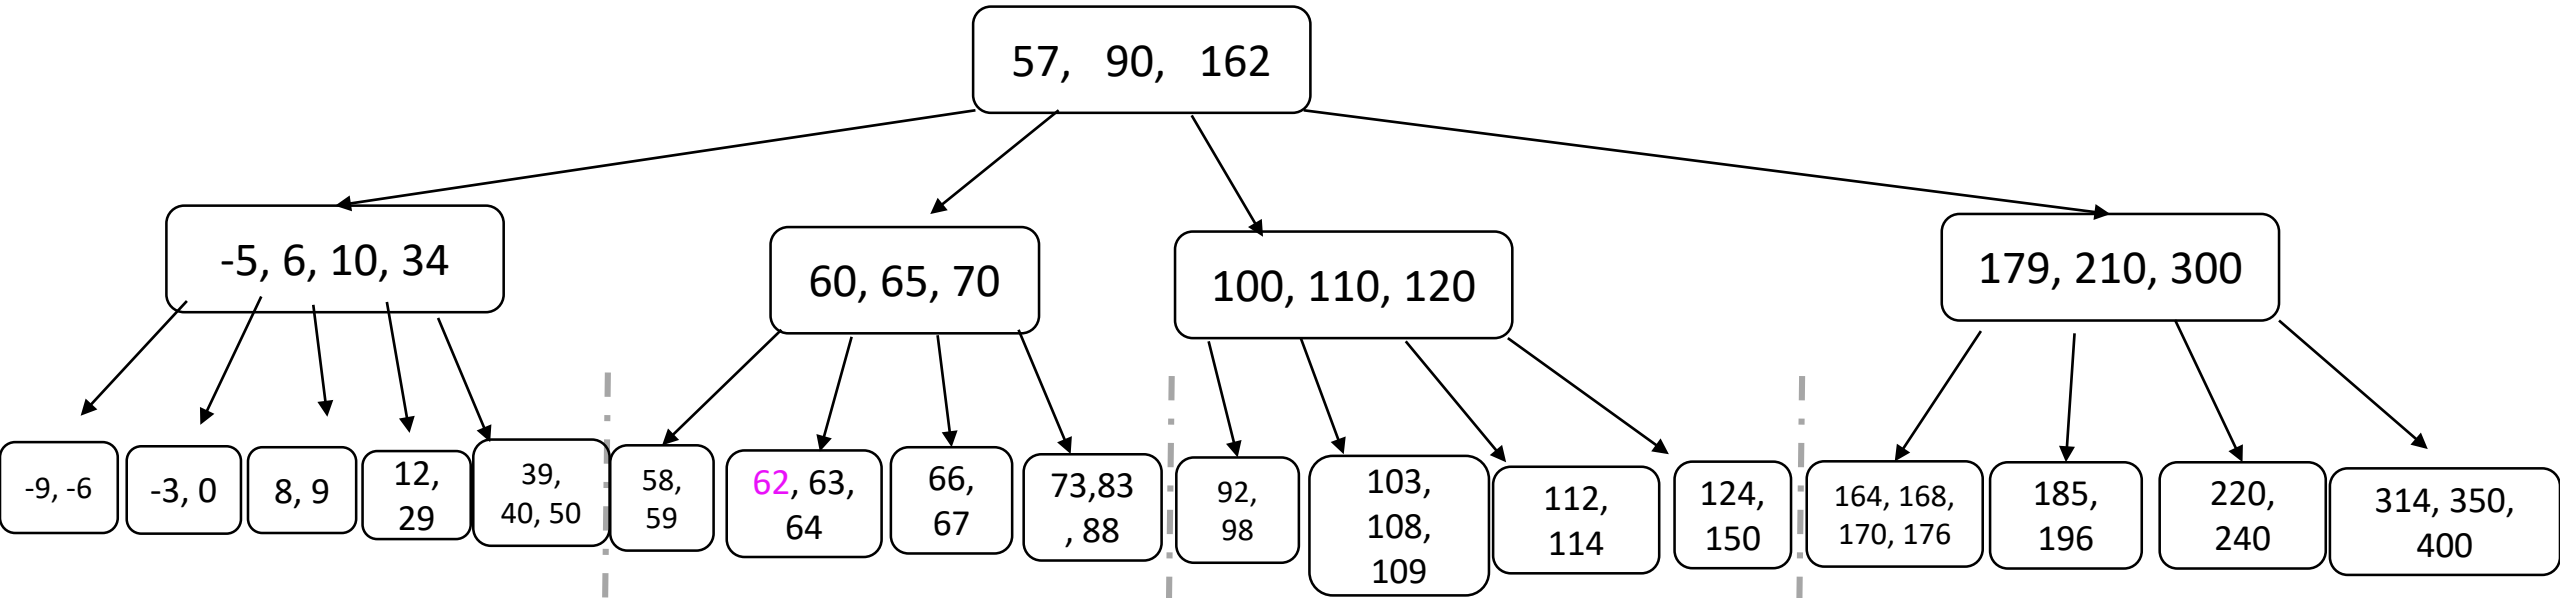
- Task: delete 62

Example #0: Leaf deletion, no underflows



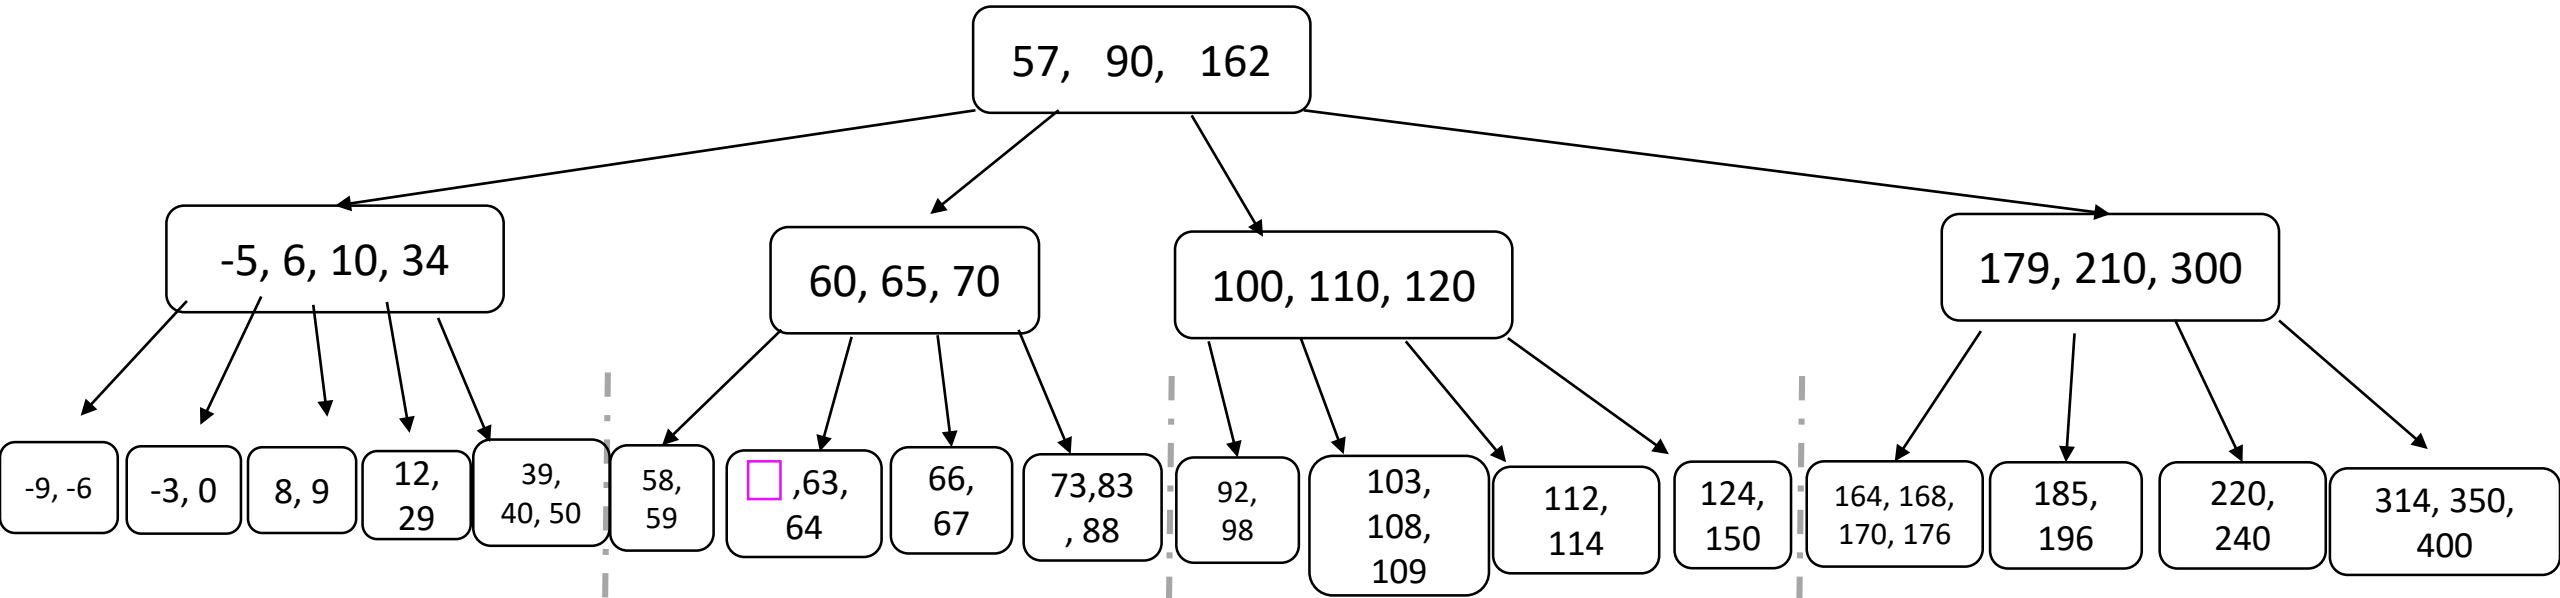
- Task: delete 62

Example #0: Leaf deletion, no underflows



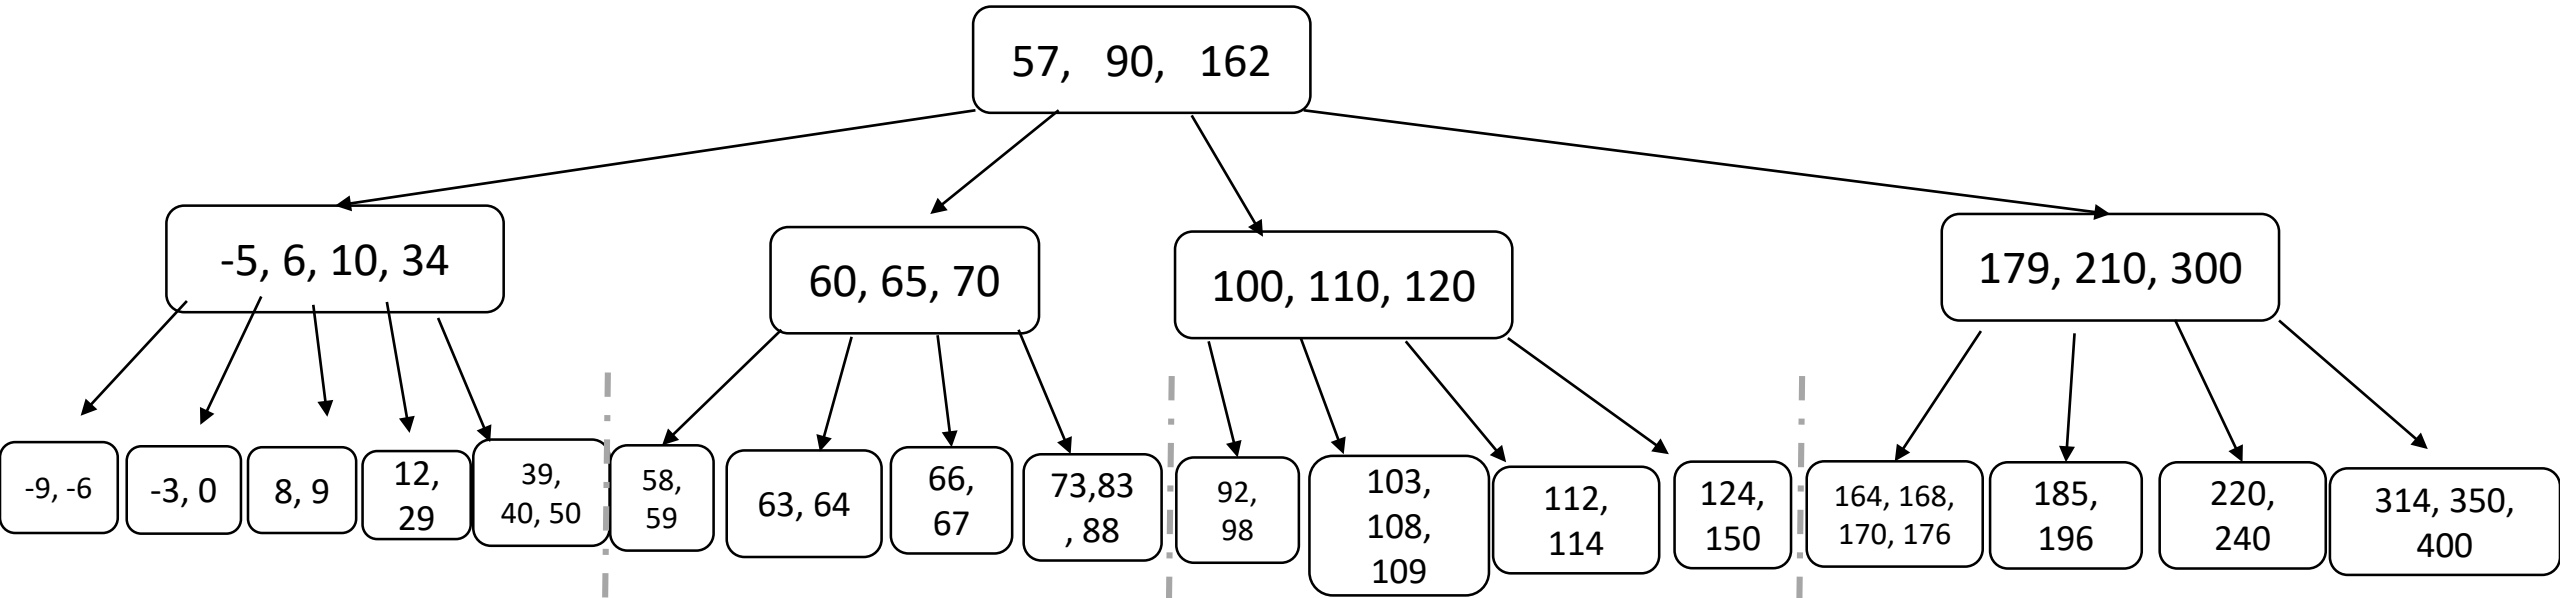
- Task: delete 62

Example #0: Leaf deletion, no underflows



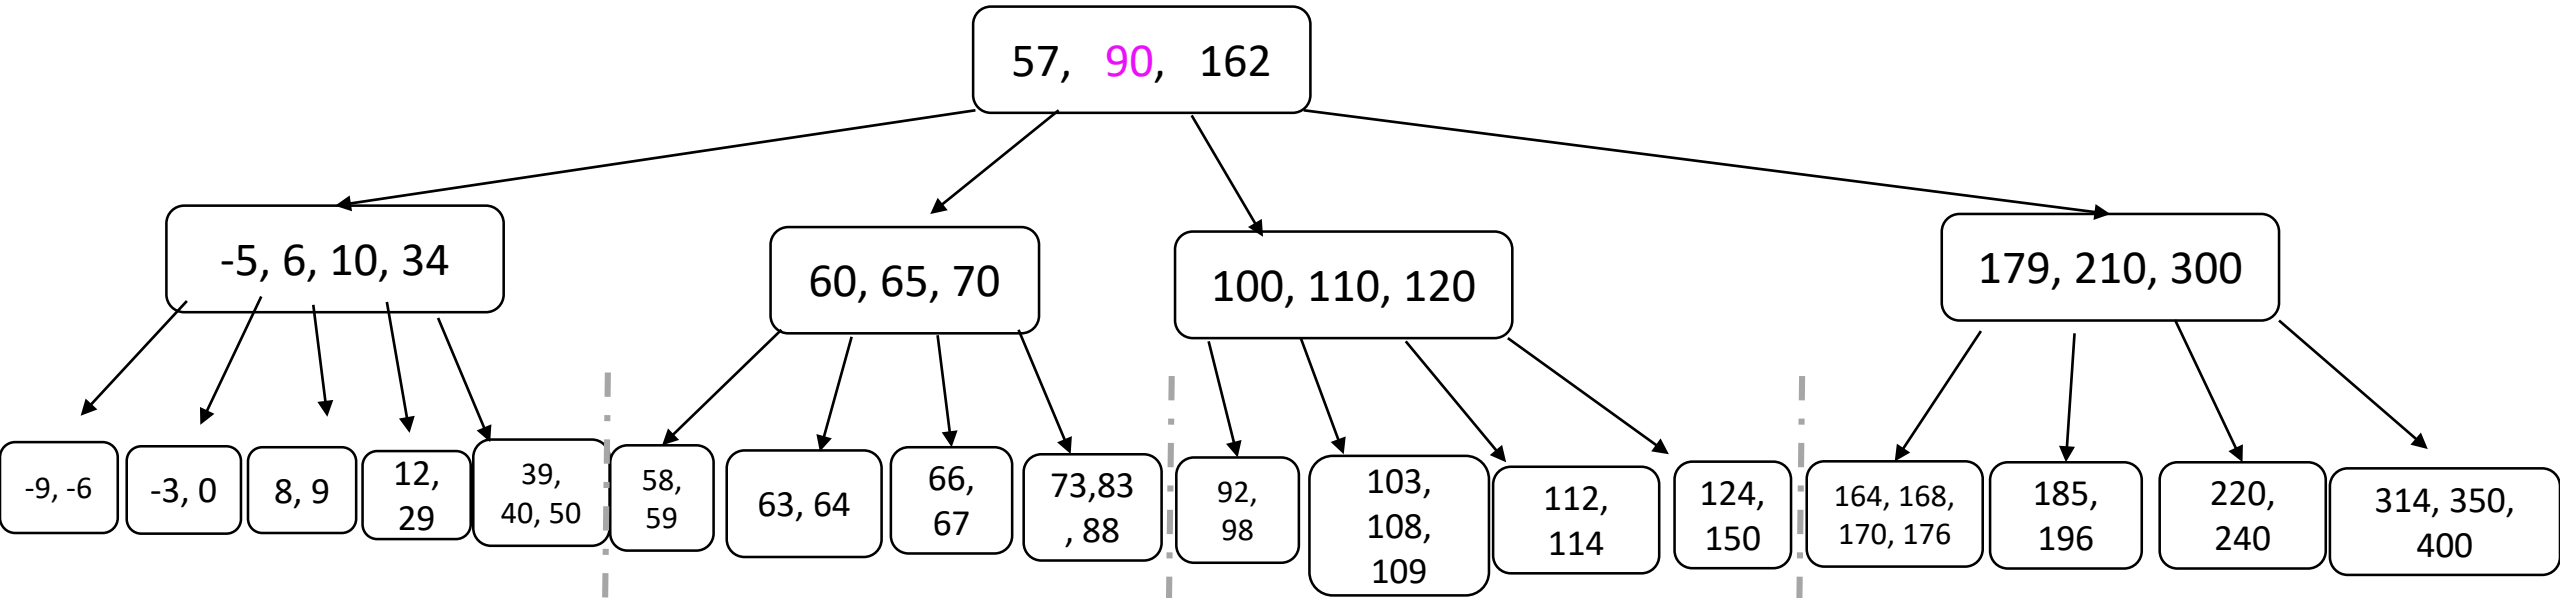
- Task: delete 62

Example #1: Inner node deletion, no underflows



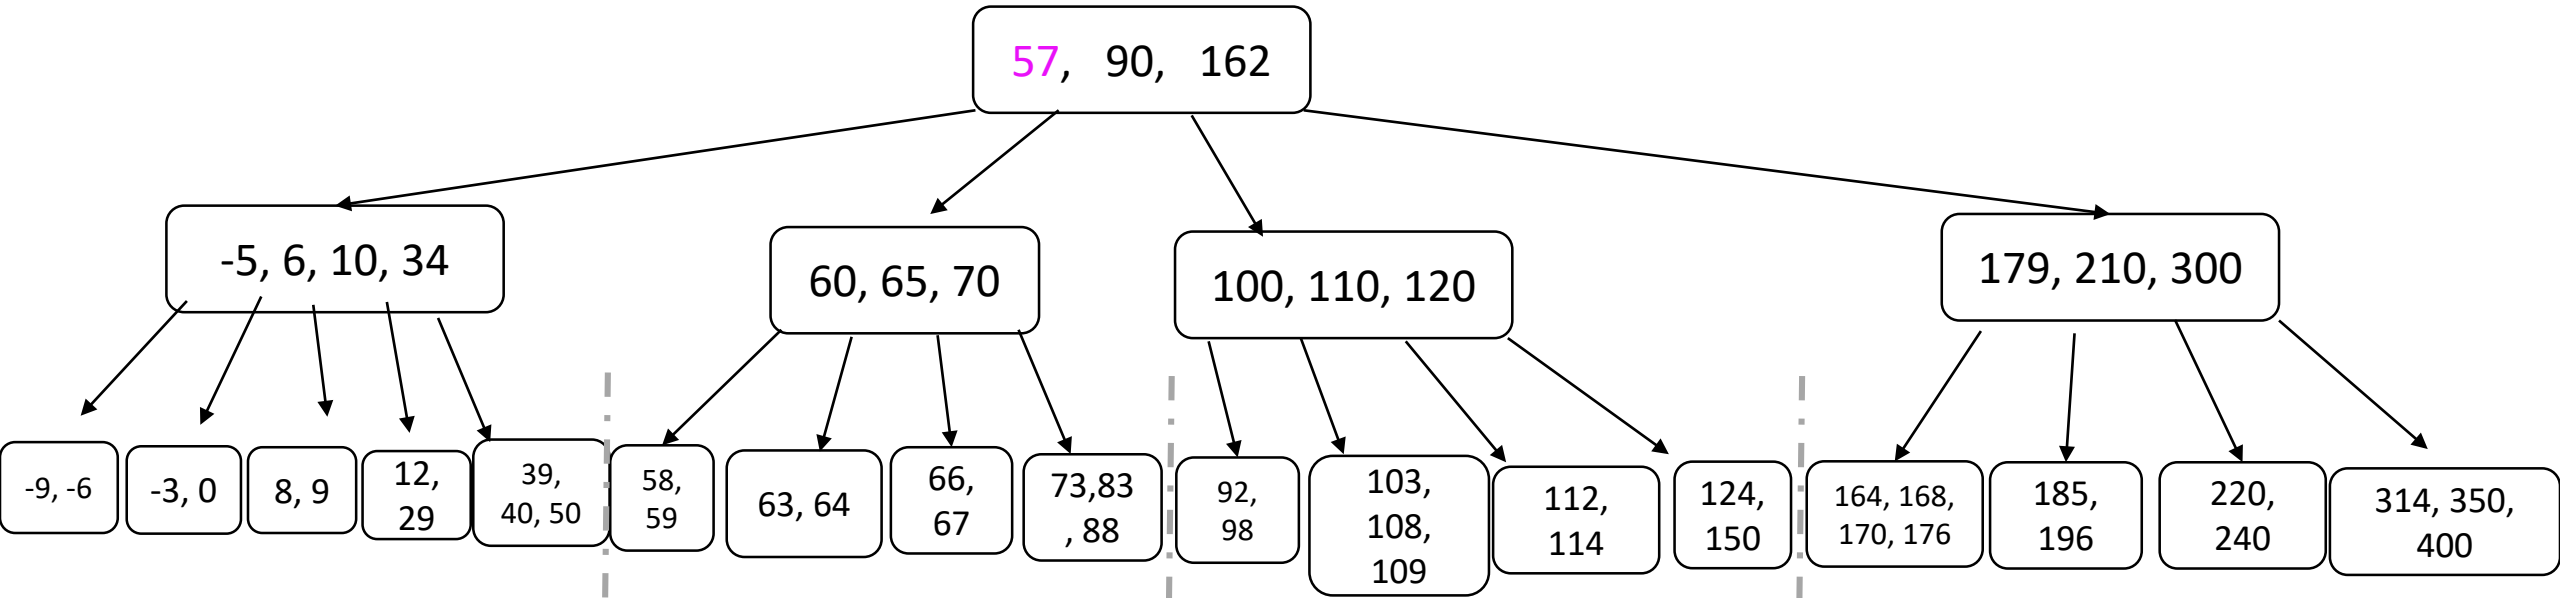
- Task: delete 34

Example #1: Inner node deletion, no underflows



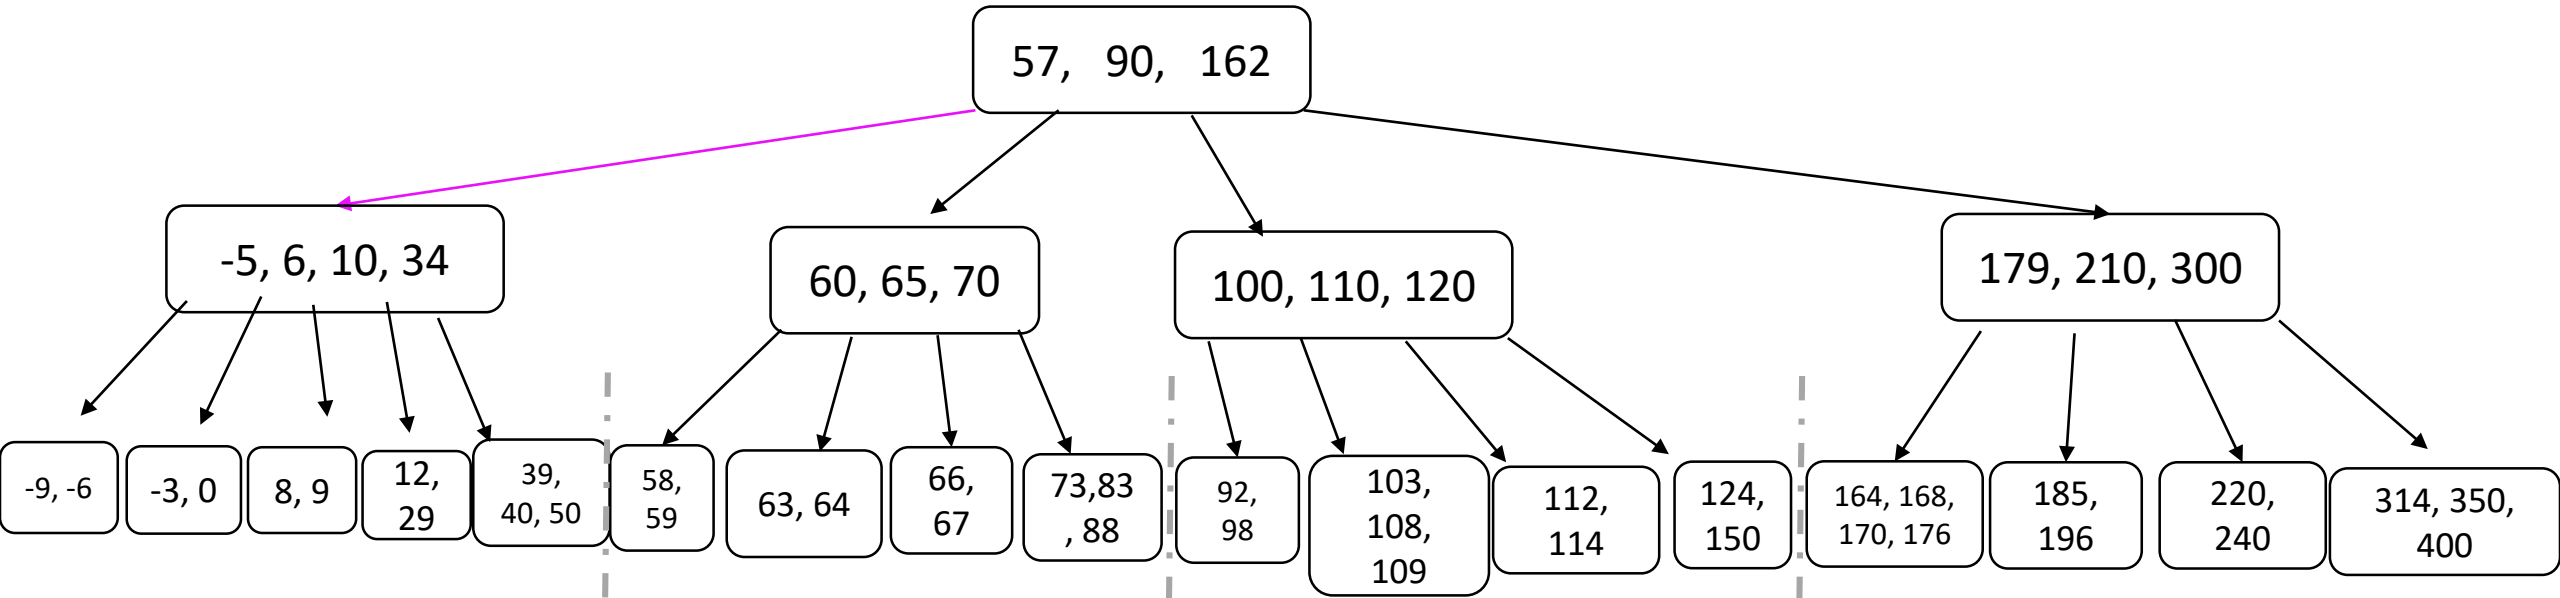
- Task: delete 34

Example #1: Inner node deletion, no underflows



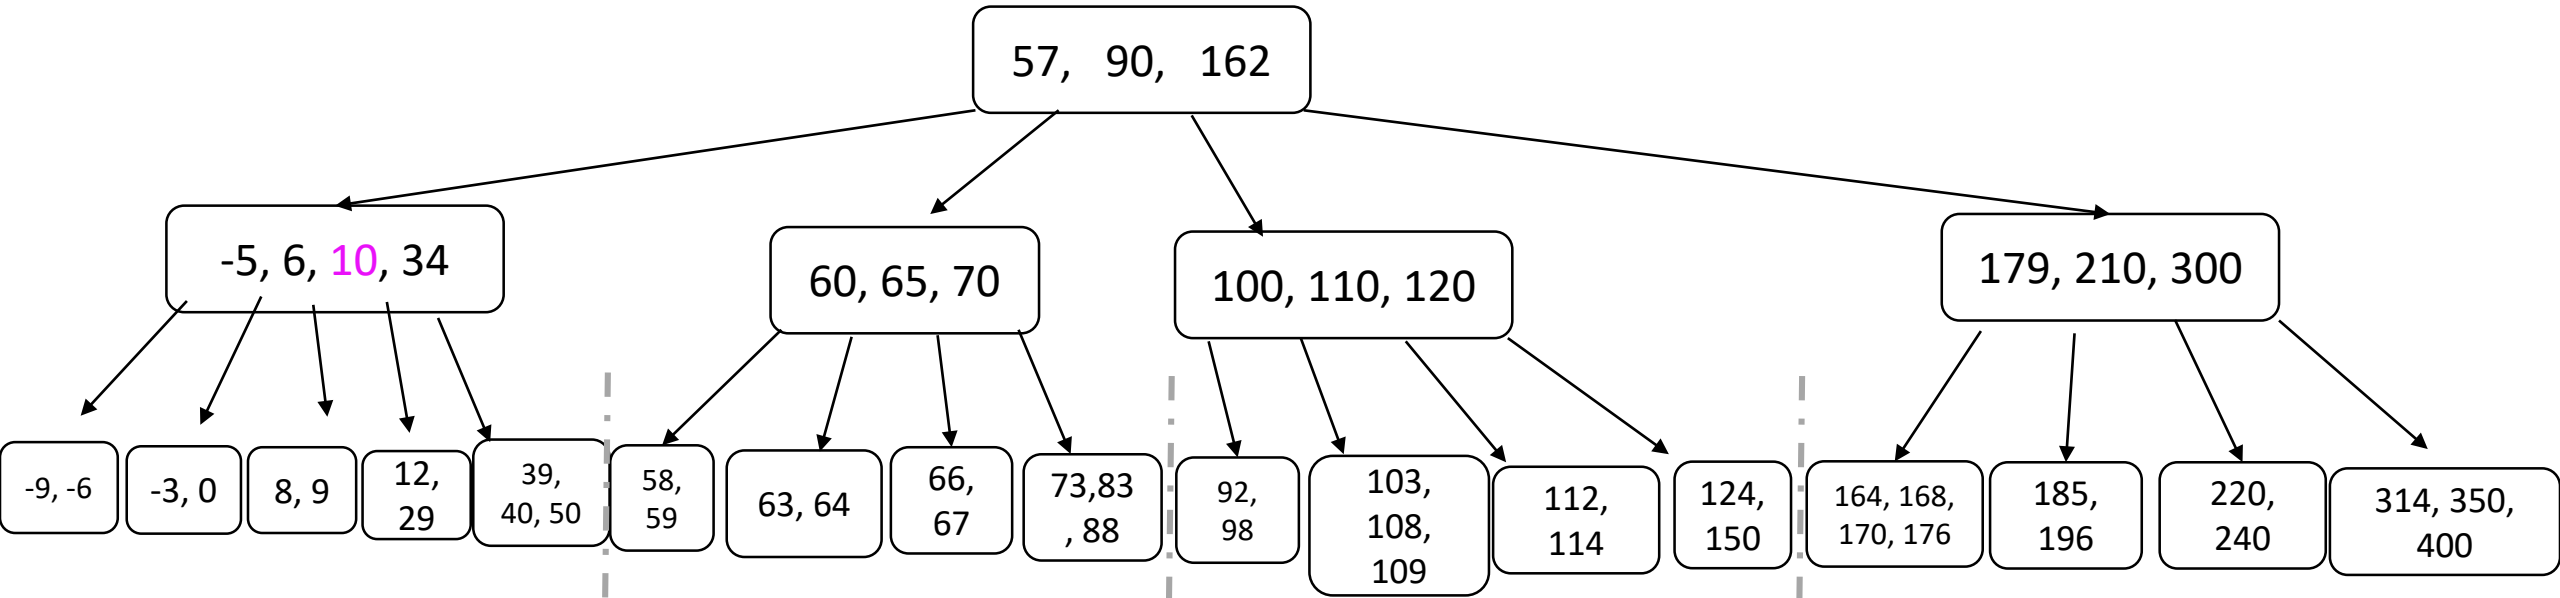
- Task: delete 34

Example #1: Inner node deletion, no underflows



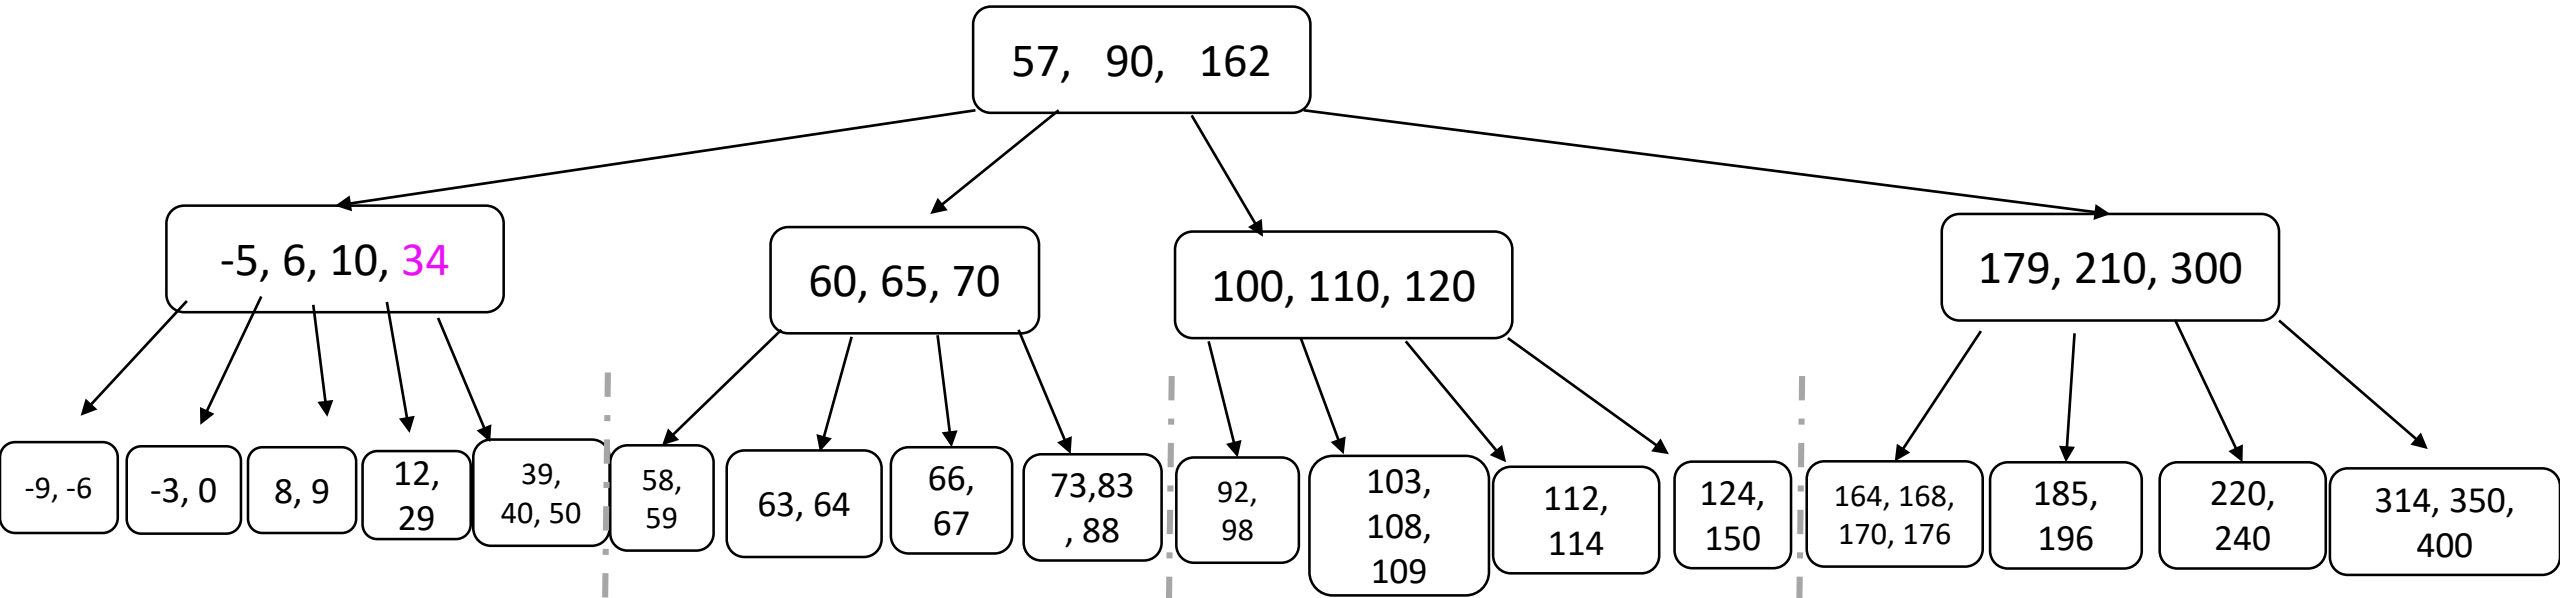
- Task: delete 34

Example #1: Inner node deletion, no underflows




- Task: delete 34

Example #1: Inner node deletion, no underflows



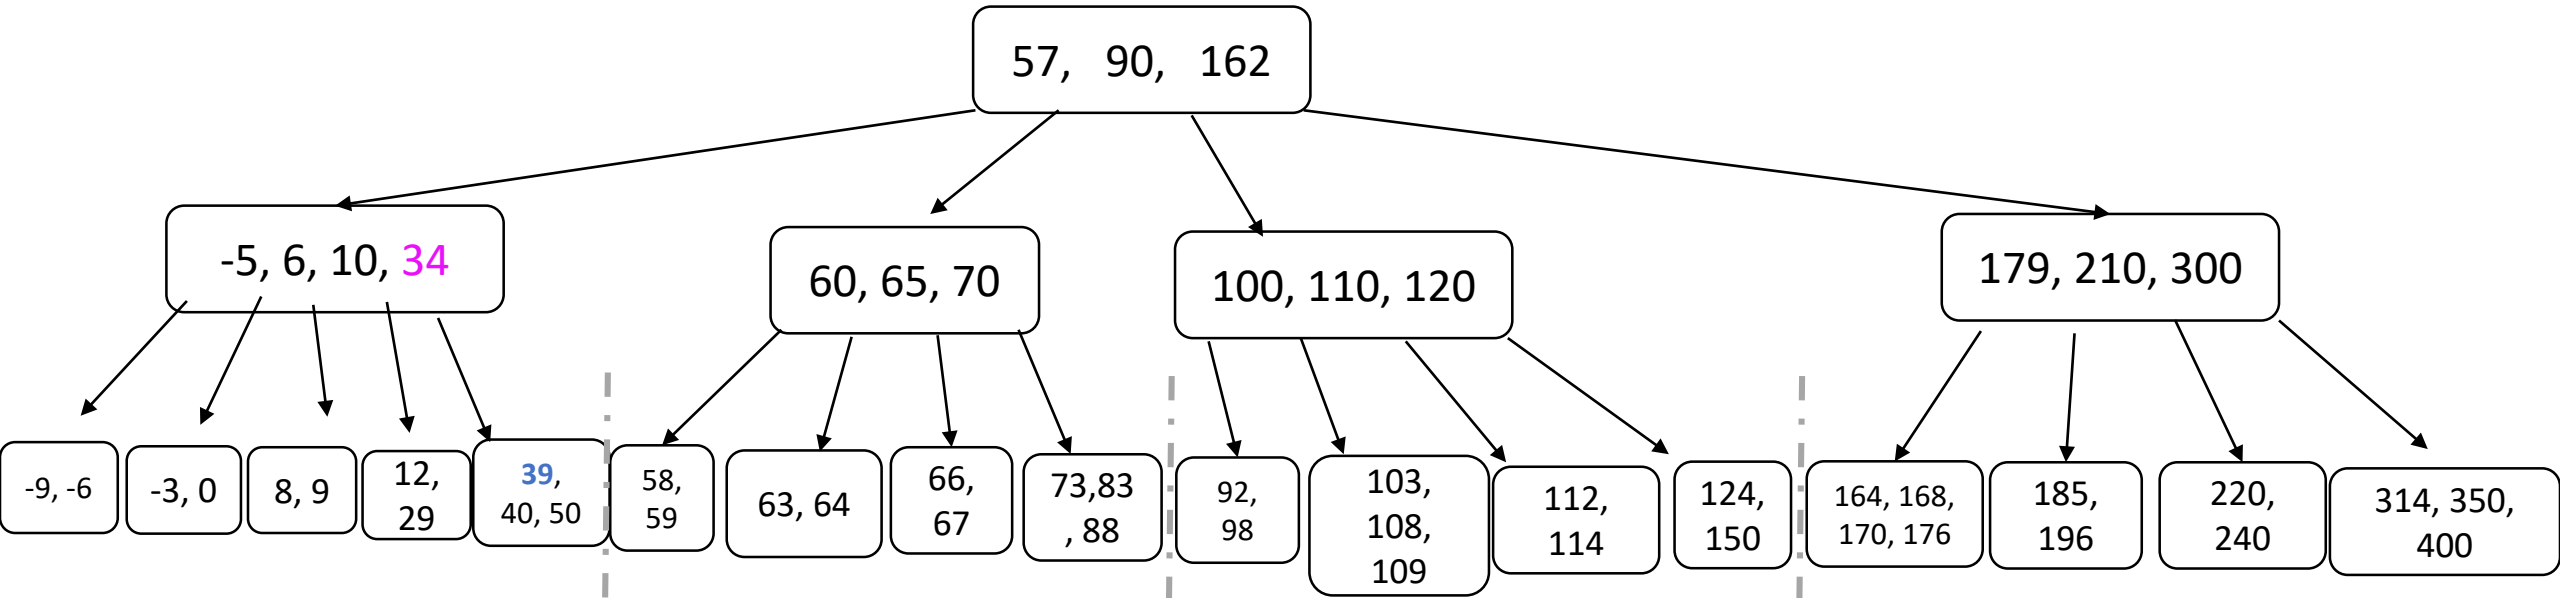
- Task: delete 34

Example #1: Inner node deletion, no underflows



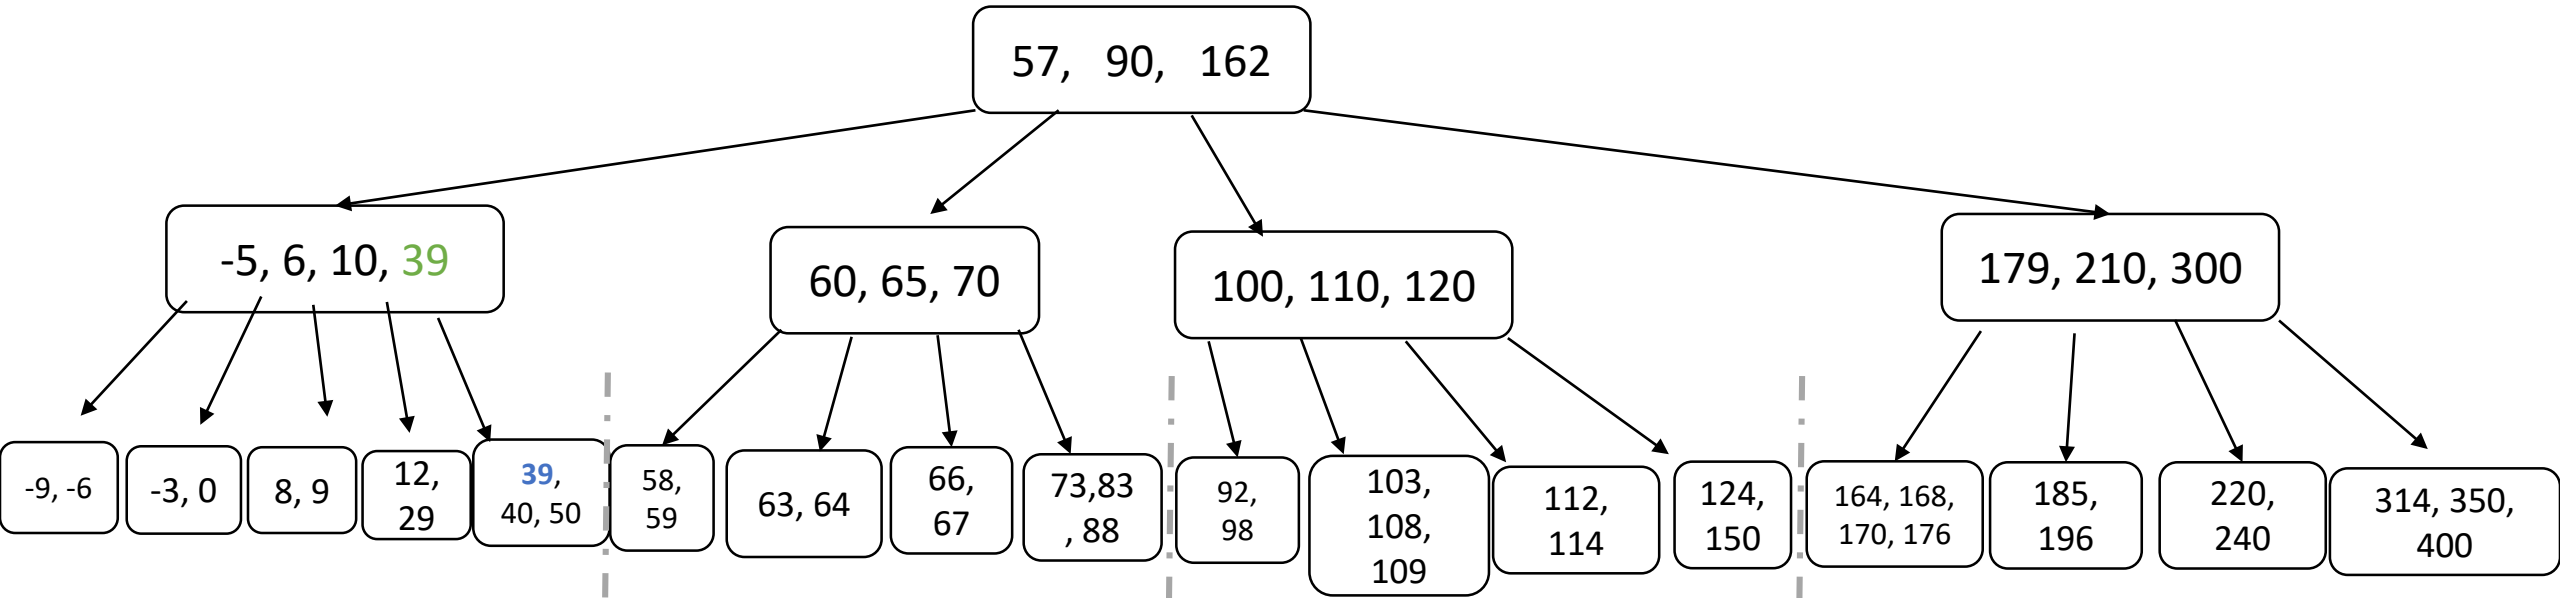
- Task: delete 34
 - Finding inorder successor...

Example #1: Inner node deletion, no underflows



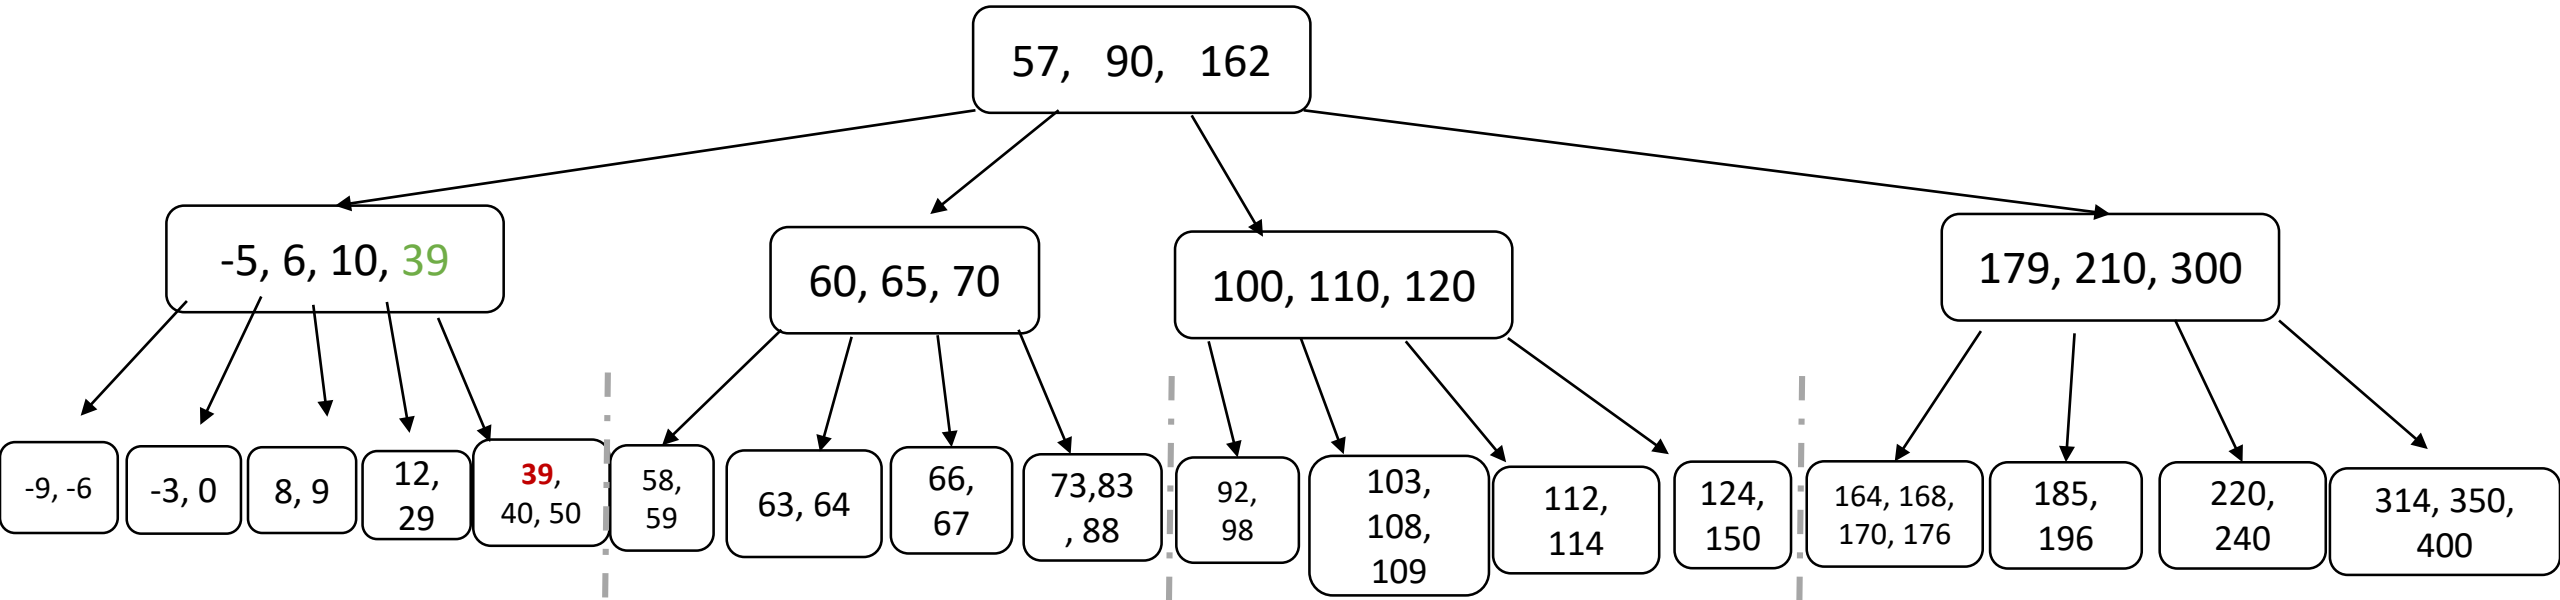
- Task: delete 34
 - Finding inorder successor...

Example #1: Inner node deletion, no underflows




- Task: delete 34
 - Finding inorder successor...
 - Copy key

Example #1: Inner node deletion, no underflows




- Task: delete 34
 - Finding inorder successor...
 - Copy key
 - Recursively delete 39

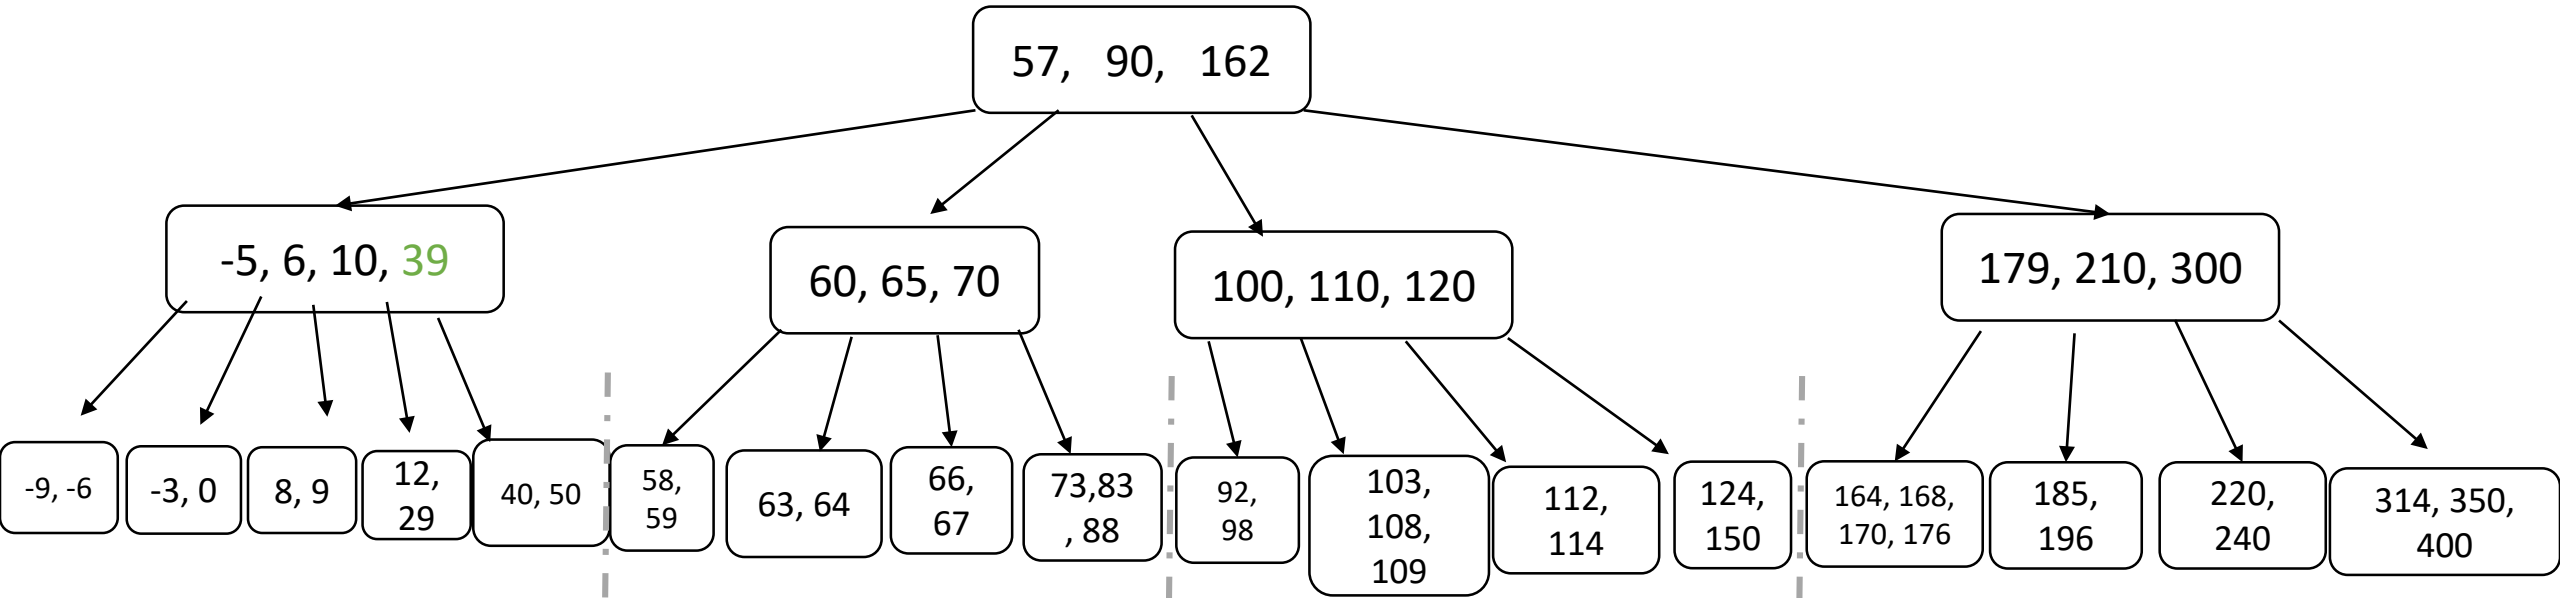
Example #1: Inner node deletion, no underflows



- Task: delete 34
 - Finding inorder successor...
 - Copy key
 - Recursively delete 39
 - No underflow, no problem!

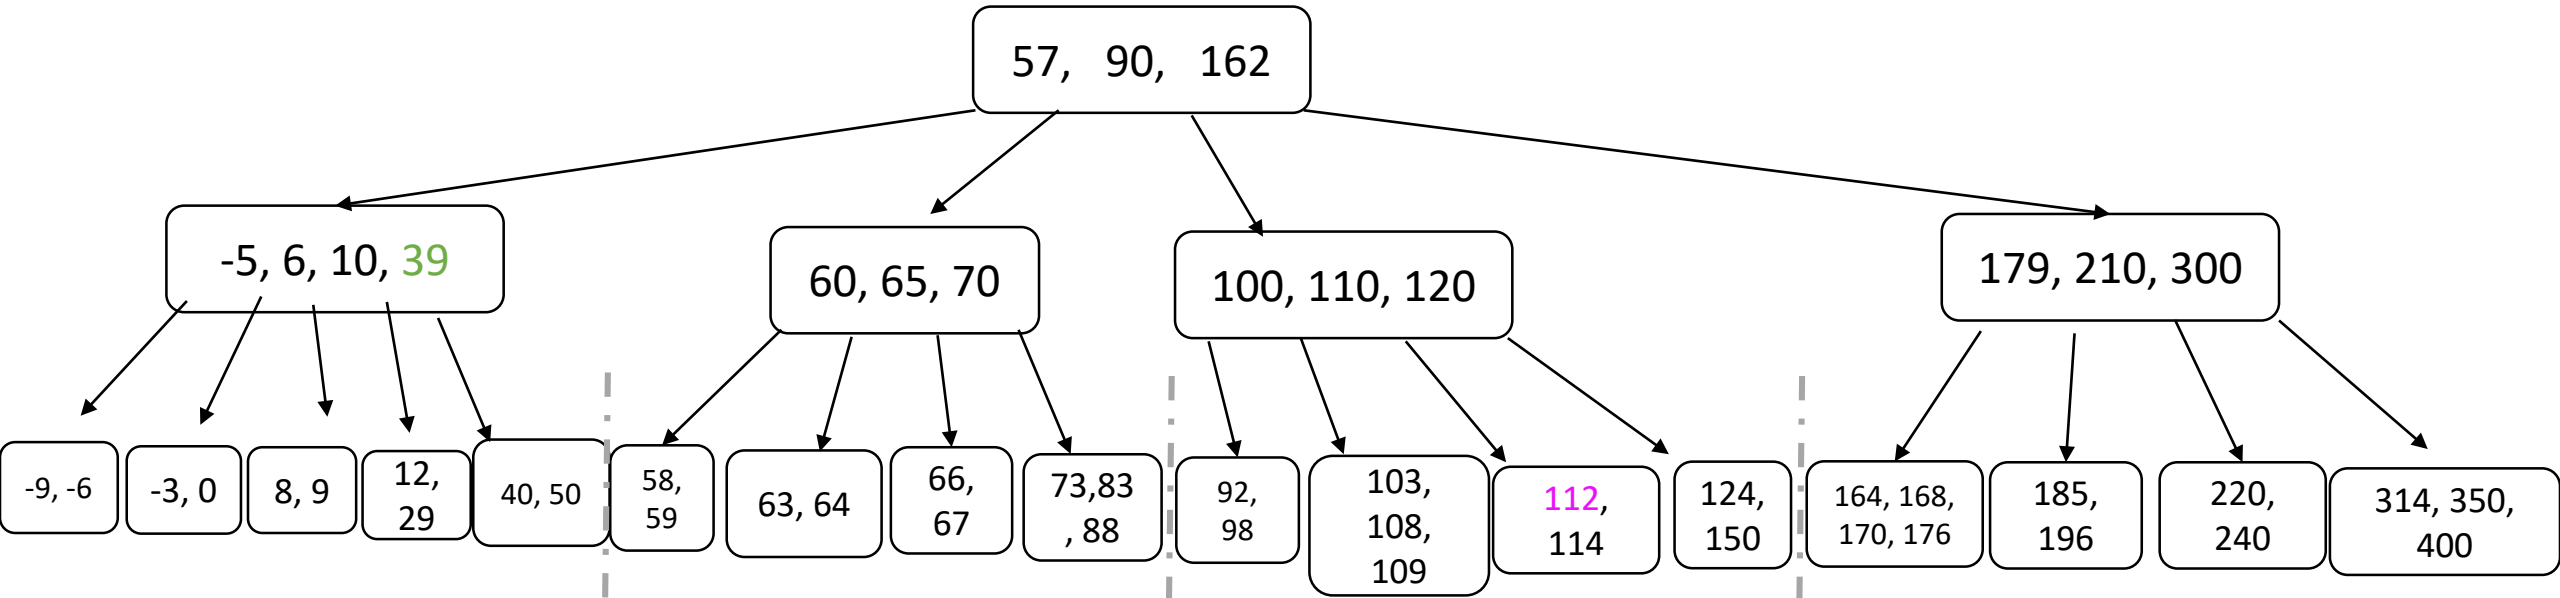


Example #2: Key rotation



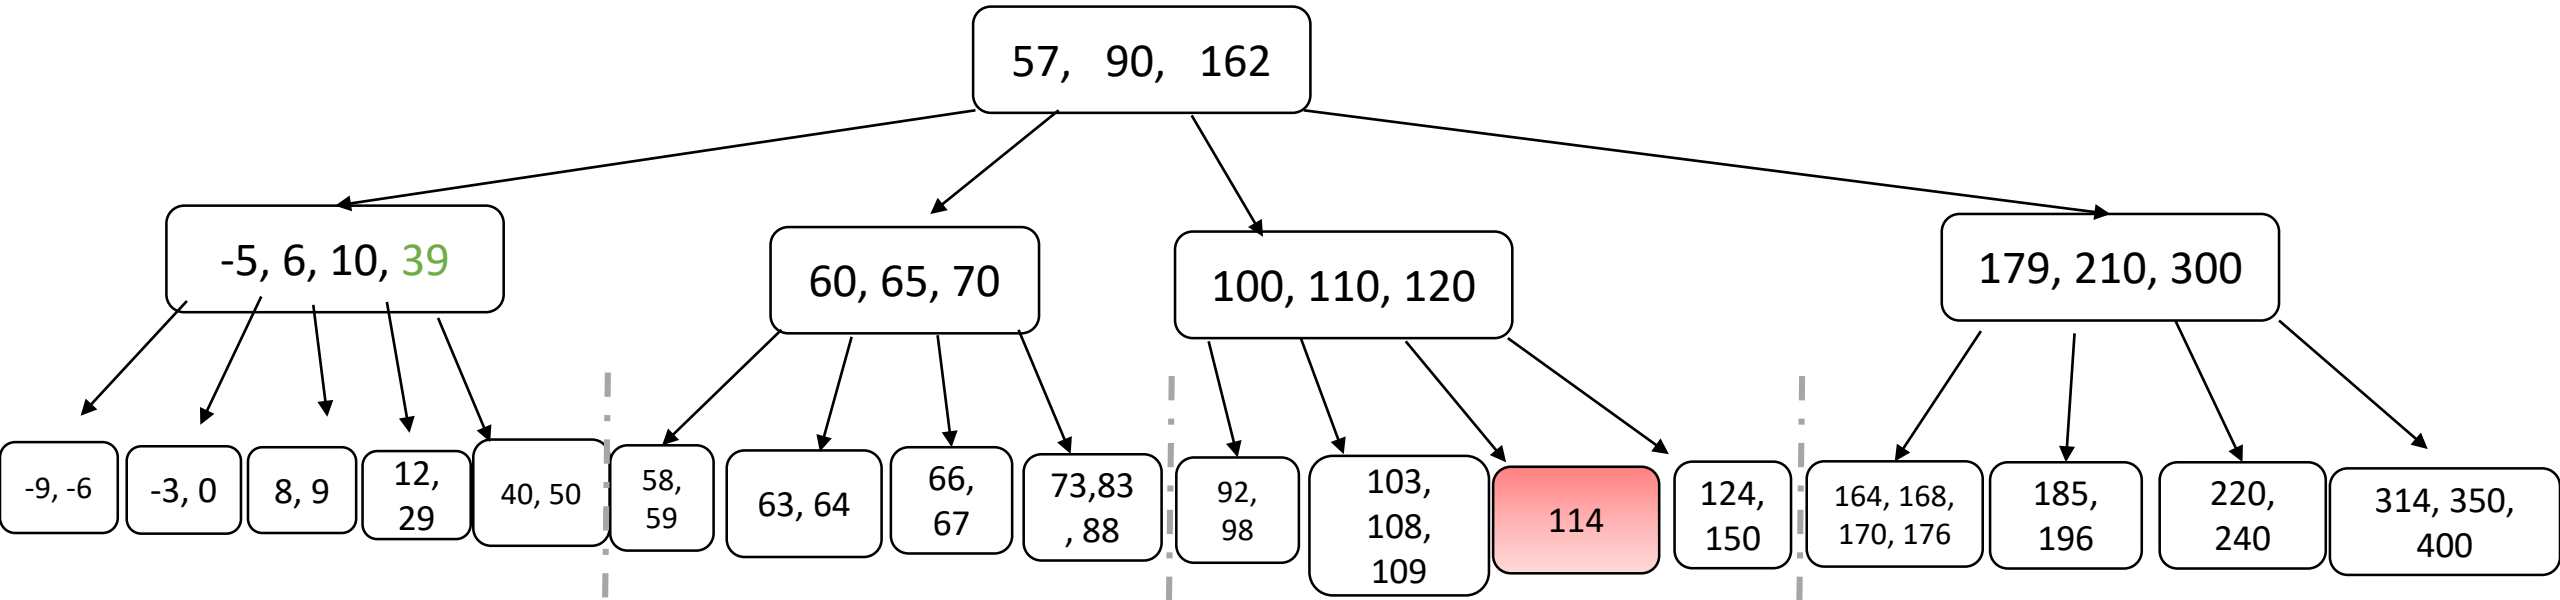
- Task: delete 112

Example #2: Key rotation



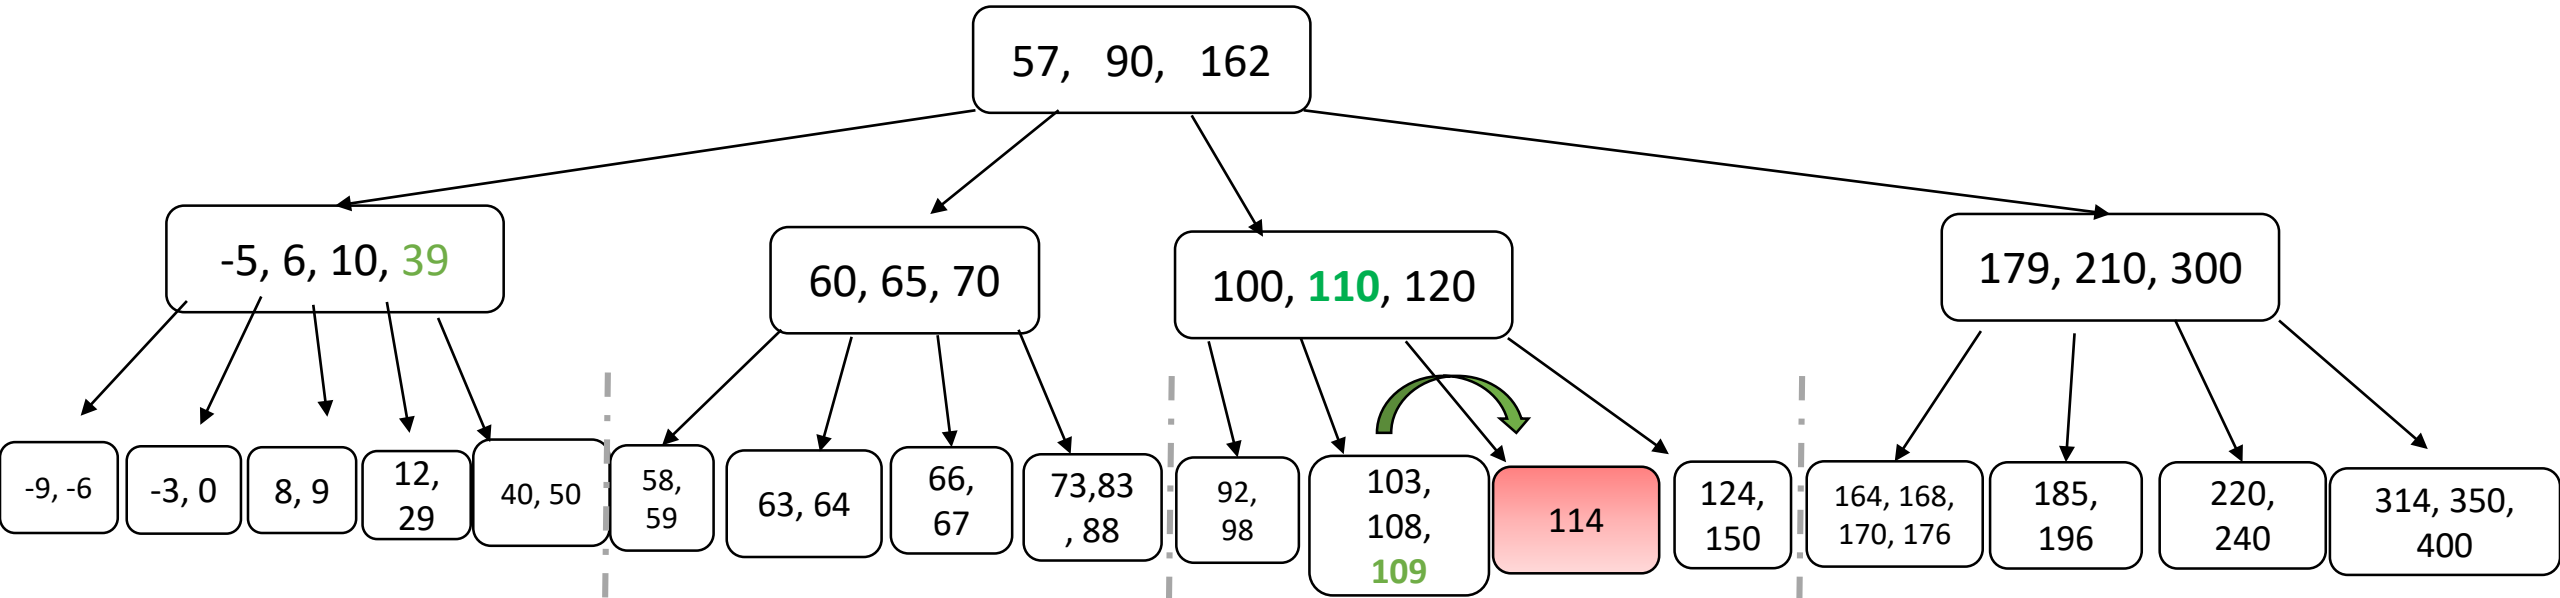
- Task: delete 112

Example #2: Key rotation




- Task: delete 112
- Underflow 😞

Example #2: Key rotation



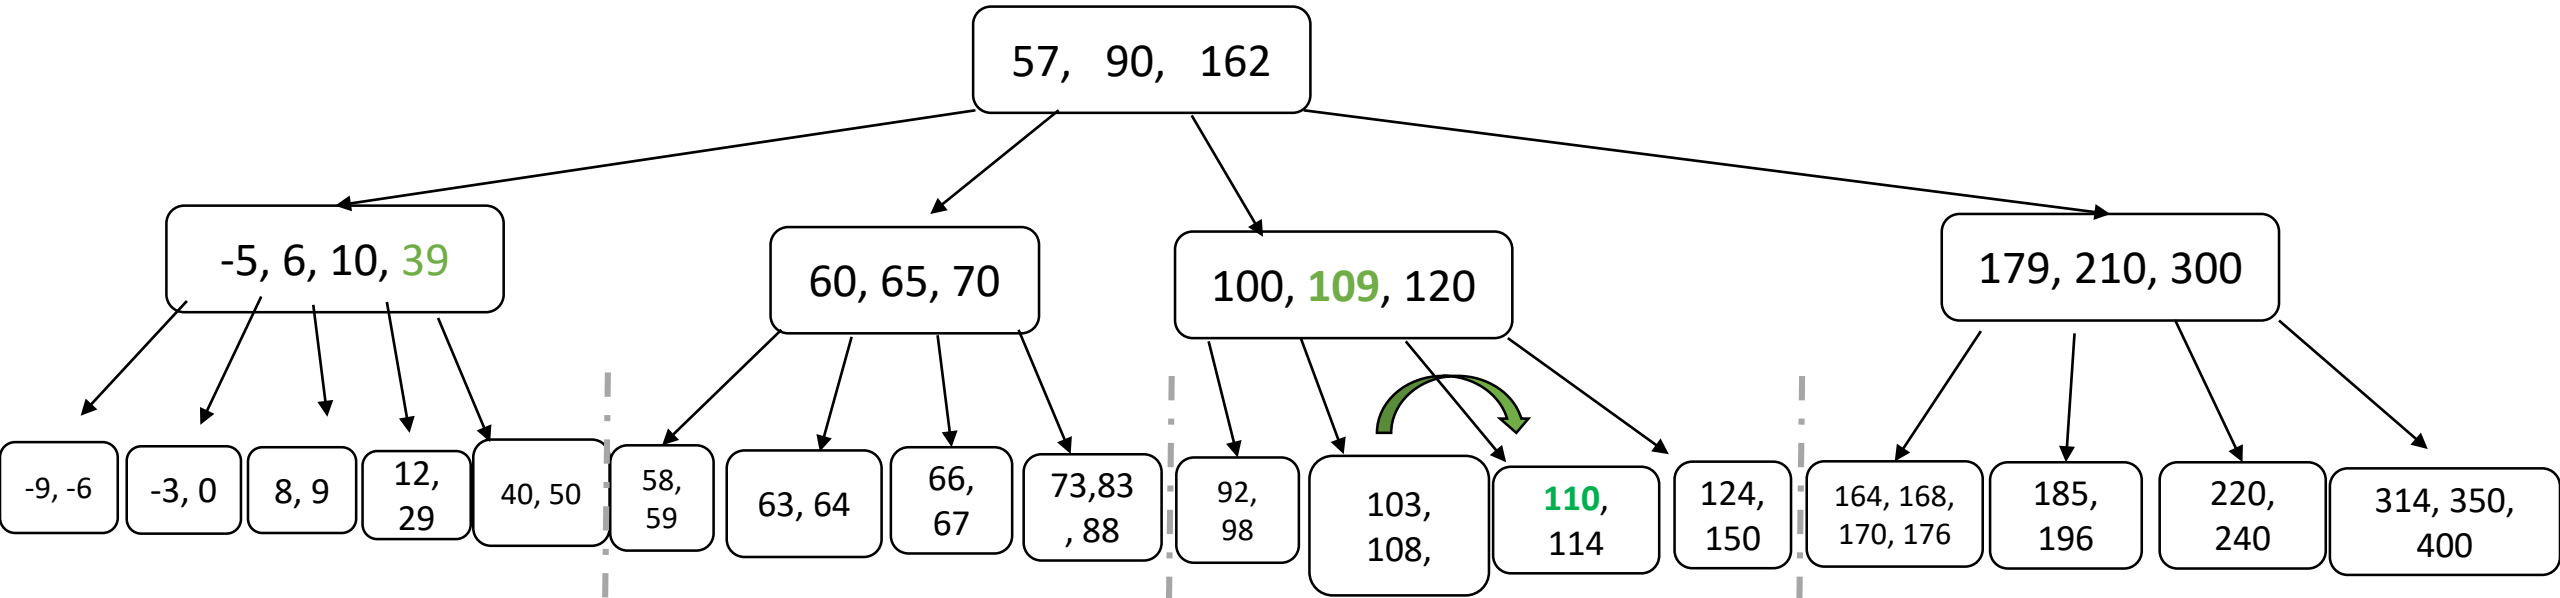
- Task: delete 112
- Underflow 😞
- Solvable via key rotation 😊

Example #2: Key rotation



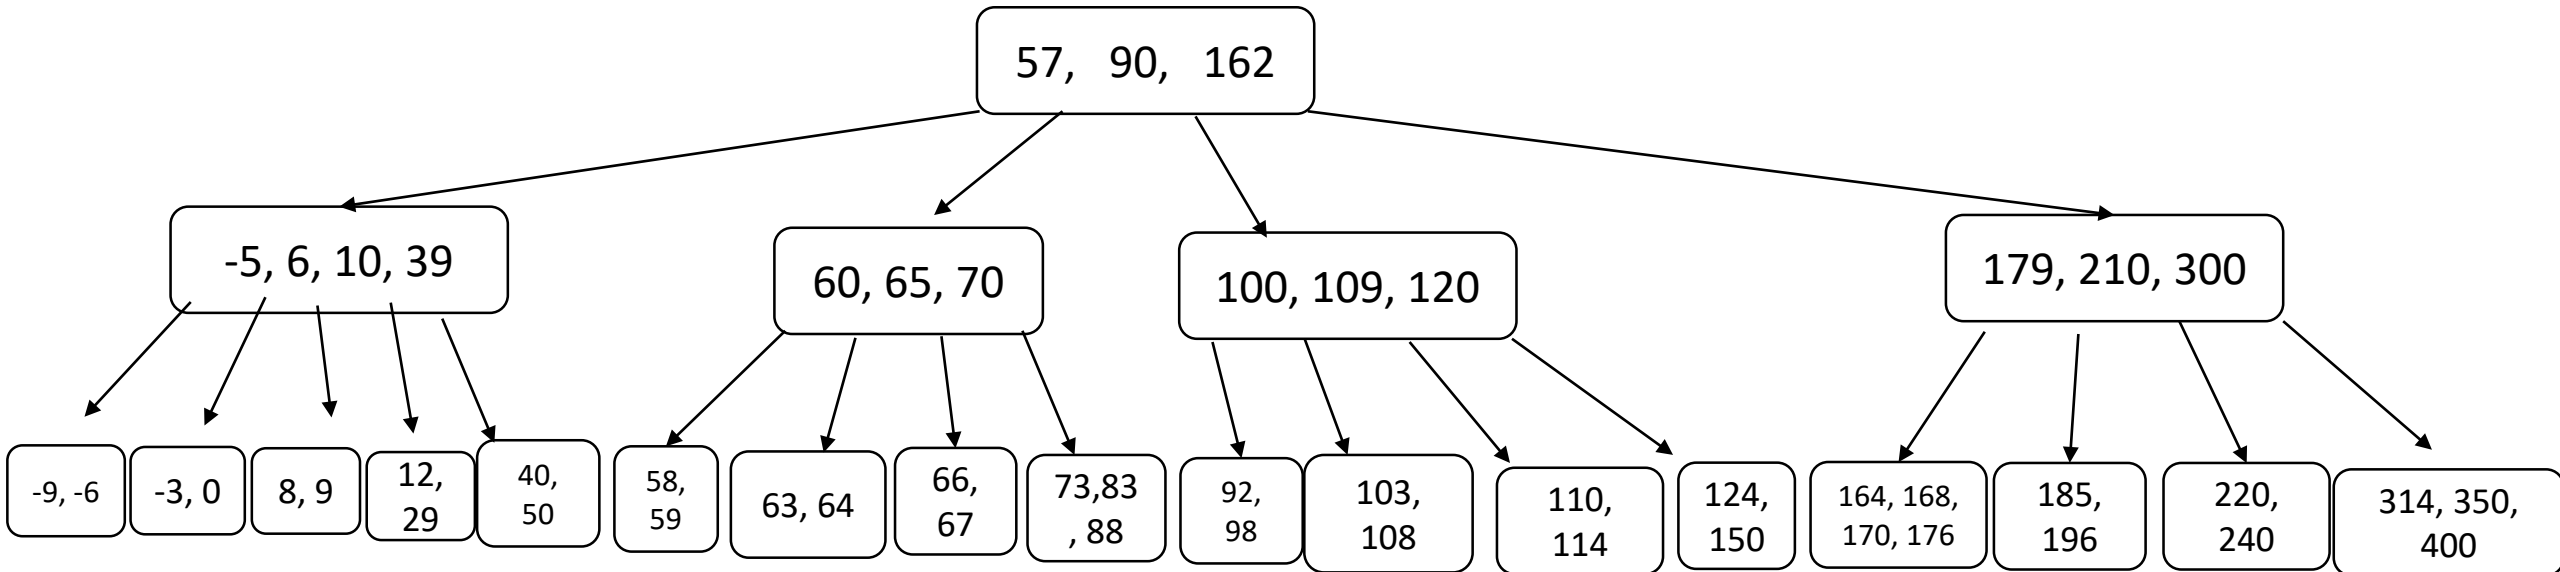
- Task: delete 112
- Underflow 😞
- Solvable via key rotation 😊

Example #2: Key rotation



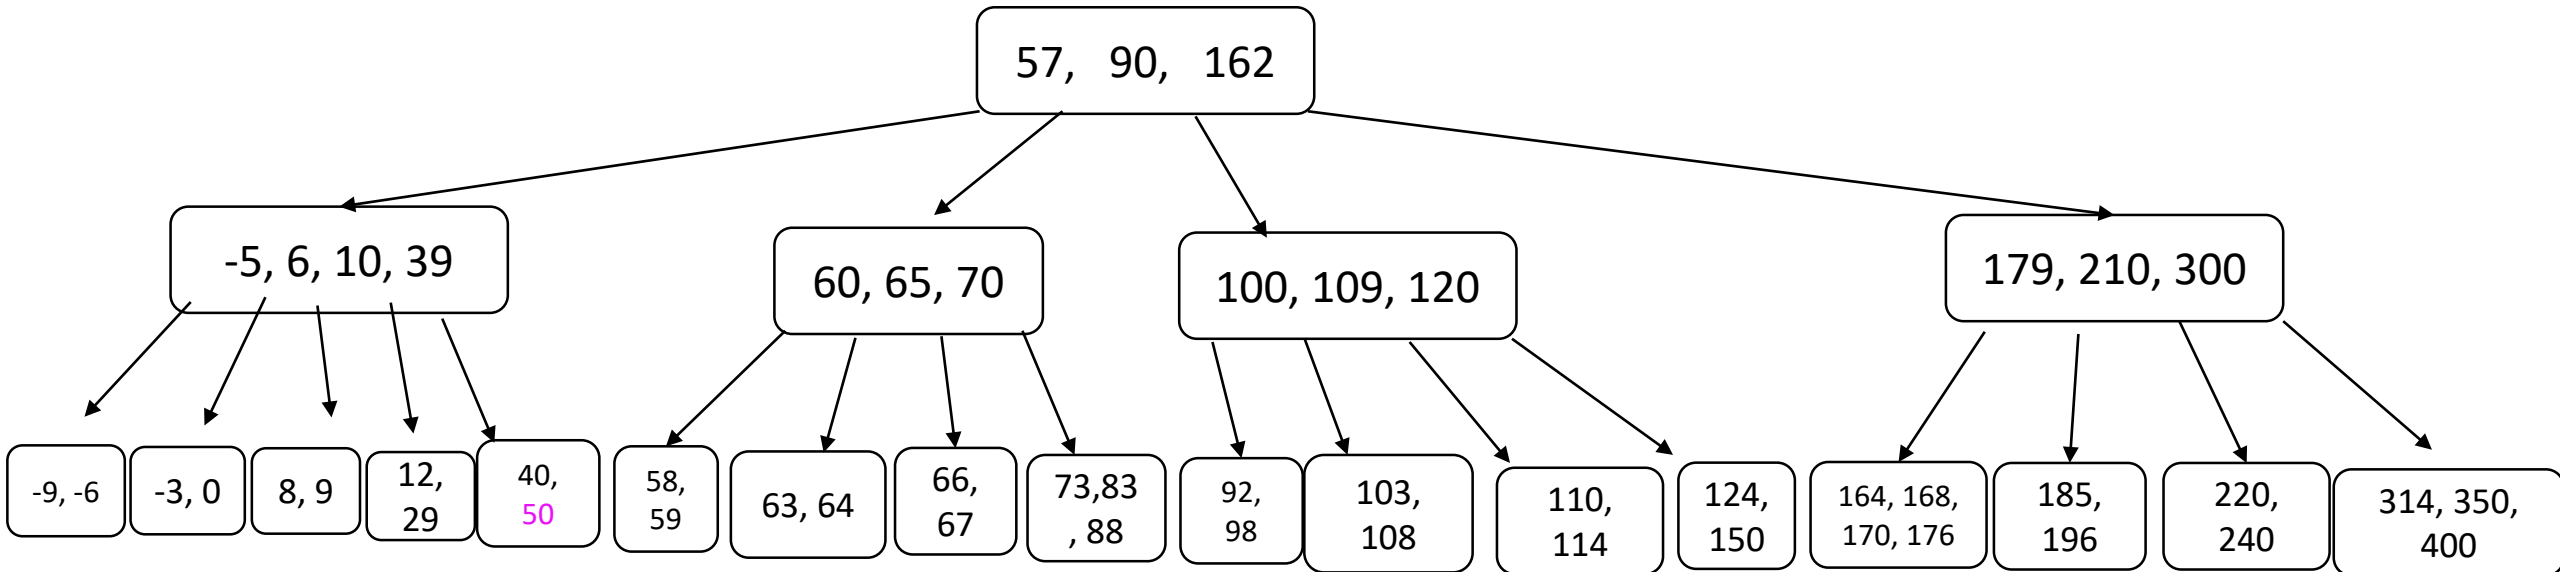
- Task: delete 112
- Underflow 😞
- Solvable via key rotation 😊
- The *hops* trick that we learned from insertion applies here as well!
 - Look as far away as *hops* tells you to.

Example #3: Merging with sibling node



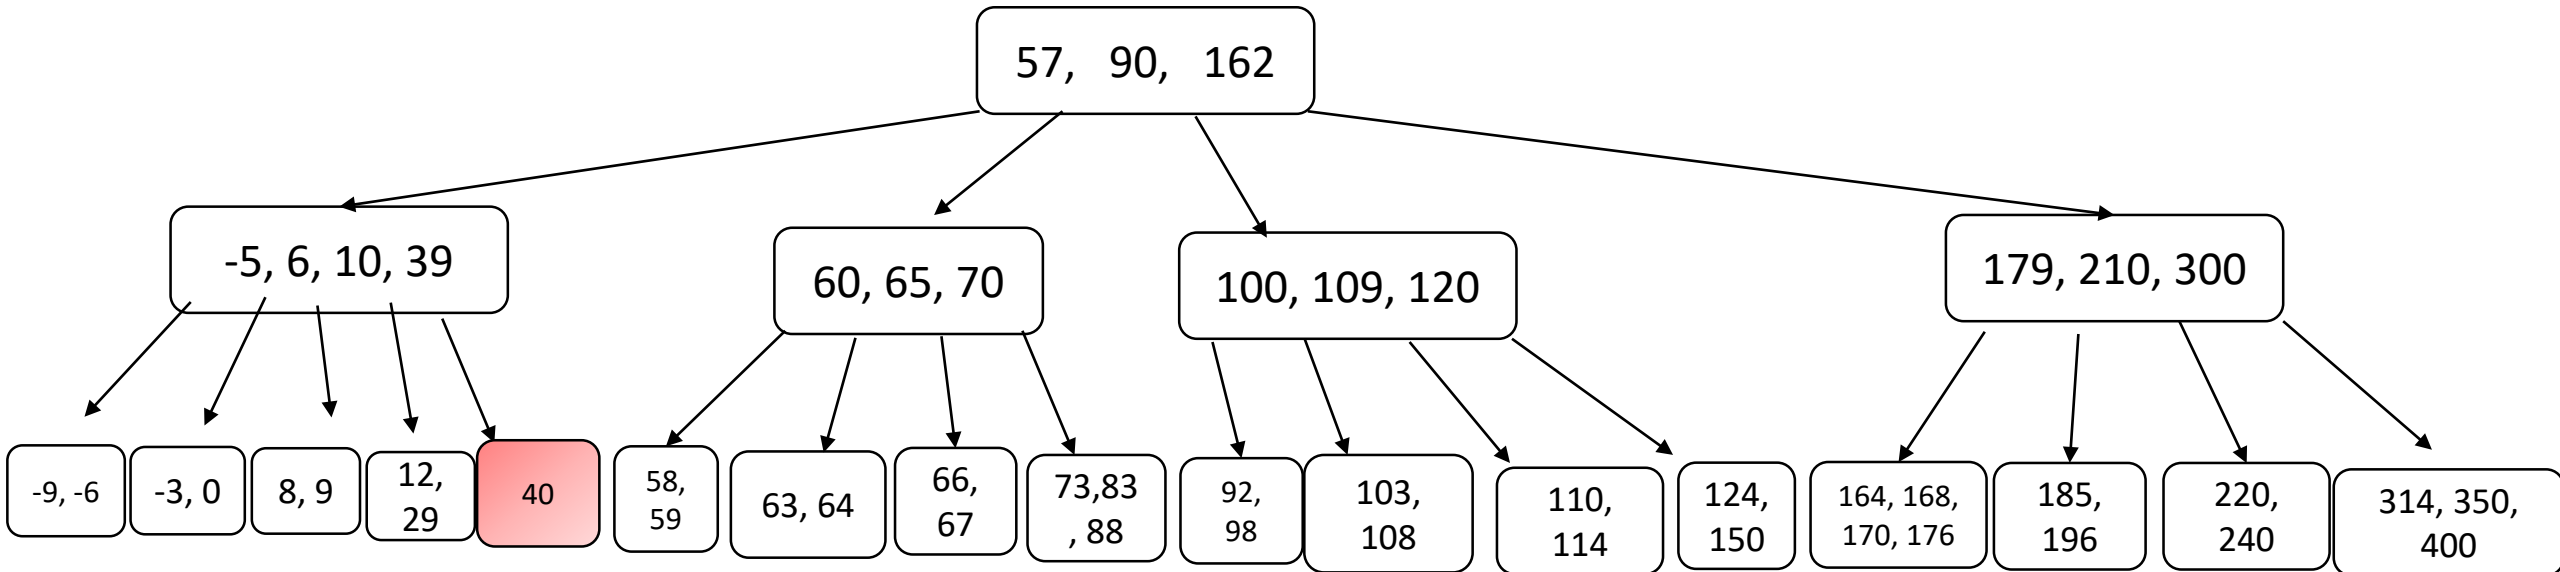
- Task: delete 50

Example #3: Merging with sibling node



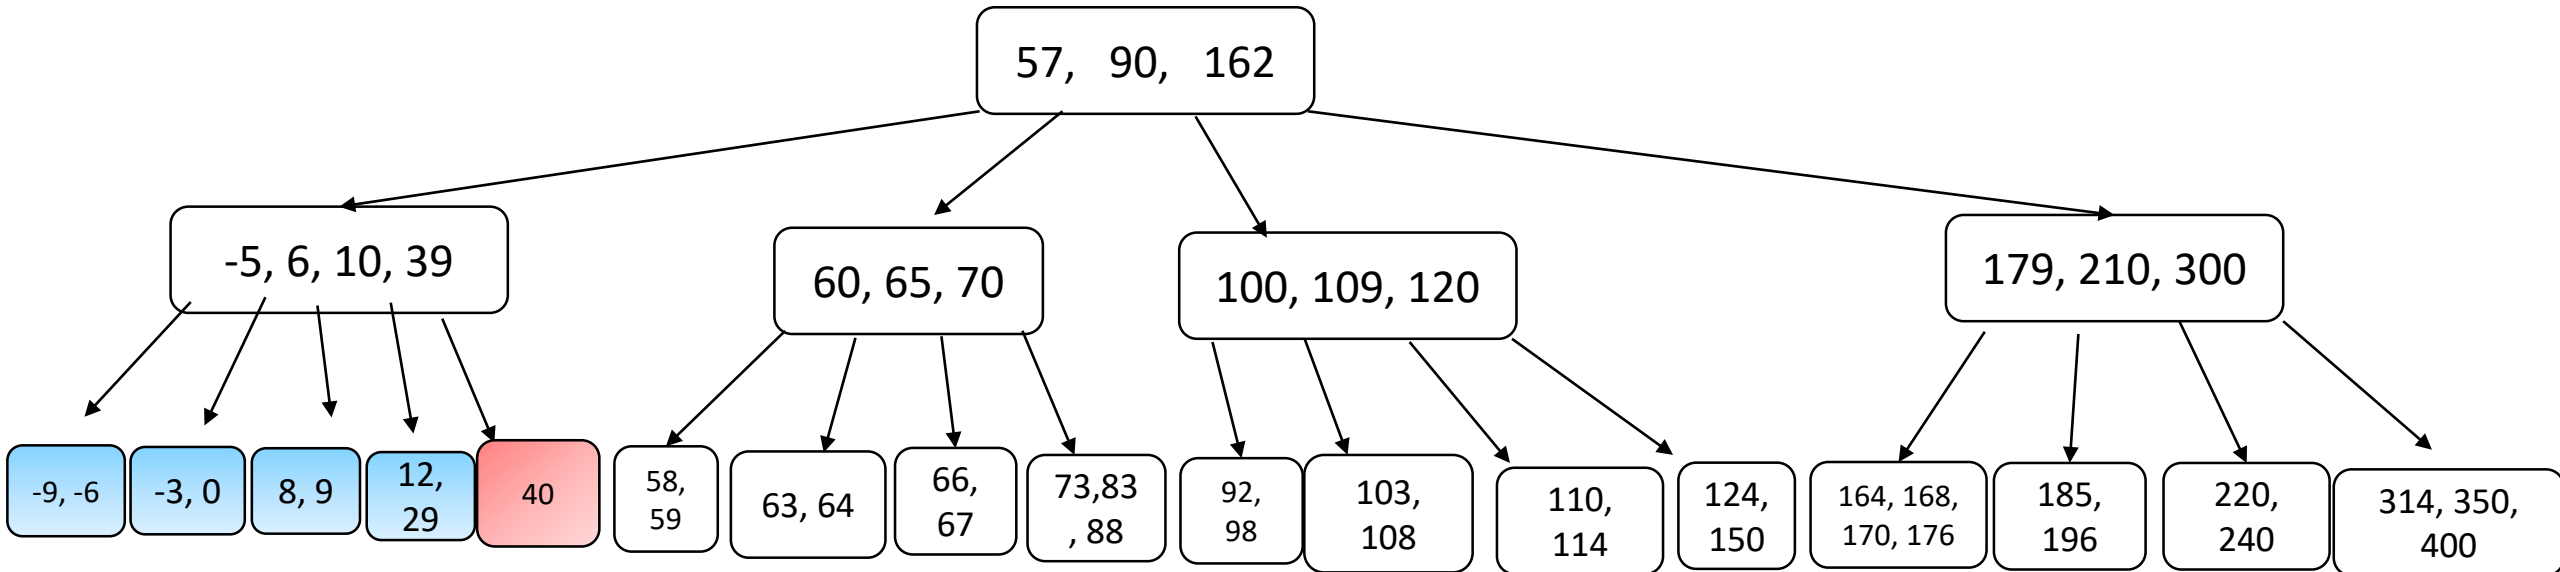
- Task: delete 50

Example #3: Merging with sibling node



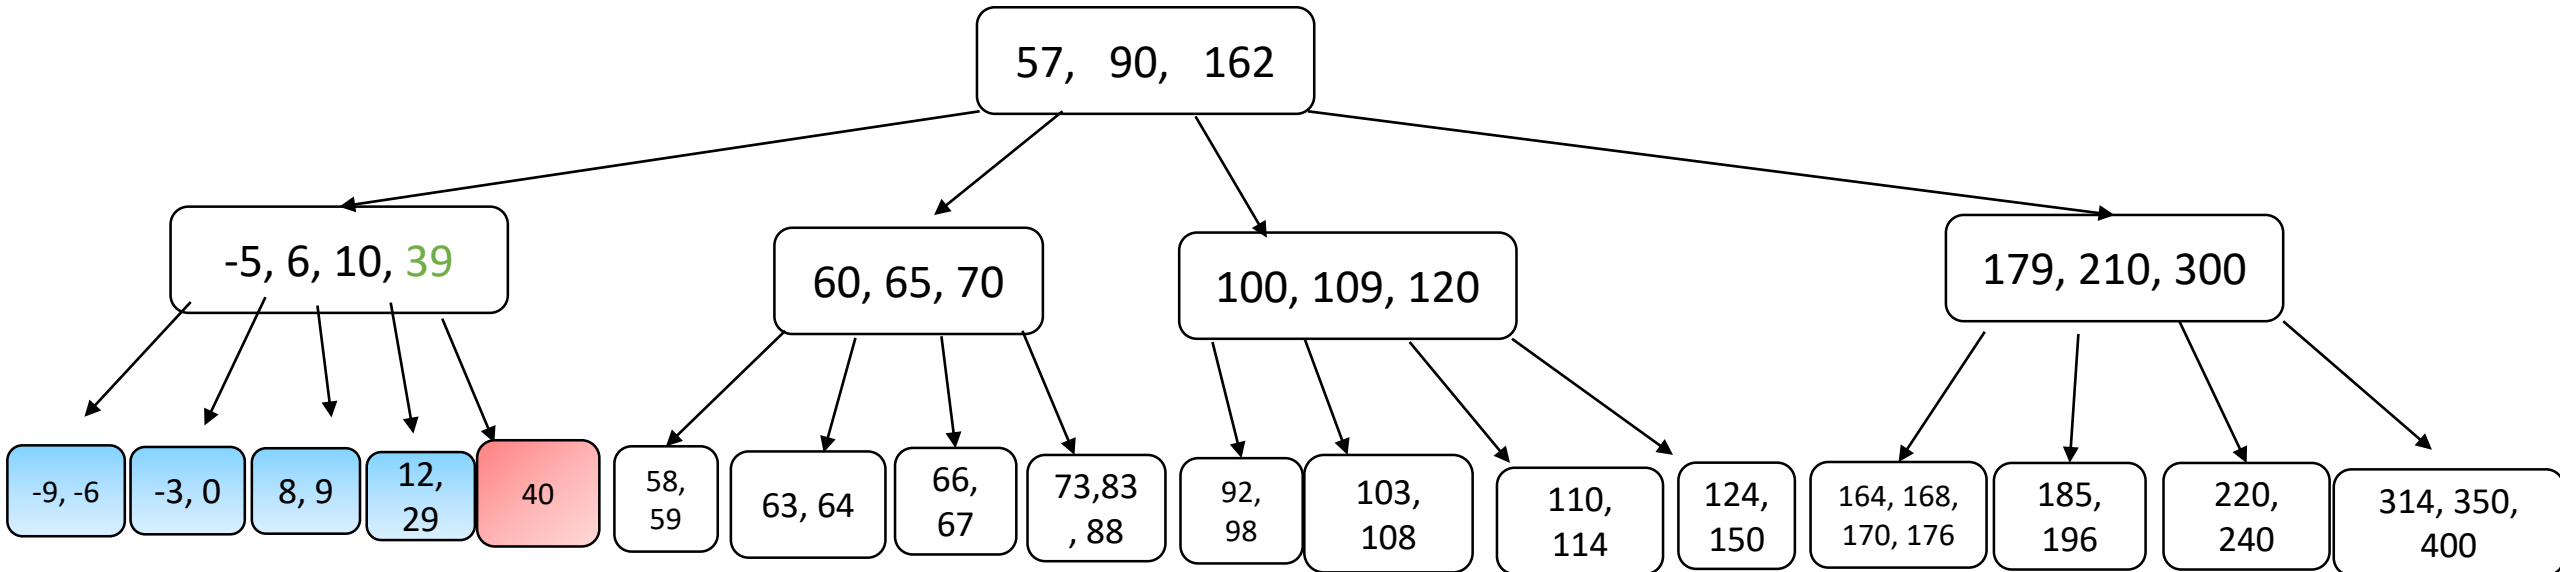
- Task: delete 50
- Underflow 😞

Example #3: Merging with sibling node



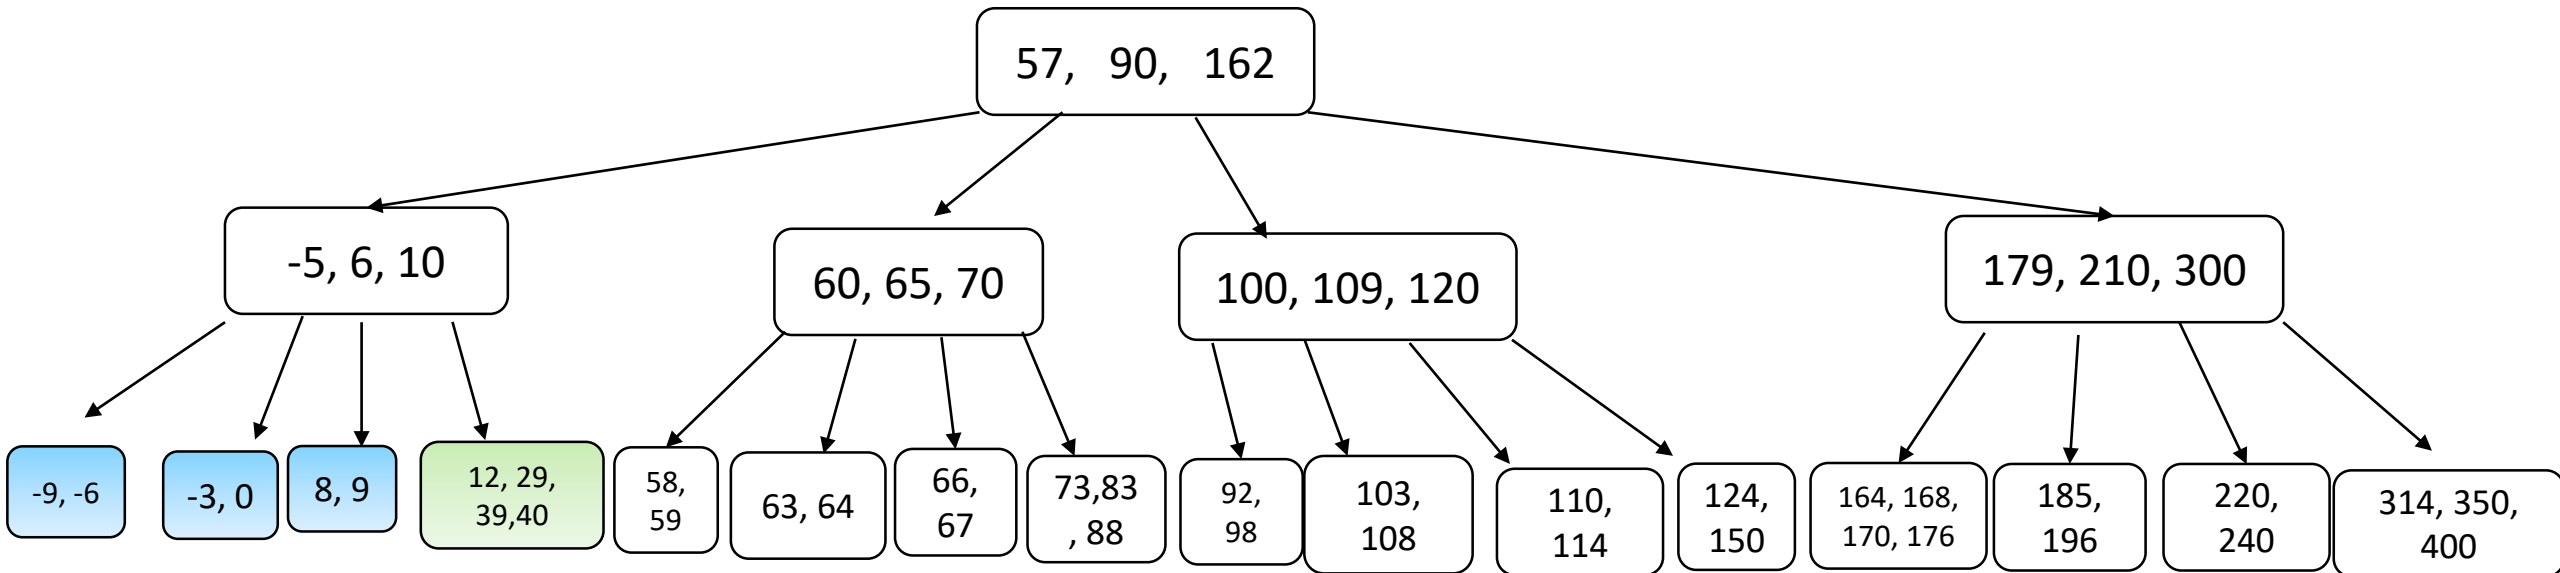
- Task: delete 50
- Underflow 😞
- Siblings can't help with key rotations 😞 (because $\min \# \text{ keys} = \left\lceil \frac{p}{2} \right\rceil - 1 = 3 - 1 = 2$)

Example #3: Merging with sibling node



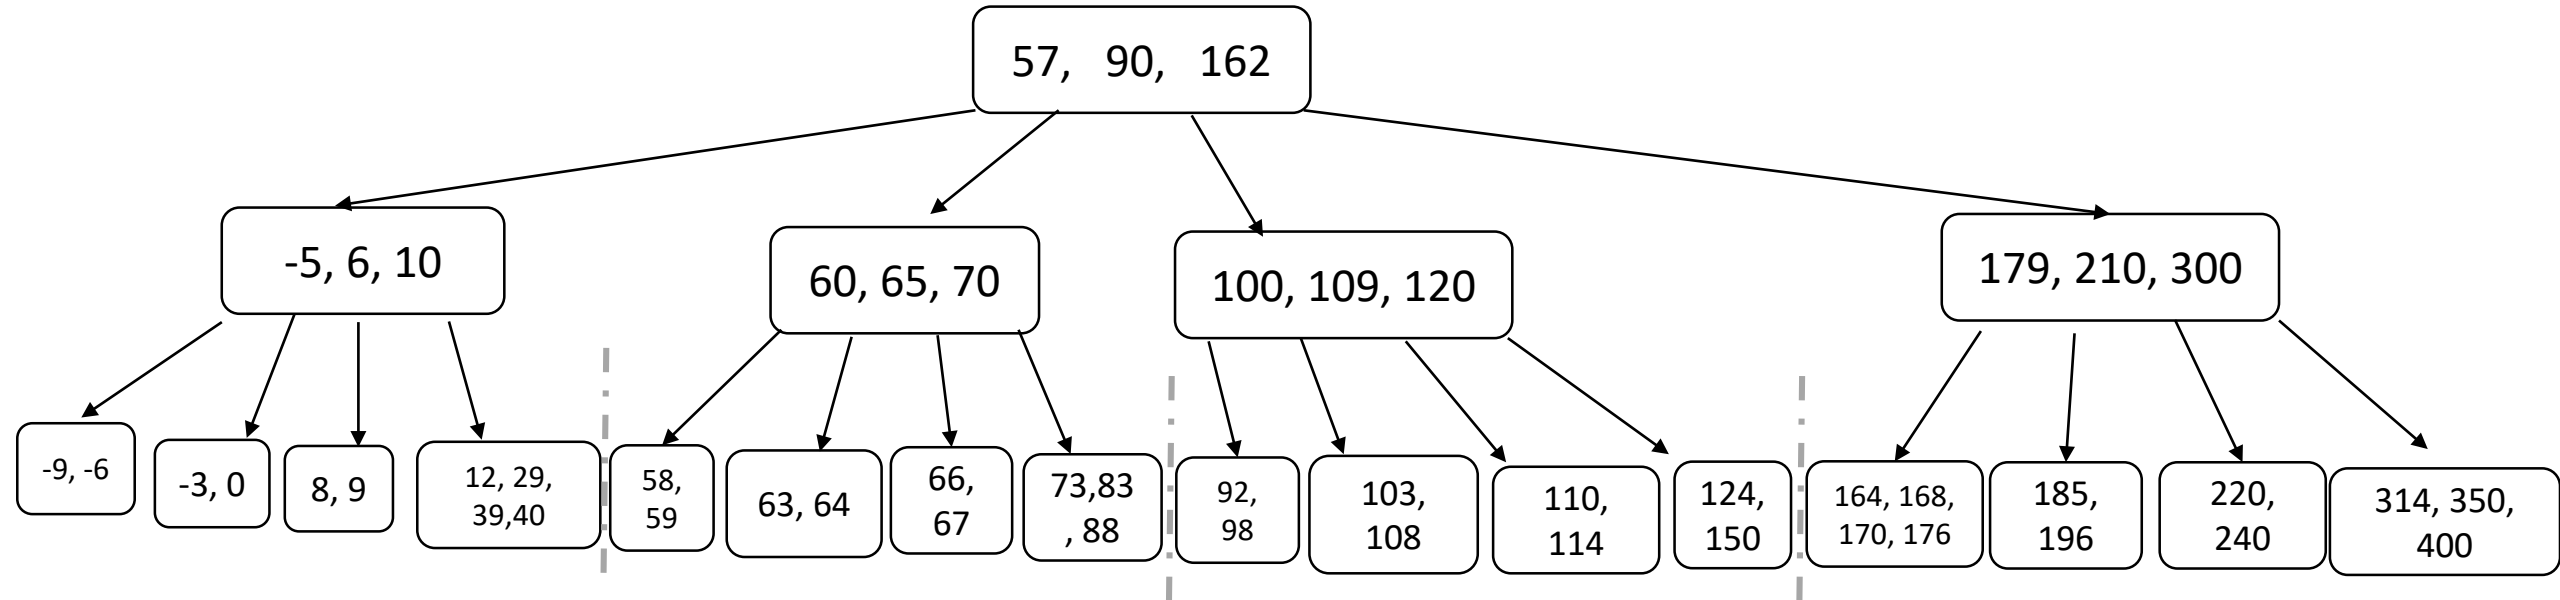
- Task: delete 50
- Underflow 😞
- Siblings can't help with key rotations 😞 (because $\min \# \text{ keys} = \left\lceil \frac{p}{2} \right\rceil - 1 = 3 - 1 = 2$)
- Solution: Merge current node with sibling node and parent key.

Example #3: Merging with parent node



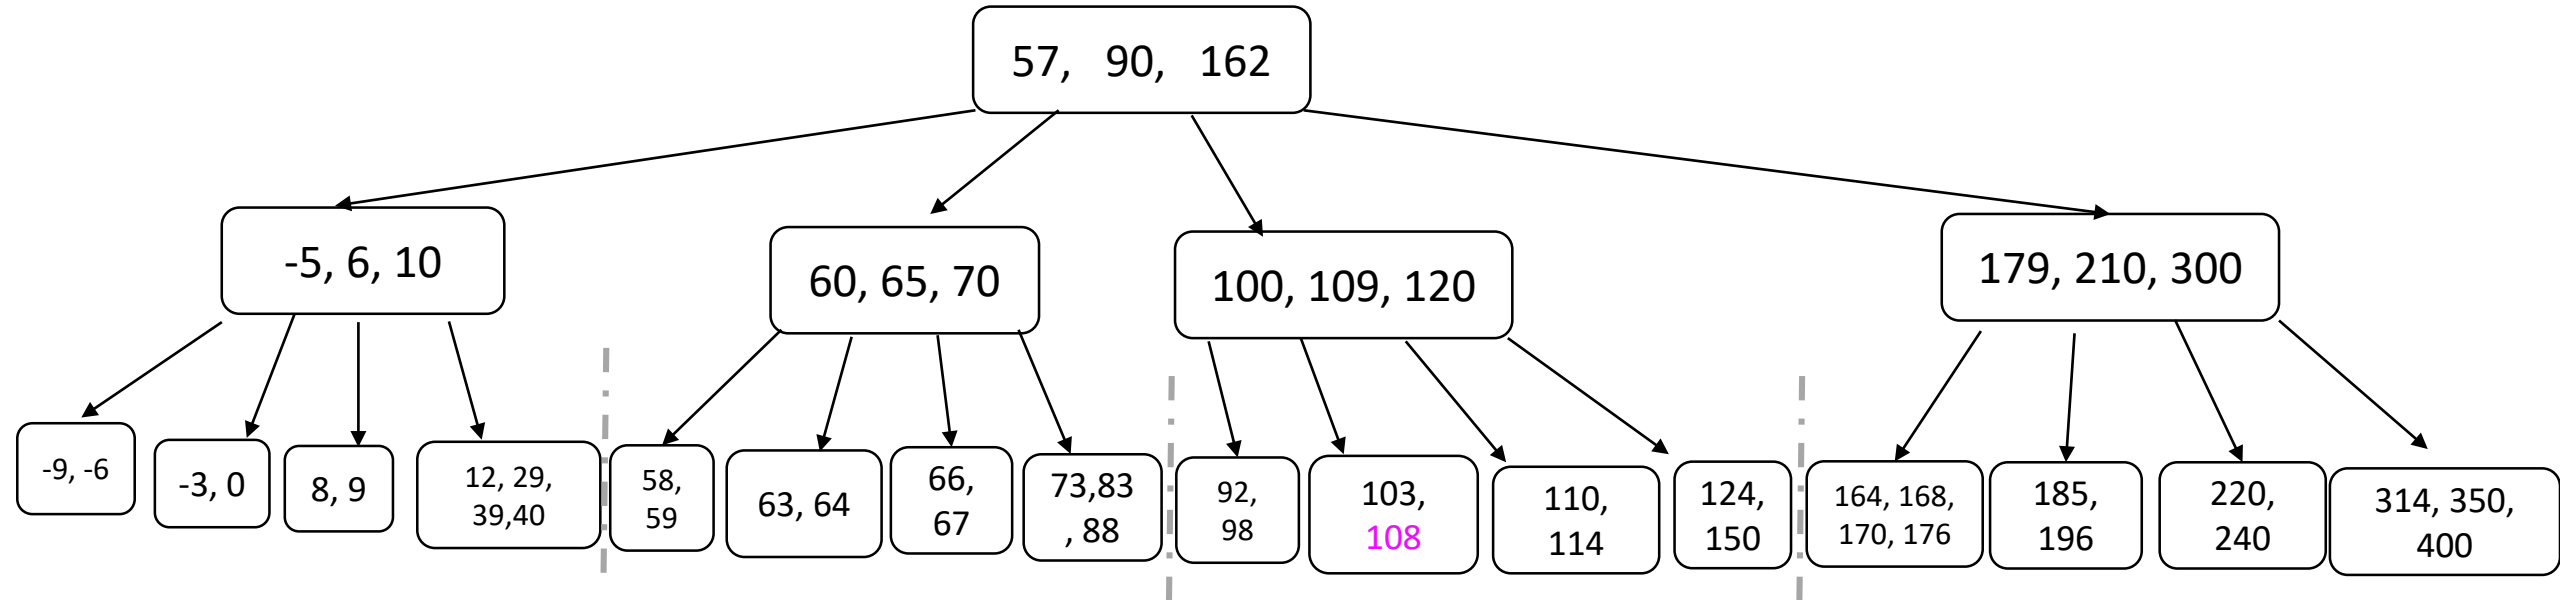
- Task: delete 50
- Underflow 😞
- Siblings can't help with key rotations 😞 (because $\min \# \text{ keys} = \left\lceil \frac{p}{2} \right\rceil - 1 = 3 - 1 = 2$)
- Solution: Merge current node with sibling node and parent key.

Example #4: Merging with either sibling



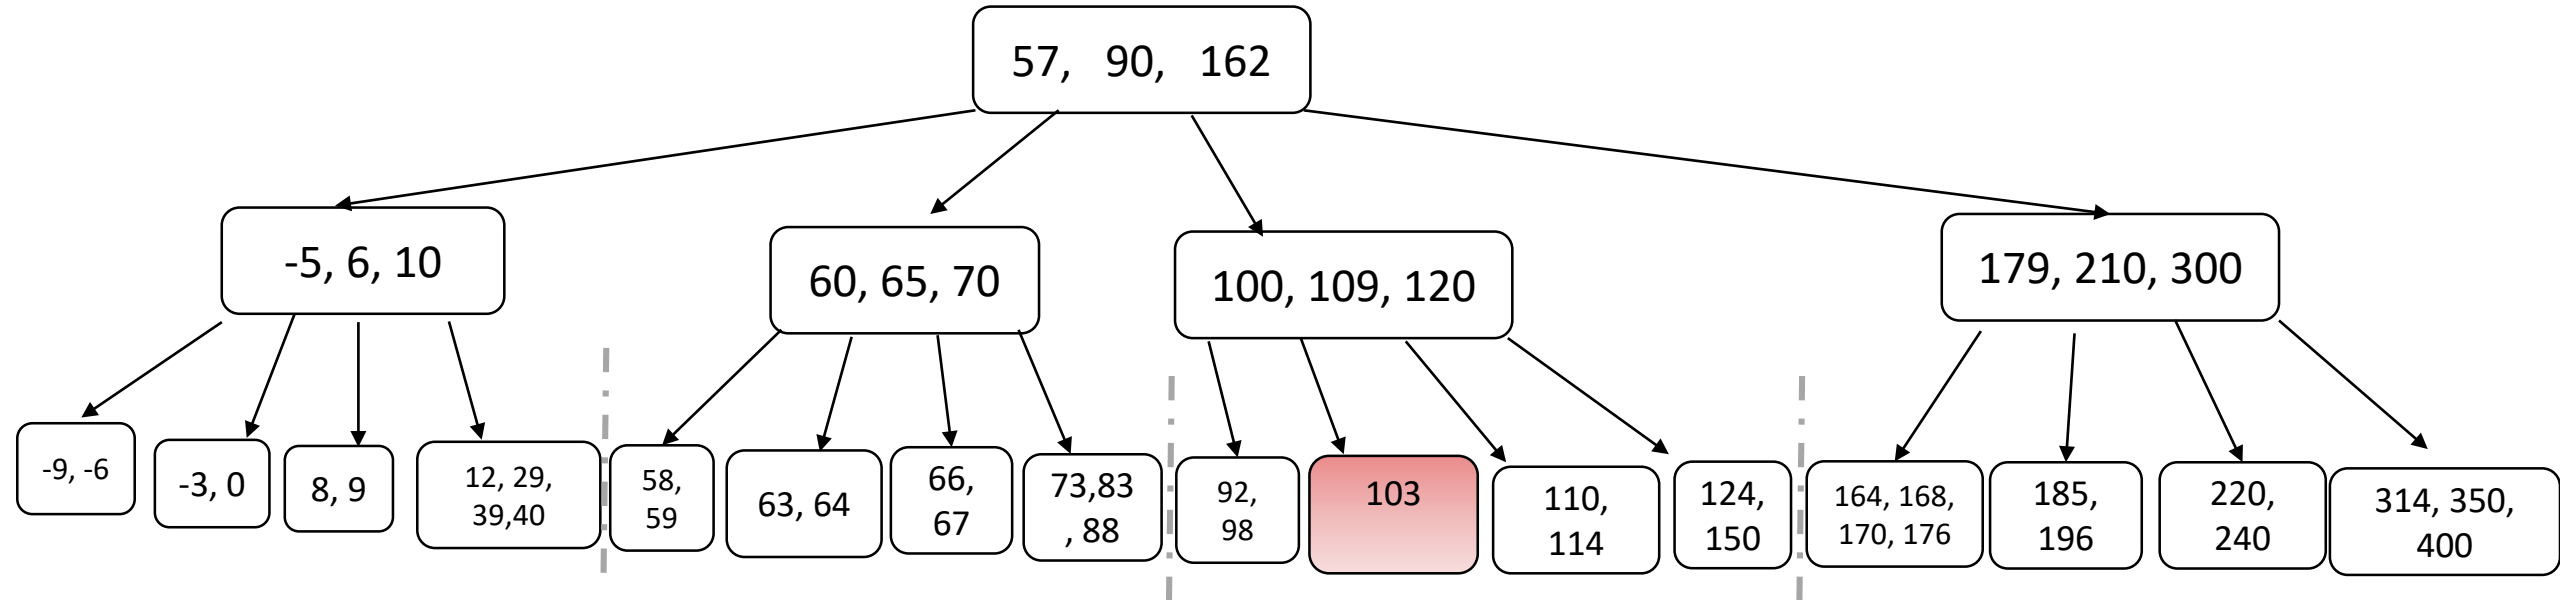
- Task: delete 108

Example #4: Merging with either sibling



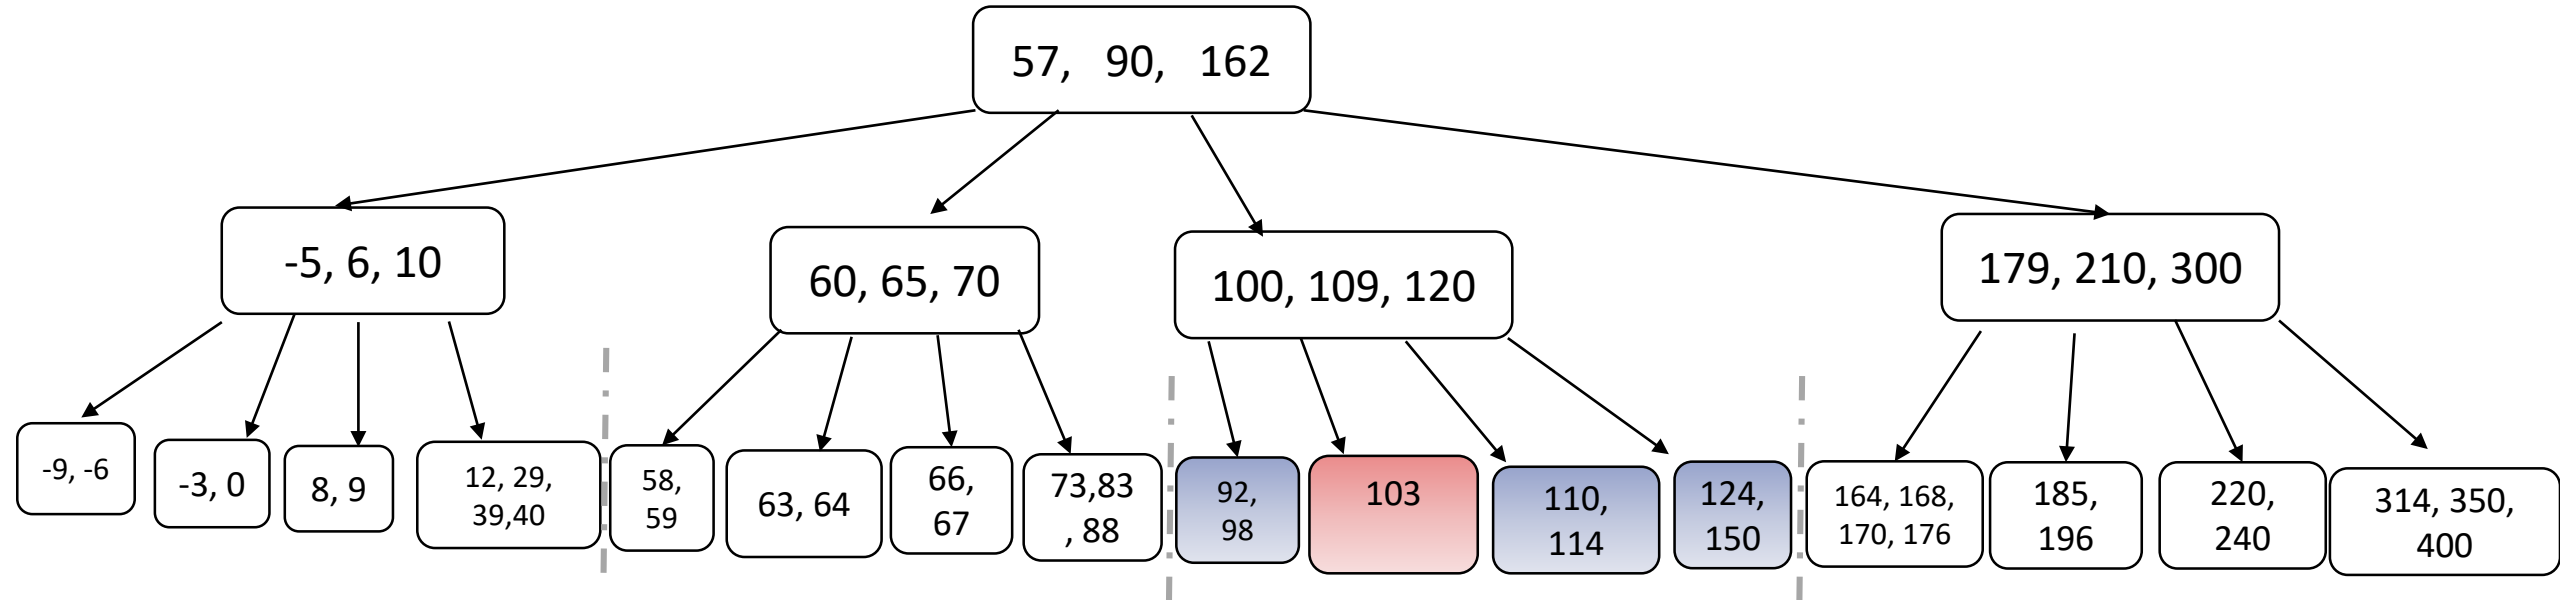
- Task: delete 108

Example #4: Merging with either sibling



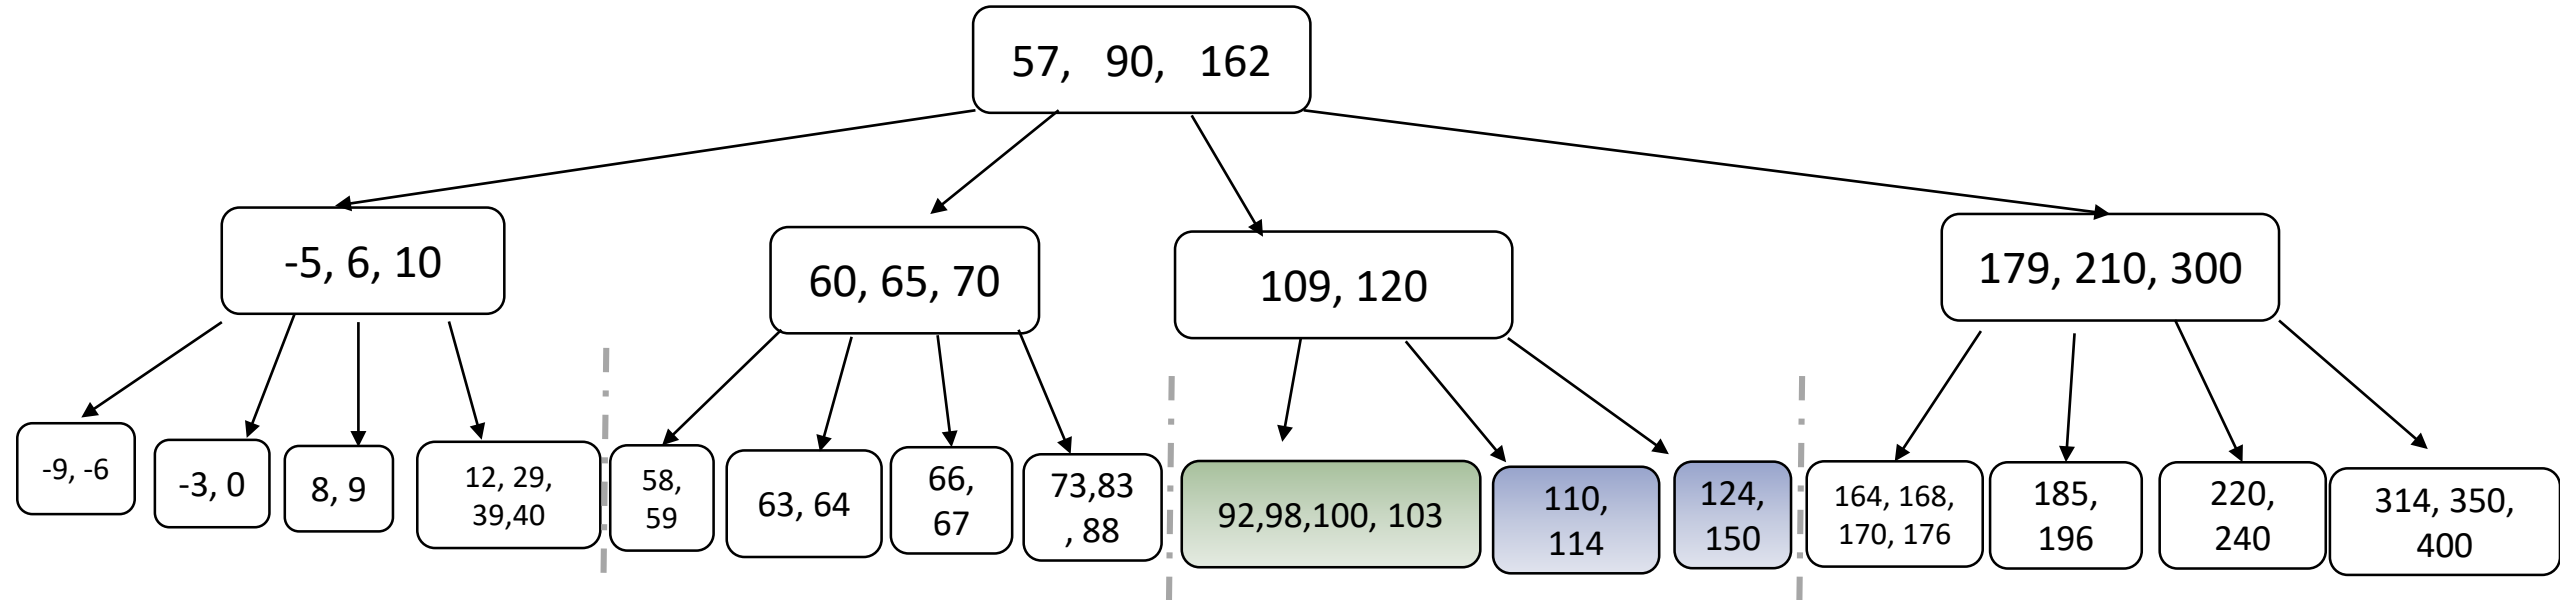
- Task: delete 108
- Underflow 😞

Example #4: Merging with either sibling



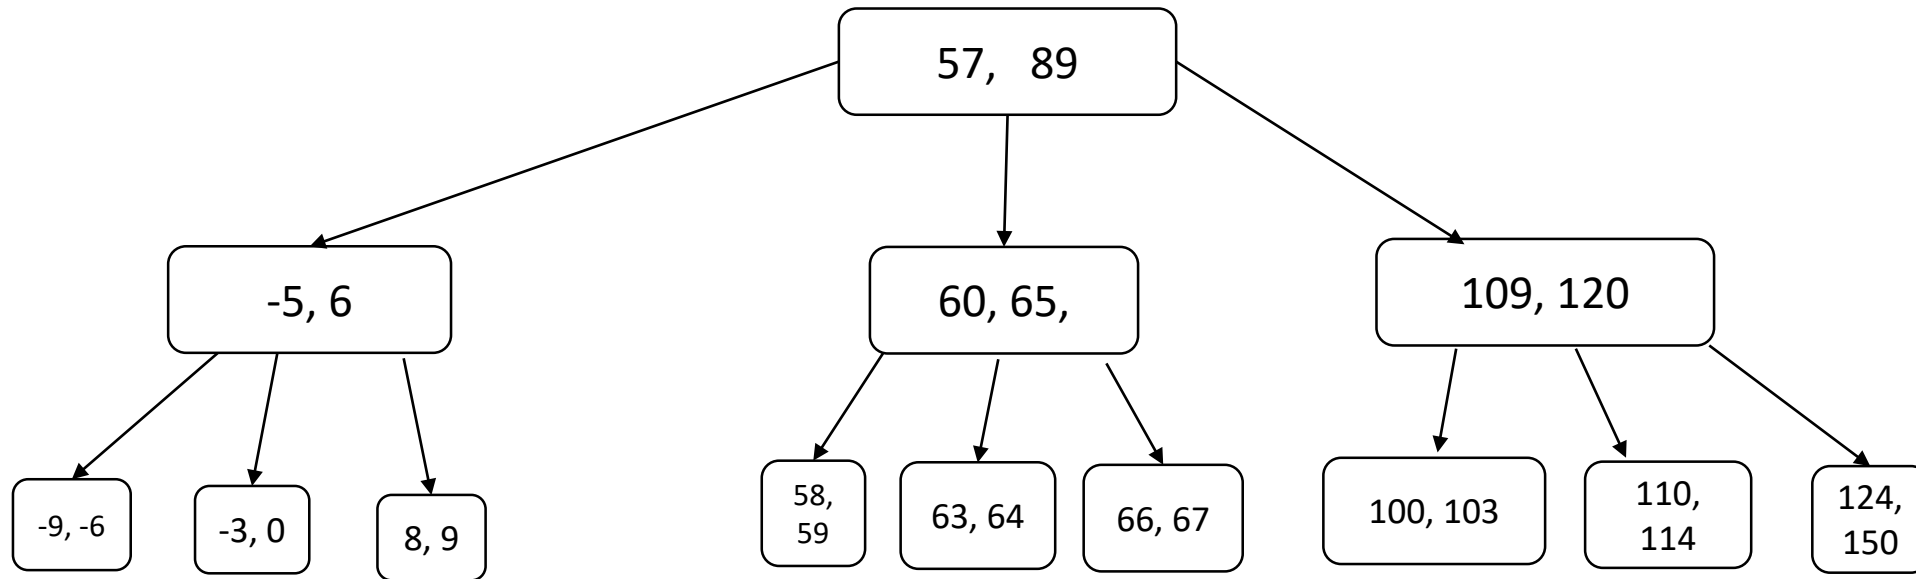
- Task: delete 108
- Underflow 😞
- Siblings can't help with key rotations 😞 (same reason)

Example #4: Merging with either sibling



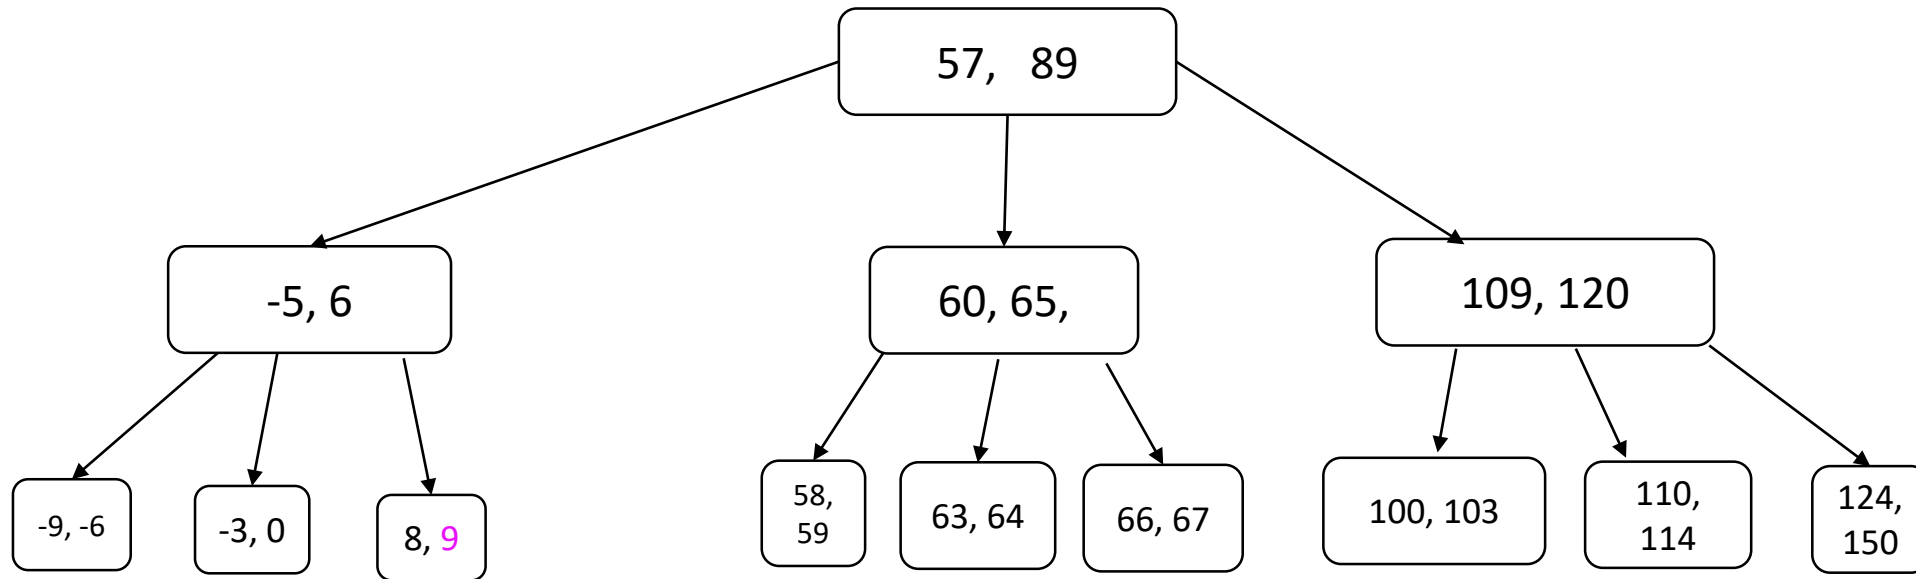
- Task: delete 108
- Underflow 😞
- Siblings can't help with key rotations 😞 (same reason)
- A relief force appears in the form of a sibling-parent key merging 😊

Example #5: Merging all the way to the root.



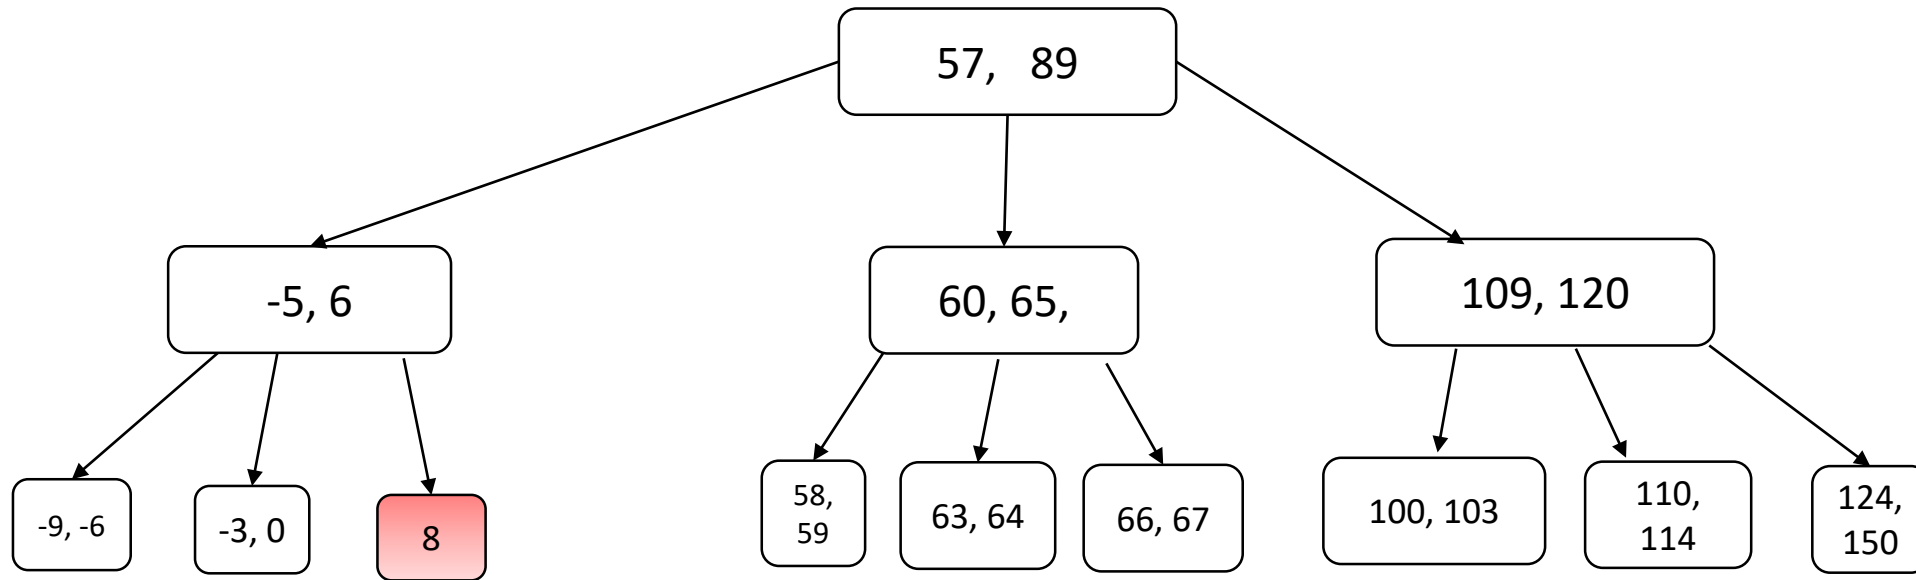
- Task: delete 9

Example #5: Merging all the way to the root.



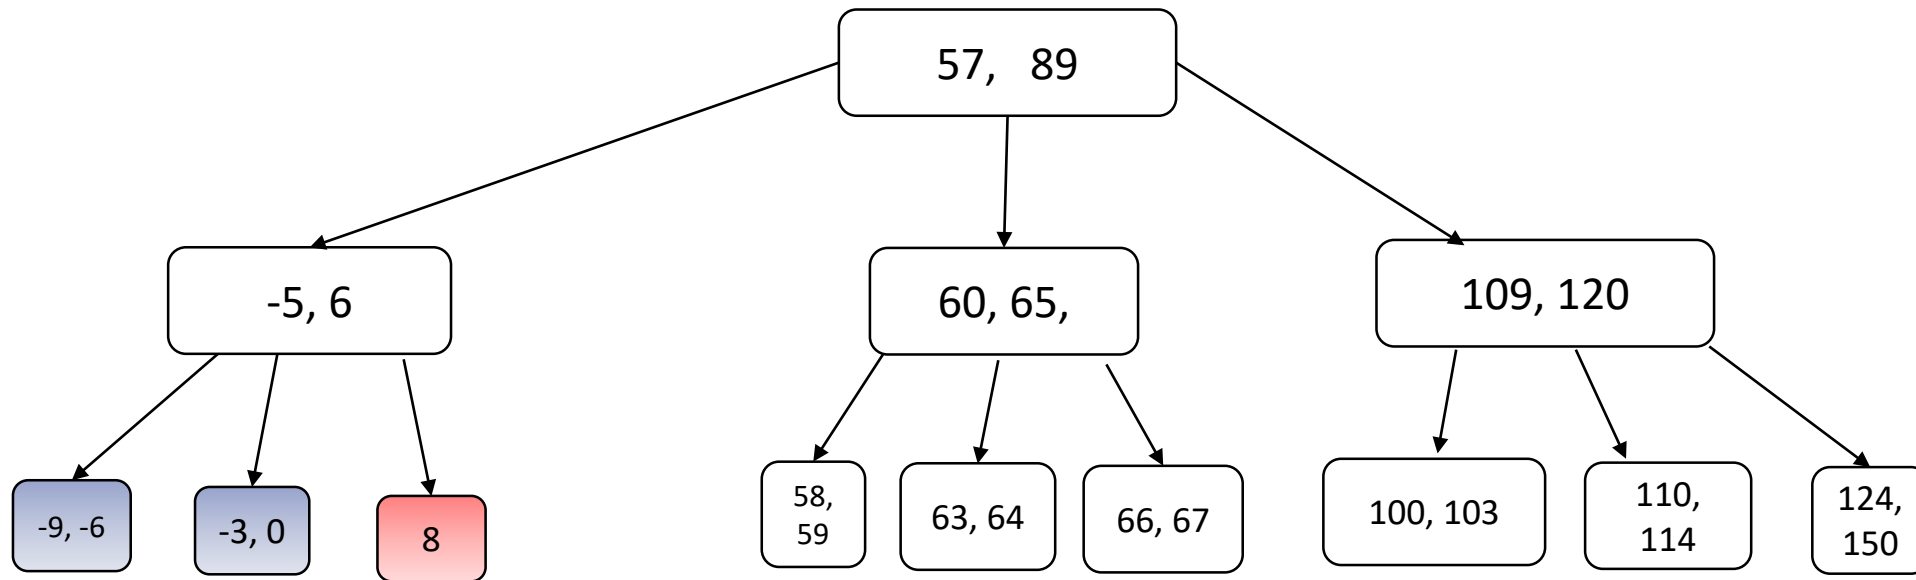
- Task: delete 9

Example #5: Merging all the way to the root.



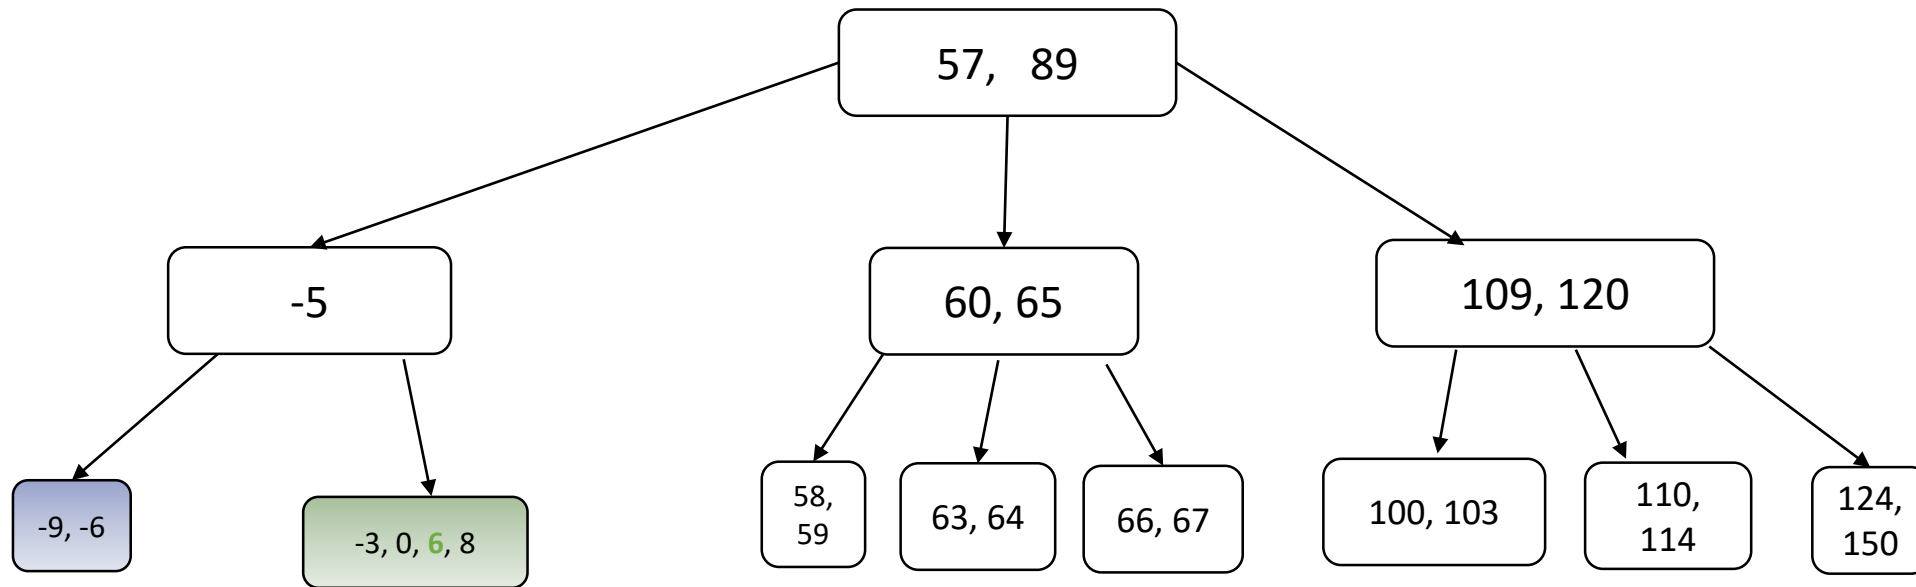
- Task: delete 9
- Underflow 😞

Example #5: Merging all the way to the root.



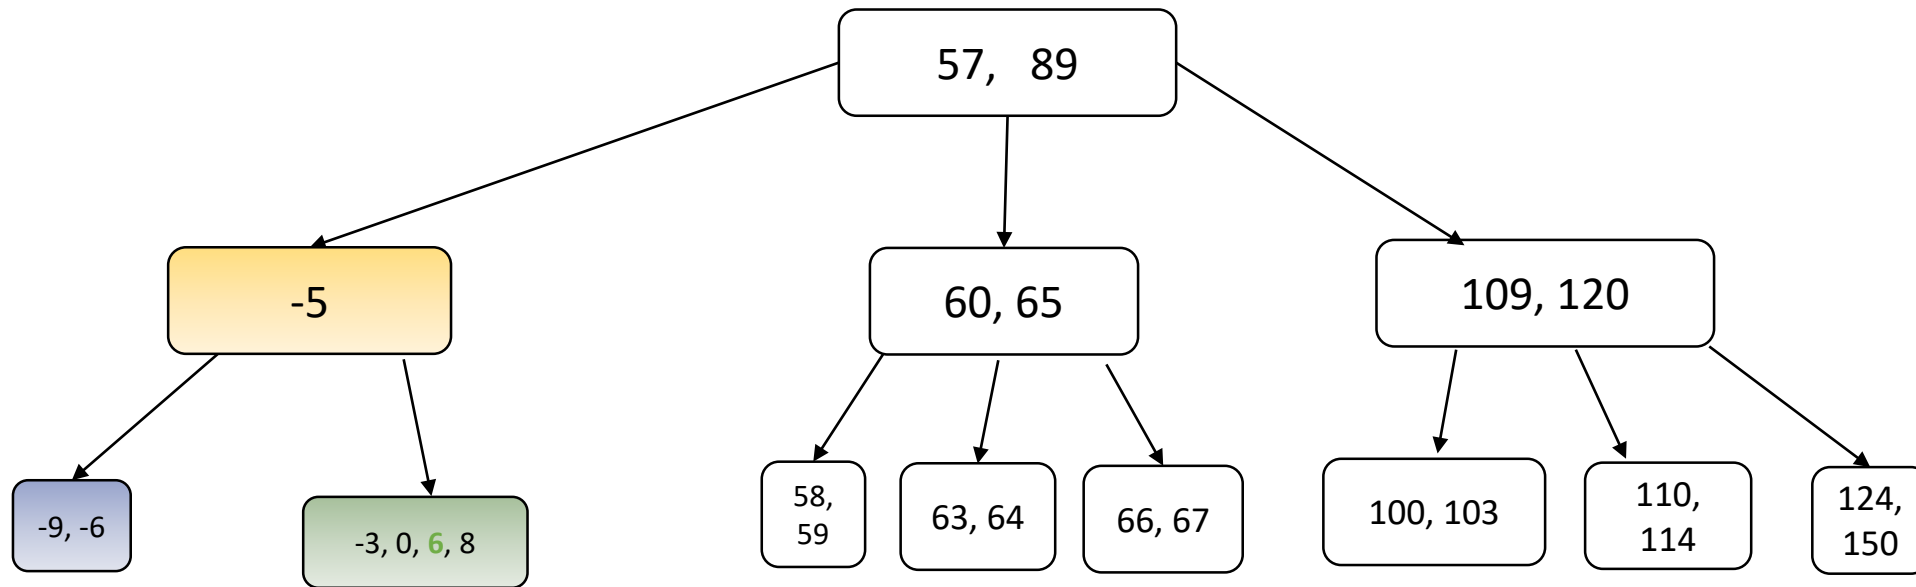
- Task: delete 9
- Underflow 😞
- Siblings can't help with key rotations! 😞

Example #5: Merging all the way to the root.



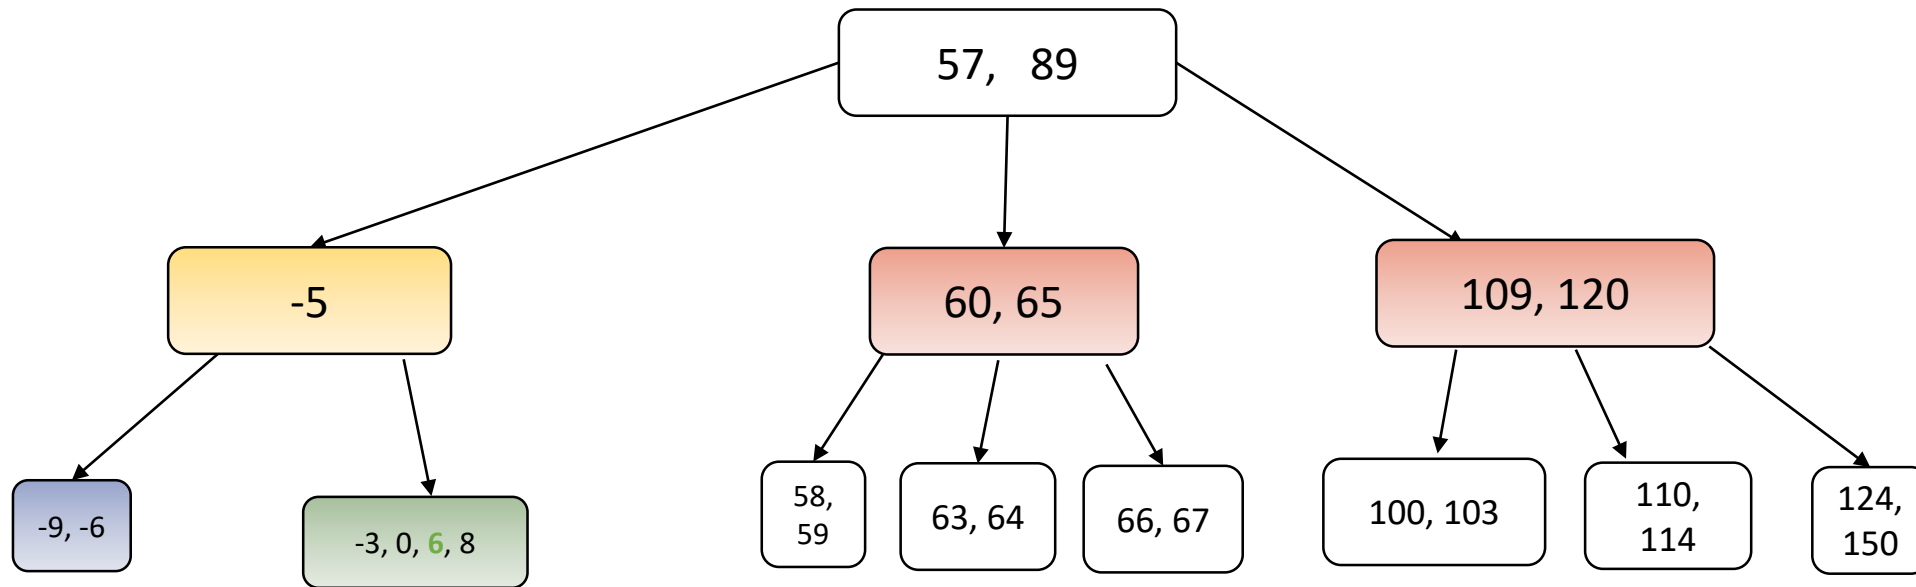
- Task: delete 9
- Underflow 😞
- Siblings can't help with key rotations! 😞
- Merging with parent key and sibling 😊

Example #5: Merging all the way to the root.



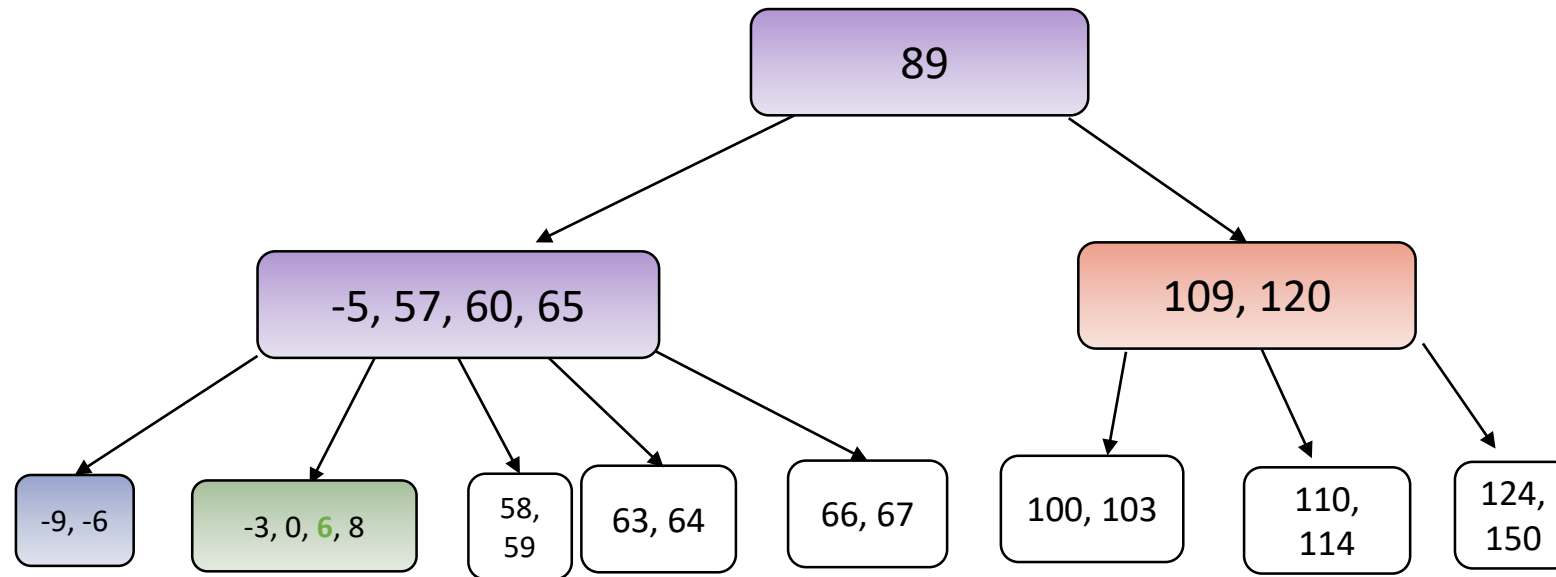
- Task: delete 9
- Underflow 😞
- Siblings can't help with key rotations! 😞
- Merging with parent key and sibling 😊
- Parent underflows 😞

Example #5: Merging all the way to the root.




- Task: delete 9
- Underflow 😞
- Siblings can't help with key rotations! 😞
- Merging with parent key and sibling 😊
- Parent underflows 😞
- Siblings can't help with rotations 😞

Example #5: Merging all the way to the root.




- Task: delete 9
- Underflow 😞
- Siblings can't help with key rotations! 😞 Solution: Merge with parent key and sibling node! 😊
- Merging with parent key and sibling 😊
- Parent underflows 😞
- Siblings can't help with rotations 😞

Example #6: Root underflow



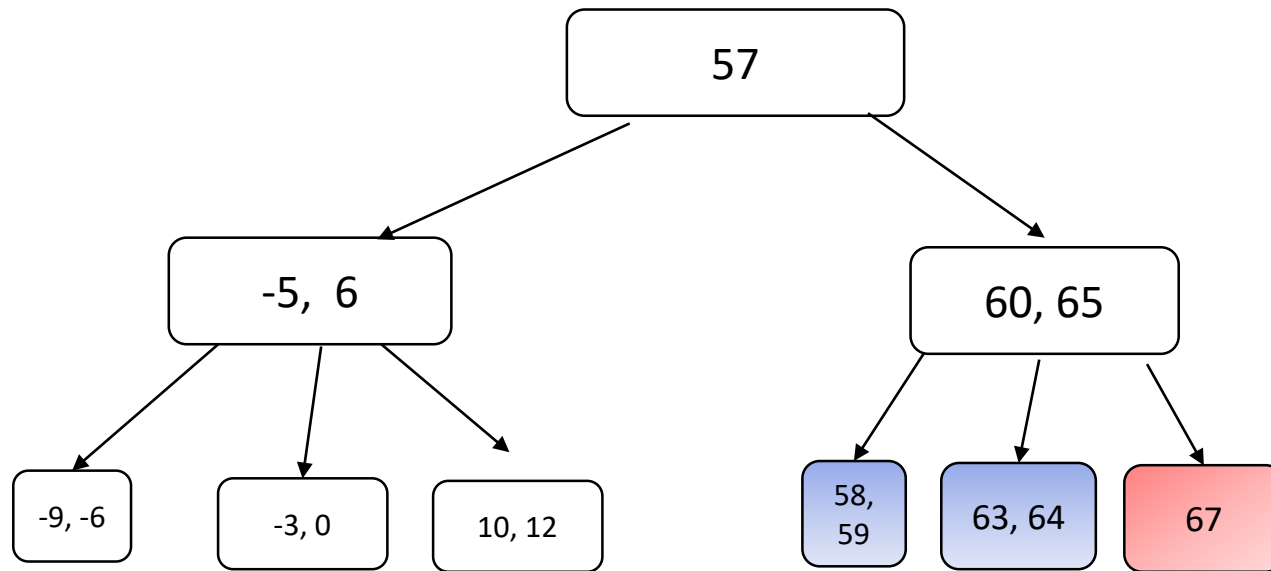
- Task: delete 66

Example #6: Root underflow



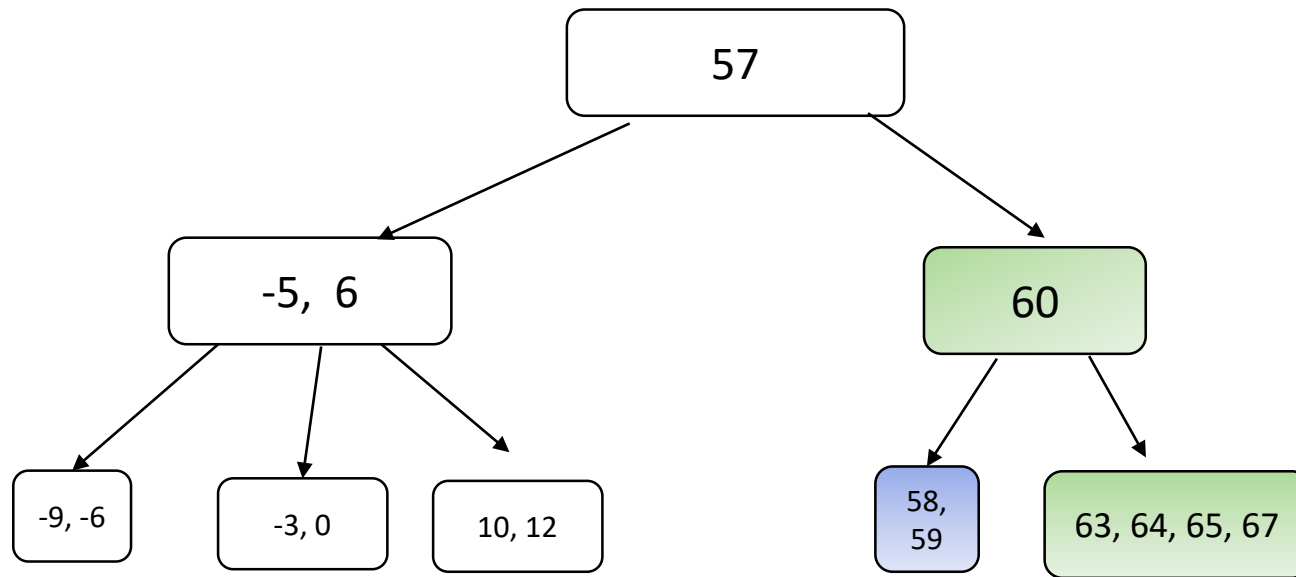
- Task: delete 66
- Underflow 😞

Example #6: Root underflow



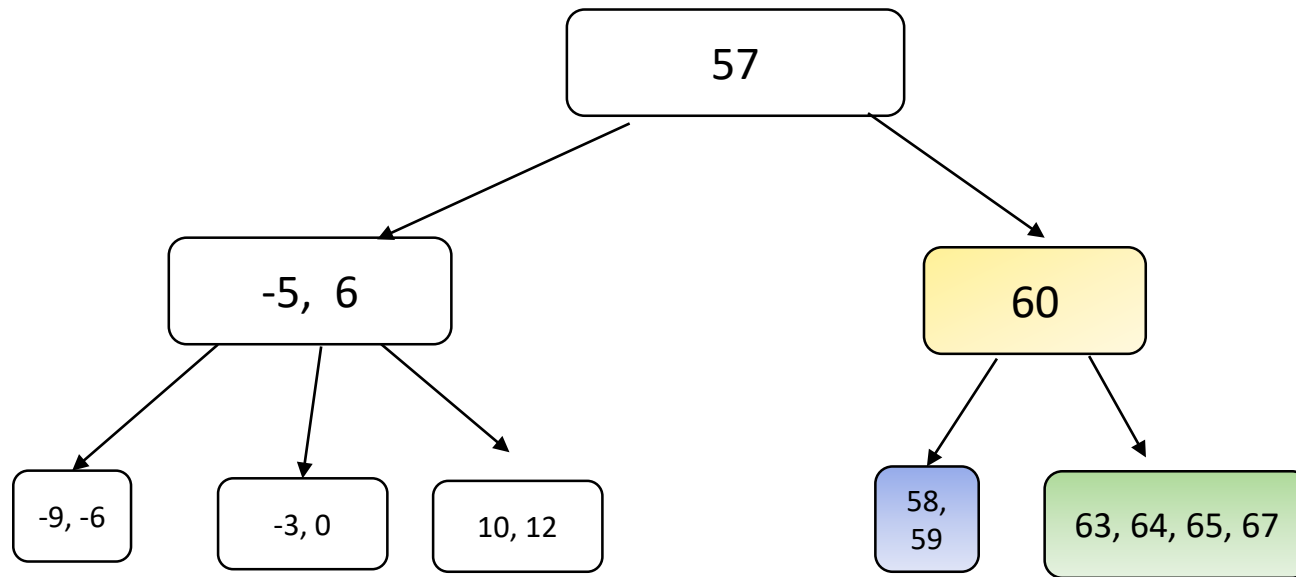
- Task: delete 66
- Underflow 😞
- Siblings can't help with key rotations 😞

Example #6: Root underflow



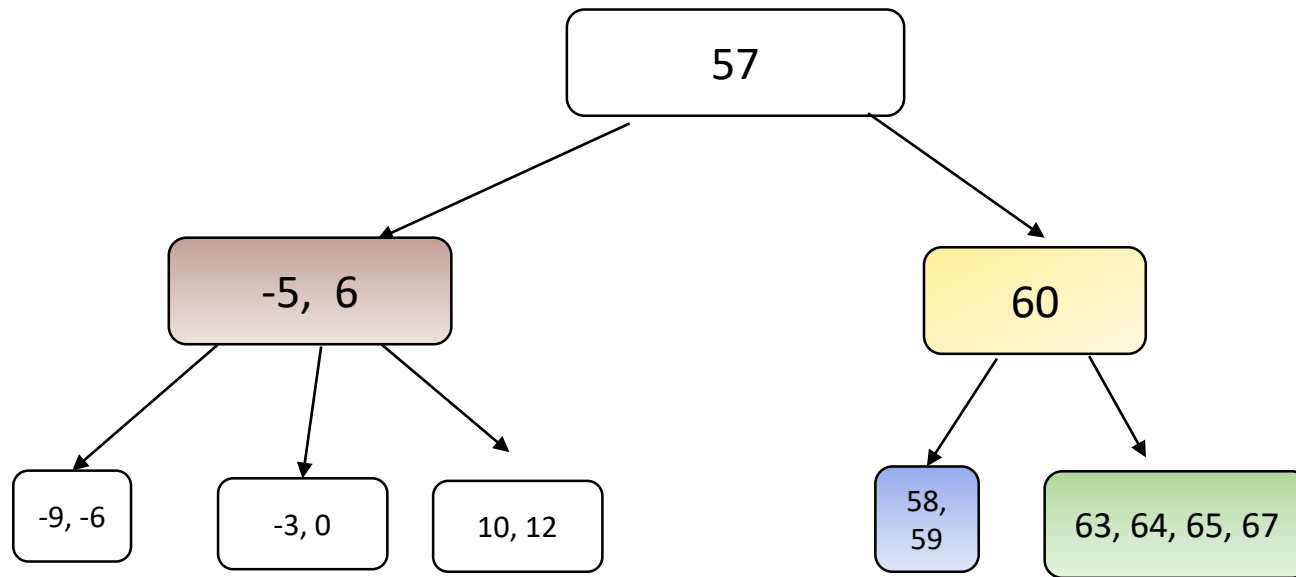
- Task: delete 66
- Underflow 😞
- Siblings can't help with key rotations 😞
- Merge 😊

Example #6: Root underflow




- Task: delete 66
- Underflow 😞
- Siblings can't help with key rotations 😞
- Merge 😊
- Underflow 😞

Example #6: Root underflow



- Task: delete 66
- Underflow 😞
- Siblings can't help with key rotations 😞
- Merge 😊
- Underflow 😞
- Siblings can't help with key rotations 😞

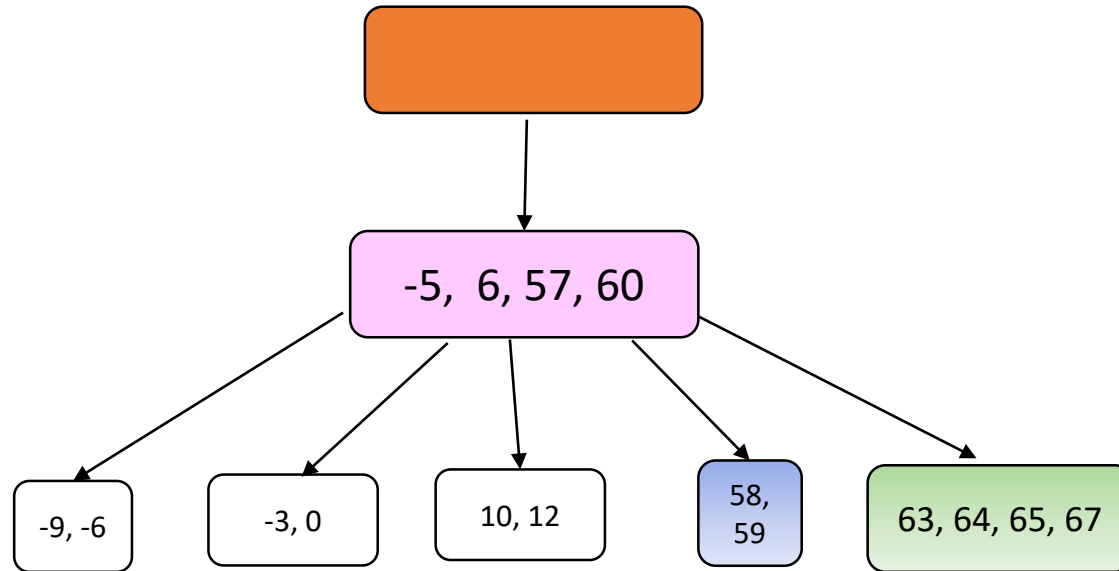
Example #6: Root underflow



- Task: delete 66
- Underflow 😞
- Siblings can't help with key rotations 😞
- Merge 😊
- Underflow 😞
- Siblings can't help with key rotations 😞

Merge 😊

Example #6: Root underflow




- Task: delete 66
- Underflow 😞
- Siblings can't help with key rotations 😞
- Merge 😊
- Underflow 😞
- Siblings can't help with key rotations 😞

Merge 😊

Root underflows 😞

Example #6: Root underflow



- Task: delete 66
- Underflow ☹️
- Siblings can't help ☹️
- Merge 😊
- Underflow ☹️
- Siblings can't help ☹️
- Root underflows ☹️
- Delete it, and set new root to its only child! 😊

Searching a B-Tree

- Searching has **effectively been covered** by presenting insertion and deletion.
- At every level, I will query the node **via binary search**. Either:
 - I will find the key, so I'm done, or
 - I will follow the relevant pointer to a subtree, and **repeat**.
- Worse case: I have to go **all the way to the leaves**.

Quiz

- I want to measure the **worst-case search cost** in a B-Tree with **branching factor p** .
- **Unit cost**: comparison of two keys via **`compareTo()`**.

Quiz

- I want to measure the **worst-case search cost** in a B-Tree with **branching factor p** .
- **Unit cost**: comparison of two keys via **`compareTo()`**.
- Question: **in the worst case**, how many comparisons will I make **for a key search** in a B-Tree with n **keys**?

Quiz

- I want to measure the **worst-case search cost** in a B-Tree with **branching factor p** .
- **Unit cost**: comparison of two keys via **`compareTo()`**.
- Question: **in the worst case**, how many comparisons will I make **for a key search** in a B-Tree with n **keys**?

$$\log_p n$$

$$p * \log_p n$$

$$(p - 1) * \log_p n$$

Something
else (what?)

Quiz

- I want to measure the **worst-case search cost** in a B-Tree with branching factor p .
- **Unit cost**: comparison of two keys via `compareTo()`.
- Question: **in the worst case**, how many comparisons will I make **for a key search** in a B-Tree with n keys?

$$\log_p n$$

$$p * \log_p n$$

$$(p - 1) * \log_p n$$

$$\log_2(p - 1) * \log_p n \approx \log_2 n$$

Something
else (what?)



With the keys **sorted by construction**, it really is dumb to not do binary search in every node! ;)

So, why even care?

- Why even care about B-Trees if their search is $\mathcal{O}(\log_2 n)$?




So, why even care?

- **Why even care** about B-Trees if their search is $\mathcal{O}(\log_2 n)$?
- Because of the **hugely successful** B+ (“B-plus”) trees, of which they are the main component!



The components of a B+-tree

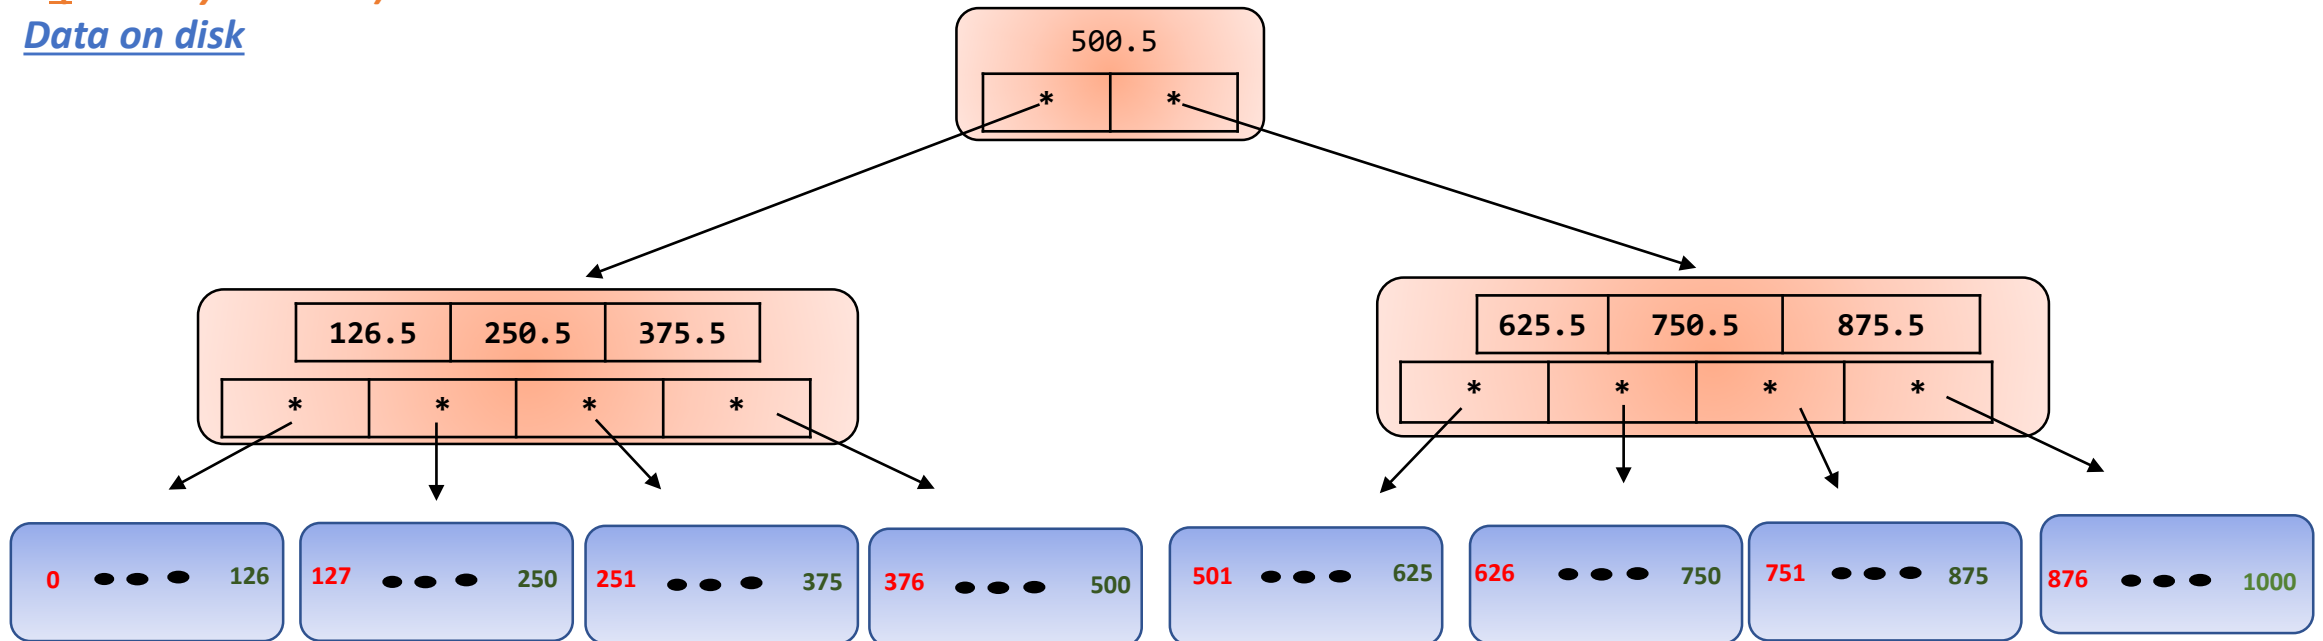
(1) A memory-resident B-Tree



The components of a B+-tree

(1) *A (memory-resident) B-Tree*

(2) *Data on disk*

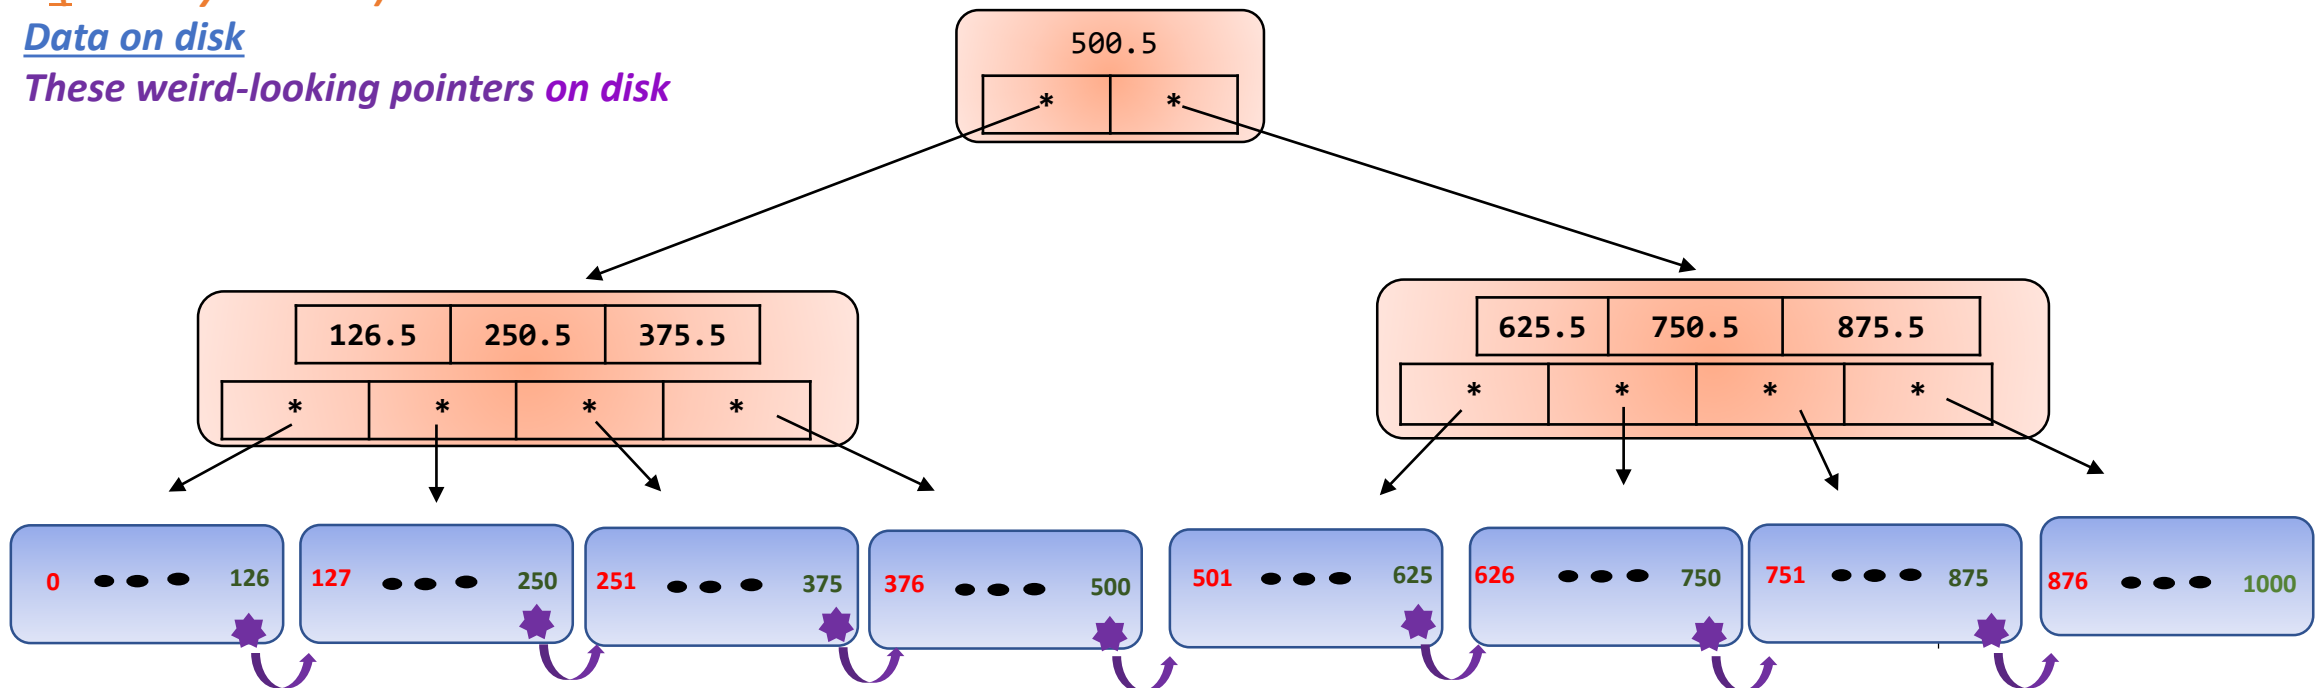


The components of a B+-tree

(1) *A (memory-resident) B-Tree*

(2) *Data on disk*

(3) *These weird-looking pointers on disk*




The components of a B+-tree

Give us a *B⁺-tree* (“B – plus tree”) index!

(1) A memory-resident B-Tree

(2) Data on disk

(3) These weird-looking pointers on disk



Are buffers a good storage choice?

- **Issue:** Static storage buffers for keys and references waste memory space for nodes.
 - Proposal: Use linked lists instead!
- This is actually not a good idea.
- Counter-arguments:
 1. Buffers waste space **MINIMAL** compared to the size of the DB (recall earlier results)
 2. Static buffers allow us to perform super-efficient binary search for seeking keys at every node
 3. In 2019, memory is cheaper than it used to be in 1970; no reason to over-optimize!

Are buffers a good storage idea?

- **Issue:** Static storage buffers for keys and references waste memory space for nodes.
 - Proposal: Use linked lists instead!
- This is actually not a good idea.
- Counter-arguments:
 1. Buffers waste space **MINIMAL** compared to the size of the DB (recall earlier results)
 2. Static buffers allow us to perform super-efficient binary search for seeking keys at every node
 3. In 2019, memory is cheaper than it used to be in 1970; no reason to over-optimize!
- Conclusion: yes, they are an excellent storage choice.

