

Red-Black Trees

CMSC 420

Resources

- We will be discussing [“left-leaning” Red-Black Binary Search Trees](#).
- Sedgwick & Wayne are a great resource (chapter 3.3).
- Slides’ programming model based on theirs.



java.util.TreeMap

compact1, compact2, compact3

java.util

Class **TreeMap**<K,V>

```
java.lang.Object
  java.util.AbstractMap<K,V>
    java.util.TreeMap<K,V>
```

Type Parameters:

K - the type of keys maintained by this map

V - the type of mapped values

All Implemented Interfaces:

`Serializable`, `Cloneable`, `Map<K,V>`, `NavigableMap<K,V>`, `SortedMap<K,V>`

```
public class TreeMap<K,V>
extends AbstractMap<K,V>
implements NavigableMap<K,V>, Cloneable, Serializable
```

A Red-Black tree based `NavigableMap` implementation. The map is sorted according to the `natural ordering` of its keys, or by a `Comparator` provided at map creation time, depending on which constructor is used.

This implementation provides guaranteed log(n) time cost for the `containsKey`, `get`, `put` and `remove` operations. Algorithms are adaptations of those in Cormen, Leiserson, and Rivest's *Introduction to Algorithms*.

The Treemap

compact1, compact2, compact3

java.util

Class `TreeMap<K,V>`

java.lang.Object
 java.util.AbstractMap<K,V>
 java.util.TreeMap<K,V>


Type Parameters:

K - the type of keys maintained by this map

V - the type of mapped values

All Implemented Interfaces:

Serializable, Cloneable, Map<K,V>, NavigableMap<K,V>, SortedMap<K,V>



```
public class TreeMap<K,V>  
extends AbstractMap<K,V>  
implements NavigableMap<K,V>, Cloneable, Serializable
```

A Red-Black tree based `NavigableMap` implementation. The map is sorted according to the natural ordering of its keys, or by a `Comparator` provided at map creation time, depending on which constructor is used.

This implementation provides guaranteed $\log(n)$ time cost for the `containsKey`, `get`, `put` and `remove` operations. Algorithms are adaptations of those in Cormen, Leiserson, and Rivest's *Introduction to Algorithms*.

Utility

- We will implement 2-3 trees with a **binary tree**!
- So we get **search** and the top-down parts of insertion and deletion for free.
 - So we also get the **top-down parts** of **insertion and deletion for free**!
- Idea: We will use **red links** to represent **3-nodes**, and **black links** to represent the **regular links between the nodes in a 2-3-Tree**.
 - It helps Jason to imagine those **as some sort of superglue**.

Caveats


1. We will be talking about so-called “Left-leaning” Red-Black Trees.
 - Those are easier to implement than “traditional” Red-Black Trees and maintain the same theoretical properties.
2. Red-Black Trees do **not** implement 2-3 tree key rotations!
 - That is, when a node overflows or underflows, the relevant 2-3 tree would “explode” (resp. “vanish”) the node **without looking at its siblings for space to hold a key.**

Caveats

3. We will **not** talk about “hard” deletion from a **Red**-Black Tree.
 - “Soft” deletion: Also known as “Mark-and-sweep”. Mark nodes “dead” by setting a bit, and at periodic intervals, build a new tree by traversing the existing tree and inserting only the “live” nodes into the new tree.
 - ArrayLists do mark-and-sweep!
 - “Hard” deletion: our familiar semantics: actual key or node removal from the structure.
 - Very hard in RBBSTs, and actually allows for key rotations (which we have learned in 2-3 trees)

Representation


2-3 Tree




Red-Black Tree

Representation

2-3 Tree




Red-Black Tree




Representation

2-3 Tree




Pointers transferred as **black** pointers!

Red-Black Tree





Representation

2-3 Tree




While 3-nodes are implemented as **left-leaning red links!**

Red-Black Tree




Representation

2-3 Tree



While 3-nodes are implemented as **left-leaning red links!**


Red-Black Tree



The result is a binary tree over keys!


Representation

2-3 Tree




null links are **also black**
(they exist, unlike in
threaded BSTs)

Red-Black Tree




Representation

2-3 Tree



And the **root** of the tree
is also **black**.

Red-Black Tree



An immediate benefit of RBBSTs

An immediate benefit of RBBSTs

- You can use existing BST code with virtually no modification!
 - Particularly for searches!
- In fact, it makes a lot of sense to override an existing BST class
 - You would then need to make any inner node classes in the BST class protected

An immediate drawback



- **Red links** make our tree taller ☹️.
- Compared to AVL Trees, which have height in $[\log_2 n, \log_2 n + 1]$
- Red-Black Trees have a height in $[\log_2 n, 2 \cdot \log_2 n]$
- Best case: A 2-3 tree **with only “two-nodes”** (so, a BST) will lead to a **perfectly balanced binary tree in the RBBST representation.**
- Worst case: **A 2-3 tree whose left links are all red.**
 - That'll lead to $2 \cdot \log_2 n$.

Quiz for you

- In a red-black tree, we are **guaranteed at least one node n who has two red links connected to it.**

True

False





Quiz for you

- In a red-black tree, we are **guaranteed at least one node n who has two red links connected to it.**

True

False

In fact, we have the **opposite guarantee**: there are **no nodes** that are linked to by **2 red links**.





Another quiz for you

- In a (left-leaning) red-black tree, it is **guaranteed that at least one red link ℓ is right-leaning instead of left-leaning**

True

False





Another quiz for you

- In a (left-leaning) red-black tree, it is **guaranteed that at least one red link ℓ is right-leaning instead of left-leaning**

True

False

No right-leaning red links in an RBBST!



Node structure

- We need:
 - **Links** to our children
 - **Data** field (some Comparable type T)
 - **A flag** that tells us what **the color of the link from our parent to the current node is!**
 - For its implementation to be **visible in subclasses**

Node structure

- We need:
 - **Links** to our children
 - **Data** field (some Comparable type T)
 - **A flag** that tells us what **the color of the link from our parent to the current node is!**
 - For its implementation to be **visible in subclasses**

```
protected enum Color{  
    RED, BLACK;  
}
```

```
protected class Node {  
    Node left, right;  
    T data;  
    Color color;  
}
```

Node structure

- We need:
 - **Links** to our children
 - **Data** field (some Comparable type T)
 - **A flag** that tells us what **the color of the link from our parent to the current node is!**
 - For its implementation to be **visible in subclasses**

```
protected enum Color{  
    RED, BLACK;  
}
```

```
protected class Node {  
    Node left, right;  
    T data;  
    Color color;  
}
```


This could work! 😊

Implement a utility

- During our operations, we will sometimes have **right-leaning red links** temporarily.

```
protected enum Color {  
    RED, BLACK;  
}
```

```
protected class Node {  
    Node left, right;  
    T data;  
    Color color;  
}
```




Implement a utility

- During our operations, we will sometimes have **right-leaning red links** temporarily.

```
protected enum Color {  
    RED, BLACK;  
}
```

```
protected class Node {  
    Node left, right;  
    T data;  
    Color color;  
}
```




In your notes, write a private method called `rotateLeft(Node n)` such that when we call `r = rotateLeft(r)`, the result is as shown on the right!

Implement a utility

```
private enum Color{  
    RED, BLACK;  
}
```

```
protected class Node {  
    Node left, right;  
    T data;  
    Color color;  
}
```

```
private Node rotateLeft(Node n){  
    Node temp = n.right;  
    n.right = temp.left; // How do I know that temp != null?  
    temp.left = n;  
    temp.left.color = RED;  
    return temp;  
}
```




Implement a utility

```
private enum Color{  
    RED, BLACK;  
}
```

```
protected class Node {  
    Node left, right;  
    T data;  
    Color color;  
}
```

```
private Node rotateLeft(Node n){  
    Node temp = n.right;  
    n.right = temp.left; // How do I know that temp != null?  
    temp.left = n;  
    temp.left.color = RED;  
    return temp;  
}
```




rotateRight() symmetric.

Implement another utility

- Sometimes, **during our operations**, we will end up with a node connected to **two red links**.

```
protected enum Color{  
    RED, BLACK;  
}
```

```
protected class Node {  
    Node left, right;  
    T data;  
    Color color;  
}
```




Implement another utility

- Sometimes, during our operations, we will end up with a node connected to two red links.

```
protected enum Color{  
    RED, BLACK;  
}
```

```
protected class Node {  
    Node left, right;  
    T data;  
    Color color;  
}
```



Write a void Java routine called `flipColors` which, when called as `flipColors(r)`, will transform the colors of the tree rooted at `r` as shown above!


Implement another utility

- Sometimes, **during our operations**, we will end up with a node connected to **two red links**.

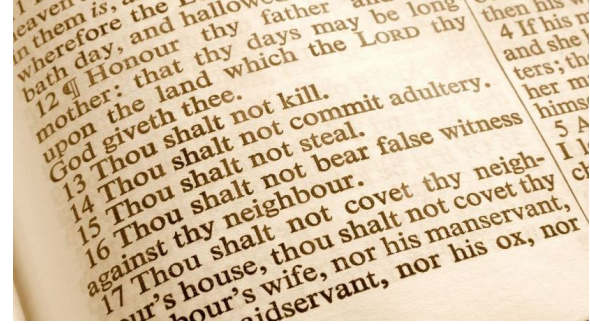
```
protected enum Color{  
    RED, BLACK;  
}
```

```
protected class Node {  
    Node left, right;  
    T data;  
    Color color;  
}
```

```
protected void flipColors(Node n){  
    n.left.color = n.right.color = BLACK;  
    n.color = RED;  
}
```




Red-Black Tree Invariants



1. Root is always **Black**
2. All **red** links point to **left**
3. Any given node has at most one **red** link only (either **to** it or **from** it)
 - Only when we model 2-3 trees.
4. Perfect **black link** balance: For **any** given leaf node ***n***, all paths from the root to one such node have the same number of **black** links.
 - This shouldn't be surprising, since **black** links correspond **exactly** to the edges of a 2-3 tree!

Maintaining our invariants




1. Root is always **Black**

- After **every** insertion, we will make sure the last call will be explicitly making the root **black**!

2. All **red** links point to **left**

- If we create a fresh node which is connected to its parent with a red link, **but it's the parent's right child, we will rotate the parent left** to make the **red** link point to the left

Maintaining our invariants



3. Any given node has one **red** link only (either **to** it or **from** it)

➤ We will see **three ways** to rectify this. Recall the following:

- a) To have a node with more than one **red** link means that we have a 4-node, which **must** be dealt with!
- b) Only way to deal with this node is to split it, since **Red-Black Trees cannot implement key rotations** (the relevant 2-3 tree “splits” aggressively).

4. Perfect **black link** balance: For **any** given leaf node ***n***, all paths from the root to one such node have the same number of **black** links.

- If this is somehow violated, **we have a bug in our implementation** ☹️

Insertion

2-3 tree

Red-Black tree

Insertion

2-3 tree




Red-Black tree



Insertion


2-3 tree

root



Red-Black tree


root



Insertion


2-3 tree

root



Red-Black tree


root



Insertion


2-3 tree

root




Red-Black tree

root




Insertion

2-3 tree




Red-Black tree



Insertion


2-3 tree

root



Red-Black tree


root




Insertion

2-3 tree

root




Red-Black tree




We always connect fresh nodes to their parents with a **red** link!

Insertion

2-3 tree




Red-Black tree




`flipcolors()` is used to exchange the link colors to represent the new 2-nodes...

Insertion

2-3 tree



Red-Black tree




`flipcolors()` is used to exchange the link colors to represent the new 2-nodes...

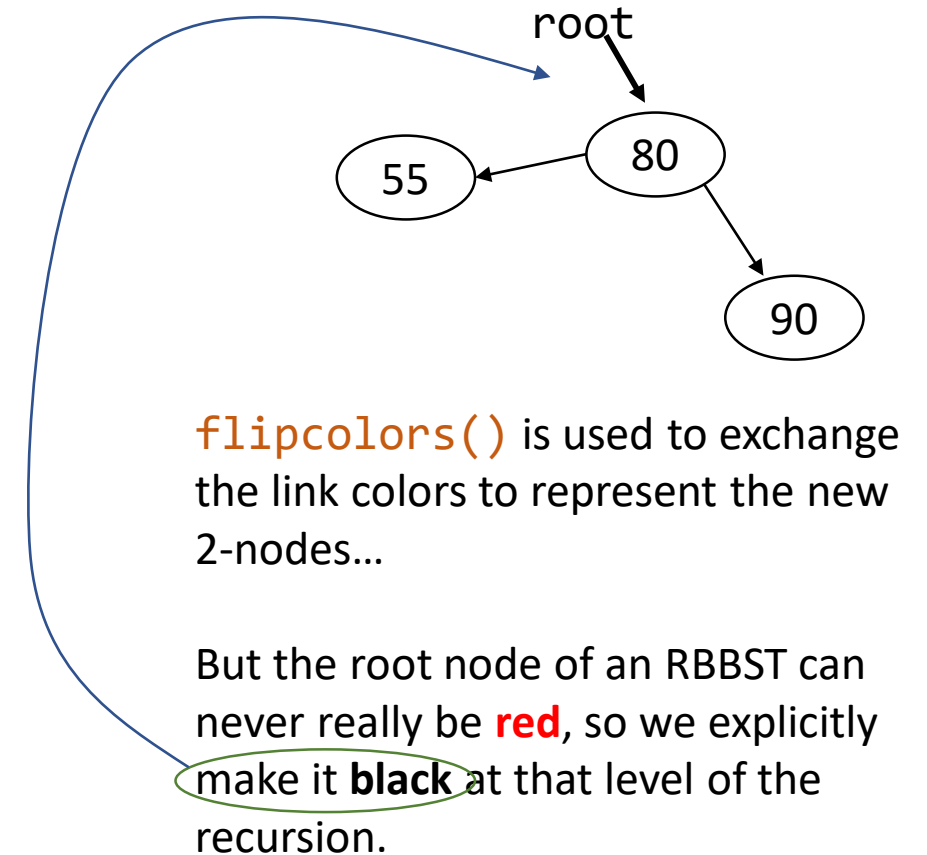
But the root node of an RBBST can never really be **red**, so we explicitly make it **black** at that level of the recursion.

Insertion

2-3 tree




Red-Black tree




Insertion

2-3 tree



```
public void insert(Key key){  
    root = insert(root, key);  
    root.color = BLACK;  
}
```

Red-Black tree




`flipcolors()` is used to exchange the link colors to represent the new 2-nodes...


But the root node of an RBBST can never really be **red**, so we explicitly make it **black** at that level of the recursion.

Insertion

2-3 tree




Red-Black tree




Insertion

2-3 tree




Red-Black tree




3-node! That's normal! 😊

Insertion

2-3 tree




Red-Black tree




Insertion

2-3 tree




Red-Black tree




Insertion

2-3 tree




Red-Black tree




Gotta
rotate!

Insertion

2-3 tree




Red-Black tree




Gotta
rotate!

Insertion

2-3 tree




Red-Black tree




Insertion

2-3 tree




Let's examine three cases for an insertion on this 3-node!

Red-Black tree




Insertion

2-3 tree



Insert 70


Red-Black tree




Insertion

2-3 tree

Insert 70




Red-Black tree



Insertion


2-3 tree



Overflow! ☹️

Insert 70


Red-Black tree



flipColors()! 😊

Insertion


2-3 tree



Insert 70

Overflow


Red-Black tree



flipColors()!


Insertion

2-3 tree




Insert 70

Red-Black tree




Insertion

2-3 tree



Insert 70


Red-Black tree



- So, this case, where we insert a key as the biggest key of a temporary 4-node, was rectified with one call to `flipColors()`.
- Notice the perfect black link balance of the Red-Black Tree?


Insertion

2-3 tree




Insert 40

Red-Black tree




Insertion

2-3 tree




Insert 40

Red-Black tree




Insertion

2-3 tree



Insert 40


Red-Black tree




Insertion

2-3 tree

Insert 40




Red-Black tree




Insertion

2-3 tree

Insert 40




Red-Black tree



So in this case (inserting as a temporary 4 node's least key), we need **1 rotation** and one **color flipping!**


Insertion

2-3 tree




Insert 60

Red-Black tree




Insertion

2-3 tree




Insert 60

Red-Black tree




Insertion

2-3 tree



Insert 60


Red-Black tree




Insertion

2-3 tree

Insert 60




Red-Black tree




Insertion

2-3 tree

Insert 60




Red-Black tree




Insertion

2-3 tree



Insert 60


Red-Black tree




In this final case (inserting as a temporary 4 node's **middle** key), we need 2 rotations and one **color flipping!**

Exercise for you!

2-3 tree




Red-Black tree




Please go ahead and insert **87** for me
please! 😊

Exercise for you!

2-3 tree




Red-Black tree




Remember: In **Red**-Black trees, we can't model key rotations! We will split the **yellow** node!

Exercise for you!

2-3 tree




Red-Black tree




Exercise for you!

2-3 tree




Red-Black tree




Exercise for you!

2-3 tree




Red-Black tree




Exercise for you!

2-3 tree




Red-Black tree




Exercise for you!

2-3 tree




Red-Black tree




Exercise for you!

2-3 tree




Don't forget to
set the root to
BLACK!

Red-Black tree



Exercise for you!

2-3 tree




Don't forget to
set the root to
BLACK!

```
public void insert(Key key){  
    root = insert(root, key);  
    root.color = BLACK;  
}
```

Speaking of...

Red-Black tree



Fill-in the insertion subroutine!

```
private Node insert(Node n, Key key){
    if(n == null) return new Node(key, RED);
    if(key.compareTo(n.key) < 0)
        n.left = insert(n.left, key);
    else if(key.compareTo(n.key) >= 0)
        n.right = insert(n.right, key);
    else return n;
    if(isRed(n.right) && _____ ) n = rotateLeft(n);
    if(isRed(n.left) && _____) n = rotateRight(n);
    if(_____ && isRed(n.right)) _____;
}
```

<pre>boolean isRed(Node n){ if(n == null) return false; else return (n.color == RED); }</pre>

Fill-in the insertion subroutine!

```
private Node insert(Node n, Key key){
    if(n == null) return new Node(key, RED);
    if(key.compareTo(n.key) < 0)
        n.left = insert(n.left, key);
    else if(key.compareTo(n.key) >= 0)
        n.right = insert(n.right, key);
    else return n;
    if(isRed(n.right) && __!isRed(n.left)__ ) n = rotateLeft(n);
    if(isRed(n.left) && __isRed(n.left.left)__ ) n = rotateRight(n);
    if(__isRed(n.left)____ && isRed(n.right)) __flipColors(n)__;
}
```

```
boolean isRed(Node n){
    if(n == null) return false;
    else return (n.color == RED);
}
```




Problem 3!



“Hard” Deletion

- By “hard” Deletion we refer to a process that actually removes the key’s wrapper, the node, from the tree.
 - We’ve already seen how to do this in threaded and non-threaded BSTs, AVL, Splay, 2-3 (conceptually).

“Hard” Deletion



- By “hard” Deletion we refer to a process that actually removes the key’s wrapper, the node, from the tree.
 - We’ve already seen how to do this in threaded and non-threaded BSTs, AVL, Splay, 2-3 (conceptually).
- Hard deletion in RBBSTs is way too hard and we will not cover it.
 - It relies on subroutines for deleting the minimum and maximum of a subtree.
 - Segdewick and Wayne show an implementation as an exercise in 3.3.

“Soft” deletion (Mark – and – Sweep)

- Instead, in practice, we perform a **Mark-and-Sweep step**.
- A key deletion does **not** lead to immediate freeing of the key’s node and re-placement of links.
 - Instead, a bit is set to indicate that the relevant object can be safely **garbage – collected**
 - Note that **since it can be garbage-collected, it can also be replaced!**
 - As long as **we are sure** we can place a node with our new key exactly where the old node used to be...
- Every T ms, a “sweeping” phase runs, **collecting only the objects with unset bits**.
 - Those objects **are re-inserted in another tree** and the root reference is assigned to that new tree **so the old one can be garbage collected**.

Some theory behind RBBSTs

- **Red**-Black Trees have perfect black link balance
 - All paths from the root to a null pointer dereference **the same # black pointers!**
 - So, if we somehow could ignore **red links** (we really shouldn't, we'd miss keys), then we'd have **excellent search all the time!**
- Remember: in classic BSTs, search can be as bad as $\mathcal{O}(n)$, for n keys and unit cost that of a pointer dereferencing.
- In RBBSTs, you are guaranteed $\mathcal{O}(\log_2 n)$ **irrespective of insertion order.**

Some theory behind RBBSTs

- **Red**-Black Trees have perfect black link balance
 - All paths from the root to a null pointer dereference **the same # black pointers!**
 - So, if we somehow could ignore **red links** (we really shouldn't, we'd miss keys), then we'd have **excellent search all the time!**
- Remember: in classic BSTs, search can be as bad as $\mathcal{O}(n)$, for n keys and unit cost that of a pointer dereferencing.
- In RBBSTs, you are guaranteed $\mathcal{O}(\log_2 n)$ **irrespective of insertion order.**



Some theory behind RBBSTs

- **Red**-Black Trees have perfect black link balance
 - All paths from the root to a null pointer dereference **the same # black pointers!**
 - So, if we somehow could ignore **red links** (we really shouldn't, we'd miss keys), then we'd have **excellent search all the time!**
- Remember: in classic BSTs, search can be as bad as $\mathcal{O}(n)$, for n keys and unit cost that of a pointer dereferencing.
- In RBBSTs, you are guaranteed $\mathcal{O}(\log_2 n)$ **irrespective of insertion order.**
 - *DIDN'T WE ALSO HAVE THIS IN AVL TREES?*




Some theory behind RBBSTs

- It turns out that it is **only rarely** that **R**BBSTs will have **anything close to $2 \log_2 n$ height!**
- To minimize the height, we need to minimize the **red** links.

Average height in RBBSTs

- Turns out that in practice, RBBSTs look more like this...



AVL vs Red-Black

AVL		Red-Black	
+	-	+	-
Height $\mathcal{O}(\log_2 n)$	Spatial overhead	Height $\mathcal{O}(\log_2 n)$	Height can be greater than $\lfloor \log_2 n \rfloor$
Easier (hard) deletions		Can help us implement both 2-3 and 2-3-4 trees	Hard deletion hard to implement