

# List arrays and Skiplists

CMSC 420

# Central question

- How can I do binary search in a sorted linked list?

# Central question

- How can I do binary search in a sorted linked list?
- Naive solution has complexity  $\mathcal{O}(n)$ , which does not improve on linear search 😞

# Central question

- How can I do **binary search in a sorted linked list?**
- Naive solution has complexity  $\mathcal{O}(n)$ , which does not improve on linear search 😞
  - When you have  $n$  nodes, you scan  $\frac{n}{2}$  nodes to find the middle element.
  - When you have  $\frac{n}{2}$ , you scan  $\frac{n}{4}$ ....
  - When you have  $\frac{n}{4}$ , you scan  $\frac{n}{8}$ ....
  - ...
  - When you have  $1 = \frac{n}{n} = \frac{n}{2^{\log_2 n}}$ , you scan  $1 = \frac{n}{\log_2 n}$  nodes....

# Central question

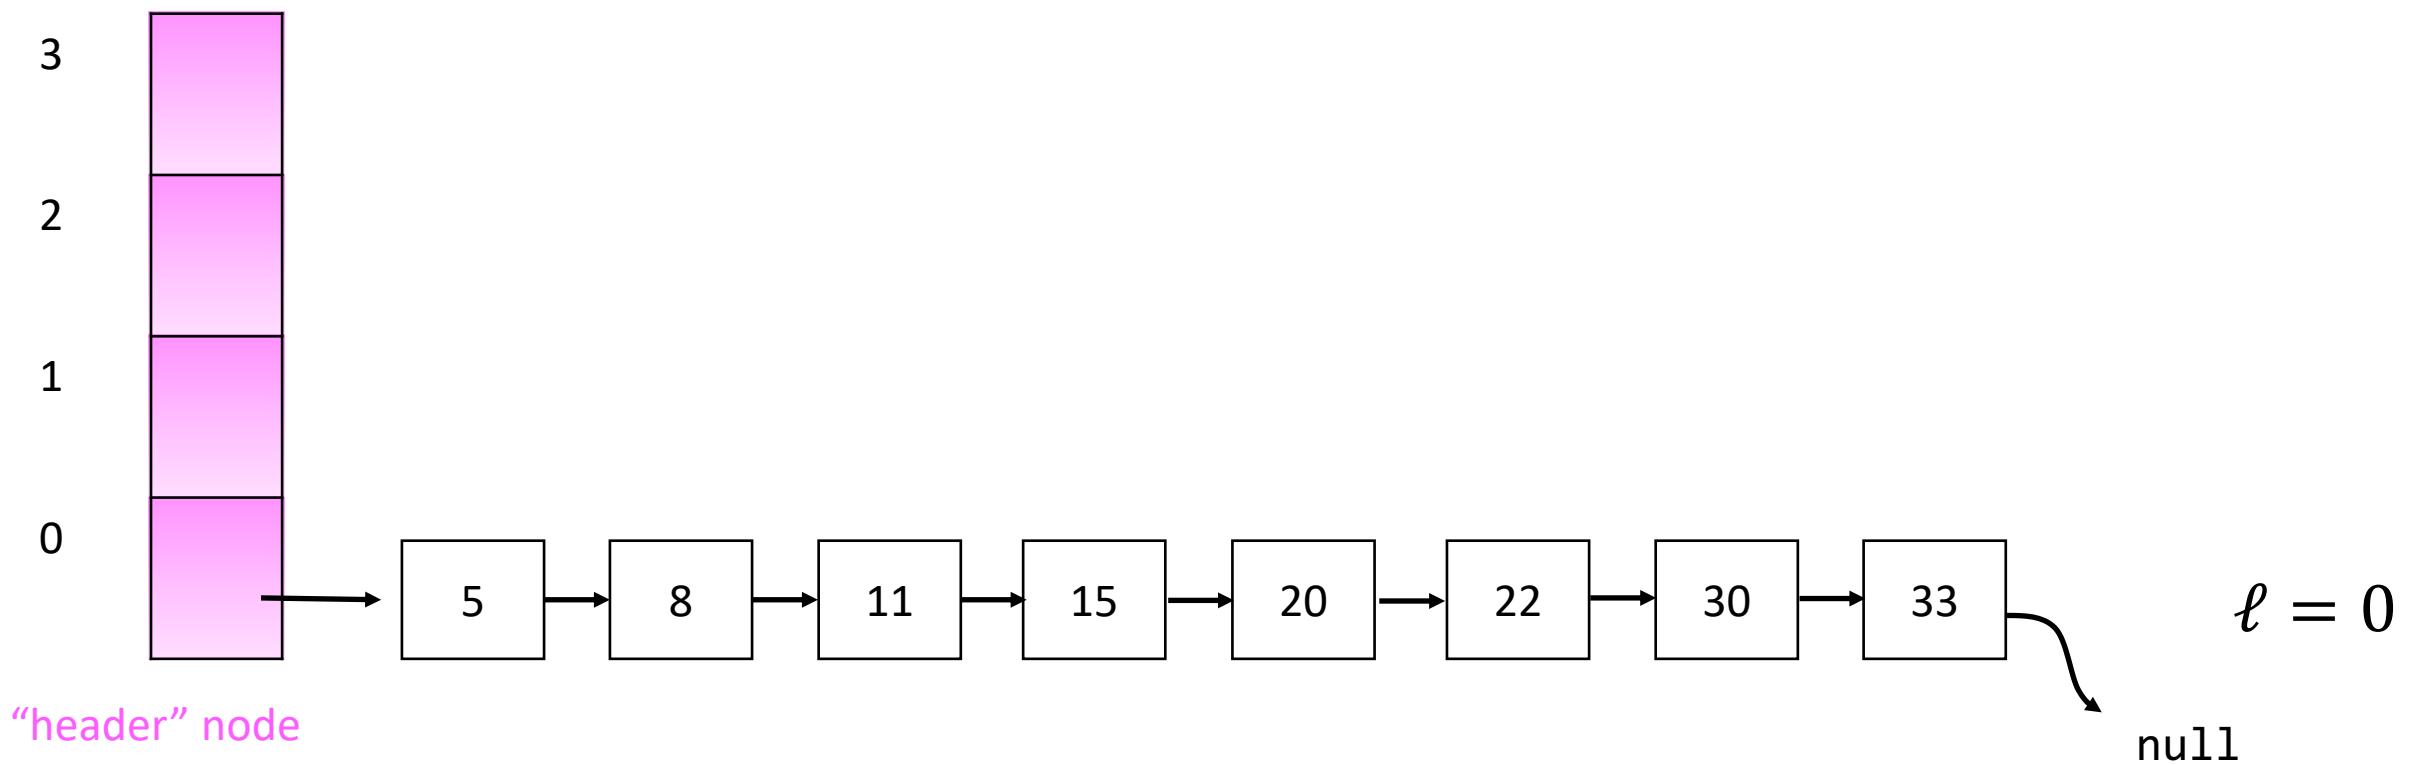
- How can I do **binary search in a sorted linked list?**
- Naive solution has complexity  $\mathcal{O}(n)$ , which does not improve on linear search 😞
  - When you have  $n$  nodes, you scan  $\frac{n}{2}$  nodes to find the middle element.
  - When you have  $\frac{n}{2}$ , you scan  $\frac{n}{4}$ ....
  - When you have  $\frac{n}{4}$ , you scan  $\frac{n}{8}$ ....
  - ...
  - When you have  $\frac{n}{2^{\log_2 n}}$ , you scan  $\frac{n}{n} = 1$  nodes....

$$\begin{aligned}\text{Total #nodes scanned} &= \frac{n}{2} + \frac{n}{4} + \dots + \frac{n}{2^{\log_2 n}} = \\ n \left( \frac{1}{2} + \frac{1}{4} + \dots + \frac{1}{2^{\log_2 n}} \right) &= \mathcal{O}(n)\end{aligned}$$

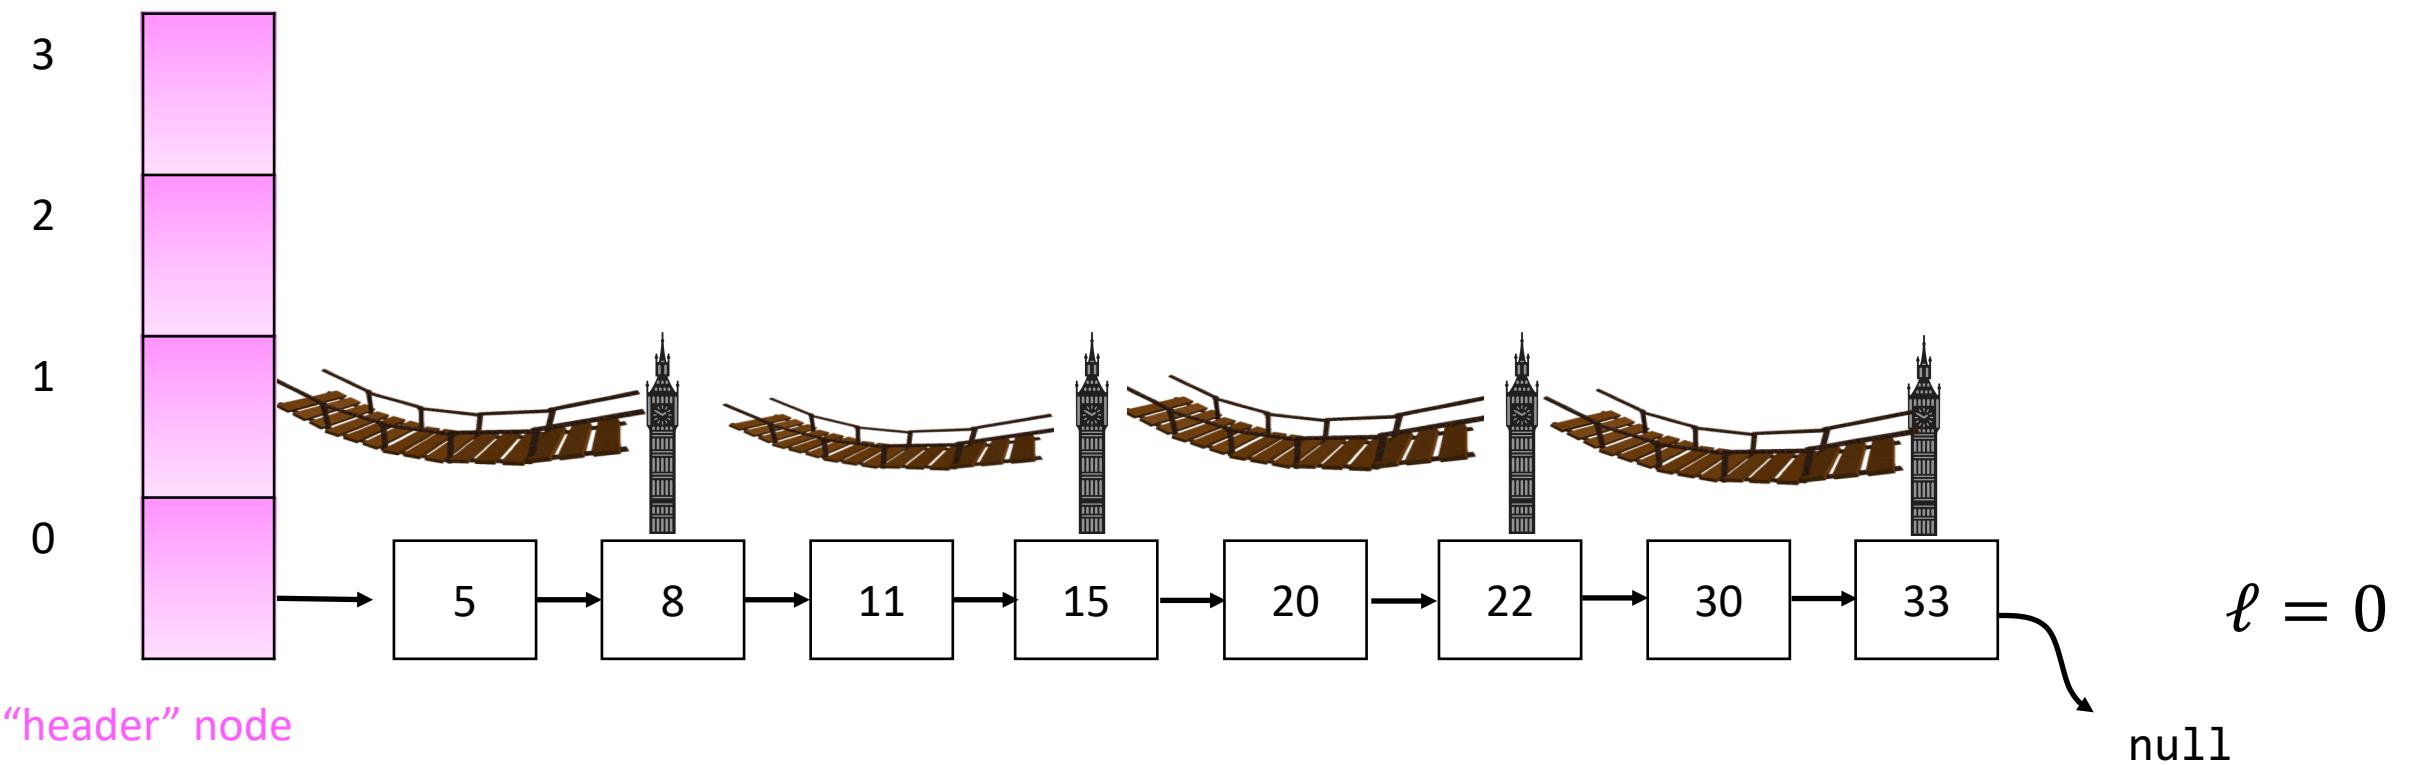
# Central question

- Idea to improve complexity: Consider every node to be a “tower” consisting of a certain number of levels, different for each node.
  - The levels themselves will be connected by linked lists
    - Think of them as bridges!
  - The “tallest” tower will have height  $\log_2 n$ !
  - First level will be the original list: connects every  $2^0 = 1$  nodes.
  - Second level will connect every  $2^1 = 2$  nodes.
  - ....
  - $i^{th}$  level will connect every  $2^i$  nodes.

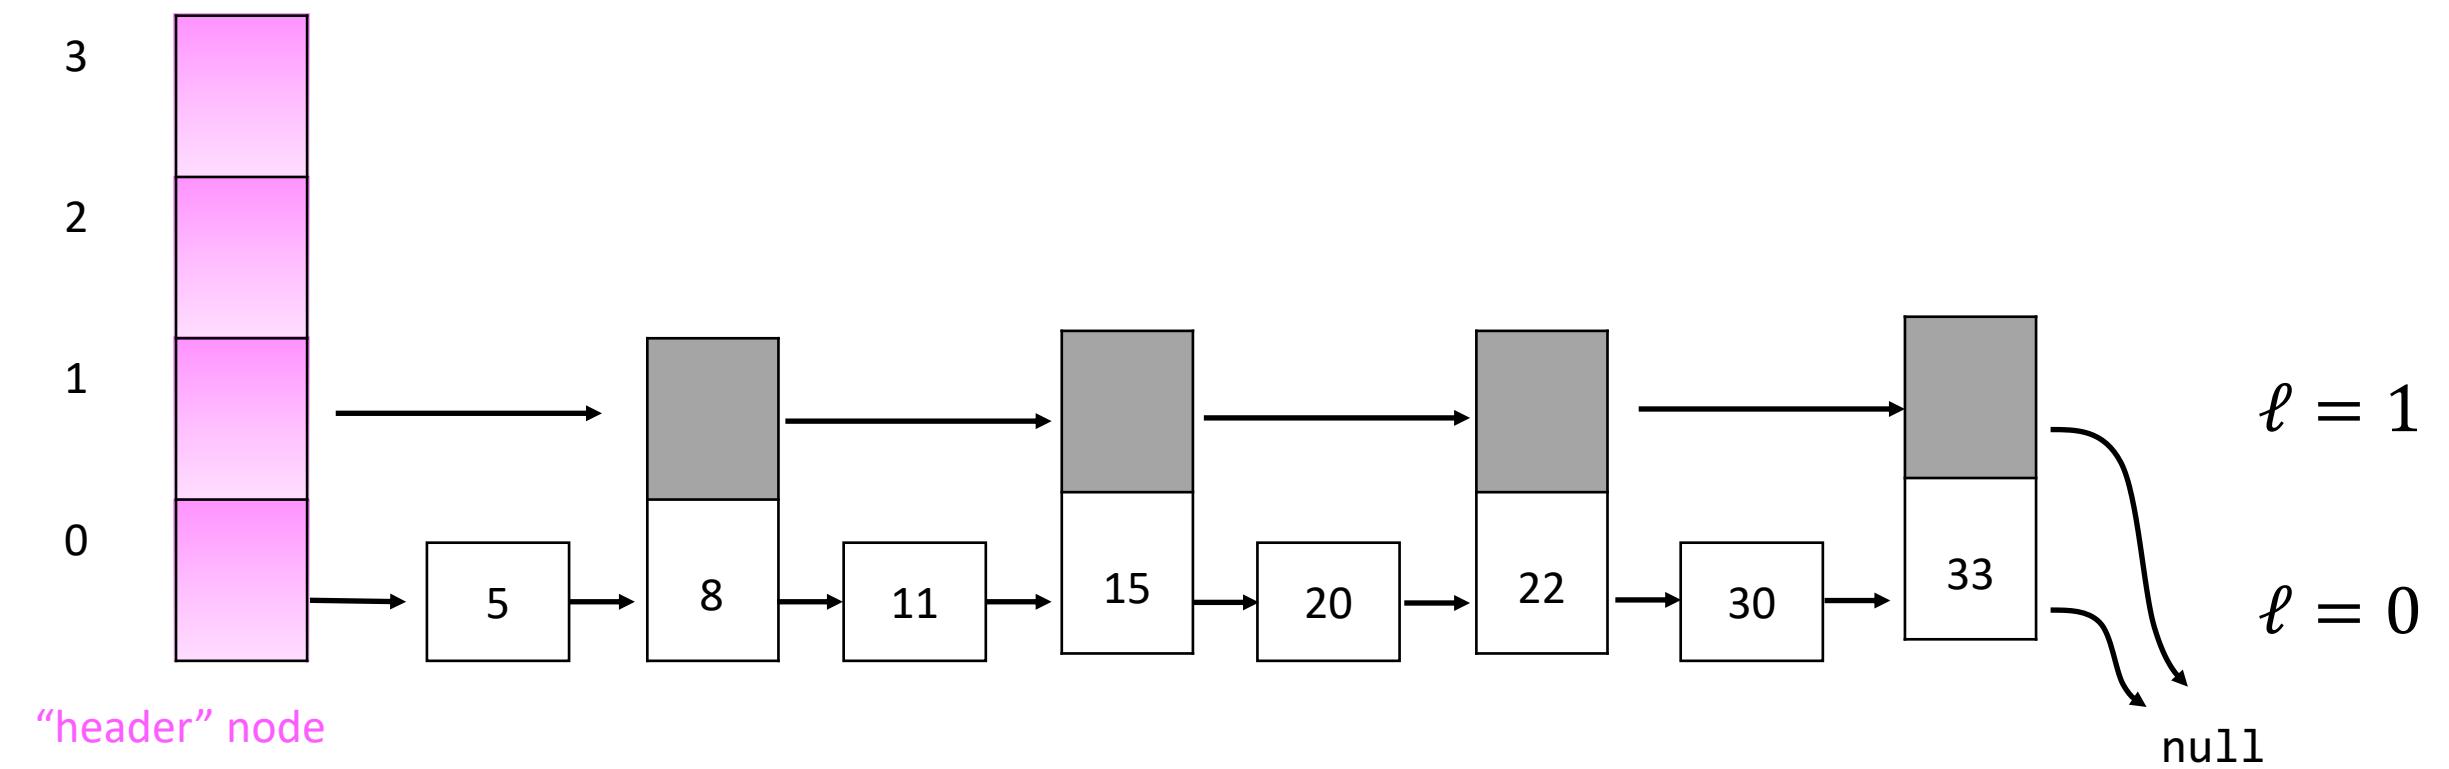
# List Arrays



# List Arrays

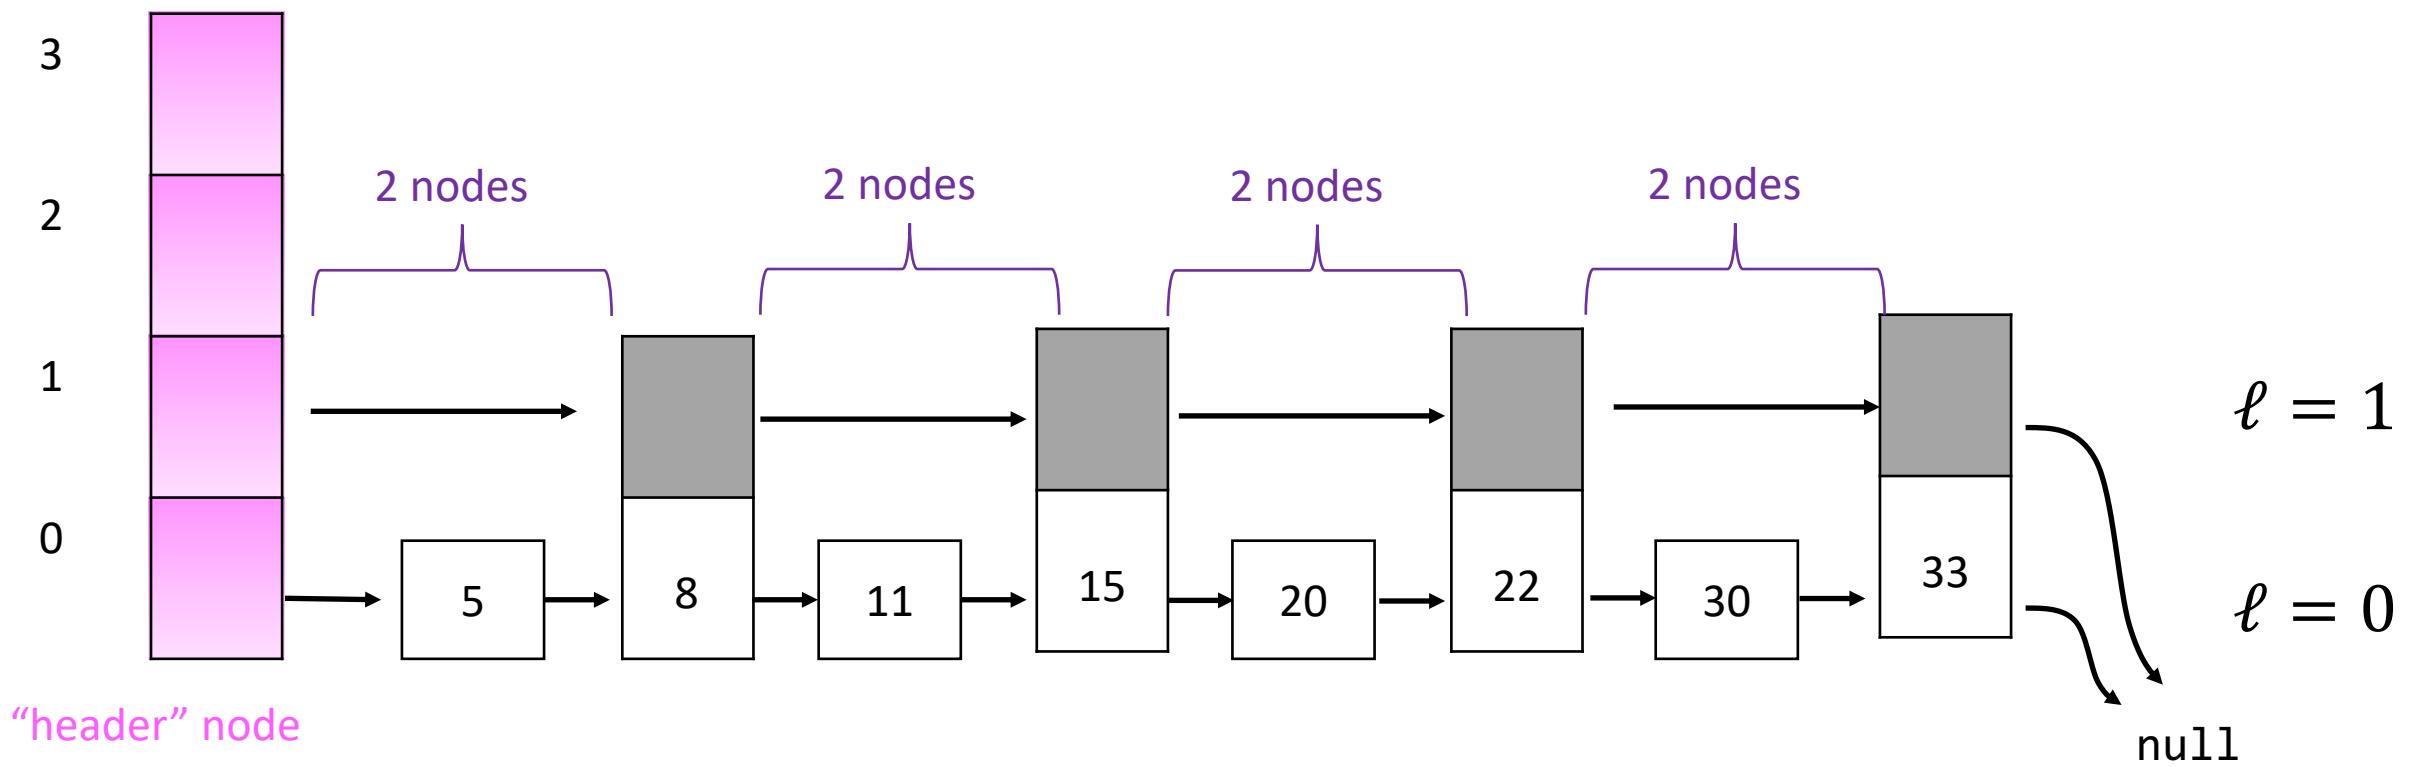


# List Arrays

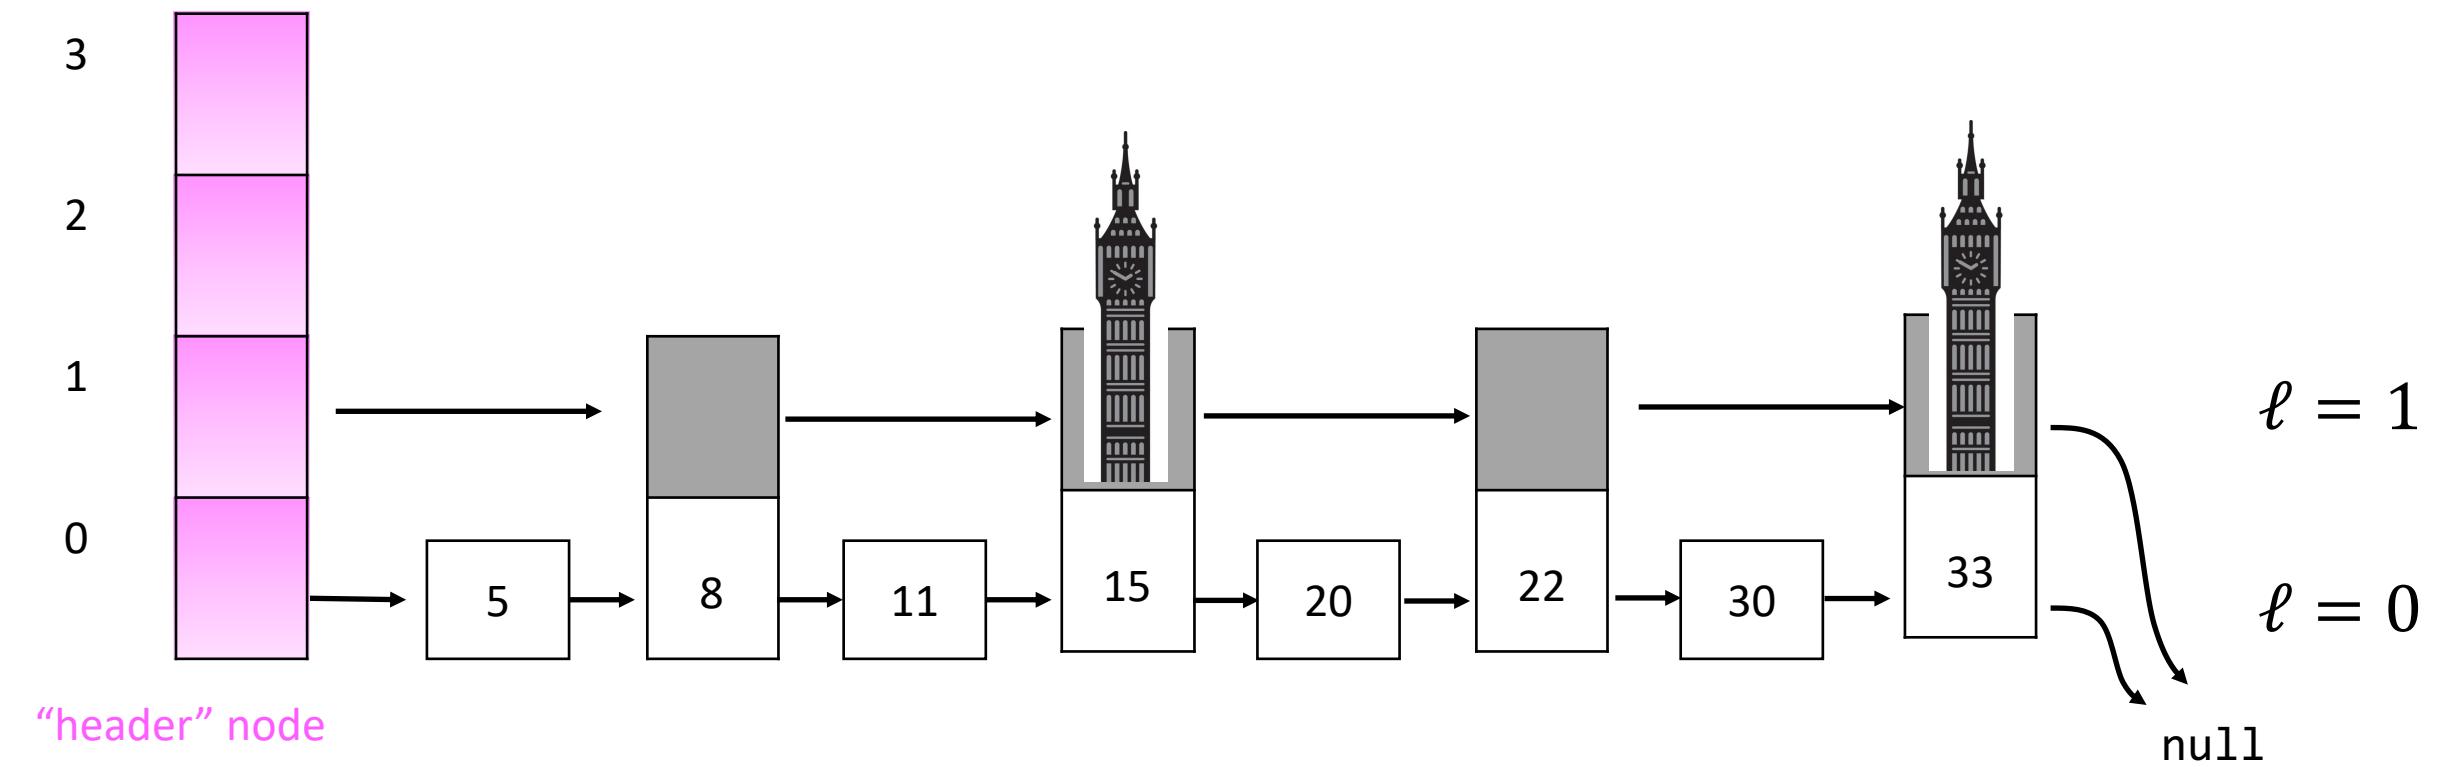


# List Arrays

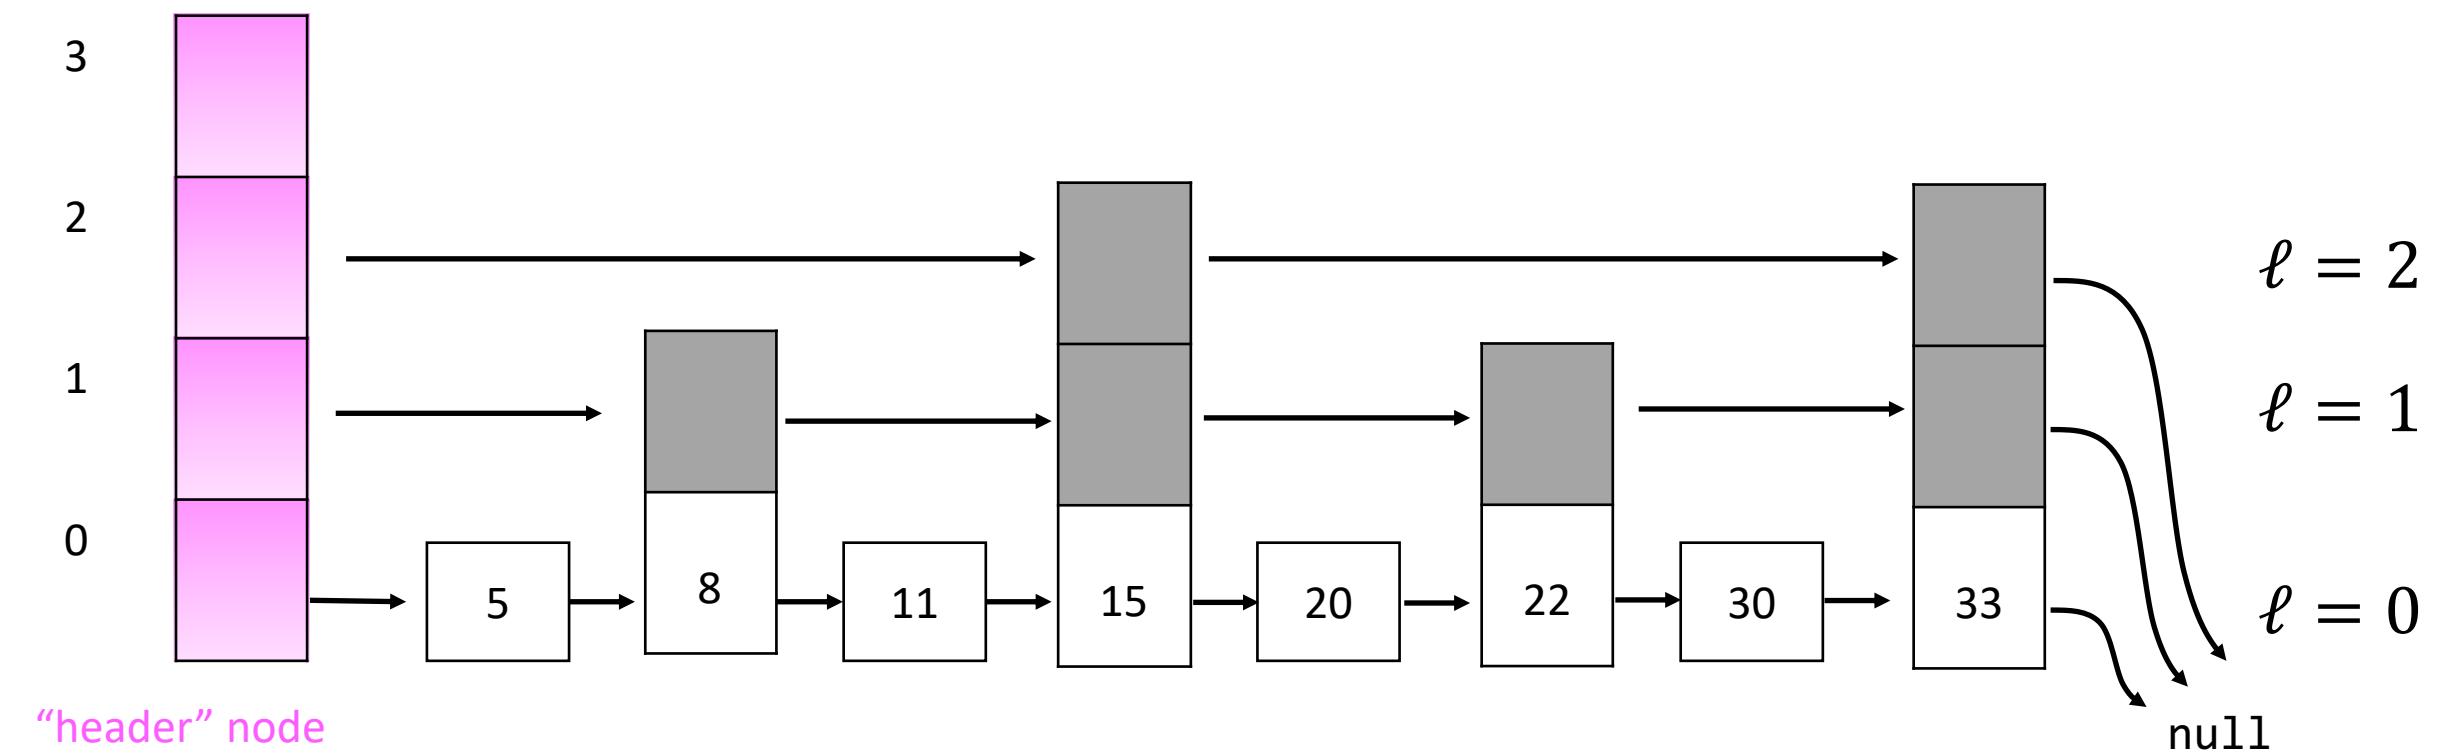
Second level ( $\ell = 1$ ) connects every  $2^1 = 2$  nodes!



# List Arrays

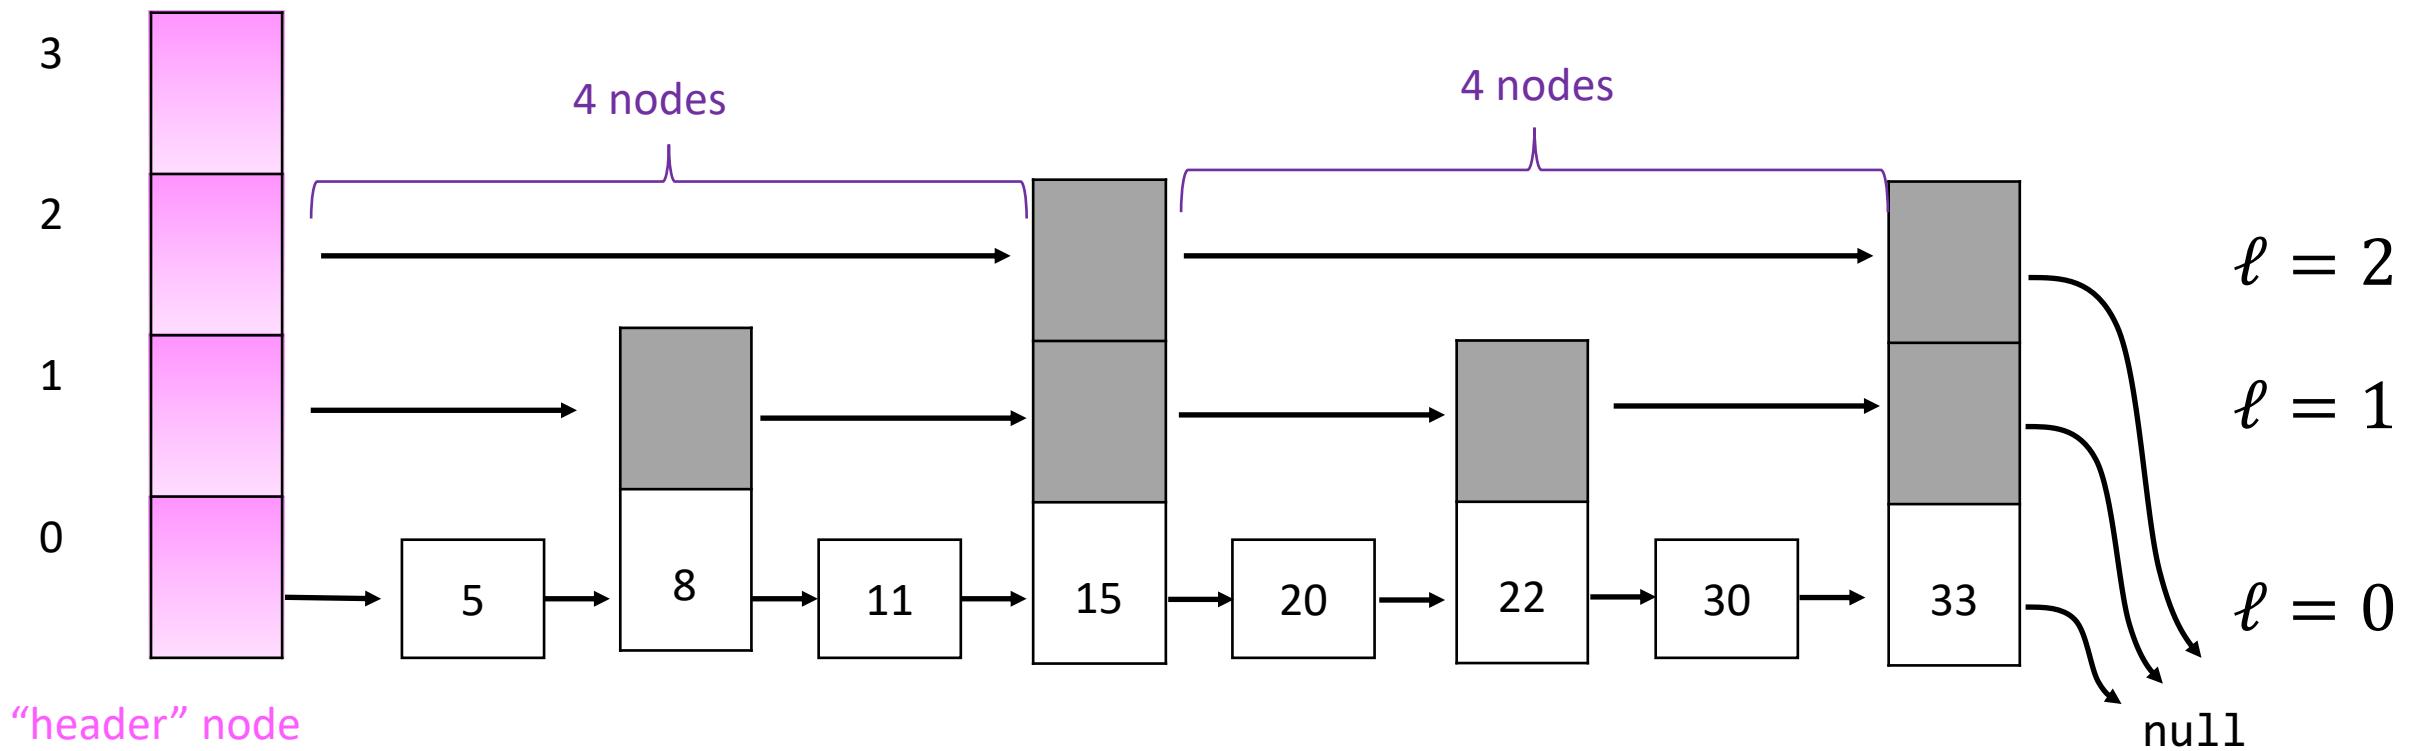


# List Arrays

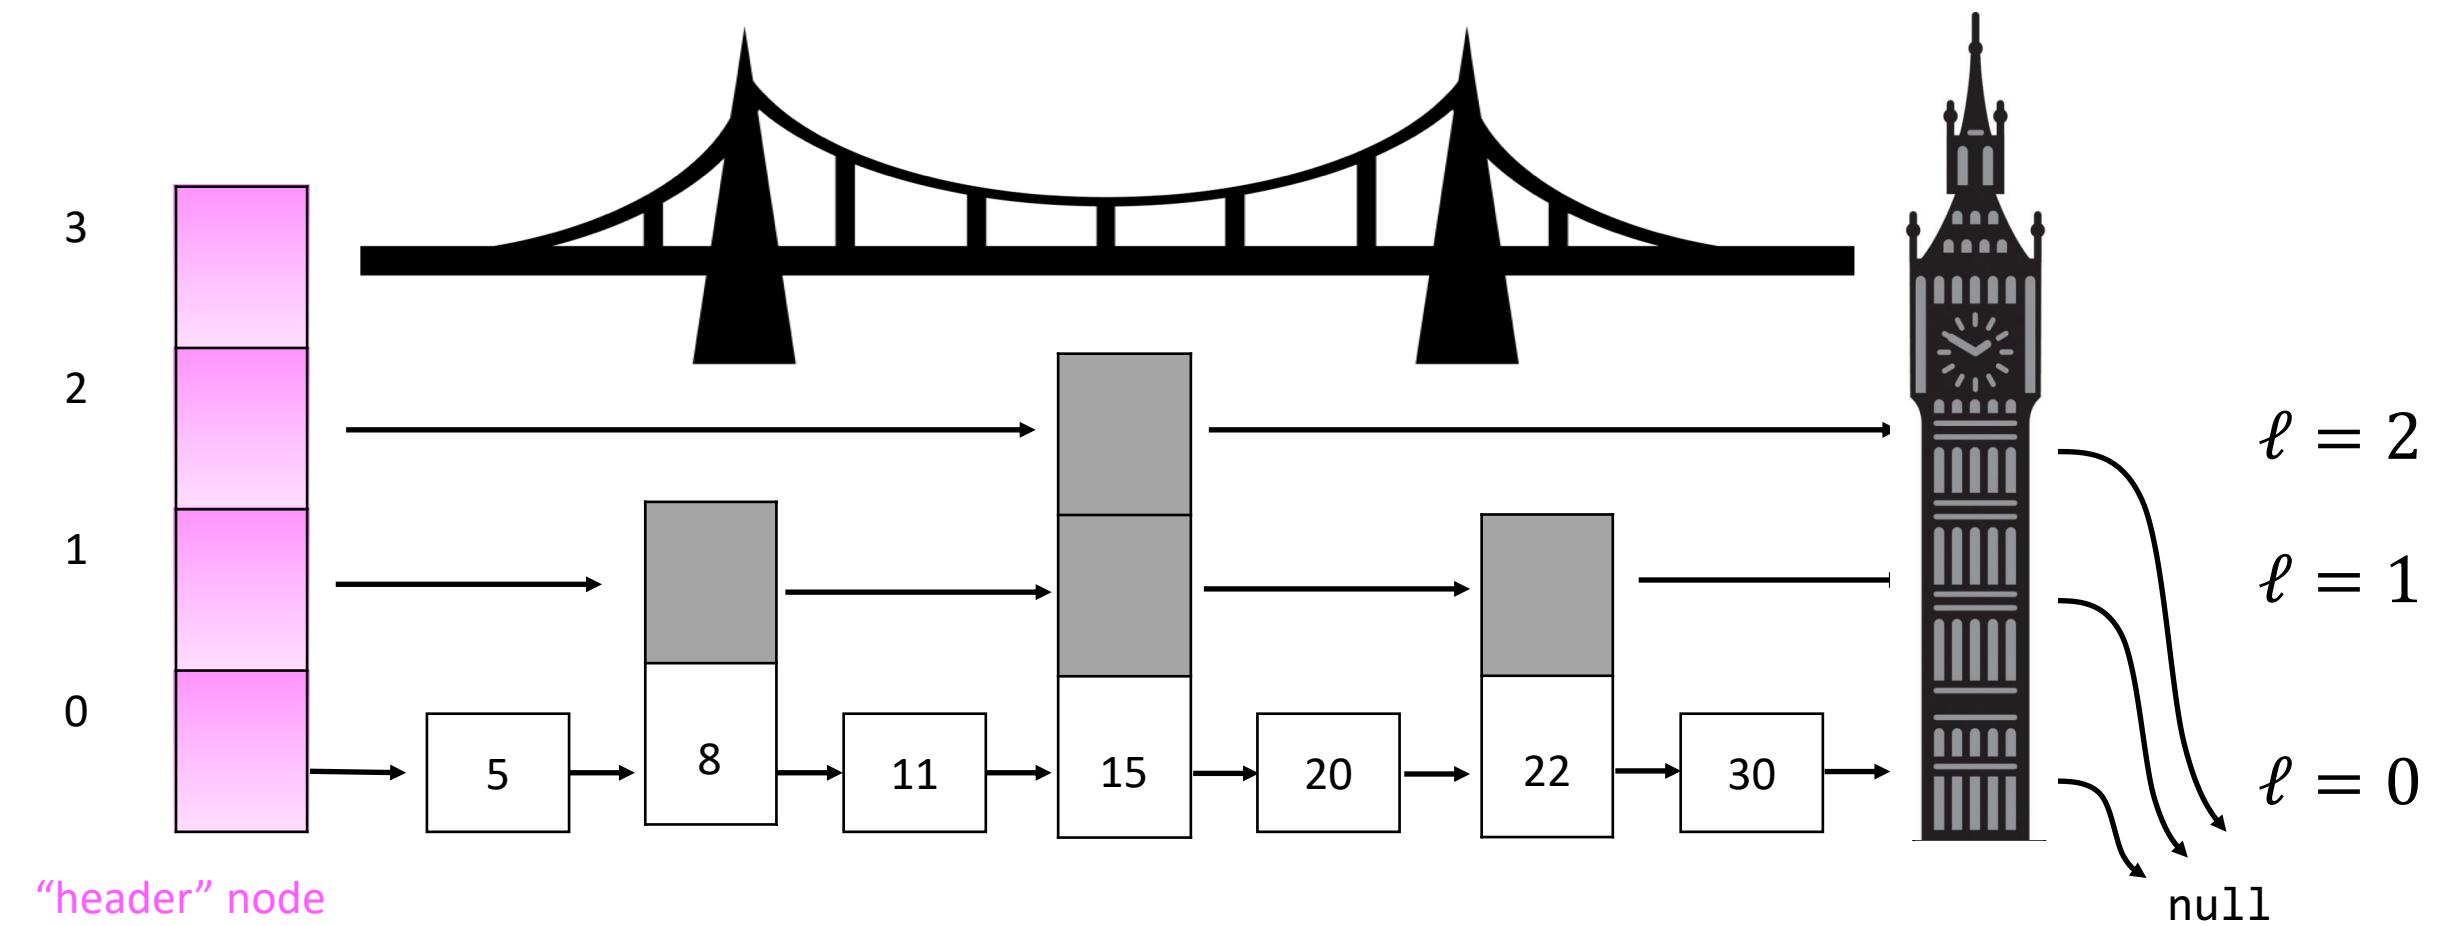


# List Arrays

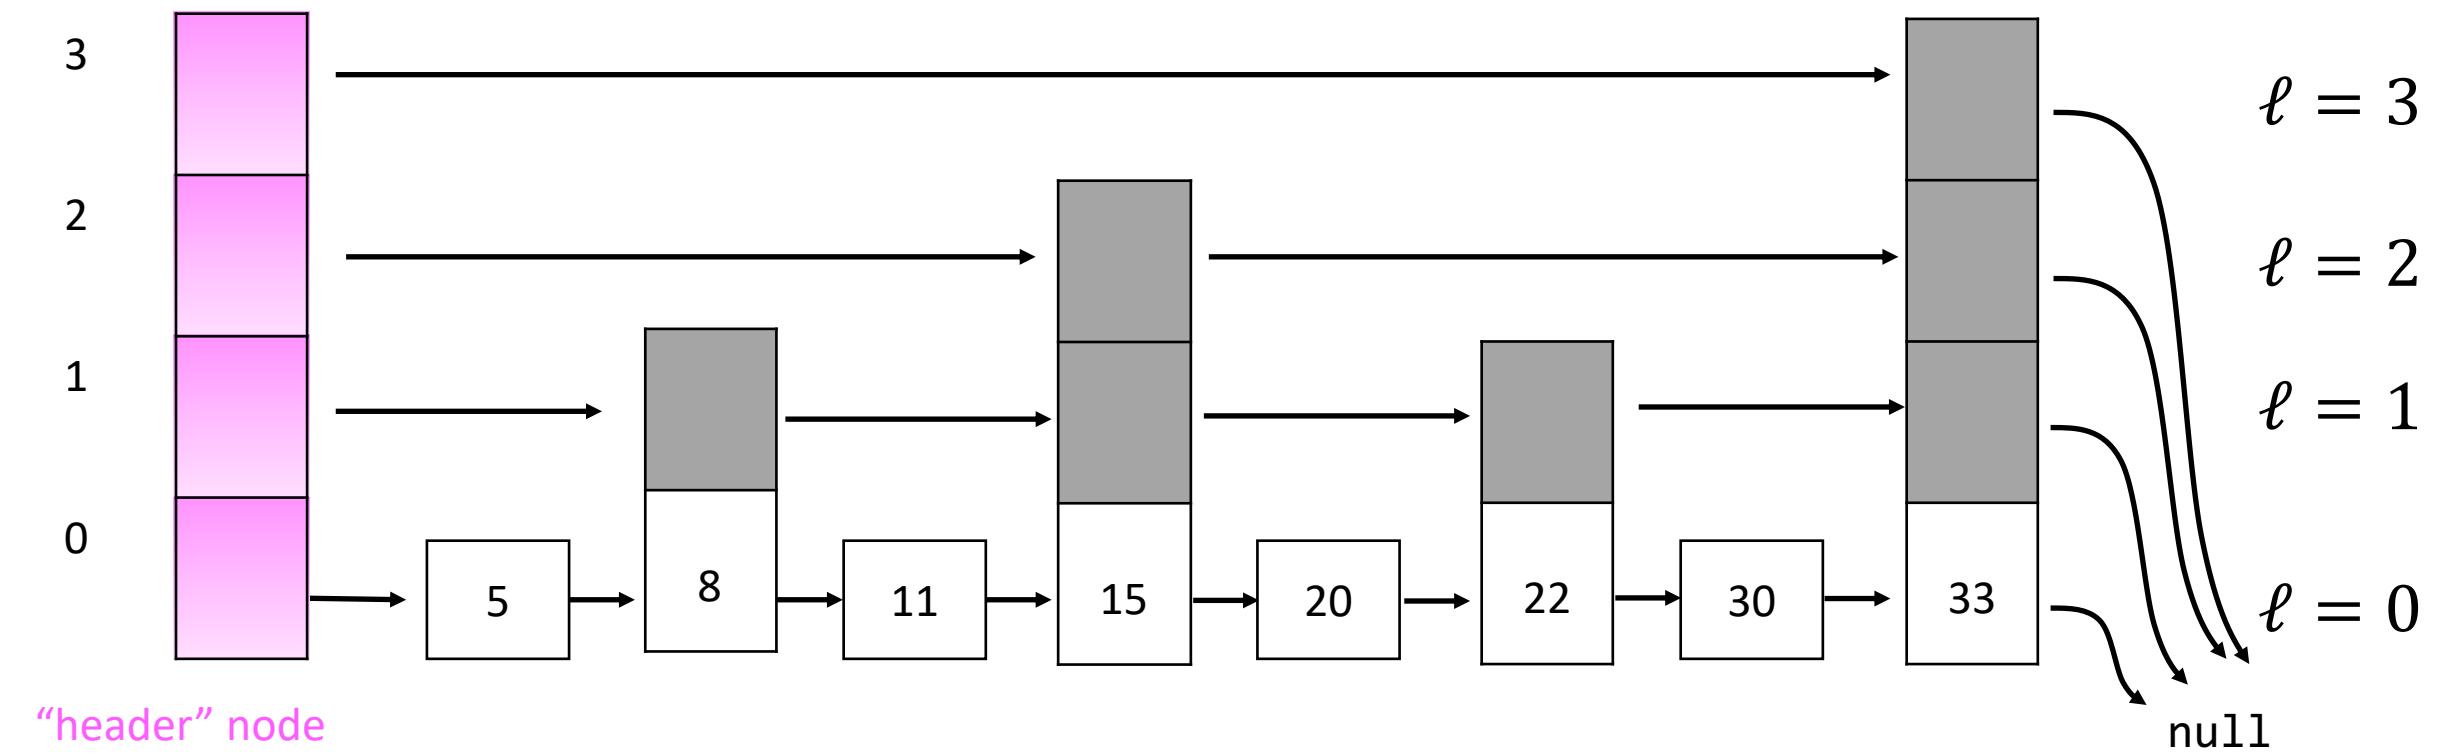
Third level ( $\ell = 2$ ) connects every  $2^2 = 4$  nodes!



# List Arrays

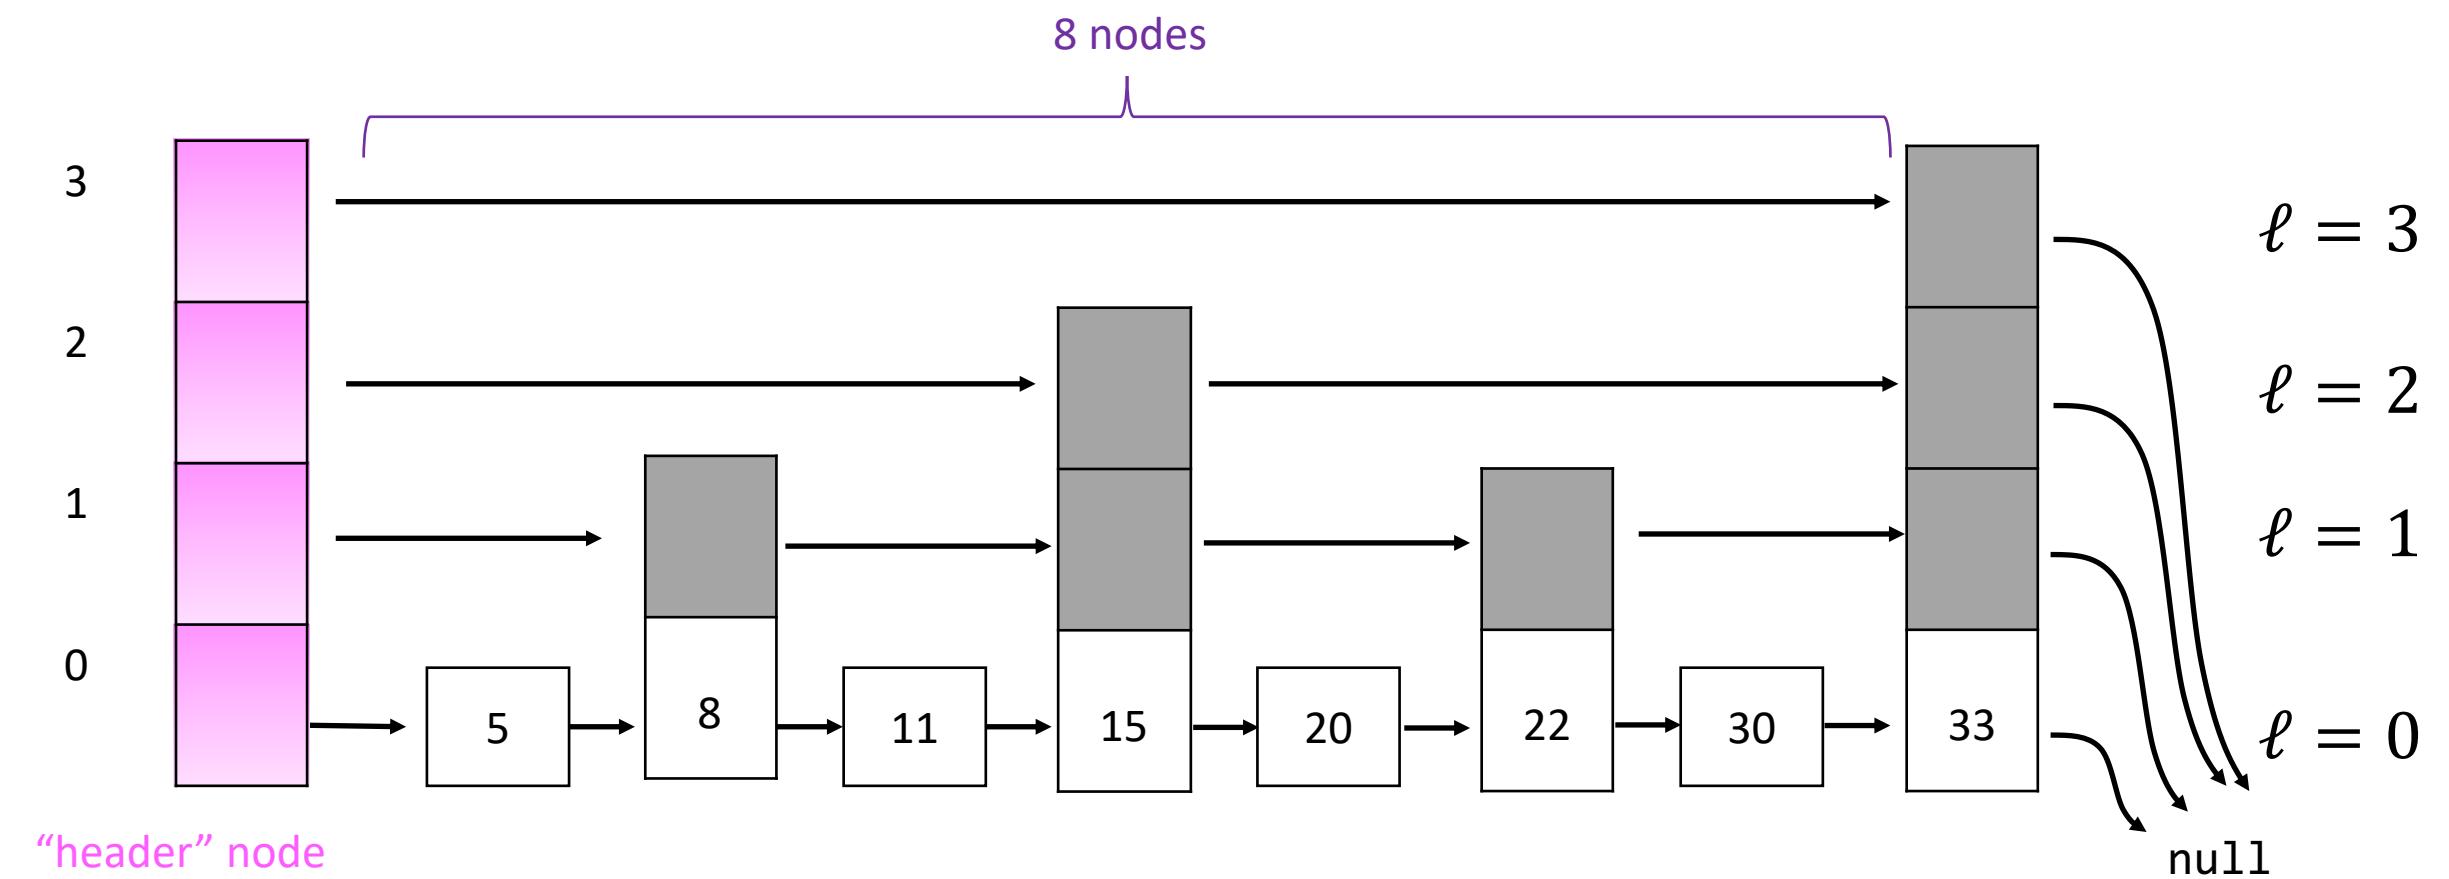


# List Arrays



# List Arrays

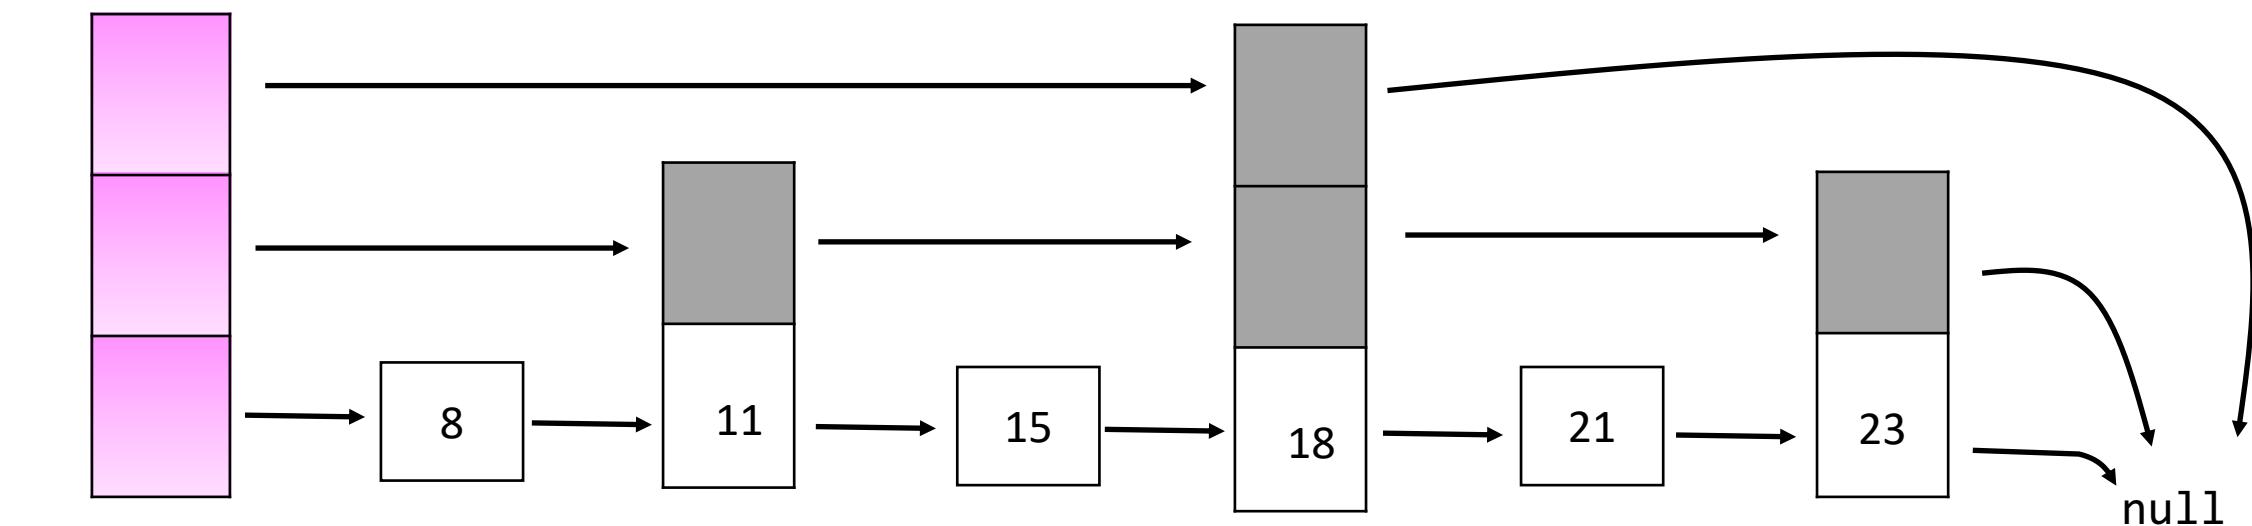
Fourth level ( $\ell = 3$ ) connects every  $2^3 = 8$  nodes!



# Extension to other lengths

- What if my list has a length that is not a power of 2?
- No problem: we then have  $\lceil \log_2 n \rceil$  levels. Example:

$$(n = 6) \Rightarrow (\lceil \log_2 n \rceil = 3)$$

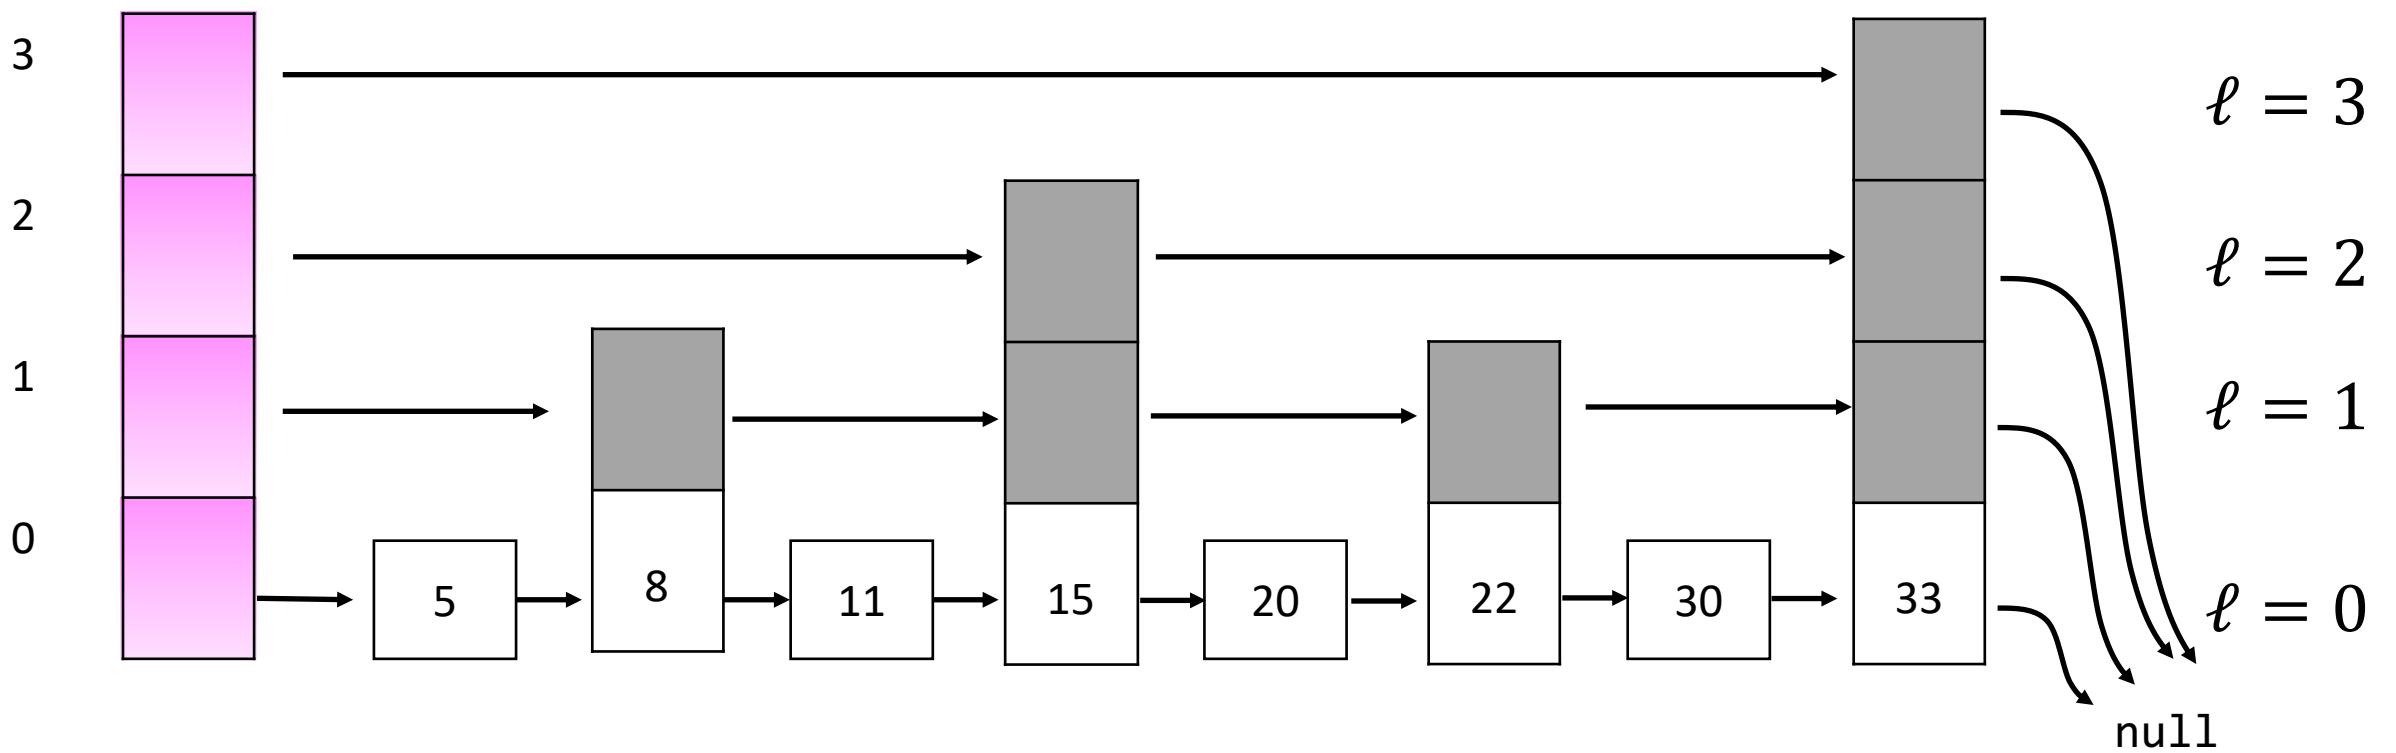


# Searching a list array

- Every one of the individual arrays has the data element stored at position 0.
- Searching works as follows:
  1. Begin at **header[levels – 1]**. Traverse the current level until:
    - a) You find the element (**search hit, congrats**)
    - b) You find a tower **whose element is greater**, or:
    - c) You hit **null**.
  2. If either (b) or (c) happens, descend to the level below and goto (1).
    - If there is no level below, you have a **search miss**. At this point you should consider how this makes you feel.

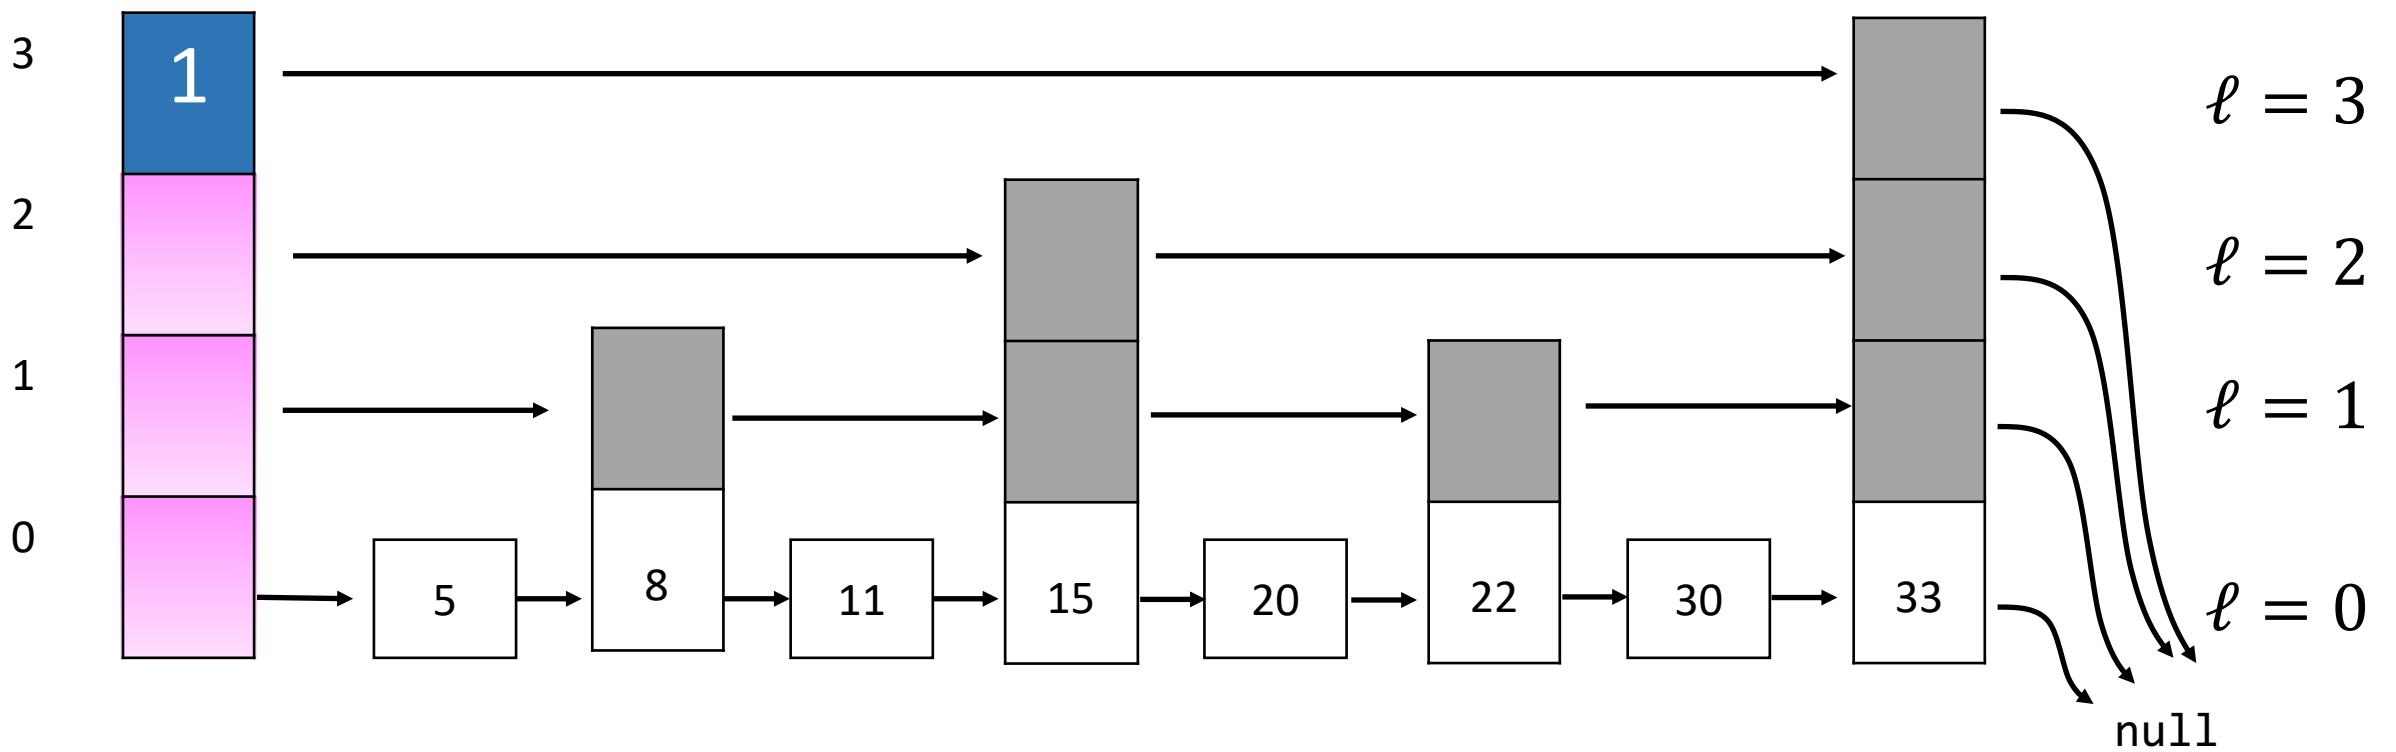
# Search example

Search for “20” – should be a successful search



# Search example

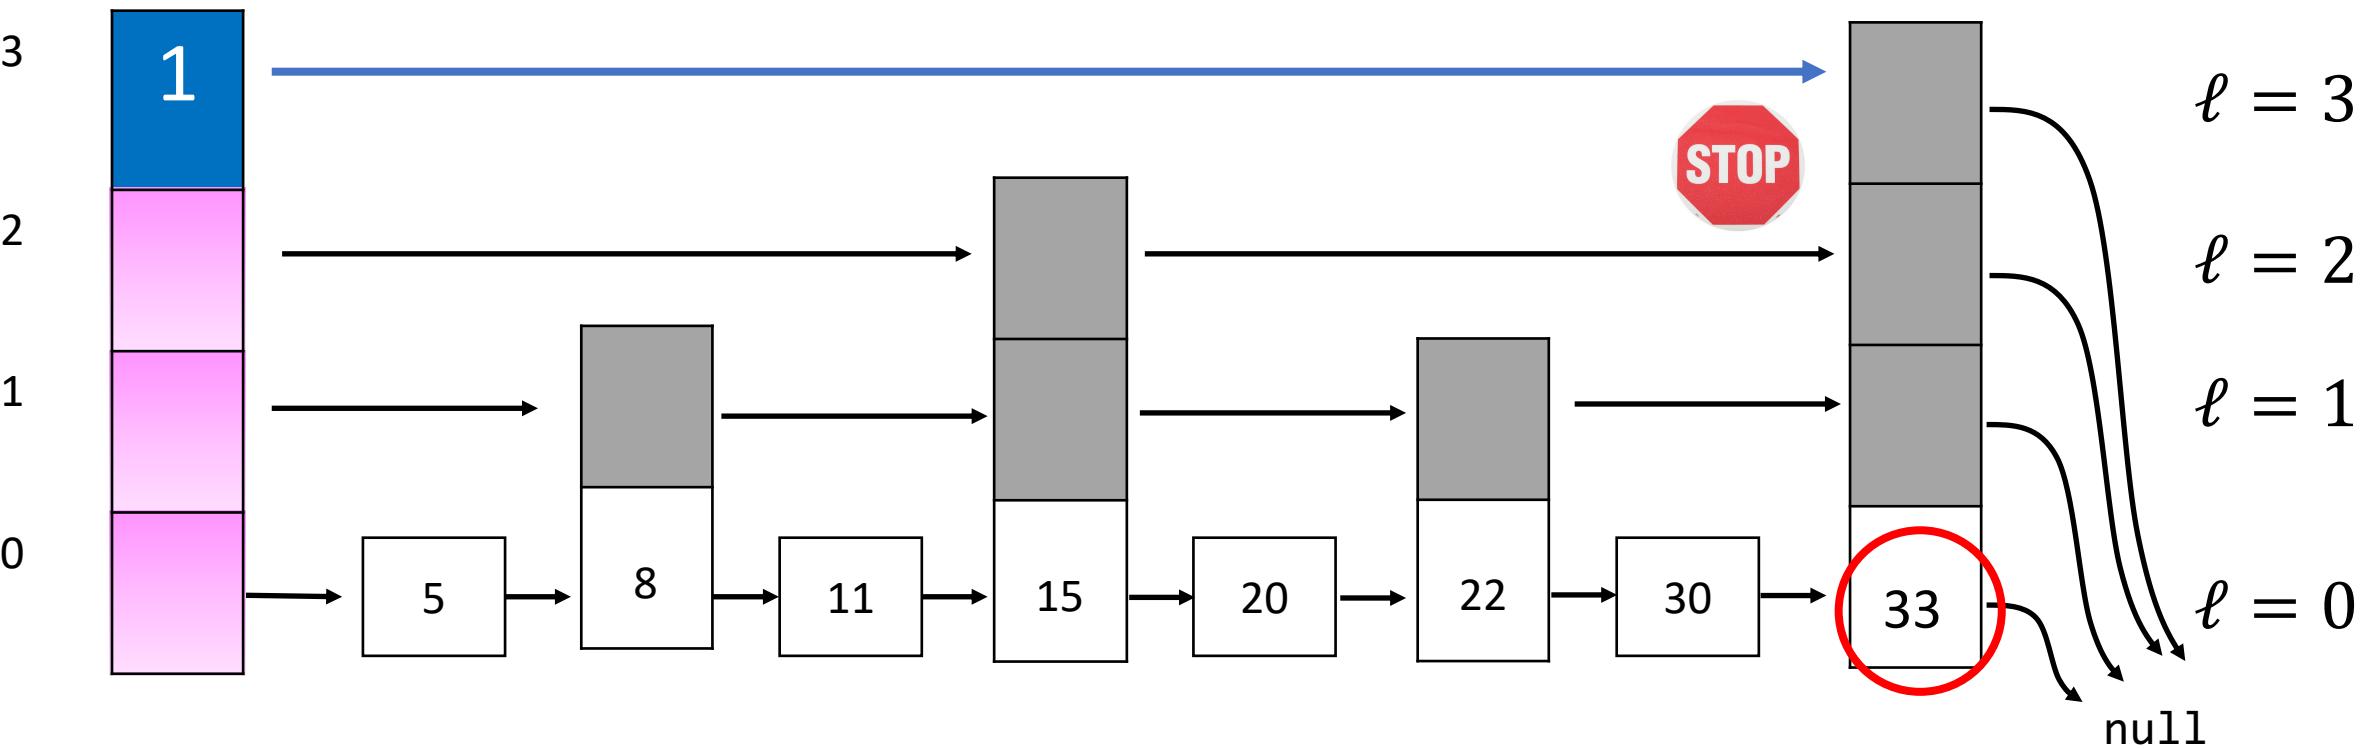
*Start at header[levels-1]...*



# Search example

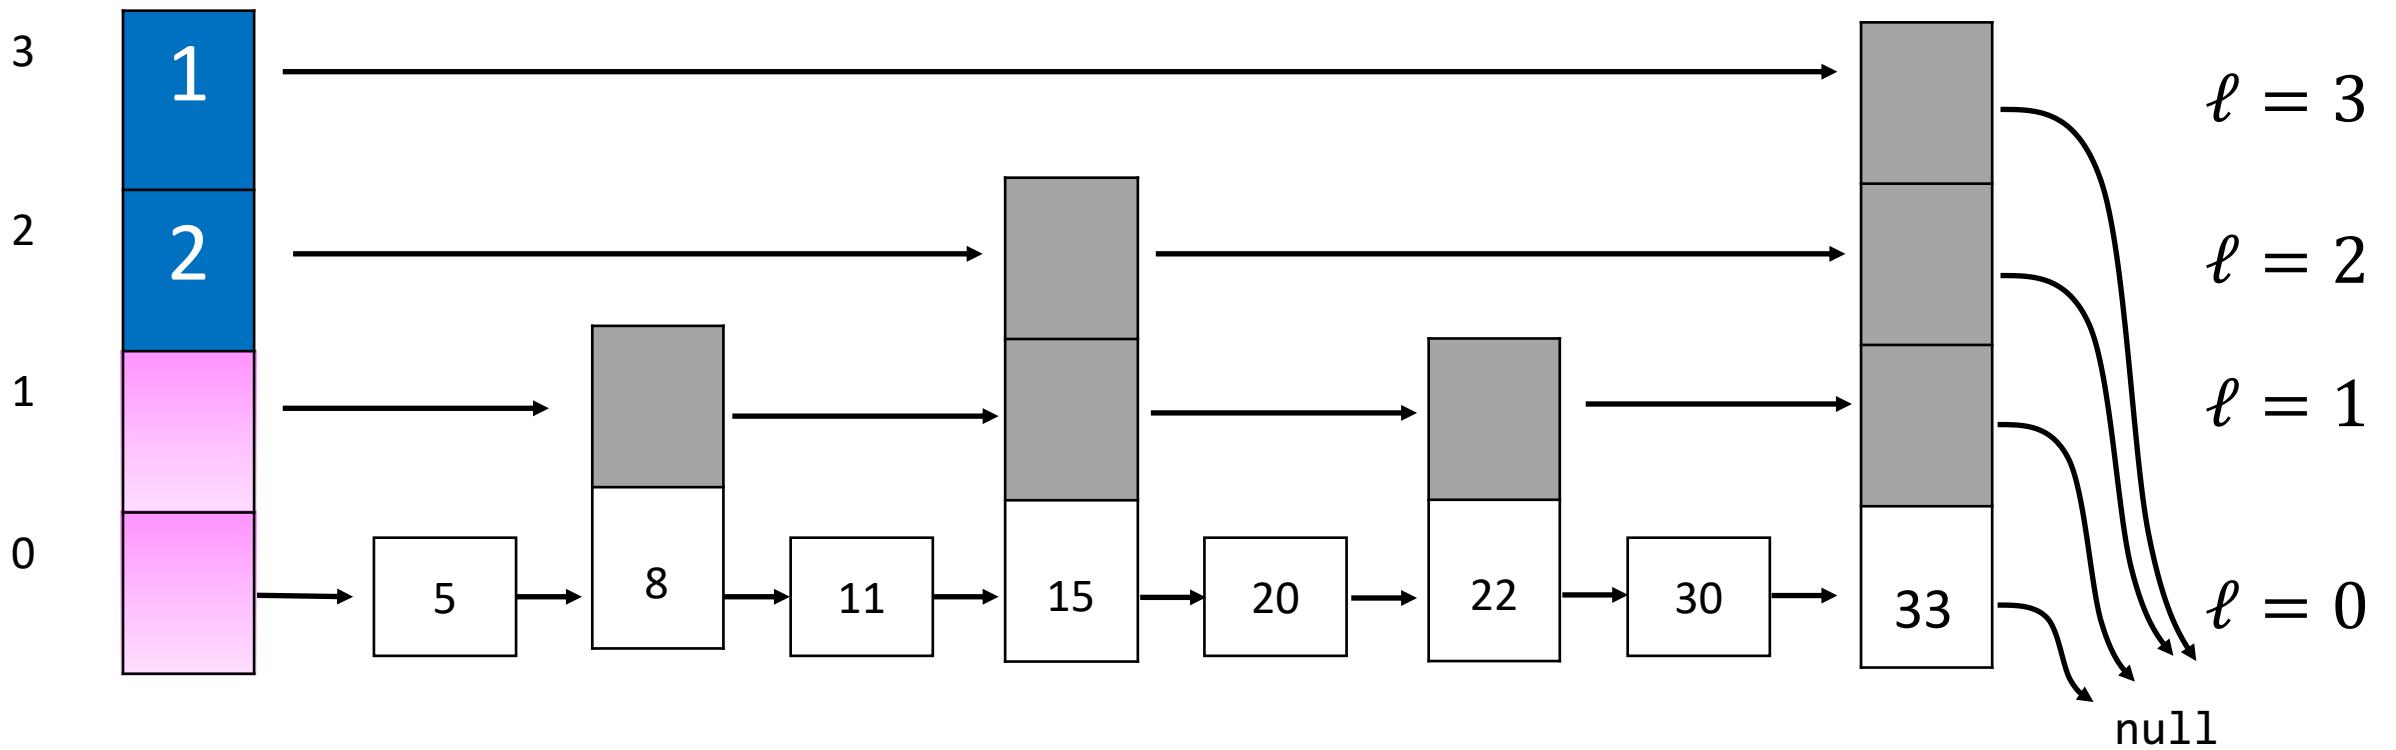


*The pointer points to a tower with element greater than 20*



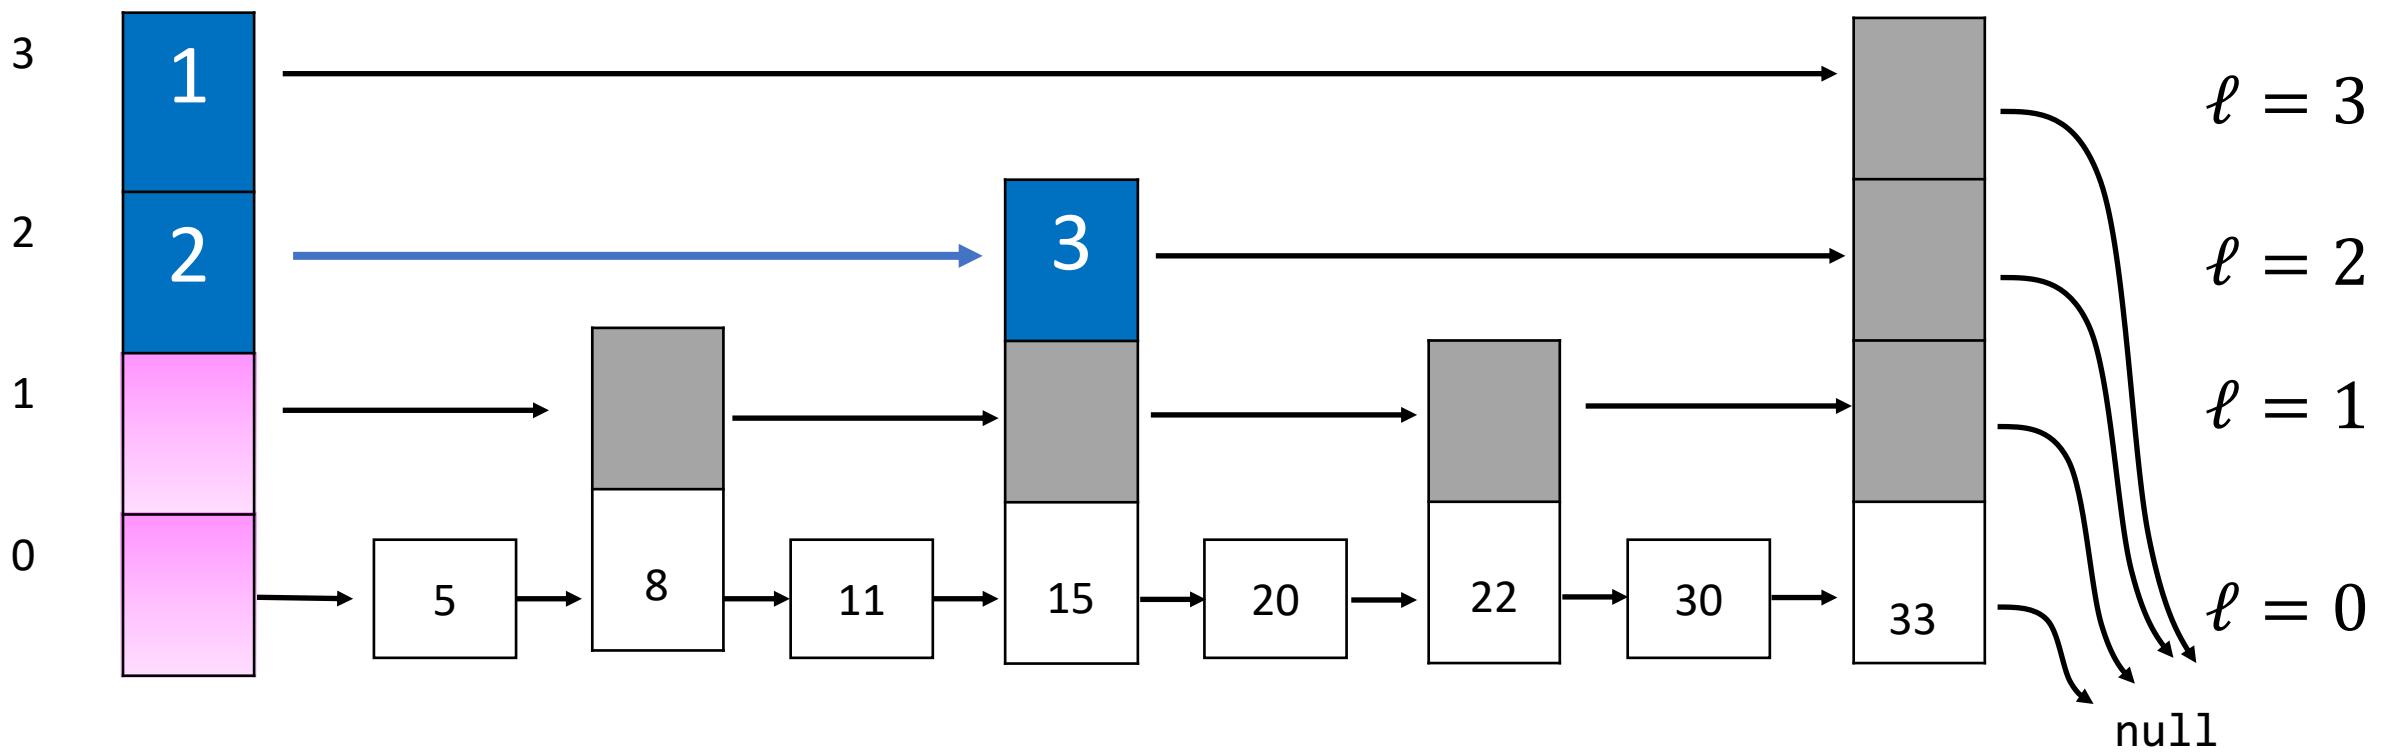
# Search example

*Solution: decrease your level!*



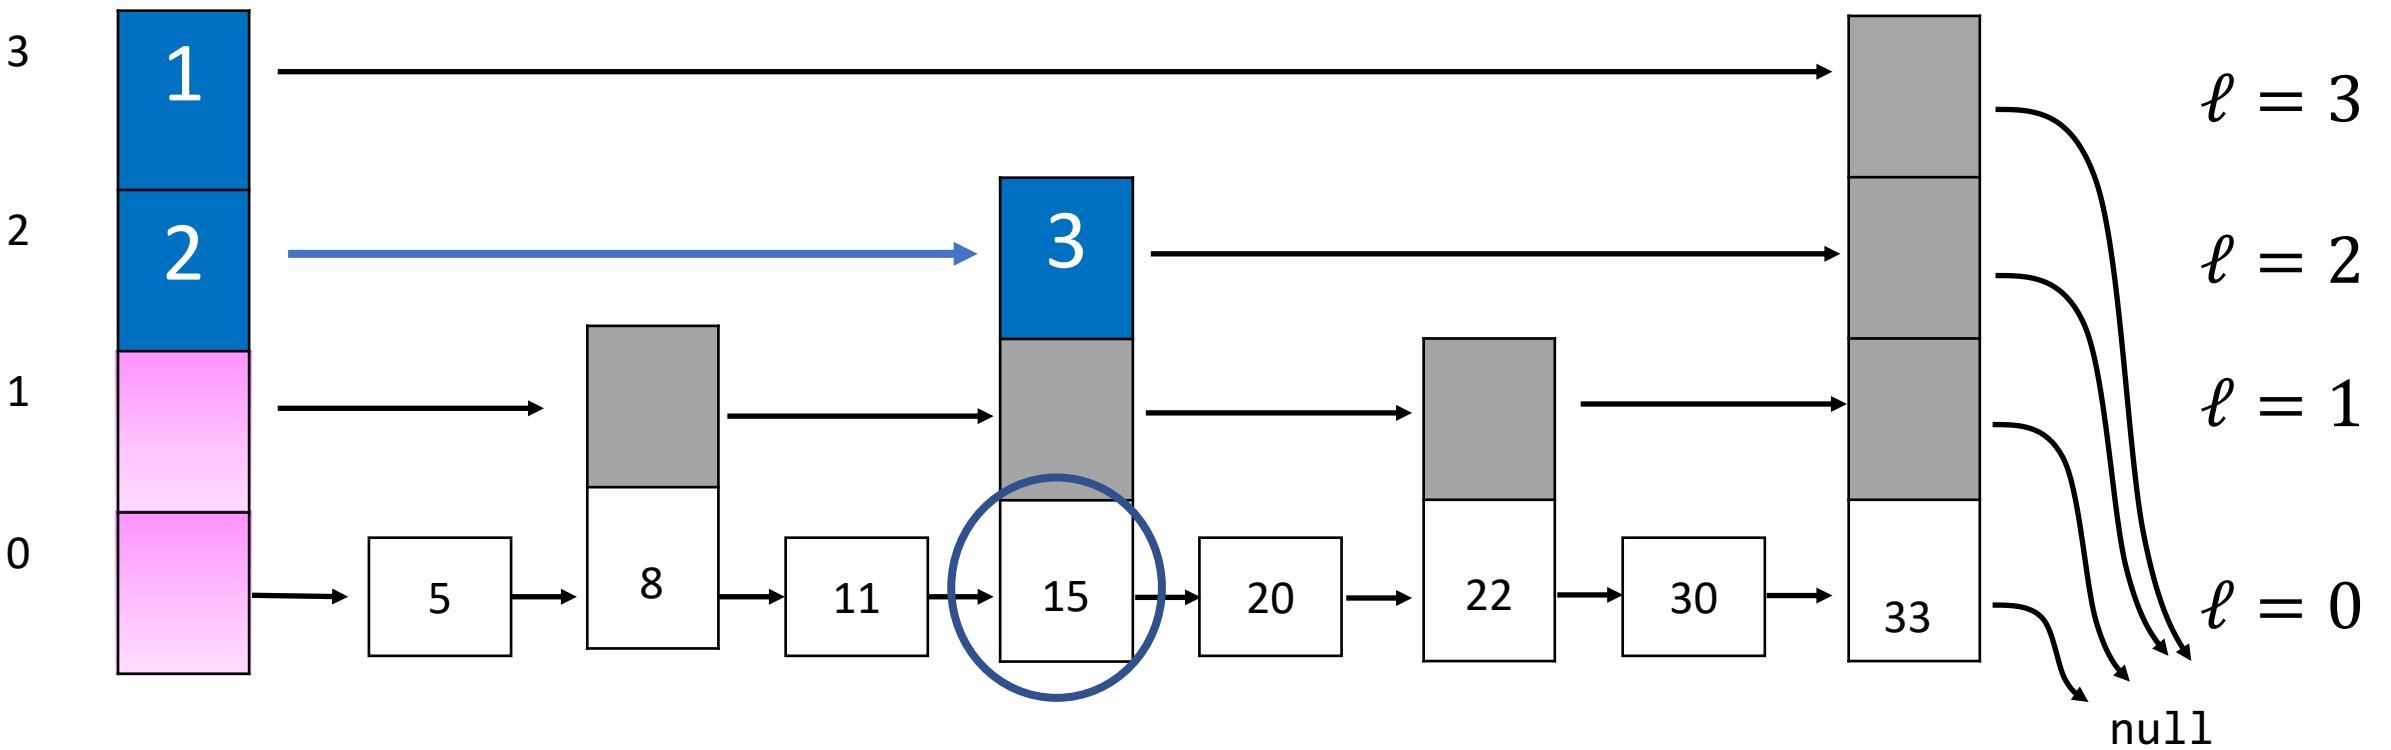
# Search example

*Dereference pointer.*



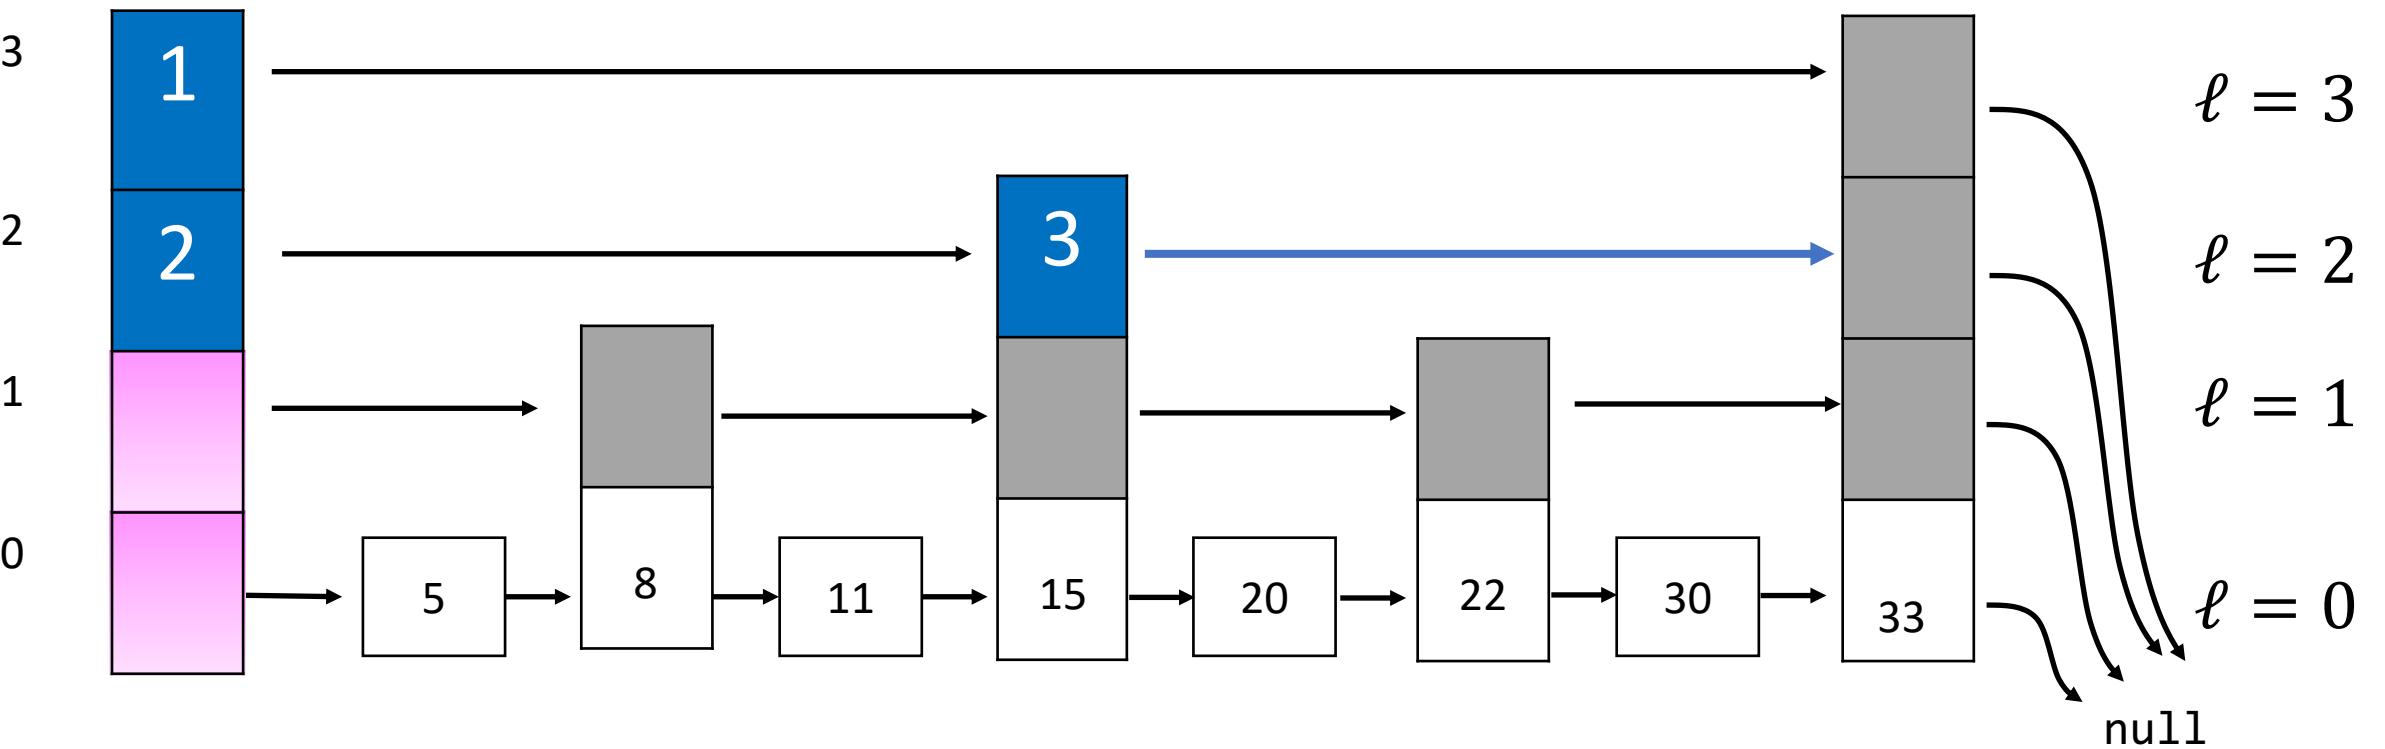
# Search example

Key (15) smaller than desired key (20)



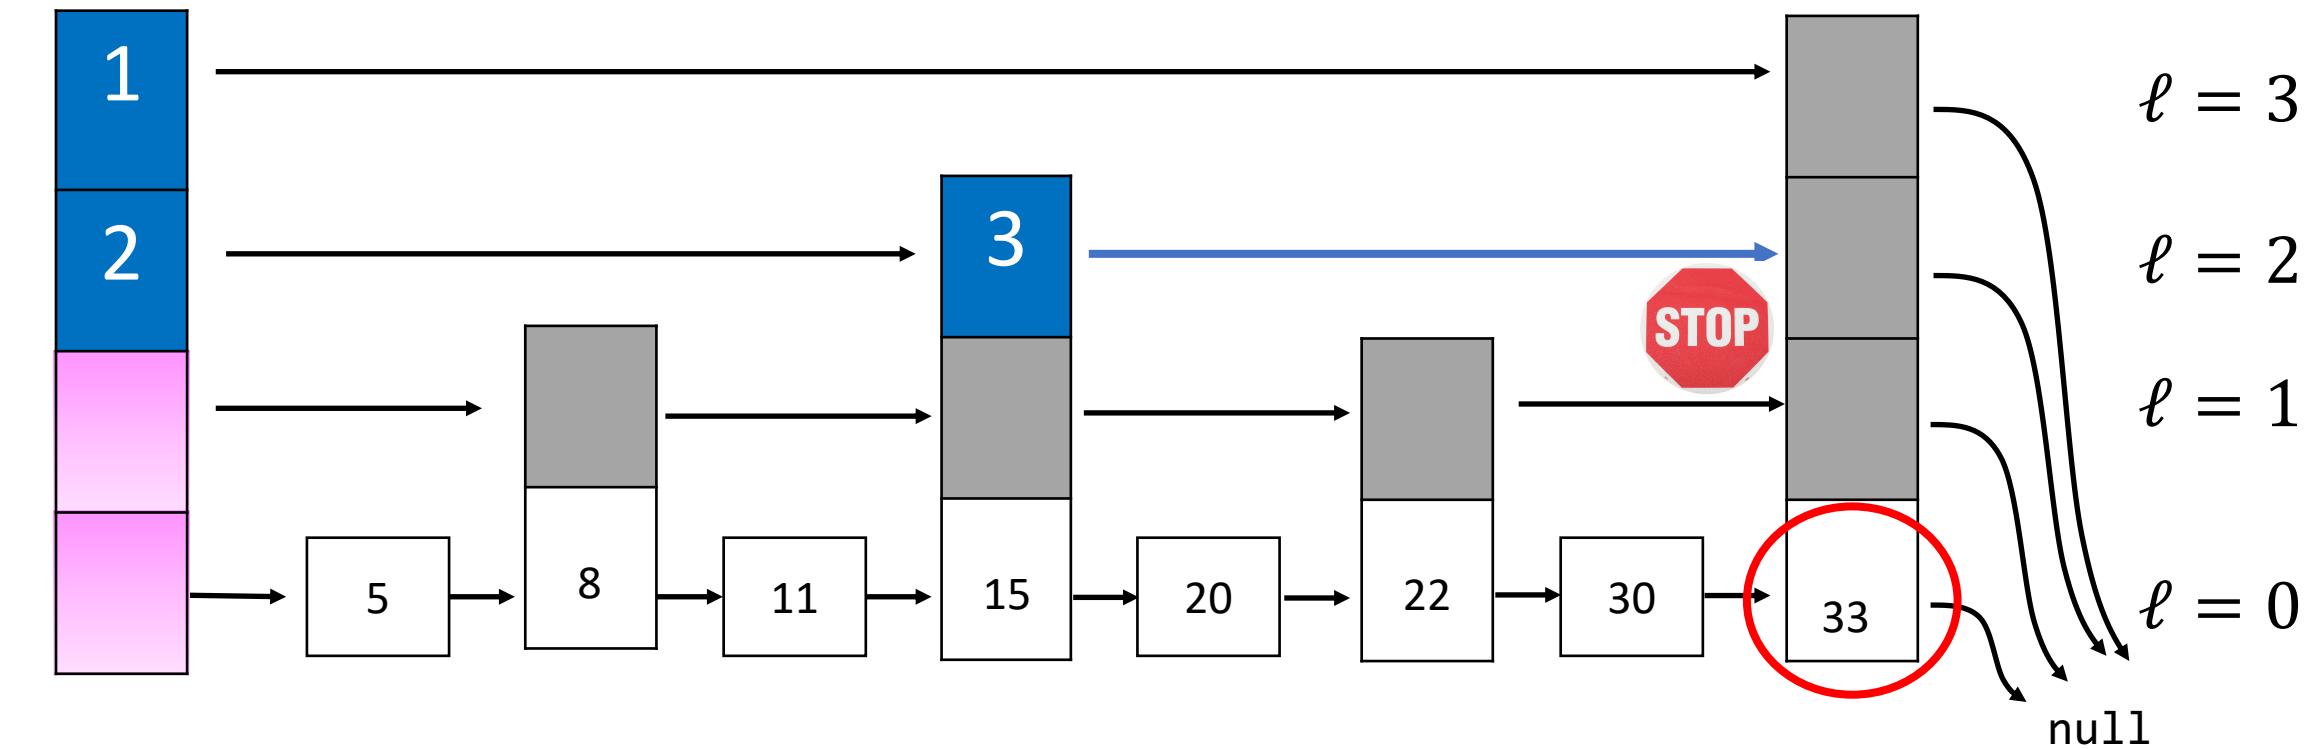
# Search example

*So keep going on the same level!*



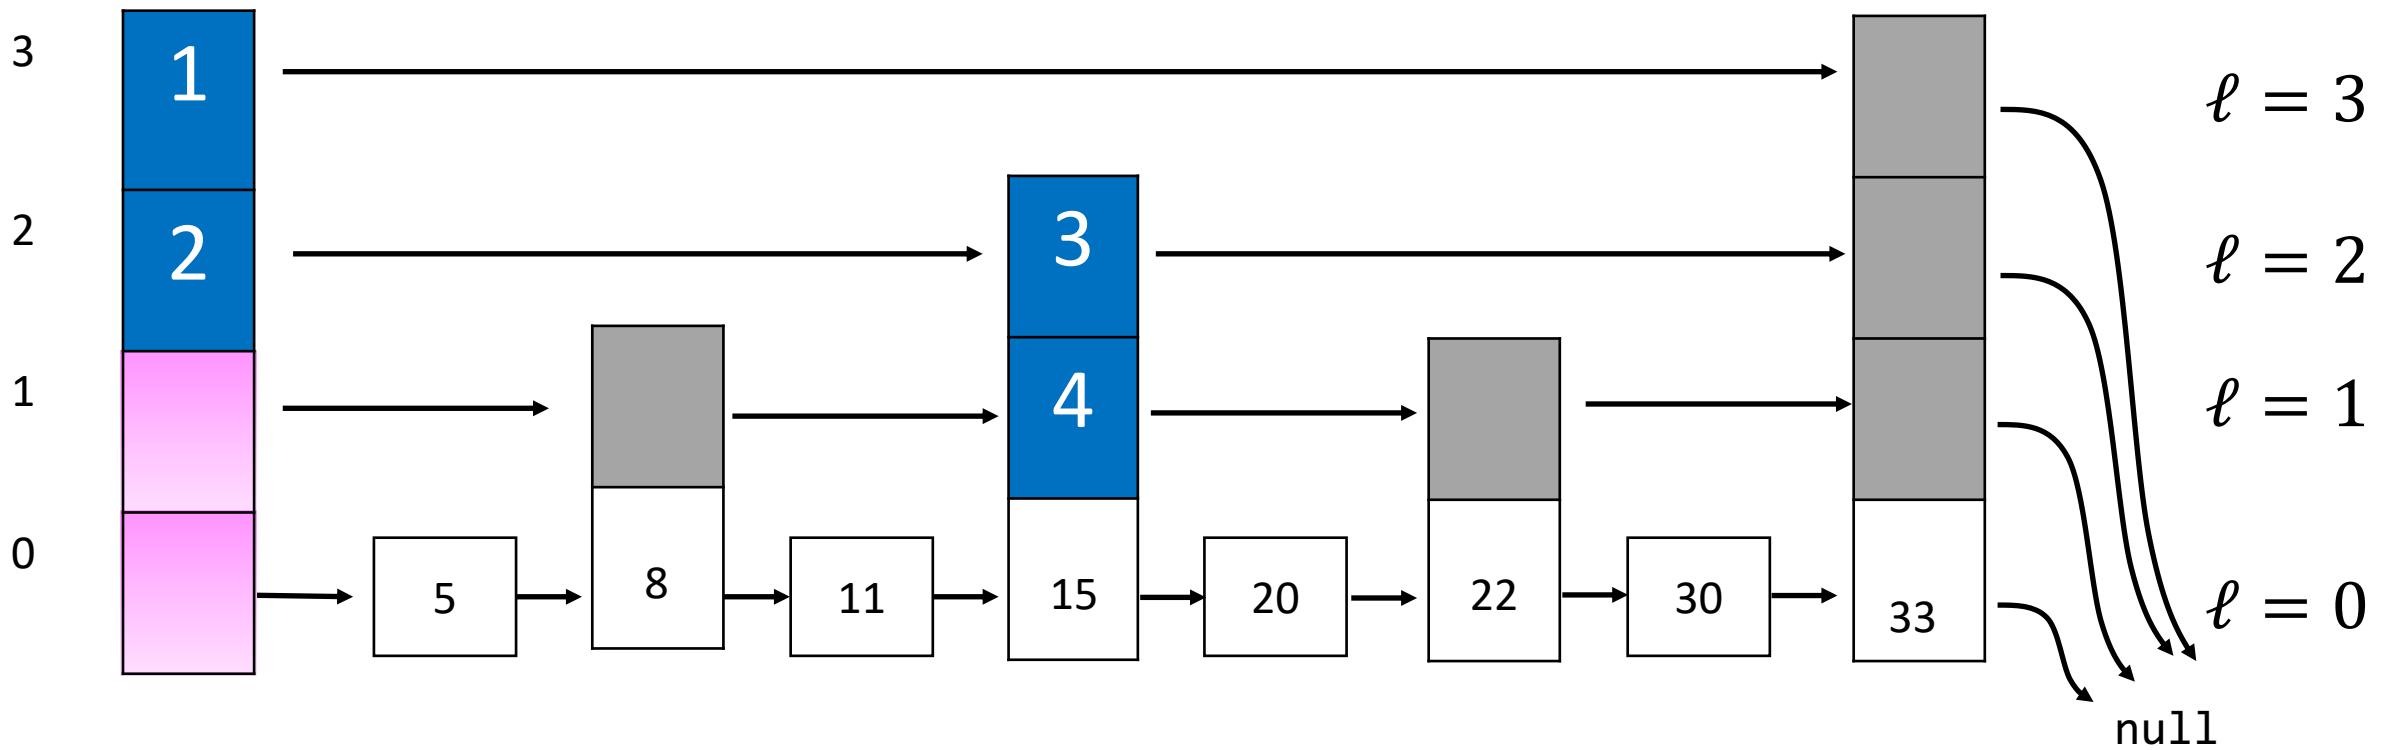
# Search example

Contained key (33) **greater** than sought key (20) again!



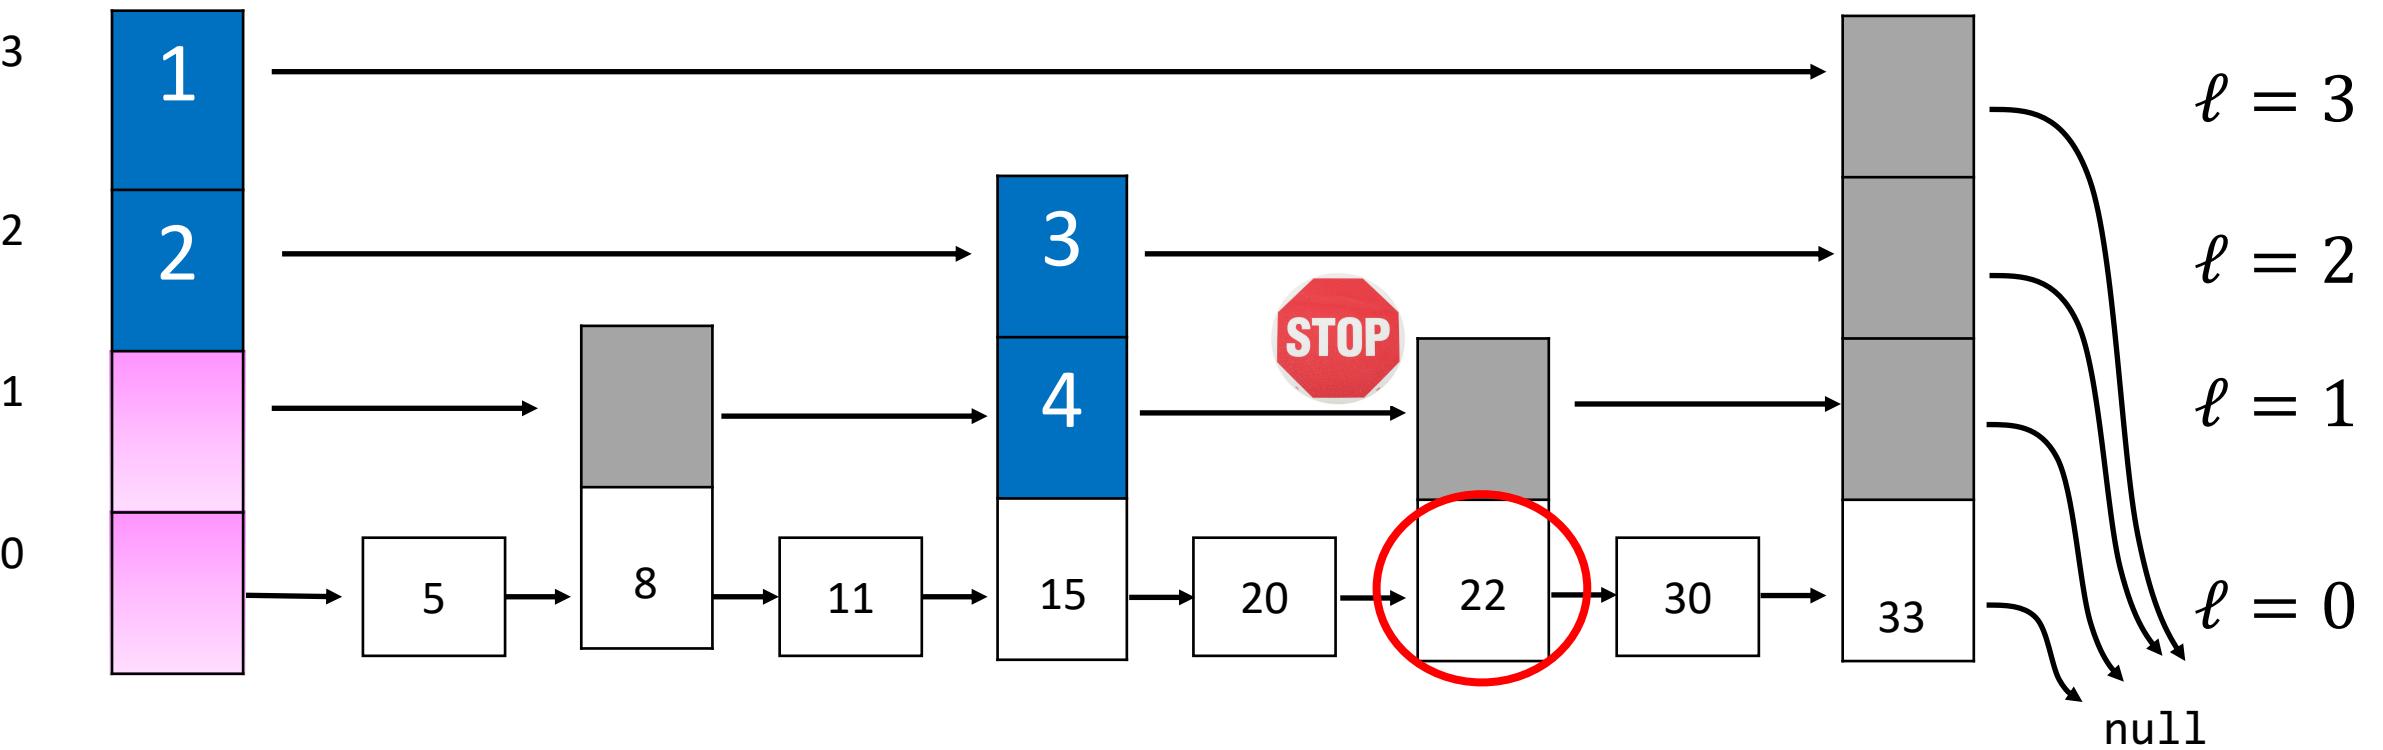
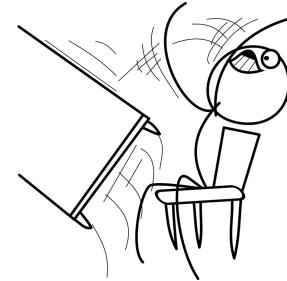
# Search example

So descend a level again!



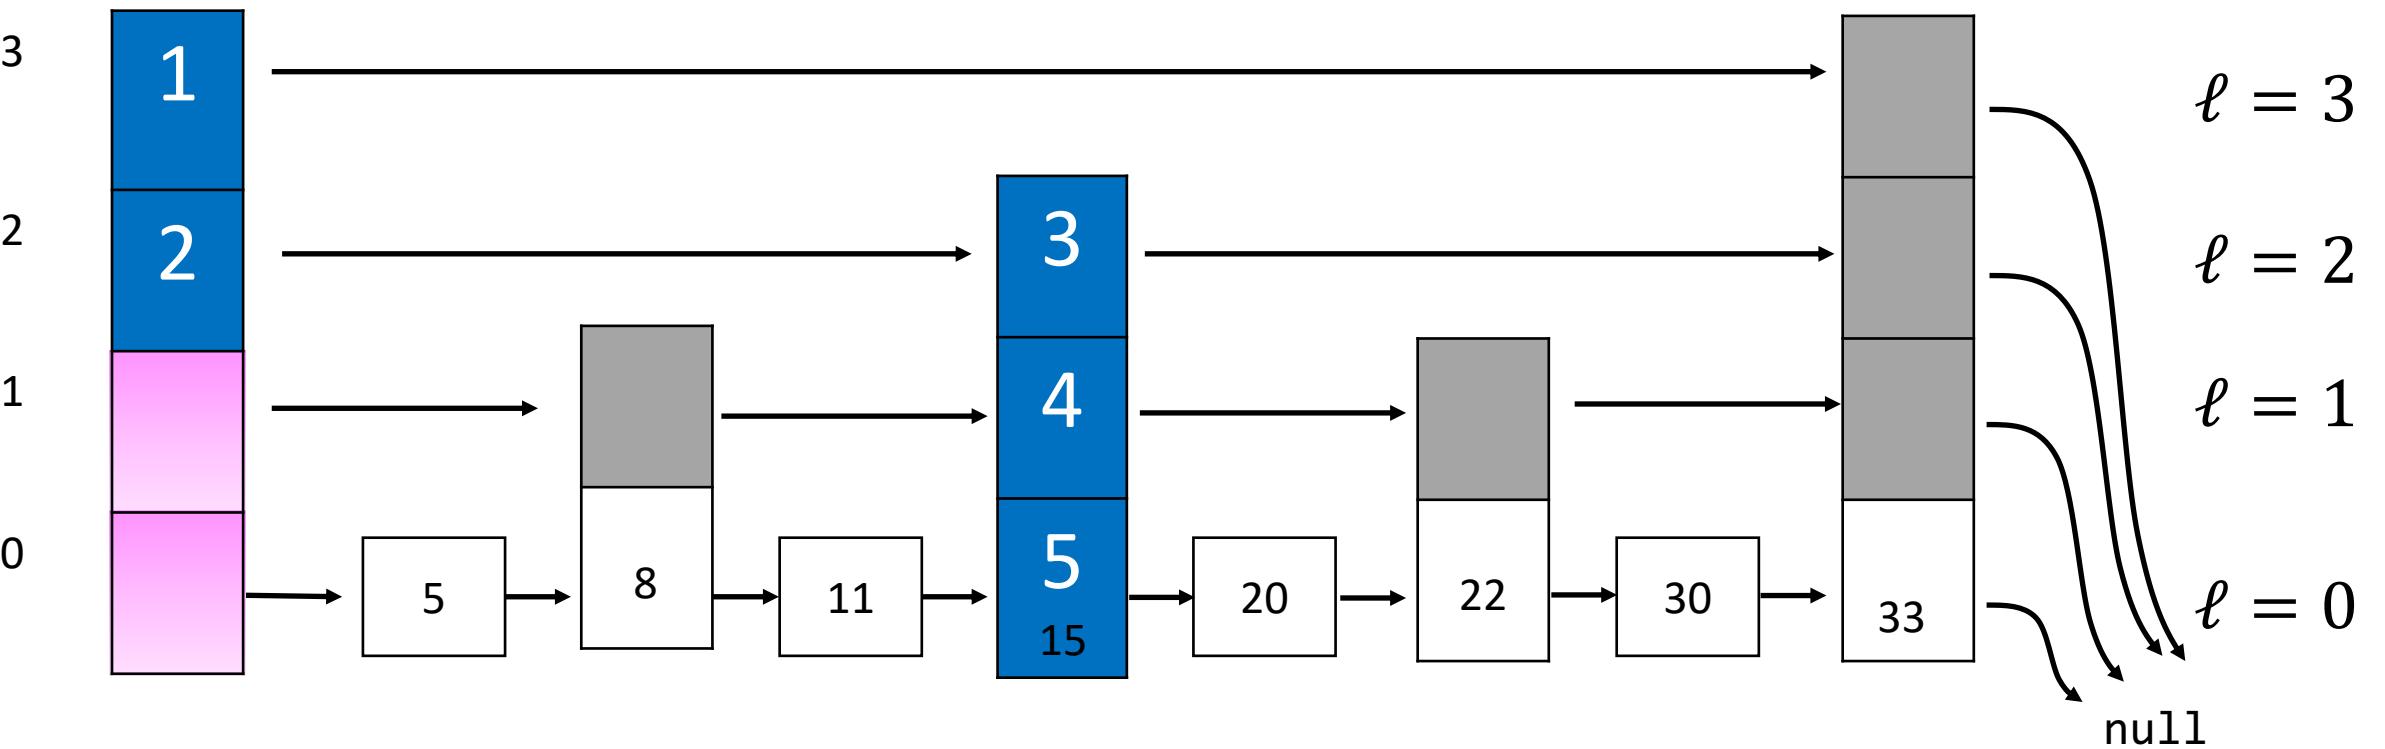
# Search example

Dereferencing pointer  
leads us to a tower with  
a key too big...



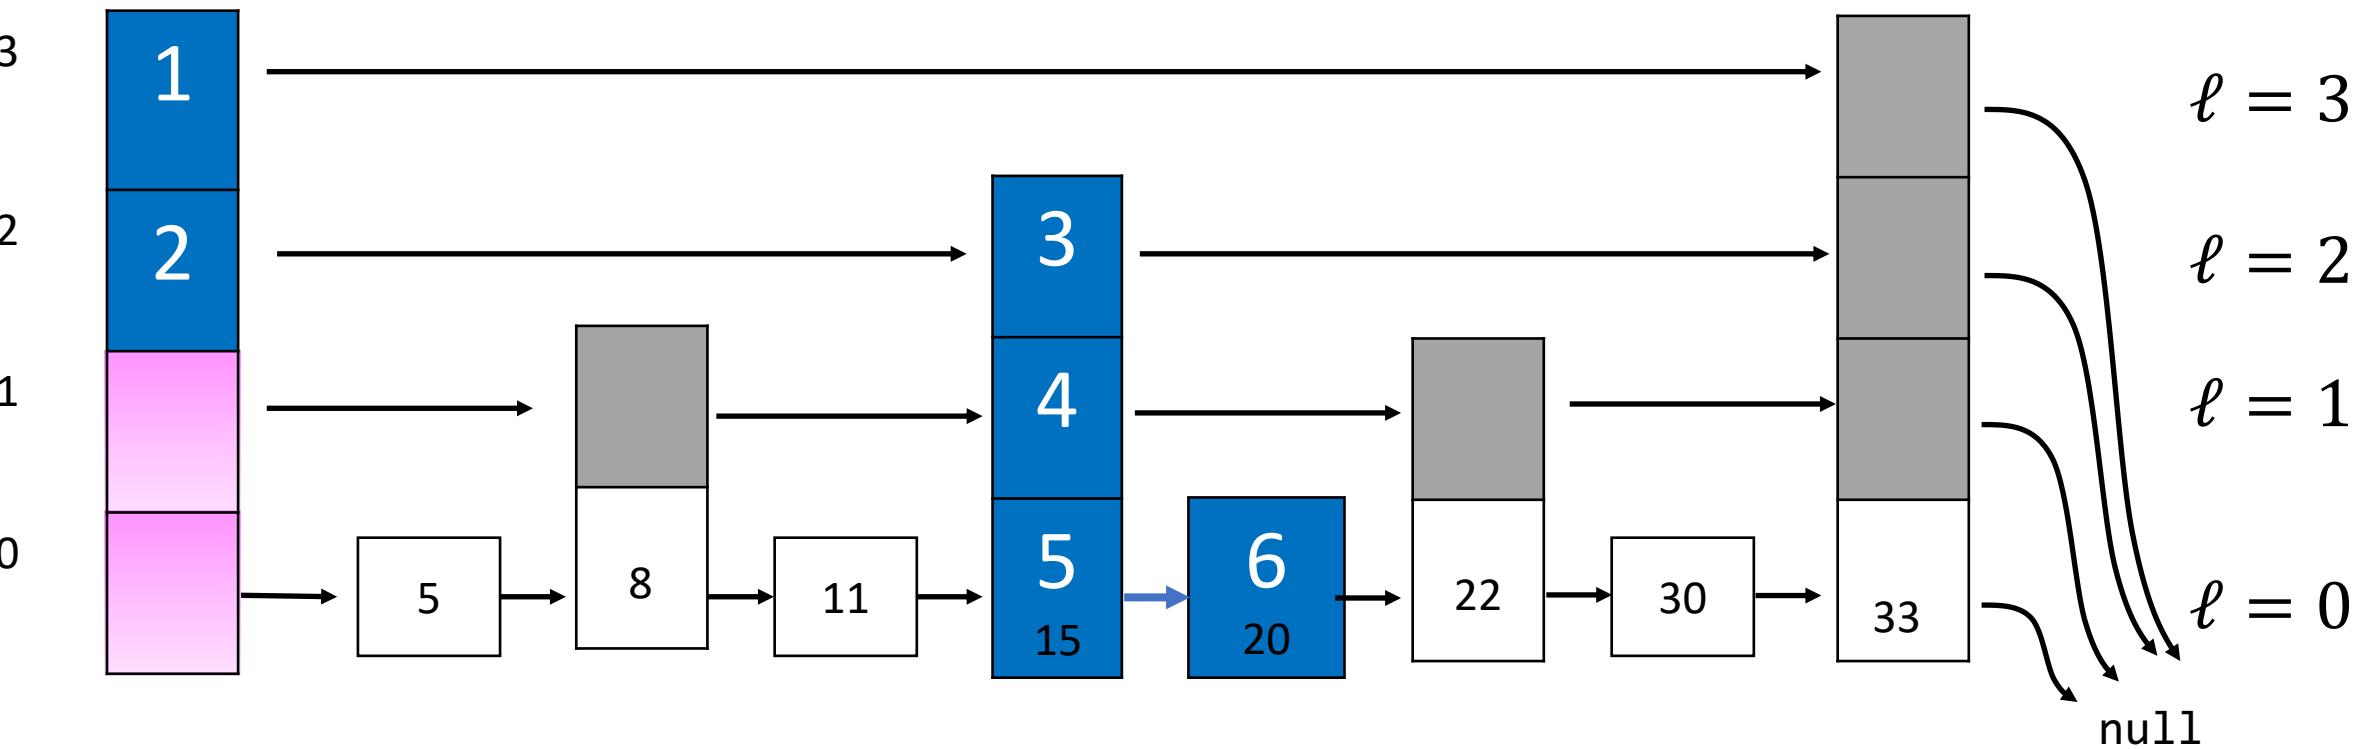
# Search example

So descend another  
level...



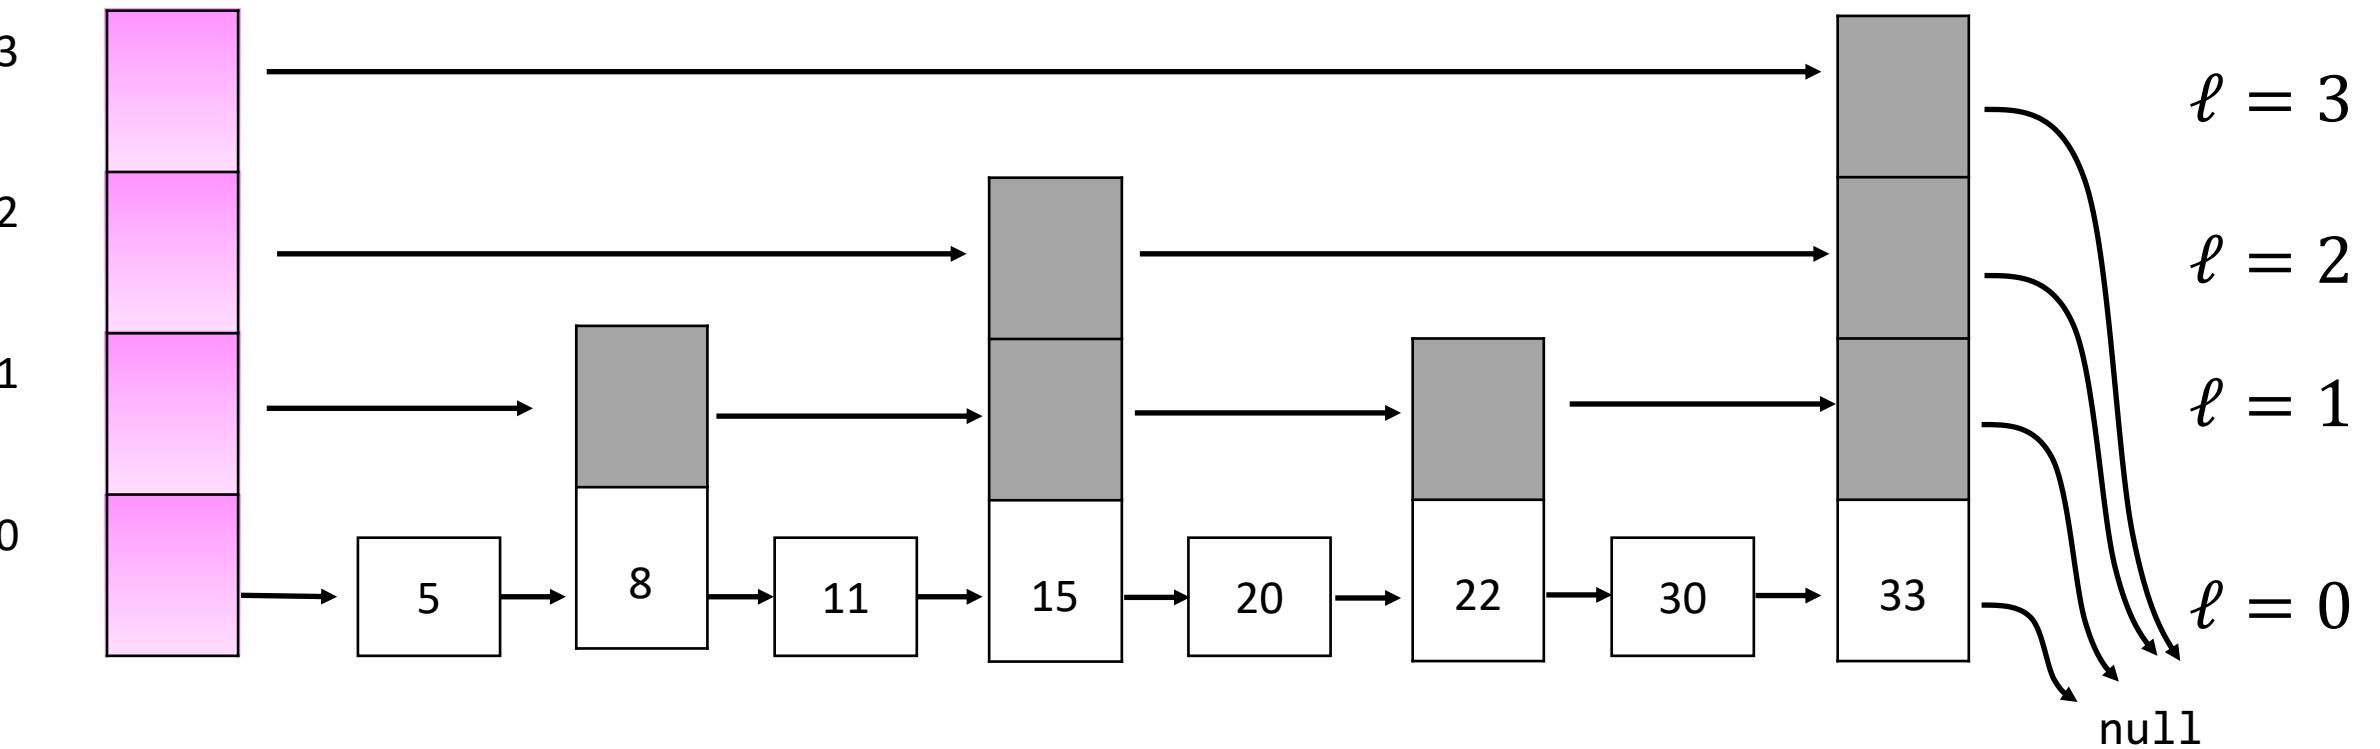
# Search example

Final pointer dereference gives us our key. Search hit.

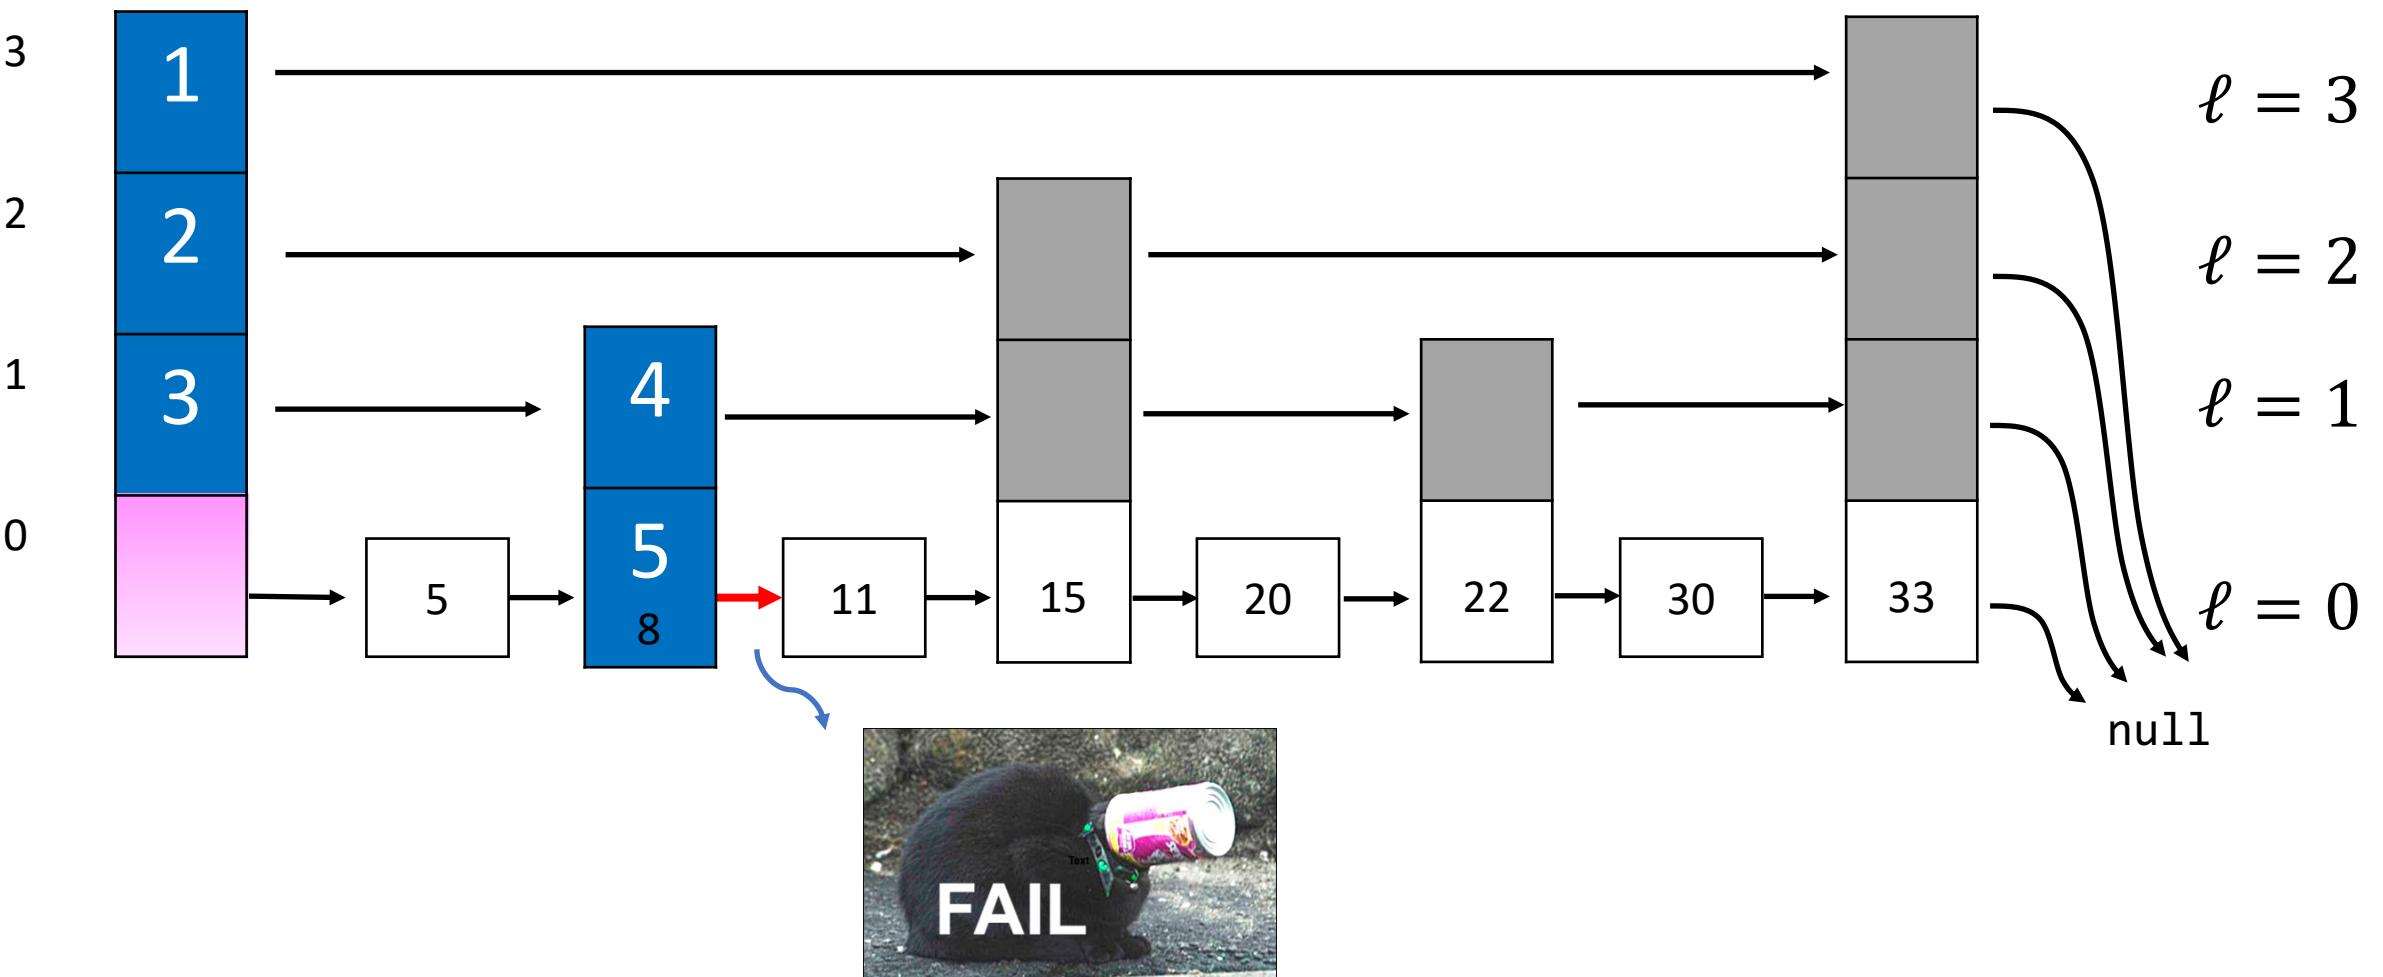


# Another search example – for you!

- For the same list array, sequentially mark the nodes traversed by searching for 9 (nine), which should fail!



# Solution



# Space analysis

- To make matters simpler, assume that  $n$  is some power of 2.
  - So,  $n/2$  will be an **integer**, as will be  $n/4 = n/2^2$ ,  $n/8 = n/2^3$ , etc  $= 2^{-(\log_2 n - 1)}$
- Alongside the list level, we have  $\log_2 n + 1$  levels.
- The list level ( $0^{\text{th}}$  level) has  $n$  nodes.
- The  $1^{\text{st}}$  level has  $n/2$  nodes.
- The  $2^{\text{nd}}$  level has  $n/4$  nodes.
- ...
- ...
- Level #  $\log_2 n$  has 1 nodes

Total # nodes:

$$n + \frac{n}{2} + \frac{n}{4} + \dots + 1 = n(1 + \frac{1}{2} + \frac{1}{4} + \dots + \frac{1}{n}) =$$
$$= n(2^0 + 2^{-1} + 2^{-2} + \dots + 2^{-\log_2 n})$$

Sum of geometric progression with  $a_0 = 2^0 = 1$ ,  
 $r = 1/2$  and  $\log_2 n + 1$  terms!

# Unpacking the sum...

- $2^0 + 2^{-1} + 2^{-2} + \dots + 2^{-(\log_2 n)} = \frac{2^0[1 - (\frac{1}{2})^{\log_2 n+1}]}{1 - \frac{1}{2}} = 2[1 - (\frac{1}{2})^{\log_2 n+1}]$
- $(\frac{1}{2})^{\log_2 n+1} = \frac{1}{2} \cdot \frac{1}{2^{\log_2 n}} = \frac{1}{2n}$
- Substituting the green stuff onto the red stuff yields:
- $2^0 + 2^{-1} + 2^{-2} + \dots + 2^{-(\log_2 n)} = 2 \left(1 - \frac{1}{2n}\right) = \frac{2n-1}{n}$
- Note that this tends to 2 as  $n$  goes to infinity, which is no surprise since the geometric series  $\sum_{i=0}^n 2^{-i}$  converges to  $\frac{1}{1-r} = \frac{1}{1 - \frac{1}{2}} = 2$

# Final answer

- Num nodes =  $n(2^0 + 2^{-1} + 2^{-2} + \dots + 2^{-\log_2 n}) =$   
 $\cancel{n} \frac{2n - 1}{\cancel{n}} = 2n - 1$
- We also have the “header” node which contributes  $\log_2 n + 1$  nodes...
- Final answer:  $2n - 1 + \log_2 n + 1 = 2n + \log_2 n = \mathcal{O}(n)$

# Search efficiency analysis

- We have a sorted linked list  $\ell$  with  $n$  elements and we build a list array  $A$  on top.
- We want to search for an element  $e$ .
- Unit cost (what  $n$  counts): dereferencing a pointer.
- What is the computational complexity of searching for  $e$  in  $\ell$  using  $A$ ?

# Search efficiency analysis

- We have a sorted linked list  $\ell$  with  $n$  elements and we build a list array  $A$  on top.
- We want to search for an element  $e$ .
- Unit cost (what  $n$  counts): dereferencing a pointer.
- What is the computational complexity of searching for  $e$  in  $\ell$  using  $A$ ?

$O(n)$

$O(\log_2 n)$

$O(\log_2 (\log_2 n))$

Something else  
(what?)

# Search efficiency analysis

- We have a sorted linked list  $\ell$  with  $n$  elements and we build a list array  $A$  on top.
- We want to search for an element  $e$ .
- Unit cost (what  $n$  counts): dereferencing a pointer.
- What is the computational complexity of searching for  $e$  in  $\ell$  using  $A$ ?

$O(n)$

$O(\log_2 n)$

$O(\log_2 (\log_2 n))$

Something else  
(what?)

Now let's determine the constant...

# Search efficiency analysis

- We have a sorted linked list  $\ell$  with  $n$  elements and we build a list array  $A$  on top.
- We want to search for an element  $e$ .
- Unit cost (what  $n$  counts): dereferencing a pointer.
- What is the maximum number of links we will scan to succeed or fail the search for  $e$  in  $\ell$  using  $A$ ?

$\log_2 n$

$1.5 \log_2 n$

$2 \log_2 n$

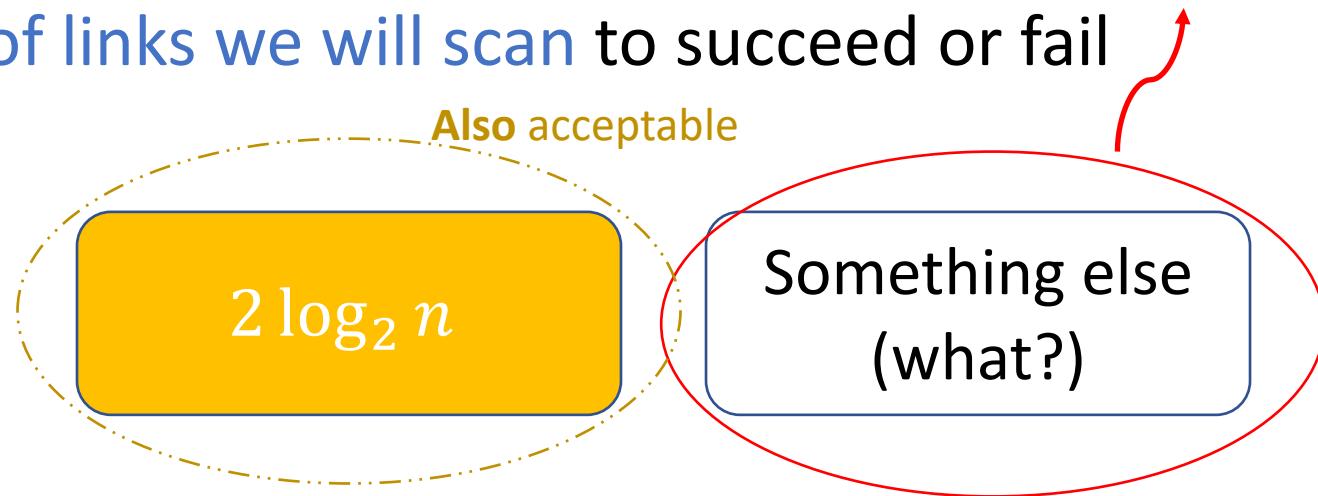
Something else  
(what?)

# Search efficiency analysis

- We have a sorted linked list  $\ell$  with  $n$  elements and we build a list array  $A$  on top.
- We want to search for an element  $e$ .
- Unit cost (what  $n$  counts): dereferencing a pointer.  $2 \log_2 n + 1$
- What is the maximum number of links we will scan to succeed or fail the search for  $e$  in  $\ell$  using  $A$ ?

$\log_2 n$

$1.5 \log_2 n$



At every level, we visit at most 2 nodes (why?) and at level 0 we might have one additional pointer dereference.

# Huge issue with list arrays

- Inserting a new element or deleting from it **requires re-building the entire list array** 😞
- Pay  $O(n)$  cost to do so.
- For the insertion and/or deletion of  $n$  elements, the cost ends up being **quadratic**, which is unacceptable.
- Conclusion: a list array is **not** a good **dictionary because of sub-par insertion and deletion**.
  - But for **immutable** sorted lists, they work pretty well!

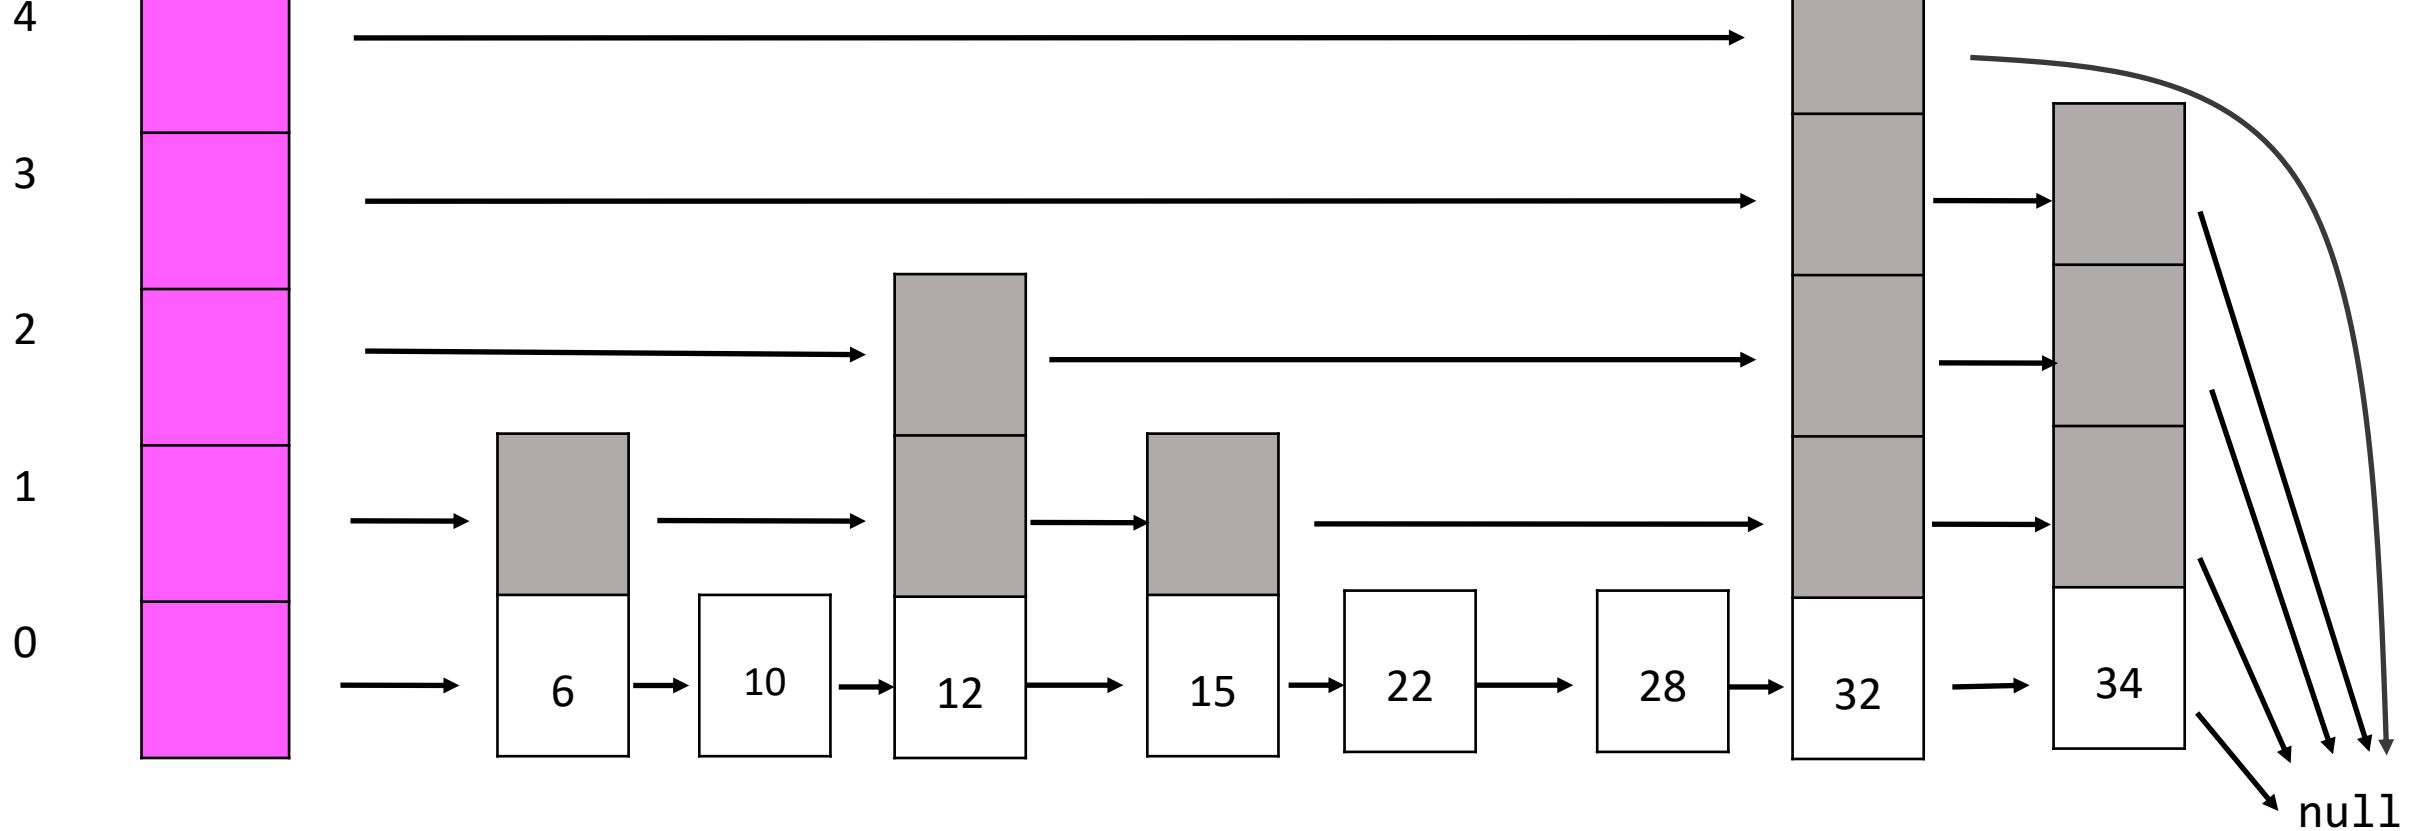
# Binary Search “vs” list arrays

Binary Search	List array
Used in contiguous storage (arrays)	Used in discontiguous storage (linked lists)
Requires up to $\log_2 n$ comparisons	Requires up to $2 \log_2 n + 1$ comparisons
Doesn't need extra space	Needs $O(n)$ extra space
Doesn't require adjustments after mutations	Needs to be re-built after every mutation

# Skiplist: relaxing the constraint

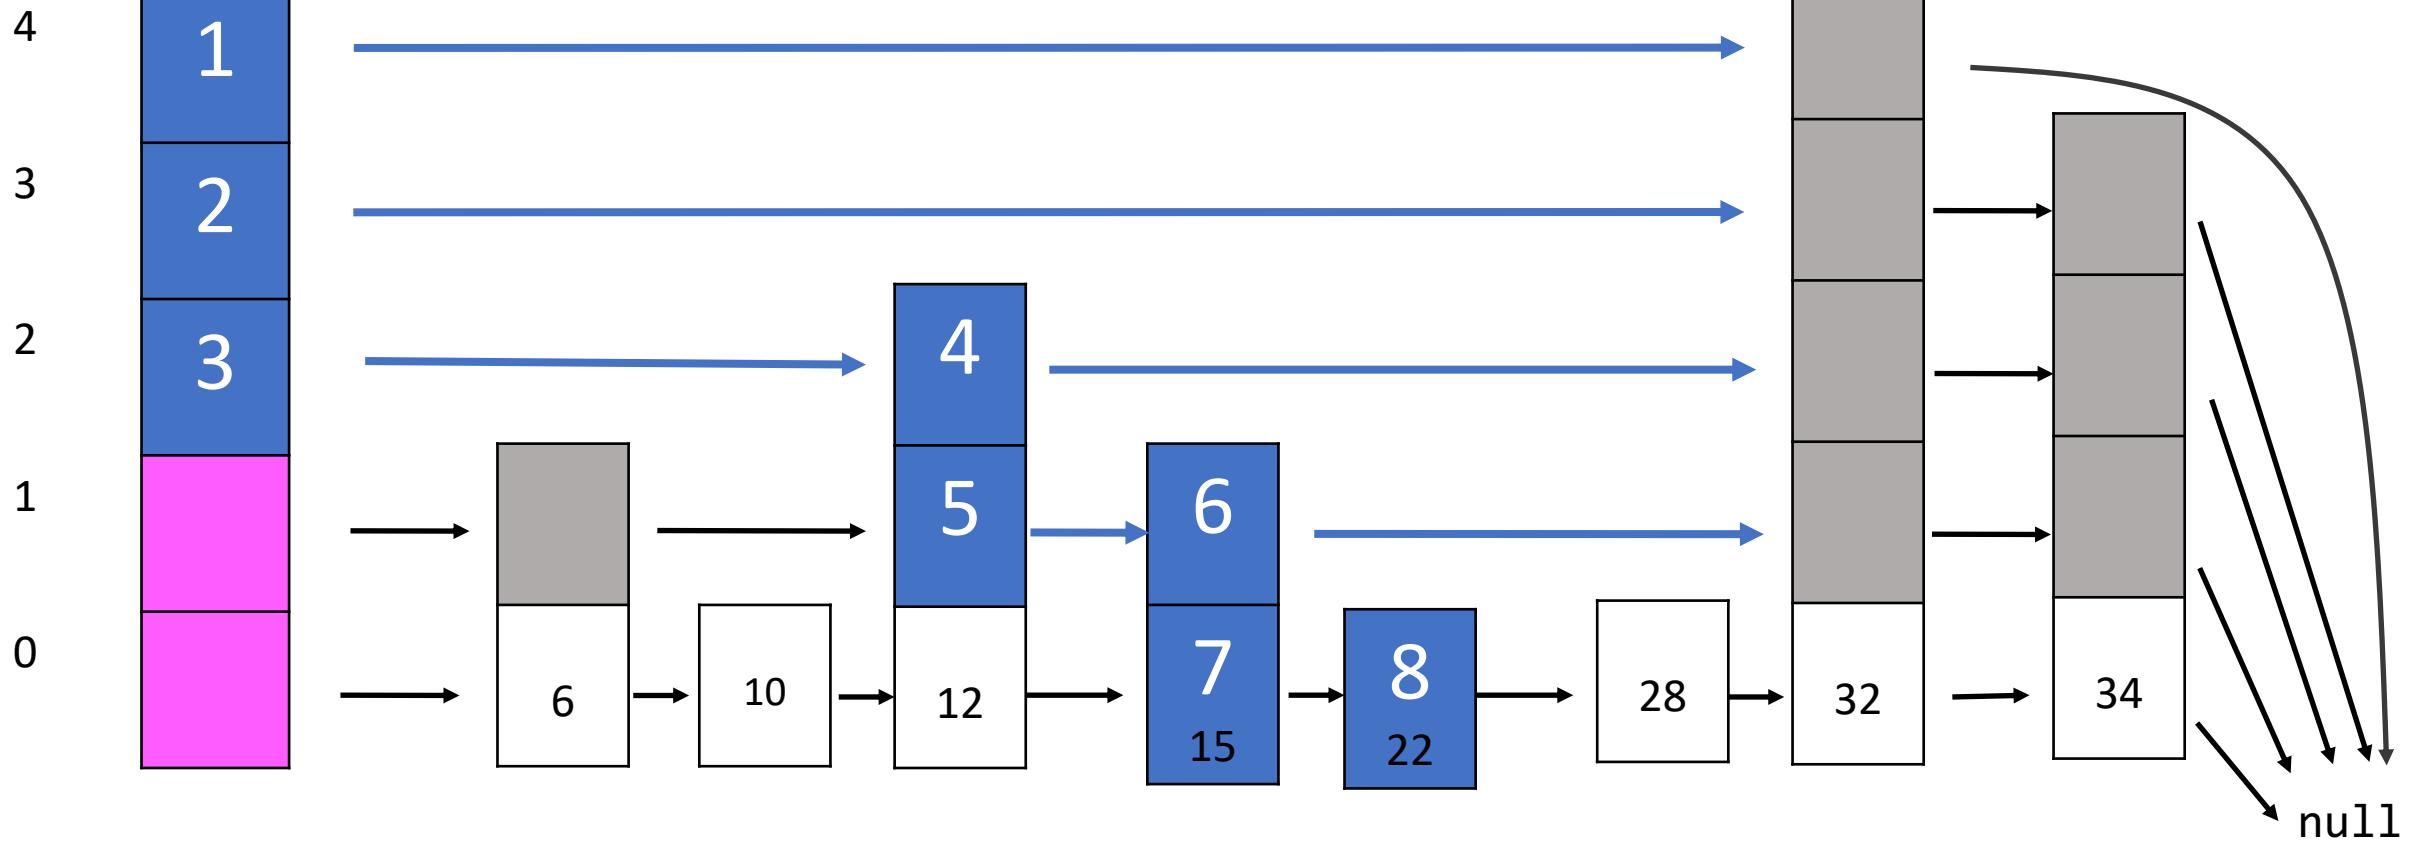
- Observation: in a list array, half ( $\frac{1}{2}$ ) of the “towers” that extend to level  $i$  also extend to level  $i + 1$ . Therefore:
  - $\frac{1}{2}$  of the list’s keys are represented in level 1
  - $\frac{1}{4}$  are represented in level 2
  - ....
  - $(\frac{1}{2})^i$  are represented in level  $i$ .
- Solution: probabilistically build the towers by randomly sampling the new height of an inserted node before actually inserting it!
- Use a geometric distribution:  $P(\text{level} = i) = \left(\frac{1}{2}\right)^{i+1}$

# Skiplist example



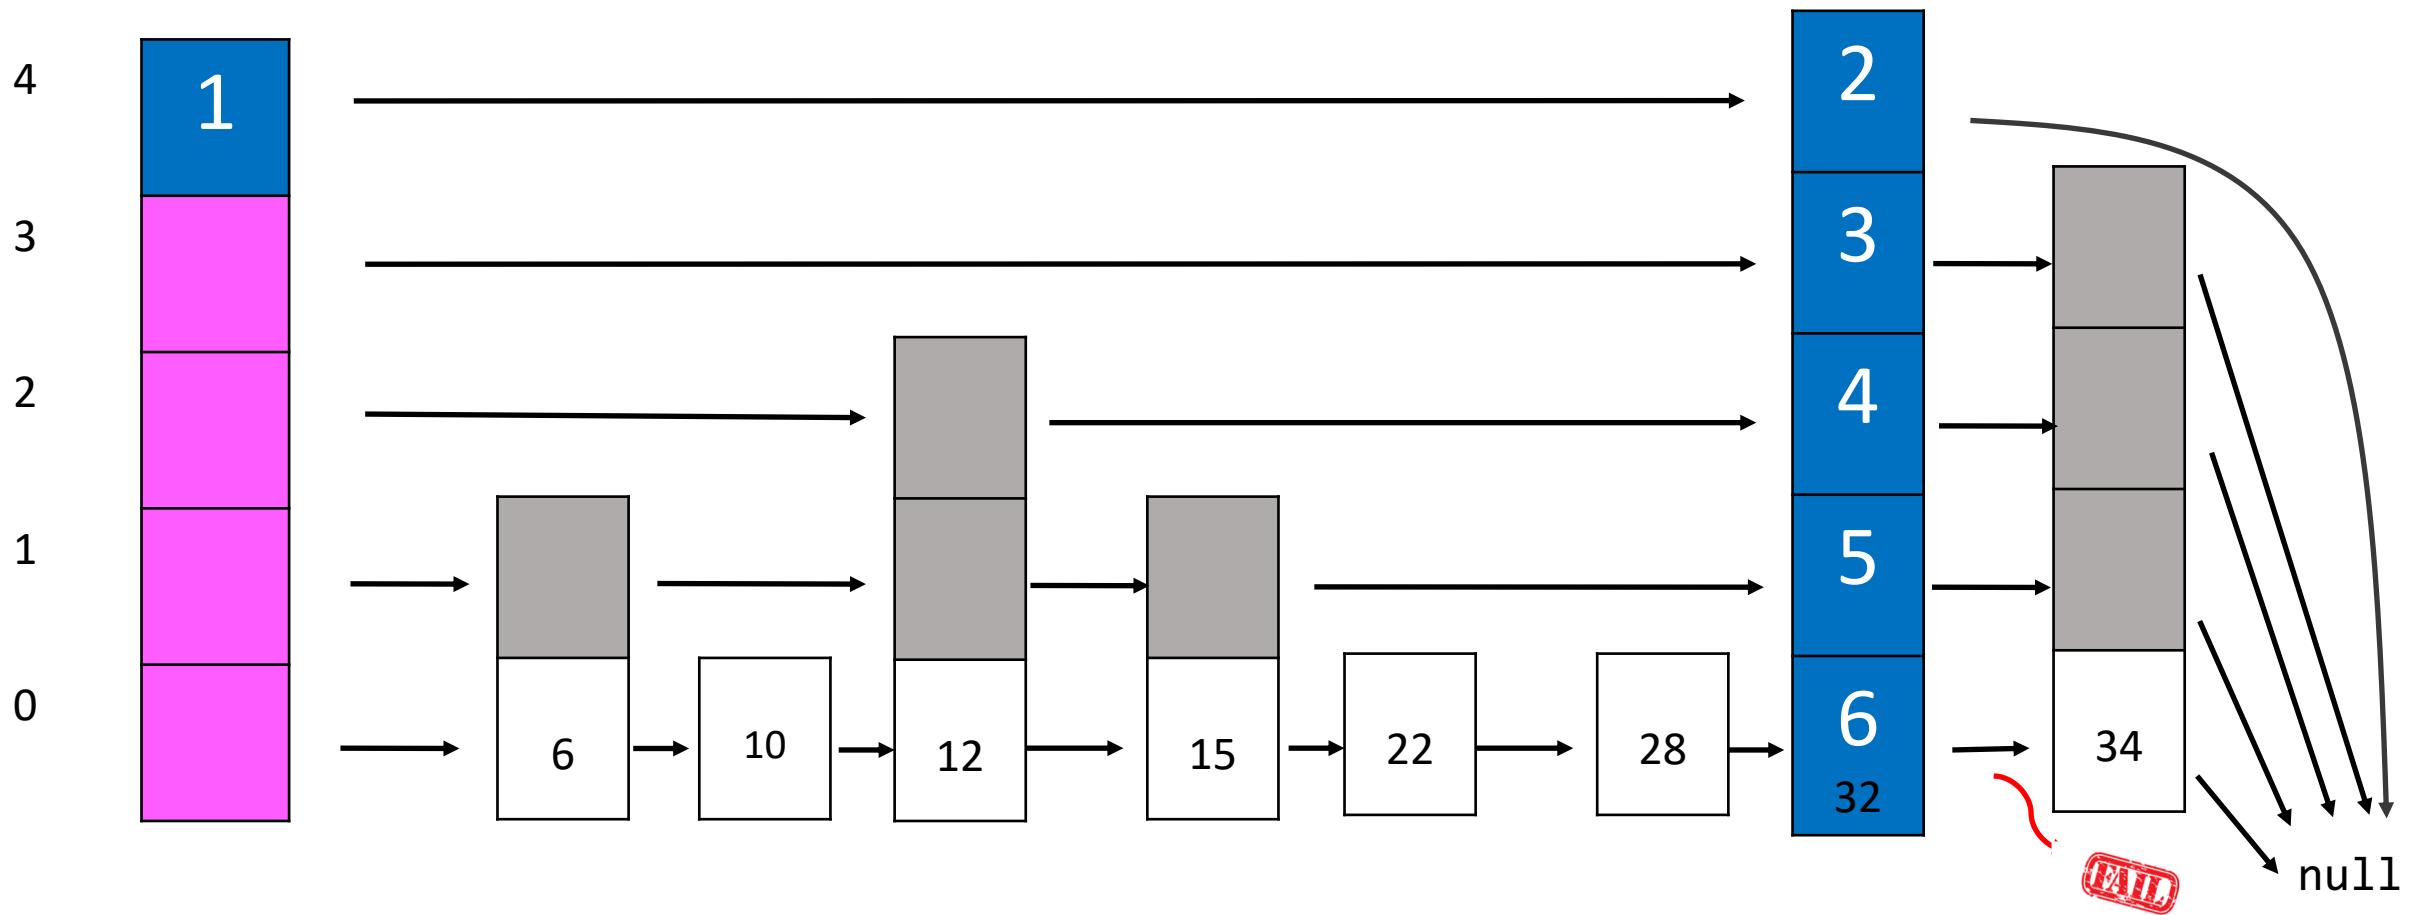
# Skiplist example

Searching for 22 (should  
be a search *hit*)



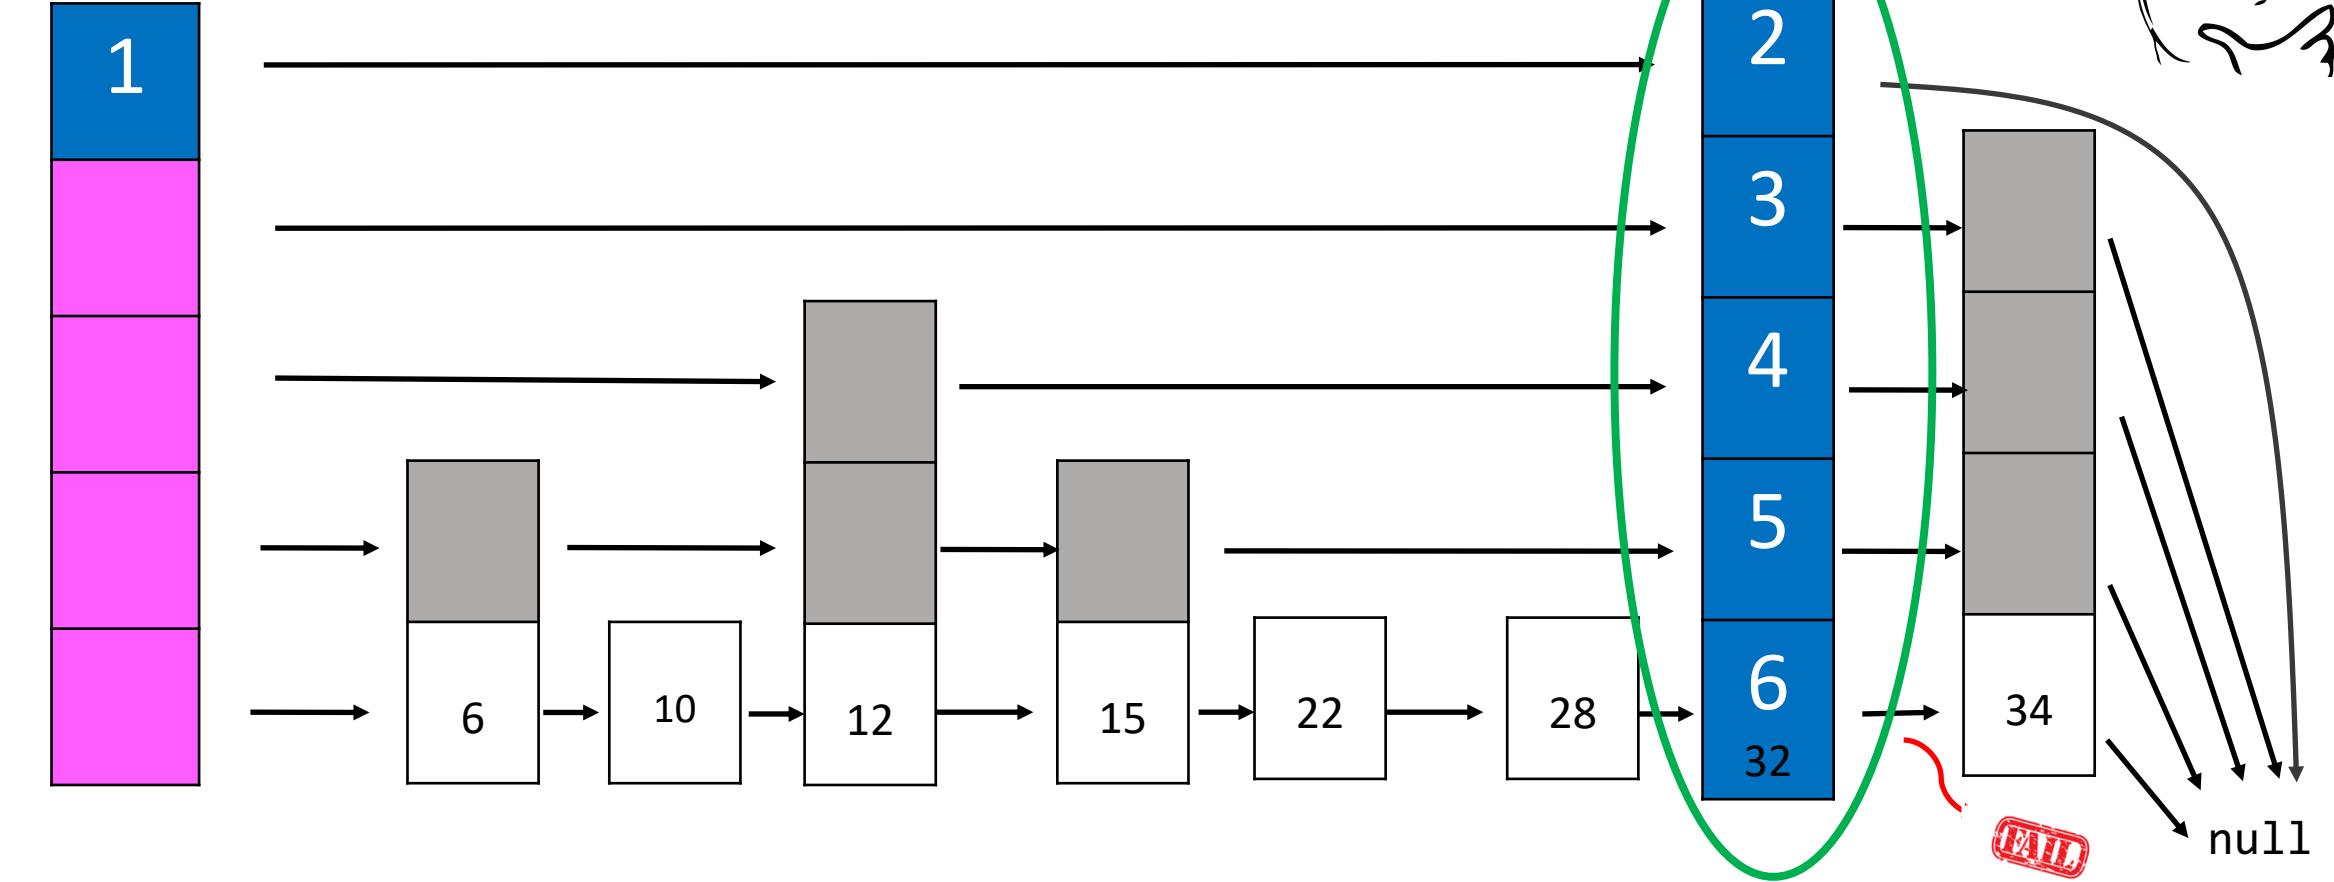
# Skiplist example

Searching for 33 (should  
be a search *miss*)

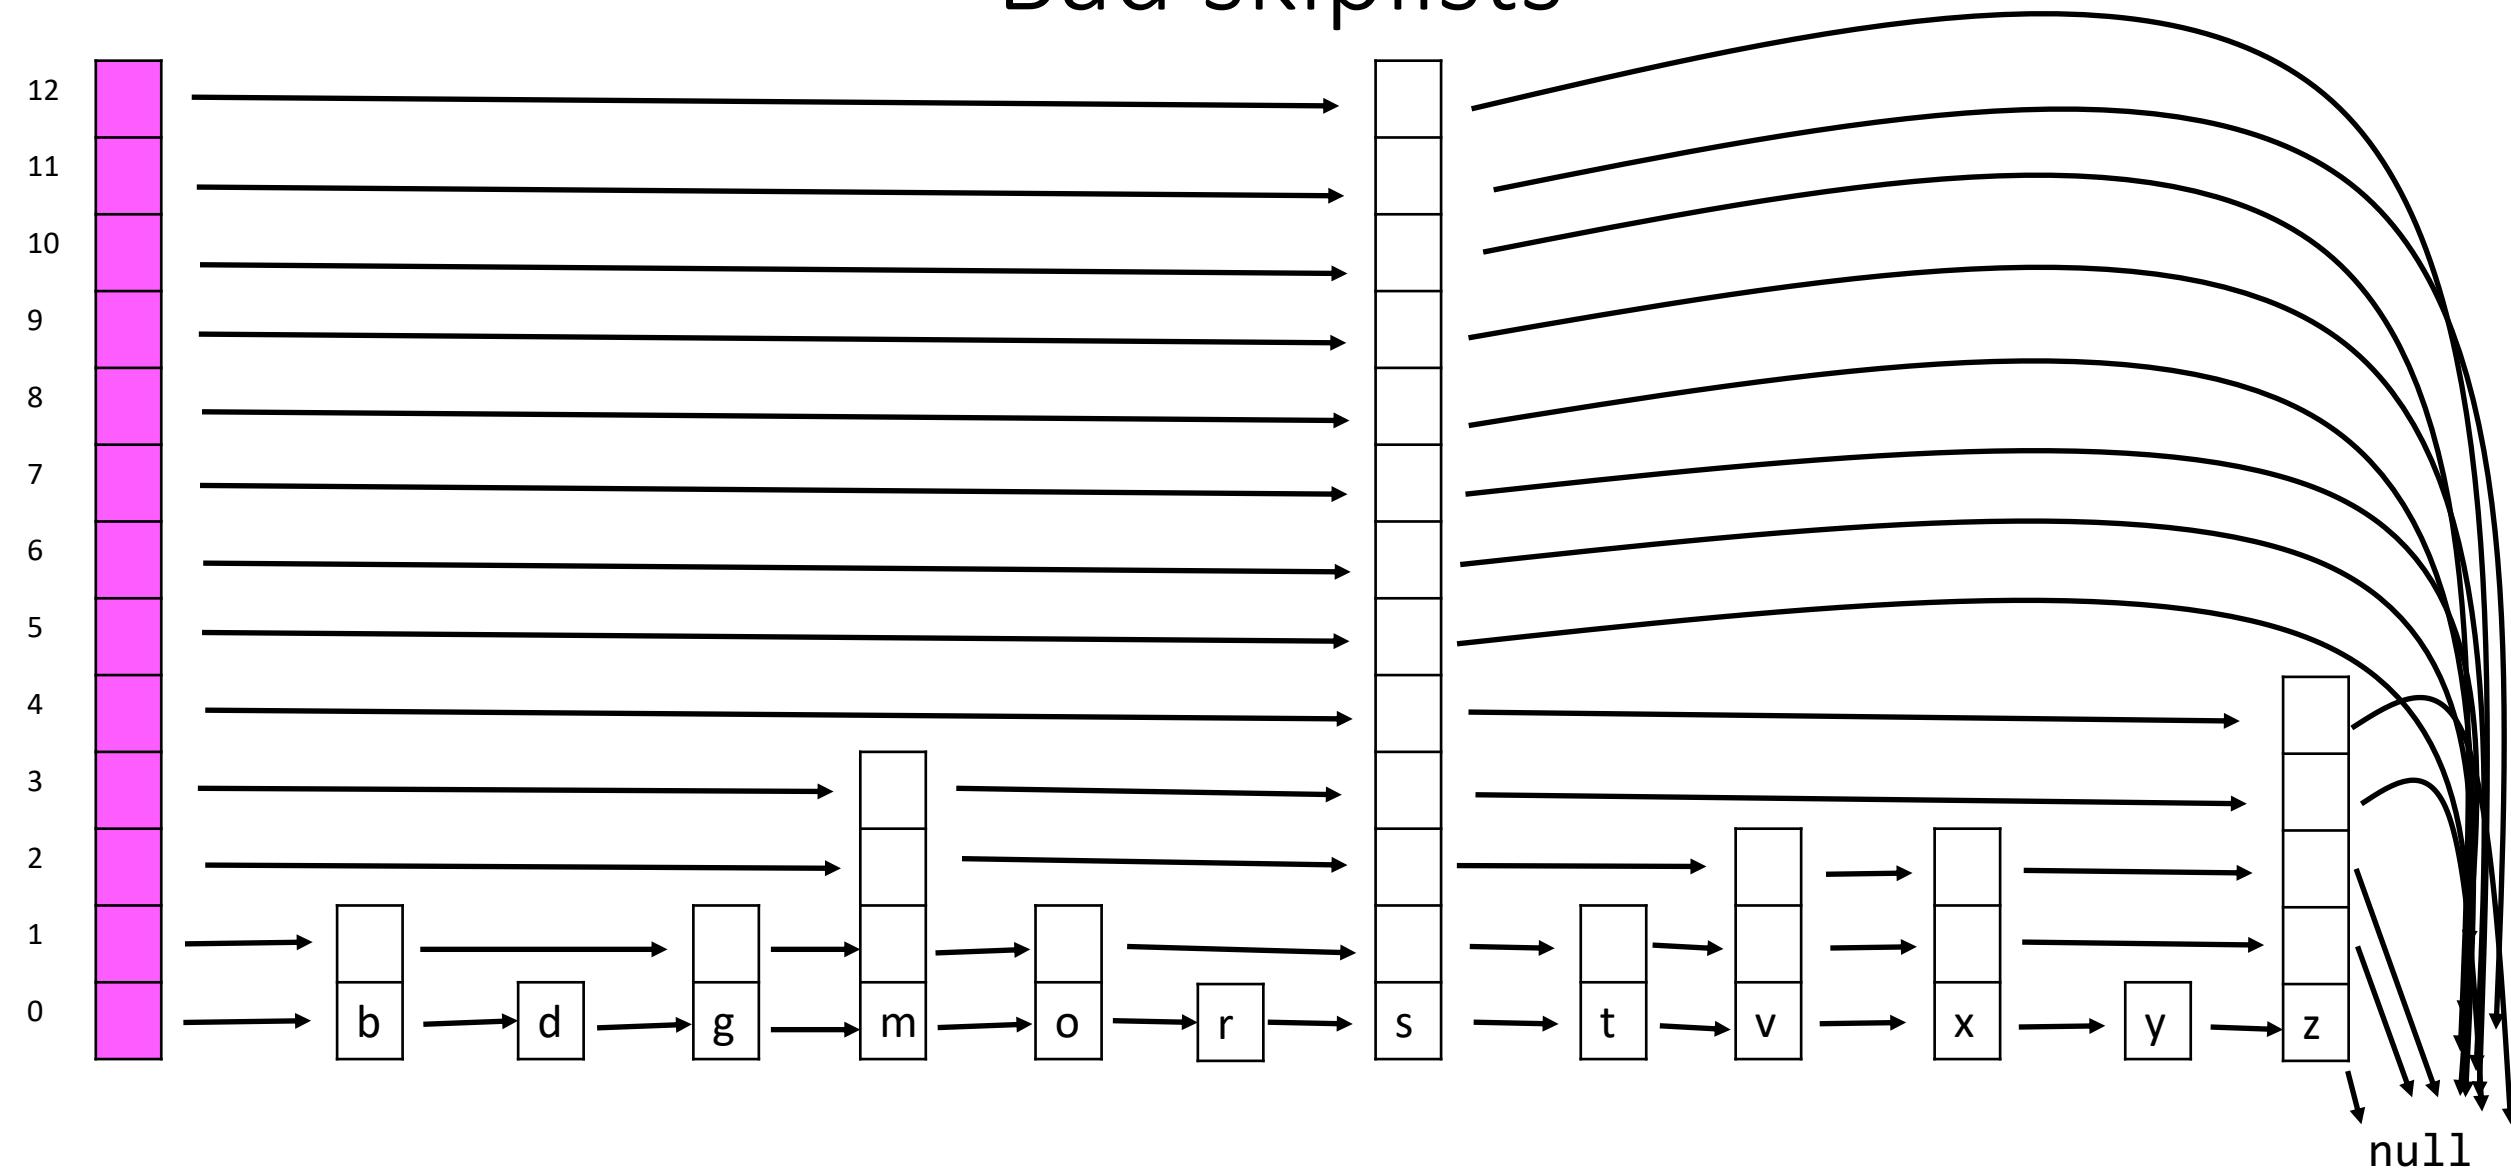


# Skiplist example

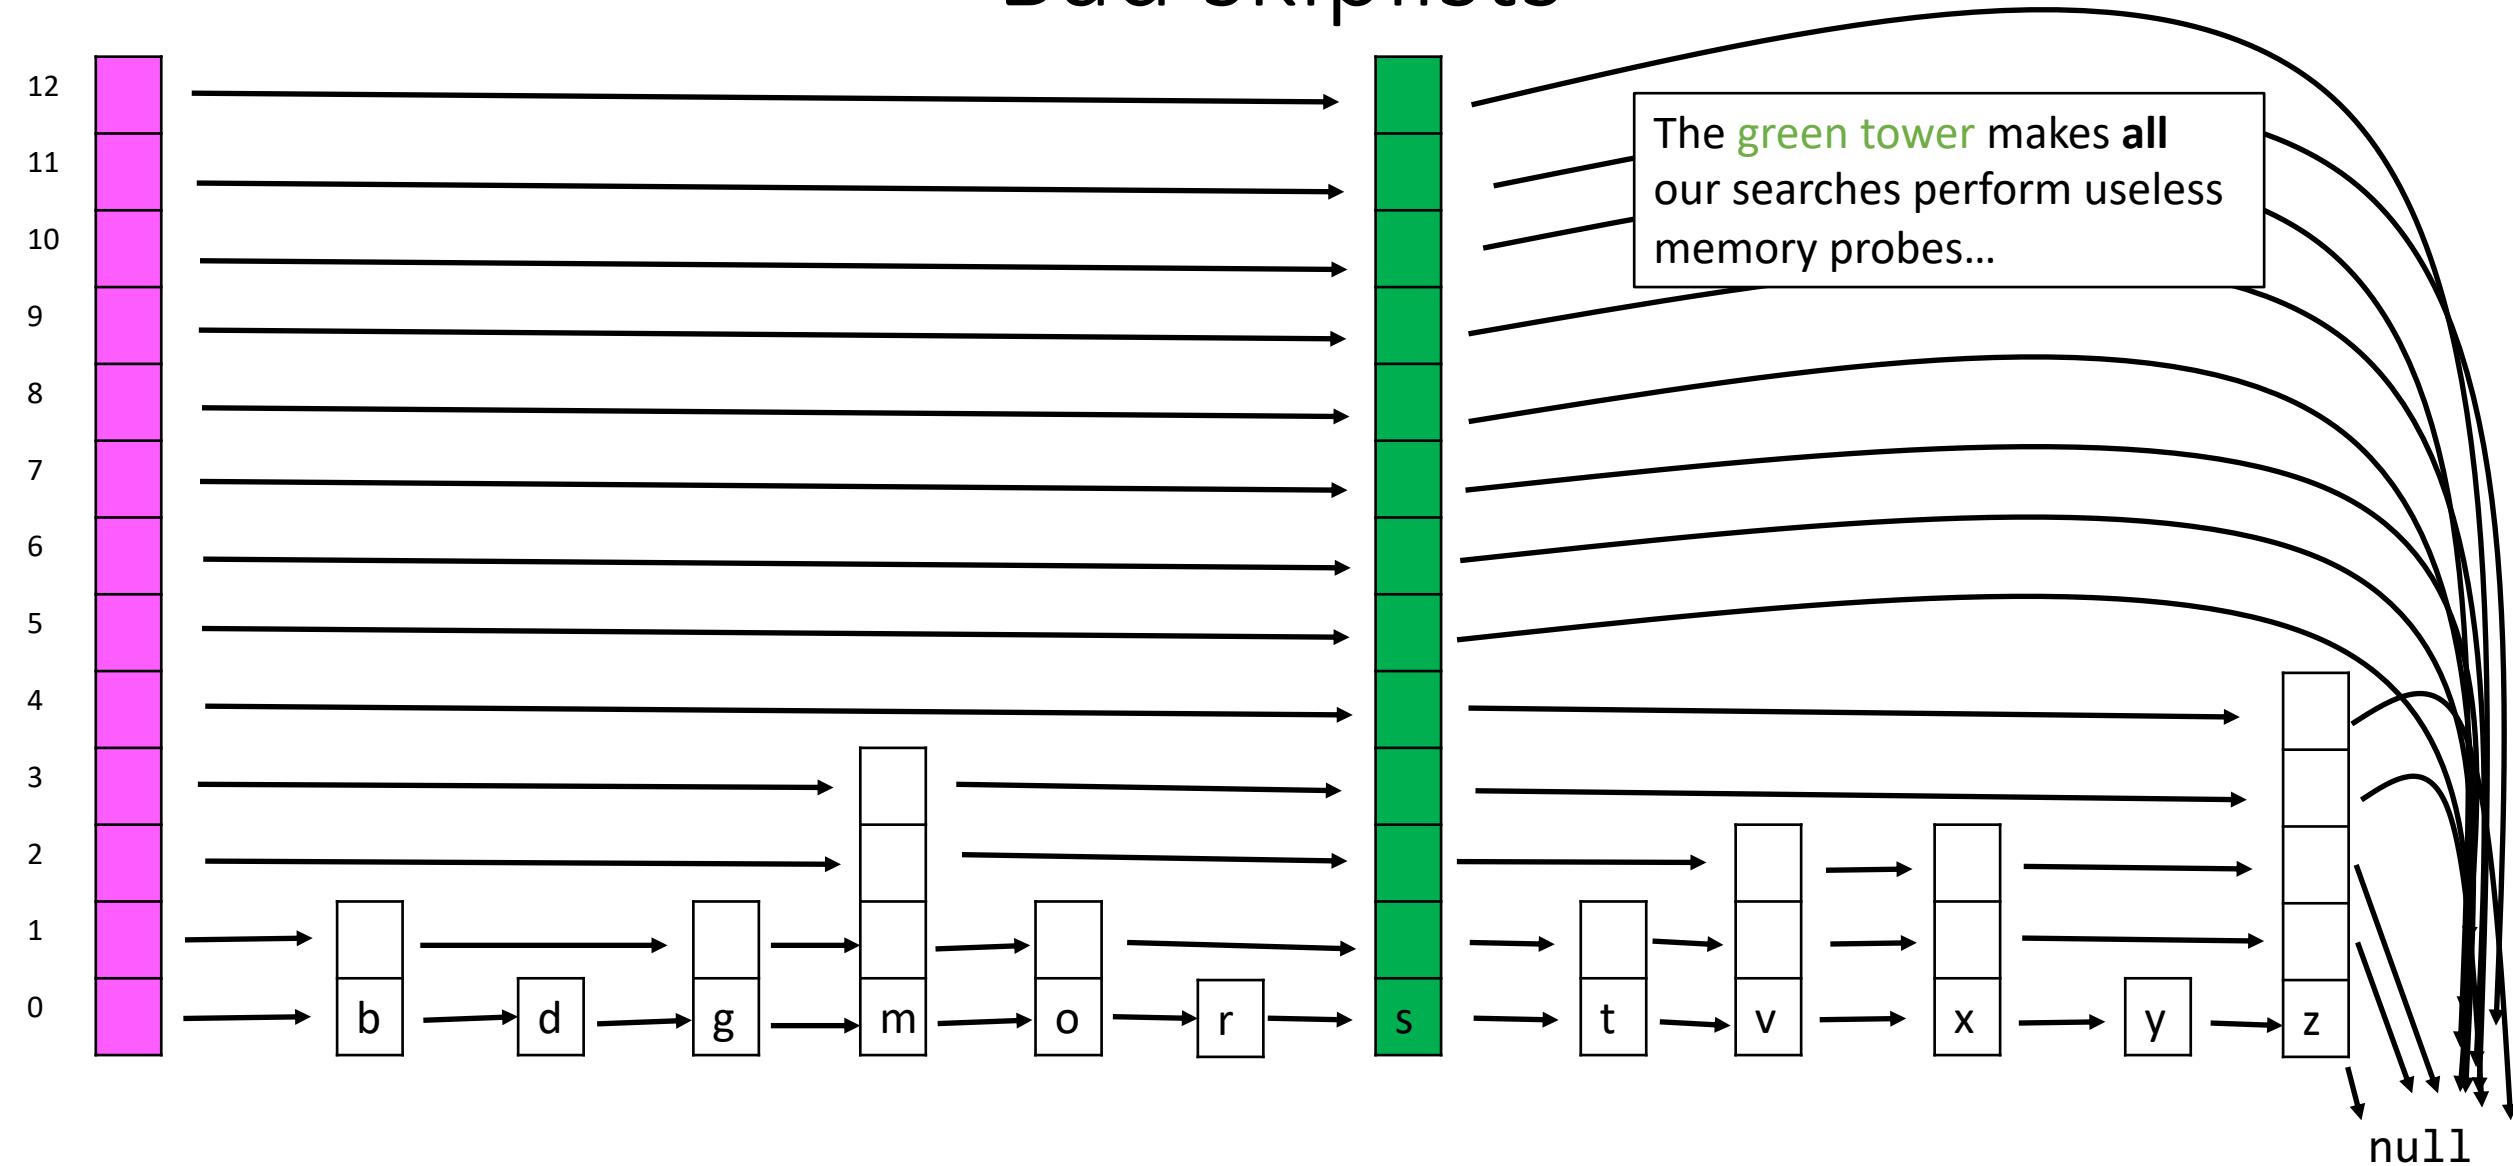
Searching for 33 (should  
be a search *miss*)



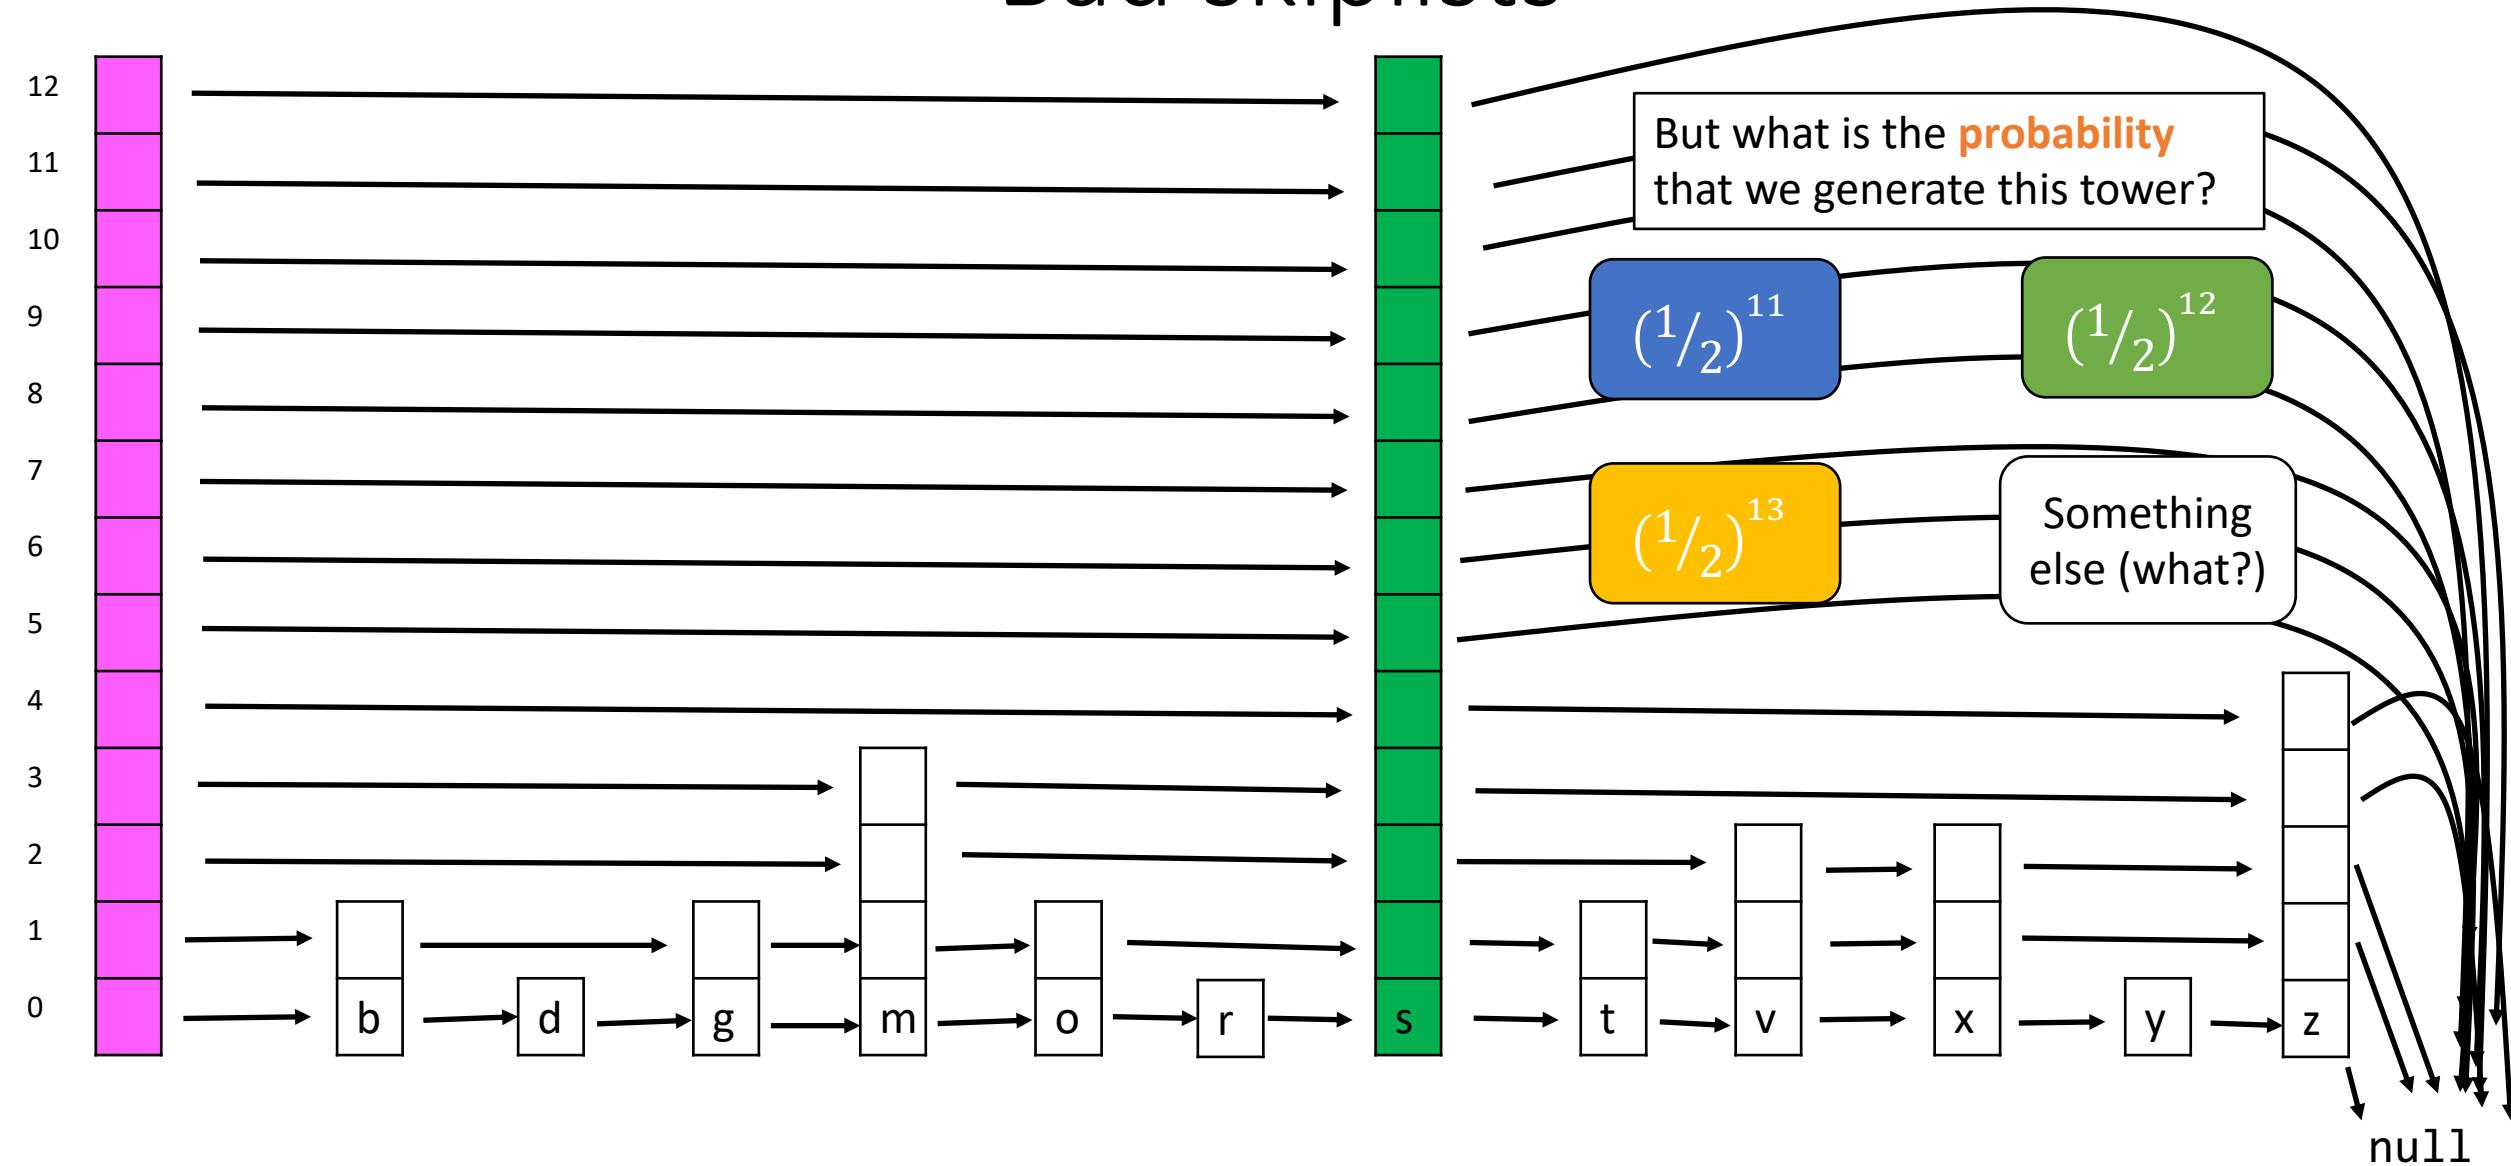
# Bad skip lists



# Bad skiplists



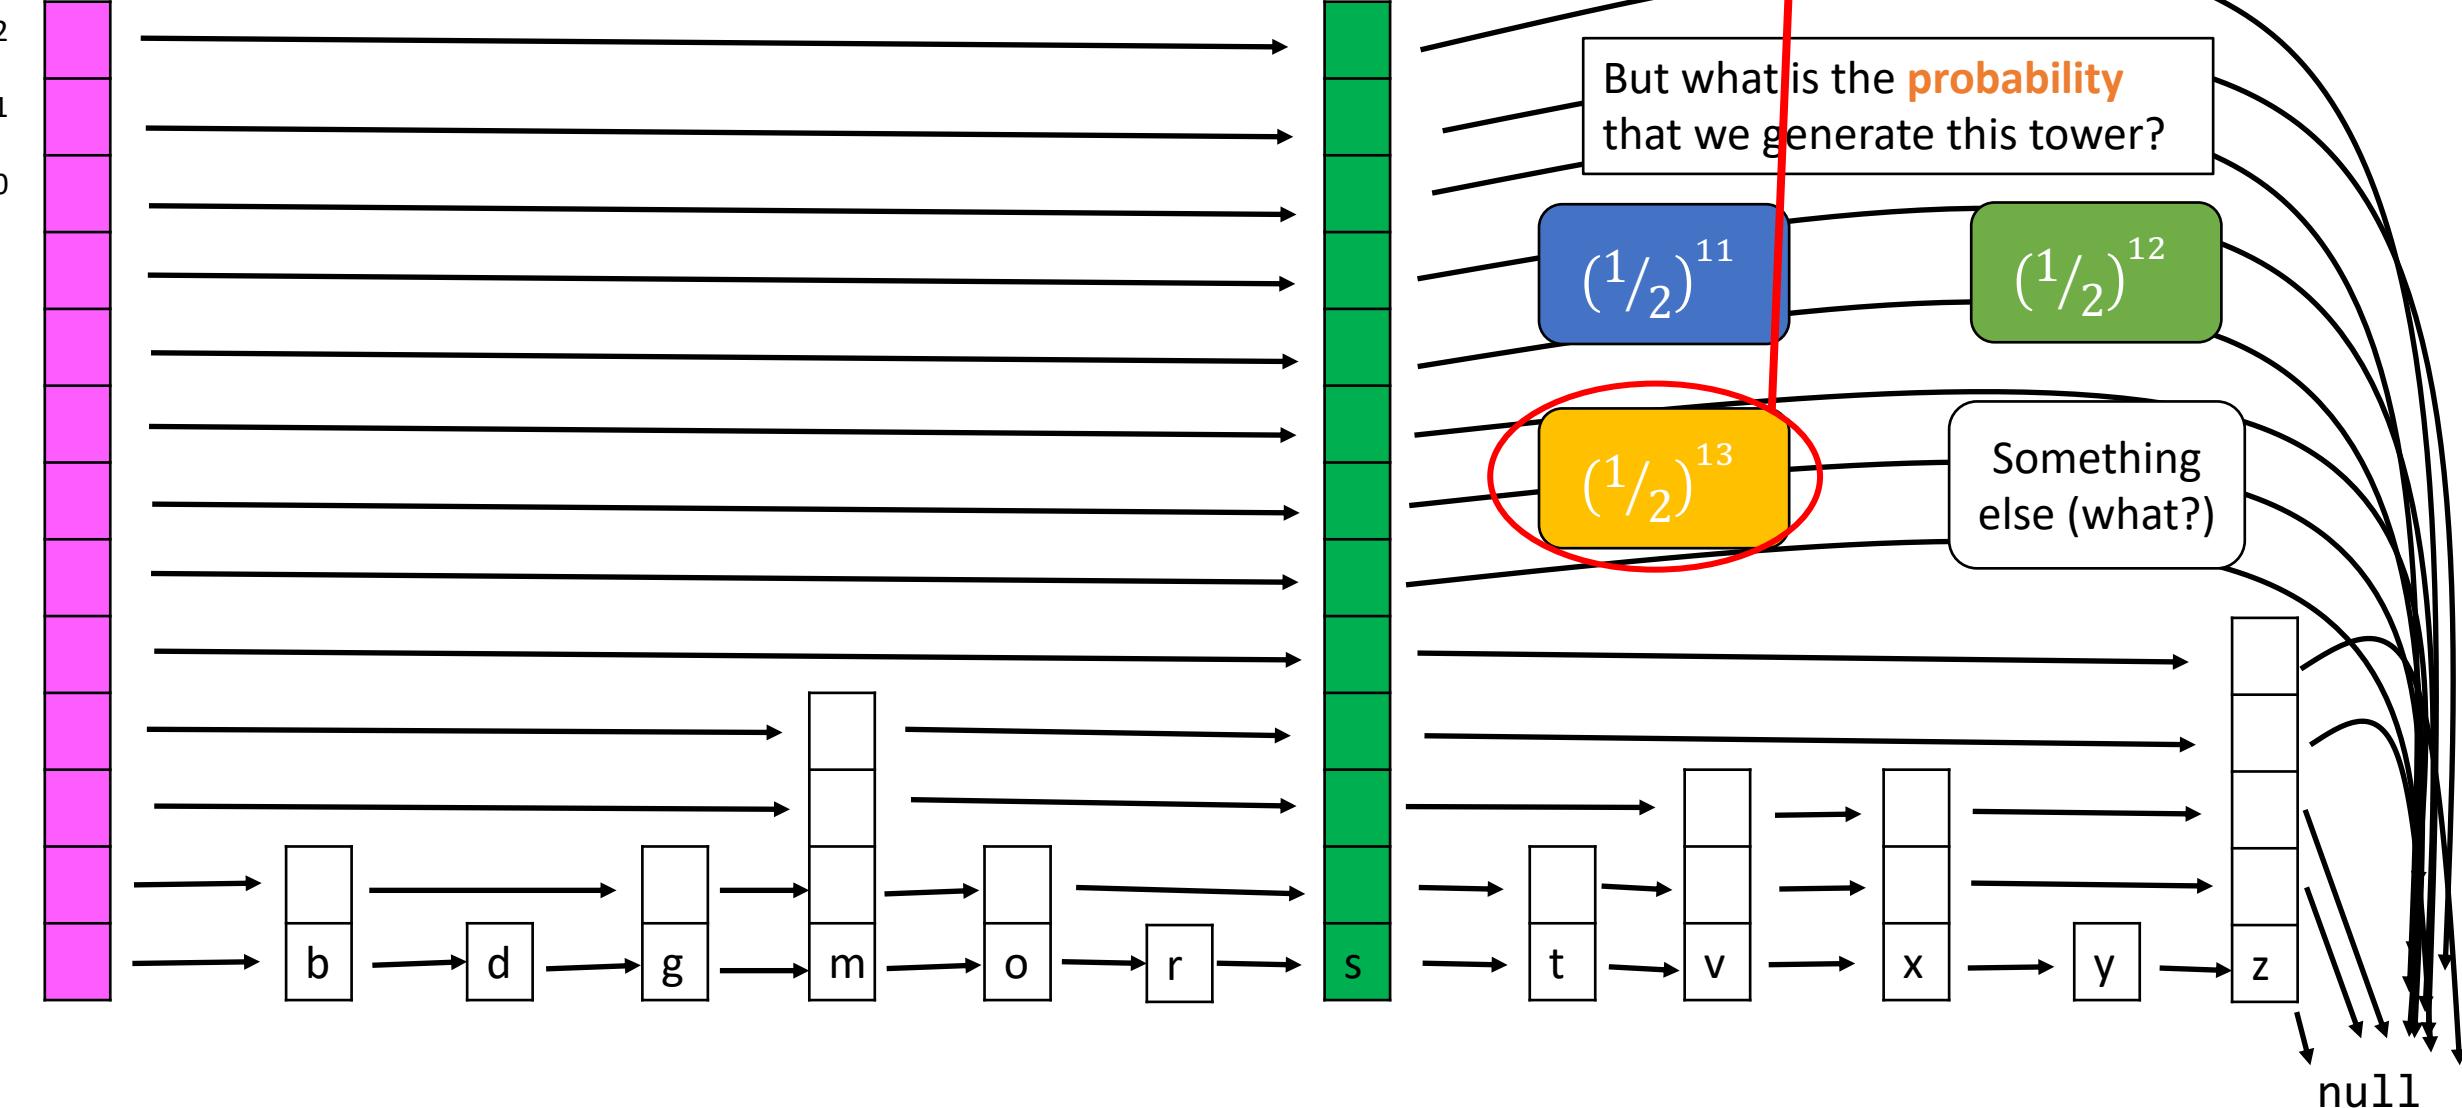
# Bad skiplists



We have 12 flips to make the tower have 13 levels  
and another to keep it there!

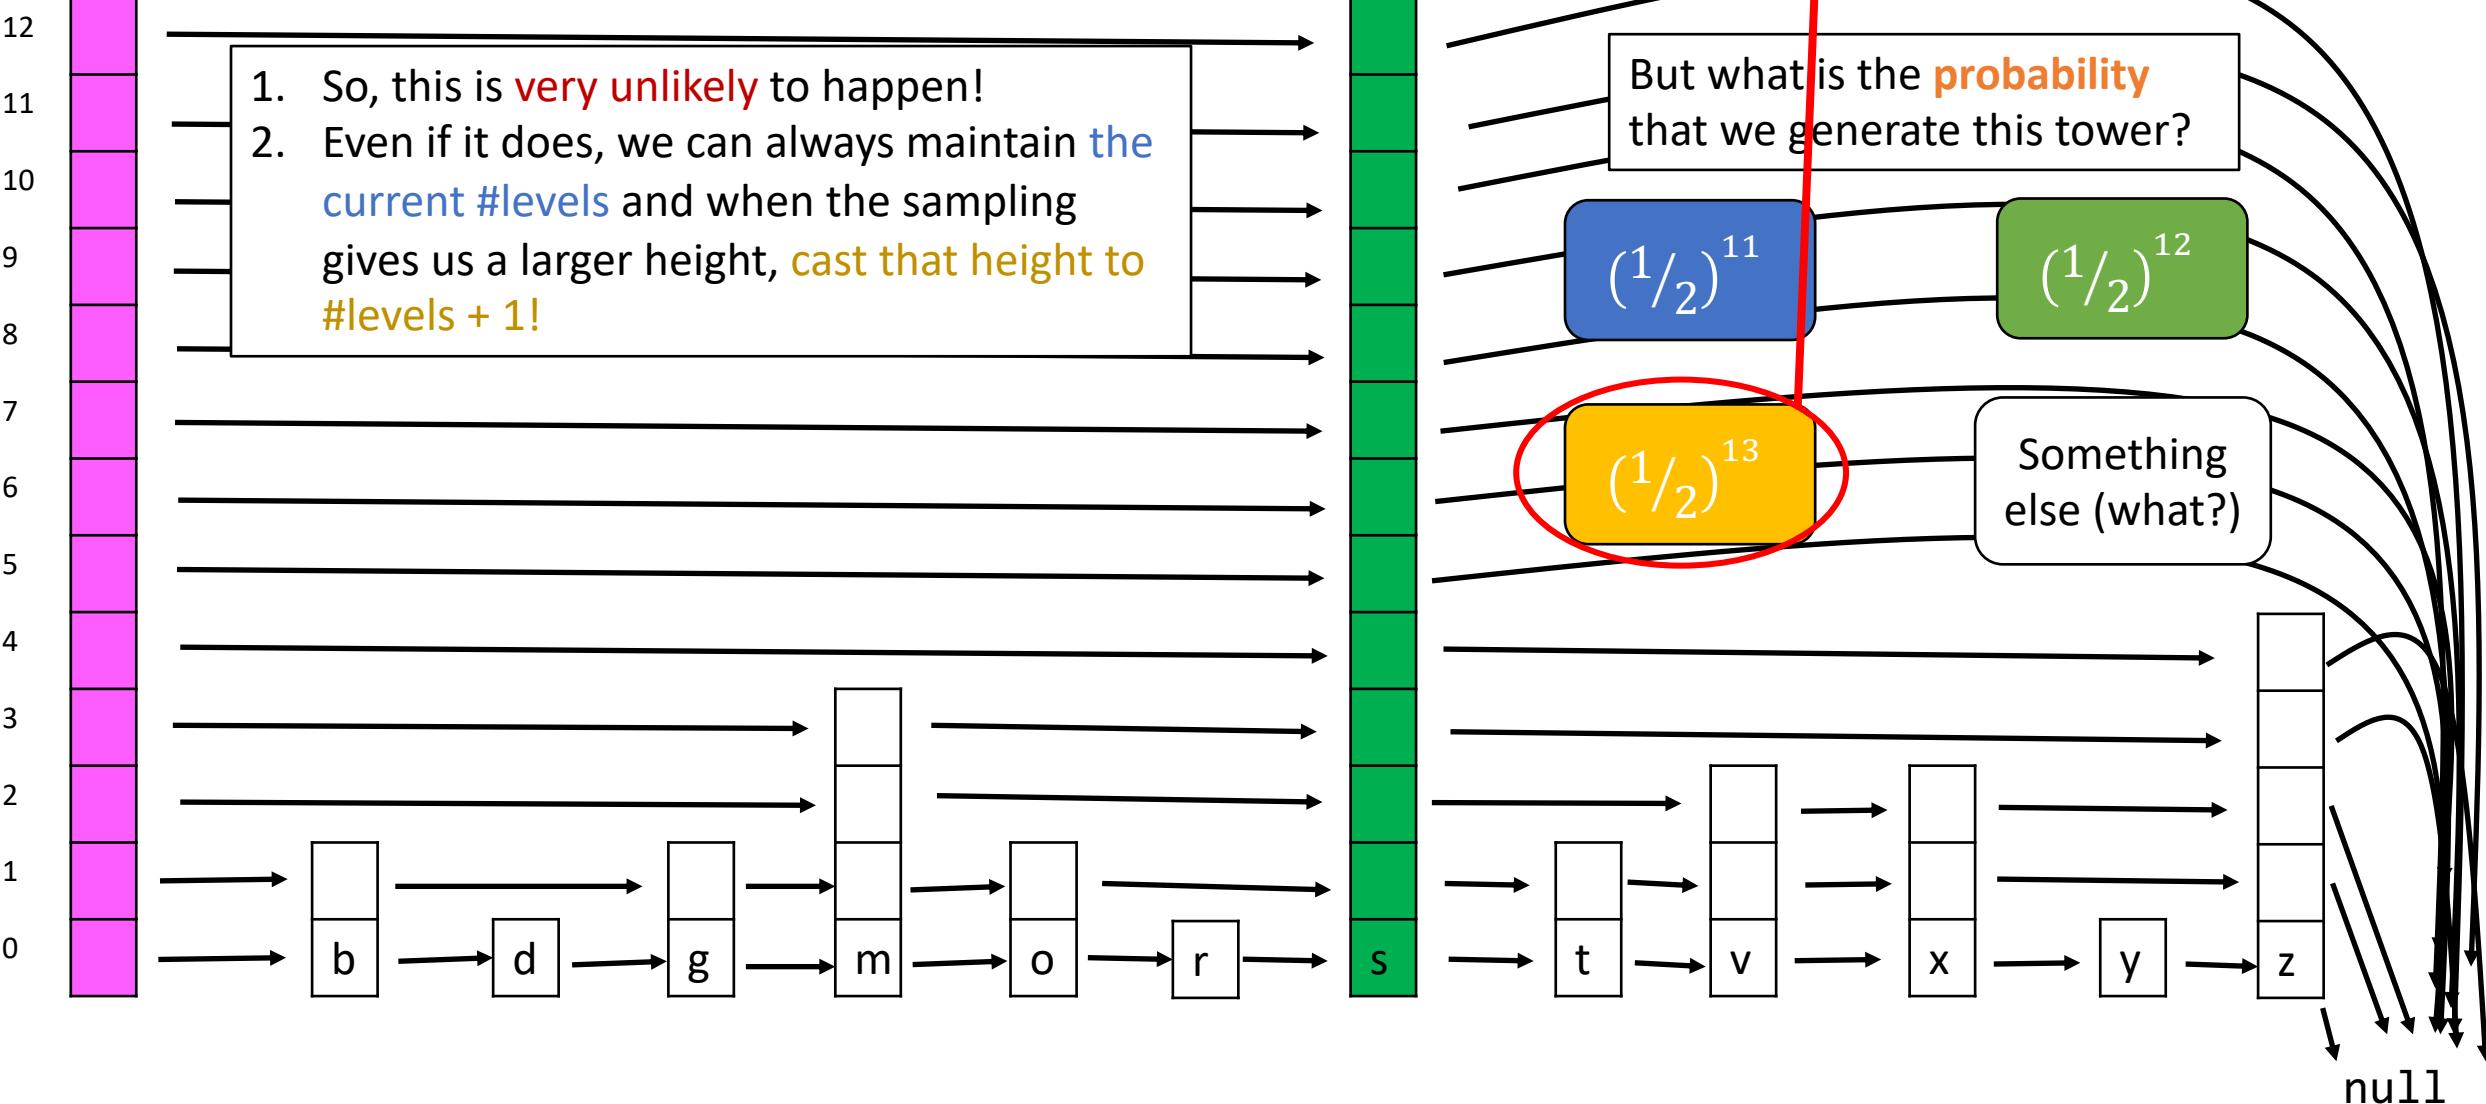
# Bad skiplists

= 0.012207031%

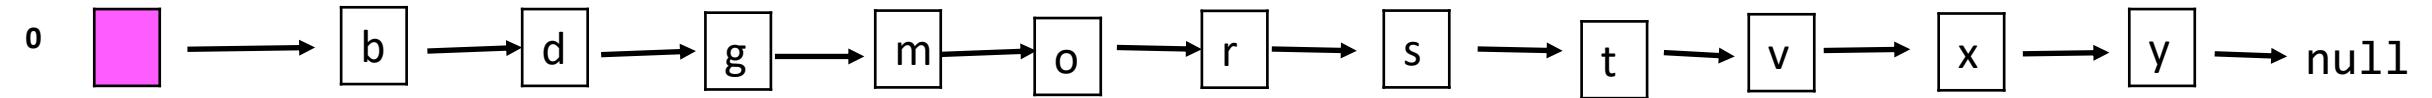


We have 12 flips to make the tower have 13 levels  
and another to keep it there!

# Bad skiplists

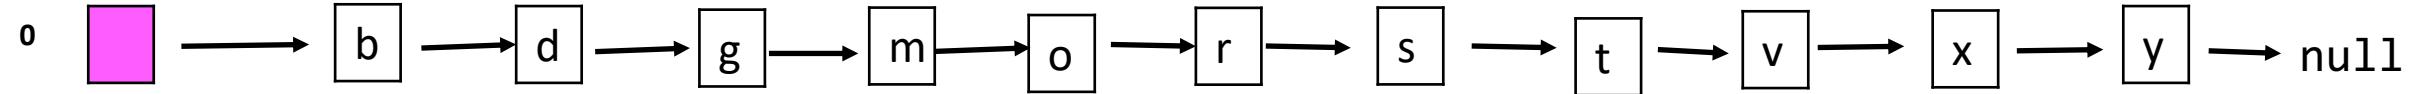


# Bad skiplists



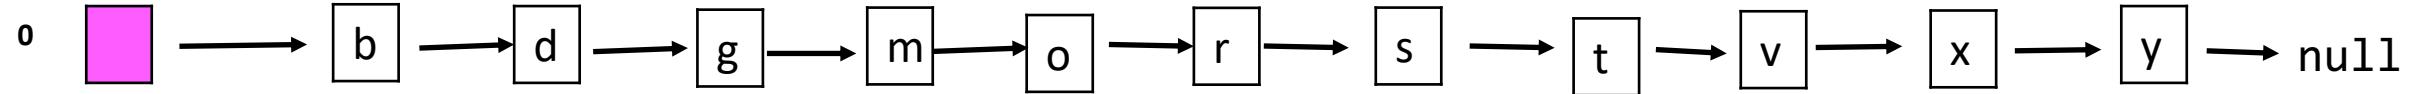
# Bad skiplists

- **Degenerate skiplist:** is equivalent to the linked list of its elements.



# Bad skiplists

- **Degenerate** skiplist: is equivalent to the linked list of its elements.
- But **how likely** is such a skiplist?



# Bad skiplists

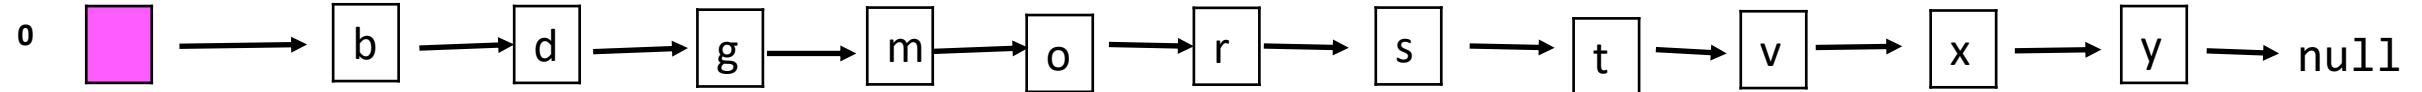
- **Degenerate skiplist:** is equivalent to the linked list of its elements.
- But **how likely** is such a skiplist?

$(1/2)^{11}$

$(1/2)^{12}$

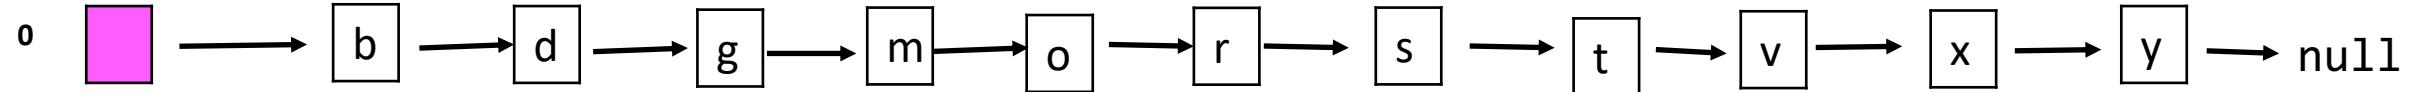
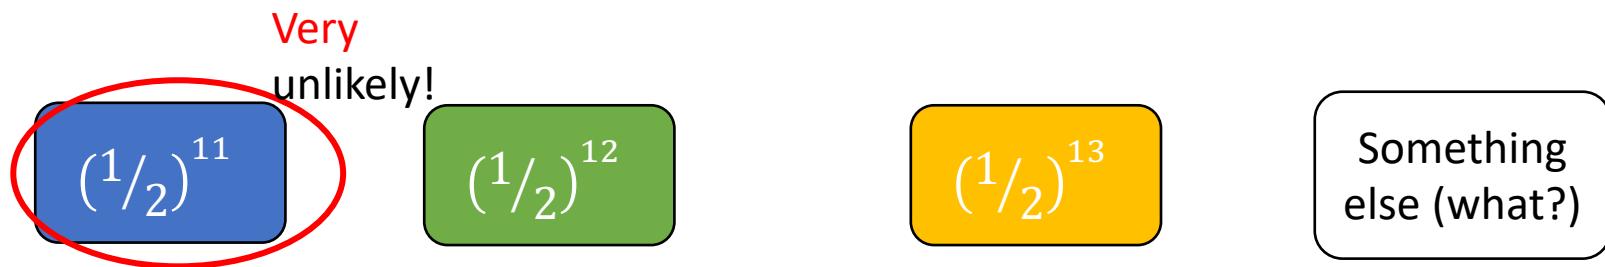
$(1/2)^{13}$

Something  
else (what?)



# Bad skiplists

- **Degenerate skiplist:** is equivalent to the linked list of its elements.
- But **how likely** is such a skiplist?



# Insertion

- Insertion into a skip list follows the “**search and splice**” rule. Steps:

# Insertion

- Insertion into a skip list follows the “**search and splice**” rule. Steps:
  1. Randomly sample a **height** for the new node using a **geometric distribution**.

# Insertion

- Insertion into a skip list follows the “**search and splice**” rule. Steps:
  1. Randomly sample a **height** for the new node using a **geometric distribution**.
  2. “**Search**” for the key in order to find the **appropriate position** to insert it into. While you’re descending into the skip list, maintain a **levels**-sized 1D array called **update** which will hold the pointers from the towers that we descended from at every level.

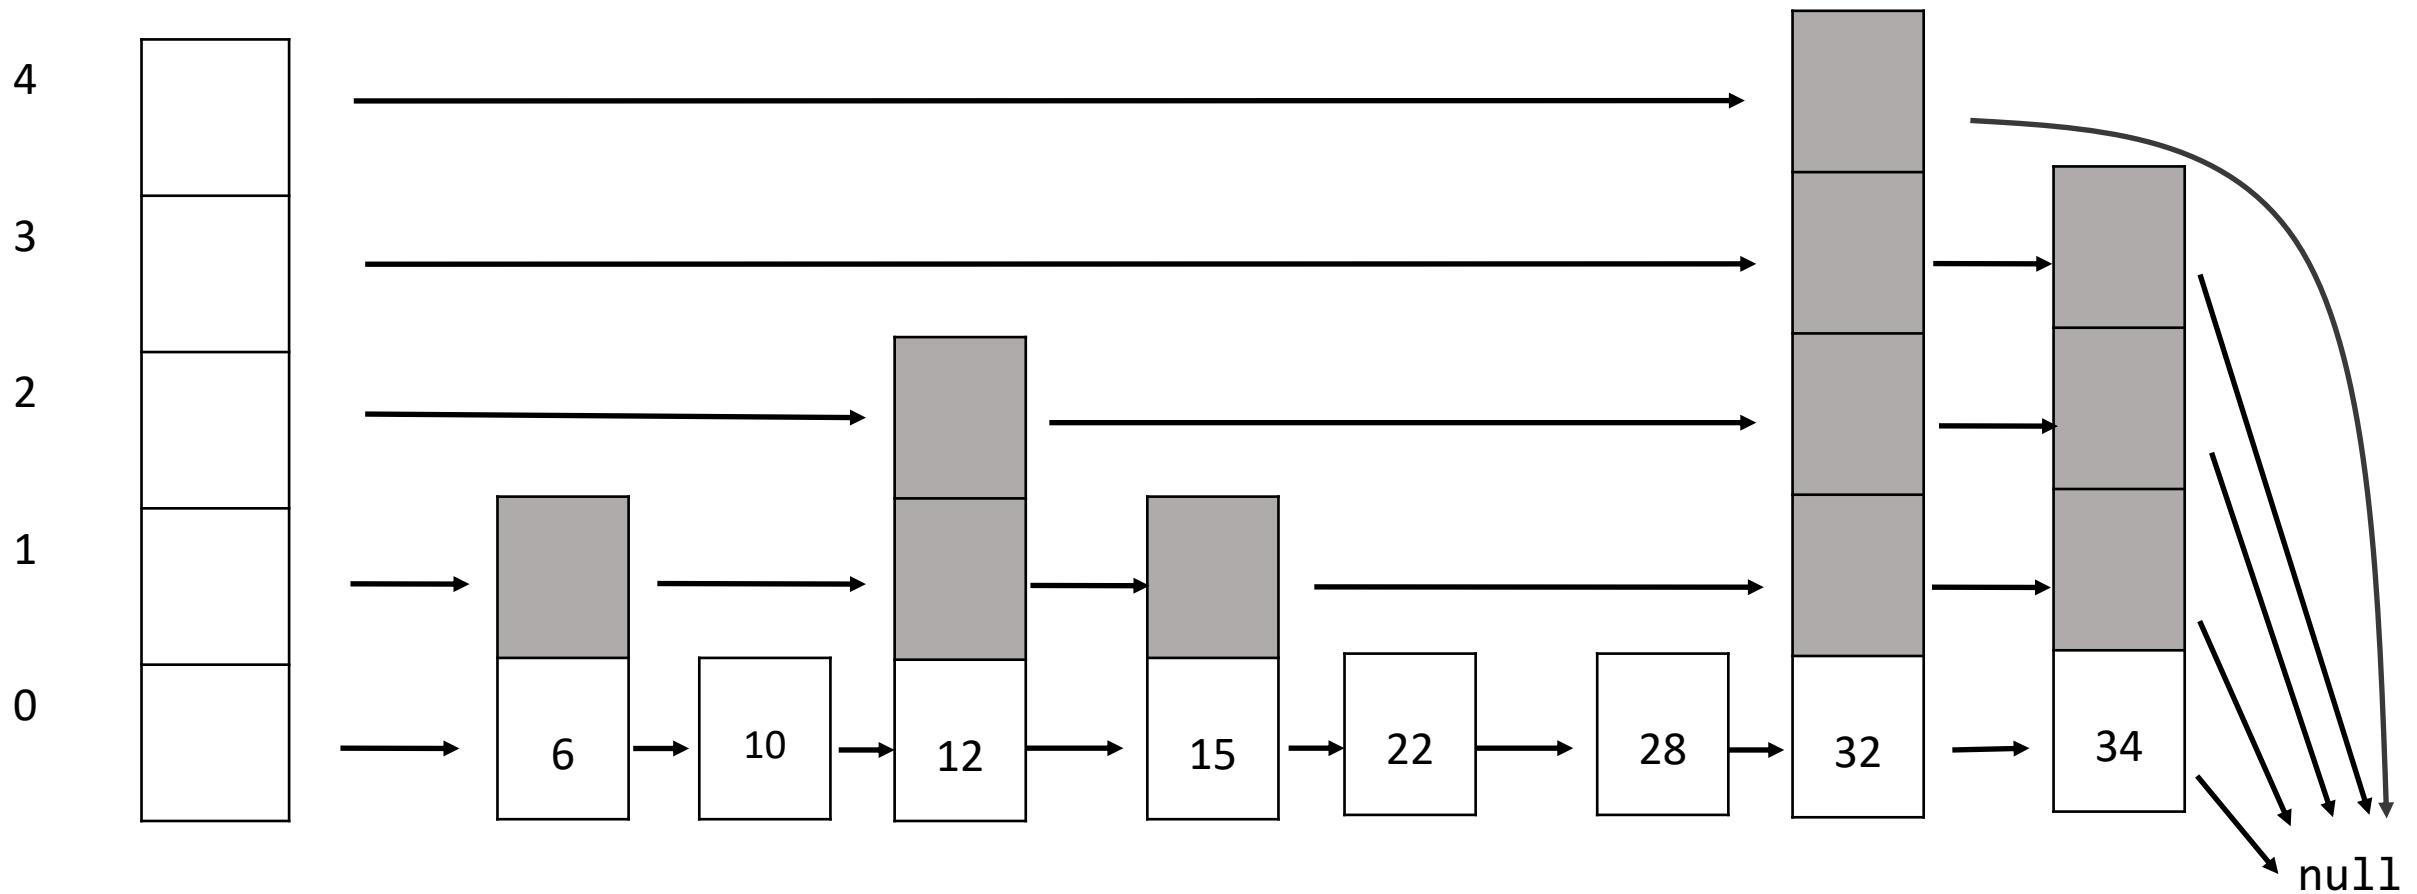
# Insertion

- Insertion into a skip list follows the “**search and splice**” rule. Steps:
  1. Randomly sample a **height** for the new node using a geometric distribution.
  2. “**Search**” for the key in order to find the **appropriate position** to insert it into. While you’re descending into the skip list, maintain a **levels**-sized 1D array called **update** which will hold the pointers from the towers that we descended from at every level.
  3. When you find the position, **splice**: insert the tower in place by using **update** to update the relevant pointers.

# Insertion

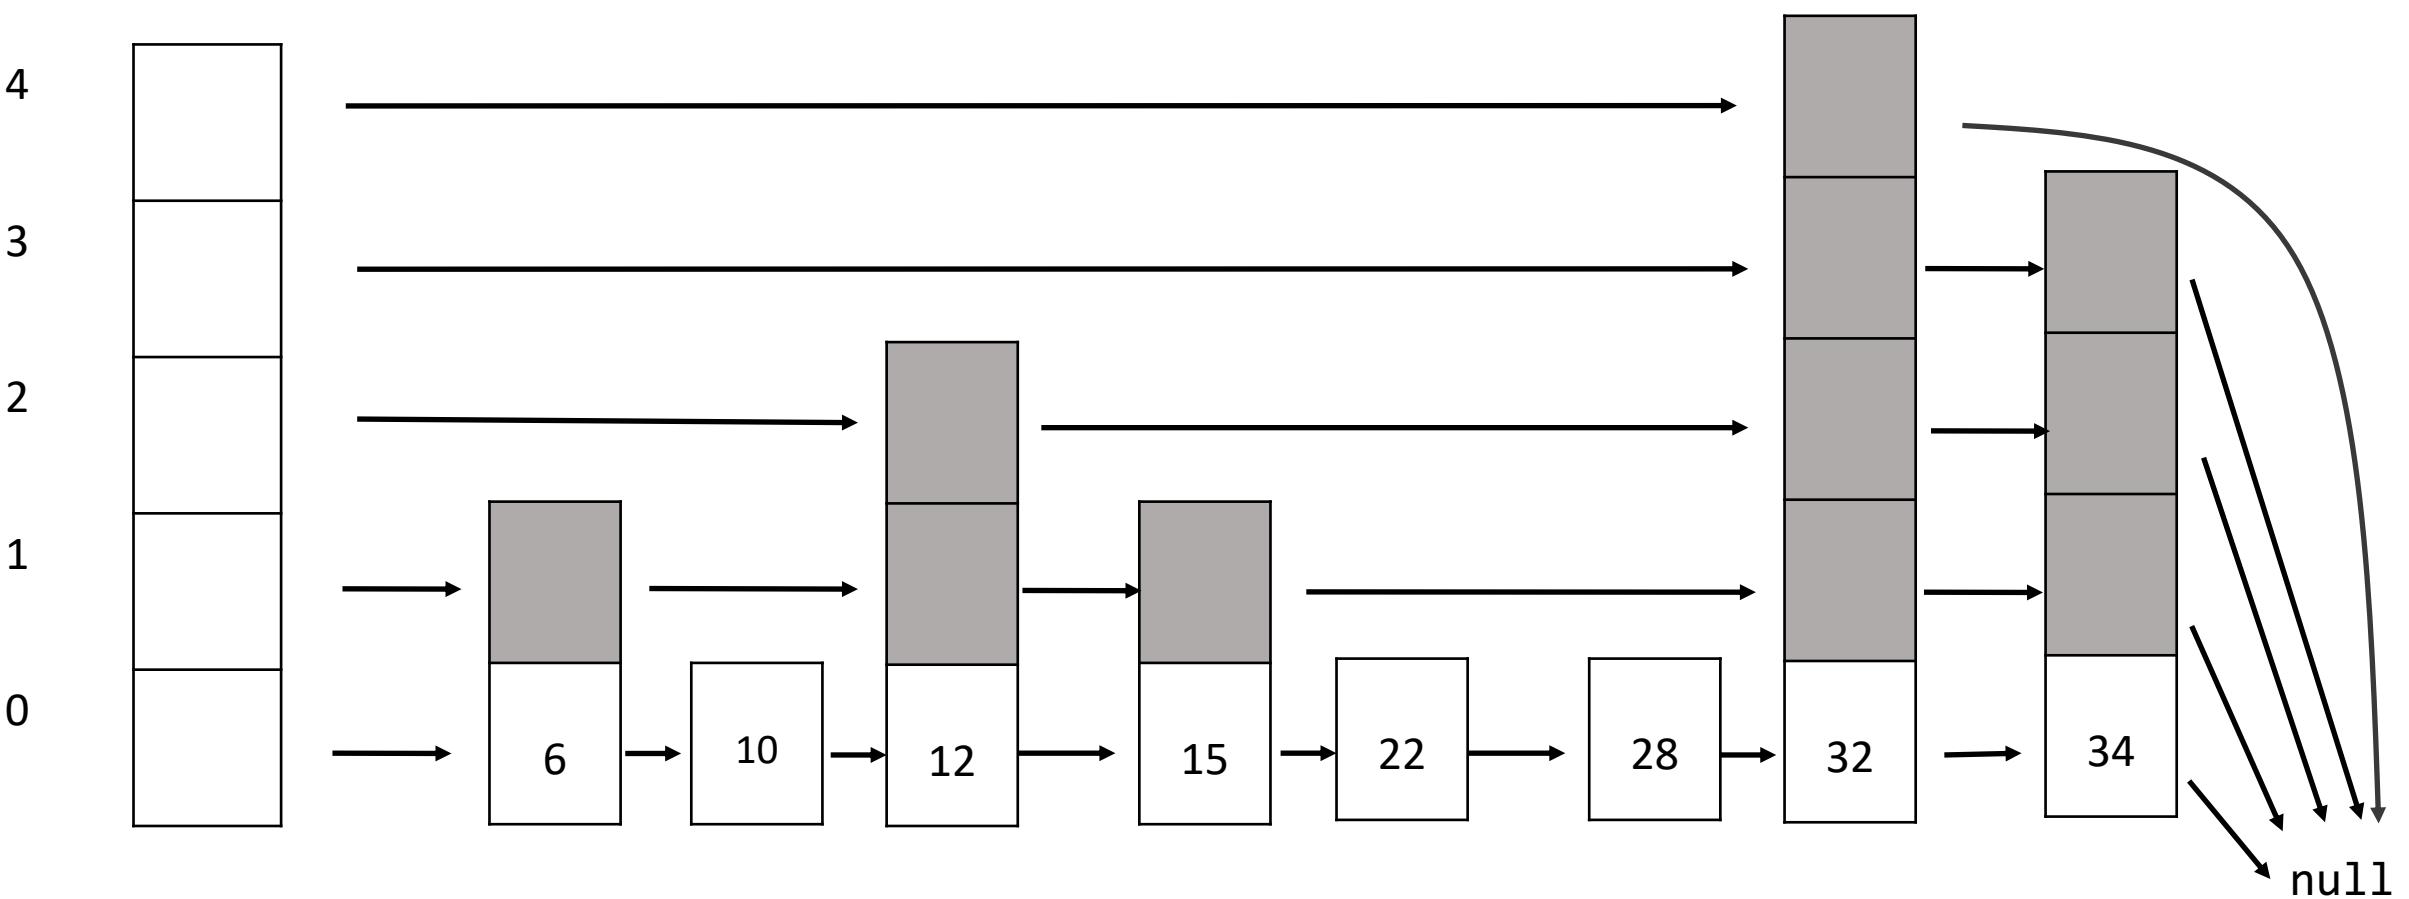
- Insertion into a skip list follows the “**search and splice**” rule. Steps:
  1. Randomly sample a **height** for the new node using a geometric distribution.
  2. “**Search**” for the key in order to find the **appropriate position** to insert it into. While you’re descending into the skip list, maintain a **levels**-sized 1D array called **update** which will hold the pointers from the towers that we descended from at every level.
  3. When you find the position, **splice**: insert the tower in place by using **update** to update the relevant pointers.
- We employ **zero-indexing**: the list level is at level 0, one level above is level 1, etc.

# Insertion Example



# Insertion Example

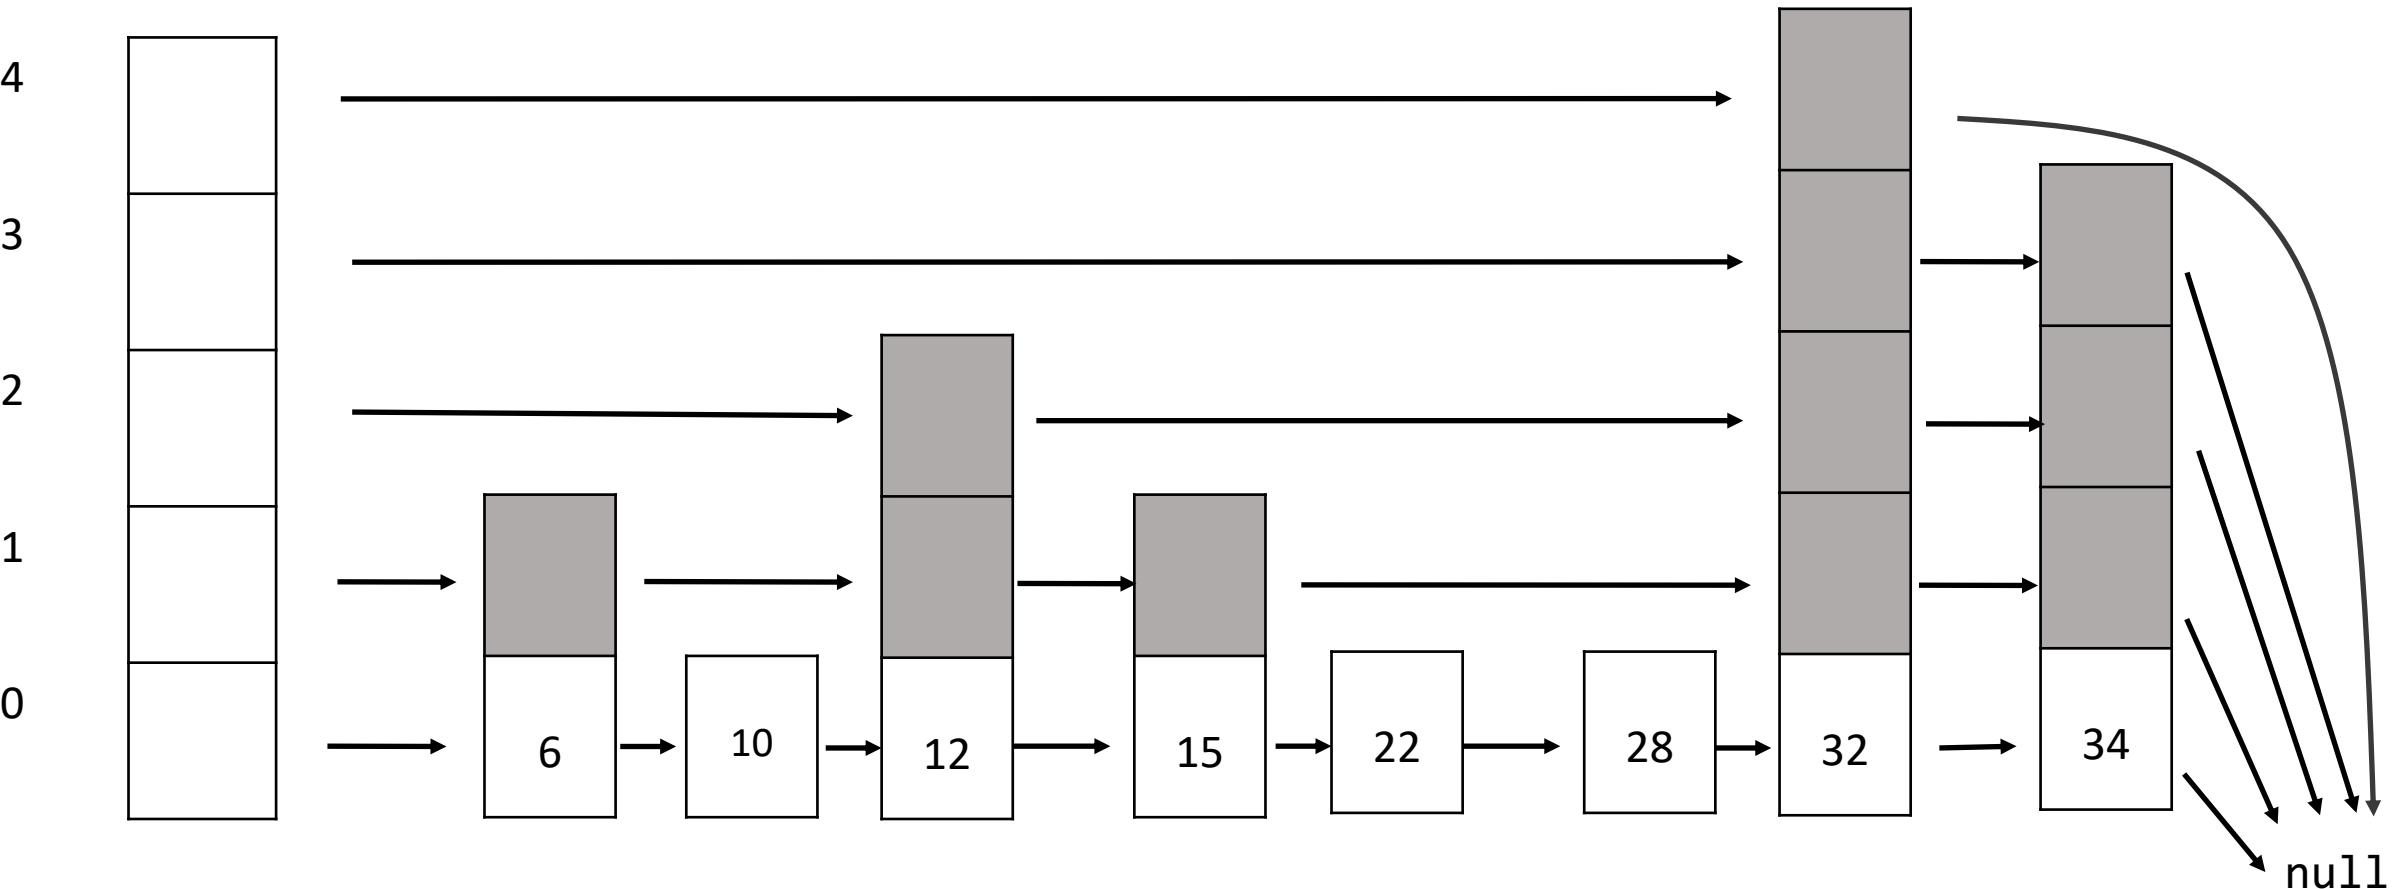
Suppose we want to insert 25.



Task: Insert 25

# Insertion Example

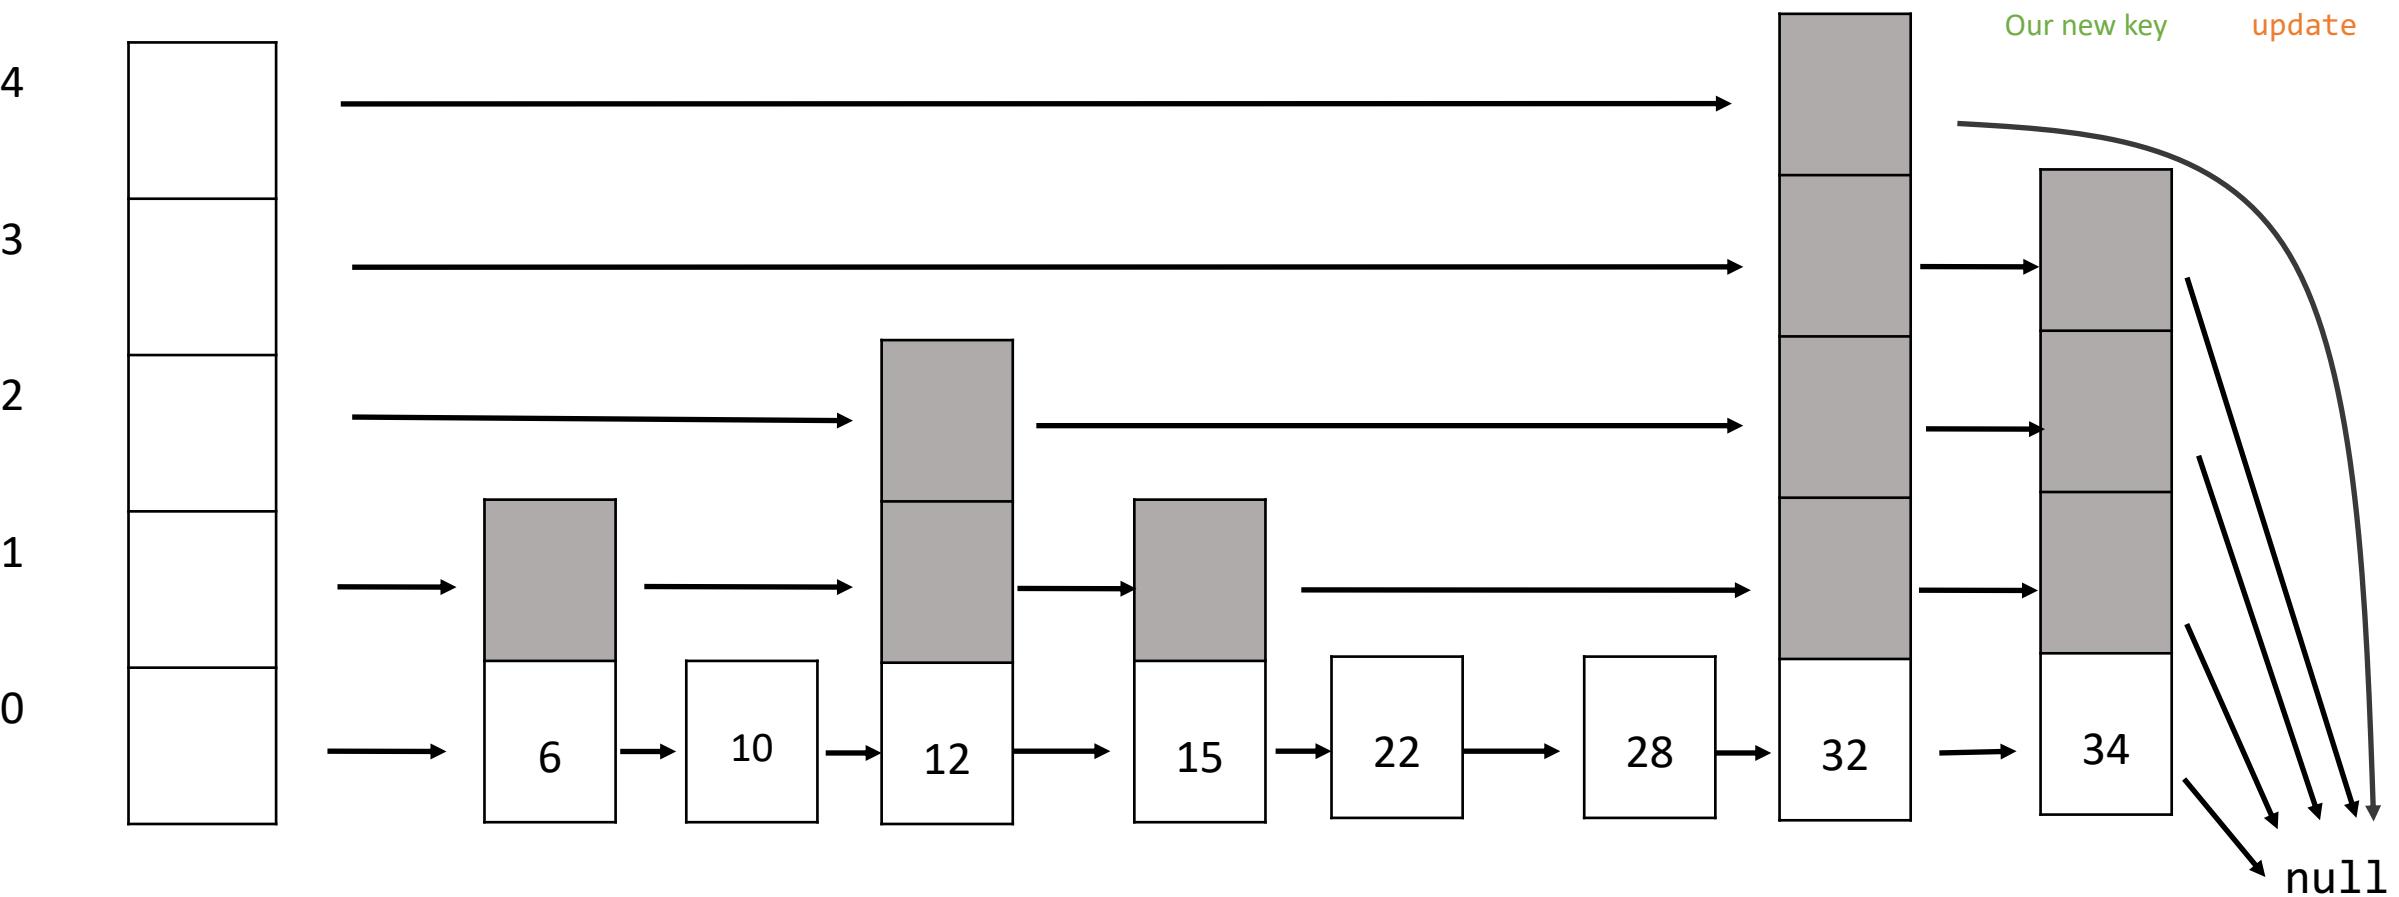
Step 1: Sample a #levels “on top” of the key. Suppose we sample 3 (probability?)



Task: Insert 25

# Insertion Example

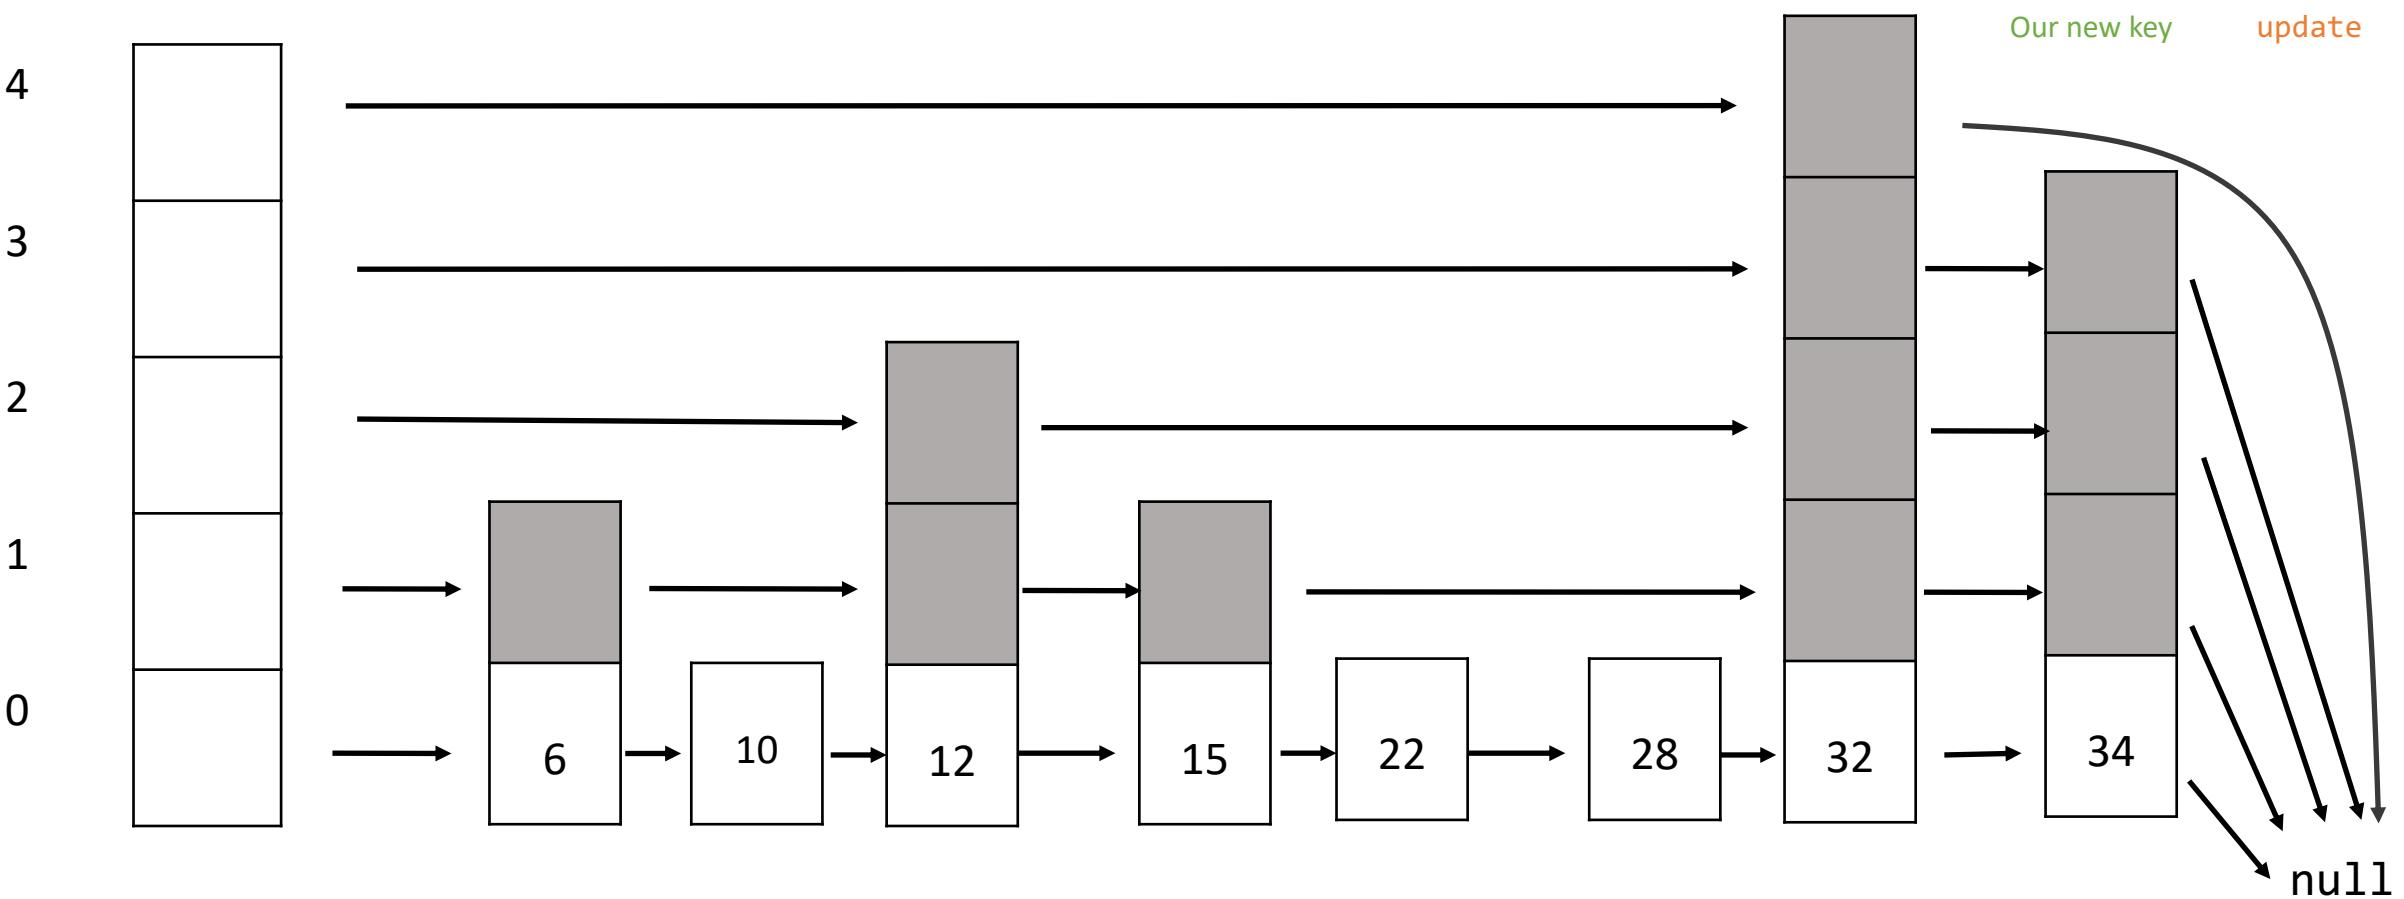
Step 1: Sample a #levels “on top” of the key. Suppose we sample 3 (probability?)



Task: Insert 25

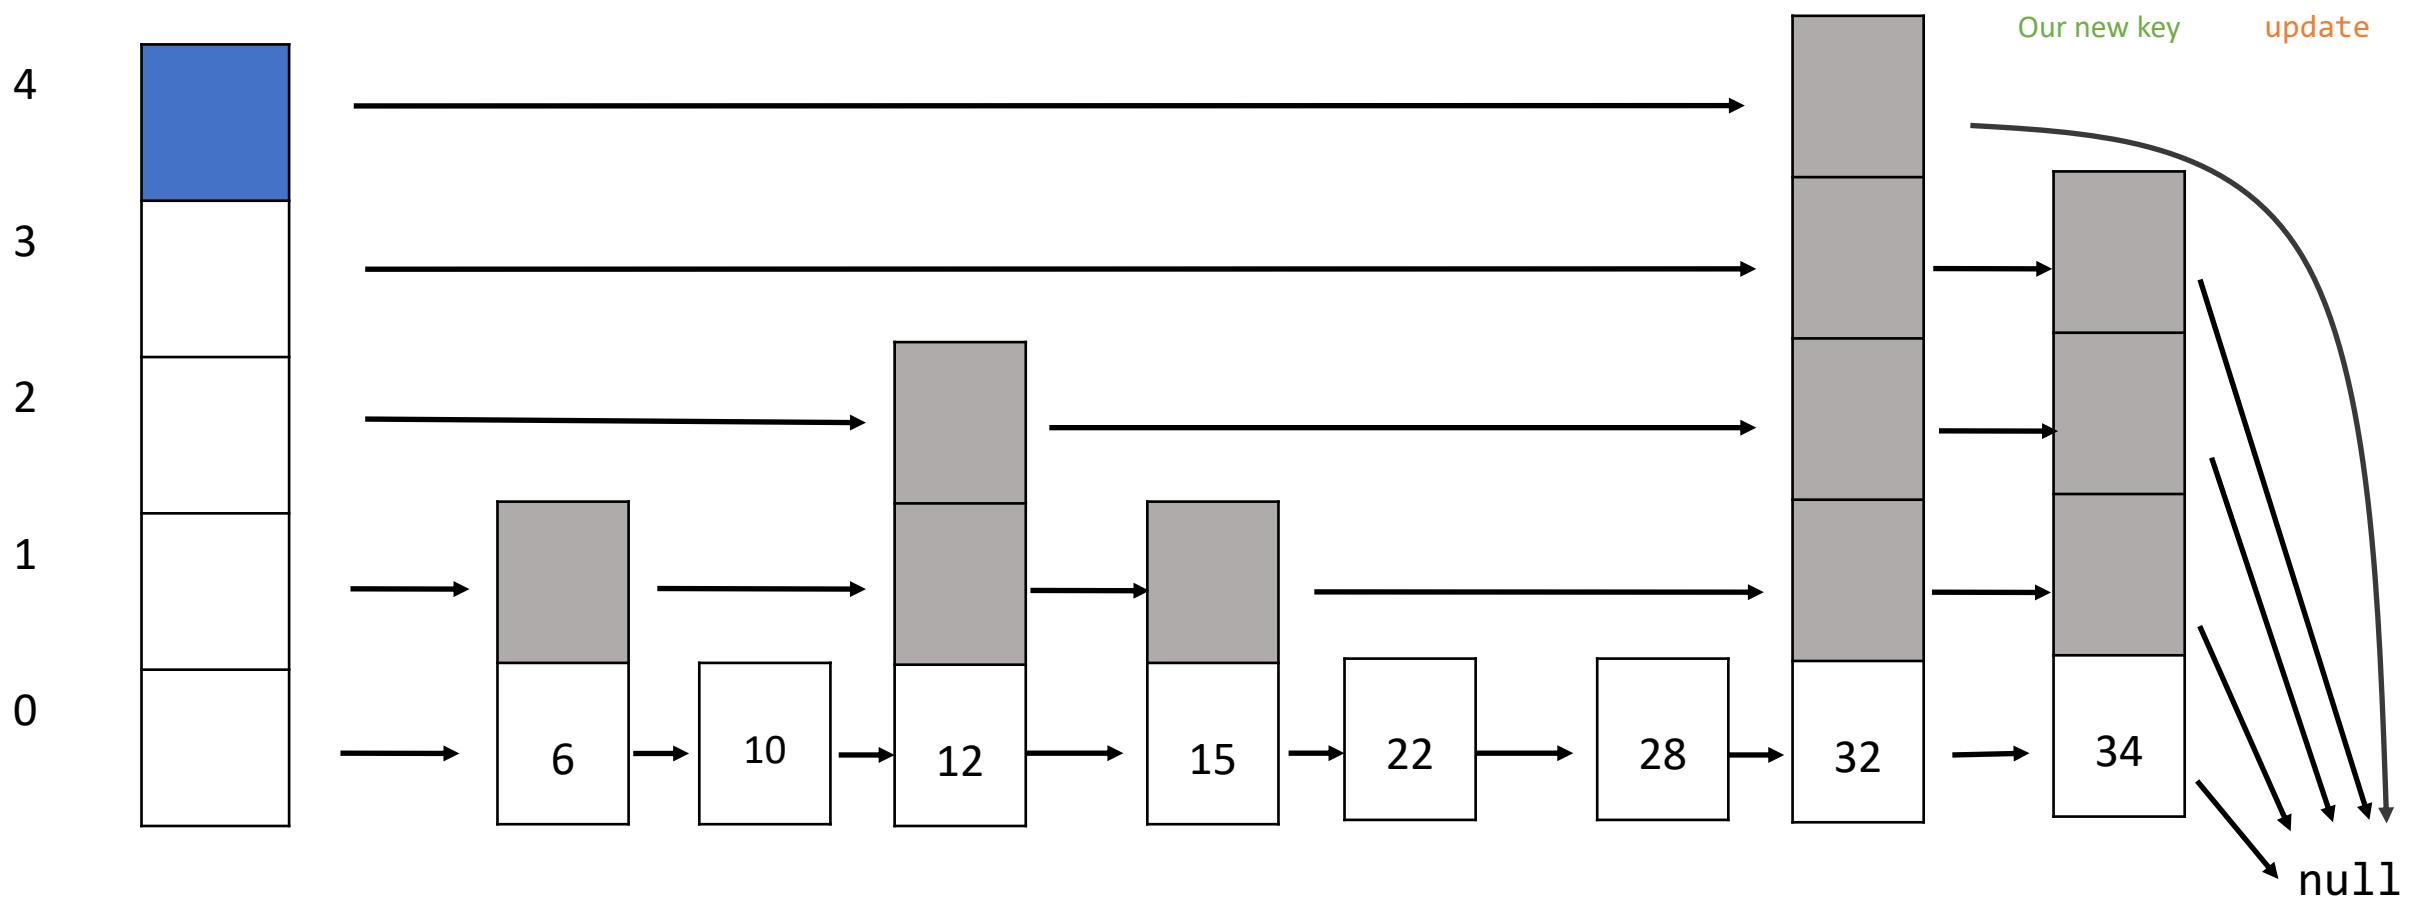
# Insertion Example

Step 2: Search for 25. We expect this search to fail (assuming no duplicates)



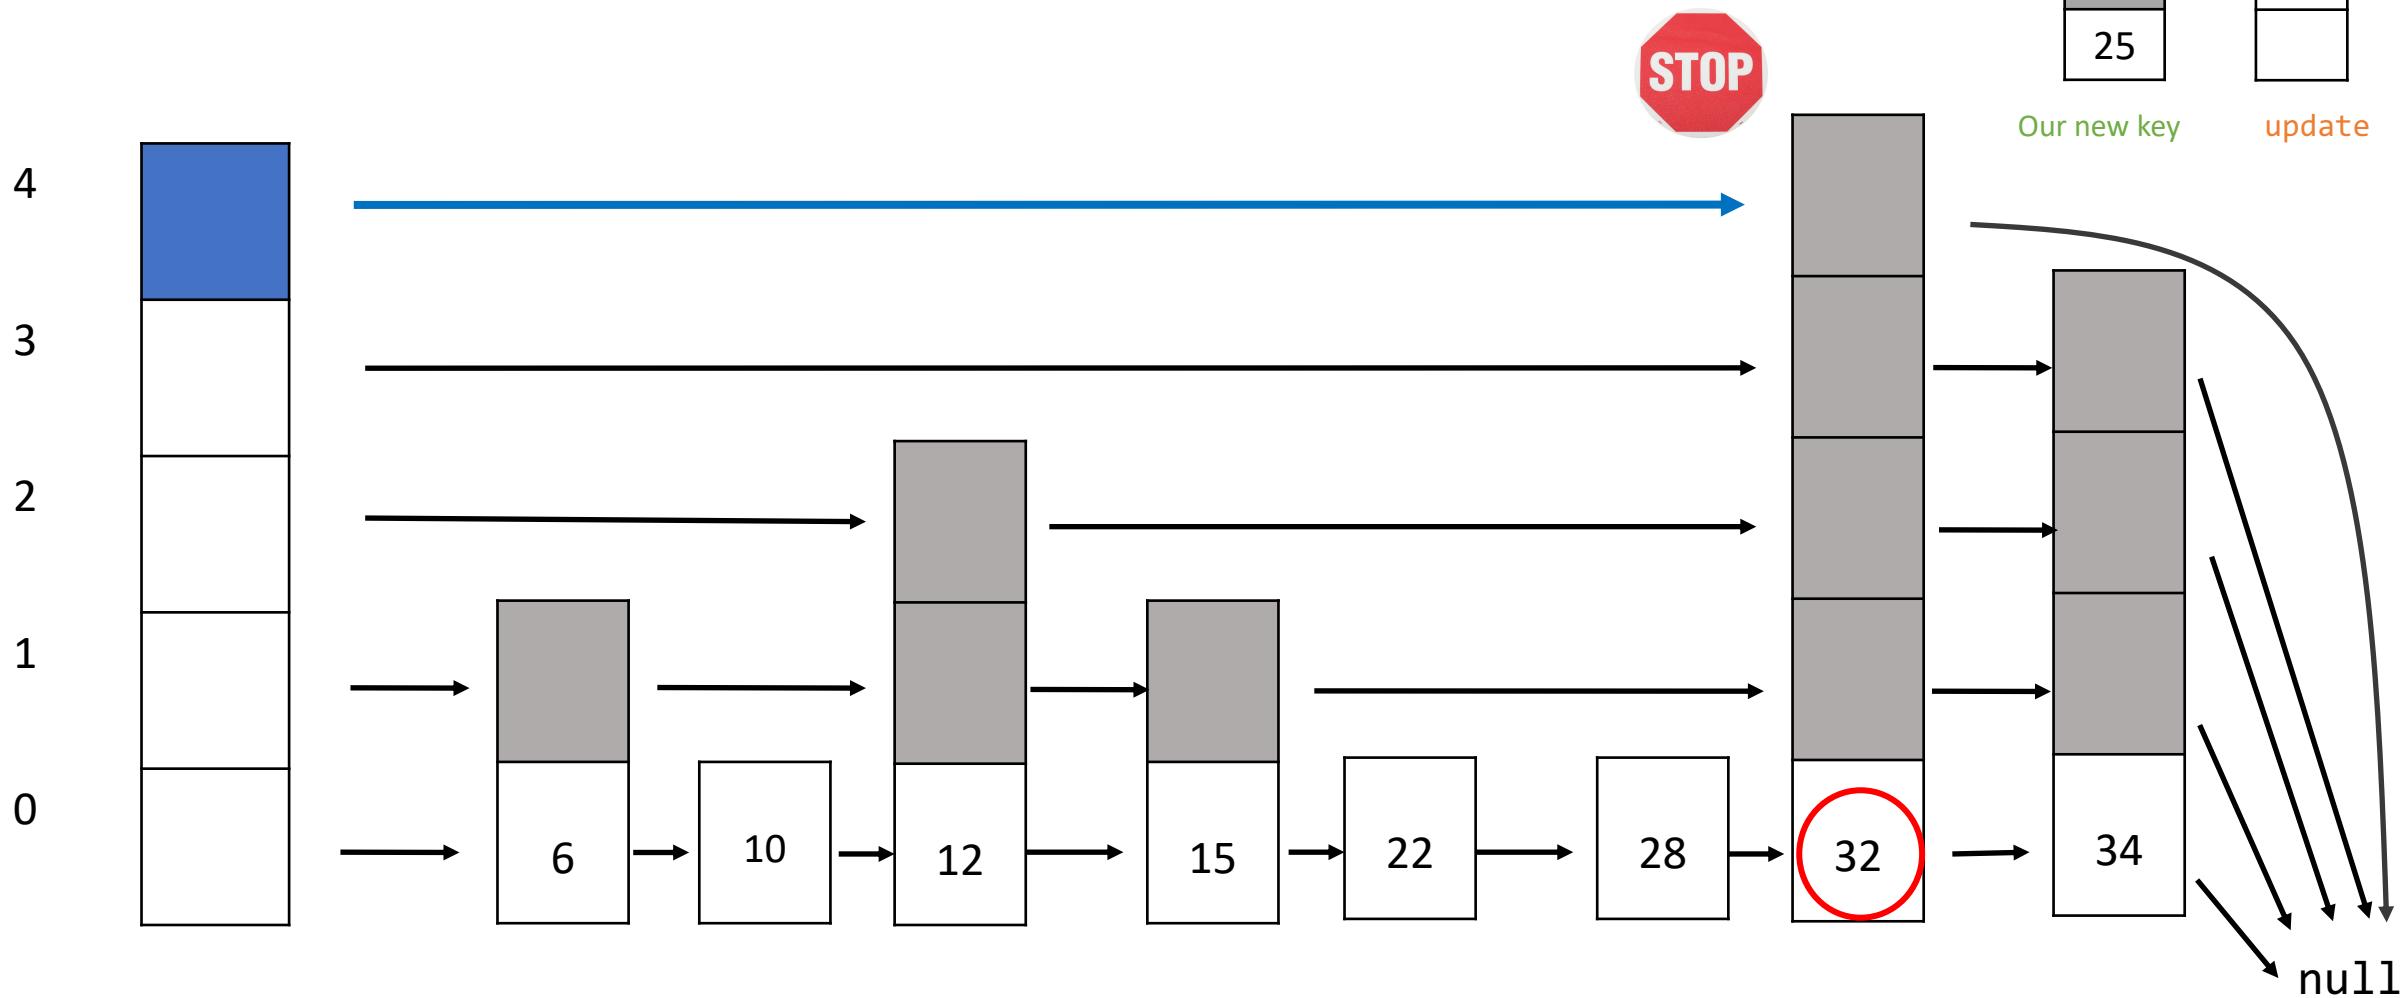
Task: Insert 25

# Insertion Example



Task: Insert 25

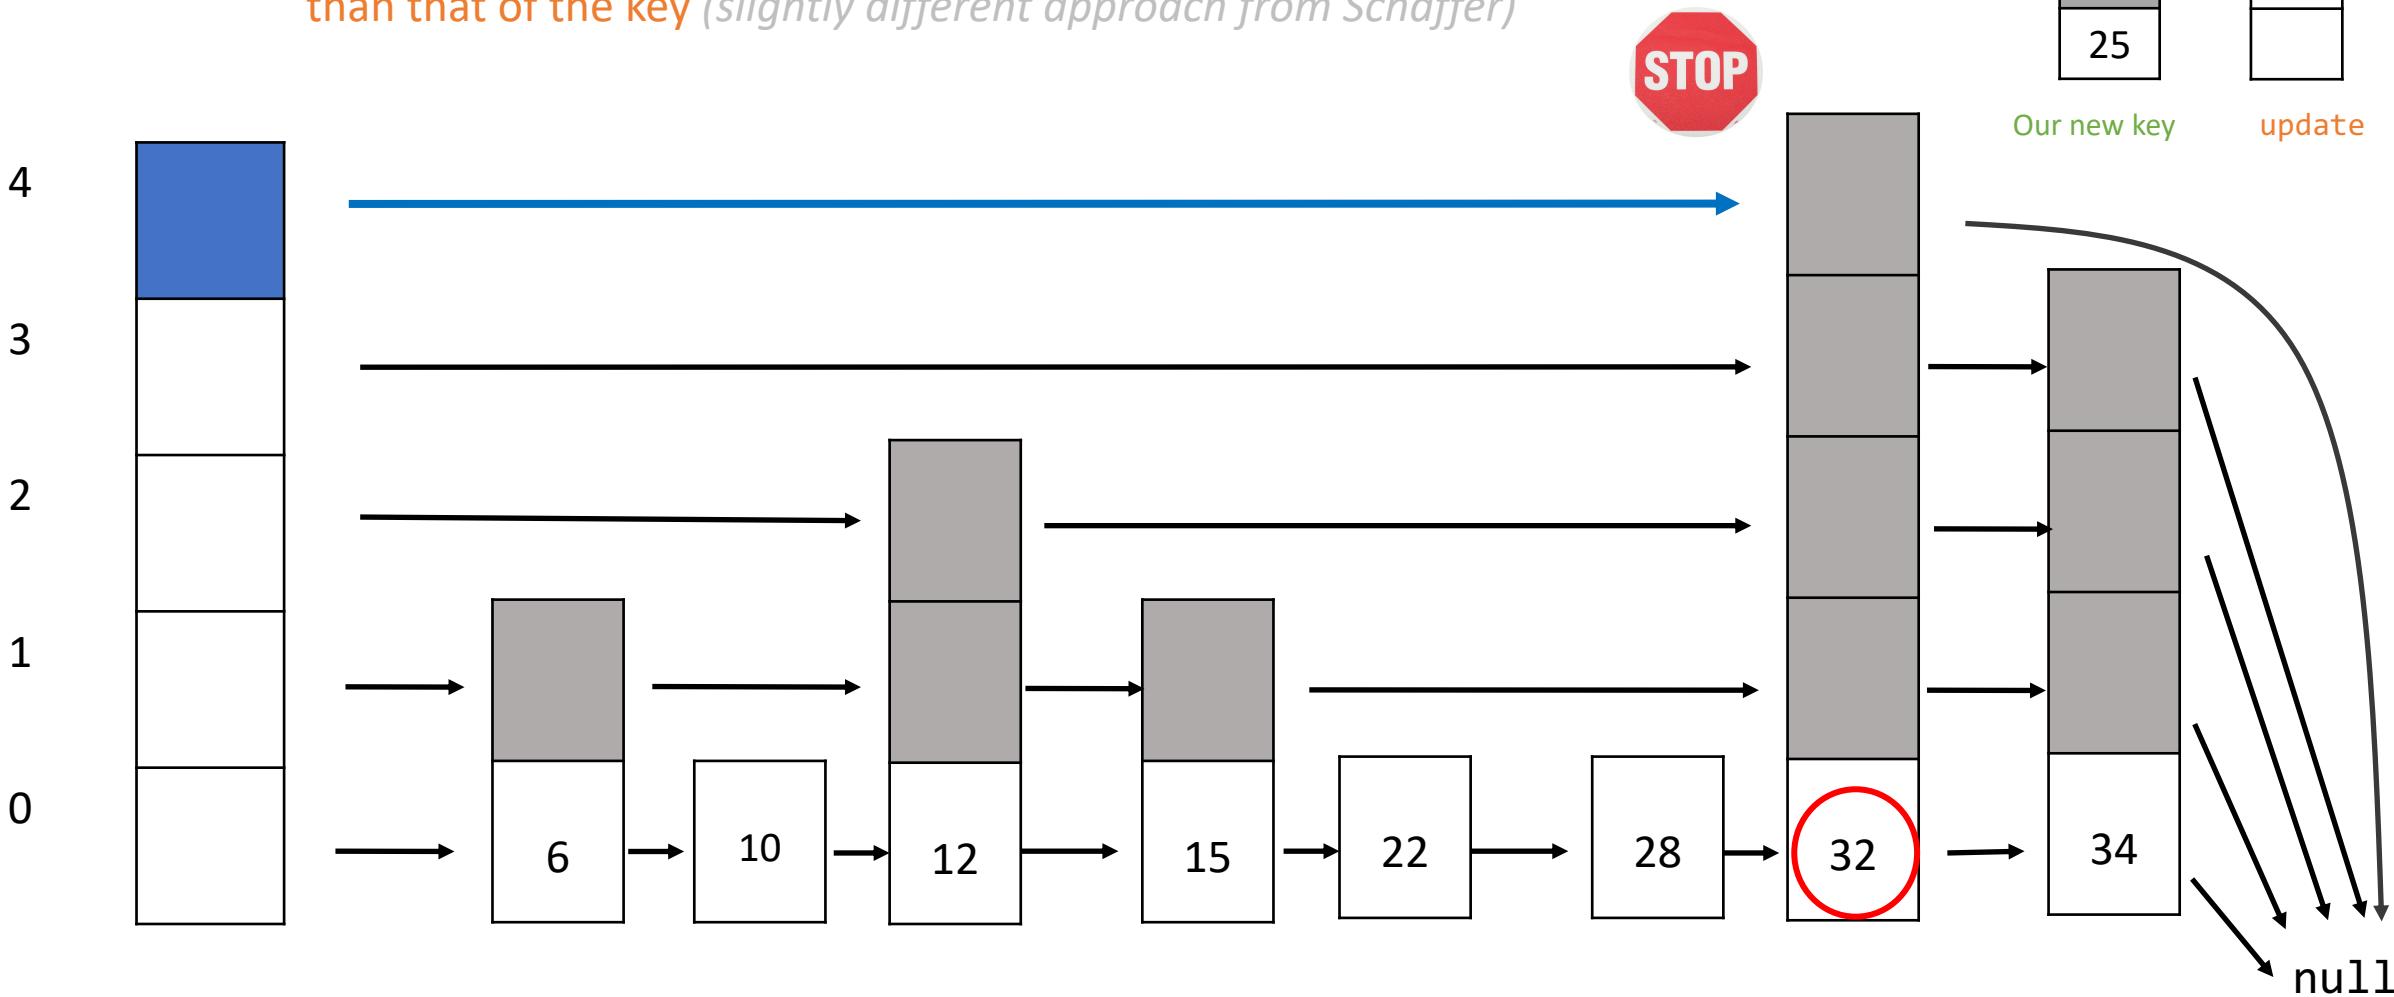
# Insertion Example



Task: Insert 25

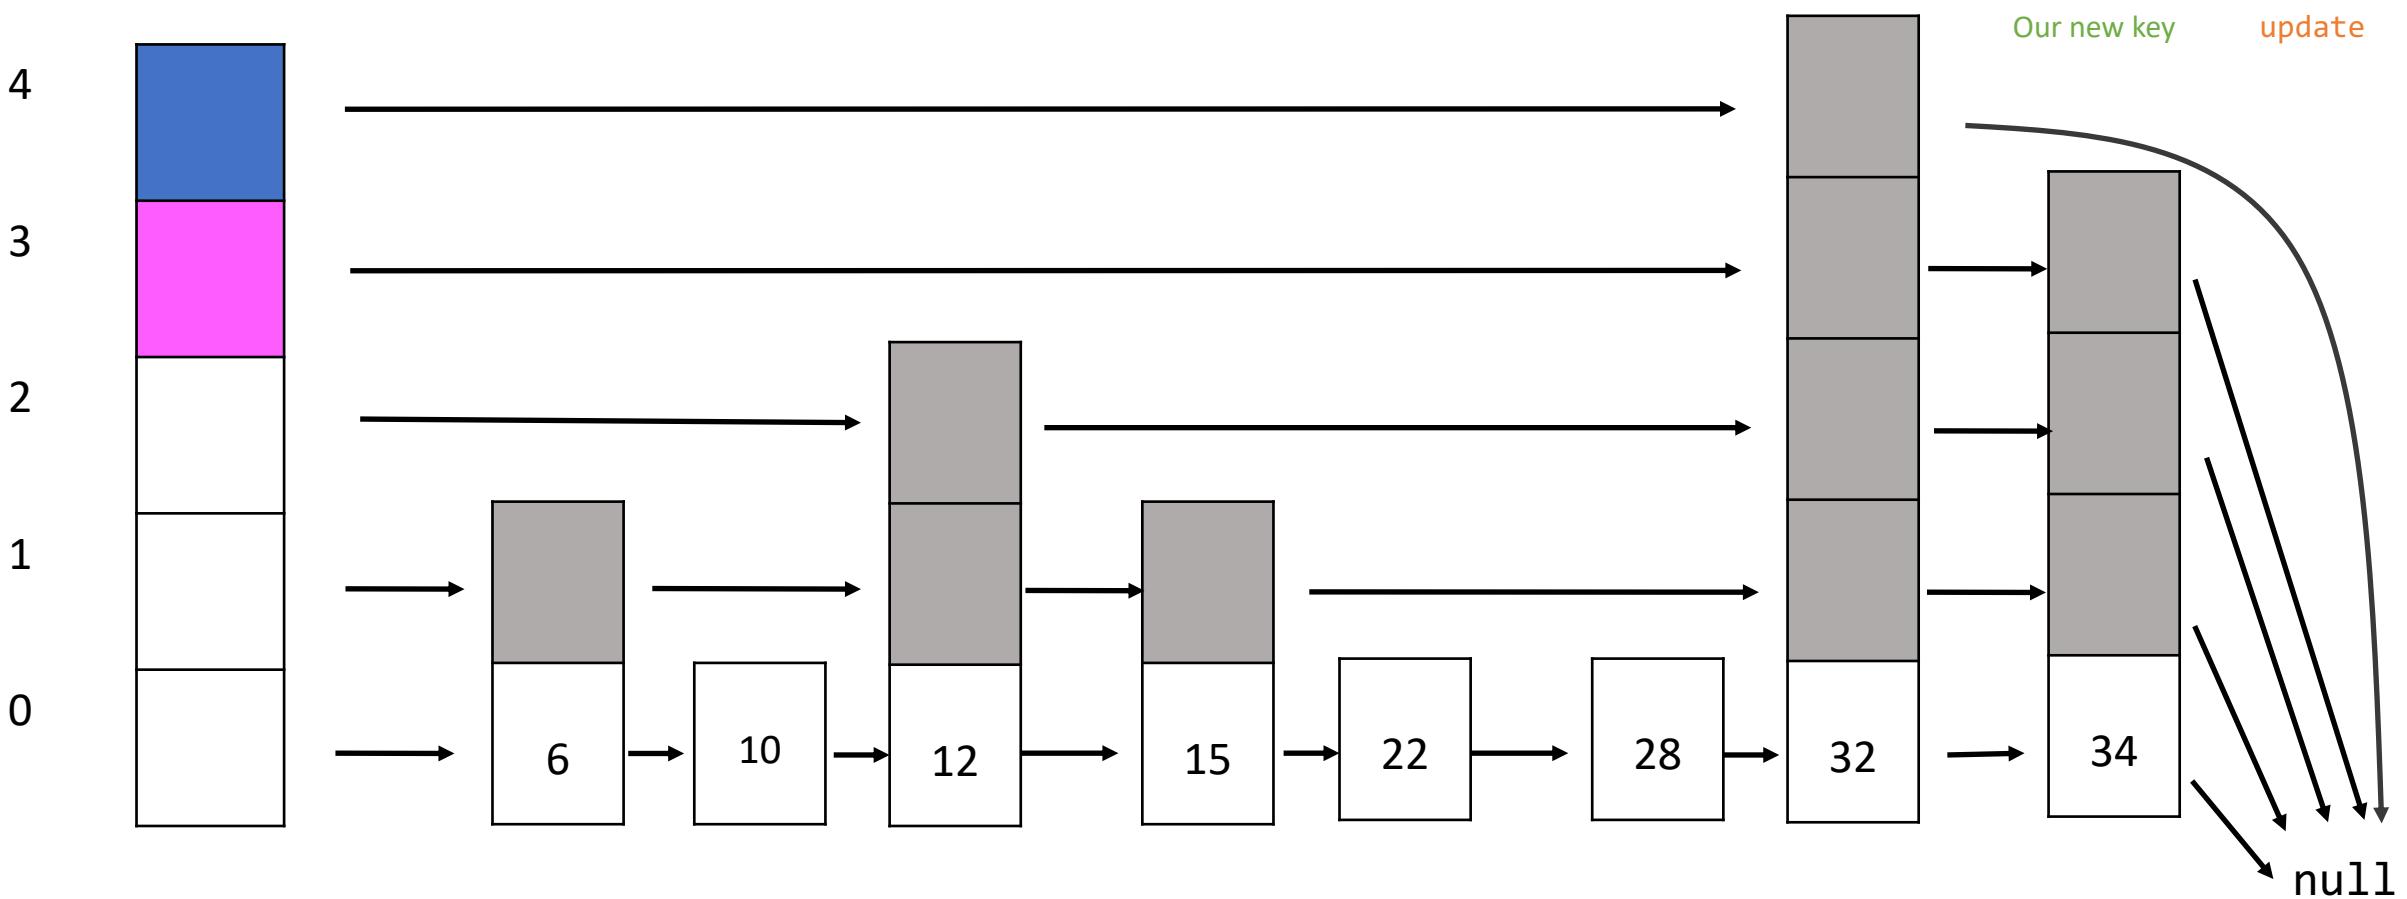
# Insertion Example

We don't have to touch update yet, since the current #levels is 4, greater than that of the key (*slightly different approach from Schaffer*)



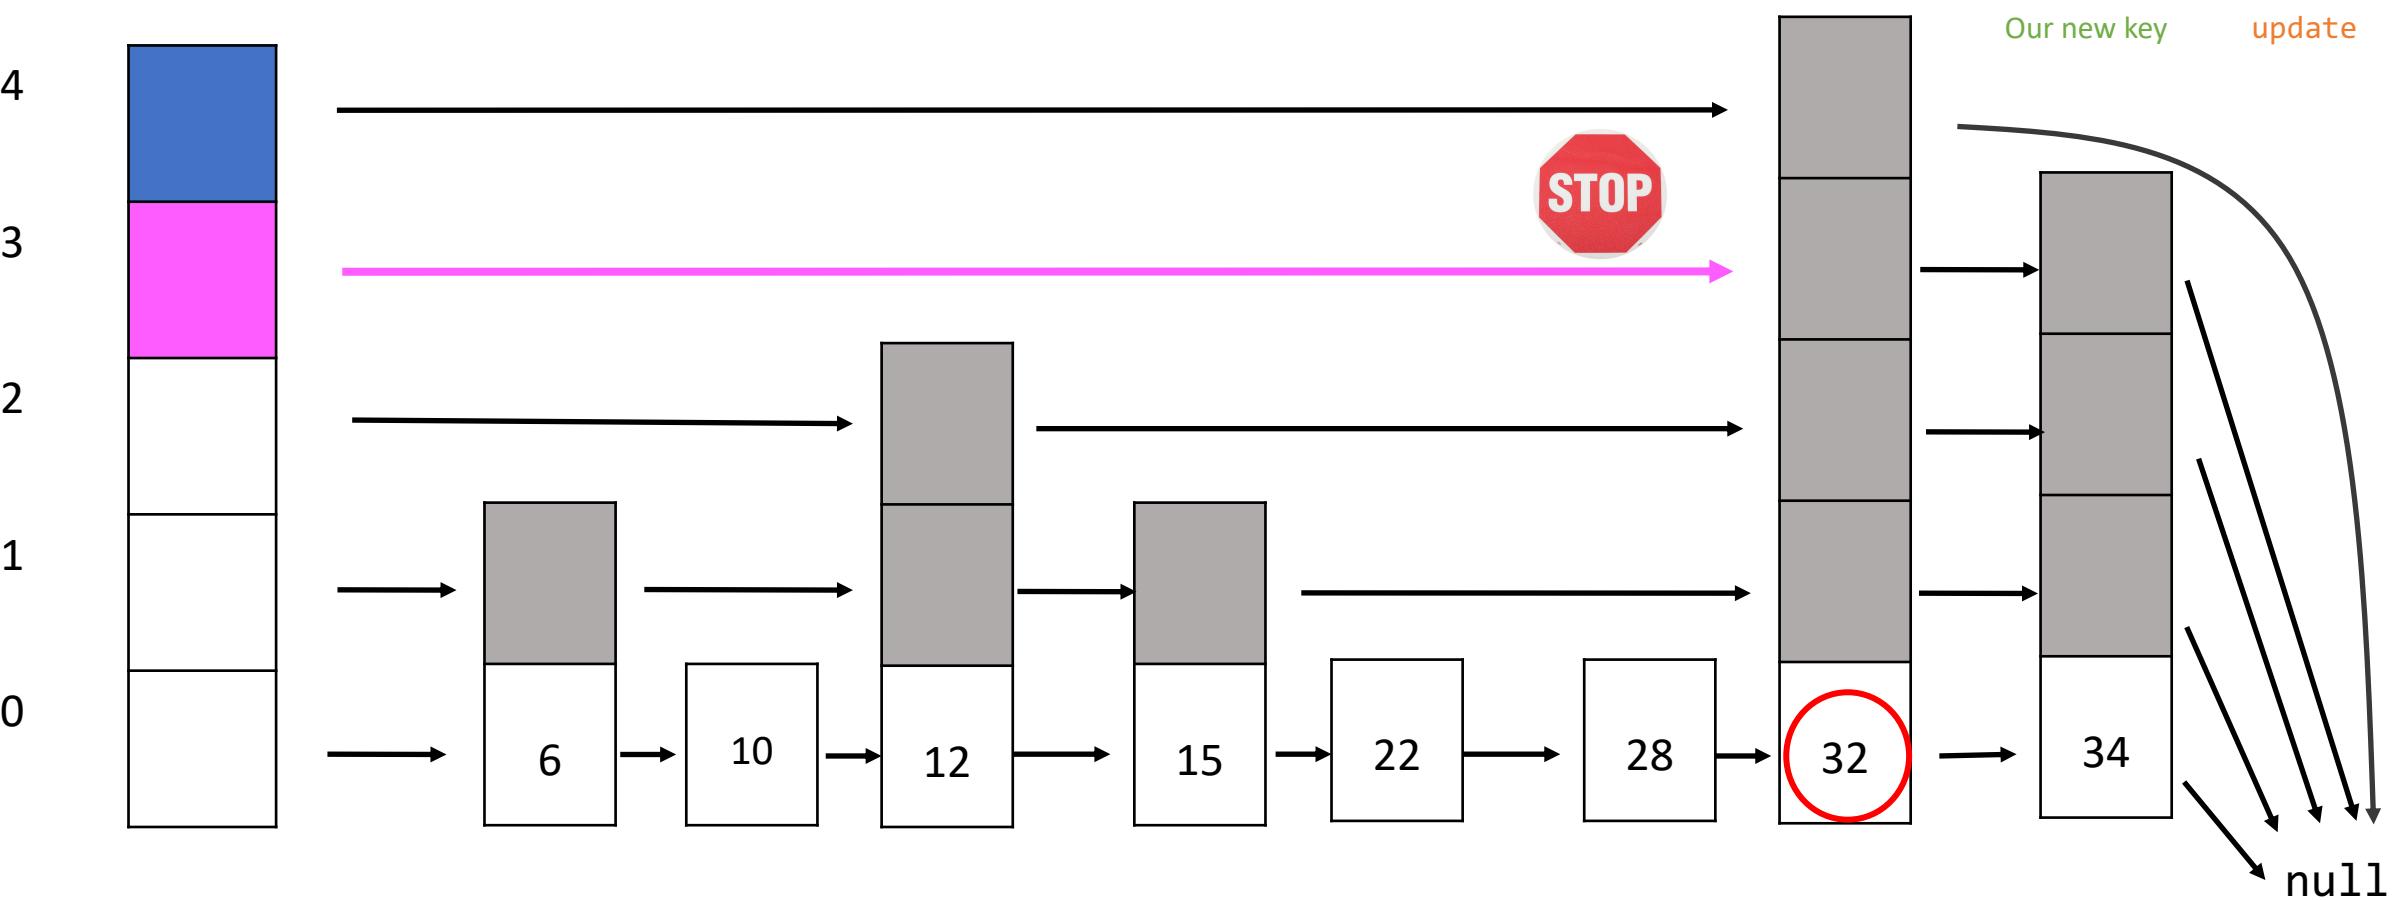
Task: Insert 25

# Insertion Example



Task: Insert 25

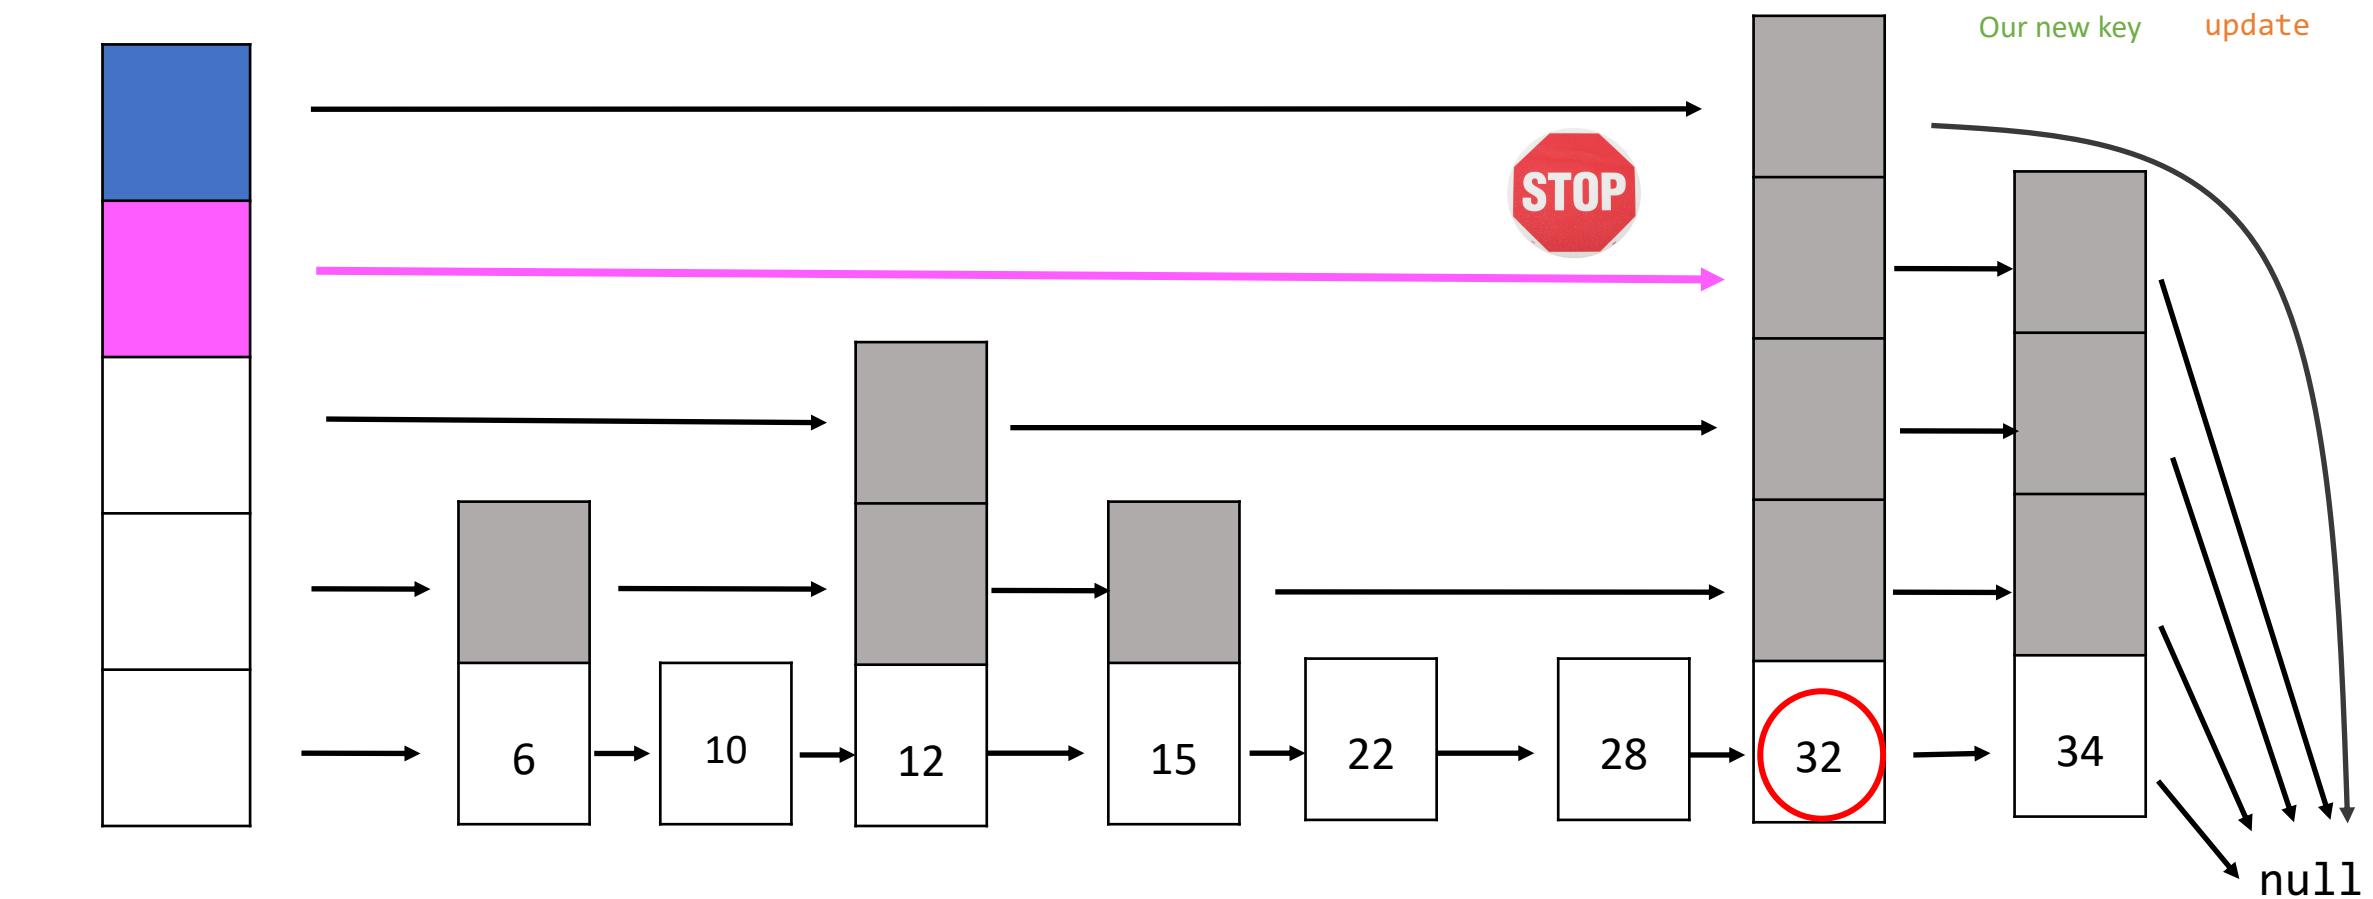
# Insertion Example



Task: Insert 25

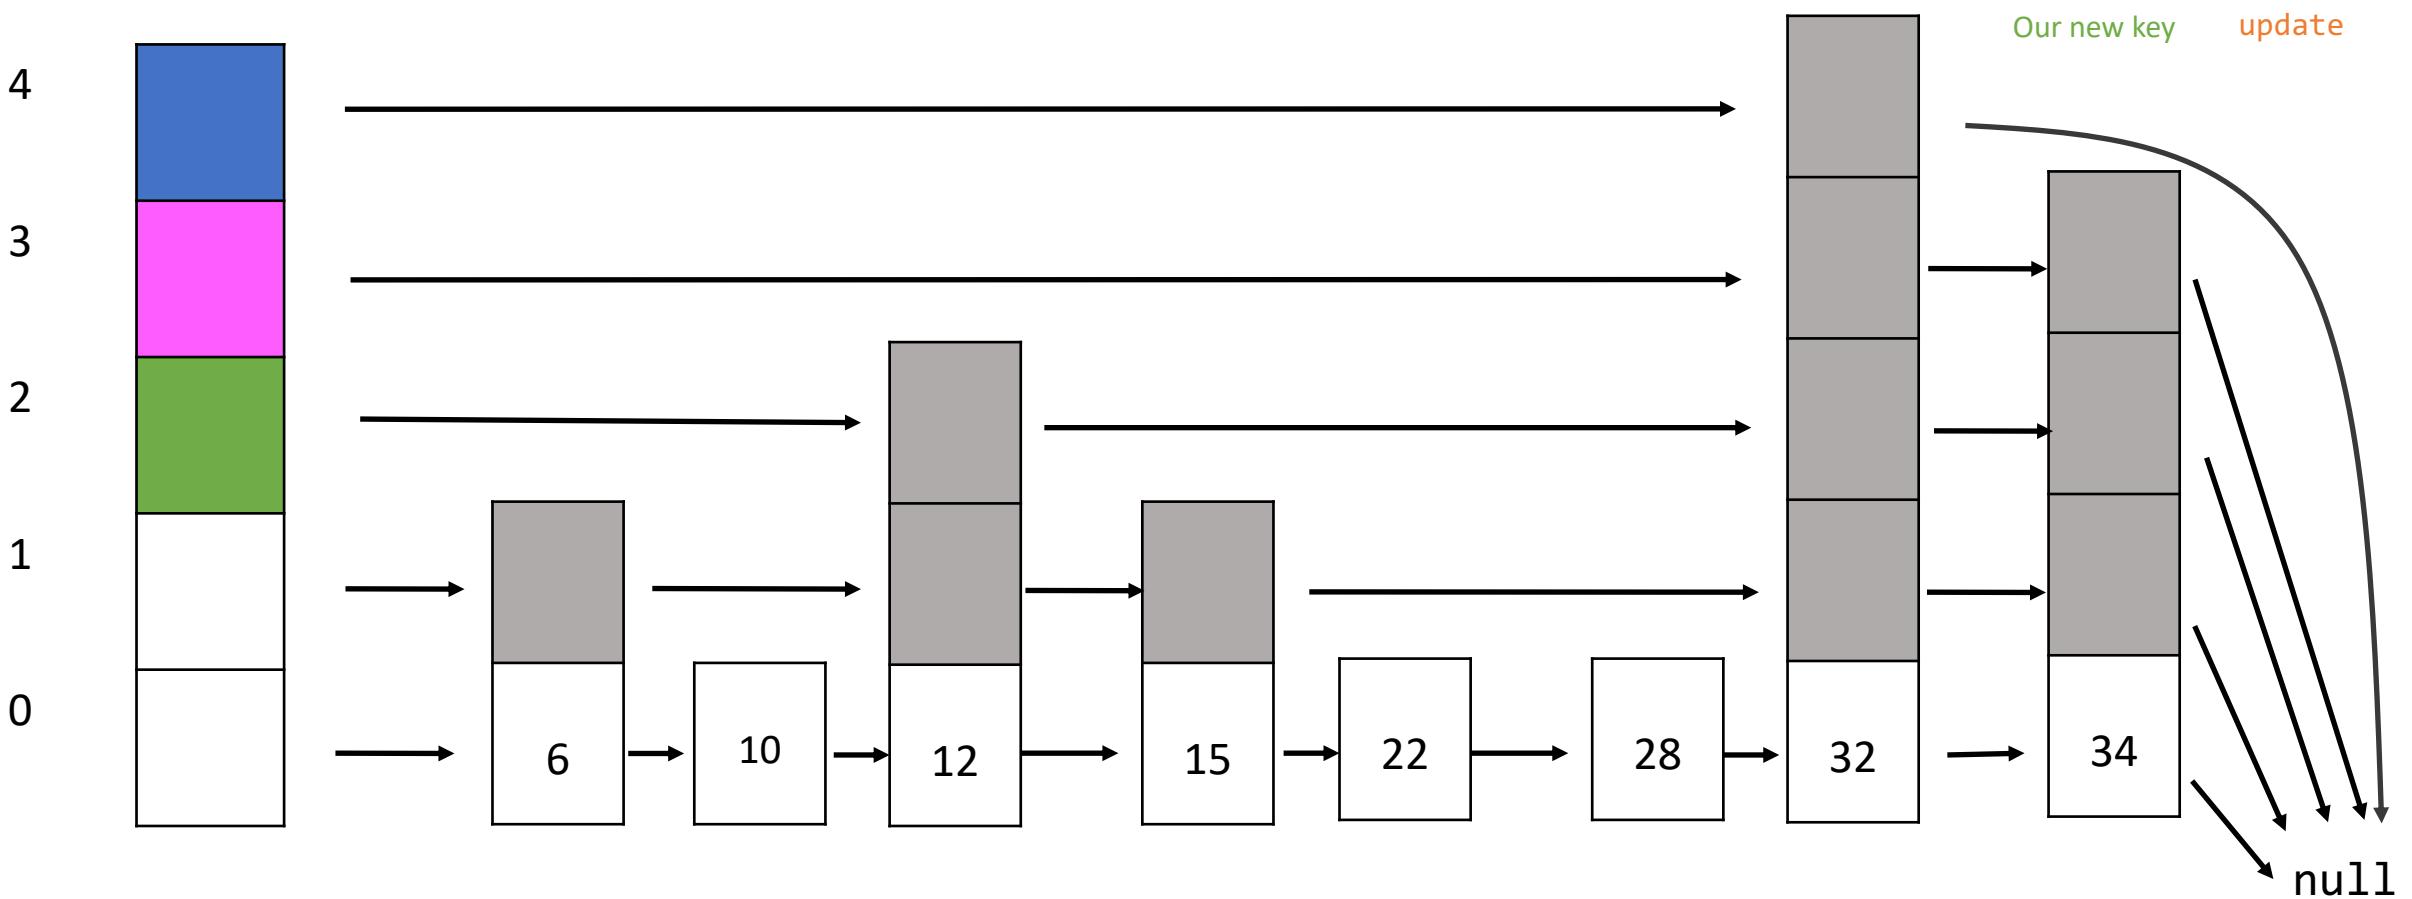
# Insertion Example

Now, however, we should touch the update array, since we are at a level (3) where pointers will need to be adjusted after insertion....



Task: Insert 25

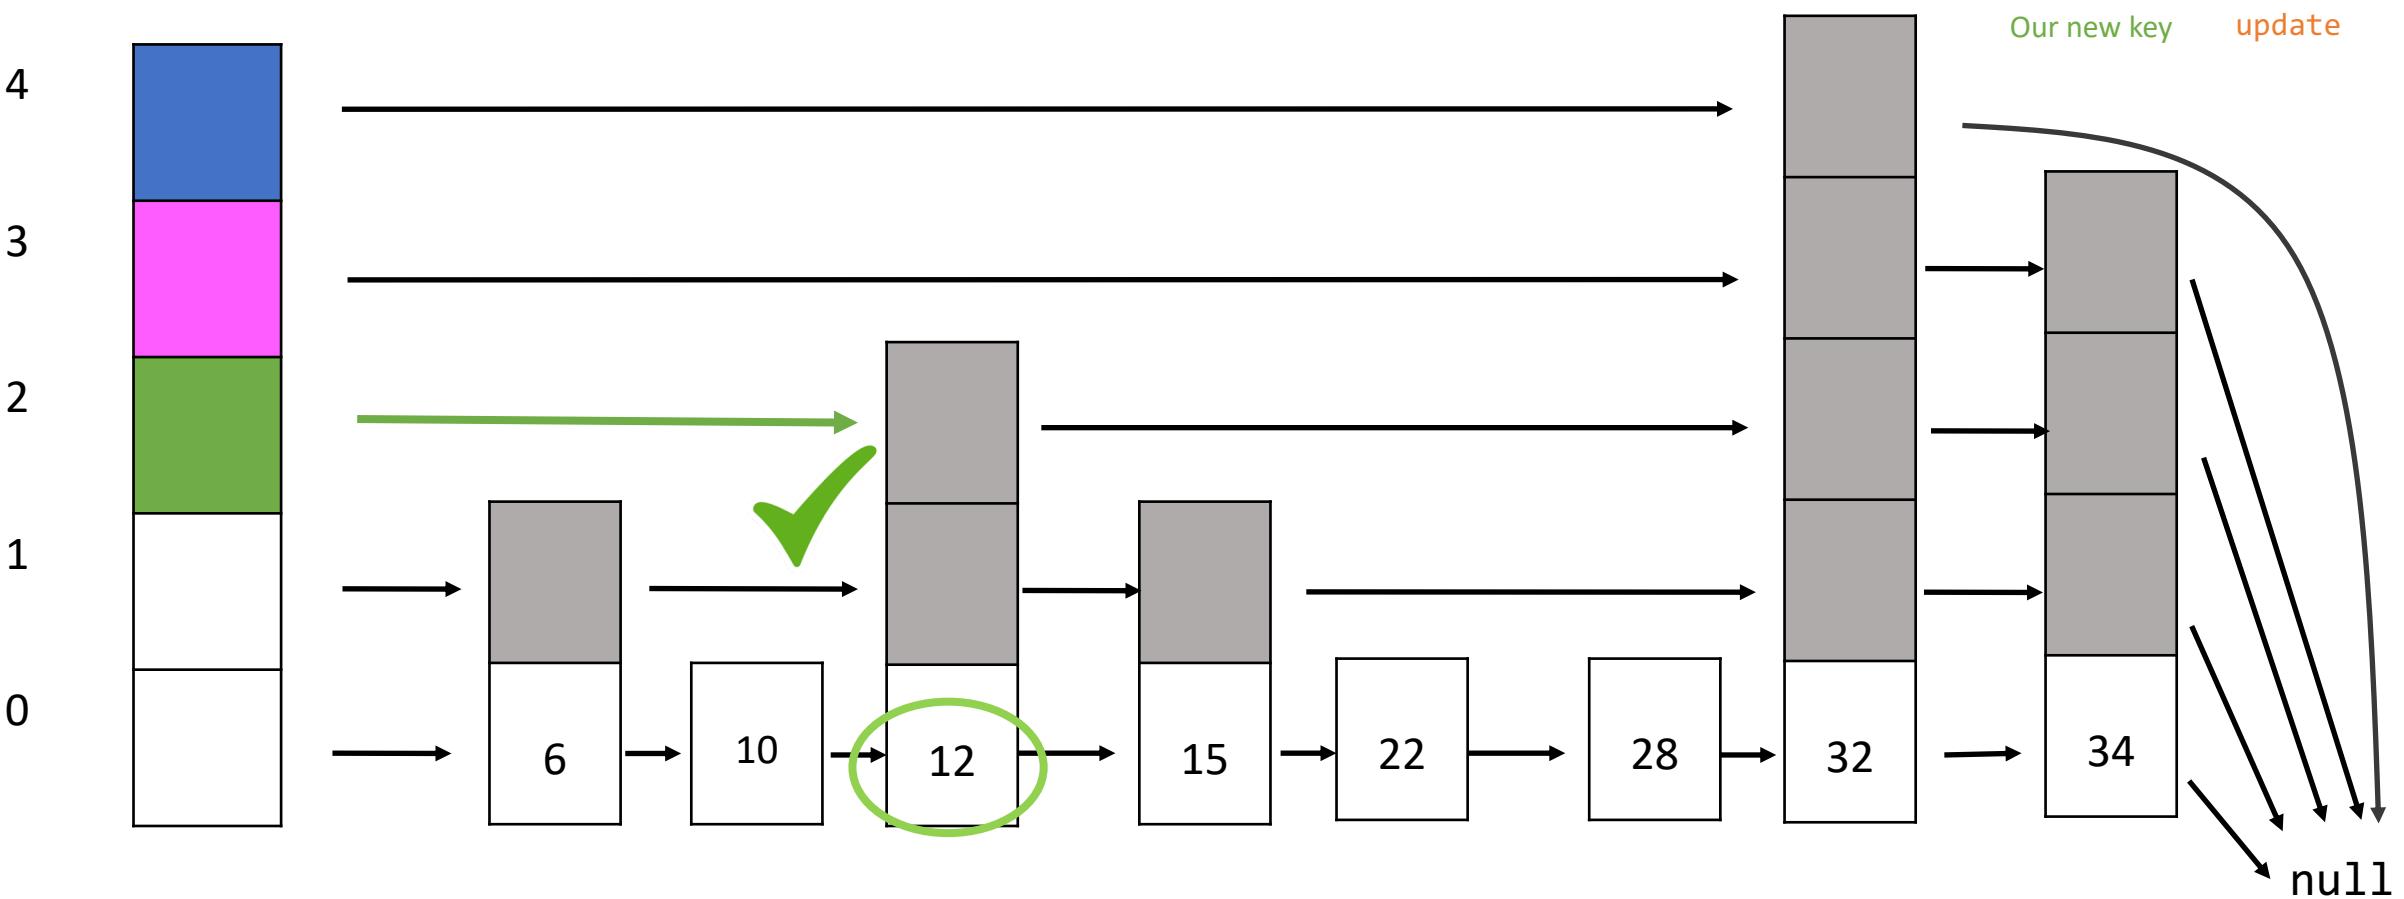
# Insertion Example



Task: Insert 25

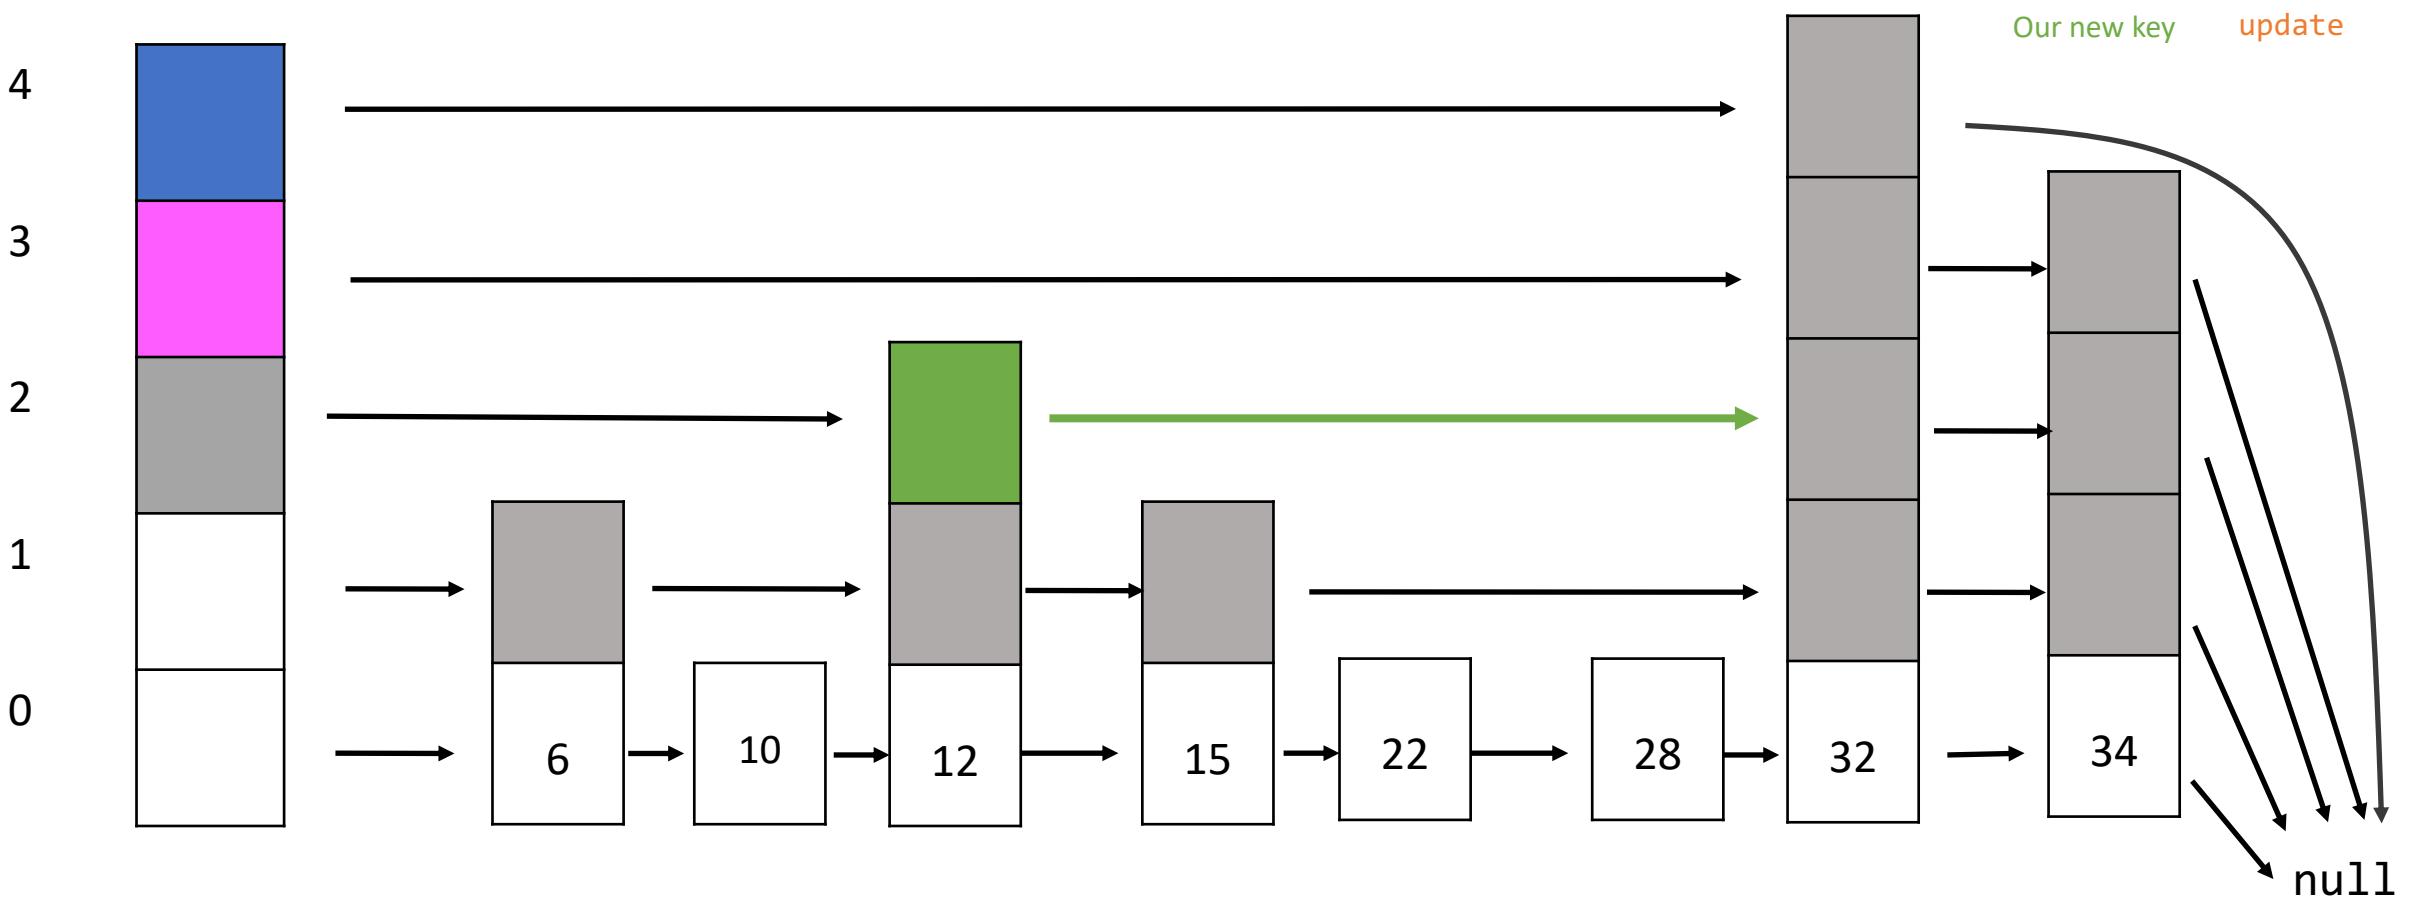
# Insertion Example

12 is smaller than 25, so we can keep going on this level: it's not time to touch "update" yet!



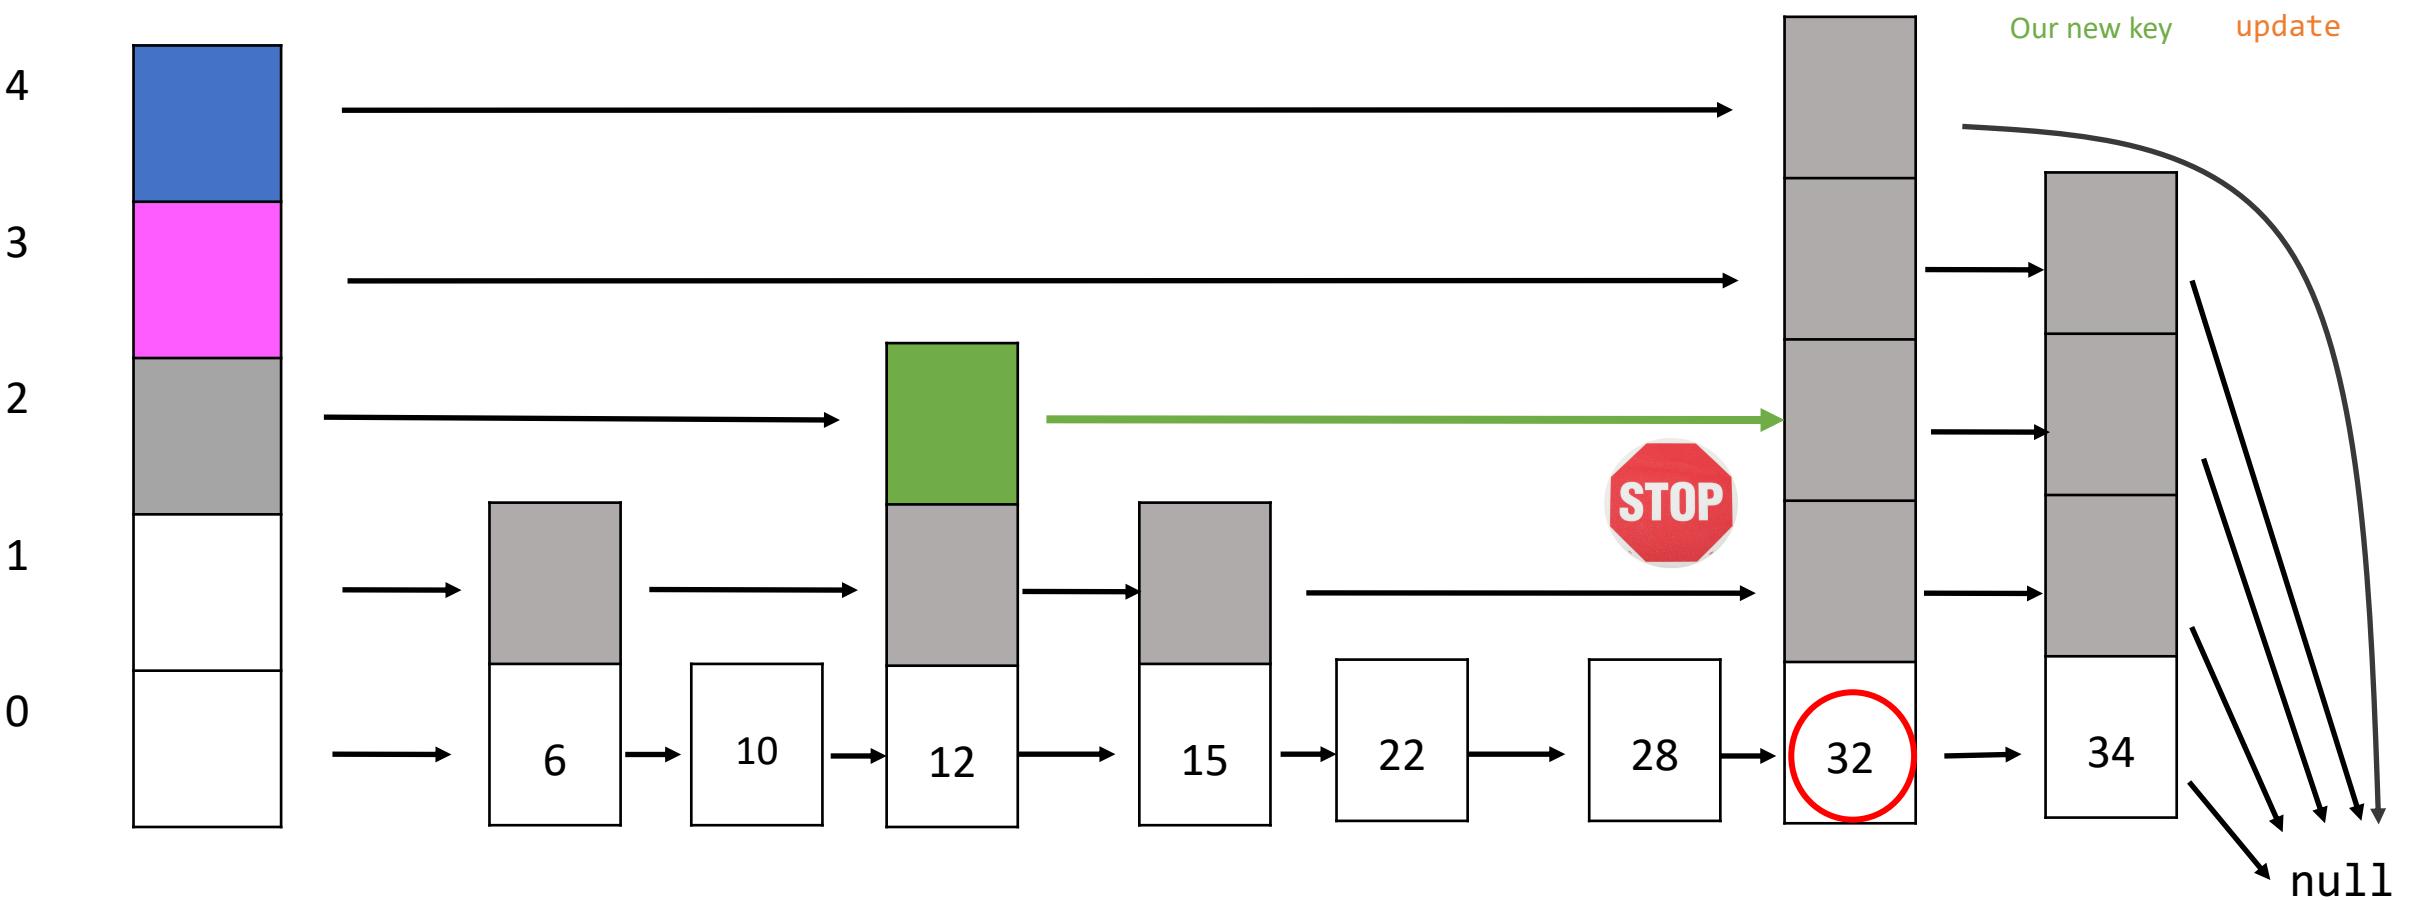
Task: Insert 25

# Insertion Example



Task: Insert 25

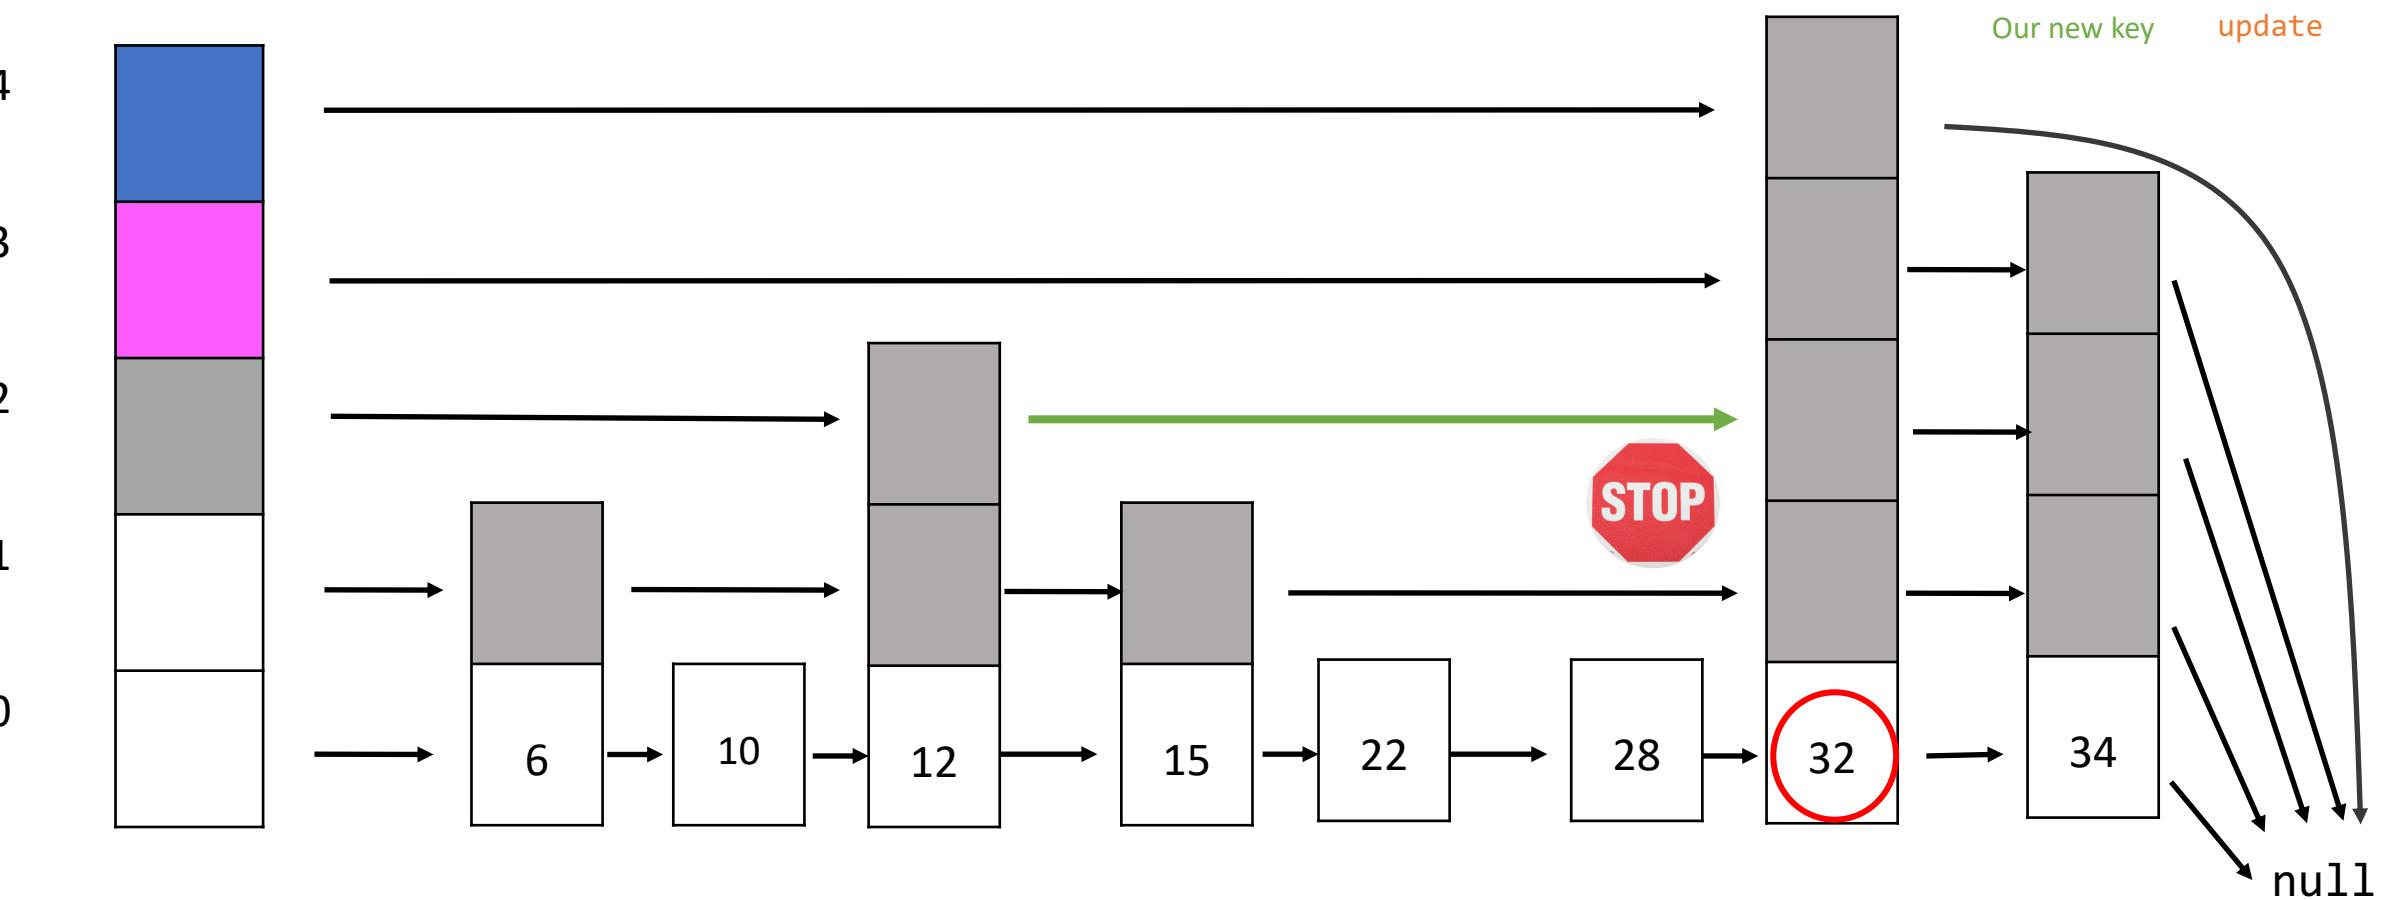
# Insertion Example



Task: Insert 25

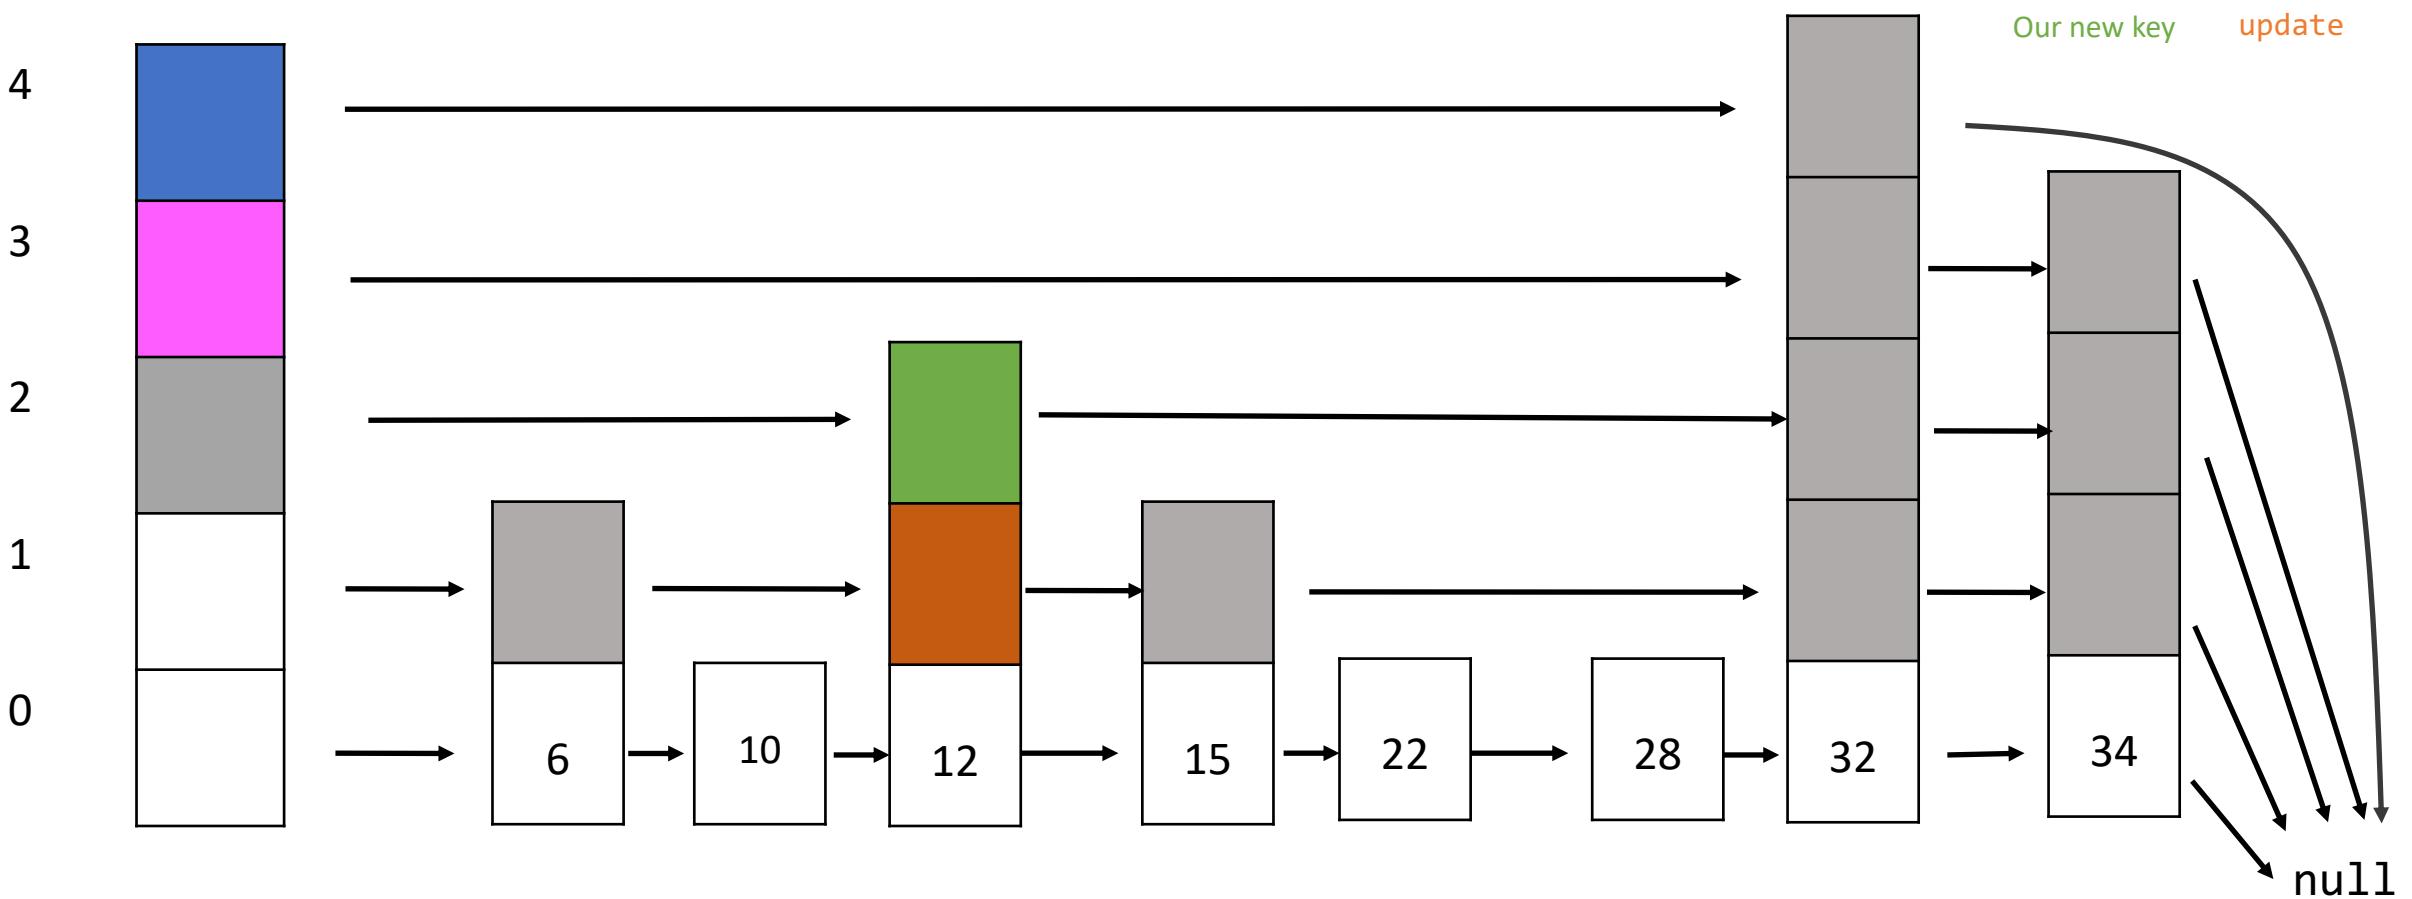
# Insertion Example

Now it's time to write into update again....



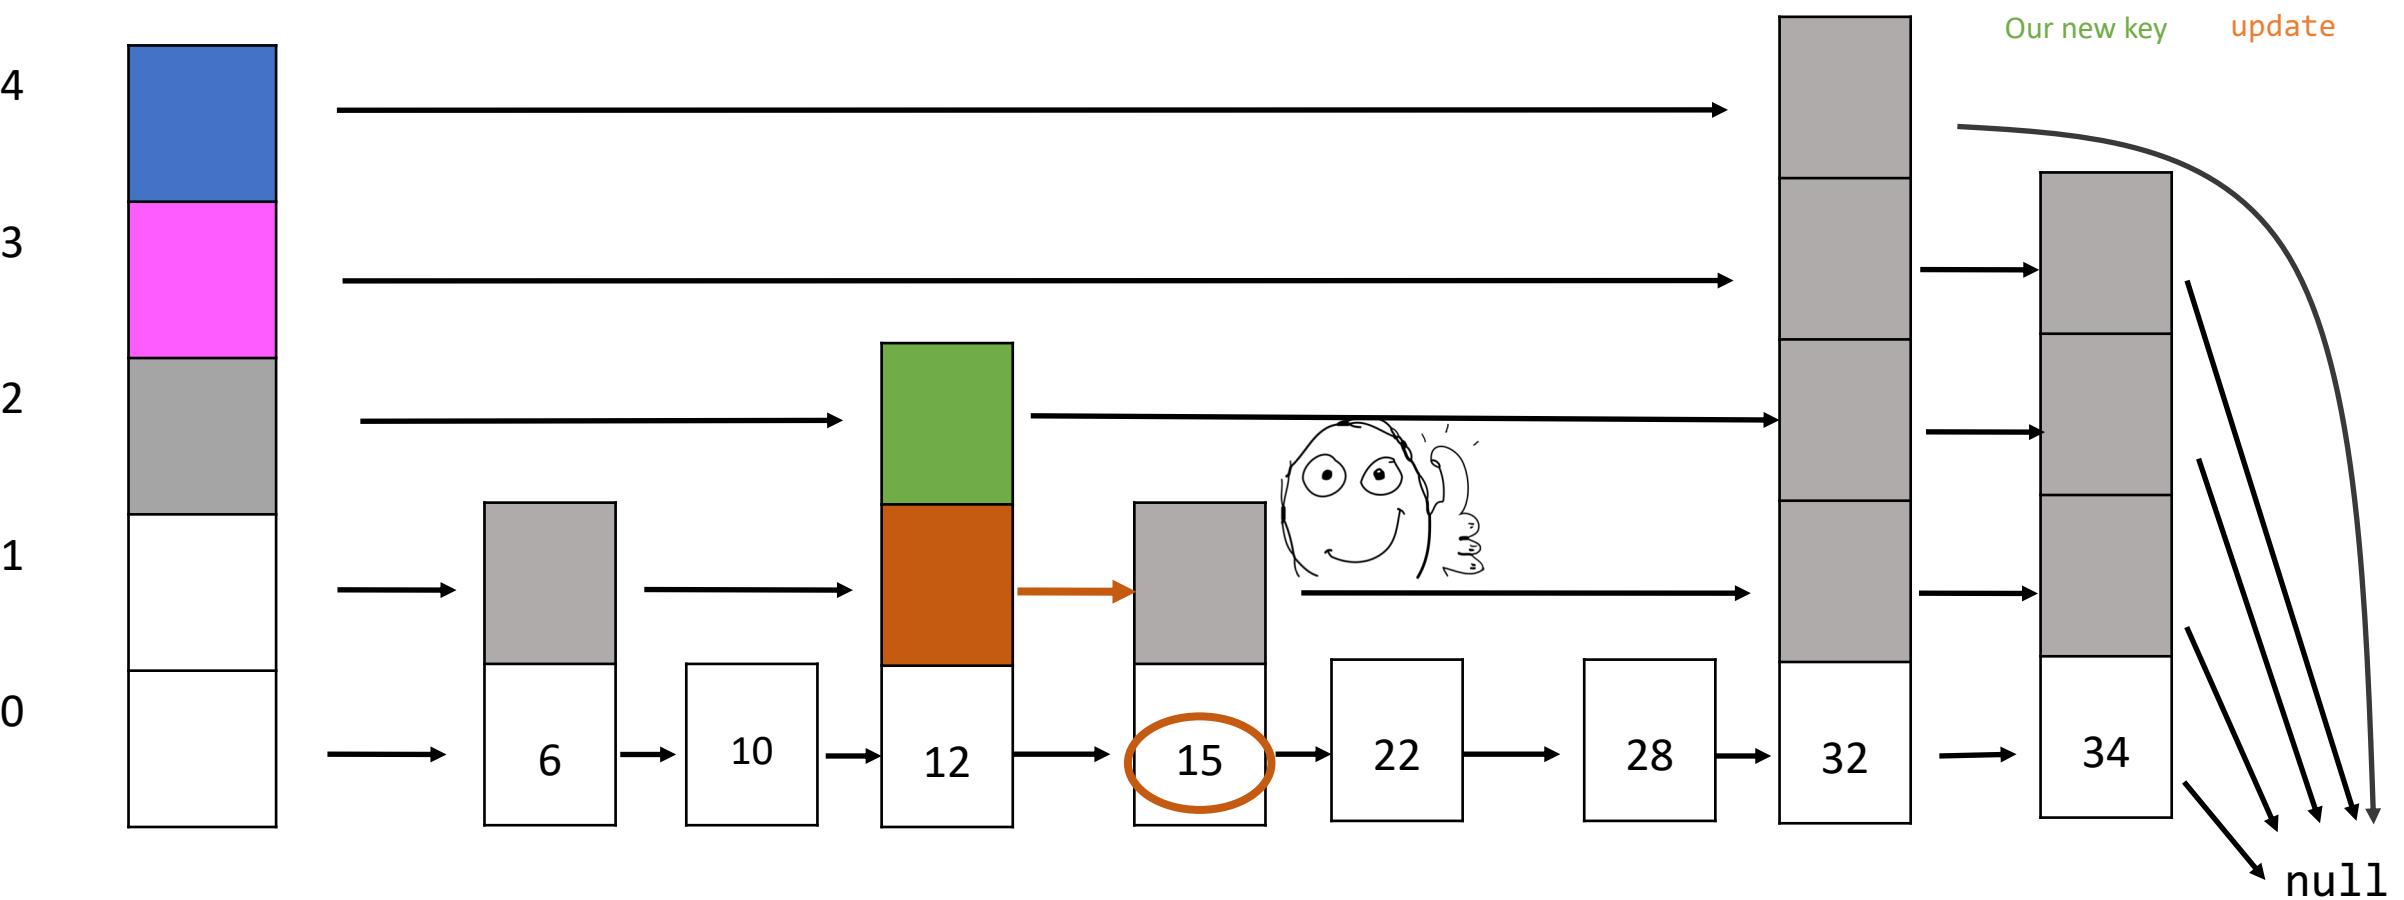
Task: Insert 25

# Insertion Example



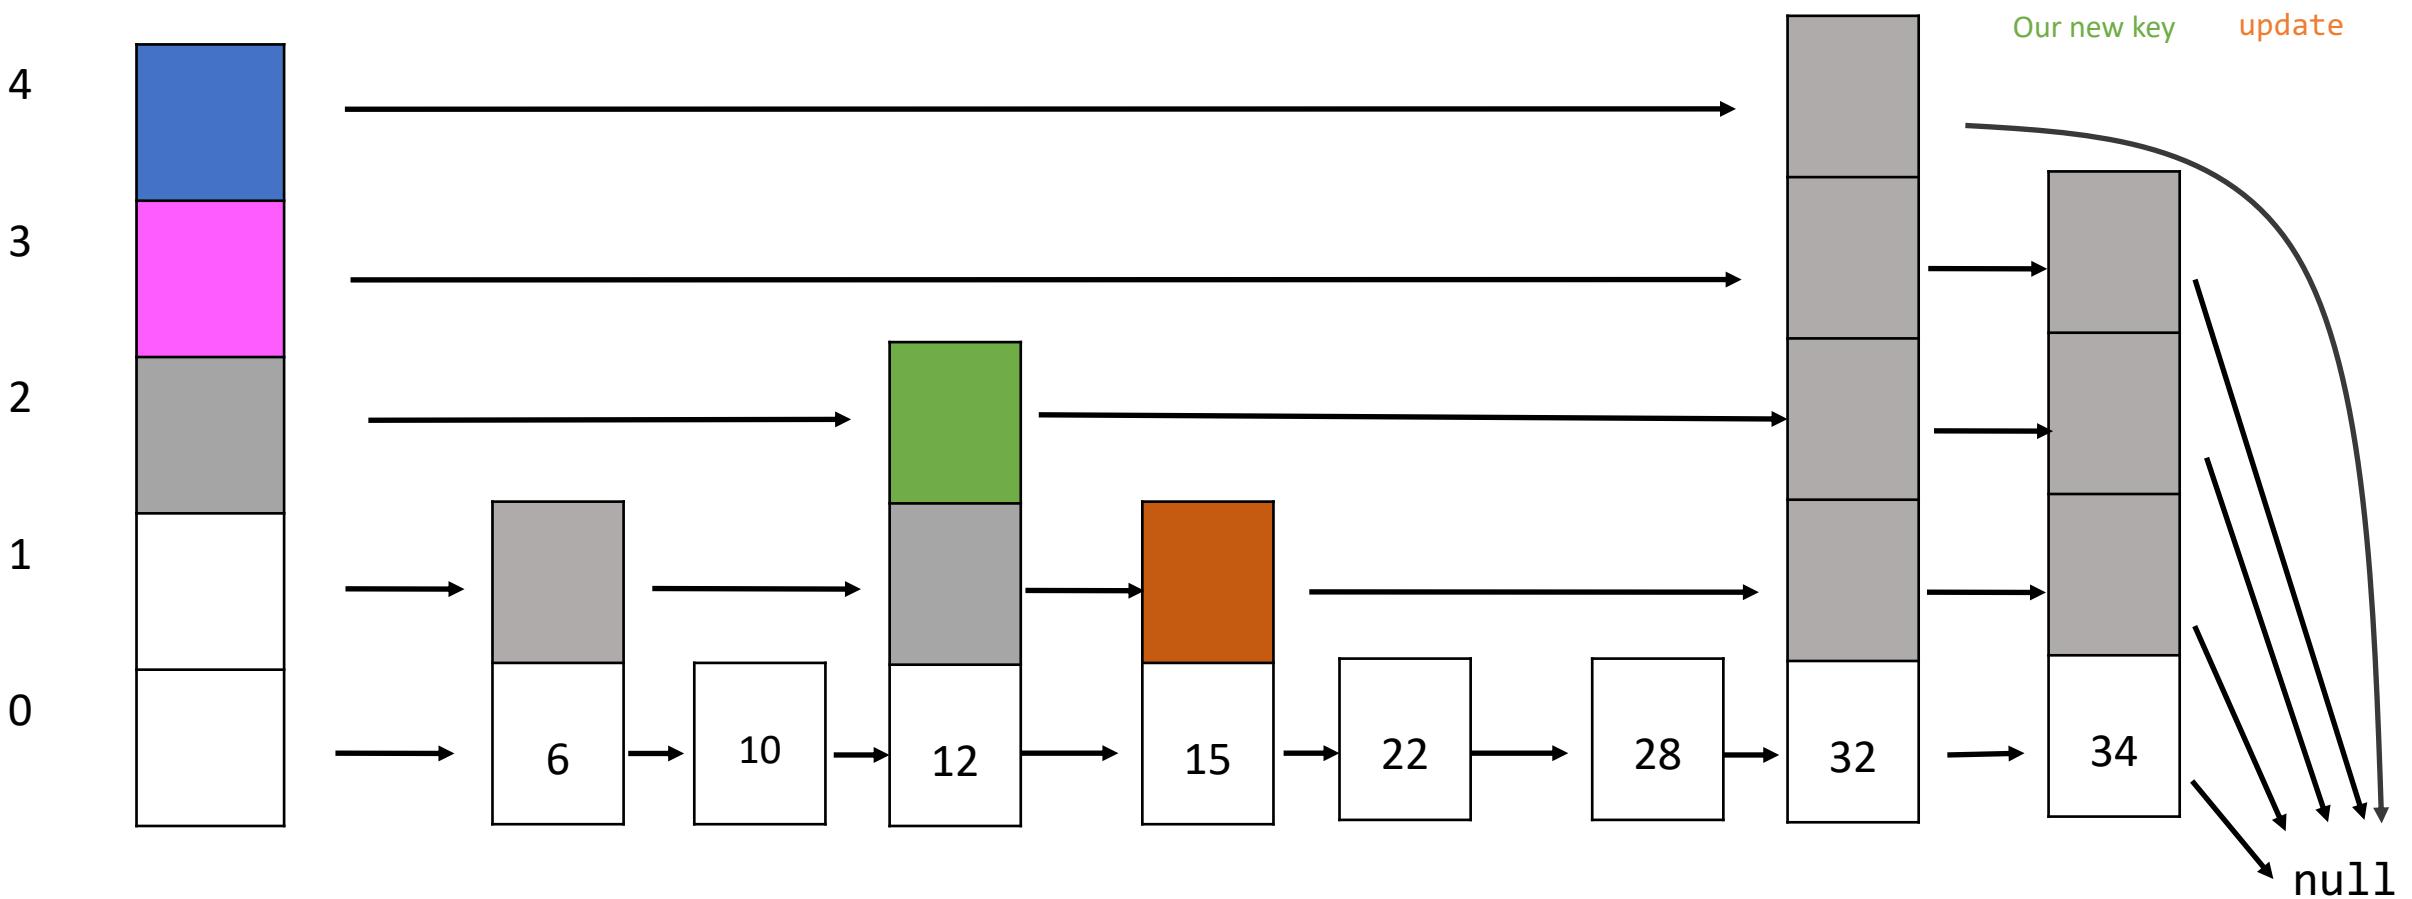
Task: Insert 25

# Insertion Example



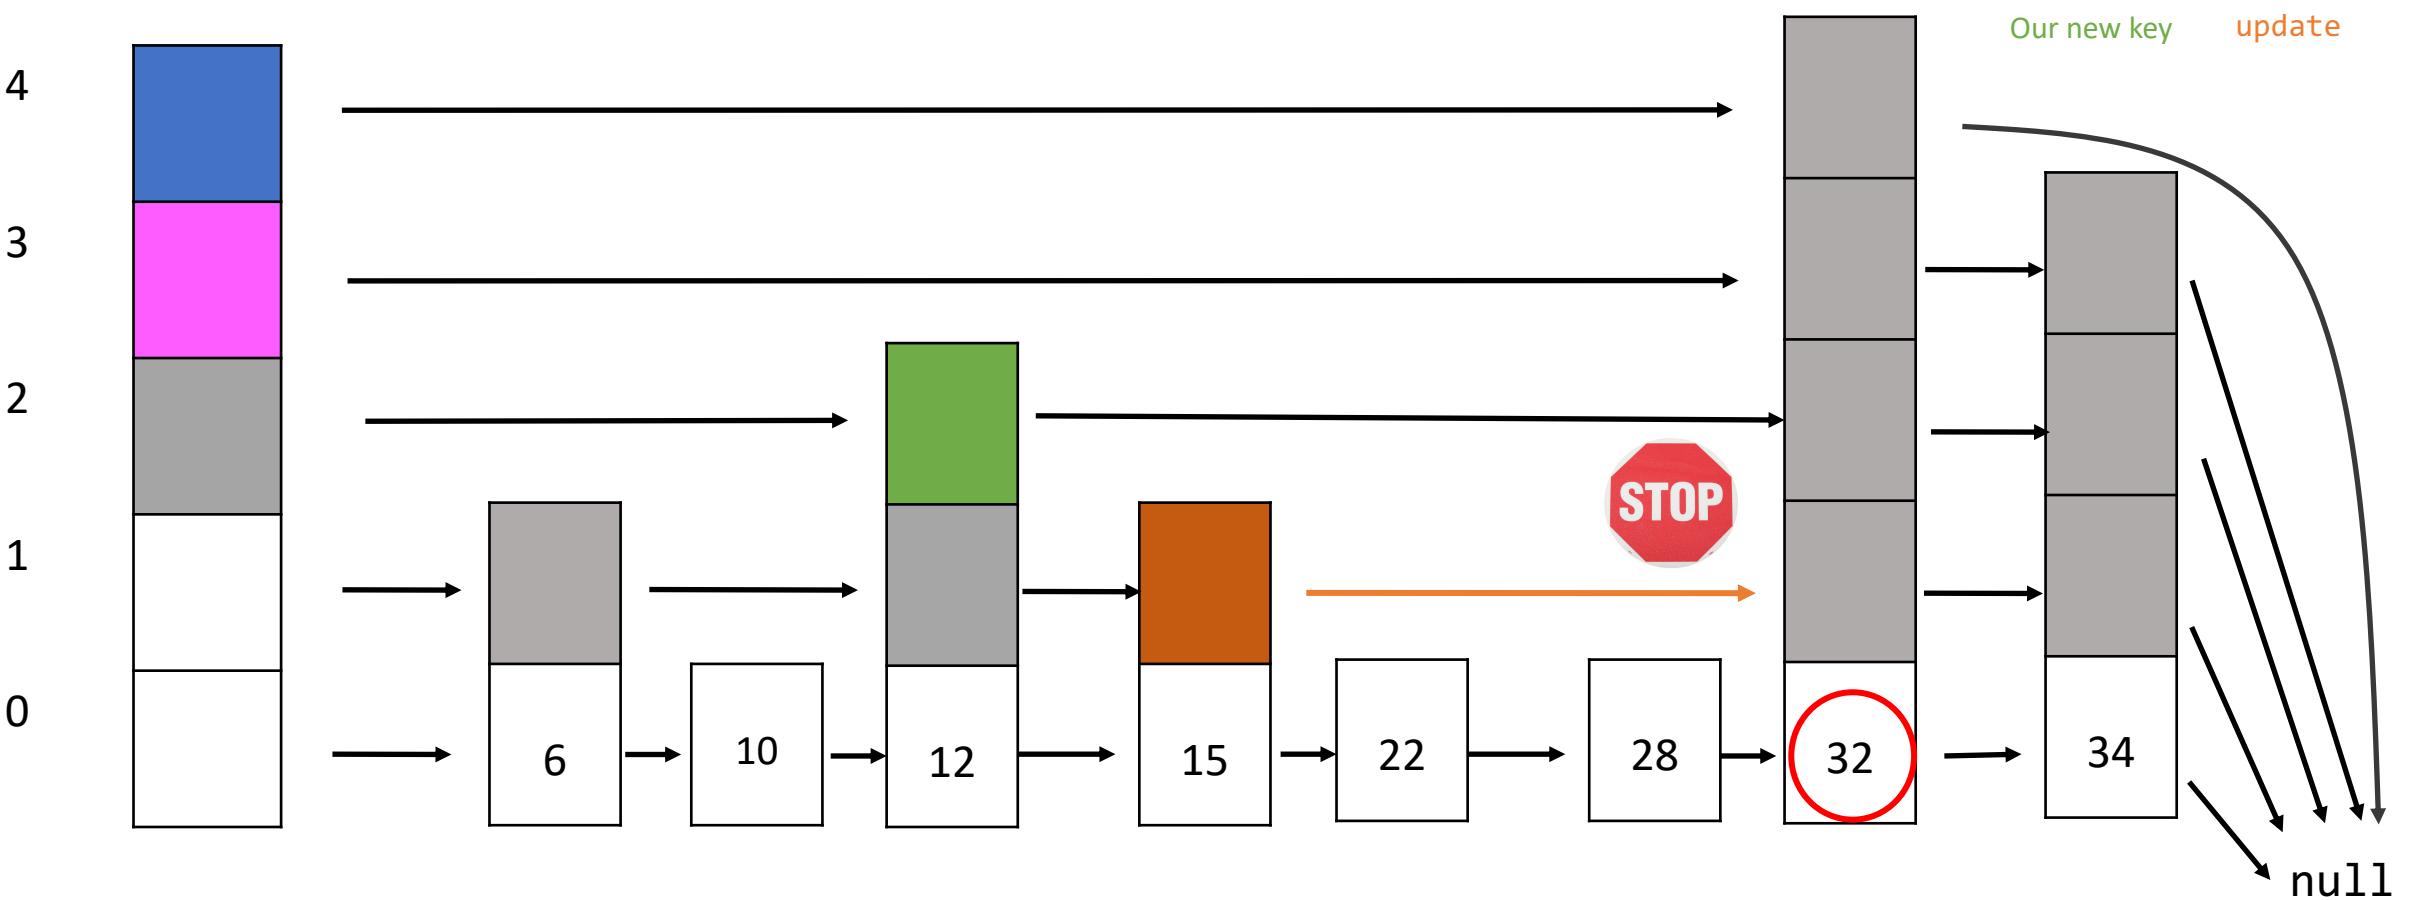
Task: Insert 25

# Insertion Example



Task: Insert 25

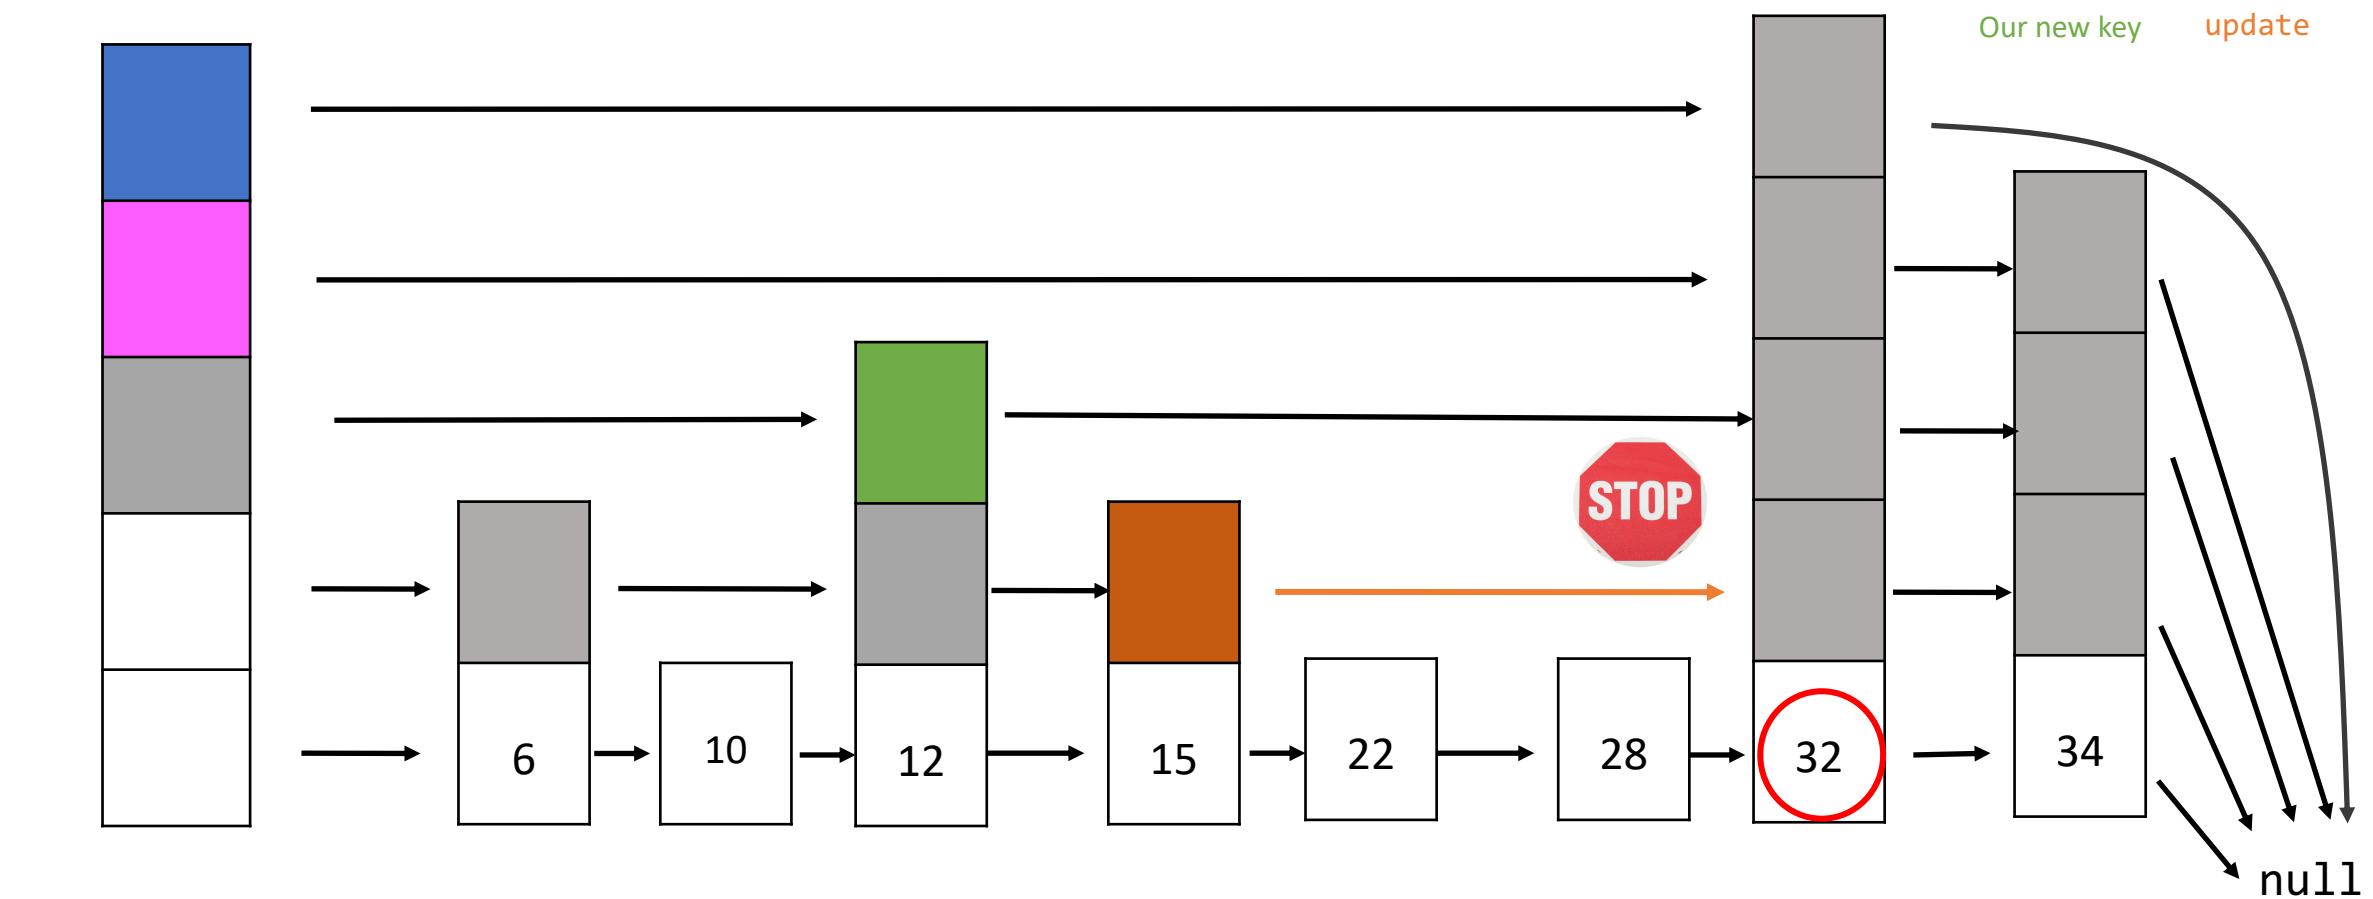
# Insertion Example



Task: Insert 25

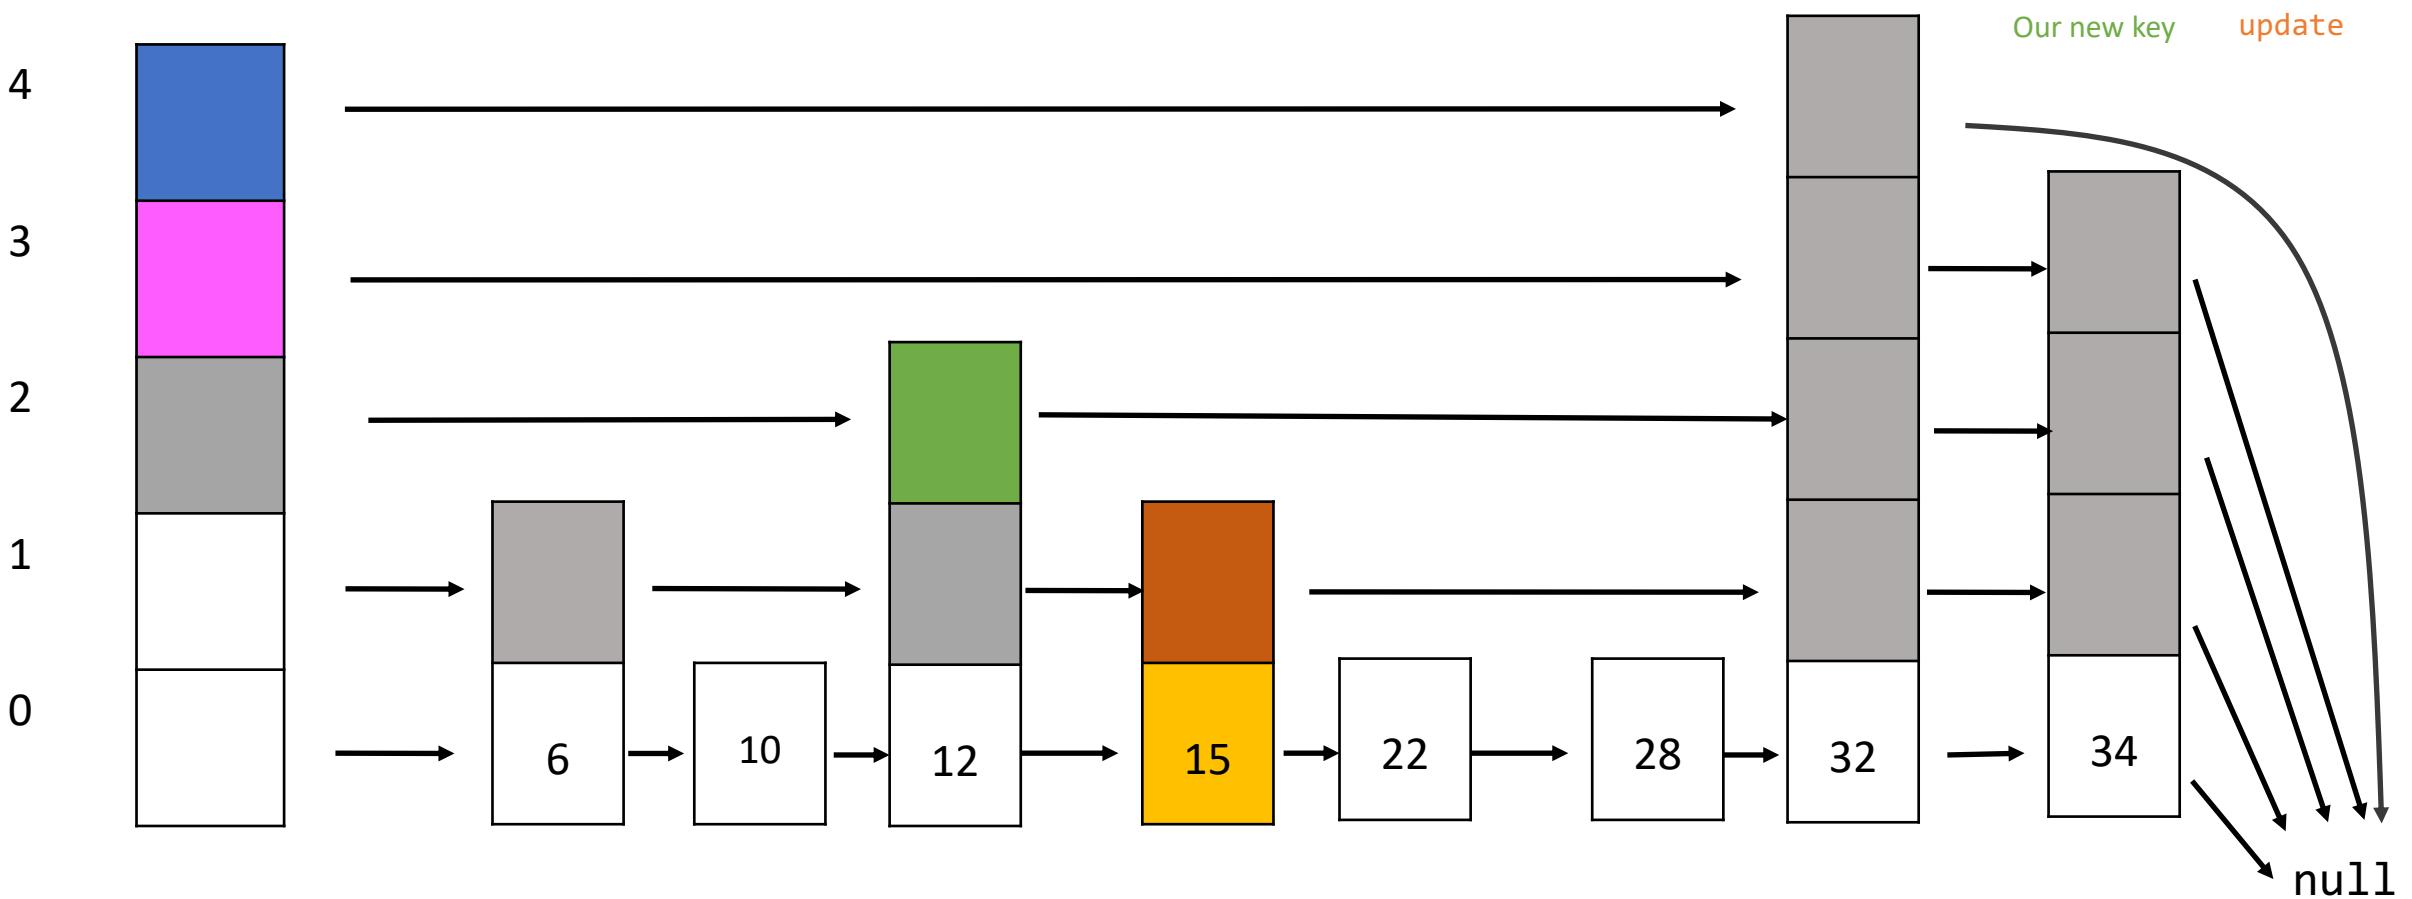
# Insertion Example

Let's not forget to write into update...



Task: Insert 25

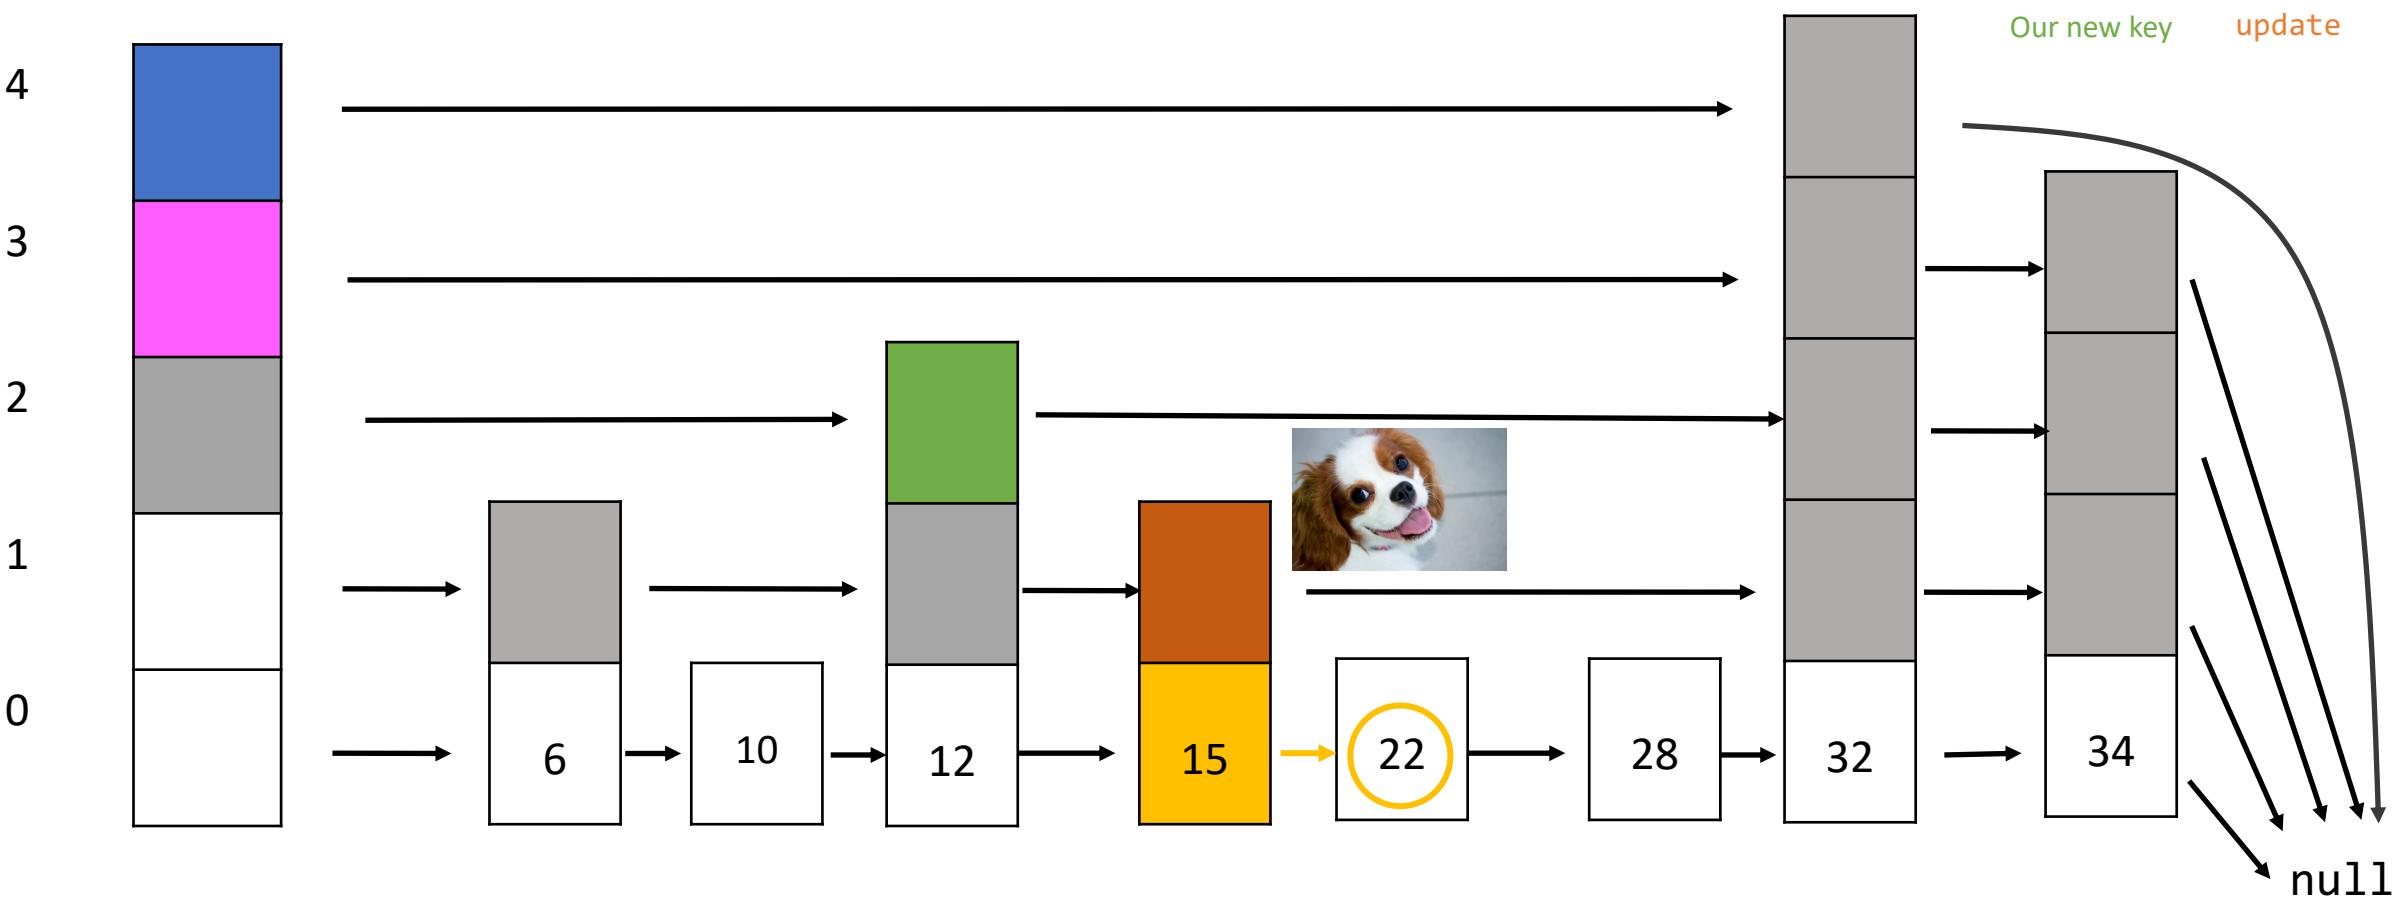
# Insertion Example



Task: Insert 25

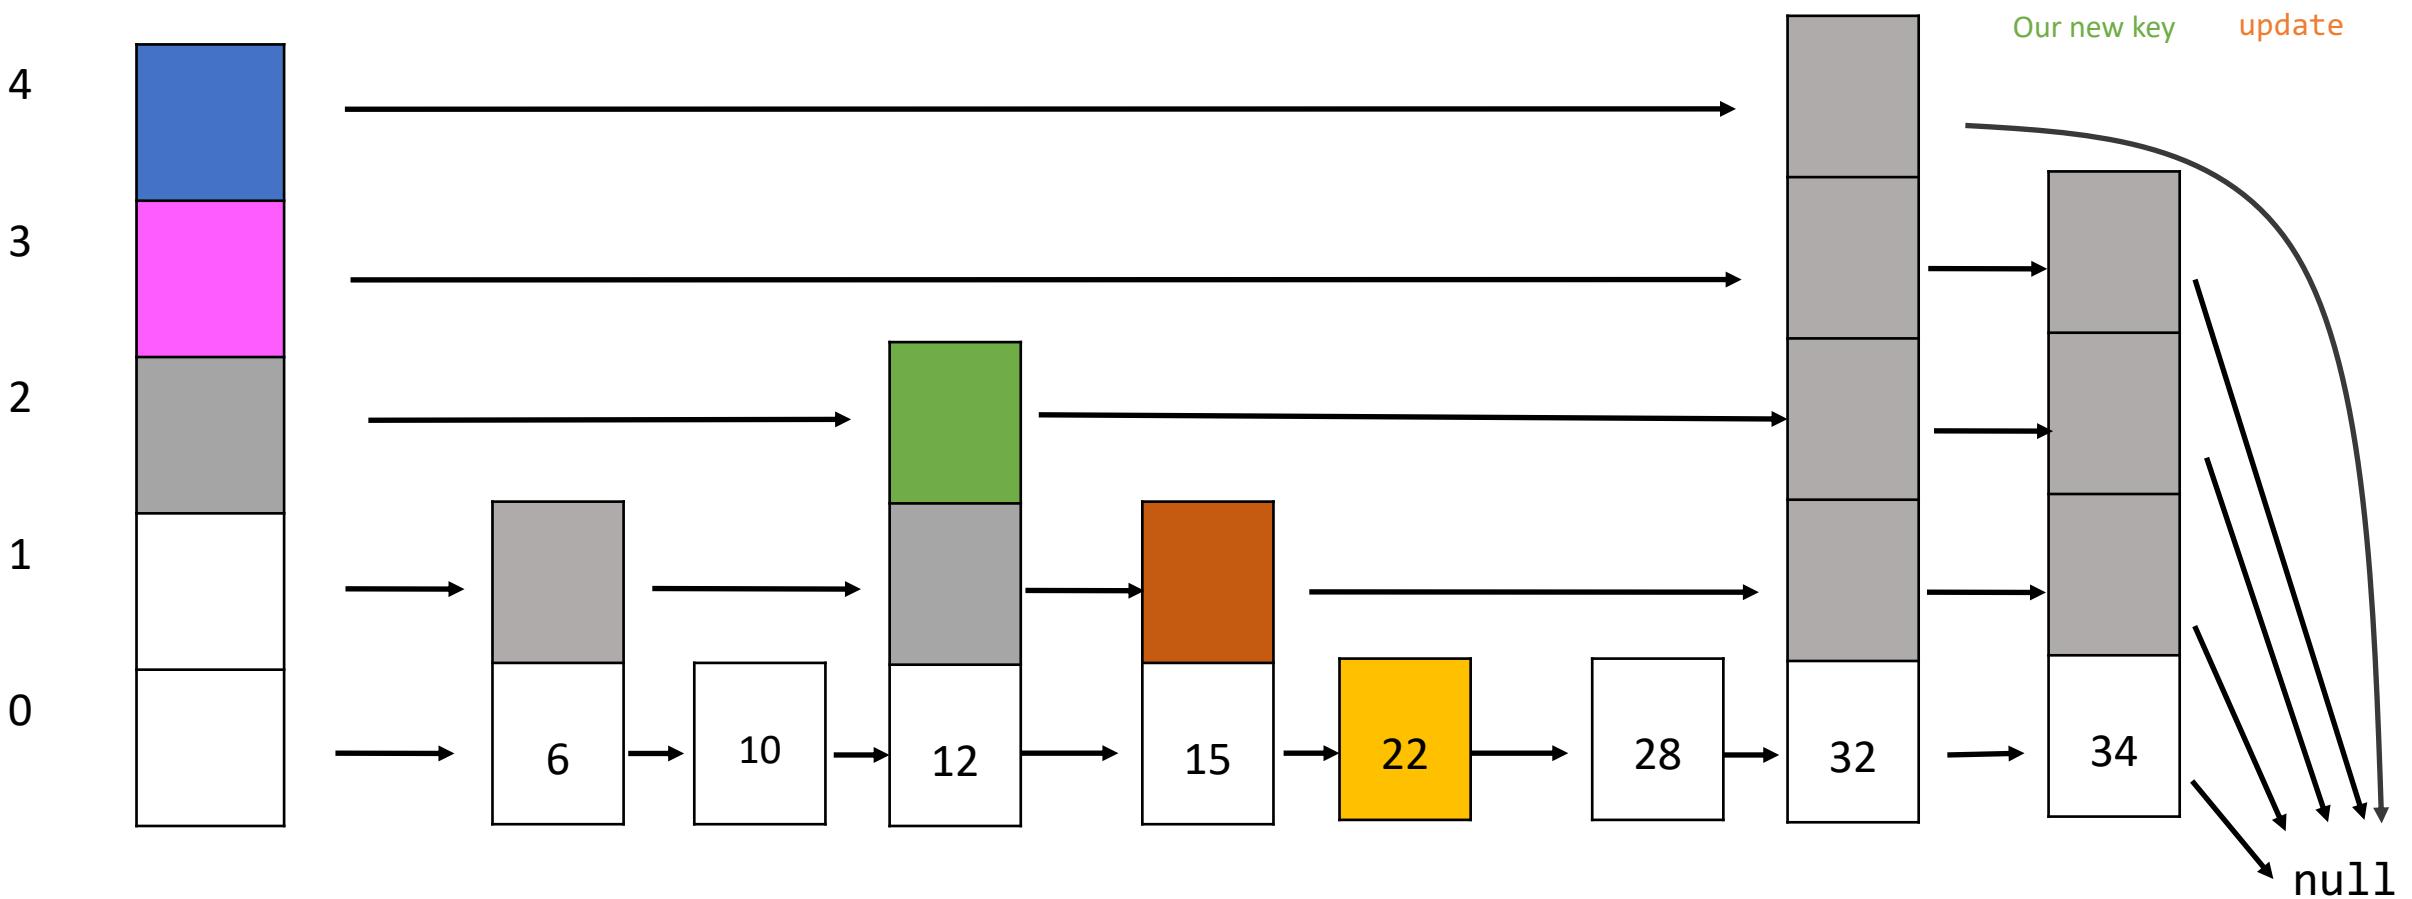
# Insertion Example

22 < 25, keep going....



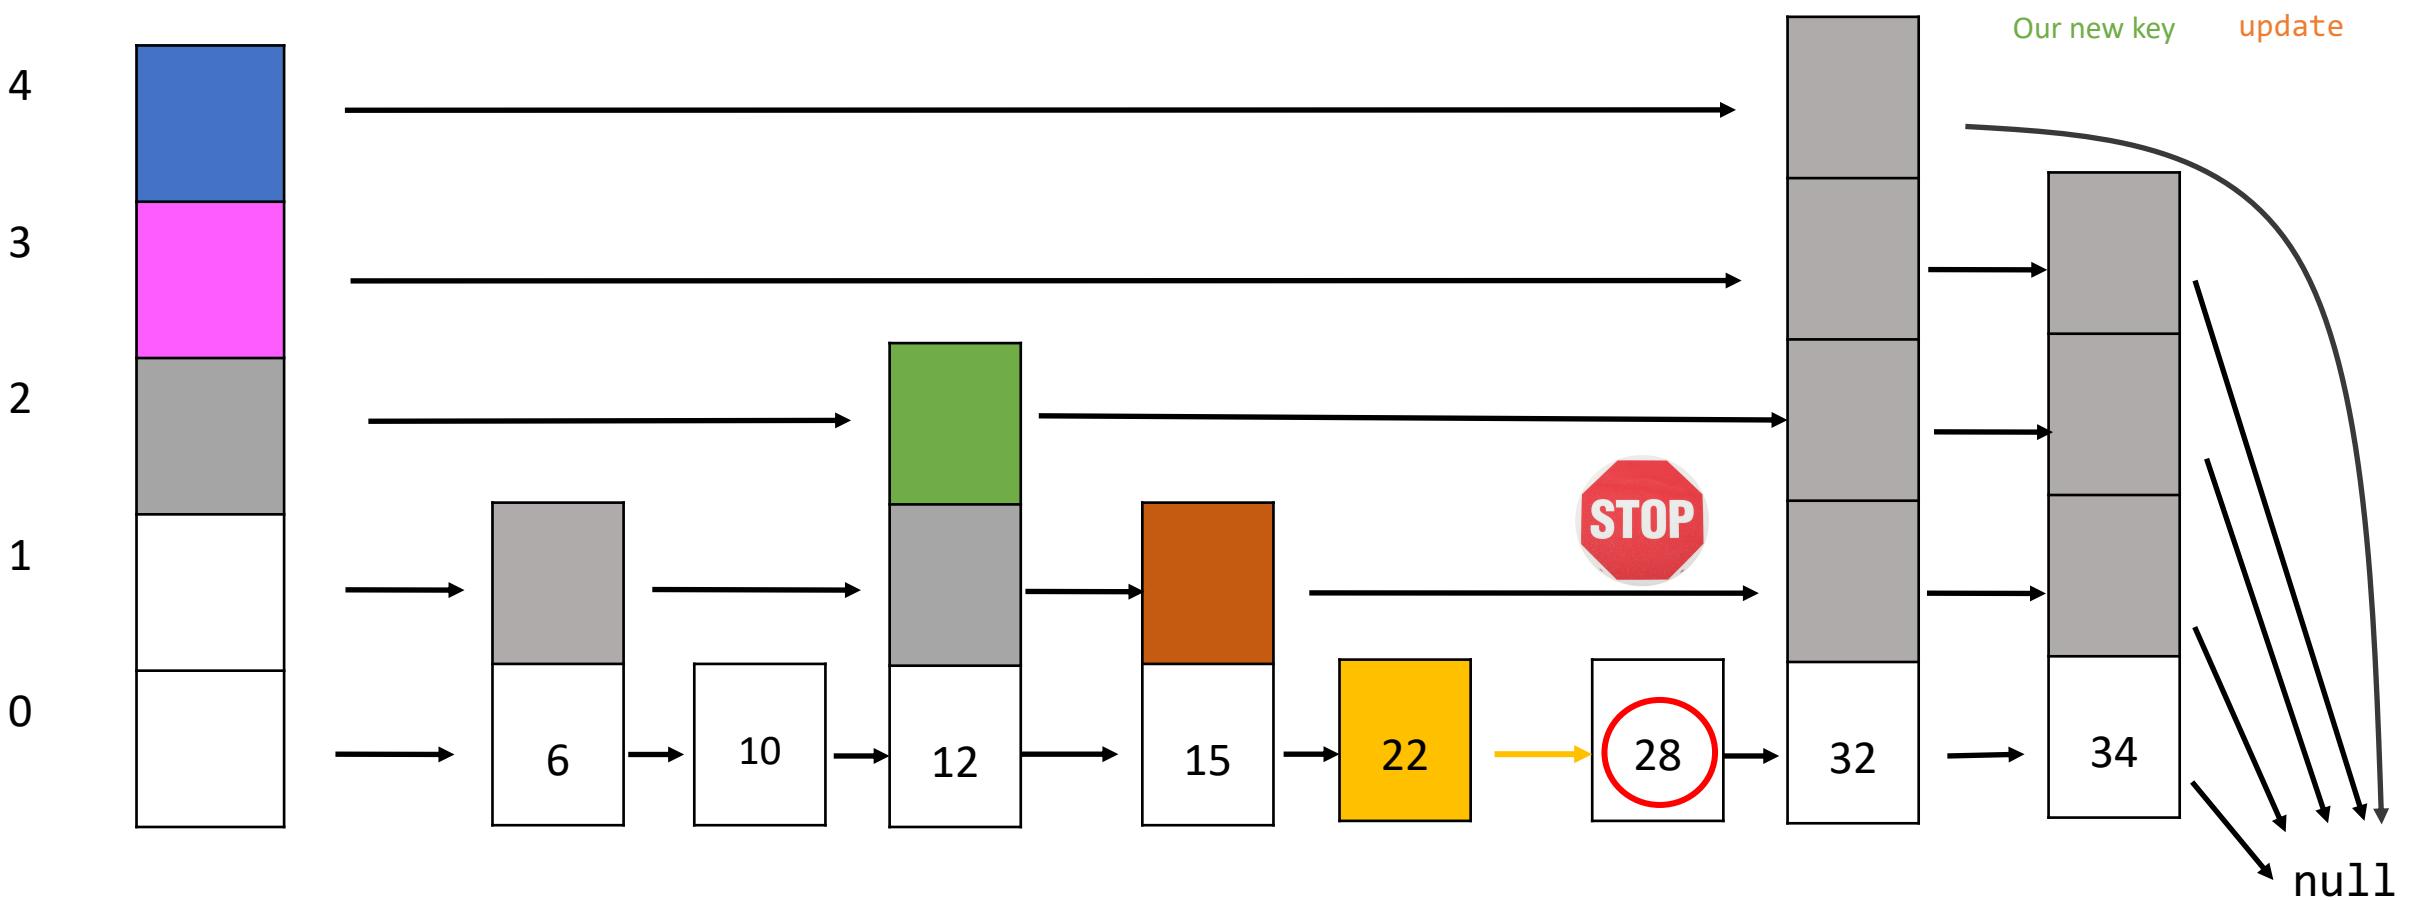
Task: Insert 25

# Insertion Example



Task: Insert 25

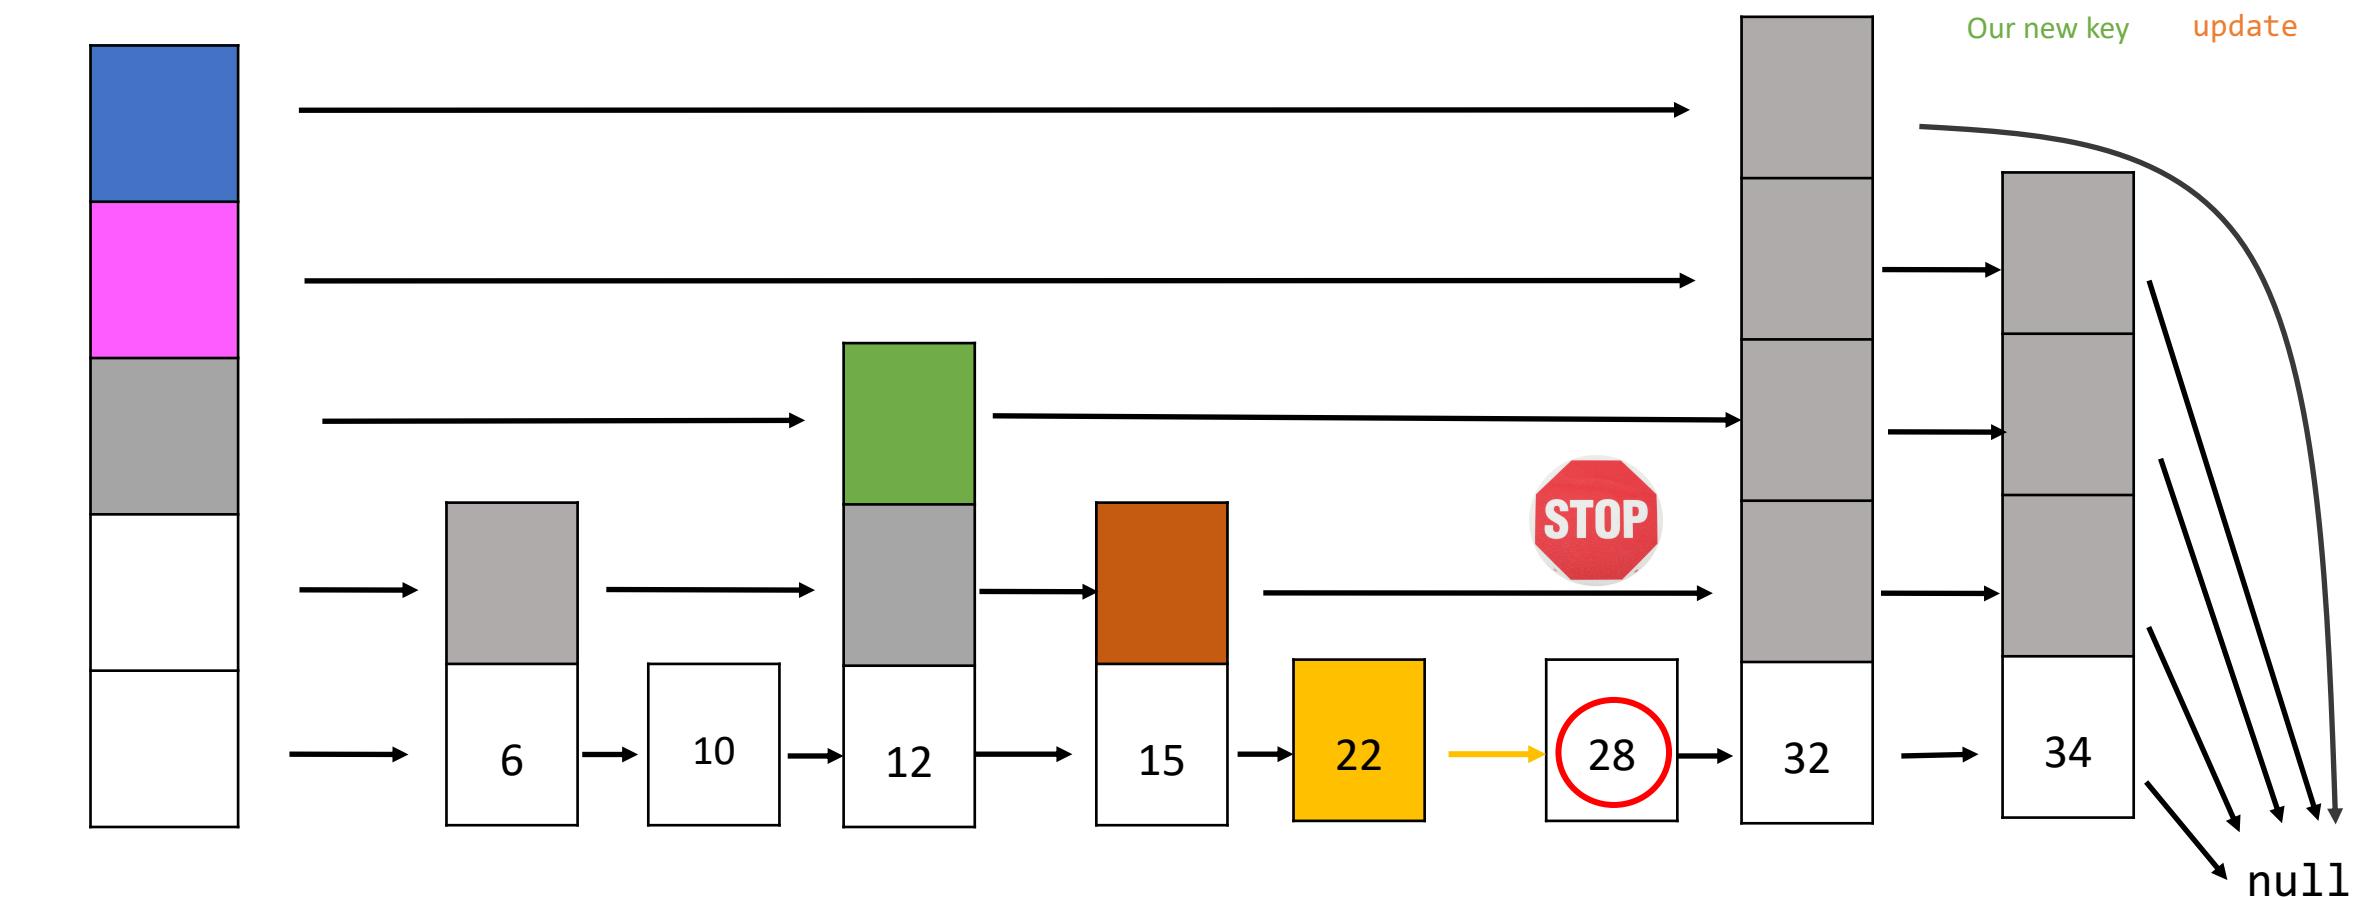
# Insertion Example



Task: Insert 25

# Insertion Example

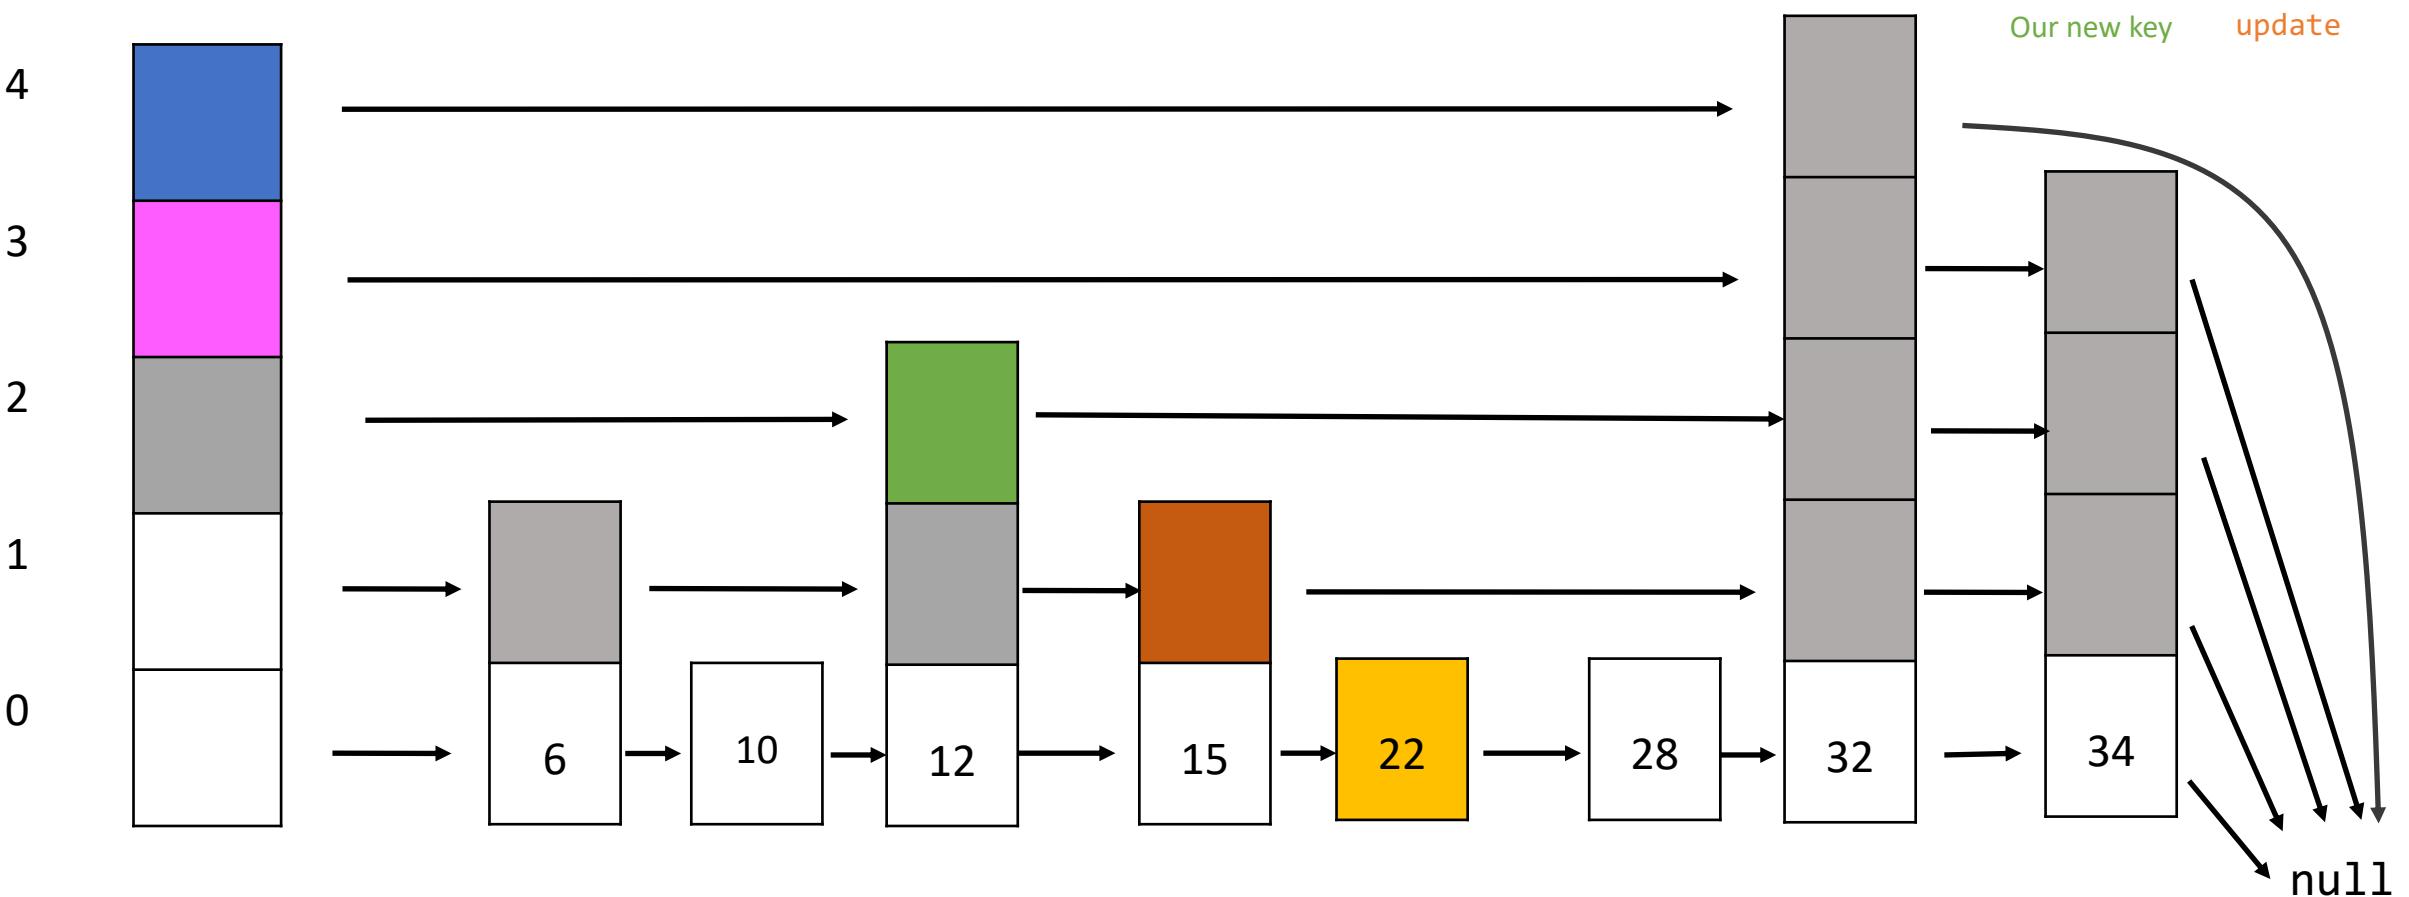
Write into update...



Task: Insert 25

# Insertion Example

Step 3: Splice the new tower in place by using update array.

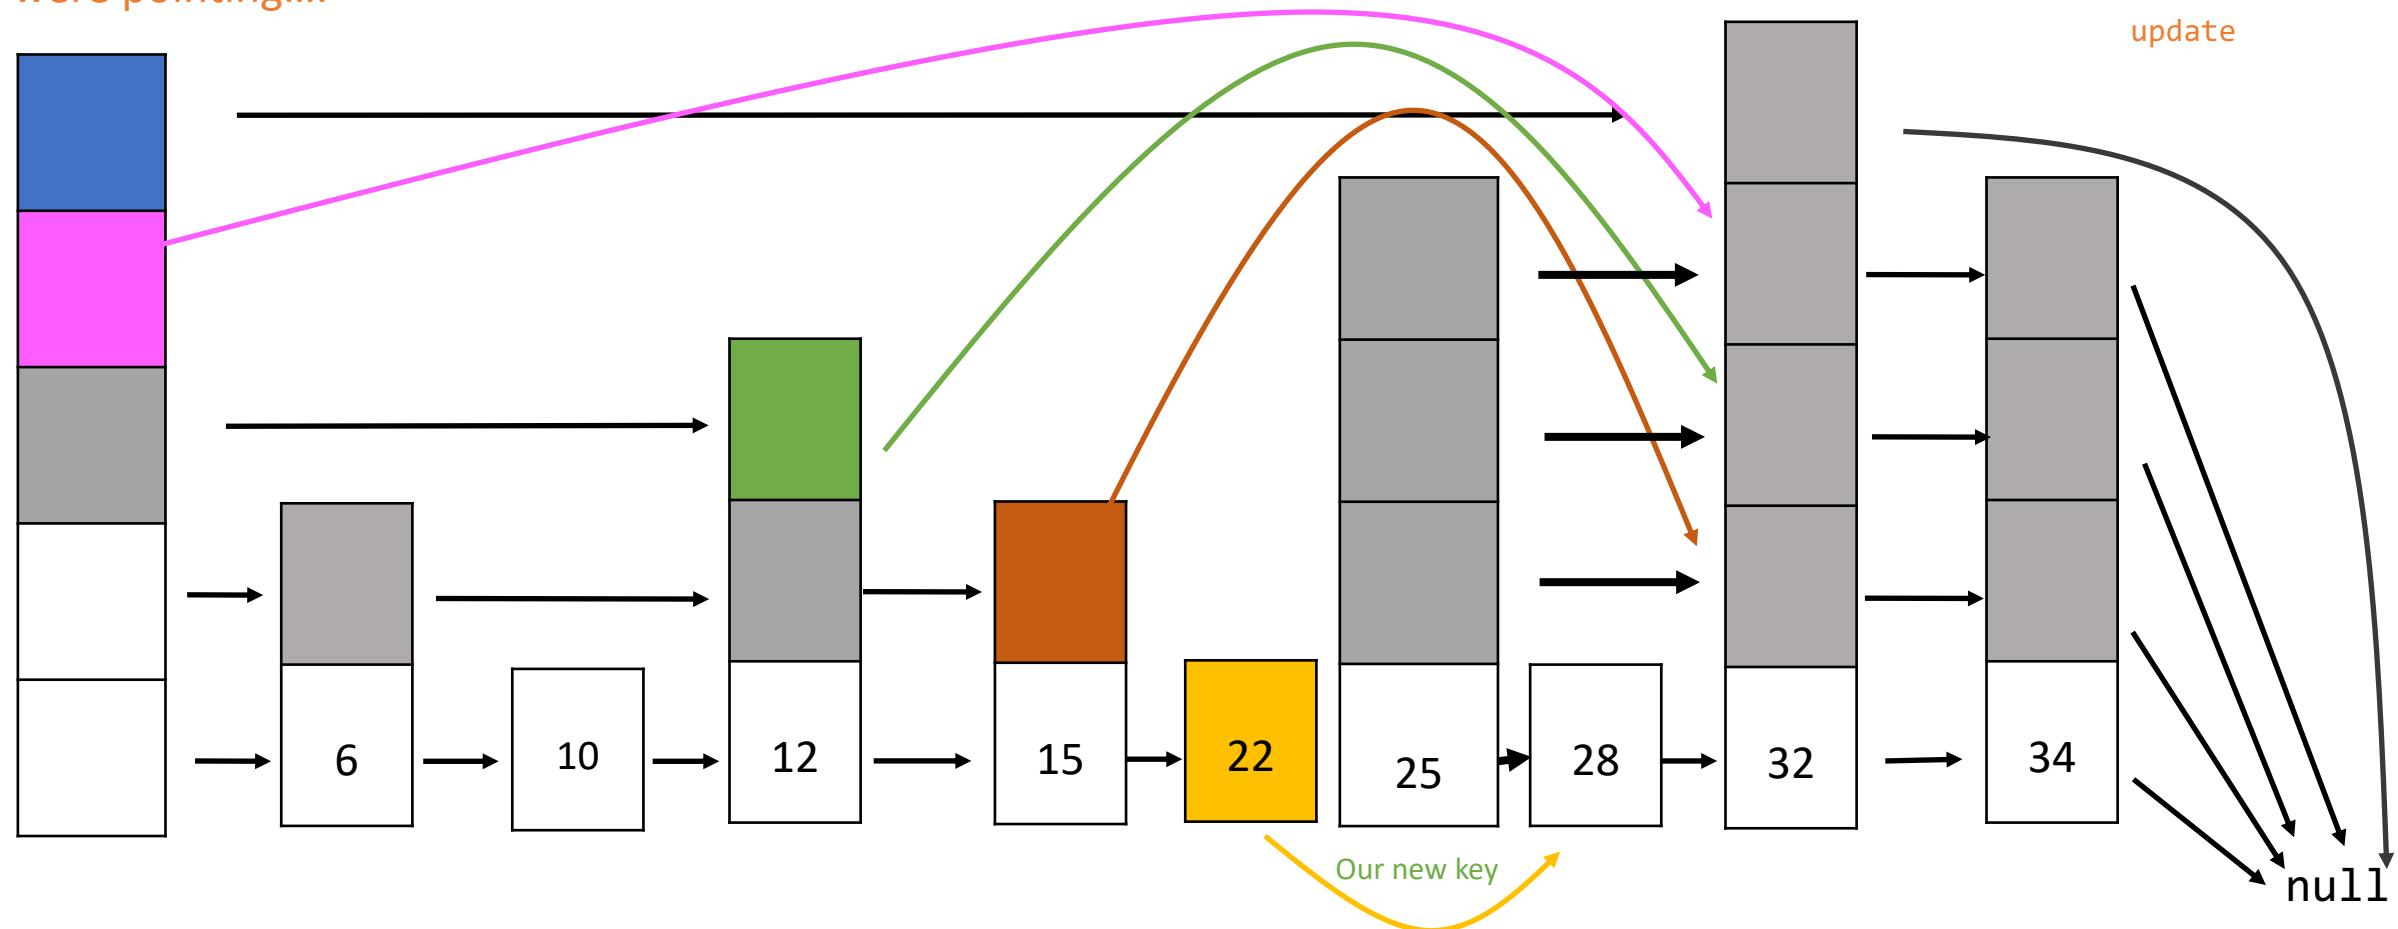


# Task: Insert 25

# Insertion Example

**Step 3: Splice the new tower in place by using update array.**

(a) Make all the forward pointers of the new key point wherever the nodes stored in update were pointing....

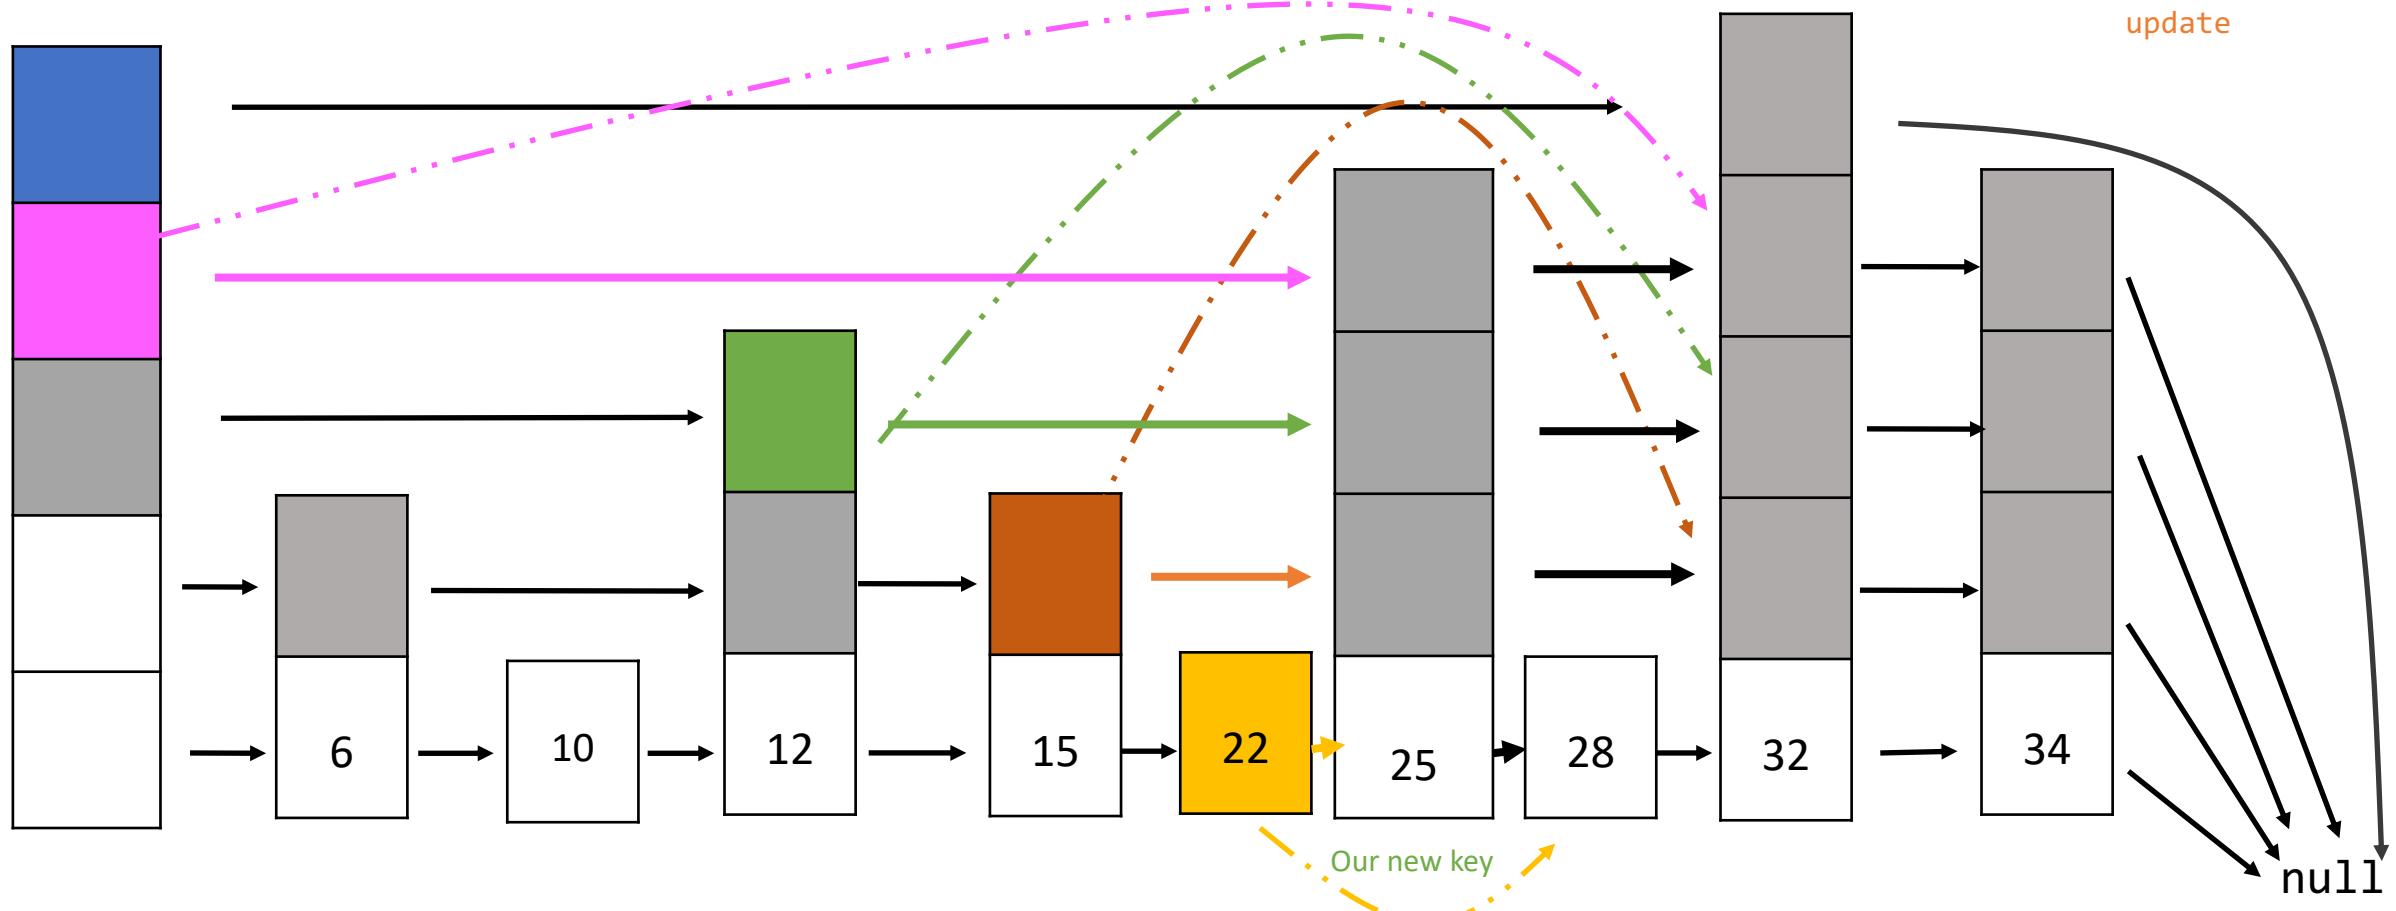
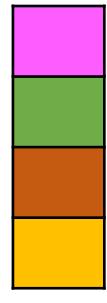


Task: Insert 25

# Insertion Example

Step 3: Splice the new tower in place by using update array.

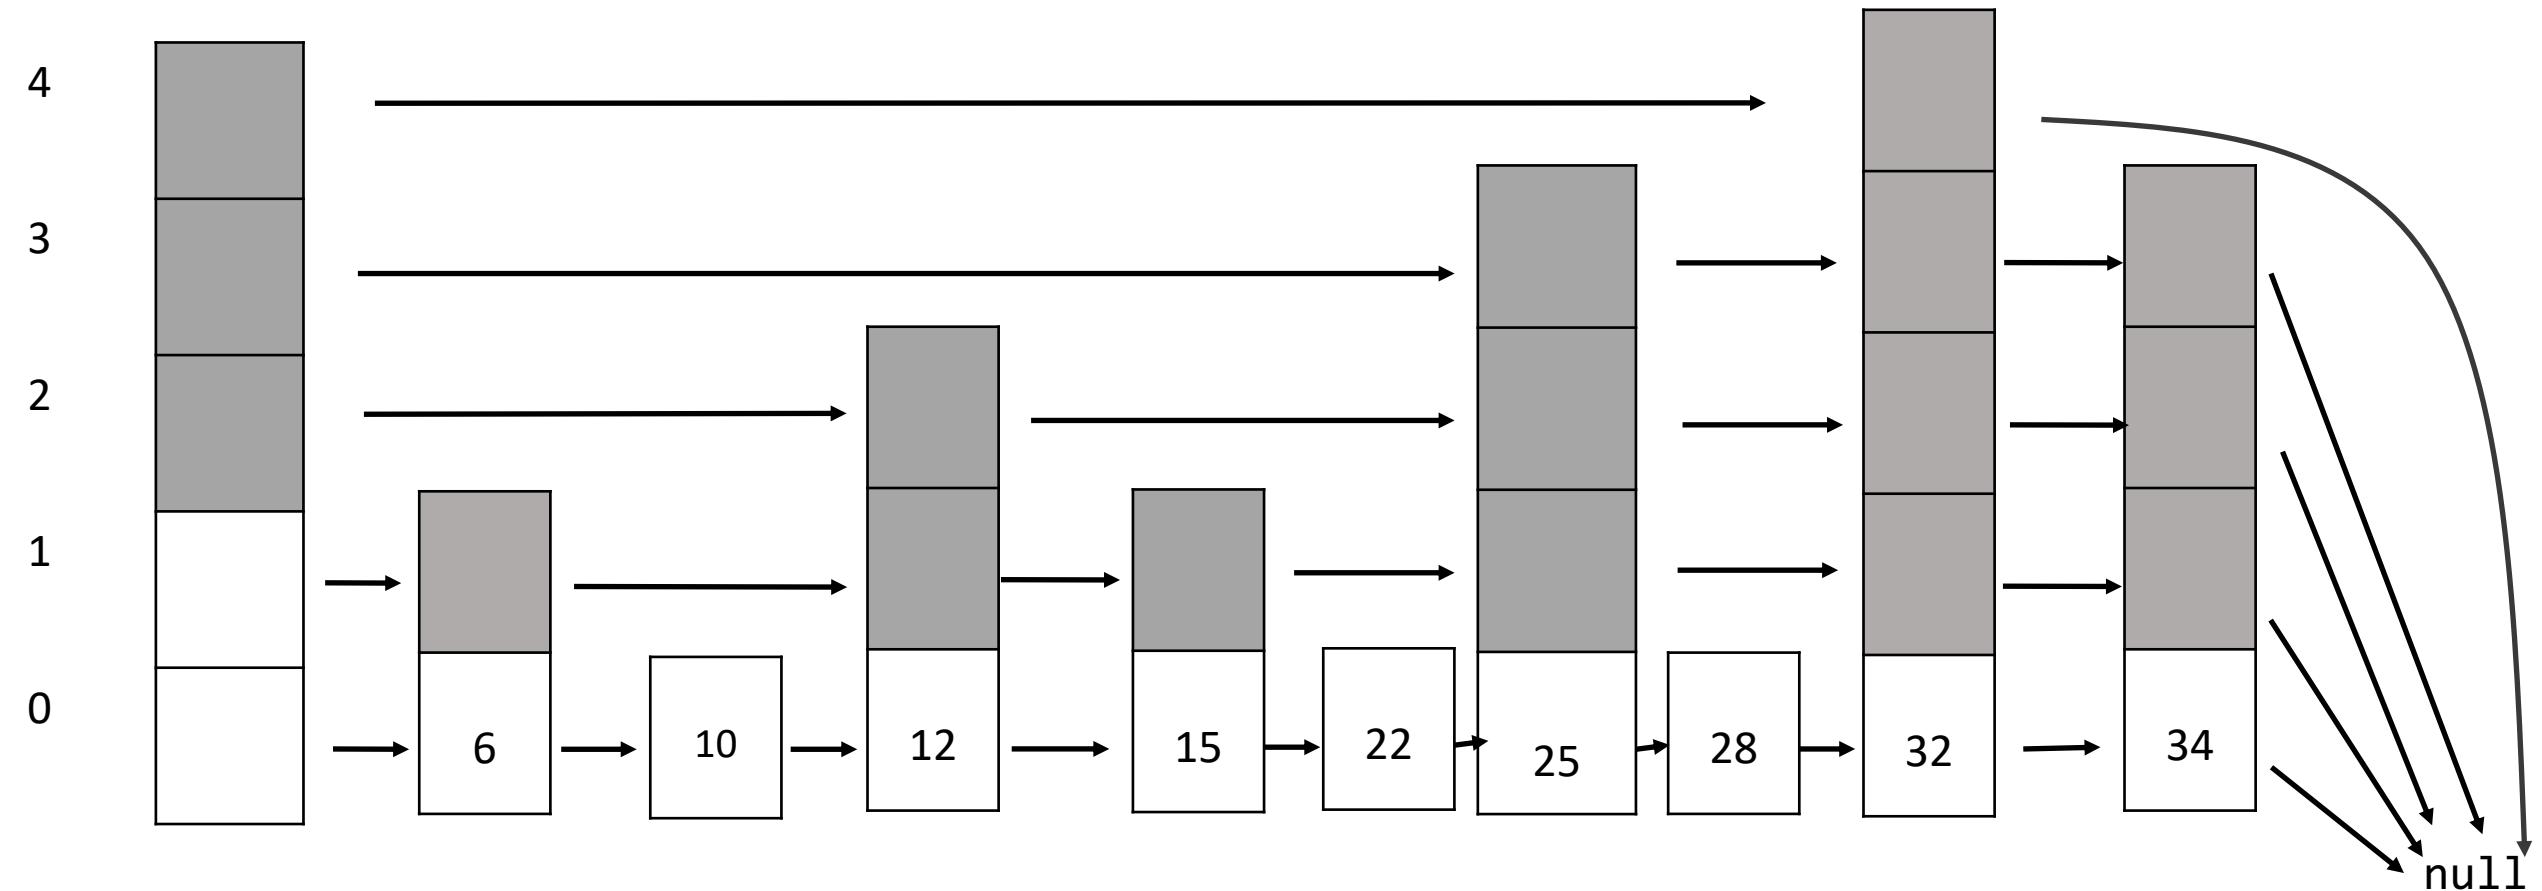
(b) ... and make sure that the nodes stored in update now point to the new key.



Task: Insert 25

# Insertion Example

Done! 😊



# Live insertion demo!

- Using **only the pseudocode** from Shaffer's book (p. 519 in the Java edition) and **a coin**, let's see if we can insert some ints into an initially empty Skiplist!

```
/** Insert a record into the skipList */
public void insert(Key k, E newValue) {
    int newLevel = randomLevel(); // New node's level
    if (newLevel > level) // If new node is deeper
        AdjustHead(newLevel); // adjust the header
    // Track end of level
    SkipNode<Key, E>[] update =
        (SkipNode<Key, E>[])new SkipNode[level+1];
    SkipNode<Key, E> x = head; // Start at header node
    for (int i=level; i>=0; i--) { // Find insert position
        while((x.forward[i] != null) &&
              (k.compareTo(x.forward[i].key()) > 0))
            x = x.forward[i];
        update[i] = x; // Track end at level i
    }
    x = new SkipNode<Key, E>(k, newValue, newLevel);
    for (int i=0; i<=newLevel; i++) { // Splice into list
        x.forward[i] = update[i].forward[i]; // Who x points to
        update[i].forward[i] = x; // Who y points to
    }
    size++; // Increment dictionary size
}

/** Pick a level using a geometric distribution */
int randomLevel() {
    int lev;
    for (lev=0; DSUtil.random(2) == 0; lev++); // Do nothing
    return lev;
}
```

Figure 16.4 Implementation for the Skip List `Insert` function.

# Skiplist insertion demo!

```
/** Insert a record into the skiplist */
public void insert(Key k, E newValue) {
    int newLevel = randomLevel(); // New node's level
    if (newLevel > level) // If new node is deeper
        AdjustHead(newLevel); // adjust the header
    // Track end of level
    SkipNode<Key, E>[] update =
        (SkipNode<Key, E>[]) new SkipNode[level+1];
    SkipNode<Key, E> x = head; // Start at header node
    for (int i=level; i>=0; i--) { // Find insert position
        while((x.forward[i] != null) &&
              (k.compareTo(x.forward[i].key()) > 0))
            x = x.forward[i];
        update[i] = x; // Track end at level i
    }
    x = new SkipNode<Key, E>(k, newValue, newLevel);
    for (int i=0; i<=newLevel; i++) { // Splice into list
        x.forward[i] = update[i].forward[i]; // Who x points to
        update[i].forward[i] = x; // Who y points to
    }
    size++; // Increment dictionary size
}

/** Pick a level using a geometric distribution */
int randomLevel() {
    int lev;
    for (lev=0; DSUtil.random(2) == 0; lev++); // Do nothing
    return lev;
}
```

# Worksheet exercise!



# Deletion

- Deletion works **very much like insertion.**
- Two differences:
  1. We don't have a need to sample a height.
  2. We use the information in the ***update*** array differently: instead of connecting nodes “before” the new tower to it and connecting the new tower to subsequent towers, **we remove the existing pointers**.

# Analysis

- Searching for an appropriate position dominates the cost for both insertions and deletions.
- Unit cost: visiting a node on our way to the appropriate position.
- Therefore, the length of the path that connects the sought key with `header[levels-1]` dominates our complexity.
- It can be proven (e.g check out “Analysis of expected search cost” [here](#)) that, with very high probability, this length is bounded above by a scalar multiple of  $\log_2 n$ .

# Analysis

- Searching for an appropriate position dominates the cost for both insertions and deletions.
- Unit cost: visiting a node on our way to the appropriate position.
- Therefore, the length of the path that connects the sought key with `header[levels-1]` dominates our complexity.
- It can be proven (e.g check out “Analysis of expected search cost” [here](#)) that, with very high probability, this length is bounded above by a scalar multiple of  $\log_2 n$ .
- So, with very high probability, searching a key  $k$  in a skip list  $\ell$  is  $\mathcal{O}(\log_2 n)$ !



# Why skip lists?

- Excellent distributed performance.
  - Locked and lock-free implementations exist
  - E.g check out Java's [ConcurrentSkipList](#)
- Compare with AVL trees and Red-Black trees where locking entire paths or even the entire tree might be necessary because of successive rotations.
- Local updates ensure that you only need to lock at most two towers per insertion, and only at the very end.
- Header node immutable; never needs to be locked.

# Why skip lists?

- **Insensitivity in insertion order.**
  - AVL + Red-Black trees might make balance guarantees, but a bad insertion order can affect performance of insertions!
- **Ease of implementation**
  - **No recursion overhead:** core of search algorithm is a double loop (while inside a for)
- Linearity of data and local updates allow **for distribution of data across nodes (very good horizontal scalability)**
  - Compare with AVL and Red-Black trees where storing different subtrees in different nodes might make it hard to **propagate rotations** from one subtree to the other.

# Disadvantages of skiplists

- $O(n)$  storage overhead.
- Probabilistic instead of deterministic guarantees
  - Compare with AVL and Red-Black trees which offer guarantees of  $\log_2 n$  and  $2 \cdot \log_2 n$  respectively.
- Worst-case search constant bigger than AVL (AVL-1) Trees and B-Trees, on par with Red-Black trees.